

HTML 5

Draft Recommendation — 16 October 2008



You can take part in this work. Join the working group's discussion list.

Web designers! We have a FAQ, a forum, and a help mailing list for you!

Multiple-page version:

<http://whatwg.org/html5>

One-page version:

<http://www.whatwg.org/specs/web-apps/current-work/>

PDF print versions:

A4: <http://www.whatwg.org/specs/web-apps/current-work/html5-a4.pdf>

Letter: <http://www.whatwg.org/specs/web-apps/current-work/html5-letter.pdf>

Version history:

Twitter messages (non-editorial changes only): <http://twitter.com/WWHATWG>

Commit-Watchers mailing list: <http://lists.whatwg.org/listinfo.cgi/commit-watchers-whatwg.org>

Interactive Web interface: <http://html5.org/tools/web-apps-tracker>

Subversion interface: <http://svn.whatwg.org/>

HTML diff with the last version in Subversion: <http://whatwg.org/specs/web-apps/current-work/index-diff>

Issues:

To send feedback: whatwg@whatwg.org

To view and vote on feedback: <http://www.whatwg.org/issues/>

Editor:

Ian Hickson, Google, ian@hixie.ch

© Copyright 2004-2008 Apple Computer, Inc., Mozilla Foundation, and Opera Software ASA.
You are granted a license to use, reproduce and create derivative works of this document.

Abstract

This specification evolves HTML and its related APIs to ease the authoring of Web-based applications. Additions include context menus, a direct-mode graphics canvas, a full duplex client-server communication channel, more semantics, audio and video, various features for offline Web applications, sandboxed iframes, and scoped styling. Heavy emphasis is placed on keeping the language backwards compatible with existing legacy user agents and on keeping user agents backwards compatible with existing legacy documents.

Status of this document

This is a work in progress! This document is changing on a daily if not hourly basis in response to comments and as a general part of its development process. Comments are very welcome, please send them to whatwg@whatwg.org. Thank you.

The current focus is in responding to the outstanding feedback. (There is a chart showing current progress.)

Implementors should be aware that this specification is not stable. **Implementors not taking part in the discussions are likely to find the specification from under them in incompatible ways.** Vendors interested in implementing this specification before it eventually reaches the call for implementations should join the WHATWG mailing list and take part in the discussions.

Watch for updates

This specification is also being produced by the W3C HTML WG. The two specifications are identical from the table of contents onwards.

This specification is intended to replace (be the new version of) what was previously the HTML4, XHTML 1.x, and DOM2 HTML specifications.

Stability

Different parts of this specification are at different levels of maturity.

- ** Some of the more major known issues are marked like this. There are many other issues that have been raised as well; the issues given in this document are not the only known issues! Also, firing of events needs to be unified (right now some bubble, some don't, they all use different text to fire events, etc).

Table of contents

- 1 Introduction (page 18)
 - 1.1 Background (page 18)
 - 1.2 Scope (page 18)
 - 1.3 History (page 19)
 - 1.4 Relationships to other specifications (page 19)
 - 1.4.1 Relationship to HTML 4.01 and DOM2 HTML (page 19)
 - 1.4.2 Relationship to XHTML 1.x (page 19)
 - 1.4.3 Relationship to XHTML2 and XForms (page 20)
 - 1.4.4 Relationship to XUL, Flash, Silverlight, and other proprietary UI languages (page 20)
 - 1.5 HTML vs XHTML (page 20)
 - 1.6 Structure of this specification (page 21)
 - 1.6.1 How to read this specification (page 22)
 - 1.6.2 Typographic conventions (page 22)
- 2 Common infrastructure (page 23)
 - 2.1 Terminology (page 23)
 - 2.1.1 XML (page 23)
 - 2.1.2 DOM trees (page 24)
 - 2.1.3 Scripting (page 24)
 - 2.1.4 Plugins (page 25)
 - 2.1.5 Character encodings (page 25)
 - 2.2 Conformance requirements (page 25)
 - 2.2.1 Dependencies (page 29)
 - 2.2.2 Features defined in other specifications (page 30)
 - 2.2.3 Common conformance requirements for APIs exposed to JavaScript (page 30)
 - 2.3 Case-sensitivity and string comparison (page 31)
 - 2.4 Common microsyntaxes (page 31)
 - 2.4.1 Common parser idioms (page 32)
 - 2.4.2 Boolean attributes (page 32)
 - 2.4.3 Numbers (page 33)
 - 2.4.3.1 Unsigned integers (page 33)
 - 2.4.3.2 Signed integers (page 33)
 - 2.4.3.3 Real numbers (page 34)
 - 2.4.3.4 Ratios (page 36)
 - 2.4.3.5 Percentages and dimensions (page 38)
 - 2.4.3.6 Lists of integers (page 38)
 - 2.4.4 Dates and times (page 40)
 - 2.4.4.1 Specific moments in time (page 40)
 - 2.4.4.2 Vaguer moments in time (page 44)
 - 2.4.4.3 UTC dates and times (page 47)
 - 2.4.4.4 Local dates and times (page 48)
 - 2.4.4.5 Dates (page 48)
 - 2.4.4.6 Months (page 48)
 - 2.4.4.7 Weeks (page 48)
 - 2.4.4.8 Times (page 49)
 - 2.4.4.9 Time offsets (page 49)
 - 2.4.5 Space-separated tokens (page 49)

- 2.4.6 Comma-separated tokens (page 50)
 - 2.4.7 Keywords and enumerated attributes (page 51)
 - 2.4.8 References (page 52)
 - 2.5 URLs (page 52)
 - 2.5.1 Terminology (page 52)
 - 2.5.2 Parsing URLs (page 53)
 - 2.5.3 Resolving URLs (page 54)
 - 2.5.4 Dynamic changes to base URLs (page 56)
 - 2.5.5 Interfaces for URL manipulation (page 57)
 - 2.6 Fetching resources (page 59)
 - 2.7 Determining the type of a resource (page 60)
 - 2.7.1 Content-Type metadata (page 60)
 - 2.7.2 Content-Type sniffing: Web pages (page 61)
 - 2.7.3 Content-Type sniffing: text or binary (page 62)
 - 2.7.4 Content-Type sniffing: unknown type (page 62)
 - 2.7.5 Content-Type sniffing: image (page 65)
 - 2.7.6 Content-Type sniffing: feed or HTML (page 65)
 - 2.8 Common DOM interfaces (page 67)
 - 2.8.1 Reflecting content attributes in DOM attributes (page 67)
 - 2.8.2 Collections (page 69)
 - 2.8.2.1 HTMLCollection (page 69)
 - 2.8.2.2 HTMLFormControlsCollection (page 70)
 - 2.8.2.3 HTMLOptionsCollection (page 71)
 - 2.8.3 DOMTokenList (page 72)
 - 2.8.4 DOMStringMap (page 74)
 - 2.8.5 DOM feature strings (page 74)
- 3 Semantics and structure of HTML documents (page 76)
- 3.1 Introduction (page 76)
 - 3.2 Documents (page 76)
 - 3.2.1 Documents in the DOM (page 76)
 - 3.2.2 Security (page 78)
 - 3.2.3 Resource metadata management (page 78)
 - 3.2.4 DOM tree accessors (page 80)
 - 3.3 Elements (page 82)
 - 3.3.1 Semantics (page 82)
 - 3.3.2 Elements in the DOM (page 84)
 - 3.3.3 Global attributes (page 86)
 - 3.3.3.1 The `id` attribute (page 87)
 - 3.3.3.2 The `title` attribute (page 87)
 - 3.3.3.3 The `lang` and `xml:lang` attributes (page 88)
 - 3.3.3.4 The `xml:base` attribute (XML only) (page 89)
 - 3.3.3.5 The `dir` attribute (page 89)
 - 3.3.3.6 The `class` attribute (page 90)
 - 3.3.3.7 The `style` attribute (page 90)
 - 3.3.3.8 Embedding custom non-visible data (page 91)
 - 3.4 Content models (page 92)
 - 3.4.1 Kinds of content (page 93)
 - 3.4.1.1 Metadata content (page 93)
 - 3.4.1.2 Flow content (page 93)
 - 3.4.1.3 Sectioning content (page 94)

- 3.4.1.4 Heading content (page 94)
 - 3.4.1.5 Phrasing content (page 94)
 - 3.4.1.6 Embedded content (page 94)
 - 3.4.1.7 Interactive content (page 95)
 - 3.4.2 Transparent content models (page 96)
 - 3.5 Paragraphs (page 96)
 - 3.6 APIs in HTML documents (page 98)
 - 3.7 Dynamic markup insertion (page 99)
 - 3.7.1 Controlling the input stream (page 99)
 - 3.7.2 Dynamic markup insertion in HTML (page 101)
 - 3.7.3 Dynamic markup insertion in XML (page 104)
- 4 The elements of HTML (page 107)
- 4.1 The root element (page 107)
 - 4.1.1 The `html` element (page 107)
 - 4.2 Document metadata (page 107)
 - 4.2.1 The `head` element (page 107)
 - 4.2.2 The `title` element (page 108)
 - 4.2.3 The `base` element (page 109)
 - 4.2.4 The `link` element (page 110)
 - 4.2.5 The `meta` element (page 113)
 - 4.2.5.1 Standard metadata names (page 115)
 - 4.2.5.2 Other metadata names (page 115)
 - 4.2.5.3 Pragma directives (page 116)
 - 4.2.5.4 Specifying the document's character encoding (page 119)
 - 4.2.6 The `style` element (page 120)
 - 4.2.7 Styling (page 122)
 - 4.3 Scripting (page 122)
 - 4.3.1 The `script` element (page 123)
 - 4.3.1.1 Scripting languages (page 129)
 - 4.3.2 The `noscript` element (page 129)
 - 4.3.3 The `eventsource` element (page 131)
 - 4.4 Sections (page 132)
 - 4.4.1 The `body` element (page 132)
 - 4.4.2 The `section` element (page 133)
 - 4.4.3 The `nav` element (page 133)
 - 4.4.4 The `article` element (page 134)
 - 4.4.5 The `aside` element (page 135)
 - 4.4.6 The `h1`, `h2`, `h3`, `h4`, `h5`, and `h6` elements (page 136)
 - 4.4.7 The `header` element (page 136)
 - 4.4.8 The `footer` element (page 138)
 - 4.4.9 The `address` element (page 139)
 - 4.4.10 Headings and sections (page 140)
 - 4.4.10.1 Creating an outline (page 142)
 - 4.4.10.2 Distinguishing site-wide headings from page headings (page 146)
 - 4.5 Grouping content (page 146)
 - 4.5.1 The `p` element (page 146)
 - 4.5.2 The `hr` element (page 147)
 - 4.5.3 The `br` element (page 148)

- 4.5.4 The `pre` element (page 149)
- 4.5.5 The `dialog` element (page 150)
- 4.5.6 The `blockquote` element (page 151)
- 4.5.7 The `ol` element (page 152)
- 4.5.8 The `ul` element (page 154)
- 4.5.9 The `li` element (page 154)
- 4.5.10 The `dl` element (page 156)
- 4.5.11 The `dt` element (page 158)
- 4.5.12 The `dd` element (page 159)
- 4.6 Text-level semantics (page 160)
 - 4.6.1 The `a` element (page 160)
 - 4.6.2 The `q` element (page 162)
 - 4.6.3 The `cite` element (page 163)
 - 4.6.4 The `em` element (page 165)
 - 4.6.5 The `strong` element (page 166)
 - 4.6.6 The `small` element (page 166)
 - 4.6.7 The `mark` element (page 167)
 - 4.6.8 The `dfn` element (page 170)
 - 4.6.9 The `abbr` element (page 171)
 - 4.6.10 The `time` element (page 172)
 - 4.6.11 The `progress` element (page 174)
 - 4.6.12 The `meter` element (page 176)
 - 4.6.13 The `code` element (page 182)
 - 4.6.14 The `var` element (page 183)
 - 4.6.15 The `samp` element (page 183)
 - 4.6.16 The `kbd` element (page 184)
 - 4.6.17 The `sub` and `sup` elements (page 185)
 - 4.6.18 The `span` element (page 186)
 - 4.6.19 The `i` element (page 186)
 - 4.6.20 The `b` element (page 187)
 - 4.6.21 The `bdo` element (page 188)
 - 4.6.22 The `ruby` element (page 189)
 - 4.6.23 The `rt` element (page 190)
 - 4.6.24 The `rp` element (page 190)
 - 4.6.25 Usage summary (page 191)
 - 4.6.26 Footnotes (page 191)
- 4.7 Edits (page 192)
 - 4.7.1 The `ins` element (page 193)
 - 4.7.2 The `del` element (page 194)
 - 4.7.3 Attributes common to `ins` and `del` elements (page 194)
 - 4.7.4 Edits and paragraphs (page 195)
 - 4.7.5 Edits and lists (page 196)
- 4.8 Embedded content (page 197)
 - 4.8.1 The `figure` element (page 197)
 - 4.8.2 The `img` element (page 199)
 - 4.8.2.1 Requirements for providing text to act as an alternative for images (page 205)
 - 4.8.2.1.1 A link or button containing nothing but the image (page 205)

- 4.8.2.1.2 A phrase or paragraph with an alternative graphical representation: charts, diagrams, graphs, maps, illustrations (page 206)
 - 4.8.2.1.3 A short phrase or label with an alternative graphical representation: icons, logos (page 207)
 - 4.8.2.1.4 Text that has been rendered to a graphic for typographical effect (page 209)
 - 4.8.2.1.5 A graphical representation of some of the surrounding text (page 209)
 - 4.8.2.1.6 A purely decorative image that doesn't add any information but is still specific to the surrounding content (page 210)
 - 4.8.2.1.7 A group of images that form a single larger picture with no links (page 211)
 - 4.8.2.1.8 A group of images that form a single larger picture with links (page 211)
 - 4.8.2.1.9 A key part of the content (page 212)
 - 4.8.2.1.10 An image not intended for the user (page 216)
 - 4.8.2.1.11 An image in an e-mail or document intended for a specific person who is known to be able to view images (page 216)
 - 4.8.2.1.12 General guidelines (page 216)
- 4.8.3 The `iframe` element (page 216)
 - 4.8.4 The `embed` element (page 221)
 - 4.8.5 The `object` element (page 224)
 - 4.8.6 The `param` element (page 228)
 - 4.8.7 The `video` element (page 228)
 - 4.8.7.1 Video and audio codecs for video elements (page 232)
 - 4.8.8 The `audio` element (page 233)
 - 4.8.8.1 Audio codecs for audio elements (page 234)
 - 4.8.9 The `source` element (page 234)
 - 4.8.10 Media elements (page 236)
 - 4.8.10.1 Error codes (page 238)
 - 4.8.10.2 Location of the media resource (page 239)
 - 4.8.10.3 Network states (page 240)
 - 4.8.10.4 Loading the media resource (page 241)
 - 4.8.10.5 Offsets into the media resource (page 246)
 - 4.8.10.6 The ready states (page 248)
 - 4.8.10.7 Cue ranges (page 249)
 - 4.8.10.8 Playing the media resource (page 250)
 - 4.8.10.9 Seeking (page 254)
 - 4.8.10.10 User interface (page 256)
 - 4.8.10.11 Time ranges (page 257)
 - 4.8.10.12 Byte ranges (page 257)
 - 4.8.10.13 Event summary (page 258)
 - 4.8.10.14 Security and privacy considerations (page 260)
 - 4.8.11 The `canvas` element (page 260)
 - 4.8.11.1 The 2D context (page 263)
 - 4.8.11.1.1 The canvas state (page 266)

- 4.8.11.1.2 Transformations (page 267)
- 4.8.11.1.3 Compositing (page 267)
- 4.8.11.1.4 Colors and styles (page 269)
- 4.8.11.1.5 Line styles (page 271)
- 4.8.11.1.6 Shadows (page 273)
- 4.8.11.1.7 Simple shapes (rectangles) (page 274)
- 4.8.11.1.8 Complex shapes (paths) (page 274)
- 4.8.11.1.9 Text (page 277)
- 4.8.11.1.10 Images (page 280)
- 4.8.11.1.11 Pixel manipulation (page 282)
- 4.8.11.1.12 Drawing model (page 286)
- 4.8.11.2 Color spaces and color correction (page 286)
- 4.8.11.3 Security with canvas elements (page 287)
- 4.8.12 The map element (page 287)
- 4.8.13 The area element (page 288)
- 4.8.14 Image maps (page 291)
 - 4.8.14.1 Authoring (page 291)
 - 4.8.14.2 Processing model (page 291)
- 4.8.15 MathML (page 294)
- 4.8.16 SVG (page 294)
- 4.8.17 Dimension attributes (page 295)
- 4.9 Tabular data (page 295)
 - 4.9.1 Introduction (page 295)
 - 4.9.2 The table element (page 295)
 - 4.9.3 The caption element (page 298)
 - 4.9.4 The colgroup element (page 299)
 - 4.9.5 The col element (page 300)
 - 4.9.6 The tbody element (page 300)
 - 4.9.7 The thead element (page 301)
 - 4.9.8 The tfoot element (page 302)
 - 4.9.9 The tr element (page 302)
 - 4.9.10 The td element (page 303)
 - 4.9.11 The th element (page 304)
 - 4.9.12 Attributes common to td and th elements (page 305)
 - 4.9.13 Processing model (page 306)
 - 4.9.13.1 Forming a table (page 307)
 - 4.9.13.2 Forming relationships between data cells and header cells (page 311)
- 4.10 Forms (page 313)
 - 4.10.1 The form element (page 314)
 - 4.10.2 The fieldset element (page 316)
 - 4.10.3 The label element (page 317)
 - 4.10.4 The input element (page 318)
 - 4.10.4.1 States of the type attribute (page 324)
 - 4.10.4.1.1 Hidden state (page 324)
 - 4.10.4.1.2 Text state (page 324)
 - 4.10.4.1.3 E-mail state (page 325)
 - 4.10.4.1.4 URL state (page 325)
 - 4.10.4.1.5 Password state (page 326)
 - 4.10.4.1.6 Date and Time state (page 326)
 - 4.10.4.1.7 Date state (page 328)

- 4.10.4.1.8 Month state (page 329)
 - 4.10.4.1.9 Week state (page 330)
 - 4.10.4.1.10 Time state (page 331)
 - 4.10.4.1.11 Local Date and Time state (page 332)
 - 4.10.4.1.12 Number state (page 333)
 - 4.10.4.1.13 Range state (page 334)
 - 4.10.4.1.14 Checkbox state (page 335)
 - 4.10.4.1.15 Radio Button state (page 336)
 - 4.10.4.1.16 File Upload state (page 337)
 - 4.10.4.1.17 Submit Button state (page 338)
 - 4.10.4.1.18 Image Button state (page 338)
 - 4.10.4.1.19 Reset Button state (page 340)
 - 4.10.4.1.20 Button state (page 341)
 - 4.10.4.2 Common input element attributes (page 341)
 - 4.10.4.2.1 The autocomplete attribute (page 341)
 - 4.10.4.2.2 The list attribute (page 342)
 - 4.10.4.2.3 The readonly attribute (page 343)
 - 4.10.4.2.4 The size attribute (page 343)
 - 4.10.4.2.5 The required attribute (page 343)
 - 4.10.4.2.6 The maxlength attribute (page 343)
 - 4.10.4.2.7 The pattern attribute (page 343)
 - 4.10.4.2.8 The min and max attributes (page 344)
 - 4.10.4.2.9 The step attribute (page 345)
 - 4.10.4.3 Common input element APIs (page 345)
 - 4.10.4.4 Common event behaviors (page 347)
 - 4.10.5 The button element (page 348)
 - 4.10.6 The select element (page 350)
 - 4.10.7 The datalist element (page 353)
 - 4.10.8 The optgroup element (page 353)
 - 4.10.9 The option element (page 354)
 - 4.10.10 The textarea element (page 356)
 - 4.10.11 The output element (page 359)
 - 4.10.12 Association of controls and forms (page 361)
 - 4.10.13 Attributes common to form controls (page 362)
 - 4.10.13.1 Naming form controls (page 362)
 - 4.10.13.2 Enabling and disabling form controls (page 362)
 - 4.10.13.3 A form control's value (page 362)
 - 4.10.13.4 Autofocusing a form control (page 362)
 - 4.10.13.5 Limiting user input length (page 363)
 - 4.10.13.6 Form submission (page 363)
 - 4.10.14 Constraints (page 364)
 - 4.10.14.1 Definitions (page 364)
 - 4.10.14.2 Constraint validation (page 365)
 - 4.10.14.3 The constraint validation API (page 366)
 - 4.10.15 Form submission (page 368)
 - 4.10.15.1 URL-encoded form data (page 373)
 - 4.10.15.2 Multipart form data (page 375)
 - 4.10.15.3 Plain text form data (page 375)
 - 4.10.16 Resetting a form (page 376)
 - 4.10.17 Event dispatch (page 376)
- 4.11 Interactive elements (page 376)

- 4.11.1 The details element (page 376)
- 4.11.2 The datagrid element (page 377)
 - 4.11.2.1 The datagrid data model (page 378)
 - 4.11.2.2 How rows are identified (page 379)
 - 4.11.2.3 The data provider interface (page 379)
 - 4.11.2.4 The default data provider (page 383)
 - 4.11.2.4.1 Common default data provider method definitions for cells (page 388)
 - 4.11.2.5 Populating the datagrid element (page 390)
 - 4.11.2.6 Updating the datagrid (page 395)
 - 4.11.2.7 Requirements for interactive user agents (page 395)
 - 4.11.2.8 The selection (page 396)
 - 4.11.2.9 Columns and captions (page 398)
- 4.11.3 The command element (page 398)
- 4.11.4 The bb element (page 401)
 - 4.11.4.1 Browser button types (page 402)
 - 4.11.4.1.1 The *make application* state (page 402)
- 4.11.5 The menu element (page 403)
 - 4.11.5.1 Introduction (page 404)
 - 4.11.5.2 Building menus and tool bars (page 404)
 - 4.11.5.3 Context menus (page 405)
 - 4.11.5.4 Toolbars (page 406)
- 4.11.6 Commands (page 406)
 - 4.11.6.1 Using the a element to define a command (page 409)
 - 4.11.6.2 Using the button element to define a command (page 409)
 - 4.11.6.3 Using the input element to define a command (page 410)
 - 4.11.6.4 Using the option element to define a command (page 410)
 - 4.11.6.5 Using the command element to define a command (page 411)
 - 4.11.6.6 Using the bb element to define a command (page 412)
- 4.12 Miscellaneous elements (page 412)
 - 4.12.1 The legend element (page 412)
 - 4.12.2 The div element (page 413)

5 Web browsers (page 414)

- 5.1 Browsing contexts (page 414)
 - 5.1.1 Nested browsing contexts (page 415)
 - 5.1.1.1 Navigating nested browsing contexts in the DOM (page 416)
 - 5.1.2 Auxiliary browsing contexts (page 416)
 - 5.1.2.1 Navigating auxiliary browsing contexts in the DOM (page 416)
 - 5.1.3 Secondary browsing contexts (page 416)
 - 5.1.4 Security (page 417)
 - 5.1.5 Groupings of browsing contexts (page 417)
 - 5.1.6 Browsing context names (page 417)

- 5.2 The default view (page 419)
 - 5.2.1 Security (page 421)
 - 5.2.2 APIs for creating and navigating browsing contexts by name (page 421)
 - 5.2.3 Accessing other browsing contexts (page 422)
- 5.3 Origin (page 423)
 - 5.3.1 Relaxing the same-origin restriction (page 426)
- 5.4 Scripting (page 427)
 - 5.4.1 Script execution contexts (page 428)
 - 5.4.2 Event loops (page 429)
 - 5.4.2.1 Generic task sources (page 430)
 - 5.4.3 Security exceptions (page 430)
 - 5.4.4 The javascript: protocol (page 430)
 - 5.4.5 Events (page 431)
 - 5.4.5.1 Event handler attributes (page 432)
 - 5.4.5.2 Event firing (page 435)
 - 5.4.5.3 Events and the Window object (page 437)
 - 5.4.5.4 Runtime script errors (page 437)
- 5.5 User prompts (page 438)
 - 5.5.1 Simple dialogs (page 438)
 - 5.5.2 Printing (page 438)
 - 5.5.3 Dialogs implemented using separate documents (page 439)
 - 5.5.4 Notifications (page 440)
- 5.6 System state and capabilities (page 443)
 - 5.6.1 Client identification (page 443)
 - 5.6.2 Custom protocol and content handlers (page 444)
 - 5.6.2.1 Security and privacy (page 445)
 - 5.6.2.2 Sample user interface (page 447)
 - 5.6.3 Client abilities (page 448)
- 5.7 Offline Web applications (page 448)
 - 5.7.1 Introduction (page 448)
 - 5.7.2 Application caches (page 448)
 - 5.7.3 The cache manifest syntax (page 450)
 - 5.7.3.1 A sample manifest (page 450)
 - 5.7.3.2 Writing cache manifests (page 450)
 - 5.7.3.3 Parsing cache manifests (page 452)
 - 5.7.4 Updating an application cache (page 455)
 - 5.7.5 Processing model (page 460)
 - 5.7.5.1 Changes to the networking model (page 462)
 - 5.7.6 Application cache API (page 463)
 - 5.7.7 Browser state (page 467)
- 5.8 Session history and navigation (page 467)
 - 5.8.1 The session history of browsing contexts (page 467)
 - 5.8.2 The History interface (page 468)
 - 5.8.3 Activating state object entries (page 470)
 - 5.8.4 The Location interface (page 471)
 - 5.8.4.1 Security (page 472)
 - 5.8.5 Implementation notes for session history (page 472)
- 5.9 Browsing the Web (page 473)
 - 5.9.1 Navigating across documents (page 473)

- 5.9.2 Page load processing model for HTML files (page 476)
- 5.9.3 Page load processing model for XML files (page 477)
- 5.9.4 Page load processing model for text files (page 478)
- 5.9.5 Page load processing model for images (page 478)
- 5.9.6 Page load processing model for content that uses plugins (page 479)
- 5.9.7 Page load processing model for inline content that doesn't have a DOM (page 479)
- 5.9.8 Navigating to a fragment identifier (page 479)
- 5.9.9 History traversal (page 480)
- 5.9.10 Closing a browsing context (page 482)
- 5.10 Structured client-side storage (page 482)
 - 5.10.1 Storing name/value pairs (page 482)
 - 5.10.1.1 Introduction (page 482)
 - 5.10.1.2 The Storage interface (page 484)
 - 5.10.1.3 The sessionStorage attribute (page 485)
 - 5.10.1.4 The localStorage attribute (page 486)
 - 5.10.1.5 The storage event (page 486)
 - 5.10.1.5.1 Event definition (page 487)
 - 5.10.1.6 Threads (page 487)
 - 5.10.2 Database storage (page 488)
 - 5.10.2.1 Introduction (page 488)
 - 5.10.2.2 Databases (page 488)
 - 5.10.2.3 Executing SQL statements (page 490)
 - 5.10.2.4 Database query results (page 491)
 - 5.10.2.5 Errors (page 492)
 - 5.10.2.6 Processing model (page 493)
 - 5.10.3 Disk space (page 494)
 - 5.10.4 Privacy (page 494)
 - 5.10.4.1 User tracking (page 494)
 - 5.10.4.2 Cookie resurrection (page 496)
 - 5.10.5 Security (page 496)
 - 5.10.5.1 DNS spoofing attacks (page 496)
 - 5.10.5.2 Cross-directory attacks (page 496)
 - 5.10.5.3 Implementation risks (page 496)
 - 5.10.5.4 SQL and user agents (page 497)
 - 5.10.5.5 SQL injection (page 497)
- 5.11 Links (page 497)
 - 5.11.1 Hyperlink elements (page 497)
 - 5.11.2 Following hyperlinks (page 498)
 - 5.11.2.1 Hyperlink auditing (page 499)
 - 5.11.3 Link types (page 500)
 - 5.11.3.1 Link type "alternate" (page 502)
 - 5.11.3.2 Link type "archives" (page 503)
 - 5.11.3.3 Link type "author" (page 503)
 - 5.11.3.4 Link type "bookmark" (page 503)
 - 5.11.3.5 Link type "external" (page 504)
 - 5.11.3.6 Link type "feed" (page 504)
 - 5.11.3.7 Link type "help" (page 505)
 - 5.11.3.8 Link type "icon" (page 505)
 - 5.11.3.9 Link type "license" (page 507)

- 5.11.3.10 Link type "nofollow" (page 507)
- 5.11.3.11 Link type "noreferrer" (page 507)
- 5.11.3.12 Link type "pingback" (page 507)
- 5.11.3.13 Link type "prefetch" (page 507)
- 5.11.3.14 Link type "search" (page 507)
- 5.11.3.15 Link type "stylesheet" (page 508)
- 5.11.3.16 Link type "sidebar" (page 508)
- 5.11.3.17 Link type "tag" (page 508)
- 5.11.3.18 Hierarchical link types (page 508)
 - 5.11.3.18.1 Link type "index" (page 509)
 - 5.11.3.18.2 Link type "up" (page 509)
- 5.11.3.19 Sequential link types (page 510)
 - 5.11.3.19.1 Link type "first" (page 510)
 - 5.11.3.19.2 Link type "last" (page 510)
 - 5.11.3.19.3 Link type "next" (page 510)
 - 5.11.3.19.4 Link type "prev" (page 511)
- 5.11.3.20 Other link types (page 511)

6 User Interaction (page 513)

- 6.1 Introduction (page 513)
- 6.2 The hidden attribute (page 513)
- 6.3 Activation (page 514)
- 6.4 Scrolling elements into view (page 514)
- 6.5 Focus (page 514)
 - 6.5.1 Sequential focus navigation (page 515)
 - 6.5.2 Focus management (page 516)
 - 6.5.3 Document-level focus APIs (page 516)
 - 6.5.4 Element-level focus APIs (page 517)
- 6.6 The text selection APIs (page 517)
 - 6.6.1 APIs for the browsing context selection (page 518)
 - 6.6.2 APIs for the text field selections (page 520)
- 6.7 The contenteditable attribute (page 521)
 - 6.7.1 User editing actions (page 522)
 - 6.7.2 Making entire documents editable (page 524)
- 6.8 Drag and drop (page 525)
 - 6.8.1 Introduction (page 525)
 - 6.8.2 The DragEvent and DataTransfer interfaces (page 525)
 - 6.8.3 Events fired during a drag-and-drop action (page 527)
 - 6.8.4 Drag-and-drop processing model (page 529)
 - 6.8.4.1 When the drag-and-drop operation starts or ends in another document (page 533)
 - 6.8.4.2 When the drag-and-drop operation starts or ends in another application (page 533)
 - 6.8.5 The draggable attribute (page 534)
 - 6.8.6 Copy and paste (page 534)
 - 6.8.6.1 Copy to clipboard (page 534)
 - 6.8.6.2 Cut to clipboard (page 535)
 - 6.8.6.3 Paste from clipboard (page 535)
 - 6.8.6.4 Paste from selection (page 535)
 - 6.8.7 Security risks in the drag-and-drop model (page 535)
- 6.9 Undo history (page 536)

- 6.9.1 The UndoManager interface (page 536)
 - 6.9.2 Undo: moving back in the undo transaction history (page 538)
 - 6.9.3 Redo: moving forward in the undo transaction history (page 539)
 - 6.9.4 The UndoManagerEvent interface and the undo and redo events (page 539)
 - 6.9.5 Implementation notes (page 540)
 - 6.10 Command APIs (page 540)
- 7 Communication (page 546)
- 7.1 Event definitions (page 546)
 - 7.2 Server-sent events (page 547)
 - 7.2.1 The RemoteEventTarget interface (page 547)
 - 7.2.2 Connecting to an event stream (page 547)
 - 7.2.3 Parsing an event stream (page 549)
 - 7.2.4 Interpreting an event stream (page 550)
 - 7.2.5 Notes (page 553)
 - 7.3 Web sockets (page 553)
 - 7.3.1 Introduction (page 553)
 - 7.3.2 The WebSocket interface (page 554)
 - 7.3.3 WebSocket Events (page 555)
 - 7.3.4 The Web Socket protocol (page 555)
 - 7.3.4.1 Client-side requirements (page 555)
 - 7.3.4.1.1 Handshake (page 555)
 - 7.3.4.1.2 Data framing (page 560)
 - 7.3.4.2 Server-side requirements (page 562)
 - 7.3.4.2.1 Minimal handshake (page 562)
 - 7.3.4.2.2 Handshake details (page 562)
 - 7.3.4.2.3 Data framing (page 563)
 - 7.3.4.3 Closing the connection (page 564)
 - 7.4 Cross-document messaging (page 564)
 - 7.4.1 Introduction (page 564)
 - 7.4.2 Security (page 565)
 - 7.4.2.1 Authors (page 565)
 - 7.4.2.2 User agents (page 565)
 - 7.4.3 Posting text (page 565)
 - 7.4.4 Posting message ports (page 566)
 - 7.4.5 Posting structured data (page 567)
 - 7.5 Channel messaging (page 567)
 - 7.5.1 Introduction (page 567)
 - 7.5.2 Message channels (page 567)
 - 7.5.3 Message ports (page 568)
 - 7.5.3.1 Ports and browsing contexts (page 571)
 - 7.5.3.2 Ports and garbage collection (page 571)
- 8 The HTML syntax (page 573)
- 8.1 Writing HTML documents (page 573)
 - 8.1.1 The DOCTYPE (page 573)
 - 8.1.2 Elements (page 574)
 - 8.1.2.1 Start tags (page 576)
 - 8.1.2.2 End tags (page 576)

- 8.1.2.3 Attributes (page 576)
 - 8.1.2.4 Optional tags (page 578)
 - 8.1.2.5 Restrictions on content models (page 579)
 - 8.1.2.6 Restrictions on the contents of CDATA and RCDATA elements (page 580)
 - 8.1.3 Text (page 581)
 - 8.1.3.1 Newlines (page 581)
 - 8.1.4 Character references (page 581)
 - 8.1.5 CDATA sections (page 582)
 - 8.1.6 Comments (page 582)
- 8.2 Parsing HTML documents (page 582)
- 8.2.1 Overview of the parsing model (page 583)
 - 8.2.2 The input stream (page 584)
 - 8.2.2.1 Determining the character encoding (page 585)
 - 8.2.2.2 Character encoding requirements (page 589)
 - 8.2.2.3 Preprocessing the input stream (page 591)
 - 8.2.2.4 Changing the encoding while parsing (page 592)
 - 8.2.3 Parse state (page 592)
 - 8.2.3.1 The insertion mode (page 592)
 - 8.2.3.2 The stack of open elements (page 594)
 - 8.2.3.3 The list of active formatting elements (page 596)
 - 8.2.3.4 The element pointers (page 597)
 - 8.2.3.5 The scripting state (page 597)
 - 8.2.4 Tokenization (page 597)
 - 8.2.4.1 Data state (page 598)
 - 8.2.4.2 Character reference data state (page 599)
 - 8.2.4.3 Tag open state (page 599)
 - 8.2.4.4 Close tag open state (page 600)
 - 8.2.4.5 Tag name state (page 601)
 - 8.2.4.6 Before attribute name state (page 601)
 - 8.2.4.7 Attribute name state (page 602)
 - 8.2.4.8 After attribute name state (page 603)
 - 8.2.4.9 Before attribute value state (page 604)
 - 8.2.4.10 Attribute value (double-quoted) state (page 604)
 - 8.2.4.11 Attribute value (single-quoted) state (page 605)
 - 8.2.4.12 Attribute value (unquoted) state (page 605)
 - 8.2.4.13 Character reference in attribute value state (page 605)
 - 8.2.4.14 After attribute value (quoted) state (page 606)
 - 8.2.4.15 Self-closing start tag state (page 606)
 - 8.2.4.16 Bogus comment state (page 606)
 - 8.2.4.17 Markup declaration open state (page 607)
 - 8.2.4.18 Comment start state (page 607)
 - 8.2.4.19 Comment start dash state (page 608)
 - 8.2.4.20 Comment state (page 608)
 - 8.2.4.21 Comment end dash state (page 608)
 - 8.2.4.22 Comment end state (page 608)
 - 8.2.4.23 DOCTYPE state (page 609)
 - 8.2.4.24 Before DOCTYPE name state (page 609)
 - 8.2.4.25 DOCTYPE name state (page 610)
 - 8.2.4.26 After DOCTYPE name state (page 610)

- 8.2.4.27 Before DOCTYPE public identifier state (page 611)
- 8.2.4.28 DOCTYPE public identifier (double-quoted) state (page 611)
- 8.2.4.29 DOCTYPE public identifier (single-quoted) state (page 611)
- 8.2.4.30 After DOCTYPE public identifier state (page 612)
- 8.2.4.31 Before DOCTYPE system identifier state (page 612)
- 8.2.4.32 DOCTYPE system identifier (double-quoted) state (page 613)
- 8.2.4.33 DOCTYPE system identifier (single-quoted) state (page 613)
- 8.2.4.34 After DOCTYPE system identifier state (page 614)
- 8.2.4.35 Bogus DOCTYPE state (page 614)
- 8.2.4.36 CDATA section state (page 614)
- 8.2.4.37 Tokenizing character references (page 615)
- 8.2.5 Tree construction (page 617)
 - 8.2.5.1 Creating and inserting elements (page 618)
 - 8.2.5.2 Closing elements that have implied end tags (page 620)
 - 8.2.5.3 Foster parenting (page 620)
 - 8.2.5.4 The "initial" insertion mode (page 620)
 - 8.2.5.5 The "before html" insertion mode (page 623)
 - 8.2.5.6 The "before head" insertion mode (page 624)
 - 8.2.5.7 The "in head" insertion mode (page 625)
 - 8.2.5.8 The "in head noscript" insertion mode (page 626)
 - 8.2.5.9 The "after head" insertion mode (page 627)
 - 8.2.5.10 The "in body" insertion mode (page 628)
 - 8.2.5.11 The "in CDATA/RCDATA" insertion mode (page 640)
 - 8.2.5.12 The "in table" insertion mode (page 642)
 - 8.2.5.13 The "in caption" insertion mode (page 644)
 - 8.2.5.14 The "in column group" insertion mode (page 645)
 - 8.2.5.15 The "in table body" insertion mode (page 646)
 - 8.2.5.16 The "in row" insertion mode (page 647)
 - 8.2.5.17 The "in cell" insertion mode (page 648)
 - 8.2.5.18 The "in select" insertion mode (page 649)
 - 8.2.5.19 The "in select in table" insertion mode (page 651)
 - 8.2.5.20 The "in foreign content" insertion mode (page 651)
 - 8.2.5.21 The "after body" insertion mode (page 653)
 - 8.2.5.22 The "in frameset" insertion mode (page 653)
 - 8.2.5.23 The "after frameset" insertion mode (page 654)
 - 8.2.5.24 The "after after body" insertion mode (page 655)
 - 8.2.5.25 The "after after frameset" insertion mode (page 656)
- 8.2.6 The end (page 656)
- 8.2.7 Coercing an HTML DOM into an infoset (page 657)
- 8.3 Namespaces (page 658)
- 8.4 Serializing HTML fragments (page 658)
- 8.5 Parsing HTML fragments (page 661)
- 8.6 Named character references (page 662)
- 9 Rendering and user-agent behavior (page 672)
 - 9.1 Rendering and the DOM (page 672)

- 9.2 Rendering and menus/toolbars (page 672)
 - 9.2.1 The 'icon' property (page 672)
- 9.3 Obsolete elements, attributes, and APIs (page 673)
 - 9.3.1 The body element (page 673)
 - 9.3.2 The applet element (page 674)

10 Things that you can't do with this specification because they are better handled using other technologies that are further described herein (page 675)

- 10.1 Localization (page 675)
- 10.2 Declarative 2D vector graphics and animation (page 675)
- 10.3 Declarative 3D scenes (page 675)
- 10.4 Timers (page 675)

Index (page 677)

References (page 678)

Acknowledgements (page 679)

1 Introduction

1.1 Background

This section is non-normative.

The World Wide Web's markup language has always been HTML. HTML was primarily designed as a language for semantically describing scientific documents, although its general design and adaptations over the years has enabled it to be used to describe a number of other types of documents.

The main area that has not been adequately addressed by HTML is a vague subject referred to as Web Applications. This specification attempts to rectify this, while at the same time updating the HTML specifications to address issues raised in the past few years.

1.2 Scope

This section is non-normative.

This specification is limited to providing a semantic-level markup language and associated semantic-level scripting APIs for authoring accessible pages on the Web ranging from static documents to dynamic applications.

The scope of this specification does not include providing mechanisms for media-specific customization of presentation (although default rendering rules for Web browsers are included at the end of this specification, and several mechanisms for hooking into CSS are provided as part of the language).

The scope of this specification does not include documenting every HTML or DOM feature supported by Web browsers. Browsers support many features that are considered to be very bad for accessibility or that are otherwise inappropriate. For example, the `blink` element is clearly presentational and authors wishing to cause text to blink should instead use CSS.

The scope of this specification is not to describe an entire operating system. In particular, hardware configuration software, image manipulation tools, and applications that users would be expected to use with high-end workstations on a daily basis are out of scope. In terms of applications, this specification is targeted specifically at applications that would be expected to be used by users on an occasional basis, or regularly but from disparate locations, with low CPU requirements. For instance online purchasing systems, searching systems, games (especially multiplayer online games), public telephone books or address books, communications software (e-mail clients, instant messaging clients, discussion software), document editing software, etc.

For sophisticated cross-platform applications, there already exist several proprietary solutions (such as Mozilla's XUL, Adobe's Flash, or Microsoft's Silverlight). These solutions are evolving faster than any standards process could follow, and the requirements are evolving even faster. These systems are also significantly more complicated to specify, and are orders of magnitude more difficult to achieve interoperability with, than the solutions described in this document. Platform-specific solutions for such sophisticated applications (for example the Mac OS X Core APIs) are even further ahead.

1.3 History

Work on HTML5 originally started in late 2003, as a proof of concept to show that it was possible to extend HTML4's forms to provide many of the features that XForms 1.0 introduced, without requiring browsers to implement rendering engines that were incompatible with existing HTML Web pages. At this early stage, while the draft was already publicly available, and input was already being solicited from all sources, the specification was only under Opera Software's copyright.

In early 2004, some of the principles that underly this effort, as well as an early draft proposal covering just forms-related features, were presented to the W3C jointly by Mozilla and Opera at a workshop discussing the future of Web Applications on the Web. The proposal was rejected on the grounds that the proposal conflicted with the previously chosen direction for the Web's evolution.

Shortly thereafter, Apple, Mozilla, and Opera jointly announced their intent to continue working on the effort. A public mailing list was created, and the drafts were moved to the WHATWG site. The copyright was subsequently amended to be jointly owned by all three vendors, and to allow reuse of the specifications.

In 2006, the W3C expressed interest in the specification, and created a working group chartered to work with the WHATWG on the development of the HTML5 specifications. The working group opened in 2007. Apple, Mozilla, and Opera allowed the W3C to publish the specifications under the W3C copyright, while keeping versions with the less restrictive license on the WHATWG site.

Since then, both groups have been working together.

1.4 Relationships to other specifications

1.4.1 Relationship to HTML 4.01 and DOM2 HTML

This section is non-normative.

This specification represents a new version of HTML4, along with a new version of the associated DOM2 HTML API. Migration from HTML4 to the format and APIs described in this specification should in most cases be straightforward, as care has been taken to ensure that backwards-compatibility is retained. [HTML4] [DOM2HTML]

1.4.2 Relationship to XHTML 1.x

This section is non-normative.

This specification is intended to replace XHTML 1.0 as the normative definition of the XML serialization of the HTML vocabulary. [XHTML10]

While this specification updates the semantics and requirements of the vocabulary defined by XHTML Modularization 1.1 and used by XHTML 1.1, it does not attempt to provide a replacement for the modularization scheme defined and used by those (and other) specifications, and therefore cannot be considered a complete replacement for them. [XHTMLMOD] [XHTML11]

Thus, authors and implementors who do not need such a modularization scheme can consider this specification a replacement for XHTML 1.x, but those who do need such a mechanism are encouraged to continue using the XHTML 1.1 line of specifications.

1.4.3 Relationship to XHTML2 and XForms

This section is non-normative.

XHTML2 defines a new vocabulary with features for hyperlinks, multimedia content, annotating document edits, rich metadata, declarative interactive forms, and describing the semantics of human literary works such as poems and scientific papers. [XHTML2]

XForms similarly defines a new vocabulary with features for complex data entry, such as tax forms or insurance forms.

However, XHTML2 and XForms lack features to express the semantics of many of the non-document types of content often seen on the Web. For instance, they are not well-suited for marking up forum sites, auction sites, search engines, online shops, mapping applications, e-mail applications, word processors, real-time strategy games, and the like.

This specification aims to extend HTML so that it is also suitable in these contexts.

XHTML2, XForms, and this specification all use different namespaces and therefore can all be implemented in the same XML processor.

1.4.4 Relationship to XUL, Flash, Silverlight, and other proprietary UI languages

This section is non-normative.

This specification is independent of the various proprietary UI languages that various vendors provide. As an open, vendor-neutral language, HTML provides for a solution to the same problems without the risk of vendor lock-in.

1.5 HTML vs XHTML

This section is non-normative.

This specification defines an abstract language for describing documents and applications, and some APIs for interacting with in-memory representations of resources that use this language.

The in-memory representation is known as "DOM5 HTML", or "the DOM" for short.

There are various concrete syntaxes that can be used to transmit resources that use this abstract language, two of which are defined in this specification.

The first such concrete syntax is "HTML5". This is the format recommended for most authors. It is compatible with all legacy Web browsers. If a document is transmitted with the MIME type `text/html`, then it will be processed as an "HTML5" document by Web browsers.

The second concrete syntax uses XML, and is known as "XHTML5". When a document is transmitted with an XML MIME type, such as application/xhtml+xml, then it is processed by an XML processor by Web browsers, and treated as an "XHTML5" document. Authors are reminded that the processing for XML and HTML differs; in particular, even minor syntax errors will prevent an XML document from being rendered fully, whereas they would be ignored in the "HTML5" syntax.

The "DOM5 HTML", "HTML5", and "XHTML5" representations cannot all represent the same content. For example, namespaces cannot be represented using "HTML5", but they are supported in "DOM5 HTML" and "XHTML5". Similarly, documents that use the noscript feature can be represented using "HTML5", but cannot be represented with "XHTML5" and "DOM5 HTML". Comments that contain the string "-->" can be represented in "DOM5 HTML" but not in "HTML5" and "XHTML5". And so forth.

1.6 Structure of this specification

This section is non-normative.

This specification is divided into the following major sections:

Common Infrastructure (page 23)

The conformance classes, algorithms, definitions, and the common underpinnings of the rest of the specification.

The DOM (page 76)

Documents are built from elements. These elements form a tree using the DOM. This section defines the features of this DOM, as well as introducing the features common to all elements, and the concepts used in defining elements.

Elements (page 107)

Each element has a predefined meaning, which is explained in this section. User agent requirements for how to handle each element are also given, along with rules for authors on how to use the element.

Web Browsers (page 414)

HTML documents do not exist in a vacuum — this section defines many of the features that affect environments that deal with multiple pages, links between pages, and running scripts.

User Interaction (page 513)

HTML documents can provide a number of mechanisms for users to interact with and modify content, which are described in this section.

The Communication APIs (page 546)

Applications written in HTML often require mechanisms to communicate with remote servers, as well as communicating with other applications from different domains running on the same client.

The Language Syntax (page 573)

All of these features would be for naught if they couldn't be represented in a serialized form and sent to other people, and so this section defines the syntax of HTML, along with rules for how to parse HTML.

There are also a couple of appendices, defining rendering rules (page 672) for Web browsers and listing areas that are out of scope (page 675) for this specification.

1.6.1 How to read this specification

This specification should be read like all other specifications. First, it should be read cover-to-cover, multiple times. Then, it should be read backwards at least once. Then it should be read by picking random sections from the contents list and following all the cross-references.

1.6.2 Typographic conventions

This is a definition, requirement, or explanation.

Note: This is a note.

|| This is an example.

** This is an open issue.

⚠Warning! This is a warning.

The defining instance of a term is marked up like **this**. Uses of that term are marked up like *this* (page 22) or like *this* (page 22).

The defining instance of an element, attribute, or API is marked up like **this**. References to that element, attribute, or API are marked up like *this*.

Other code fragments are marked up like *this*.

Variables are marked up like *this*.

```
interface Example {  
    // this is an IDL definition  
};
```

2 Common infrastructure

2.1 Terminology

This specification refers to both HTML and XML attributes and DOM attributes, often in the same context. When it is not clear which is being referred to, they are referred to as **content attributes** for HTML and XML attributes, and **DOM attributes** for those from the DOM. Similarly, the term "properties" is used for both ECMAScript object properties and CSS properties. When these are ambiguous they are qualified as object properties and CSS properties respectively.

The term **HTML documents** (page 76) is sometimes used in contrast with XML documents (page 76) to specifically mean documents that were parsed using an HTML parser (page 582) (as opposed to using an XML parser or created purely through the DOM).

Generally, when the specification states that a feature applies to HTML or XHTML, it also includes the other. When a feature specifically only applies to one of the two languages, it is called out by explicitly stating that it does not apply to the other format, as in "for HTML, ... (this does not apply to XHTML)".

This specification uses the term *document* to refer to any use of HTML, ranging from short static documents to long essays or reports with rich multimedia, as well as to fully-fledged interactive applications.

For simplicity, terms such as *shown*, *displayed*, and *visible* might sometimes be used when referring to the way a document is rendered to the user. These terms are not meant to imply a visual medium; they must be considered to apply to other media in equivalent ways.

Some of the algorithms in this specification, for historical reasons, require the user agent to **pause** until some condition has been met. While a user agent is paused, it must ensure that no scripts execute (e.g. no event handlers, no timers, etc). User agents should remain responsive to user input while paused, however, albeit without letting the user interact with Web pages where that would involve invoking any script.

2.1.1 XML

To ease migration from HTML to XHTML, UAs conforming to this specification will place elements in HTML in the `http://www.w3.org/1999/xhtml` namespace, at least for the purposes of the DOM and CSS. The term "**elements in the HTML namespace**", or "**HTML elements**" for short, when used in this specification, thus refers to both HTML and XHTML elements.

Unless otherwise stated, all elements defined or mentioned in this specification are in the `http://www.w3.org/1999/xhtml` namespace, and all attributes defined or mentioned in this specification have no namespace (they are in the per-element partition).

When an XML name, such as an attribute or element name, is referred to in the form `prefix:localName`, as in `xml:id` or `svg:rect`, it refers to a name with the local name `localName` and the namespace given by the prefix, as defined by the following table:

`xml`

`http://www.w3.org/XML/1998/namespace`

html

<http://www.w3.org/1999/xhtml>

svg

<http://www.w3.org/2000/svg>

Attribute names are said to be **XML-compatible** if they match the Name production defined in XML, they contain no U+003A COLON (:) characters, and their first three characters are not an ASCII case-insensitive (page 31) match for the string "xml". [XML]

2.1.2 DOM trees

The term **root element**, when not explicitly qualified as referring to the document's root element, means the furthest ancestor element node of whatever node is being discussed, or the node itself if it has no ancestors. When the node is a part of the document, then that is indeed the document's root element; however, if the node is not currently part of the document tree, the root element will be an orphaned node.

The Document of a Node (such as an element) is the Document that the Node's ownerDocument DOM attribute returns.

An element is said to have been **inserted into a document** when its root element (page 24) changes and is now the document's root element (page 24). If a Node is in a Document then that Document is always the Node's Document, and the Node's ownerDocument DOM attribute thus always returns that Document.

The term **tree order** means a pre-order, depth-first traversal of DOM nodes involved (through the parentNode/childNodes relationship).

When it is stated that some element or attribute is **ignored**, or treated as some other value, or handled as if it was something else, this refers only to the processing of the node after it is in the DOM. A user agent must not mutate the DOM in such situations.

The term **text node** refers to any Text node, including CDATASection nodes; specifically, any Node with node type TEXT_NODE (3) or CDATA_SECTION_NODE (4). [DOM3CORE]

2.1.3 Scripting

The construction "a Foo object", where Foo is actually an interface, is sometimes used instead of the more accurate "an object implementing the interface Foo".

A DOM attribute is said to be *getting* when its value is being retrieved (e.g. by author script), and is said to be *setting* when a new value is assigned to it.

If a DOM object is said to be **live**, then that means that any attributes returning that object must always return the same object (not a new object each time), and the attributes and methods on that object must operate on the actual underlying data, not a snapshot of the data.

The terms *fire* and *dispatch* are used interchangeably in the context of events, as in the DOM Events specifications. [DOM3EVENTS]

2.1.4 Plugins

The term **plugin** is used to mean any content handler, typically a third-party content handler, for Web content types that are not supported by the user agent natively, or for content types that do not expose a DOM, that supports rendering the content as part of the user agent's interface.

One example of a plugin would be a PDF viewer that is instantiated in a browsing context (page 414) when the user navigates to a PDF file. This would count as a plugin regardless of whether the party that implemented the PDF viewer component was the same as that which implemented the user agent itself. However, a PDF viewer application that launches separate from the user agent (as opposed to using the same interface) is not a plugin by this definition.

Note: *This specification does not define a mechanism for interacting with plugins, as it is expected to be user-agent- and platform-specific. Some UAs might opt to support a plugin mechanism such as the Netscape Plugin API; others might use remote content converters or have built-in support for certain types. [NPAPI]*

⚠ Warning! *Browsers should take extreme care when interacting with external content intended for plugins (page 25). When third-party software is run with the same privileges as the user agent itself, vulnerabilities in the third-party software become as dangerous as those in the user agent.*

2.1.5 Character encodings

An **ASCII-compatible character encoding** is one that is a superset of US-ASCII (specifically, ANSI_X3.4-1968) for bytes in the set 0x09, 0x0A, 0x0C, 0x0D, 0x20 - 0x22, 0x26, 0x27, 0x2C - 0x3F, 0x41 - 0x5A, and 0x61 - 0x7A.

2.2 Conformance requirements

All diagrams, examples, and notes in this specification are non-normative, as are all sections explicitly marked non-normative. Everything else in this specification is normative.

The key words "MUST", "MUST NOT", "REQUIRED", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in the normative parts of this document are to be interpreted as described in RFC2119. For readability, these words do not appear in all uppercase letters in this specification. [RFC2119]

Requirements phrased in the imperative as part of algorithms (such as "strip any leading space characters" or "return false and abort these steps") are to be interpreted with the meaning of the key word ("must", "should", "may", etc) used in introducing the algorithm.

This specification describes the conformance criteria for user agents (relevant to implementors) and documents (relevant to authors and authoring tool implementors).

Note: *There is no implied relationship between document conformance requirements and implementation conformance requirements. User agents*

are not free to handle non-conformant documents as they please; the processing model described in this specification applies to implementations regardless of the conformity of the input documents.

User agents fall into several (overlapping) categories with different conformance requirements.

Web browsers and other interactive user agents

Web browsers that support XHTML (page 28) must process elements and attributes from the HTML namespace (page 658) found in XML documents (page 76) as described in this specification, so that users can interact with them, unless the semantics of those elements have been overridden by other specifications.

A conforming XHTML processor would, upon finding an XHTML script element in an XML document, execute the script contained in that element. However, if the element is found within an XSLT transformation sheet (assuming the UA also supports XSLT), then the processor would instead treat the script element as an opaque element that forms part of the transform.

Web browsers that support HTML (page 29) must process documents labeled as text/html as described in this specification, so that users can interact with them.

Non-interactive presentation user agents

User agents that process HTML and XHTML documents purely to render non-interactive versions of them must comply to the same conformance criteria as Web browsers, except that they are exempt from requirements regarding user interaction.

Note: Typical examples of non-interactive presentation user agents are printers (static UAs) and overhead displays (dynamic UAs). It is expected that most static non-interactive presentation user agents will also opt to lack scripting support (page 26).

A non-interactive but dynamic presentation UA would still execute scripts, allowing forms to be dynamically submitted, and so forth. However, since the concept of "focus" is irrelevant when the user cannot interact with the document, the UA would not need to support any of the focus-related DOM APIs.

User agents with no scripting support

Implementations that do not support scripting (or which have their scripting features disabled entirely) are exempt from supporting the events and DOM interfaces mentioned in this specification. For the parts of this specification that are defined in terms of an events model or in terms of the DOM, such user agents must still act as if events and the DOM were supported.

Note: Scripting can form an integral part of an application. Web browsers that do not support scripting, or that have scripting disabled, might be unable to fully convey the author's intent.

Conformance checkers

Conformance checkers must verify that a document conforms to the applicable conformance criteria described in this specification. Automated conformance checkers

are exempt from detecting errors that require interpretation of the author's intent (for example, while a document is non-conforming if the content of a `blockquote` element is not a quote, conformance checkers running without the input of human judgement do not have to check that `blockquote` elements only contain quoted material).

Conformance checkers must check that the input document conforms when parsed without a browsing context (page 414) (meaning that no scripts are run, and that the parser's scripting flag (page 597) is disabled), and should also check that the input document conforms when parsed with a browsing context (page 414) in which scripts execute, and that the scripts never cause non-conforming states to occur other than transiently during script execution itself. (This is only a "SHOULD" and not a "MUST" requirement because it has been proven to be impossible. [HALTINGPROBLEM])

The term "HTML5 validator" can be used to refer to a conformance checker that itself conforms to the applicable requirements of this specification.

XML DTDs cannot express all the conformance requirements of this specification. Therefore, a validating XML processor and a DTD cannot constitute a conformance checker. Also, since neither of the two authoring formats defined in this specification are applications of SGML, a validating SGML system cannot constitute a conformance checker either.

To put it another way, there are three types of conformance criteria:

- 1. Criteria that can be expressed in a DTD.***
- 2. Criteria that cannot be expressed by a DTD, but can still be checked by a machine.***
- 3. Criteria that can only be checked by a human.***

A conformance checker must check for the first two. A simple DTD-based validator only checks for the first class of errors and is therefore not a conforming conformance checker according to this specification.

Data mining tools

Applications and tools that process HTML and XHTML documents for reasons other than to either render the documents or check them for conformance should act in accordance to the semantics of the documents that they process.

A tool that generates document outlines (page 142) but increases the nesting level for each paragraph and does not increase the nesting level for each section would not be conforming.

Authoring tools and markup generators

Authoring tools and markup generators must generate conforming documents. Conformance criteria that apply to authors also apply to authoring tools, where appropriate.

Authoring tools are exempt from the strict requirements of using elements only for their specified purpose, but only to the extent that authoring tools are not yet able to determine author intent.

For example, it is not conforming to use an address element for arbitrary contact information; that element can only be used for marking up contact information for the author of the document or section. However, since an authoring tool is likely unable to determine the difference, an authoring tool is exempt from that requirement.

Note: In terms of conformance checking, an editor is therefore required to output documents that conform to the same extent that a conformance checker will verify.

When an authoring tool is used to edit a non-conforming document, it may preserve the conformance errors in sections of the document that were not edited during the editing session (i.e. an editing tool is allowed to round-trip erroneous content). However, an authoring tool must not claim that the output is conformant if errors have been so preserved.

Authoring tools are expected to come in two broad varieties: tools that work from structure or semantic data, and tools that work on a What-You-See-Is-What-You-Get media-specific editing basis (WYSIWYG).

The former is the preferred mechanism for tools that author HTML, since the structure in the source information can be used to make informed choices regarding which HTML elements and attributes are most appropriate.

However, WYSIWYG tools are legitimate. WYSIWYG tools should use elements they know are appropriate, and should not use elements that they do not know to be appropriate. This might in certain extreme cases mean limiting the use of flow elements to just a few elements, like div, b, i, and span and making liberal use of the style attribute.

All authoring tools, whether WYSIWYG or not, should make a best effort attempt at enabling users to create well-structured, semantically rich, media-independent content.

Some conformance requirements are phrased as requirements on elements, attributes, methods or objects. Such requirements fall into two categories: those describing content model restrictions, and those describing implementation behavior. The former category of requirements are requirements on documents and authoring tools. The second category are requirements on user agents.

Conformance requirements phrased as algorithms or specific steps may be implemented in any manner, so long as the end result is equivalent. (In particular, the algorithms defined in this specification are intended to be easy to follow, and not intended to be performant.)

User agents may impose implementation-specific limits on otherwise unconstrained inputs, e.g. to prevent denial of service attacks, to guard against running out of memory, or to work around platform-specific limitations.

For compatibility with existing content and prior specifications, this specification describes two authoring formats: one based on XML (referred to as **XHTML5**), and one using a custom

format (page 582) inspired by SGML (referred to as **HTML5**). Implementations may support only one of these two formats, although supporting both is encouraged.

XHTML (page 28) documents (XML documents (page 76) using elements from the HTML namespace (page 658)) that use the new features described in this specification and that are served over the wire (e.g. by HTTP) must be sent using an XML MIME type such as application/xml or application/xhtml+xml and must not be served as text/html. [RFC3023]

Such XML documents may contain a DOCTYPE if desired, but this is not required to conform to this specification.

Note: According to the XML specification, XML processors are not guaranteed to process the external DTD subset referenced in the DOCTYPE. This means, for example, that using entity references for characters in XHTML documents is unsafe (except for <, >, &, " and ').

HTML documents (page 29), if they are served over the wire (e.g. by HTTP) must be labeled with the text/html MIME type.

The language in this specification assumes that the user agent expands all entity references, and therefore does not include entity reference nodes in the DOM. If user agents do include entity reference nodes in the DOM, then user agents must handle them as if they were fully expanded when implementing this specification. For example, if a requirement talks about an element's child text nodes, then any text nodes that are children of an entity reference that is a child of that element would be used as well. Entity references to unknown entities must be treated as if they contained just an empty text node for the purposes of the algorithms defined in this specification.

2.2.1 Dependencies

This specification relies on several other underlying specifications.

XML

Implementations that support XHTML5 must support some version of XML, as well as its corresponding namespaces specification, because XHTML5 uses an XML serialization with namespaces. [XML] [XMLNAMES]

DOM

The Document Object Model (DOM) is a representation — a model — of a document and its content. The DOM is not just an API; the conformance criteria of HTML implementations are defined, in this specification, in terms of operations on the DOM. [DOM3CORE]

Implementations must support some version of DOM Core and DOM Events, because this specification is defined in terms of the DOM, and some of the features are defined as extensions to the DOM Core interfaces. [DOM3CORE] [DOM3EVENTS]

ECMAScript

Implementations that use ECMAScript to implement the APIs defined in this specification must implement them in a manner consistent with the ECMAScript Bindings defined in

the Web IDL specification, as this specification uses that specification's terminology. [WebIDL]

Media Queries

Implementations must support some version of the Media Queries language. [MQ]

This specification does not require support of any particular network transport protocols, style sheet language, scripting language, or any of the DOM and WebAPI specifications beyond those described above. However, the language described by this specification is biased towards CSS as the styling language, ECMAScript as the scripting language, and HTTP as the network protocol, and several features assume that those languages and protocols are in use.

Note: This specification might have certain additional requirements on character encodings, image formats, audio formats, and video formats in the respective sections.

2.2.2 Features defined in other specifications

** this section will be removed at some point

Some elements are defined in terms of their DOM **textContent** attribute. This is an attribute defined on the Node interface in DOM3 Core. [DOM3CORE]

The interface **DOMTimeStamp** is defined in DOM3 Core. [DOM3CORE]

The rules for handling alternative style sheets are defined in the CSS object model specification. [CSSOM]

** See <http://dev.w3.org/cvsweb/~checkout~/csswg/cssom/Overview.html?content-type=text/html;%20charset=utf-8>

2.2.3 Common conformance requirements for APIs exposed to JavaScript

** This section will eventually be removed in favour of WebIDL.

** A lot of arrays/lists/collection (page 69)s in this spec assume zero-based indexes but use the term "*indexth*" liberally. We should define those to be zero-based and be clearer about this.

Unless otherwise specified, if a DOM attribute that is a floating point number type (`float`) is

** assigned an Infinity or Not-a-Number value, a **NOT_SUPPORTED_ERR** exception must be raised.

Unless otherwise specified, if a method with an argument that is a floating point number type

- ** (float) is passed an Infinity or Not-a-Number value, a **NOT_SUPPORTED_ERR** exception must be raised.

Unless otherwise specified, if a method is passed fewer arguments than is defined for that

- ** method in its IDL definition, a **NOT_SUPPORTED_ERR** exception must be raised.

Unless otherwise specified, if a method is passed more arguments than is defined for that method in its IDL definition, the excess arguments must be ignored.

2.3 Case-sensitivity and string comparison

This specification defines several comparison operators for strings.

Comparing two strings in a **case-sensitive** manner means comparing them exactly, codepoint for codepoint.

Comparing two strings in a **ASCII case-insensitive** manner means comparing them exactly, codepoint for codepoint, except that the characters in the range U+0041 .. U+005A (i.e. LATIN CAPITAL LETTER A to LATIN CAPITAL LETTER Z) and the corresponding characters in the range U+0061 .. U+007A (i.e. LATIN SMALL LETTER A to LATIN SMALL LETTER Z) are considered to also match.

Comparing two strings in a **compatibility caseless** manner means using the Unicode *compatibility caseless match* operation to compare the two strings. [UNICODECASE]

Converting a string to uppercase means replacing all characters in the range U+0061 .. U+007A (i.e. LATIN SMALL LETTER A to LATIN SMALL LETTER Z) with the corresponding characters in the range U+0041 .. U+005A (i.e. LATIN CAPITAL LETTER A to LATIN CAPITAL LETTER Z).

Converting a string to lowercase means replacing all characters in the range U+0041 .. U+005A (i.e. LATIN CAPITAL LETTER A to LATIN CAPITAL LETTER Z) with the corresponding characters in the range U+0061 .. U+007A (i.e. LATIN SMALL LETTER A to LATIN SMALL LETTER Z).

A string *pattern* is a **prefix match** for a string *s* when *pattern* is not longer than *s* and truncating *s* to *pattern*'s length leaves the two strings as matches of each other.

2.4 Common microsyntaxes

There are various places in HTML that accept particular data types, such as dates or numbers. This section describes what the conformance criteria for content in those formats is, and how to parse them.

- ** Need to go through the whole spec and make sure all the attribute values are clearly defined either in terms of microsyntaxes or in terms of other specs, or as "Text" or some such.

2.4.1 Common parser idioms

The **space characters**, for the purposes of this specification, are U+0020 SPACE, U+0009 CHARACTER TABULATION (tab), U+000A LINE FEED (LF), U+000C FORM FEED (FF), and U+000D CARRIAGE RETURN (CR).

The **White_Space characters** are those that have the Unicode property "White_Space".
[UNICODE]

Some of the micro-parsers described below follow the pattern of having an *input* variable that holds the string being parsed, and having a *position* variable pointing at the next character to parse in *input*.

For parsers based on this pattern, a step that requires the user agent to **collect a sequence of characters** means that the following algorithm must be run, with *characters* being the set of characters that can be collected:

1. Let *input* and *position* be the same variables as those of the same name in the algorithm that invoked these steps.
2. Let *result* be the empty string.
3. While *position* doesn't point past the end of *input* and the character at *position* is one of the *characters*, append that character to the end of *result* and advance *position* to the next character in *input*.
4. Return *result*.

The step **skip whitespace** means that the user agent must collect a sequence of characters (page 32) that are space characters (page 32). The step **skip White_Space characters** means that the user agent must collect a sequence of characters (page 32) that are White_Space (page 32) characters. In both cases, the collected characters are not used.
[UNICODE]

When a user agent is to **strip line breaks** from a string, the user agent must remove any U+000A LINE FEED (LF) and U+000D CARRIAGE RETURN (CR) characters from that string.

The **codepoint length** of a string is the number of Unicode codepoints in that string.

2.4.2 Boolean attributes

A number of attributes in HTML5 are **boolean attributes**. The presence of a boolean attribute on an element represents the true value, and the absence of the attribute represents the false value.

If the attribute is present, its value must either be the empty string or a value that is an ASCII case-insensitive (page 31) match for the attribute's canonical name, with no leading or trailing whitespace.

2.4.3 Numbers

2.4.3.1 Unsigned integers

A string is a **valid non-negative integer** if it consists of one or more characters in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9).

The **rules for parsing non-negative integers** are as given in the following algorithm. When invoked, the steps must be followed in the order given, aborting at the first step that returns a value. This algorithm will either return zero, a positive integer, or an error. Leading spaces are ignored. Trailing spaces and indeed any trailing garbage characters are ignored.

1. Let *input* be the string being parsed.
2. Let *position* be a pointer into *input*, initially pointing at the start of the string.
3. Let *value* have the value 0.
4. Skip whitespace (page 32).
5. If *position* is past the end of *input*, return an error.
6. If the next character is not one of U+0030 DIGIT ZERO (0) .. U+0039 DIGIT NINE (9), then return an error.
7. If the next character is one of U+0030 DIGIT ZERO (0) .. U+0039 DIGIT NINE (9):
 1. Multiply *value* by ten.
 2. Add the value of the current character (0..9) to *value*.
 3. Advance *position* to the next character.
 4. If *position* is not past the end of *input*, return to the top of step 7 in the overall algorithm (that's the step within which these substeps find themselves).
8. Return *value*.

2.4.3.2 Signed integers

A string is a **valid integer** if it consists of one or more characters in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9), optionally prefixed with a U+002D HYPHEN-MINUS ("−") character.

The **rules for parsing integers** are similar to the rules for non-negative integers, and are as given in the following algorithm. When invoked, the steps must be followed in the order given, aborting at the first step that returns a value. This algorithm will either return an integer or an error. Leading spaces are ignored. Trailing spaces and trailing garbage characters are ignored.

1. Let *input* be the string being parsed.
2. Let *position* be a pointer into *input*, initially pointing at the start of the string.
3. Let *value* have the value 0.

4. Let *sign* have the value "positive".
5. Skip whitespace (page 32).
6. If *position* is past the end of *input*, return an error.
7. If the character indicated by *position* (the first character) is a U+002D HYPHEN-MINUS ("−") character:
 1. Let *sign* be "negative".
 2. Advance *position* to the next character.
 3. If *position* is past the end of *input*, return an error.
8. If the next character is not one of U+0030 DIGIT ZERO (0) .. U+0039 DIGIT NINE (9), then return an error.
9. If the next character is one of U+0030 DIGIT ZERO (0) .. U+0039 DIGIT NINE (9):
 1. Multiply *value* by ten.
 2. Add the value of the current character (0..9) to *value*.
 3. Advance *position* to the next character.
 4. If *position* is not past the end of *input*, return to the top of step 9 in the overall algorithm (that's the step within which these substeps find themselves).
10. If *sign* is "positive", return *value*, otherwise return 0-*value*.

2.4.3.3 Real numbers

A string is a **valid floating point number** if it consists of:

1. Optionally, a U+002D HYPHEN-MINUS ("−") character.
2. A series of one or more characters in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9).
3. Optionally:
 1. A single U+002E FULL STOP (".") character.
 2. A series of one or more characters in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9).
4. Optionally:
 1. Either a U+0065 LATIN SMALL LETTER E character or a U+0045 LATIN CAPITAL LETTER E character.
 2. Optionally, a U+002D HYPHEN-MINUS ("−") character.
 3. A series of characters in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9).

The **rules for parsing floating point number values** are as given in the following algorithm. As with the previous algorithms, when this one is invoked, the steps must be followed in the order given, aborting at the first step that returns something. This algorithm will either return a number or an error. Leading spaces are ignored. Trailing spaces and garbage characters are ignored.

1. Let *input* be the string being parsed.

2. Let *position* be a pointer into *input*, initially pointing at the start of the string.
3. Let *value* have the value 1.
4. Let *divisor* have the value 1.
5. Let *exponent* have the value 1.
6. Skip whitespace (page 32).
7. If *position* is past the end of *input*, return an error.
8. If the character indicated by *position* is a U+002D HYPHEN-MINUS ("−") character:
 1. Change *value* and *divisor* to −1.
 2. Advance *position* to the next character.
 3. If *position* is past the end of *input*, return an error.
9. If the character indicated by *position* is not one of U+0030 DIGIT ZERO (0) .. U+0039 DIGIT NINE (9), then return an error.
10. Collect a sequence of characters (page 32) in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9), and interpret the resulting sequence as a base-ten integer. Multiply *value* by that integer.
11. If *position* is past the end of *input*, return *value*.
12. If the character indicated by *position* is a U+002E FULL STOP ("."), run these substeps:
 1. Advance *position* to the next character.
 2. If *position* is past the end of *input*, or if the character indicated by *position* is not one of U+0030 DIGIT ZERO (0) .. U+0039 DIGIT NINE (9), then return *value*.
 3. *Fraction loop*: Multiply *divisor* by ten.
 4. Add the value of the current character interpreted as a base-ten digit (0..9) divided by *divisor*, to *value*.
 5. Advance *position* to the next character.
 6. If *position* is past the end of *input*, then return *value*.
 7. If the character indicated by *position* is one of U+0030 DIGIT ZERO (0) .. U+0039 DIGIT NINE (9), return to the step labeled *fraction loop* in these substeps.
13. If the character indicated by *position* is a U+0065 LATIN SMALL LETTER E character or a U+0045 LATIN CAPITAL LETTER E character, run these substeps:
 1. Advance *position* to the next character.
 2. If *position* is past the end of *input*, then return *value*.

3. If the character indicated by *position* is a U+002D HYPHEN-MINUS ("–") character:
 1. Change *exponent* to –1.
 2. Advance *position* to the next character.
 3. If *position* is past the end of *input*, then return *value*.
 4. If the character indicated by *position* is not one of U+0030 DIGIT ZERO (0) .. U+0039 DIGIT NINE (9), then return *value*.
 5. Collect a sequence of characters (page 32) in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9), and interpret the resulting sequence as a base-ten integer. Multiply *exponent* by that integer.
 6. Multiply *value* by ten raised to the *exponent*th power.
14. Return *value*.

2.4.3.4 Ratios

Note: The algorithms described in this section are used by the progress and meter elements.

A **valid denominator punctuation character** is one of the characters from the table below. There is a **value associated with each denominator punctuation character**, as shown in the table below.

Denominator Punctuation Character		Value
U+0025 PERCENT SIGN	%	100
U+066A ARABIC PERCENT SIGN	%	100
U+FE6A SMALL PERCENT SIGN	?	100
U+FF05 FULLWIDTH PERCENT SIGN	?	100
U+2030 PER MILLE SIGN	‰	1000
U+2031 PER TEN THOUSAND SIGN	‰‰	10000

The **steps for finding one or two numbers of a ratio in a string** are as follows:

1. If the string is empty, then return nothing and abort these steps.
2. Find a number (page 37) in the string according to the algorithm below, starting at the start of the string.
3. If the sub-algorithm in step 2 returned nothing or returned an error condition, return nothing and abort these steps.
4. Set *number1* to the number returned by the sub-algorithm in step 2.
5. Starting with the character immediately after the last one examined by the sub-algorithm in step 2, skip all White_Space (page 32) characters in the string (this might match zero characters).

6. If there are still further characters in the string, and the next character in the string is a valid denominator punctuation character (page 36), set *denominator* to that character.
7. If the string contains any other characters in the range U+0030 DIGIT ZERO to U+0039 DIGIT NINE, but *denominator* was given a value in the step 6, return nothing and abort these steps.
8. Otherwise, if *denominator* was given a value in step 6, return *number1* and *denominator* and abort these steps.
9. Find a number (page 37) in the string again, starting immediately after the last character that was examined by the sub-algorithm in step 2.
10. If the sub-algorithm in step 9 returned nothing or an error condition, return *number1* and abort these steps.
11. Set *number2* to the number returned by the sub-algorithm in step 9.
12. Starting with the character immediately after the last one examined by the sub-algorithm in step 9, skip all White_Space (page 32) characters in the string (this might match zero characters).
13. If there are still further characters in the string, and the next character in the string is a valid denominator punctuation character (page 36), return nothing and abort these steps.
14. If the string contains any other characters in the range U+0030 DIGIT ZERO to U+0039 DIGIT NINE, return nothing and abort these steps.
15. Otherwise, return *number1* and *number2*.

The algorithm to **find a number** is as follows. It is given a string and a starting position, and returns either nothing, a number, or an error condition.

1. Starting at the given starting position, ignore all characters in the given string until the first character that is either a U+002E FULL STOP or one of the ten characters in the range U+0030 DIGIT ZERO to U+0039 DIGIT NINE.
2. If there are no such characters, return nothing and abort these steps.
3. Starting with the character matched in step 1, collect all the consecutive characters that are either a U+002E FULL STOP or one of the ten characters in the range U+0030 DIGIT ZERO to U+0039 DIGIT NINE, and assign this string of one or more characters to *string*.
4. If *string* consists of just a single U+002E FULL STOP character or if it contains more than one U+002E FULL STOP character then return an error condition and abort these steps.
5. Parse *string* according to the rules for parsing floating point number values (page 34), to obtain *number*. This step cannot fail (*string* is guaranteed to be a valid floating point number (page 34)).
6. Return *number*.

2.4.3.5 Percentages and dimensions

- ** **valid positive non-zero integers rules for parsing dimension values** (only used by height/width on img, embed, object — lengths in css pixels or percentages)

2.4.3.6 Lists of integers

A **valid list of integers** is a number of valid integers (page 33) separated by U+002C COMMA characters, with no other characters (e.g. no space characters (page 32)). In addition, there might be restrictions on the number of integers that can be given, or on the range of values allowed.

The **rules for parsing a list of integers** are as follows:

1. Let *input* be the string being parsed.
2. Let *position* be a pointer into *input*, initially pointing at the start of the string.
3. Let *numbers* be an initially empty list of integers. This list will be the result of this algorithm.
4. If there is a character in the string *input* at position *position*, and it is either a U+0020 SPACE, U+002C COMMA, or U+003B SEMICOLON character, then advance *position* to the next character in *input*, or to beyond the end of the string if there are no more characters.
5. If *position* points to beyond the end of *input*, return *numbers* and abort.
6. If the character in the string *input* at position *position* is a U+0020 SPACE, U+002C COMMA, or U+003B SEMICOLON character, then return to step 4.
7. Let *negated* be false.
8. Let *value* be 0.
9. Let *started* be false. This variable is set to true when the parser sees a number or a "-" character.
10. Let *got number* be false. This variable is set to true when the parser sees a number.
11. Let *finished* be false. This variable is set to true to switch parser into a mode where it ignores characters until the next separator.
12. Let *bogus* be false.
13. *Parser:* If the character in the string *input* at position *position* is:
 - ↪ **A U+002D HYPHEN-MINUS character**
Follow these substeps:
 1. If *got number* is true, let *finished* be true.
 2. If *finished* is true, skip to the next step in the overall set of steps.
 3. If *started* is true, let *negated* be false.

4. Otherwise, if *started* is false and if *bogus* is false, let *negated* be true.
5. Let *started* be true.

↪ **A character in the range U+0030 DIGIT ZERO .. U+0039 DIGIT NINE**

Follow these substeps:

1. If *finished* is true, skip to the next step in the overall set of steps.
2. Multiply *value* by ten.
3. Add the value of the digit, interpreted in base ten, to *value*.
4. Let *started* be true.
5. Let *got number* be true.

↪ **A U+0020 SPACE character**

↪ **A U+002C COMMA character**

↪ **A U+003B SEMICOLON character**

Follow these substeps:

1. If *got number* is false, return the *numbers* list and abort. This happens if an entry in the list has no digits, as in "1,2,x,4".
2. If *negated* is true, then negate *value*.
3. Append *value* to the *numbers* list.
4. Jump to step 4 in the overall set of steps.

↪ **A U+002E FULL STOP character**

Follow these substeps:

1. If *got number* is true, let *finished* be true.
2. If *finished* is true, skip to the next step in the overall set of steps.
3. Let *negated* be false.

↪ **Any other character**

Follow these substeps:

1. If *finished* is true, skip to the next step in the overall set of steps.
2. Let *negated* be false.
3. Let *bogus* be true.
4. If *started* is true, then return the *numbers* list, and abort. (The value in *value* is not appended to the list first; it is dropped.)

14. Advance *position* to the next character in *input*, or to beyond the end of the string if there are no more characters.

15. If *position* points to a character (and not to beyond the end of *input*), jump to the big *Parser* step above.
16. If *negated* is true, then negate *value*.
17. If *got number* is true, then append *value* to the *numbers* list.
18. Return the *numbers* list and abort.

2.4.4 Dates and times

In the algorithms below, the **number of days in month month of year year** is: 31 if *month* is 1, 3, 5, 7, 8, 10, or 12; 30 if *month* is 4, 6, 9, or 11; 29 if *month* is 2 and *year* is a number divisible by 400, or if *year* is a number divisible by 4 but not by 100; and 28 otherwise. This takes into account leap years in the Gregorian calendar. [GREGORIAN]

2.4.4.1 Specific moments in time

- ** This syntax is going to be tightened up and made almost exactly the same as the valid UTC date and time string (page 47) syntax, with the exception of allowing time zones. In fact what we might do is allow time zones in general, use the same parser, etc, but require that UAs always use the UTC time zone when synthesizing datetimes for form submission.

A string is a **valid datetime** if it has four digits (representing the year), a literal hyphen, two digits (representing the month), a literal hyphen, two digits (representing the day), optionally some spaces, either a literal T or a space, optionally some more spaces, two digits (for the hour), a colon, two digits (the minutes), optionally the seconds (which, if included, must consist of another colon, two digits (the integer part of the seconds), and optionally a decimal point followed by one or more digits (for the fractional part of the seconds)), optionally some spaces, and finally either a literal Z (indicating the time zone is UTC), or, a plus sign or a minus sign followed by two digits, a colon, and two digits (for the sign, the hours and minutes of the timezone offset respectively); with the month-day combination being a valid date in the given year according to the Gregorian calendar, the hour values (*h*) being in the range $0 \leq h \leq 23$, the minute values (*m*) in the range $0 \leq m \leq 59$, and the second value (*s*) being in the range $0 \leq s < 60$. [GREGORIAN]

The digits must be characters in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9), the hyphens must be a U+002D HYPHEN-MINUS characters, the T must be a U+0054 LATIN CAPITAL LETTER T, the colons must be U+003A COLON characters, the decimal point must be a U+002E FULL STOP, the Z must be a U+005A LATIN CAPITAL LETTER Z, the plus sign must be a U+002B PLUS SIGN, and the minus U+002D (same as the hyphen).

The following are some examples of dates written as valid datetimes (page 40).

"0037-12-13 00:00 Z"

Midnight UTC on the birthday of Nero (the Roman Emperor).

"1979-10-14T12:00:00.001-04:00"

One millisecond after noon on October 14th 1979, in the time zone in use on the east coast of North America during daylight saving time.

"8592-01-01 T 02:09 +02:09"

Midnight UTC on the 1st of January, 8592. The time zone associated with that time is two hours and nine minutes ahead of UTC.

Several things are notable about these dates:

- Years with fewer than four digits have to be zero-padded. The date "37-12-13" would not be a valid date.
- To unambiguously identify a moment in time prior to the introduction of the Gregorian calendar, the date has to be first converted to the Gregorian calendar from the calendar in use at the time (e.g. from the Julian calendar). The date of Nero's birth is the 15th of December 37, in the Julian Calendar, which is the 13th of December 37 in the Gregorian Calendar.
- The time and timezone components are not optional.
- Dates before the year 0 or after the year 9999 can't be represented as a datetime in this version of HTML.
- Time zones differ based on daylight savings time.

Note: Conformance checkers can use the algorithm below to determine if a datetime is a valid datetime or not.

To **parse a string as a datetime value**, a user agent must apply the following algorithm to the string. This will either return a time in UTC, with associated timezone information for round tripping or display purposes, or nothing, indicating the value is not a valid datetime (page 40). If at any point the algorithm says that it "fails", this means that it returns nothing.

1. Let *input* be the string being parsed.
2. Let *position* be a pointer into *input*, initially pointing at the start of the string.
3. Collect a sequence of characters (page 32) in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9). If the collected sequence is not exactly four characters long, then fail. Otherwise, interpret the resulting sequence as a base-ten integer. Let that number be the *year*.
4. If *position* is beyond the end of *input* or if the character at *position* is not a U+002D HYPHEN-MINUS character, then fail. Otherwise, move *position* forwards one character.
5. Collect a sequence of characters (page 32) in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9). If the collected sequence is not exactly two characters long, then fail. Otherwise, interpret the resulting sequence as a base-ten integer. Let that number be the *month*.
6. If *month* is not a number in the range $1 \leq month \leq 12$, then fail.
7. Let *maxday* be the number of days in month *month* of year *year* (page 40).

8. If *position* is beyond the end of *input* or if the character at *position* is not a U+002D HYPHEN-MINUS character, then fail. Otherwise, move *position* forwards one character.
9. Collect a sequence of characters (page 32) in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9). If the collected sequence is not exactly two characters long, then fail. Otherwise, interpret the resulting sequence as a base-ten integer. Let that number be the *day*.
10. If *day* is not a number in the range $1 \leq \text{month} \leq \text{maxday}$, then fail.
11. Collect a sequence of characters (page 32) that are either U+0054 LATIN CAPITAL LETTER T characters or space characters (page 32). If the collected sequence is zero characters long, or if it contains more than one U+0054 LATIN CAPITAL LETTER T character, then fail.
12. Collect a sequence of characters (page 32) in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9). If the collected sequence is not exactly two characters long, then fail. Otherwise, interpret the resulting sequence as a base-ten integer. Let that number be the *hour*.
13. If *hour* is not a number in the range $0 \leq \text{hour} \leq 23$, then fail.
14. If *position* is beyond the end of *input* or if the character at *position* is not a U+003A COLON character, then fail. Otherwise, move *position* forwards one character.
15. Collect a sequence of characters (page 32) in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9). If the collected sequence is not exactly two characters long, then fail. Otherwise, interpret the resulting sequence as a base-ten integer. Let that number be the *minute*.
16. If *minute* is not a number in the range $0 \leq \text{minute} \leq 59$, then fail.
17. Let *second* be a string with the value "0".
18. If *position* is beyond the end of *input*, then fail.
19. If the character at *position* is a U+003A COLON, then:
 1. Advance *position* to the next character in *input*.
 2. If *position* is beyond the end of *input*, or at the last character in *input*, or if the next two characters in *input* starting at *position* are not two characters both in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9), then fail.
 3. Collect a sequence of characters (page 32) that are either characters in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9) or U+002E FULL STOP characters. If the collected sequence has more than one U+002E FULL STOP characters, or if the last character in the sequence is a U+002E FULL STOP character, then fail. Otherwise, let the collected string be *second* instead of its previous value.
20. Interpret *second* as a base-ten number (possibly with a fractional part). Let that number be *second* instead of the string version.

21. If *second* is not a number in the range $0 \leq second < 60$, then fail. (The values 60 and 61 are not allowed: leap seconds cannot be represented by datetime values.)
22. If *position* is beyond the end of *input*, then fail.
23. Skip whitespace (page 32).
24. If the character at *position* is a U+005A LATIN CAPITAL LETTER Z, then:
 1. Let *timezone_{hours}* be 0.
 2. Let *timezone_{minutes}* be 0.
 3. Advance *position* to the next character in *input*.
25. Otherwise, if the character at *position* is either a U+002B PLUS SIGN ("+") or a U+002D HYPHEN-MINUS ("-"), then:
 1. If the character at *position* is a U+002B PLUS SIGN ("+"), let *sign* be "positive". Otherwise, it's a U+002D HYPHEN-MINUS ("-"); let *sign* be "negative".
 2. Advance *position* to the next character in *input*.
 3. Collect a sequence of characters (page 32) in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9). If the collected sequence is not exactly two characters long, then fail. Otherwise, interpret the resulting sequence as a base-ten integer. Let that number be the *timezone_{hours}*.
 4. If *timezone_{hours}* is not a number in the range $0 \leq timezone_{hours} \leq 23$, then fail.
 5. If *sign* is "negative", then negate *timezone_{hours}*.
 6. If *position* is beyond the end of *input* or if the character at *position* is not a U+003A COLON character, then fail. Otherwise, move *position* forwards one character.
 7. Collect a sequence of characters (page 32) in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9). If the collected sequence is not exactly two characters long, then fail. Otherwise, interpret the resulting sequence as a base-ten integer. Let that number be the *timezone_{minutes}*.
 8. If *timezone_{minutes}* is not a number in the range $0 \leq timezone_{minutes} \leq 59$, then fail.
 9. If *sign* is "negative", then negate *timezone_{minutes}*.
26. If *position* is not beyond the end of *input*, then fail.
27. Let *time* be the moment in time at year *year*, month *month*, day *day*, hours *hour*, minute *minute*, second *second*, subtracting *timezone_{hours}* hours and *timezone_{minutes}* minutes. That moment in time is a moment in the UTC timezone.
28. Let *timezone* be *timezone_{hours}* hours and *timezone_{minutes}* minutes from UTC.

29. Return *time* and *timezone*.

2.4.4.2 Vaguer moments in time

This section defines **date or time strings**. There are two kinds, **date or time strings in content**, and **date or time strings in attributes**. The only difference is in the handling of whitespace characters.

To parse a date or time string (page 44), user agents must use the following algorithm. A date or time string (page 44) is a *valid* date or time string if the following algorithm, when run on the string, doesn't say the string is invalid.

The algorithm may return nothing (in which case the string will be invalid), or it may return a date, a time, a date and a time, or a date and a time and a timezone. Even if the algorithm returns one or more values, the string can still be invalid.

1. Let *input* be the string being parsed.
2. Let *position* be a pointer into *input*, initially pointing at the start of the string.
3. Let *results* be the collection of results that are to be returned (one or more of a date, a time, and a timezone), initially empty. If the algorithm aborts at any point, then whatever is currently in *results* must be returned as the result of the algorithm.
4. For the "in content" variant: skip White_Space characters (page 32); for the "in attributes" variant: skip whitespace (page 32).
5. Collect a sequence of characters (page 32) in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9). If the collected sequence is empty, then the string is invalid; abort these steps.
6. Let the sequence of characters collected in the last step be *s*.
7. If *position* is past the end of *input*, the string is invalid; abort these steps.
8. If the character at *position* is not a U+003A COLON character, then:
 1. If the character at *position* is not a U+002D HYPHEN-MINUS ("−") character either, then the string is invalid, abort these steps.
 2. If the sequence *s* is not exactly four digits long, then the string is invalid. (This does not stop the algorithm, however.)
 3. Interpret the sequence of characters collected in step 5 as a base-ten integer, and let that number be *year*.
 4. Advance *position* past the U+002D HYPHEN-MINUS ("−") character.
 5. Collect a sequence of characters (page 32) in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9). If the collected sequence is empty, then the string is invalid; abort these steps.
 6. If the sequence collected in the last step is not exactly two digits long, then the string is invalid.

7. Interpret the sequence of characters collected two steps ago as a base-ten integer, and let that number be *month*.
8. If *month* is not a number in the range $1 \leq month \leq 12$, then the string is invalid, abort these steps.
9. Let *maxday* be the number of days in month *month* of year *year* (page 40).
10. If *position* is past the end of *input*, or if the character at *position* is not a U+002D HYPHEN-MINUS ("−") character, then the string is invalid, abort these steps. Otherwise, advance *position* to the next character.
11. Collect a sequence of characters (page 32) in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9). If the collected sequence is empty, then the string is invalid; abort these steps.
12. If the sequence collected in the last step is not exactly two digits long, then the string is invalid.
13. Interpret the sequence of characters collected two steps ago as a base-ten integer, and let that number be *day*.
14. If *day* is not a number in the range $1 \leq day \leq maxday$, then the string is invalid, abort these steps.
15. Add the date represented by *year*, *month*, and *day* to the *results*.
16. For the "in content" variant: skip White_Space characters (page 32); for the "in attributes" variant: skip whitespace (page 32).
17. If the character at *position* is a U+0054 LATIN CAPITAL LETTER T, then move *position* forwards one character.
18. For the "in content" variant: skip White_Space characters (page 32); for the "in attributes" variant: skip whitespace (page 32).
19. Collect a sequence of characters (page 32) in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9). If the collected sequence is empty, then the string is invalid; abort these steps.
20. Let *s* be the sequence of characters collected in the last step.
9. If *s* is not exactly two digits long, then the string is invalid.
10. Interpret the sequence of characters collected two steps ago as a base-ten integer, and let that number be *hour*.
11. If *hour* is not a number in the range $0 \leq hour \leq 23$, then the string is invalid, abort these steps.
12. If *position* is past the end of *input*, or if the character at *position* is not a U+003A COLON character, then the string is invalid, abort these steps. Otherwise, advance *position* to the next character.

13. Collect a sequence of characters (page 32) in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9). If the collected sequence is empty, then the string is invalid; abort these steps.
14. If the sequence collected in the last step is not exactly two digits long, then the string is invalid.
15. Interpret the sequence of characters collected two steps ago as a base-ten integer, and let that number be *minute*.
16. If *minute* is not a number in the range $0 \leq \text{minute} \leq 59$, then the string is invalid, abort these steps.
17. Let *second* be 0. It might be changed to another value in the next step.
18. If *position* is not past the end of *input* and the character at *position* is a U+003A COLON character, then:
 1. Collect a sequence of characters (page 32) that are either characters in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9) or are U+002E FULL STOP. If the collected sequence is empty, or contains more than one U+002E FULL STOP character, then the string is invalid; abort these steps.
 2. If the first character in the sequence collected in the last step is not in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9), then the string is invalid.
 3. Interpret the sequence of characters collected two steps ago as a base-ten number (possibly with a fractional part), and let that number be *second*.
 4. If *second* is not a number in the range $0 \leq \text{second} < 60$, then the string is invalid, abort these steps.
19. Add the time represented by *hour*, *minute*, and *second* to the *results*.
20. If *results* has both a date and a time, then:
 1. For the "in content" variant: skip White_Space characters (page 32); for the "in attributes" variant: skip whitespace (page 32).
 2. If *position* is past the end of *input*, then skip to the next step in the overall set of steps.
 3. Otherwise, if the character at *position* is a U+005A LATIN CAPITAL LETTER Z, then:
 1. Add the timezone corresponding to UTC (zero offset) to the *results*.
 2. Advance *position* to the next character in *input*.
 3. Skip to the next step in the overall set of steps.
 4. Otherwise, if the character at *position* is either a U+002B PLUS SIGN ("+") or a U+002D HYPHEN-MINUS ("-"), then:

1. If the character at *position* is a U+002B PLUS SIGN ("+"), let *sign* be "positive". Otherwise, it's a U+002D HYPHEN-MINUS ("-"); let *sign* be "negative".
2. Advance *position* to the next character in *input*.
3. Collect a sequence of characters (page 32) in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9). If the collected sequence is not exactly two characters long, then the string is invalid.
4. Interpret the sequence collected in the last step as a base-ten number, and let that number be *timezone_{hours}*.
5. If *timezone_{hours}* is not a number in the range $0 \leq \text{timezone}_{\text{hours}} \leq 23$, then the string is invalid; abort these steps.
6. If *sign* is "negative", then negate *timezone_{hours}*.
7. If *position* is beyond the end of *input* or if the character at *position* is not a U+003A COLON character, then the string is invalid; abort these steps. Otherwise, move *position* forwards one character.
8. Collect a sequence of characters (page 32) in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9). If the collected sequence is not exactly two characters long, then the string is invalid.
9. Interpret the sequence collected in the last step as a base-ten number, and let that number be *timezone_{minutes}*.
10. If *timezone_{minutes}* is not a number in the range $0 \leq \text{timezone}_{\text{minutes}} \leq 59$, then the string is invalid; abort these steps.
11. Add the timezone corresponding to an offset of *timezone_{hours}* hours and *timezone_{minutes}* minutes to the *results*.
12. Skip to the next step in the overall set of steps.
5. Otherwise, the string is invalid; abort these steps.
21. For the "in content" variant: skip White_Space characters (page 32); for the "in attributes" variant: skip whitespace (page 32).
22. If *position* is not past the end of *input*, then the string is invalid.
23. Abort these steps (the string is parsed).

2.4.4.3 UTC dates and times

A **UTC date and time** consists of a specific Gregorian date expressed relative to the UTC timezone, consisting of a year, a month, a day, an hour, a minute, a second, and a fraction of a second. [GREGORIAN]

** ... valid UTC date and time string

** ... rules to **parse a UTC date and time string**

2.4.4.4 Local dates and times

A **local date and time** consists of a specific Gregorian date with no timezone information, consisting of a year, a month, a day, an hour, a minute, a second, and a fraction of a second. [GREGORIAN]

** ... **valid local date and time string**

** ... rules to **parse a local date and time string**

2.4.4.5 Dates

A **date** consists of a specific Gregorian date with no timezone information, consisting of a year, a month, and a day. [GREGORIAN]

** ... **valid date string**

** ... rules to **parse a date string**

2.4.4.6 Months

A **month** consists of a specific Gregorian date with no timezone information and no date information beyond a year and a month. [GREGORIAN]

** ... **valid month string**

** ... rules to **parse a month string**

2.4.4.7 Weeks

A **week** consists of a specific Gregorian date with no timezone information and no date information beyond a year and a week. [GREGORIAN]

** ... **valid week string**

** ... rules to **parse a week string**

2.4.4.8 Times

A **time** consists of a specific time with no timezone information, consisting of an hour, a minute, a second, and a fraction of a second.

** ... **valid time string**

** ... **rules to parse a time string**

2.4.4.9 Time offsets

** **valid time offset, rules for parsing time offsets, time offset serialization rules;**
probably in the format "5d4h3m2.1s" or similar, with all components being optional, and
** the last component's unit suffix being optional if it's in seconds.

2.4.5 Space-separated tokens

A **set of space-separated tokens** is a set of zero or more words separated by one or more space characters (page 32), where words consist of any string of one or more characters, none of which are space characters (page 32).

A string containing a set of space-separated tokens (page 49) may have leading or trailing space characters (page 32).

An **unordered set of unique space-separated tokens** is a set of space-separated tokens (page 49) where none of the words are duplicated.

An **ordered set of unique space-separated tokens** is a set of space-separated tokens (page 49) where none of the words are duplicated but where the order of the tokens is meaningful.

Sets of space-separated tokens (page 49) sometimes have a defined set of allowed values. When a set of allowed values is defined, the tokens must all be from that list of allowed values; other values are non-conforming. If no such set of allowed values is provided, then all values are conforming.

When a user agent has to **split a string on spaces**, it must use the following algorithm:

1. Let *input* be the string being parsed.
2. Let *position* be a pointer into *input*, initially pointing at the start of the string.
3. Let *tokens* be a list of tokens, initially empty.
4. Skip whitespace (page 32)
5. While *position* is not past the end of *input*:
 1. Collect a sequence of characters (page 32) that are not space characters (page 32).

2. Add the string collected in the previous step to *tokens*.
3. Skip whitespace (page 32)
6. Return *tokens*.

When a user agent has to **remove a token from a string**, it must use the following algorithm:

1. Let *input* be the string being modified.
2. Let *token* be the token being removed. It will not contain any space characters (page 32).
3. Let *output* be the output string, initially empty.
4. Let *position* be a pointer into *input*, initially pointing at the start of the string.
5. If *position* is beyond the end of *input*, set the string being modified to *output*, and abort these steps.
6. If the character at *position* is a space character (page 32):
 1. Append the character at *position* to the end of *output*.
 2. Increment *position* so it points at the next character in *input*.
 3. Return to step 5 in the overall set of steps.
7. Otherwise, the character at *position* is the first character of a token. Collect a sequence of characters (page 32) that are not space characters (page 32), and let that be *s*.
8. If *s* is exactly equal to *token*, then:
 1. Skip whitespace (page 32) (in *input*).
 2. Remove any space characters (page 32) currently at the end of *output*.
 3. If *position* is not past the end of *input*, and *output* is not the empty string, append a single U+0020 SPACE character at the end of *output*.
9. Otherwise, append *s* to the end of *output*.
10. Return to step 6 in the overall set of steps.

Note: This causes any occurrences of the token to be removed from the string, and any spaces that were surrounding the token to be collapsed to a single space, except at the start and end of the string, where such spaces are removed.

2.4.6 Comma-separated tokens

** We should allow whitespace around commas, and leading/trailing whitespace.

A **set of comma-separated tokens** is a set of zero or more tokens each separated from the next by a single U+002C COMMA character (,), where tokens consist of any string of zero or more characters, none of which are U+002C COMMA characters (,).

Sets of comma-separated tokens (page 51) sometimes have further restrictions on what consists a valid token. When such restrictions are defined, the tokens must all fit within those restrictions; other values are non-conforming. If no such restrictions are specified, then all values are conforming.

When a user agent has to **split a string on commas**, it must use the following algorithm:

1. Let *input* be the string being parsed.
2. Let *position* be a pointer into *input*, initially pointing at the start of the string.
3. Let *tokens* be a list of tokens, initially empty.
4. *Token*: If *position* is past the end of *input*, jump to the last step.
5. Collect a sequence of characters (page 32) that are not U+002C COMMA characters (,).
6. Add the string collected in the previous step (which might be the empty string) to *tokens*.
7. If *position* is not past the end of *input*, then the character at *position* is a U+002C COMMA character (,); advance *position* past that character.
8. Jump back to the step labeled *token*.
9. Return *tokens*.

2.4.7 Keywords and enumerated attributes

Some attributes are defined as taking one of a finite set of keywords. Such attributes are called **enumerated attributes**. The keywords are each defined to map to a particular *state* (several keywords might map to the same state, in which case some of the keywords are synonyms of each other; additionally, some of the keywords can be said to be non-conforming, and are only in the specification for historical reasons). In addition, two default states can be given. The first is the *invalid value default*, the second is the *missing value default*.

If an enumerated attribute is specified, the attribute's value must be an ASCII case-insensitive (page 31) match for one of the given keywords that are not said to be non-conforming, with no leading or trailing whitespace.

When the attribute is specified, if its value is an ASCII case-insensitively match for one of the given keywords then that keyword's state is the state that the attribute represents. If the attribute value matches none of the given keywords, but the attribute has an *invalid value default*, then the attribute represents that state. Otherwise, if the attribute value matches none of the keywords but there is a *missing value default* state defined, then *that* is the state represented by the attribute. Otherwise, there is no default, and invalid values must be ignored.

When the attribute is *not* specified, if there is a *missing value default* state defined, then that is the state represented by the (missing) attribute. Otherwise, the absence of the attribute means that there is no state represented.

Note: *The empty string can be one of the keywords in some cases. For example the contenteditable attribute has two states: true, matching the true keyword and the empty string, false, matching false and all other keywords (it's the invalid value default). It could further be thought of as having a third state inherit, which would be the default when the attribute is not specified at all (the missing value default), but for various reasons that isn't the way this specification actually defines it.*

2.4.8 References

A **valid hash-name reference** to an element of type *type* is a string consisting of a U+0023 NUMBER SIGN (#) character followed by a string which exactly matches the value of the name attribute of an element in the document with type *type*.

The **rules for parsing a hash-name reference** to an element of type *type* are as follows:

1. If the string being parsed does not contain a U+0023 NUMBER SIGN character, or if the first such character in the string is the last character in the string, then return null and abort these steps.
2. Let *s* be the string from the character immediately after the first U+0023 NUMBER SIGN character in the string being parsed up to the end of that string.
3. Return the first element of type *type* that has an id or name attribute whose value is a compatibility caseless (page 31) match for *s*.

2.5 URLs

This specification defines the term URL (page 52), and defines various algorithms for dealing with URLs, because for historical reasons the rules defined by the URI and IRI specifications are not a complete description of what HTML user agents need to implement to be compatible with Web content.

2.5.1 Terminology

A **URL** is a string used to identify a resource. A URL (page 52) is always associated with a Document, either explicitly when the URL is created or defined; or through a DOM node, in which case the associated Document is the node's Document; or through a script, in which case the associated Document is the script's script document context (page 428).

A URL (page 52) is a **valid URL** if at least one of the following conditions holds:

- The URL (page 52) is a valid URI reference [RFC3986].
- The URL (page 52) is a valid IRI reference and it has no query component. [RFC3987]

- The URL (page 52) is a valid IRI reference and its query component contains no unescaped non-ASCII characters. [RFC3987]
- The URL (page 52) is a valid IRI reference and the character encoding (page 79) of the URL's Document is UTF-8 or UTF-16. [RFC3987]

Note: The term "URL" in this specification is used in a manner distinct from the precise technical meaning it is given in RFC 3986. Readers familiar with that RFC will find it easier to read this specification if they pretend the term "URL" as used herein is really called something else altogether.

2.5.2 Parsing URLs

To parse a URL *url* into its component parts, the user agent must use the following steps:

1. Strip leading and trailing space characters (page 32) from *url*.
2. Parse *url* in the manner defined by RFC 3986, with the following exceptions:
 - Add all characters with codepoints less than or equal to U+0020 or greater than or equal to U+007F to the <unreserved> production.
 - Add the characters U+0022, U+003C, U+003E, U+005B .. U+005E, U+0060, and U+007B .. U+007D to the <unreserved> production.
 - Add a single U+0025 PERCENT SIGN character as a second alternative way of matching the <pct-encoded> production, except when the <pct-encoded> is used in the <reg-name> production.
 - Add the U+0023 NUMBER SIGN character to the characters allowed in the <fragment> production.
3. If *url* doesn't match the <URI-reference> production, even after the above changes are made to the ABNF definitions, then parsing the URL fails with an error. [RFC3986]

Otherwise, parsing *url* was successful; the components of the URL are substrings of *url* defined as follows:

<scheme>

The substring matched by the <scheme> production, if any.

<host>

The substring matched by the <host> production, if any.

<port>

The substring matched by the <port> production, if any.

<hostport>

If there is a <scheme> component and a <port> component and the port given by the <port> component is different than the default port defined for the protocol given by the <scheme> component, then <hostport> is the substring that starts with the substring matched by the <host> production and ends with

the substring matched by the `<port>` production, and includes the colon in between the two. Otherwise, it is the same as the `<host>` component.

`<path>`

The substring matched by one of the following productions, if one of them was matched:

- `<path-abempty>`
- `<path-absolute>`
- `<path-noscheme>`
- `<path-rootless>`
- `<path-empty>`

`<query>`

The substring matched by the `<query>` production, if any.

`<fragment>`

The substring matched by the `<fragment>` production, if any.

2.5.3 Resolving URLs

Relative URLs are resolved relative to a base URL. The **base URL** of a URL (page 52) is the absolute URL (page 56) obtained as follows:

↪ If the URL to be resolved was passed to an API

The base URL is the document base URL (page 54) of the script's script document context (page 428).

↪ If the URL to be resolved is from the value of a content attribute

The base URL is the *base URI of the element* that the attribute is on, as defined by the XML Base specification, with the *base URI of the document entity* being defined as the document base URL (page 54) of the Document that owns the element.

For the purposes of the XML Base specification, user agents must act as if all Document objects represented XML documents.

Note: It is possible for `xml:base` attributes to be present even in HTML fragments, as such attributes can be added dynamically using script. (Such scripts would not be conforming, however, as `xml:base` attributes are not allowed in HTML documents (page 76).)

↪ If the URL to be resolved was found in an offline application cache manifest

The base URL is the URL of the application cache (page 448) manifest (page 449).

The **document base URL** of a Document is the absolute URL (page 56) obtained by running these steps:

1. If there is no base element that is both a child of the head element (page 80) and has an href attribute, then the document base URL (page 54) is the document's address.
2. Otherwise, let *url* be the value of the href attribute of the first such element.

3. Resolve (page 55) the *url* URL, using the document's address as the base URL (page 54) (thus, the base href attribute isn't affected by xml:base attributes).
4. The document base URL (page 54) is the result of the previous step if it was successful; otherwise it is the document's address.

To **resolve a URL** to an absolute URL (page 56) the user agent must use the following steps. Resolving a URL can result in an error, in which case the URL is not resolvable.

1. Let *url* be the URL (page 52) being resolved.
2. Let *document* be the Document associated with (page 52) *url*.
3. Let *encoding* be the character encoding (page 79) of *document*.
4. If *encoding* is UTF-16, then change it to UTF-8.
5. Let *base* be the base URL (page 54) for *url*. (This is an absolute URL (page 56).)
6. Parse (page 53) *url* into its component parts.
7. If parsing *url* resulted in a <host> (page 53) component, then replace the matching substring of *url* with the string that results from expanding any sequences of percent-encoded octets in that component that are valid UTF-8 sequences into Unicode characters as defined by UTF-8.

If any percent-encoded octets in that component are not valid UTF-8 sequences, then return an error and abort these steps.

Apply the IDNA ToASCII algorithm to the matching substring, with both the AllowUnassigned and UseSTD3ASCIIRules flags set. Replace the matching substring with the result of the ToASCII algorithm.

If ToASCII fails to convert one of the components of the string, e.g. because it is too long or because it contains invalid characters, then return an error and abort these steps. [RFC3490]

8. If parsing *url* resulted in a <path> (page 54) component, then replace the matching substring of *url* with the string that results from applying the following steps to each character other than U+0025 PERCENT SIGN (%) that doesn't match the original <path> production defined in RFC 3986:
 1. Encode the character into a sequence of octets as defined by UTF-8.
 2. Replace the character with the percent-encoded form of those octets. [RFC3986]

For instance if *url* was "://example.com/a^b@c%FFd%z/?e", then the <path> (page 54) component's substring would be "/a^b@c%FFd%z/" and the two characters that would have to be escaped would be "^" and "@". The result after this step was applied would therefore be that *url* now had the value "://example.com/a%5Eb%E2%98%BAc%FFd%z/?e".

9. If parsing *url* resulted in a <query> (page 54) component, then replace the matching substring of *url* with the string that results from applying the following steps to each

character other than U+0025 PERCENT SIGN (%) that doesn't match the original `<query>` production defined in RFC 3986:

1. If the character in question cannot be expressed in the encoding *encoding*, then replace it with a single 0x3F octet (an ASCII question mark) and skip the remaining substeps for this character.
 2. Encode the character into a sequence of octets as defined by the encoding *encoding*.
 3. Replace the character with the percent-encoded form of those octets. [RFC3986]
10. Apply the algorithm described in RFC 3986 section 5.2 Relative Resolution, using *url* as the potentially relative URI reference (*R*), and *base* as the base URI (*Base*). [RFC3986]
11. Apply any relevant conformance criteria of RFC 3986 and RFC 3987, returning an error and aborting these steps if appropriate. [RFC3986] [RFC3987]
- For instance, if an absolute URI that would be returned by the above algorithm violates the restrictions specific to its scheme, e.g. a data: URI using the "://" server-based naming authority syntax, then user agents are to treat this as an error instead.
12. Let *result* be the target URI (*T*) returned by the Relative Resolution algorithm.
 13. If *result* uses a scheme with a server-based naming authority, replace all U+005C REVERSE SOLIDUS (\) characters in *result* with U+002F SOLIDUS (/) characters.
 14. Return *result*.

A URL (page 52) is an **absolute URL** if resolving (page 55) it results in the same URL without an error.

2.5.4 Dynamic changes to base URLs

When an `xml:base` attribute changes, the attribute's element, and all descendant elements, are affected by a base URL change (page 56).

When a document's document base URL (page 54) changes, all elements in that document are affected by a base URL change (page 56).

When an element is moved from one document to another, if the two documents have different base URLs (page 54), then that element and all its descendants are affected by a base URL change (page 56).

When an element is **affected by a base URL change**, it must act as described in the following list:

↪ If the element is a hyperlink element (page 497)

If the absolute URL (page 56) identified by the hyperlink is being shown to the user, or if any data derived from that URL is affecting the display, then the `href` attribute should be reresolved (page 55) and the UI updated appropriately.

For example, the CSS `:link`/`:visited` pseudo-classes might have been affected.

If the hyperlink has a `ping` attribute and its absolute URL(s) (page 56) are being shown to the user, then the `ping` attribute's tokens should be reresolved (page 55) and the UI updated appropriately.

↪ **If the element is a `blockquote`, `q`, `ins`, or `del` element with a `cite` attribute**

If the absolute URL (page 56) identified by the `cite` attribute is being shown to the user, or if any data derived from that URL is affecting the display, then the UI should be reresolved (page 55) and the UI updated appropriately.

↪ **Otherwise**

The element is not directly affected.

Changing the base URL doesn't affect the image displayed by `img` elements, although subsequent accesses of the `src` DOM attribute from script will return a new absolute URL (page 56) that might no longer correspond to the image being shown.

2.5.5 Interfaces for URL manipulation

An interface that has a complement of **URL decomposition attributes** will have seven attributes with the following definitions:

```
attribute DOMString protocol;  
attribute DOMString host;  
attribute DOMString hostname;  
attribute DOMString port;  
attribute DOMString pathname;  
attribute DOMString search;  
attribute DOMString hash;
```

The attributes defined to be URL decomposition attributes must act as described for the attributes with the same corresponding names in this section.

In addition, an interface with a complement of URL decomposition attributes will define an **input**, which is a URL (page 52) that the attributes act on, and a **common setter action**, which is a set of steps invoked when any of the attributes' setters are invoked.

The seven URL decomposition attributes have similar requirements.

On getting, if the input (page 57) fulfills the condition given in the "getter condition" column corresponding to the attribute in the table below, the user agent must return the part of the input (page 57) URL given in the "component" column, with any prefixes specified in the "prefix" column appropriately added to the start of the string and any suffixes specified in the "suffix" column appropriately added to the end of the string. Otherwise, the attribute must return the empty string.

On setting, the new value must first be mutated as described by the "setter preprocessor" column, then mutated by %-escaping any characters in the new value that are not valid in the relevant component as given by the "component" column. Then, if the resulting new

value fulfills the condition given in the "setter condition" column, the user agent must make a new string *output* by replacing the component of the URL given by the "component" column in the input (page 57) URL with the new value; otherwise, the user agent must let *output* be equal to the input (page 57). Finally, the user agent must invoke the common setter action (page 57) with the value of *output*.

When replacing a component in the URL, if the component is part of an optional group in the URL syntax consisting of a character followed by the component, the component (including its prefix character) must be included even if the new value is the empty string.

Note: The previous paragraph applies in particular to the ":" before a <port> component, the "?" before a <query> component, and the "#" before a <fragment> component.

For the purposes of the above definitions, URLs must be parsed using the URL parsing rules (page 53) defined in this specification.

Attribute	Component	Getter Condition	Prefix	Suffix	Setter Preprocessor	Setter Condition
protocol	<scheme> (page 53)	—	—	U+003A COLON (":")	Remove all trailing U+003A COLON (":") characters	The new value is not the empty string
host	<hostport> (page 53)	input (page 57) is hierarchical and uses a server-based naming authority	—	—	—	—
hostname	<host> (page 53)	input (page 57) is hierarchical and uses a server-based naming authority	—	—	Remove all leading U+002F SOLIDUS ("/") characters	—
port	<port> (page 53)	input (page 57) is hierarchical, uses a server-based naming authority, and contained a <port> (page 53) component (possibly an empty one)	—	—	Remove any characters in the new value that are not in the range U+0030 DIGIT ZERO .. U+0039 DIGIT NINE. If the resulting string is empty, set it to a single U+0030 DIGIT ZERO character ('0').	—
pathname	<path> (page 54)	input (page 57) is hierarchical	—	—	If it has no leading U+002F SOLIDUS ("/") character, prepend a U+002F SOLIDUS ("/") character to the new value	—
search	<query> (page 54)	input (page 57) is hierarchical, and contained a <query> (page 54) component (possibly an empty one)	U+003F QUESTION MARK ("?")	—	Remove one leading U+003F QUESTION MARK ("?") character, if any	—
hash	<fragment> (page 54)	input (page 57) contained a <fragment> (page 54) component (possibly an empty one)	U+0023 NUMBER SIGN ("#")	—	Remove one leading U+0023 NUMBER SIGN ("#") character, if any	—

The table below demonstrates how the getter condition for search results in different results depending on the exact original syntax of the URL:

Input URL	search value	Explanation
http://example.com/	empty string	No <query> (page 54) component in input URL.
http://example.com/?	?	There is a <query> (page 54) component, but it is empty. The question mark in the resulting value is the prefix.
http://example.com/?test	?test	The <query> (page 54) component has the value "test".
http://example.com/?test#	?test	The (empty) <fragment> (page 54) component is not part of the <query> (page 54) component.

2.6 Fetching resources

When a user agent is to **fetch** a resource, the following steps must be run:

1. If the resource is identified by a URL (page 52), then immediately resolve that URL (page 55).
2. If the resulting absolute URL (page 56) is **about:blank**, then return the empty string and abort these steps.
3. Perform the remaining steps asynchronously.
4. If the resource identified by the resulting absolute URL (page 56) is already being downloaded for other reasons (e.g. another invocation of this algorithm), and the resource is to be obtained using a idempotent action (such as an HTTP GET or equivalent), and the user agent is configured such that it is to reuse the data from the existing download instead of initiating a new one, then use the results of the existing download instead of starting a new one.

Otherwise, at a time convenient to the user and the user agent, download the resource, applying the semantics of the relevant specifications (e.g. performing an HTTP GET or POST operation, or reading the file from disk, following redirects, dereferencing javascript: URLs (page 430), etc).

5. When the resource is available, queue a task (page 429) that uses the resource as appropriate. If the resource can be processed incrementally, as, for instance, with a progressively interlaced JPEG or an HTML file, multiple tasks may be queued to process the data as it is downloaded. The task source (page 429) for these tasks is the networking task source (page 430).

Note: The offline application cache processing model introduces some changes to the networking model (page 462) to handle the returning of cached resources.

Note: The navigation (page 473) processing model handles redirects itself, overriding the redirection handling that would be done by the fetching algorithm.

Note: Whether the type sniffing rules (page 60) apply to the fetched resource depends on the algorithm that invokes the rules — they are not always applicable.

2.7 Determining the type of a resource

⚠Warning! It is imperative that the rules in this section be followed exactly. When a user agent uses different heuristics for content type detection than the server expects, security problems can occur. For example, if a server believes that the client will treat a contributed file as an image (and thus treat it as benign), but a Web browser believes the content to be HTML (and thus execute any scripts contained therein), the end user can be exposed to malicious content, making the user vulnerable to cookie theft attacks and other cross-site scripting attacks.

2.7.1 Content-Type metadata

What explicit **Content-Type metadata** is associated with the resource (the resource's type information) depends on the protocol that was used to fetch (page 59) the resource.

For HTTP resources, only the first Content-Type HTTP header, if any, contributes any type information; the explicit type of the resource is then the value of that header, interpreted as described by the HTTP specifications. If the Content-Type HTTP header is present but the value of the first such header cannot be interpreted as described by the HTTP specifications (e.g. because its value doesn't contain a U+002F SOLIDUS ('/') character), then the resource has no type information (even if there are multiple Content-Type HTTP headers and one of the other ones is syntactically correct). [HTTP]

For resources fetched from the file system, user agents should use platform-specific conventions, e.g. operating system extension/type mappings.

Extensions must not be used for determining resource types for resources fetched over HTTP.

For resources fetched over most other protocols, e.g. FTP, there is no type information.

The **algorithm for extracting an encoding from a Content-Type**, given a string s , is as follows. It either returns an encoding or nothing.

1. Find the first seven characters in s that are an ASCII case-insensitive (page 31) match for the word "charset". If no such match is found, return nothing.
2. Skip any U+0009, U+000A, U+000C, U+000D, or U+0020 characters that immediately follow the word 'charset' (there might not be any).
3. If the next character is not a U+003D EQUALS SIGN ('='), return nothing.
4. Skip any U+0009, U+000A, U+000C, U+000D, or U+0020 characters that immediately follow the equals sign (there might not be any).
5. Process the next character as follows:

- ↪ If it is a U+0022 QUOTATION MARK ("") and there is a later U+0022 QUOTATION MARK ("") in *s*
- ↪ If it is a U+0027 APOSTROPHE (') and there is a later U+0027 APOSTROPHE (') in *s*

Return the string between this character and the next earliest occurrence of this character.

- ↪ If it is an unmatched U+0022 QUOTATION MARK ("")
- ↪ If it is an unmatched U+0027 APOSTROPHE (')
- ↪ If there is no next character

Return nothing.

- ↪ Otherwise

Return the string from this character to the first U+0009, U+000A, U+000C, U+000D, U+0020, or U+003B character or the end of *s*, whichever comes first.

Note: The above algorithm is a willful violation of the HTTP specification.
[RFC2616]

2.7.2 Content-Type sniffing: Web pages

The **sniffed type of a resource** must be found as follows:

1. If the user agent is configured to strictly obey Content-Type headers for this resource, then jump to the last step in this set of steps.
2. If the resource was fetched over an HTTP protocol and there is an HTTP Content-Type header and the value of the first such header has bytes that exactly match one of the following lines:

Bytes in Hexadecimal	Textual representation
74 65 78 74 2f 70 6c 61 69 6e	text/plain
74 65 78 74 2f 70 6c 61 69 6e 3b 20 63 68 61 72 73 65 74 3d 49 53 4f 2d 38 38 35 39 2d 31	text/ plain; charset=ISO-8859-1
74 65 78 74 2f 70 6c 61 69 6e 3b 20 63 68 61 72 73 65 74 3d 69 73 6f 2d 38 38 35 39 2d 31	text/ plain; charset=iso-8859-1
74 65 78 74 2f 70 6c 61 69 6e 3b 20 63 68 61 72 73 65 74 3d 55 54 46 2d 38	text/plain; charset=UTF-8

...then jump to the *text or binary* (page 62) section below.

3. Let *official type* be the type given by the Content-Type metadata (page 60) for the resource, ignoring parameters. If there is no such type, jump to the *unknown type* (page 62) step below. Comparisons with this type, as defined by MIME specifications, are done in an ASCII case-insensitive (page 31) manner. [RFC2046]
4. If *official type* is "unknown/unknown" or "application/unknown", jump to the *unknown type* (page 62) step below.

5. If *official type* ends in "+xml", or if it is either "text/xml" or "application/xml", then the sniffed type of the resource is *official type*; return that and abort these steps.
6. If *official type* is an image type supported by the user agent (e.g. "image/png", "image/gif", "image/jpeg", etc), then jump to the *images* (page 65) section below, passing it the *official type*.
7. If *official type* is "text/html", then jump to the *feed or HTML* (page 65) section below.
8. The sniffed type of the resource is *official type*.

2.7.3 Content-Type sniffing: text or binary

1. The user agent may wait for 512 or more bytes of the resource to be available.
2. Let n be the smaller of either 512 or the number of bytes already available.
3. If n is 4 or more, and the first bytes of the resource match one of the following byte sets:

Bytes in Hexadecimal	Description
FE FF	UTF-16BE BOM
FF FE	UTF-16LE BOM
EF BB BF	UTF-8 BOM

...then the sniffed type of the resource is "text/plain". Abort these steps.

4. If none of the first n bytes of the resource are binary data bytes (page 62) then the sniffed type of the resource is "text/plain". Abort these steps.
5. If the first bytes of the resource match one of the byte sequences in the "pattern" column of the table in the *unknown type* (page 62) section below, ignoring any rows whose cell in the "security" column says "scriptable" (or "n/a"), then the sniffed type of the resource is the type given in the corresponding cell in the "sniffed type" column on that row; abort these steps.

⚠Warning! It is critical that this step not ever return a scriptable type (e.g. text/html), as otherwise that would allow a privilege escalation attack.

6. Otherwise, the sniffed type of the resource is "application/octet-stream".

Bytes covered by the following ranges are **binary data bytes**:

- 0x00 - 0x08
- 0x0B
- 0x0E - 0x1A
- 0x1C - 0x1F

2.7.4 Content-Type sniffing: unknown type

1. The user agent may wait for 512 or more bytes of the resource to be available.
2. Let *stream length* be the smaller of either 512 or the number of bytes already available.

3. For each row in the table below:

↪ If the row has no "WS" bytes:

1. Let *pattern length* be the length of the pattern (number of bytes described by the cell in the second column of the row).
2. If *stream length* is smaller than *pattern length* then skip this row.
3. Apply the "and" operator to the first *pattern length* bytes of the resource and the given mask (the bytes in the cell of first column of that row), and let the result be the *data*.
4. If the bytes of the *data* matches the given pattern bytes exactly, then the sniffed type of the resource is the type given in the cell of the third column in that row; abort these steps.

↪ If the row has a "WS" byte:

1. Let $\text{index}_{\text{pattern}}$ be an index into the mask and pattern byte strings of the row.
2. Let $\text{index}_{\text{stream}}$ be an index into the byte stream being examined.
3. Loop: If $\text{index}_{\text{stream}}$ points beyond the end of the byte stream, then this row doesn't match, skip this row.
4. Examine the $\text{index}_{\text{stream}}$ th byte of the byte stream as follows:

↪ If the $\text{index}_{\text{pattern}}$ th byte of the pattern is a normal hexadecimal byte and not a "WS" byte:

If the "and" operator, applied to the $\text{index}_{\text{stream}}$ th byte of the stream and the $\text{index}_{\text{pattern}}$ th byte of the mask, yield a value different than the $\text{index}_{\text{pattern}}$ th byte of the pattern, then skip this row.

Otherwise, increment $\text{index}_{\text{pattern}}$ to the next byte in the mask and pattern and $\text{index}_{\text{stream}}$ to the next byte in the byte stream.

↪ Otherwise, if the $\text{index}_{\text{pattern}}$ th byte of the pattern is a "WS" byte:

"WS" means "whitespace", and allows insignificant whitespace to be skipped when sniffing for a type signature.

If the $\text{index}_{\text{stream}}$ th byte of the stream is one of 0x09 (ASCII TAB), 0x0A (ASCII LF), 0x0C (ASCII FF), 0x0D (ASCII CR), or 0x20 (ASCII space), then increment only the $\text{index}_{\text{stream}}$ to the next byte in the byte stream.

Otherwise, increment only the $\text{index}_{\text{pattern}}$ to the next byte in the mask and pattern.

5. If $index_{pattern}$ does not point beyond the end of the mask and pattern byte strings, then jump back to the *loop* step in this algorithm.
6. Otherwise, the sniffed type of the resource is the type given in the cell of the third column in that row; abort these steps.
4. If none of the first n bytes of the resource are binary data bytes (page 62) then the sniffed type of the resource is "text/plain". Abort these steps.
5. Otherwise, the sniffed type of the resource is "application/octet-stream".

The table used by the above algorithm is:

Bytes in Hexadecimal		Sniffed type	Security	Comment
Mask	Pattern			
FF FF DF DF DF	3C 21 44 4F 43	text/html	Scriptable	The string "<!DOCTYPE HTML" in US-ASCII or compatible encodings, case-insensitively.
DF DF DF DF FF	54 59 50 45 20			
DF DF DF DF	48 54 4D 4C			
FF FF DF DF DF	WS 3C 48 54 4D	text/html	Scriptable	The string "<HTML" in US-ASCII or compatible encodings, case-insensitively, possibly with leading spaces.
DF	4C			
FF FF DF DF DF	WS 3C 48 45 41	text/html	Scriptable	The string "<HEAD" in US-ASCII or compatible encodings, case-insensitively, possibly with leading spaces.
DF	44			
FF FF DF DF DF	WS 3C 53 43 52	text/html	Scriptable	The string "<SCRIPT" in US-ASCII or compatible encodings, case-insensitively, possibly with leading spaces.
DF DF DF	49 50 54			
FF FF FF FF FF	25 50 44 46 2D	application/pdf	Scriptable	The string "%PDF-", the PDF signature.
FF FF FF FF FF FF	25 21 50 53 2D			
FF FF FF FF FF	41 64 6F 62 65			
2D				
FF FF 00 00	FE FF 00 00	text/plain	n/a	UTF-16BE BOM
FF FF 00 00	FF FF 00 00	text/plain	n/a	UTF-16LE BOM
FF FF FF 00	EF BB BF 00	text/plain	n/a	UTF-8 BOM
FF FF FF FF FF FF	47 49 46 38 37	image/gif	Safe	The string "GIF87a", a GIF signature.
61				
FF FF FF FF FF FF	47 49 46 38 39	image/gif	Safe	The string "GIF89a", a GIF signature.
61				
FF FF FF FF FF FF	89 50 4E 47 0D	image/png	Safe	The PNG signature.
FF FF	0A 1A 0A			
FF FF FF	FF D8 FF	image/jpeg	Safe	A JPEG SOI marker followed by the first byte of another marker.
FF FF	42 4D	image/bmp	Safe	The string "BM", a BMP signature.
FF FF FF FF	00 00 01 00	image/vnd.microsoft.icon	Safe	A 0 word following by a 1 word, a Windows Icon file format signature.

** I'd like to add types like MPEG, AVI, Flash, Java, etc, to the above table.

User agents may support further types if desired, by implicitly adding to the above table. However, user agents should not use any other patterns for types already mentioned in the table above, as this could then be used for privilege escalation (where, e.g., a server uses the

above table to determine that content is not HTML and thus safe from XSS attacks, but then a user agent detects it as HTML anyway and allows script to execute).

The column marked "security" is used by the algorithm in the "text or binary" section, to avoid sniffing text/plain content as a type that can be used for a privilege escalation attack.

2.7.5 Content-Type sniffing: image

If the resource's *official type* is "image/svg+xml", then the sniffed type of the resource is its *official type* (an XML type).

Otherwise, if the first bytes of the resource match one of the byte sequences in the first column of the following table, then the sniffed type of the resource is the type given in the corresponding cell in the second column on the same row:

Bytes in Hexadecimal	Sniffed type	Comment
47 49 46 38 37 61	image/gif	The string "GIF87a", a GIF signature.
47 49 46 38 39 61	image/gif	The string "GIF89a", a GIF signature.
89 50 4E 47 0D 0A 1A 0A	image/png	The PNG signature.
FF D8 FF	image/jpeg	A JPEG SOI marker followed by the first byte of another marker.
42 4D	image/bmp	The string "BM", a BMP signature.
00 00 01 00	image/vnd.microsoft.icon	A 0 word following by a 1 word, a Windows Icon file format signature.

Otherwise, the sniffed type of the resource is the same as its *official type*.

2.7.6 Content-Type sniffing: feed or HTML

1. The user agent may wait for 512 or more bytes of the resource to be available.
2. Let s be the stream of bytes, and let $s[i]$ represent the byte in s with position i , treating s as zero-indexed (so the first byte is at $i=0$).
3. If at any point this algorithm requires the user agent to determine the value of a byte in s which is not yet available, or which is past the first 512 bytes of the resource, or which is beyond the end of the resource, the user agent must stop this algorithm, and assume that the sniffed type of the resource is "text/html".

Note: User agents are allowed, by the first step of this algorithm, to wait until the first 512 bytes of the resource are available.

4. Initialize pos to 0.
5. If $s[0]$ is 0xEF, $s[1]$ is 0xBB, and $s[2]$ is 0xBF, then set pos to 3. (This skips over a leading UTF-8 BOM, if any.)
6. *Loop start:* Examine $s[pos]$.

↪ If it is 0x09 (ASCII tab), 0x20 (ASCII space), 0x0A (ASCII LF), or 0x0D (ASCII CR)

Increase *pos* by 1 and repeat this step.

↪ If it is 0x3C (ASCII "<")

Increase *pos* by 1 and go to the next step.

↪ If it is anything else

The sniffed type of the resource is "text/html". Abort these steps.

7. If the bytes with positions *pos* to *pos*+2 in *s* are exactly equal to 0x21, 0x2D, 0x2D respectively (ASCII for "!--"), then:

1. Increase *pos* by 3.
2. If the bytes with positions *pos* to *pos*+2 in *s* are exactly equal to 0x2D, 0x2D, 0x3E respectively (ASCII for "-->"), then increase *pos* by 3 and jump back to the previous step (the step labeled *loop start*) in the overall algorithm in this section.
3. Otherwise, increase *pos* by 1.
4. Return to step 2 in these substeps.

8. If *s[pos]* is 0x21 (ASCII "!"):

1. Increase *pos* by 1.
2. If *s[pos]* equal 0x3E, then increase *pos* by 1 and jump back to the step labeled *loop start* in the overall algorithm in this section.
3. Otherwise, return to step 1 in these substeps.

9. If *s[pos]* is 0x3F (ASCII "?"):

1. Increase *pos* by 1.
2. If *s[pos]* and *s[pos+1]* equal 0x3F and 0x3E respectively, then increase *pos* by 1 and jump back to the step labeled *loop start* in the overall algorithm in this section.
3. Otherwise, return to step 1 in these substeps.

10. Otherwise, if the bytes in *s* starting at *pos* match any of the sequences of bytes in the first column of the following table, then the user agent must follow the steps given in the corresponding cell in the second column of the same row.

Bytes in Hexadecimal	Requirement	Comment
72 73 73	The sniffed type of the resource is "application/rss+xml"; abort these steps	The three ASCII characters "rss"
66 65 65 64	The sniffed type of the resource is "application/atom+xml"; abort these steps	The four ASCII characters "feed"
72 64 66 3A 52 44 46	Continue to the next step in this algorithm	The ASCII characters "rdf:RDF"

If none of the byte sequences above match the bytes in *s* starting at *pos*, then the sniffed type of the resource is "text/html". Abort these steps.

- ** 11. If, before the next ">", you find two xmlns* attributes with <http://www.w3.org/1999/02/22-rdf-syntax-ns#> and <http://purl.org/rss/1.0/> as the namespaces, then the sniffed type of the resource is "application/rss+xml", abort these steps.
** (maybe we only need to check for <http://purl.org/rss/1.0/> actually)

- 12. Otherwise, the sniffed type of the resource is "text/html".

Note: For efficiency reasons, implementations may wish to implement this algorithm and the algorithm for detecting the character encoding of HTML documents in parallel.

2.8 Common DOM interfaces

2.8.1 Reflecting content attributes in DOM attributes

Some DOM attributes are defined to **reflect** a particular content attribute. This means that on getting, the DOM attribute returns the current value of the content attribute, and on setting, the DOM attribute changes the value of the content attribute to the given value.

If a reflecting DOM attribute is a DOMString attribute whose content attribute is defined to contain a URL (page 52), then on getting, the DOM attribute must resolve (page 55) the value of the content attribute and return the resulting absolute URL (page 56) if that was successful, or the empty string otherwise; and on setting, must set the content attribute to the specified literal value. If the content attribute is absent, the DOM attribute must return the default value, if the content attribute has one, or else the empty string.

If a reflecting DOM attribute is a DOMString attribute whose content attribute is defined to contain one or more URLs (page 52), then on getting, the DOM attribute must split the content attribute on spaces and return the concatenation of resolving (page 55) each token URL to an absolute URL (page 56), with a single U+0020 SPACE character between each URL, ignoring any tokens that did not resolve successfully. If the content attribute is absent, the DOM attribute must return the default value, if the content attribute has one, or else the empty string. On setting, the DOM attribute must set the content attribute to the specified literal value.

If a reflecting DOM attribute is a DOMString whose content attribute is an enumerated attribute (page 51), and the DOM attribute is **limited to only known values**, then, on getting, the DOM attribute must return the conforming value associated with the state the attribute is in (in its canonical case), or the empty string if the attribute is in a state that has no associated keyword value; and on setting, if the new value is an ASCII case-insensitive (page 31) match for one of the keywords given for that attribute, then the content attribute must be set to the conforming value associated with the state that the attribute would be in if set to the given new value, otherwise, if the new value is the empty string, then the content attribute must be removed, otherwise, the setter must raise a SYNTAX_ERR exception.

If a reflecting DOM attribute is a DOMString but doesn't fall into any of the above categories, then the getting and setting must be done in a transparent, case-preserving manner.

If a reflecting DOM attribute is a boolean attribute, then on getting the DOM attribute must return true if the attribute is set, and false if it is absent. On setting, the content attribute must be removed if the DOM attribute is set to false, and must be set to have the same value as its name if the DOM attribute is set to true. (This corresponds to the rules for boolean content attributes (page 32).)

If a reflecting DOM attribute is a signed integer type (`long`) then, on getting, the content attribute must be parsed according to the rules for parsing signed integers (page 33), and if that is successful, the resulting value must be returned. If, on the other hand, it fails, or if the attribute is absent, then the default value must be returned instead, or 0 if there is no default value. On setting, the given value must be converted to the shortest possible string representing the number as a valid integer (page 33) in base ten and then that string must be used as the new content attribute value.

If a reflecting DOM attribute is an *unsigned* integer type (`unsigned long`) then, on getting, the content attribute must be parsed according to the rules for parsing unsigned integers (page 33), and if that is successful, the resulting value must be returned. If, on the other hand, it fails, or if the attribute is absent, the default value must be returned instead, or 0 if there is no default value. On setting, the given value must be converted to the shortest possible string representing the number as a valid non-negative integer (page 33) in base ten and then that string must be used as the new content attribute value.

If a reflecting DOM attribute is an *unsigned* integer type (`unsigned long`) that is **limited to only positive non-zero numbers**, then the behavior is similar to the previous case, but zero is not allowed. On getting, the content attribute must first be parsed according to the rules for parsing unsigned integers (page 33), and if that is successful, the resulting value must be returned. If, on the other hand, it fails, or if the attribute is absent, the default value must be returned instead, or 1 if there is no default value. On setting, if the value is zero, the user agent must fire an `INDEX_SIZE_ERR` exception. Otherwise, the given value must be converted to the shortest possible string representing the number as a valid non-negative integer (page 33) in base ten and then that string must be used as the new content attribute value.

If a reflecting DOM attribute is a floating point number type (`float`) and the content attribute is defined to contain a time offset, then, on getting, the content attribute must be parsed according to the rules for parsing time offsets (page 49), and if that is successful, the resulting value, in seconds, must be returned. If that fails, or if the attribute is absent, the default value must be returned, or the not-a-number value (`Nan`) if there is no default value. On setting, the given value, interpreted as a time offset in seconds, must be converted to a string using the time offset serialization rules (page 49), and that string must be used as the new content attribute value.

If a reflecting DOM attribute is a floating point number type (`float`) and it doesn't fall into one of the earlier categories, then, on getting, the content attribute must be parsed according to the rules for parsing floating point number values (page 34), and if that is successful, the resulting value must be returned. If, on the other hand, it fails, or if the attribute is absent, the default value must be returned instead, or 0.0 if there is no default value. On setting, the given value must be converted to the shortest possible string representing the number as a valid floating point number (page 34) in base ten and then that string must be used as the new content attribute value.

If a reflecting DOM attribute is of the type `DOMTokenList`, then on getting it must return a `DOMTokenList` object whose underlying string is the element's corresponding content

attribute. When the `DOMTokenList` object mutates its underlying string, the `content` attribute must itself be immediately mutated. When the attribute is absent, then the string represented by the `DOMTokenList` object is the empty string; when the object mutates this empty string, the user agent must first add the corresponding `content` attribute, and then mutate that attribute instead. `DOMTokenList` attributes are always read-only. The same `DOMTokenList` object must be returned every time for each attribute.

If a reflecting DOM attribute has the type `HTMLElement`, or an interface that descends from `HTMLElement`, then, on getting, it must run the following algorithm (stopping at the first point where a value is returned):

1. If the corresponding `content` attribute is absent, then the DOM attribute must return null.
2. Let `candidate` be the element that the `document.getElementById()` method would find if it was passed as its argument the current value of the corresponding `content` attribute.
3. If `candidate` is null, or if it is not type-compatible with the DOM attribute, then the DOM attribute must return null.
4. Otherwise, it must return `candidate`.

On setting, if the given element has an `id` attribute, then the `content` attribute must be set to the value of that `id` attribute. Otherwise, the DOM attribute must be set to the empty string.

2.8.2 Collections

The `HTMLCollection`, `HTMLFormControlsCollection`, and `HTMLOptionsCollection` interfaces represent various lists of DOM nodes. Collectively, objects implementing these interfaces are called **collections**.

When a collection (page 69) is created, a filter and a root are associated with the collection.

For example, when the `HTMLCollection` object for the `document.images` attribute is created, it is associated with a filter that selects only `img` elements, and rooted at the root of the document.

The collection (page 69) then **represents** a live (page 24) view of the subtree rooted at the collection's root, containing only nodes that match the given filter. The view is linear. In the absence of specific requirements to the contrary, the nodes within the collection must be sorted in tree order (page 24).

Note: *The rows list is not in tree order.*

An attribute that returns a collection must return the same object every time it is retrieved.

2.8.2.1 `HTMLCollection`

The `HTMLCollection` interface represents a generic collection (page 69) of elements.

```
interface HTMLCollection {
    readonly attribute unsigned long length;
    [IndexGetter] Element item(in unsigned long index);
    [NameGetter] Element namedItem(in DOMString name);
};
```

The **length** attribute must return the number of nodes represented by the collection (page 69).

The **item(index)** method must return the *index*th node in the collection. If there is no *index*th node in the collection, then the method must return null.

The **namedItem(key)** method must return the first node in the collection that matches the following requirements:

- It is an a, applet, area, form, img, or object element with a name attribute equal to key, or,
- It is an HTML element (page 23) of any kind with an id attribute equal to key. (Non-HTML elements, even if they have IDs, are not searched for the purposes of namedItem().)

If no such elements are found, then the method must return null.

2.8.2.2 HTMLFormControlsCollection

The HTMLFormControlsCollection interface represents a collection (page 69) of listed (page 314) elements in form and fieldset elements.

```
interface HTMLFormControlsCollection {
    readonly attribute unsigned long length;
    [IndexGetter] HTMLElement item(in unsigned long index);
    [NameGetter] Object namedItem(in DOMString name);
};
```

The **length** attribute must return the number of nodes represented by the collection (page 69).

The **item(index)** method must return the *index*th node in the collection. If there is no *index*th node in the collection, then the method must return null.

The **namedItem(key)** method must act according to the following algorithm:

1. If, at the time the method is called, there is exactly one node in the collection that has either an id attribute or a name attribute equal to key, then return that node and stop the algorithm.
2. Otherwise, if there are no nodes in the collection that have either an id attribute or a name attribute equal to key, then return null and stop the algorithm.
3. Otherwise, create a NodeList object representing a live view of the HTMLFormControlsCollection object, further filtered so that the only nodes in the

`NodeList` object are those that have either an `id` attribute or a `name` attribute equal to `key`. The nodes in the `NodeList` object must be sorted in tree order (page 24).

4. Return that `NodeList` object.

2.8.2.3 `HTMLOptionsCollection`

The `HTMLOptionsCollection` interface represents a list of option elements. It is always rooted on a `select` element and has attributes and methods that manipulate that element's descendants.

```
interface HTMLOptionsCollection {  
    attribute unsigned long length;  
    [IndexGetter] HTMLOptionElement item(in unsigned long index);  
    [NameGetter] Object namedItem(in DOMString name);  
    void add(in HTMLElement element, in HTMLElement before);  
    void add(in HTMLElement element, in long before);  
    void remove(in long index);  
};
```

On getting, the `length` attribute must return the number of nodes represented by the collection (page 69).

On setting, the behavior depends on whether the new value is equal to, greater than, or less than the number of nodes represented by the collection (page 69) at that time. If the number is the same, then setting the attribute must do nothing. If the new value is greater, then n new option elements with no attributes and no child nodes must be appended to the `select` element on which the `HTMLOptionsCollection` is rooted, where n is the difference between the two numbers (new value minus old value). If the new value is lower, then the last n nodes in the collection must be removed from their parent nodes, where n is the difference between the two numbers (old value minus new value).

Note: Setting `length` never removes or adds any optgroup elements, and never adds new children to existing optgroup elements (though it can remove children from them).

The `item(index)` method must return the `index`th node in the collection. If there is no `index`th node in the collection, then the method must return null.

The `namedItem(key)` method must act according to the following algorithm:

1. If, at the time the method is called, there is exactly one node in the collection that has either an `id` attribute or a `name` attribute equal to `key`, then return that node and stop the algorithm.
2. Otherwise, if there are no nodes in the collection that have either an `id` attribute or a `name` attribute equal to `key`, then return null and stop the algorithm.
3. Otherwise, create a `NodeList` object representing a live view of the `HTMLOptionsCollection` object, further filtered so that the only nodes in the `NodeList` object are those that have either an `id` attribute or a `name` attribute equal to `key`. The nodes in the `NodeList` object must be sorted in tree order (page 24).

4. Return that NodeList object.

The **add(element, before)** method must act according to the following algorithm:

1. If *element* is not an option or optgroup element, then return and abort these steps.
2. If *element* is an ancestor of the select element element on which the HTMLOptionsCollection is rooted, then throw a HIERARCHY_REQUEST_ERR exception.
3. If *before* is an element, but that element isn't a descendant of the select element element on which the HTMLOptionsCollection is rooted, then throw a NOT_FOUND_ERR exception.
4. If *element* and *before* are the same element, then return and abort these steps.
5. If *before* is a node, then let *reference* be that node. Otherwise, if *before* is an integer, and there is a *beforeth* node in the collection, let *reference* be that node. Otherwise, let *reference* be null.
6. If *reference* is not null, let *parent* be the parent node of *reference*. Otherwise, let *parent* be the select element element on which the HTMLOptionsCollection is rooted.
7. Act as if the DOM Core `insertBefore()` method was invoked on the *parent* node, with *element* as the first argument and *reference* as the second argument.

The **remove(index)** method must act according to the following algorithm:

1. If the number of nodes represented by the collection (page 69) is zero, abort these steps.
2. If *index* is not a number greater than or equal to 0 and less than the number of nodes represented by the collection (page 69), let *element* be the first element in the collection. Otherwise, let *element* be the *indexth* element in the collection.
3. Remove *element* from its parent node.

2.8.3 DOMTokenList

The DOMTokenList interface represents an interface to an underlying string that consists of an unordered set of unique space-separated tokens (page 49).

Which string underlies a particular DOMTokenList object is defined when the object is created. It might be a content attribute (e.g. the string that underlies the classList object is the class attribute), or it might be an anonymous string (e.g. when a DOMTokenList object is passed to an author-implemented callback in the datagrid APIs).

```
[Stringifies] interface DOMTokenList {
  readonly attribute unsigned long length;
  [IndexGetter] DOMString item(in unsigned long index);
  boolean has(in DOMString token);
  void add(in DOMString token);
  void remove(in DOMString token);
```

```
    boolean toggle(in DOMString token);  
};
```

The **length** attribute must return the number of *unique* tokens that result from splitting the underlying string on spaces (page 49).

The **item(index)** method must split the underlying string on spaces (page 49), sort the resulting list of tokens by Unicode codepoint, remove exact duplicates, and then return the *index*th item in this list. If *index* is equal to or greater than the number of tokens, then the method must return null.

The **has(token)** method must run the following algorithm:

1. If the *token* argument contains any space characters (page 32), then raise an INVALID_CHARACTER_ERR exception and stop the algorithm.
2. Otherwise, split the underlying string on spaces (page 49) to get the list of tokens in the object's underlying string.
3. If the token indicated by *token* is one of the tokens in the object's underlying string then return true and stop this algorithm.
4. Otherwise, return false.

The **add(token)** method must run the following algorithm:

1. If the *token* argument contains any space characters (page 32), then raise an INVALID_CHARACTER_ERR exception and stop the algorithm.
2. Otherwise, split the underlying string on spaces (page 49) to get the list of tokens in the object's underlying string.
3. If the given *token* is already one of the tokens in the DOMTokenList object's underlying string then stop the algorithm.
4. Otherwise, if the DOMTokenList object's underlying string is not the empty string and the last character of that string is not a space character (page 32), then append a U+0020 SPACE character to the end of that string.
5. Append the value of *token* to the end of the DOMTokenList object's underlying string.

The **remove(token)** method must run the following algorithm:

1. If the *token* argument contains any space characters (page 32), then raise an INVALID_CHARACTER_ERR exception and stop the algorithm.
2. Otherwise, remove the given *token* from the underlying string (page 50).

The **toggle(token)** method must run the following algorithm:

1. If the *token* argument contains any space characters (page 32), then raise an INVALID_CHARACTER_ERR exception and stop the algorithm.
2. Otherwise, split the underlying string on spaces (page 49) to get the list of tokens in the object's underlying string.

3. If the given *token* is already one of the tokens in the DOMTokenList object's underlying string then remove the given *token* from the underlying string (page 50), and stop the algorithm, returning false.
4. Otherwise, if the DOMTokenList object's underlying string is not the empty string and the last character of that string is not a space character (page 32), then append a U+0020 SPACE character to the end of that string.
5. Append the value of *token* to the end of the DOMTokenList object's underlying string.
6. Return true.

Objects implementing the DOMTokenList interface must **stringify** to the object's underlying string representation.

2.8.4 DOMStringMap

The DOMStringMap interface represents a set of name-value pairs. When a DOMStringMap object is instanced, it is associated with three algorithms, one for getting values from names, one for setting names to certain values, and one for deleting names.

****** The names of the methods on this interface are temporary and will be fixed when the Web IDL / "Language Bindings for DOM Specifications" spec is ready to handle this case.

```
interface DOMStringMap {
  [NameGetter] DOMString XXX1(in DOMString name);
  [NameSetter] void XXX2(in DOMString name, in DOMString value);
  [XXX] boolean XXX3(in DOMString name);
};
```

The **XXX1(*name*)** method must call the algorithm for getting values from names, passing *name* as the name, and must return the corresponding value, or null if *name* has no corresponding value.

The **XXX2(*name*, *value*)** method must call the algorithm for setting names to certain values, passing *name* as the name and *value* as the value.

The **XXX3(*name*)** method must call the algorithm for deleting names, passing *name* as the name, and must return true.

2.8.5 DOM feature strings

DOM3 Core defines mechanisms for checking for interface support, and for obtaining implementations of interfaces, using feature strings. [DOM3CORE]

A DOM application can use the **hasFeature(*feature*, *version*)** method of the DOMImplementation interface with parameter values "HTML" and "5.0" (respectively) to determine whether or not this module is supported by the implementation. In addition to the feature string "HTML", the feature string "XHTML" (with version string "5.0") can be used to check if the implementation supports XHTML. User agents should respond with a true value when the hasFeature method is queried with these values. Authors are cautioned, however,

that UAs returning true might not be perfectly compliant, and that UAs returning false might well have support for features in this specification; in general, therefore, use of this method is discouraged.

The values "HTML" and "XHTML" (both with version "5.0") should also be supported in the context of the `getFeature()` and `isSupported()` methods, as defined by DOM3 Core.

Note: The interfaces defined in this specification are not always supersets of the interfaces defined in DOM2 HTML; some features that were formerly deprecated, poorly supported, rarely used or considered unnecessary have been removed. Therefore it is not guaranteed that an implementation that supports "HTML" "5.0" also supports "HTML" "2.0".

3 Semantics and structure of HTML documents

3.1 Introduction

This section is non-normative.

** An introduction to marking up a document.

3.2 Documents

Every XML and HTML document in an HTML UA is represented by a Document object.
[DOM3CORE]

3.2.1 Documents in the DOM

Document objects are assumed to be **XML documents** unless they are flagged as being **HTML documents** when they are created. Whether a document is an HTML document (page 76) or an XML document (page 76) affects the behavior of certain APIs, as well as a few CSS rendering rules. [CSS21]

Note: A Document object created by the `createDocument()` API on the `DOMImplementation` object is initially an XML document (page 76), but can be made into an HTML document (page 76) by calling `document.open()` on it.

All Document objects (in user agents implementing this specification) must also implement the `HTMLDocument` interface, available using binding-specific methods. (This is the case whether or not the document in question is an HTML document (page 76) or indeed whether it contains any HTML elements (page 23) at all.) Document objects must also implement the document-level interface of any other namespaces found in the document that the UA supports. For example, if an HTML implementation also supports SVG, then the Document object must implement `HTMLDocument` and `SVGDocument`.

Note: Because the `HTMLDocument` interface is now obtained using binding-specific casting methods instead of simply being the primary interface of the document object, it is no longer defined as inheriting from `Document`.

```
interface HTMLDocument {
    // resource metadata management
    [PutForwards:href] readonly attribute Location location;
    readonly attribute DOMString URL;
        attribute DOMString domain;
    readonly attribute DOMString referrer;
        attribute DOMString cookie;
    readonly attribute DOMString lastModified;
    readonly attribute DOMString compatMode;
        attribute DOMString charset;
    readonly attribute DOMString characterSet;
```

```

readonly attribute DOMString defaultCharset;
readonly attribute DOMString readyState;

// DOM tree accessors
    attribute DOMString title;
    attribute DOMString dir;
    attribute HTMLElement body;
readonly attribute HTMLCollection images;
readonly attribute HTMLCollection embeds;
readonly attribute HTMLCollection plugins;
readonly attribute HTMLCollection links;
readonly attribute HTMLCollection forms;
readonly attribute HTMLCollection anchors;
readonly attribute HTMLCollection scripts;
 NodeList getElementsByName(in DOMString elementName);
 NodeList getElementsByClassName(in DOMString classNames);

// dynamic markup insertion
    attribute DOMString innerHTML;
HTMLDocument open();
HTMLDocument open(in DOMString type);
HTMLDocument open(in DOMString type, in DOMString replace);
Window open(in DOMString url, in DOMString name, in DOMString
features);
    Window open(in DOMString url, in DOMString name, in DOMString
features, in boolean replace);
    void close();
    void write([Variadic] in DOMString text);
    void writeln([Variadic] in DOMString text);

// user interaction
Selection getSelection();
readonly attribute Element activeElement;
boolean hasFocus();
    attribute boolean designMode;
boolean execCommand(in DOMString commandId);
boolean execCommand(in DOMString commandId, in boolean showUI);
boolean execCommand(in DOMString commandId, in boolean showUI, in
DOMString value);
boolean queryCommandEnabled(in DOMString commandId);
boolean queryCommandIndeterm(in DOMString commandId);
boolean queryCommandState(in DOMString commandId);
boolean queryCommandSupported(in DOMString commandId);
DOMString queryCommandValue(in DOMString commandId);
readonly attribute HTMLCollection commands;
};


```

Since the `HTMLDocument` interface holds methods and attributes related to a number of disparate features, the members of this interface are described in various different sections.

3.2.2 Security

User agents must raise a security exception (page 430) whenever any of the members of an `HTMLDocument` object are accessed by scripts whose effective script origin (page 423) is not the same (page 426) as the Document's effective script origin (page 423).

3.2.3 Resource metadata management

The `URL` attribute must return the document's address.

The `referrer` attribute must return either the address of the active document (page 414) of the source browsing context (page 473) at the time the navigation was started (that is, the page which navigated (page 473) the browsing context (page 414) to the current document), or the empty string if there is no such originating page, or if the UA has been configured not to report referrers in this case, or if the navigation was initiated for a hyperlink (page 497) with a `noreferrer` keyword.

Note: In the case of HTTP, the `referrer` DOM attribute will match the `Referer` (sic) header that was sent when fetching (page 59) the current page.

Note: Typically user agents are configured to not report referrers in the case where the referrer uses an encrypted protocol and the current page does not (e.g. when navigating from an `https:` page to an `http:` page).

The `cookie` attribute represents the cookies of the resource.

On getting, if the sandboxed origin browsing context flag (page 218) is set on the browsing context (page 414) of the document, the user agent must raise a security exception (page 430). Otherwise, it must return the same string as the value of the `Cookie` HTTP header it would include if fetching (page 59) the resource indicated by the document's address over HTTP, as per RFC 2109 section 4.3.4 or later specifications. [RFC2109] [RFC2965]

On setting, if the sandboxed origin browsing context flag (page 218) is set on the browsing context (page 414) of the document, the user agent must raise a security exception (page 430). Otherwise, the user agent must act as it would when processing cookies if it had just attempted to fetch (page 59) the document's address over HTTP, and had received a response with a `Set-Cookie` header whose value was the specified value, as per RFC 2109 sections 4.3.1, 4.3.2, and 4.3.3 or later specifications. [RFC2109] [RFC2965]

Note: Since the `cookie` attribute is accessible across frames, the path restrictions on cookies are only a tool to help manage which cookies are sent to which parts of the site, and are not in any way a security feature.

The `lastModified` attribute, on getting, must return the date and time of the Document's source file's last modification, in the user's local timezone, in the following format:

1. The month component of the date.
2. A U+002F SOLIDUS character ('/').

3. The day component of the date.
4. A U+002F SOLIDUS character ('/').
5. The year component of the date.
6. A U+0020 SPACE character.
7. The hours component of the time.
8. A U+003A COLON character (:).
9. The minutes component of the time.
10. A U+003A COLON character (:).
11. The seconds component of the time.

All the numeric components above, other than the year, must be given as two digits in the range U+0030 DIGIT ZERO to U+0039 DIGIT NINE representing the number in base ten, zero-padded if necessary.

The Document's source file's last modification date and time must be derived from relevant features of the networking protocols used, e.g. from the value of the HTTP Last-Modified header of the document, or from metadata in the file system for local files. If the last modification date and time are not known, the attribute must return the string 01/01/1970 00:00:00.

A Document is always set to one of three modes: **no quirks mode**, the default; **quirks mode**, used typically for legacy documents; and **limited quirks mode**, also known as "almost standards" mode. The mode is only ever changed from the default by the HTML parser (page 582), based on the presence, absence, or value of the DOCTYPE string.

The **compatMode** DOM attribute must return the literal string "CSS1Compat" unless the document has been set to quirks mode (page 79) by the HTML parser (page 582), in which case it must instead return the literal string "BackCompat".

** As far as parsing goes, the quirks I know of are:

- Comment parsing is different.
- p can contain table
- Safari and IE have special parsing rules for <% ... %> (even in standards mode, though clearly this should be quirks-only).

**

Documents have an associated **character encoding**. When a Document object is created, the document's character encoding (page 79) must be initialized to UTF-16. Various algorithms during page loading affect this value, as does the charset setter. [IANACHARSET]

The **charset** DOM attribute must, on getting, return the preferred MIME name of the document's character encoding (page 79). On setting, if the new value is an IANA-registered alias for a character encoding, the document's character encoding (page 79) must be set to that character encoding. (Otherwise, nothing happens.)

The **characterSet** DOM attribute must, on getting, return the preferred MIME name of the document's character encoding (page 79).

The **defaultCharset** DOM attribute must, on getting, return the preferred MIME name of a character encoding, possibly the user's default encoding, or an encoding associated with the user's current geographical location, or any arbitrary encoding name.

Each document has a **current document readiness**. When a Document object is created, it must have its current document readiness (page 80) set to the string "loading". Various algorithms during page loading affect this value. When the value is set, the user agent must fire a simple event (page 436) called `readystatechange` at the Document object.

The **readyState** DOM attribute must, on getting, return the current document readiness (page 80).

3.2.4 DOM tree accessors

The `html` element of a document is the document's root element, if there is one and it's an `html` element, or null otherwise.

The `head` element of a document is the first `head` element that is a child of the `html` element (page 80), if there is one, or null otherwise.

The `title` element of a document is the first `title` element in the document (in tree order), if there is one, or null otherwise.

The **title** attribute must, on getting, run the following algorithm:

1. If the root element (page 24) is an `svg` element in the "`http://www.w3.org/2000/svg`" namespace, and the user agent supports SVG, then the getter must return the value that would have been returned by the DOM attribute of the same name on the `SVGDocument` interface.
2. Otherwise, it must return a concatenation of the data of all the child text nodes (page 24) of the `title` element (page 80), in tree order, or the empty string if the `title` element (page 80) is null.

On setting, the following algorithm must be run:

1. If the root element (page 24) is an `svg` element in the "`http://www.w3.org/2000/svg`" namespace, and the user agent supports SVG, then the setter must defer to the setter for the DOM attribute of the same name on the `SVGDocument` interface (if it is `readonly`, then this will raise an exception). Stop the algorithm here.
2. If the `title` element (page 80) is null and the `head` element (page 80) is null, then the attribute must do nothing. Stop the algorithm here.
3. If the `title` element (page 80) is null, then a new `title` element must be created and appended to the `head` element (page 80).
4. The children of the `title` element (page 80) (if any) must all be removed.
5. A single Text node whose data is the new value being assigned must be appended to the `title` element (page 80).

The `title` attribute on the `HTMLDocument` interface should shadow the attribute of the same name on the `SVGDocument` interface when the user agent supports both HTML and SVG.

The `body` element of a document is the first child of the `html` element (page 80) that is either a `body` element or a `frameset` element. If there is no such element, it is null. If the `body` element is null, then when the specification requires that events be fired at "the `body` element", they must instead be fired at the `Document` object.

The `body` attribute, on getting, must return the `body` element (page 81) of the document (either a `body` element, a `frameset` element, or null). On setting, the following algorithm must be run:

1. If the new value is not a `body` or `frameset` element, then raise a `HIERARCHY_REQUEST_ERR` exception and abort these steps.
2. Otherwise, if the new value is the same as the `body` element (page 81), do nothing. Abort these steps.
3. Otherwise, if the `body` element (page 81) is not null, then replace that element with the new value in the DOM, as if the root element's `replaceChild()` method had been called with the new value and the incumbent `body` element (page 81) as its two arguments respectively, then abort these steps.
4. Otherwise, the `body` element (page 81) is null. Append the new value to the root element.

The `images` attribute must return an `HTMLCollection` rooted at the `Document` node, whose filter matches only `img` elements.

The `embeds` attribute must return an `HTMLCollection` rooted at the `Document` node, whose filter matches only `embed` elements.

The `plugins` attribute must return the same object as that returned by the `embeds` attribute.

The `links` attribute must return an `HTMLCollection` rooted at the `Document` node, whose filter matches only `a` elements with `href` attributes and `area` elements with `href` attributes.

The `forms` attribute must return an `HTMLCollection` rooted at the `Document` node, whose filter matches only `form` elements.

The `anchors` attribute must return an `HTMLCollection` rooted at the `Document` node, whose filter matches only `a` elements with `name` attributes.

The `scripts` attribute must return an `HTMLCollection` rooted at the `Document` node, whose filter matches only `script` elements.

The `getElementsByName(name)` method takes a string `name`, and must return a live `NodeList` containing all the `a`, `applet`, `button`, `form`, `iframe`, `img`, `input`, `map`, `meta`, `object`, `select`, and `textarea` elements in that document that have a `name` attribute whose value is equal to the `name` argument (in a case-sensitive (page 31) manner), in tree order (page 24).

The `getElementsByClassName(classNames)` method takes a string that contains an unordered set of unique space-separated tokens (page 49) representing classes. When called, the method must return a live `NodeList` object containing all the elements in the document, in tree order (page 24), that have all the classes specified in that argument,

having obtained the classes by splitting a string on spaces (page 49). If there are no tokens specified in the argument, then the method must return an empty NodeList. If the document is in quirks mode (page 79), then the comparisons for the classes must be done in an ASCII case-insensitive (page 31) manner, otherwise, the comparisons must be done in a case-sensitive (page 31) manner.

The `getElementsByClassName()` method on the `HTMLElement` interface must return a live NodeList with the nodes that the `HTMLDocument getElementsByClassName()` method would return when passed the same argument(s), excluding any elements that are not descendants of the `HTMLElement` object on which the method was invoked.

HTML, SVG, and MathML elements define which classes they are in by having an attribute in the per-element partition with the name `class` containing a space-separated list of classes to which the element belongs. Other specifications may also allow elements in their namespaces to be labeled as being in specific classes.

Given the following XHTML fragment:

```
<div id="example">
  <p id="p1" class="aaa bbb"/>
  <p id="p2" class="aaa ccc"/>
  <p id="p3" class="bbb ccc"/>
</div>
```

A call to `document.getElementById('example').getElementsByClassName('aaa')` would return a NodeList with the two paragraphs p1 and p2 in it.

A call to `getElementsByClassName('ccc bbb')` would only return one node, however, namely p3. A call to
`document.getElementById('example').getElementsByClassName('bbb ccc ')` would return the same thing.

A call to `getElementsByClassName('aaa,bbb')` would return no nodes; none of the elements above are in the "aaa,bbb" class.

Note: The `dir` attribute on the `HTMLDocument` interface is defined along with the `dir` content attribute.

3.3 Elements

3.3.1 Semantics

Elements, attributes, and attribute values in HTML are defined (by this specification) to have certain meanings (semantics). For example, the `ol` element represents an ordered list, and the `lang` attribute represents the language of the content.

Authors must not use elements, attributes, and attribute values for purposes other than their appropriate intended semantic purpose.

For example, the following document is non-conforming, despite being syntactically correct:

```

<!DOCTYPE html>
<html lang="en-GB">
  <head> <title> Demonstration </title> </head>
  <body>
    <table>
      <tr> <td> My favourite animal is the cat. </td> </tr>
      <tr>
        <td>
          <a href="http://example.org/~ernest/"><cite>Ernest</cite></a>,
          in an essay from 1992
        </td>
      </tr>
    </table>
  </body>
</html>

```

...because the data placed in the cells is clearly not tabular data (and the cite element mis-used). A corrected version of this document might be:

```

<!DOCTYPE html>
<html lang="en-GB">
  <head> <title> Demonstration </title> </head>
  <body>
    <blockquote>
      <p> My favourite animal is the cat. </p>
    </blockquote>
    <p>
      <a href="http://example.org/~ernest/">Ernest</a>,
      in an essay from 1992
    </p>
  </body>
</html>

```

This next document fragment, intended to represent the heading of a corporate site, is similarly non-conforming because the second line is not intended to be a heading of a subsection, but merely a subheading or subtitle (a subordinate heading for the same section).

```

<body>
  <h1>ABC Company</h1>
  <h2>Leading the way in widget design since 1432</h2>
  ...

```

The header element should be used in these kinds of situations:

```

<body>
  <header>
    <h1>ABC Company</h1>
    <h2>Leading the way in widget design since 1432</h2>
  </header>
  ...

```

Through scripting and using other mechanisms, the values of attributes, text, and indeed the entire structure of the document may change dynamically while a user agent is processing it. The semantics of a document at an instant in time are those represented by the state of the document at that instant in time, and the semantics of a document can therefore change over time. User agents must update their presentation of the document as this occurs.

HTML has a progress element that describes a progress bar. If its "value" attribute is dynamically updated by a script, the UA would update the rendering to show the progress changing.

3.3.2 Elements in the DOM

The nodes representing HTML elements (page 23) in the DOM must implement, and expose to scripts, the interfaces listed for them in the relevant sections of this specification. This includes HTML elements (page 23) in XML documents (page 76), even when those documents are in another context (e.g. inside an XSLT transform).

Elements in the DOM represent things; that is, they have intrinsic *meaning*, also known as semantics.

For example, an ol element represents an ordered list.

The basic interface, from which all the HTML elements (page 23)' interfaces inherit, and which must be used by elements that have no additional requirements, is the HTMLElement interface.

```
interface HTMLElement : Element {
    // DOM tree accessors
    NodeList getElementsByTagName(in DOMString classNames);

    // dynamic markup insertion
    attribute DOMString innerHTML;
    attribute DOMString outerHTML;
    void insertAdjacentHTML(in DOMString position, in DOMString text);

    // metadata attributes
    attribute DOMString id;
    attribute DOMString title;
    attribute DOMString lang;
    attribute DOMString dir;
    attribute DOMString className;
    readonly attribute DOMTokenList classList;
    readonly attribute DOMStringMap dataset;

    // user interaction
    attribute boolean hidden;
    void click();
    void scrollIntoView();
    void scrollIntoView(in boolean top);
        attribute long tabIndex;
    void focus();
    void blur();
```

```

        attribute boolean draggable;
        attribute DOMString contentEditable;
readonly attribute boolean isContentEditable;
        attribute HTMLElement contextMenu;

// styling
readonly attribute CSSStyleDeclaration style;

// event handler DOM attributes
        attribute EventListener onabort;
        attribute EventListener onbeforeunload;
        attribute EventListener onblur;
        attribute EventListener onchange;
        attribute EventListener onclick;
        attribute EventListener oncontextmenu;
        attribute EventListener ondblclick;
        attribute EventListener ondrag;
        attribute EventListener ondragend;
        attribute EventListener ondragenter;
        attribute EventListener ondragleave;
        attribute EventListener ondragover;
        attribute EventListener ondragstart;
        attribute EventListener ondrop;
        attribute EventListener onerror;
        attribute EventListener onfocus;
        attribute EventListener onhashchange;
        attribute EventListener onkeydown;
        attribute EventListener onkeypress;
        attribute EventListener onkeyup;
        attribute EventListener onload;
        attribute EventListener onmessage;
        attribute EventListener onmousedown;
        attribute EventListener onmousemove;
        attribute EventListener onmouseout;
        attribute EventListener onmouseover;
        attribute EventListener onmouseup;
        attribute EventListener onmousewheel;
        attribute EventListener onresize;
        attribute EventListener onscroll;
        attribute EventListener onselect;
        attribute EventListener onstorage;
        attribute EventListener onsubmit;
        attribute EventListener onunload;

};


```

The `HTMLElement` interface holds methods and attributes related to a number of disparate features, and the members of this interface are therefore described in various different sections of this specification.

3.3.3 Global attributes

The following attributes are common to and may be specified on all HTML elements (page 23) (even those not defined in this specification):

Global attributes:

```
class  
contenteditable  
contextmenu  
dir  
draggable  
id  
hidden  
lang  
ref  
registrationmark  
style  
tabindex  
title
```

In addition, the following event handler content attributes (page 432) may be specified on any HTML element:

Event handler content attributes:

```
onabort  
onbeforeunload  
onblur  
onchange  
onclick  
oncontextmenu  
ondblclick  
ondrag  
ondragend  
ondragenter  
ondragleave  
ondragover  
ondragstart  
ondrop  
onerror  
onfocus  
onhashchange  
onkeydown  
onkeypress  
onkeyup  
onload  
onmessage  
onmousedown  
onmousemove  
onmouseout  
onmouseover  
onmouseup  
onmousewheel
```

```
onresize  
onscroll  
onselect  
onstorage  
onsubmit  
onunload
```

Also, custom data attributes (page 91) (e.g. data-foldername or data-msgid) can be specified on any HTML element, to store custom data specific to the page.

In HTML documents (page 76), elements in the HTML namespace (page 658) may have an `xmlns` attribute specified, if, and only if, it has the exact value "`http://www.w3.org/1999/xhtml`". This does not apply to XML documents (page 76).

Note: In HTML, the `xmlns` attribute has absolutely no effect. It is basically a talisman. It is allowed merely to make migration to and from XHTML mildly easier. When parsed by an HTML parser (page 582), the attribute ends up in no namespace, not the "`http://www.w3.org/2000/xmlns/`" namespace like namespace declaration attributes in XML do.

Note: In XML, an `xmlns` attribute is part of the namespace declaration mechanism, and an element cannot actually have an `xmlns` attribute in no namespace specified.

3.3.3.1 The `id` attribute

The `id` attribute represents its element's unique identifier. The value must be unique in the subtree within which the element finds itself and must contain at least one character. The value must not contain any space characters (page 32).

If the value is not the empty string, user agents must associate the element with the given value (exactly, including any space characters) for the purposes of ID matching within the subtree the element finds itself (e.g. for selectors in CSS or for the `getElementById()` method in the DOM).

Identifiers are opaque strings. Particular meanings should not be derived from the value of the `id` attribute.

This specification doesn't preclude an element having multiple IDs, if other mechanisms (e.g. DOM Core methods) can set an element's ID in a way that doesn't conflict with the `id` attribute.

The `id` DOM attribute must reflect (page 67) the `id` content attribute.

3.3.3.2 The `title` attribute

The `title` attribute represents advisory information for the element, such as would be appropriate for a tooltip. On a link, this could be the title or a description of the target resource; on an image, it could be the image credit or a description of the image; on a paragraph, it could be a footnote or commentary on the text; on a citation, it could be further information about the source; and so forth. The value is text.

If this attribute is omitted from an element, then it implies that the `title` attribute of the nearest ancestor HTML element (page 23) with a `title` attribute set is also relevant to this element. Setting the attribute overrides this, explicitly stating that the advisory information of any ancestors is not relevant to this element. Setting the attribute to the empty string indicates that the element has no advisory information.

If the `title` attribute's value contains U+000A LINE FEED (LF) characters, the content is split into multiple lines. Each U+000A LINE FEED (LF) character represents a line break.

Some elements, such as `link` and `abbr`, define additional semantics for the `title` attribute beyond the semantics described above.

The `title` DOM attribute must reflect (page 67) the `title` content attribute.

3.3.3.3 The `lang` and `xml:lang` attributes

The `lang` attribute specifies the primary **language** for the element's contents and for any of the element's attributes that contain text. Its value must be a valid RFC 3066 language code, or the empty string. [RFC3066]

The `xml:lang` attribute (that is, the `lang` attribute with the `xml` prefix in the <http://www.w3.org/XML/1998/namespace> namespace) is defined in XML. [XML]

If these attributes are omitted from an element, then it implies that the language of this element is the same as the language of the parent element. Setting the attribute to the empty string indicates that the primary language is unknown.

The `lang` attribute may be used on any HTML element (page 23).

The `xml:lang` attribute may be used on HTML elements (page 23) in XML documents (page 76), as well as elements in other namespaces if the relevant specifications allow it (in particular, MathML and SVG allow `xml:lang` attributes to be specified on their elements). If both the `lang` attribute and the `xml:lang` attribute are specified on the same element, they must have exactly the same value when compared in an ASCII case-insensitive (page 31) manner.

Authors must not use the `xml:lang` attribute (that is, the `lang` attribute with the `xml` prefix in the <http://www.w3.org/XML/1998/namespace> namespace) in HTML documents (page 76). To ease migration to and from XHTML, authors may specify an attribute in no namespace with no prefix and with the localname `xml:lang` on HTML elements (page 23) in HTML documents (page 76), but such attributes must only be specified if a `lang` attribute is also specified, and both attributes must have the same value when compared in an ASCII case-insensitive (page 31) manner.

To determine the language of a node, user agents must look at the nearest ancestor element (including the element itself if the node is an element) that has an `xml:lang` attribute set or is an HTML element (page 23) and has a `lang` attribute set. That attribute specifies the language of the node.

If both the `xml:lang` attribute and the `lang` attribute are set on an element, user agents must use the `xml:lang` attribute, and the `lang` attribute must be ignored (page 24) for the purposes of determining the element's language.

If no explicit language is given for the root element (page 24), but there is a document-wide default language (page 116) set, then that is the language of the node.

If there is no document-wide default language (page 116), then language information from a higher-level protocol (such as HTTP), if any, must be used as the final fallback language. In the absence of any language information, the default value is unknown (the empty string).

If the resulting value is not a recognised language code, then it must be treated as an unknown language (as if the value was the empty string).

User agents may use the element's language to determine proper processing or rendering (e.g. in the selection of appropriate fonts or pronunciations, or for dictionary selection).

The **lang** DOM attribute must reflect (page 67) the lang content attribute.

3.3.3.4 The **xml:base** attribute (XML only)

The **xml:base** attribute is defined in XML Base. [XMLBASE]

The **xml:base** attribute may be used on elements of XML documents (page 76). Authors must not use the **xml:base** attribute in HTML documents (page 76).

3.3.3.5 The **dir** attribute

The **dir** attribute specifies the element's text directionality. The attribute is an enumerated attribute (page 51) with the keyword *ltr* mapping to the state *ltr*, and the keyword *rtl* mapping to the state *rtl*. The attribute has no defaults.

The processing of this attribute is primarily performed by the presentation layer. For example, CSS 2.1 defines a mapping from this attribute to the CSS 'direction' and 'unicode-bidi' properties, and defines rendering in terms of those properties.

The directionality of an element, which is used in particular by the canvas element's text rendering API, is either '*ltr*' or '*rtl*'. If the user agent supports CSS and the 'direction' property on this element has a computed value of either '*ltr*' or '*rtl*', then that is the directionality (page 89) of the element. Otherwise, if the element is being rendered, then the directionality (page 89) of the element is the directionality used by the presentation layer, potentially determined from the value of the **dir** attribute on the element. Otherwise, if the element's **dir** attribute has the state *ltr*, the element's directionality is '*ltr*' (left-to-right); if the attribute has the state *rtl*, the element's directionality is '*rtl*' (right-to-left); and otherwise, the element's directionality is the same as its parent element, or '*ltr*' if there is no parent element.

The **dir** DOM attribute on an element must reflect (page 67) the **dir** content attribute of that element, limited to only known values (page 67).

The **dir** DOM attribute on **HTMLDocument** objects must reflect (page 67) the **dir** content attribute of the **html** element (page 80), if any, limited to only known values (page 67). If there is no such element, then the attribute must return the empty string and do nothing on setting.

3.3.3.6 The class attribute

Every HTML element (page 23) may have a `class` attribute specified.

The attribute, if specified, must have a value that is an unordered set of unique space-separated tokens (page 49) representing the various classes that the element belongs to.

The classes that an HTML element (page 23) has assigned to it consists of all the classes returned when the value of the `class` attribute is split on spaces (page 49).

Note: *Assigning classes to an element affects class matching in selectors in CSS, the `getElementsByClassName()` method in the DOM, and other such features.*

Authors may use any value in the `class` attribute, but are encouraged to use the values that describe the nature of the content, rather than values that describe the desired presentation of the content.

The `className` and `classList` DOM attributes must both reflect (page 67) the `class` content attribute.

3.3.3.7 The style attribute

All elements may have the `style` content attribute set. If specified, the attribute must contain only a list of zero or more semicolon-separated (;) CSS declarations. [CSS21]

The attribute, if specified, must be parsed and treated as the body (the part inside the curly brackets) of a declaration block in a rule whose selector matches just the element on which the attribute is set. For the purposes of the CSS cascade, the attribute must be considered to be a 'style' attribute at the author level.

Documents that use `style` attributes on any of their elements must still be comprehensible and usable if those attributes were removed.

Note: *In particular, using the `style` attribute to hide and show content, or to convey meaning that is otherwise not included in the document, is non-conforming.*

The `style` DOM attribute must return a `CSSStyleDeclaration` whose value represents the declarations specified in the attribute, if present. Mutating the `CSSStyleDeclaration` object must create a `style` attribute on the element (if there isn't one already) and then change its value to be a value representing the serialized form of the `CSSStyleDeclaration` object. [CSSOM]

In the following example, the words that refer to colors are marked up using the `span` element and the `style` attribute to make those words show up in the relevant colors in visual media.

```
<p>My sweat suit is <span style="color: green; background: transparent">green</span> and my eyes are <span style="color: blue; background: transparent">blue</span>. </p>
```

3.3.3.8 Embedding custom non-visible data

A **custom data attribute** is an attribute whose name starts with the string "**data-**", has at least one character after the hyphen, is XML-compatible (page 24), has no namespace, and contains no characters in the range U+0041 .. U+005A (LATIN CAPITAL LETTER A .. LATIN CAPITAL LETTER Z).

Note: All attributes in HTML documents (page 76) get lowercased automatically, so the restriction on uppercase letters doesn't affect such documents.

Custom data attributes (page 91) are intended to store custom data private to the page or application, for which there are no more appropriate attributes or elements.

Every HTML element (page 23) may have any number of custom data attributes (page 91) specified, with any value.

The **dataset** DOM attribute provides convenient accessors for all the data-* attributes on an element. On getting, the dataset DOM attribute must return a `DOMStringMap` object, associated with the following three algorithms, which expose these attributes on their element:

The algorithm for getting values from names

1. Let *name* be the concatenation of the string `data-` and the name passed to the algorithm, converted to lowercase (page 31).
2. If the element does not have an attribute with the name *name*, then the name has no corresponding value, abort.
3. Otherwise, return the value of the attribute with the name *name*.

The algorithm for setting names to certain values

1. Let *name* be the concatenation of the string `data-` and the name passed to the algorithm, converted to lowercase (page 31).
2. Let *value* be the value passed to the algorithm.
3. Set the value of the attribute with the name *name*, to the value *value*, replacing any previous value if the attribute already existed. If `setAttribute()` would have raised an exception when setting an attribute with the name *name*, then this must raise the same exception.

The algorithm for deleting names

1. Let *name* be the concatenation of the string `data-` and the name passed to the algorithm, converted to lowercase (page 31).
2. Remove the attribute with the name *name*, if such an attribute exists. Do nothing otherwise.

If a Web page wanted an element to represent a space ship, e.g. as part of a game, it would have to use the `class` attribute along with data-* attributes:

```
<div class="spaceship" data-id="92432"
    data-weapons="laser 2" data-shields="50%"
    data-x="30" data-y="10" data-z="90">
    <button class="fire"
        onclick="spaceships[this.parentNode.dataset.id].fire()">
        Fire
    </button>
</div>
```

Authors should carefully design such extensions so that when the attributes are ignored and any associated CSS dropped, the page is still usable.

User agents must not derive any implementation behavior from these attributes or values. Specifications intended for user agents must not define these attributes to have any meaningful values.

3.4 Content models

All the elements in this specification have a defined content model, which describes what nodes are allowed inside the elements, and thus what the structure of an HTML document or fragment must look like.

Note: As noted in the conformance and terminology sections, for the purposes of determining if an element matches its content model or not, CDATASection nodes in the DOM are treated as equivalent to Text nodes (page 24), and entity reference nodes are treated as if they were expanded in place (page 29).

The space characters (page 32) are always allowed between elements. User agents represent these characters between elements in the source markup as text nodes in the DOM. Empty text nodes (page 24) and text nodes (page 24) consisting of just sequences of those characters are considered **inter-element whitespace**.

Inter-element whitespace (page 92), comment nodes, and processing instruction nodes must be ignored when establishing whether an element matches its content model or not, and must be ignored when following algorithms that define document and element semantics.

An element *A* is said to be **preceded or followed** by a second element *B* if *A* and *B* have the same parent node and there are no other element nodes or text nodes (other than inter-element whitespace (page 92)) between them.

Authors must not use elements in the HTML namespace (page 23) anywhere except where they are explicitly allowed, as defined for each element, or as explicitly required by other specifications. For XML compound documents, these contexts could be inside elements from other namespaces, if those elements are defined as providing the relevant contexts.

The SVG specification defines the SVG foreignObject element as allowing foreign namespaces to be included, thus allowing compound documents to be created by inserting subdocument content under that element. This specification defines the XHTML html element as being allowed where subdocument fragments are allowed in a

compound document. Together, these two definitions mean that placing an XHTML `html` element as a child of an SVG `foreignObject` element is conforming. [SVG]

The Atom specification defines the Atom content element, when its type attribute has the value `xhtml`, as requiring that it contains a single HTML `div` element. Thus, a `div` element is allowed in that context, even though this is not explicitly normatively stated by this specification. [ATOM]

In addition, elements in the HTML namespace (page 23) may be orphan nodes (i.e. without a parent node).

For example, creating a `td` element and storing it in a global variable in a script is conforming, even though `td` elements are otherwise only supposed to be used inside `tr` elements.

```
var data = {  
    name: "Banana",  
    cell: document.createElement('td'),  
};
```

3.4.1 Kinds of content

Each element in HTML falls into zero or more categories that group elements with similar characteristics together. The following categories are used in this specification:

- Metadata content (page 93)
- Flow content (page 93)
- Sectioning content (page 94)
- Heading content (page 94)
- Phrasing content (page 94)
- Embedded content (page 94)
- Form control content
- Interactive content (page 95)

Some elements have unique requirements and do not fit into any particular category.

3.4.1.1 Metadata content

Metadata content is content that sets up the presentation or behavior of the rest of the content, or that sets up the relationship of the document with other documents, or that conveys other "out of band" information.

Elements from other namespaces whose semantics are primarily metadata-related (e.g. RDF) are also metadata content (page 93).

3.4.1.2 Flow content

Most elements that are used in the body of documents and applications are categorized as **flow content**.

As a general rule, elements whose content model allows any flow content (page 93) should have either at least one descendant text node that is not inter-element whitespace (page 92), or at least one descendant element node that is embedded content (page 94). For the

purposes of this requirement, `del` elements and their descendants must not be counted as contributing to the ancestors of the `del` element.

This requirement is not a hard requirement, however, as there are many cases where an element can be empty legitimately, for example when it is used as a placeholder which will later be filled in by a script, or when the element is part of a template and would on most pages be filled in but on some pages is not relevant.

3.4.1.3 Sectioning content

Sectioning content is content that defines the scope of headers (page 94), footers (page 138), and contact information (page 139).

Each sectioning content (page 94) element potentially has a heading. See the section on headings and sections (page 140) for further details.

3.4.1.4 Heading content

Heading content defines the header of a section (whether explicitly marked up using sectioning content (page 94) elements, or implied by the heading content itself).

3.4.1.5 Phrasing content

Phrasing content is the text of the document, as well as elements that mark up that text at the intra-paragraph level. Runs of phrasing content (page 94) form paragraphs (page 96).

All phrasing content (page 94) is also flow content (page 93). Any content model that expects flow content (page 93) also expects phrasing content (page 94).

As a general rule, elements whose content model allows any phrasing content (page 94) should have either at least one descendant text node that is not inter-element whitespace (page 92), or at least one descendant element node that is embedded content (page 94). For the purposes of this requirement, nodes that are descendants of `del` elements must not be counted as contributing to the ancestors of the `del` element.

Note: *Most elements that are categorized as phrasing content can only contain elements that are themselves categorized as phrasing content, not any flow content.*

Text nodes that are not inter-element whitespace (page 92) are phrasing content (page 94).

3.4.1.6 Embedded content

Embedded content is content that imports another resource into the document, or content from another vocabulary that is inserted into the document.

All embedded content (page 94) is also phrasing content (page 94) (and flow content (page 93)). Any content model that expects phrasing content (page 94) (or flow content (page 93)) also expects embedded content (page 94).

Elements that are from namespaces other than the HTML namespace (page 658) and that convey content but not metadata, are embedded content (page 94) for the purposes of the content models defined in this specification. (For example, MathML, or SVG.)

Some embedded content elements can have **fallback content**: content that is to be used when the external resource cannot be used (e.g. because it is of an unsupported format). The element definitions state what the fallback is, if any.

3.4.1.7 Interactive content

Interactive content is content that is specifically intended for user interaction.

Certain elements in HTML have an activation behavior (page 96), which means the user agent should allow the user to manually trigger them in some way, for instance using keyboard or voice input (though not mouse clicks, which are handled above). When the user triggers an element with a defined activation behavior (page 96), the default action of the interaction event must be to run synthetic click activation steps (page 95) on the element.

When a user agent is to **run synthetic click activation steps** on an element, the user agent must run pre-click activation steps (page 96) on the element, then fire a **click** event (page 436) at the element. The default action of this click event must be to run post-click activation steps (page 96) on the element. If the event is canceled, the user agent must run canceled activation steps (page 96) on the element instead.

Given an element *target*, the **nearest activatable element** is the element returned by the following algorithm:

1. If *target* has a defined activation behavior (page 96), then return *target* and abort these steps.
2. If *target* has a parent element, then set *target* to that parent element and return to the first step.
3. Otherwise, there is no nearest activatable element (page 95).

When a pointing device is clicked, the user agent must run these steps:

1. Let *e* be the nearest activatable element of the element designated by the user, if any.
2. If there is an element *e*, run pre-click activation steps (page 96) on it.
3. Dispatching the required **click** event.

**

Another specification presumably requires the firing of the **click** event?

If there is an element *e*, then the default action of the **click** event must be to run post-click activation steps (page 96) on element *e*.

If there is an element *e* but the event is canceled, the user agent must run canceled activation steps (page 96) on element *e*.

Note: The above doesn't happen for arbitrary synthetic events dispatched by author script. However, the click() method can be used to make it happen programmatically.

When a user agent is to **run post-click activation steps** on an element, the user agent must fire a simple event (page 436) called DOMActivate at that element. The default action of this event must be to run final activation steps (page 96) on that element. If the event is canceled, the user agent must run canceled activation steps (page 96) on the element instead.

When a user agent is to **run pre-click activation steps** on an element, it must run the **pre-click activation steps** defined for that element, if any.

When a user agent is to **run canceled activation steps** on an element, it must run the **canceled activation steps** defined for that element, if any.

When a user agent is to **run final activation steps** on an element, it must run the **activation behavior** defined for that element. Activation behaviors can refer to the click and DOMActivate events that were fired by the steps above leading up to this point.

3.4.2 Transparent content models

Some elements are described as **transparent**; they have "transparent" as their content model. Some elements are described as **semi-transparent**; this means that part of their content model is "transparent" but that is not the only part of the content model that must be satisfied.

When a content model includes a part that is "transparent", those parts must not contain content that would not be conformant if all transparent and semi-transparent elements in the tree were replaced, in their parent element, by the children in the "transparent" part of their content model, retaining order.

When a transparent or semi-transparent element has no parent, then the part of its content model that is "transparent" must instead be treated as accepting any flow content (page 93).

3.5 Paragraphs

A **paragraph** is typically a block of text with one or more sentences that discuss a particular topic, as in typography, but can also be used for more general thematic grouping. For instance, an address is also a paragraph, as is a part of a form, a byline, or a stanza in a poem.

Paragraphs in flow content (page 93) are defined relative to what the document looks like without the a, ins and del elements complicating matters, since those elements, with their hybrid content models, can straddle paragraph boundaries.

Let *view* be a view of the DOM that replaces all a, ins and del elements in the document with their contents. Then, in *view*, for each run of phrasing content (page 94) uninterrupted by other types of content, in an element that accepts content other than phrasing content (page 94), let *first* be the first node of the run, and let *last* be the last node of the run. For

each run, a paragraph exists in the original DOM from immediately before *first* to immediately after *last*. (Paragraphs can thus span across a, ins and del elements.)

A paragraph (page 96) is also formed explicitly by p elements.

Note: The p element can be used to wrap individual paragraphs when there would otherwise not be any content other than phrasing content to separate the paragraphs from each other.

In the following example, there are two paragraphs in a section. There is also a header, which contains phrasing content that is not a paragraph. Note how the comments and intra-element whitespace do not form paragraphs.

```
<section>
  <h1>Example of paragraphs</h1>
  This is the <em>first</em> paragraph in this example.
  <p>This is the second.</p>
  <!-- This is not a paragraph. -->
</section>
```

The following example takes that markup and puts ins and del elements around some of the markup to show that the text was changed (though in this case, the changes don't really make much sense, admittedly). Notice how this example has exactly the same paragraphs as the previous one, despite the ins and del elements.

```
<section>
  <ins><h1>Example of paragraphs</h1>
  This is the <em>first</em> paragraph in</ins> this example<del>.
  <p>This is the second.</p></del>
  <!-- This is not a paragraph. -->
</section>
```

In the following example, the link spans half of the first paragraph, all of the header separating the two paragraphs, and half of the second paragraph.

```
<aside>
  Welcome!
  <a href="about.html">
    This is home of...
    <h1>The Falcons!</h1>
    The Lockheed Martin multirole jet fighter aircraft!
  </a>
  This page discusses the F-16 Fighting Falcon's innermost secrets.
</aside>
```

Here is another way of marking this up, this time showing the paragraphs explicitly, and splitting the one link element into three:

```
<aside>
  <p>Welcome! <a href="about.html">This is home of...</a></p>
  <h1><a href="about.html">The Falcons!</a></h1>
  <p><a href="about.html">The Lockheed Martin multirole jet
  fighter aircraft!</a> This page discusses the F-16 Fighting
```

```
Falcon's innermost secrets.</p>
</aside>
```

Note: Generally, having elements straddle paragraph boundaries is best avoided. Maintaining such markup can be difficult.

3.6 APIs in HTML documents

For HTML documents (page 76), and for HTML elements (page 23) in HTML documents (page 76), certain APIs defined in DOM3 Core become case-insensitive or case-changing, as sometimes defined in DOM3 Core, and as summarized or required below. [DOM3CORE].

This does not apply to XML documents (page 76) or to elements that are not in the HTML namespace (page 658) despite being in HTML documents (page 76).

`Element.tagName` and `Node.nodeName`

These attributes must return element names converted to uppercase (page 31), regardless of the case with which they were created.

`Document.createElement()`

The canonical form of HTML markup is all-lowercase; thus, this method will lowercase (page 31) the argument before creating the requisite element. Also, the element created must be in the HTML namespace (page 658).

Note: This doesn't apply to `Document.createElementNS()`. Thus, it is possible, by passing this last method a tag name in the wrong case, to create an element that claims to have the tag name of an element defined in this specification, but doesn't support its interfaces, because it really has another tag name not accessible from the DOM APIs.

`Element.setAttributeNode()`

When an Attr node is set on an HTML element (page 23), it must have its name converted to lowercase (page 31) before the element is affected.

Note: This doesn't apply to `Document.setAttributeNS()`.

`Element.setAttribute()`

When an attribute is set on an HTML element (page 23), the name argument must be converted to lowercase before the element is affected.

Note: This doesn't apply to `Document.setAttributeNS()`.

`Document.getElementsByTagName()` and `Element.getElementsByTagName()`

These methods (but not their namespaced counterparts) must compare the given argument in an ASCII case-insensitive (page 31) manner when looking at HTML elements (page 23), and in a case-sensitive (page 31) manner otherwise.

Note: Thus, in an HTML document (page 76) with nodes in multiple namespaces, these methods will be both case-sensitive and case-insensitive at the same time.

Document.renameNode()

If the new namespace is the HTML namespace (page 658), then the new qualified name must be converted to lowercase (page 31) before the rename takes place.

3.7 Dynamic markup insertion

APIs for dynamically inserting markup into the document interact with the parser, and thus their behavior varies depending on whether they are used with HTML documents (page 76) (and the HTML parser (page 582)) or XHTML in XML documents (page 76) (and the XML parser). The following table cross-references the various versions of these APIs.

	For documents that are HTML documents (page 76)	For documents that are XML documents (page 76)
document.open()	document.open() (page 99)	
document.write()	document.write() in HTML (page 101)	not supported (page 106)
innerHTML	innerHTML in HTML (page 102)	innerHTML in XML (page 104)
outerHTML	outerHTML in HTML (page 102)	not supported (page 106)
insertAdjacentHTML()	insertAdjacentHTML() in HTML (page 103)	not supported (page 106)

Regardless of the parsing mode, the `document.writeln(...)` method must call the `document.write()` method with the same argument(s), plus an extra argument consisting of a string containing a single line feed character (U+000A).

Note: The `innerHTML` attribute applies to both Element nodes as well as Document nodes. The `outerHTML` and `insertAdjacentHTML()` members, on the other hand, only apply to Element nodes.

Note: When inserted using the `document.write()` method, script elements execute (typically synchronously), but when inserted using `innerHTML` and `outerHTML` attributes, they do not execute at all.

3.7.1 Controlling the input stream

The `open()` method comes in several variants with different numbers of arguments.

When called with two or fewer arguments, the method must act as follows:

1. Let `type` be the value of the first argument, if there is one, or "text/html" otherwise.
2. Let `replace` be true if there is a second argument and it is an ASCII case-insensitive (page 31) match for the value "replace", and false otherwise.
3. If the document has an active parser that isn't a script-created parser (page 100), and the insertion point (page 591) associated with that parser's input stream (page

584) is not undefined (that is, it *does* point to somewhere in the input stream), then the method does nothing. Abort these steps and return the Document object on which the method was invoked.

Note: This basically causes document.open() to be ignored when it's called in an inline script found during the parsing of data sent over the network, while still letting it have an effect when called asynchronously or on a document that is itself being spoon-fed using these APIs.

- ** 4. onbeforeunload, onunload, reset timers, empty event queue, kill any pending transactions, XMLHttpRequests, etc
- ** 5. If the document has an active parser, then stop that parser, and throw away any pending content in the input stream. what about if it doesn't, because it's either
- ** like a text/plain, or Atom, or PDF, or XHTML, or image document, or something?
- 6. Remove all child nodes of the document.
- 7. Change the document's character encoding (page 79) to UTF-16.
- 8. Create a new HTML parser (page 582) and associate it with the document. This is a **script-created parser** (meaning that it can be closed by the document.open() and document.close() methods, and that the tokeniser will wait for an explicit call to document.close() before emitting an end-of-file token).
- 9. Mark the document as being an HTML document (page 76) (it might already be so-marked).
- 10. If the type string contains a U+003B SEMICOLON (;) character, remove the first such character and all characters from it up to the end of the string.
Strip all leading and trailing space characters (page 32) from type.

If type is *not* now an ASCII case-insensitive (page 31) match for the string "text/html", then act as if the tokeniser had emitted a start tag token with the tag name "pre", then set the HTML parser (page 582)'s tokenization (page 597) stage's content model flag (page 597) to PLAINTEXT.

Note: All other values are treated as text/html.

- 11. If replace is false, then:
 1. Remove all the entries in the browsing context (page 414)'s session history (page 467) after the current entry (page 468) in its Document's History object
 2. Remove any earlier entries that share the same Document

3. Add a new entry just before the last entry that is associated with the text that was parsed by the previous parser associated with the Document object, as well as the state of the document at the start of these steps. (This allows the user to step backwards in the session history to see the page before it was blown away by the document.open() call.)
12. Finally, set the insertion point (page 591) to point at just before the end of the input stream (page 584) (which at this point will be empty).
13. Return the Document on which the method was invoked.

When called with three or more arguments, the open() method on the HTMLDocument object must call the open() method on the Window interface of the object returned by the defaultView attribute of the DocumentView interface of the HTMLDocument object, with the same arguments as the original call to the open() method, and return whatever that method returned. If the defaultView attribute of the DocumentView interface of the HTMLDocument object is null, then the method must raise an INVALID_ACCESS_ERR exception.

The **close()** method must do nothing if there is no script-created parser (page 100) associated with the document. If there is such a parser, then, when the method is called, the user agent must insert an explicit "EOF" character (page 591) at the insertion point (page 591) of the parser's input stream (page 584).

3.7.2 Dynamic markup insertion in HTML

In HTML, the **document.write(...)** method must act as follows:

1. If the insertion point (page 591) is undefined, the open() method must be called (with no arguments) on the document object. The insertion point (page 591) will point at just before the end of the (empty) input stream (page 584).
2. The string consisting of the concatenation of all the arguments to the method must be inserted into the input stream (page 584) just before the insertion point (page 591).
3. If there is a pending external script, then the method must now return without further processing of the input stream (page 584).
4. Otherwise, the tokeniser must process the characters that were inserted, one at a time, processing resulting tokens as they are emitted, and stopping when the tokeniser reaches the insertion point or when the processing of the tokeniser is aborted by the tree construction stage (this can happen if a script start tag token is emitted by the tokeniser).

Note: If the document.write() method was called from script executing inline (i.e. executing because the parser parsed a set of script tags), then this is a reentrant invocation of the parser (page 584).

5. Finally, the method must return.

On getting, the **innerHTML** DOM attribute must return the result of running the HTML fragment serialization algorithm (page 658) on the node.

On setting, **if the node is a document**, the innerHTML DOM attribute must run the following algorithm:

1. If the document has an active parser, then stop that parser, and throw away any pending content in the input stream. what about if it doesn't, because it's either like a text/plain, or Atom, or PDF, or XHTML, or image document, or something?
2. Remove the children nodes of the Document whose innerHTML attribute is being set.
3. Create a new HTML parser (page 582), in its initial state, and associate it with the Document node.
4. Place into the input stream (page 584) for the HTML parser (page 582) just created the string being assigned into the innerHTML attribute.
5. Start the parser and let it run until it has consumed all the characters just inserted into the input stream. (The Document node will have been populated with elements and a load event will have fired on its body element (page 81).)

Otherwise, if the node is an element, then setting the innerHTML DOM attribute must cause the following algorithm to run instead:

1. Invoke the HTML fragment parsing algorithm (page 661), with the element whose innerHTML attribute is being set as the *context* element, and the string being assigned into the innerHTML attribute as the *input*. Let *new children* be the result of this algorithm.
2. Remove the children of the element whose innerHTML attribute is being set.
3. Let *target document* be the ownerDocument of the Element node whose innerHTML attribute is being set.
4. Set the ownerDocument of all the nodes in *new children* to the *target document*.
5. Append all the *new children* nodes to the node whose innerHTML attribute is being set, preserving their order.

On getting, the **outerHTML** DOM attribute must return the result of running the HTML fragment serialization algorithm (page 658) on a fictional node whose only child is the node on which the attribute was invoked.

On setting, the outerHTML DOM attribute must cause the following algorithm to run:

1. Let *target* be the element whose outerHTML attribute is being set.
2. If *target* has no parent node, then abort these steps. There would be no way to obtain a reference to the nodes created even if the remaining steps were run.

3. If *target*'s parent node is a Document object, throw a NO_MODIFICATION_ALLOWED_ERR exception and abort these steps.
4. Let *parent* be *target*'s parent node, unless that is a DocumentFragment node, in which case let *parent* be an arbitrary body element.
5. Invoke the HTML fragment parsing algorithm (page 661), with *parent* as the *context* element and the string being assigned into the outerHTML attribute as the *input*. Let *new children* be the result of this algorithm.
6. Let *target document* be the ownerDocument of *target*.
7. Set the ownerDocument of all the nodes in *new children* to the *target document*.
8. Remove *target* from its parent node and insert in its place all the *new children* nodes, preserving their order.

The **insertAdjacentHTML(*position*, *text*)** method, when invoked, must run the following steps:

1. Let *position* and *text* be the method's first and second arguments, respectively.
2. Let *target* be the element on which the method was invoked.
3. Use the first matching item from this list:

If *position* is an ASCII case-insensitive (page 31) match for the string "beforebegin"

If *position* is an ASCII case-insensitive (page 31) match for the string "afterend"

If *target* has no parent node, then abort these steps.

If *target*'s parent node is a Document object, then throw a NO_MODIFICATION_ALLOWED_ERR exception and abort these steps.

Otherwise, let *context* be the parent node of *target*.

If *position* is an ASCII case-insensitive (page 31) match for the string "afterbegin"

If *position* is an ASCII case-insensitive (page 31) match for the string "beforeend"

Let *context* be the same as *target*.

Otherwise

Throw a SYNTAX_ERR exception.

4. Invoke the HTML fragment parsing algorithm (page 661), with the *context* element being that selected by the previous step, and *input* being the method's *text* argument. Let *new children* be the result of this algorithm.
5. Let *target document* be the ownerDocument of *target*.
6. Set the ownerDocument of all the nodes in *new children* to the *target document*.

7. Use the first matching item from this list:

If *position* is an ASCII case-insensitive (page 31) match for the string "beforebegin"

Insert all the *new children* nodes immediately before *target*, preserving their order.

If *position* is an ASCII case-insensitive (page 31) match for the string "afterbegin"

Insert all the *new children* nodes before the first child of *target*, if there is one, preserving their order. If there is no such child, append them all to *target*, preserving their order.

If *position* is an ASCII case-insensitive (page 31) match for the string "beforeend"

Append all the *new children* nodes to *target*, preserving their order.

If *position* is an ASCII case-insensitive (page 31) match for the string "afterend"

Insert all the *new children* nodes immediately after *target*, preserving their order.

3.7.3 Dynamic markup insertion in XML

In an XML context, the `innerHTML` DOM attribute on `HTMLElements` must return a string in the form of an internal general parsed entity, and on `HTMLDocuments` must return a string in the form of a document entity. The string returned must be XML namespace-well-formed and must be an isomorphic serialization of all of that node's child nodes, in document order. User agents may adjust prefixes and namespace declarations in the serialization (and indeed might be forced to do so in some cases to obtain namespace-well-formed XML). For the `innerHTML` attribute on `HTMLElement` objects, if any of the elements in the serialization are in no namespace, the default namespace in scope for those elements must be explicitly declared as the empty string. (This doesn't apply to the `innerHTML` attribute on `HTMLDocument` objects.) [XML] [XMLNS]

If any of the following cases are found in the DOM being serialized, the user agent must raise an `INVALID_STATE_ERR` exception:

- A Document node with no child element nodes.
- A DocumentType node that has an external subset public identifier or an external subset system identifier that contains both a U+0022 QUOTATION MARK ("") and a U+0027 APOSTROPHE ("").
- A node with a prefix or local name containing a U+003A COLON (:).
- An Attr node, Text node, CDATASection node, Comment node, or ProcessingInstruction node whose data contains characters that are not matched by the XML Char production. [XML]
- A CDATASection node whose data contains the string "]]>".

- A Comment node whose data contains two adjacent U+002D HYPHEN-MINUS (-) characters or ends with such a character.
- A ProcessingInstruction node whose target name is an ASCII case-insensitive (page 31) match for the string "xml".
- A ProcessingInstruction node whose target name contains a U+003A COLON (":").
- A ProcessingInstruction node whose data contains the string "?>".

Note: These are the only ways to make a DOM unserializable. The DOM enforces all the other XML constraints; for example, trying to set an attribute with a name that contains an equals sign (=) will raise an INVALID_CHARACTER_ERR exception.

On setting, in an XML context, the innerHTML DOM attribute on HTMLElements and HTMLDocuments must run the following algorithm:

1. The user agent must create a new XML parser.
2. If the innerHTML attribute is being set on an element, the user agent must feed the parser just created the string corresponding to the start tag of that element, declaring all the namespace prefixes that are in scope on that element in the DOM, as well as declaring the default namespace (if any) that is in scope on that element in the DOM.

A namespace prefix is in scope if the DOM Core `lookupNamespaceURI()` method on the element would return a non-null value for that prefix.

The default namespace is the namespace for which the DOM Core `isDefaultNamespace()` method on the element would return true.

3. The user agent must feed the parser just created the string being assigned into the innerHTML attribute.
4. If the innerHTML attribute is being set on an element, the user agent must feed the parser the string corresponding to the end tag of that element.
5. If the parser found an XML well-formedness or XML namespace well-formedness error, the attribute's setter must raise a SYNTAX_ERR exception and abort these steps.
6. The user agent must remove the children nodes of the node whose innerHTML attribute is being set.
7. If the attribute is being set on a Document node, let *new children* be the children of the document, preserving their order. Otherwise, the attribute is being set on an Element node; let *new children* be the children of the document's root element, preserving their order.
8. If the attribute is being set on a Document node, let *target document* be that Document node. Otherwise, the attribute is being set on an Element node; let *target document* be the ownerDocument of that Element.
9. Set the ownerDocument of all the nodes in *new children* to the *target document*.

10. Append all the *new children* nodes to the node whose `innerHTML` attribute is being set, preserving their order.

In an XML context, the `document.write()` and `insertAdjacentHTML()` methods, and the `outerHTML` attribute on both getting and setting, must raise an `INVALID_ACCESS_ERR` exception.

4 The elements of HTML

4.1 The root element

4.1.1 The html element

Categories

None.

Contexts in which this element may be used:

As the root element of a document.

Wherever a subdocument fragment is allowed in a compound document.

Content model:

A head element followed by a body element.

Element-specific attributes:

manifest

DOM interface:

Uses HTMLElement.

The html element represents the root of an HTML document.

The **manifest** attribute gives the address of the document's application cache (page 448) manifest (page 449), if there is one. If the attribute is present, the attribute's value must be a valid URL (page 52).

The manifest attribute only has an effect (page 460) during the early stages of document load. Changing the attribute dynamically thus has no effect (and thus, no DOM API is provided for this attribute).

Note: *Later base elements don't affect the resolving of relative URLs (page 55) in manifest attributes, as the attributes are processed before those elements are seen.*

4.2 Document metadata

4.2.1 The head element

Categories

None.

Contexts in which this element may be used:

As the first element in an html element.

Content model:

One or more elements of metadata content (page 93), of which exactly one is a title element.

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The head element collects the document's metadata.

4.2.2 The title element

Categories

Metadata content (page 93).

Contexts in which this element may be used:

In a head element containing no other title elements.

Content model:

Text.

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The title element represents the document's title or name. Authors should use titles that identify their documents even when they are used out of context, for example in a user's history or bookmarks, or in search results. The document's title is often different from its first header, since the first header does not have to stand alone when taken out of context.

There must be no more than one title element per document.

The title element must not contain any elements.

Here are some examples of appropriate titles, contrasted with the top-level headers that might be used on those same pages.

```
<title>Introduction to The Mating Rituals of Bees</title>
...
<h1>Introduction</h1>
<p>This companion guide to the highly successful
<cite>Introduction to Medieval Bee-Keeping</cite> book is...
```

The next page might be a part of the same site. Note how the title describes the subject matter unambiguously, while the first header assumes the reader knows what the context is and therefore won't wonder if the dances are Salsa or Waltz:

```
<title>Dances used during bee mating rituals</title>
...
<h1>The Dances</h1>
```

The string to use as the document's title is given by the `document.title` DOM attribute. User agents should use the document's title when referring to the document in their user interface.

4.2.3 The base element

Categories

Metadata content (page 93).

Contexts in which this element may be used:

In a head element containing no other base elements.

Content model:

Empty.

Element-specific attributes:

`href`
`target`

DOM interface:

```
interface HTMLBaseElement : HTMLElement {  
    attribute DOMString href;  
    attribute DOMString target;  
};
```

The base element allows authors to specify the document base URL (page 54) for the purposes of resolving relative URLs (page 55), and the name of the default browsing context (page 414) for the purposes of following hyperlinks (page 498).

There must be no more than one base element per document.

A base element must have either an `href` attribute, a `target` attribute, or both.

The `href` content attribute, if specified, must contain a valid URL (page 52).

A base element, if it has an `href` attribute, must come before any other elements in the tree that have attributes defined as taking URLs (page 52), except the `html` element (its `manifest` attribute isn't affected by base elements).

Note: If there are multiple base elements with href attributes, all but the first are ignored.

The `target` attribute, if specified, must contain a valid browsing context name or keyword (page 418). User agents use this name when following hyperlinks (page 498).

A base element, if it has a `target` attribute, must come before any elements in the tree that represent hyperlinks (page 497).

Note: If there are multiple base elements with target attributes, all but the first are ignored.

The **href** and **target** DOM attributes must reflect (page 67) the respective content attributes of the same name.

4.2.4 The link element

Categories

Metadata content (page 93).

Contexts in which this element may be used:

Where metadata content (page 93) is expected.

In a noscript element that is a child of a head element.

Content model:

Empty.

Element-specific attributes:

href
rel
media
hreflang
type
sizes

Also, the title attribute has special semantics on this element.

DOM interface:

```
interface HTMLLinkElement : HTMLElement {  
    attribute boolean disabled;  
    attribute DOMString href;  
    attribute DOMString rel;  
    readonly attribute DOMTokenList relList;  
    attribute DOMString media;  
    attribute DOMString hreflang;  
    attribute DOMString type;  
    attribute DOMString sizes;  
};
```

The LinkStyle interface must also be implemented by this element, the styling processing model (page 122) defines how. [CSSOM]

The link element allows authors to link their document to other resources.

The destination of the link is given by the **href** attribute, which must be present and must contain a valid URL (page 52). If the href attribute is absent, then the element does not define a link.

The type of link indicated (the relationship) is given by the value of the **rel** attribute, which must be present, and must have a value that is a set of space-separated tokens (page 49). The allowed values and their meanings (page 500) are defined in a later section. If the rel attribute is absent, or if the value used is not allowed according to the definitions in this specification, then the element does not define a link.

Two categories of links can be created using the link element. **Links to external resources** are links to resources that are to be used to augment the current document, and **hyperlink links** are links to other documents (page 497). The link types section (page 500) defines whether a particular link type is an external resource or a hyperlink. One element can create multiple links (of which some might be external resource links and some might be hyperlinks); exactly which and how many links are created depends on the keywords given in the rel attribute. User agents must process the links on a per-link basis, not a per-element basis.

The exact behavior for links to external resources depends on the exact relationship, as defined for the relevant link type. Some of the attributes control whether or not the external resource is to be applied (as defined below). For external resources that are represented in the DOM (for example, style sheets), the DOM representation must be made available even if the resource is not applied. (However, user agents may opt to only fetch (page 59) such resources when they are needed, instead of pro-actively fetching (page 59) all the external resources that are not applied.)

The semantics of the protocol used (e.g. HTTP) must be followed when fetching external resources. (For example, redirects must be followed and 404 responses must cause the external resource to not be applied.)

Interactive user agents should provide users with a means to follow the hyperlinks (page 498) created using the link element, somewhere within their user interface. The exact interface is not defined by this specification, but it should include the following information (obtained from the element's attributes, again as defined below), in some form or another (possibly simplified), for each hyperlink created with each link element in the document:

- The relationship between this document and the resource (given by the rel attribute)
- The title of the resource (given by the title attribute).
- The address of the resource (given by the href attribute).
- The language of the resource (given by the hreflang attribute).
- The optimum media for the resource (given by the media attribute).

User agents may also include other information, such as the type of the resource (as given by the type attribute).

Note: Hyperlinks created with the link element and its rel attribute apply to the whole page. This contrasts with the rel attribute of a and area elements, which indicates the type of a link whose context is given by the link's location within the document.

The media attribute says which media the resource applies to. The value must be a valid media query (page 30). [MQ]

If the link is a hyperlink (page 111) then the media attribute is purely advisory, and describes for which media the document in question was designed.

However, if the link is an external resource link (page 111), then the media attribute is prescriptive. The user agent must apply the external resource to views while their state

match the listed media and the other relevant conditions apply, and must not apply them otherwise.

The default, if the `media` attribute is omitted, is `all`, meaning that by default links apply to all media.

The `hreflang` attribute on the `link` element has the same semantics as the `hreflang` attribute on hyperlink elements (page 498).

The `type` attribute gives the MIME type of the linked resource. It is purely advisory. The value must be a valid MIME type, optionally with parameters. [RFC2046]

For external resource links (page 111), the `type` attribute is used as a hint to user agents so that they can avoid fetching resources they do not support. If the attribute is present, then the user agent must assume that the resource is of the given type. If the attribute is omitted, but the external resource link type has a default type defined, then the user agent must assume that the resource is of that type. If the UA does not support the given MIME type for the given link relationship, then the UA should not fetch the resource; if the UA does support the given MIME type for the given link relationship, then the UA should fetch (page 59) the resource. If the attribute is omitted, and the external resource link type does not have a default type defined, but the user agent would fetch the resource if the type was known and supported, then the user agent should fetch (page 59) the resource under the assumption that it will be supported.

User agents must not consider the `type` attribute authoritative — upon fetching the resource, user agents must not use the `type` attribute to determine its actual type. Only the actual type (as defined in the next paragraph) is used to determine whether to *apply* the resource, not the aforementioned assumed type.

If the resource is expected to be an image, user agents may apply the image sniffing rules (page 65), with the *official* type being the type determined from the resource's Content-Type metadata (page 60), and use the resulting sniffed type of the resource as if it was the actual type. Otherwise, if the resource is not expected to be an image, or if the user agent opts not to apply those rules, then the user agent must use the resource's Content-Type metadata (page 60) to determine the type of the resource. If there is no type metadata, but the external resource link type has a default type defined, then the user agent must assume that the resource is of that type.

Once the user agent has established the type of the resource, the user agent must apply the resource if it is of a supported type and the other relevant conditions apply, and must ignore the resource otherwise.

If a document contains style sheet links labeled as follows:

```
<link rel="stylesheet" href="A" type="text/plain">
<link rel="stylesheet" href="B" type="text/css">
<link rel="stylesheet" href="C">
```

...then a compliant UA that supported only CSS style sheets would fetch the B and C files, and skip the A file (since `text/plain` is not the MIME type for CSS style sheets).

For files B and C, it would then check the actual types returned by the server. For those that are sent as `text/css`, it would apply the styles, but for those labeled as `text/plain`, or any other type, it would not.

If one the two files was returned without a Content-Type (page 60) metadata, or with a syntactically incorrect type like Content-Type: "null", then the default type for stylesheet links would kick in. Since that default type is text/css, the style sheet would nonetheless be applied.

The **title** attribute gives the title of the link. With one exception, it is purely advisory. The value is text. The exception is for style sheet links, where the **title** attribute defines alternative style sheet sets (page 122).

Note: *The title attribute on link elements differs from the global title attribute of most other elements in that a link without a title does not inherit the title of the parent element: it merely has no title.*

The **sizes** attribute is used with the icon link type. The attribute must not be specified on link elements that do not have a **rel** attribute that specifies the icon keyword.

Some versions of HTTP defined a Link: header, to be processed like a series of link elements. If supported, for the purposes of ordering links defined by HTTP headers must be assumed to come before any links in the document, in the order that they were given in the HTTP entity header. (URIs in these headers are to be processed and resolved according to the rules given in HTTP; the rules of *this* specification don't apply.) [RFC2616] [RFC2068]

The DOM attributes **href**, **rel**, **media**, **hreflang**, and **type**, and **sizes** each must reflect (page 67) the respective content attributes of the same name.

The DOM attribute **relList** must reflect (page 67) the **rel** content attribute.

The DOM attribute **disabled** only applies to style sheet links. When the **link** element defines a style sheet link, then the disabled attribute behaves as defined for the alternative style sheets DOM (page 122). For all other link elements it always return false and does nothing on setting.

4.2.5 The **meta** element

Categories

Metadata content (page 93).

Contexts in which this element may be used:

If the charset attribute is present, or if the element is in the Encoding declaration state (page 117): as the first element in a head element.

If the http-equiv attribute is present, and the element is not in the Encoding declaration state (page 117): in a head element.

If the http-equiv attribute is present, and the element is not in the Encoding declaration state (page 117): in a noscript element that is a child of a head element.

If the name attribute is present: where metadata content (page 93) is expected.

Content model:

Empty.

Element-specific attributes:

name

http-equiv

content
charset (HTML (page 76) only)

DOM interface:

```
interface HTMLMetaElement : HTMLElement {  
    attribute DOMString content;  
    attribute DOMString name;  
    attribute DOMString httpEquiv;  
};
```

The meta element represents various kinds of metadata that cannot be expressed using the title, base, link, style, and script elements.

The meta element can represent document-level metadata with the name attribute, pragma directives with the http-equiv attribute, and the file's character encoding declaration (page 119) when an HTML document is serialized to string form (e.g. for transmission over the network or for disk storage) with the charset attribute.

Exactly one of the name, http-equiv, and charset attributes must be specified.

If either name or http-equiv is specified, then the content attribute must also be specified. Otherwise, it must be omitted.

The **charset** attribute specifies the character encoding used by the document. This is called a character encoding declaration (page 119).

The charset attribute may be specified in HTML documents (page 29) only, it must not be used in XML documents (page 28). If the charset attribute is specified, the element must be the first element in the head element (page 80) of the file.

The **content** attribute gives the value of the document metadata or pragma directive when the element is used for those purposes. The allowed values depend on the exact context, as described in subsequent sections of this specification.

If a meta element has a **name** attribute, it sets document metadata. Document metadata is expressed in terms of name/value pairs, the name attribute on the meta element giving the name, and the content attribute on the same element giving the value. The name specifies what aspect of metadata is being set; valid names and the meaning of their values are described in the following sections. If a meta element has no content attribute, then the value part of the metadata name/value pair is the empty string.

If a meta element has the http-equiv attribute specified, it must be either in a head element or in a noscript element that itself is in a head element. If a meta element does not have the http-equiv attribute specified, it must be in a head element.

The DOM attributes **name** and **content** must reflect (page 67) the respective content attributes of the same name. The DOM attribute **httpEquiv** must reflect (page 67) the content attribute http-equiv.

4.2.5.1 Standard metadata names

This specification defines a few names for the name attribute of the meta element.

application-name

The value must be a short free-form string that giving the name of the Web application that the page represents. If the page is not a Web application, the application-name metadata name must not be used. User agents may use the application name in UI in preference to the page's title, since the title might include status messages and the like relevant to the status of the page at a particular moment in time instead of just being the name of the application.

description

The value must be a free-form string that describes the page. The value must be appropriate for use in a directory of pages, e.g. in a search engine.

generator

The value must be a free-form string that identifies the software used to generate the document. This value must not be used on hand-authored pages.

4.2.5.2 Other metadata names

Extensions to the predefined set of metadata names may be registered in the WHATWG Wiki MetaExtensions page.

Anyone is free to edit the WHATWG Wiki MetaExtensions page at any time to add a type. These new names must be specified with the following information:

Keyword

The actual name being defined. The name should not be confusingly similar to any other defined name (e.g. differing only in case).

Brief description

A short description of what the metadata name's meaning is, including the format the value is required to be in.

Link to more details

A link to a more detailed description of the metadata name's semantics and requirements. It could be another page on the Wiki, or a link to an external page.

Synonyms

A list of other names that have exactly the same processing requirements. Authors should not use the names defined to be synonyms, they are only intended to allow user agents to support legacy content.

Status

One of the following:

Proposal

The name has not received wide peer review and approval. Someone has proposed it and is using it.

Accepted

The name has received wide peer review and approval. It has a specification that unambiguously defines how to handle pages that use the name, including when they use it in incorrect ways.

Unendorsed

The metadata name has received wide peer review and it has been found wanting. Existing pages are using this keyword, but new pages should avoid it. The "brief description" and "link to more details" entries will give details of what authors should use instead, if anything.

If a metadata name is added with the "proposal" status and found to be redundant with existing values, it should be removed and listed as a synonym for the existing value.

Conformance checkers must use the information given on the WHATWG Wiki MetaExtensions page to establish if a value not explicitly defined in this specification is allowed or not. When an author uses a new type not defined by either this specification or the Wiki page, conformance checkers should offer to add the value to the Wiki, with the details described above, with the "proposal" status.

This specification does not define how new values will get approved. It is expected that the Wiki will have a community that addresses this.

Metadata names whose values are to be URLs (page 52) must not be proposed or accepted. Links must be represented using the link element, not the meta element.

4.2.5.3 Pragma directives

When the **http-equiv** attribute is specified on a meta element, the element is a pragma directive.

The http-equiv attribute is an enumerated attribute (page 51). The following table lists the keywords defined for this attribute. The states given in the first cell of the rows with keywords give the states to which those keywords map.

State	Keywords
Content Language (page 116)	Content-Language
Encoding declaration (page 117)	Content-Type
Default style (page 117)	default-style
Refresh (page 117)	refresh

When a meta element is inserted into the document (page 24), if its http-equiv attribute is present and represents one of the above states, then the user agent must run the algorithm appropriate for that state, as described in the following list:

Content language

This pragma sets the **document-wide default language**. Until the pragma is successfully processed, there is no document-wide default language (page 116).

1. If another meta element in the Content Language state (page 116) has already been successfully processed (i.e. when it was inserted the user agent processed it and reached the last step of this list of steps), then abort these steps.

2. If the meta element has no content attribute, or if that attribute's value is the empty string, then abort these steps.
3. Let *input* be the value of the element's content attribute.
4. Let *position* point at the first character of *input*.
5. Skip whitespace (page 32).
6. Collect a sequence of characters (page 32) that are neither space characters (page 32) nor a U+002C COMMA character (",").
7. Let the document-wide default language (page 116) be the string that resulted from the previous step.

For meta elements in the Content Language state (page 116), the content attribute must have a value consisting of a valid RFC 3066 language code. [RFC3066]

Note: This pragma is not exactly equivalent to the HTTP Content-Language header, for instance it only supports one language. [RFC2616]

Encoding declaration state

The Encoding declaration state's (page 117) user agent requirements are all handled by the parsing section of the specification. The state is just an alternative form of setting the charset attribute: it is a character encoding declaration (page 119).

For meta elements in the Encoding declaration state (page 117), the content attribute must have a value that is an ASCII case-insensitive (page 31) match for a string that consists of: the literal string "text/html;", optionally followed by any number of space characters (page 32), followed by the literal string "charset=", followed by the character encoding name of the character encoding declaration (page 119).

If the document contains a meta element in the Encoding declaration state (page 117) then that element must be the first element in the document's head element, and the document must not contain a meta element with the charset attribute present.

The Encoding declaration state (page 117) may be used in HTML documents (page 29) only, elements in that state must not be used in XML documents (page 28).

Default style state

**

- | | |
|----|-----|
| 1. | ... |
|----|-----|

Refresh state

1. If another meta element in the Refresh state (page 117) has already been successfully processed (i.e. when it was inserted the user agent processed it and reached the last step of this list of steps), then abort these steps.
2. If the meta element has no content attribute, or if that attribute's value is the empty string, then abort these steps.
3. Let *input* be the value of the element's content attribute.

4. Let *position* point at the first character of *input*.
5. Skip whitespace (page 32).
6. Collect a sequence of characters (page 32) in the range U+0030 DIGIT ZERO to U+0039 DIGIT NINE, and parse the resulting string using the rules for parsing non-negative integers (page 33). If the sequence of characters collected is the empty string, then no number will have been parsed; abort these steps. Otherwise, let *time* be the parsed number.
7. Collect a sequence of characters (page 32) in the range U+0030 DIGIT ZERO to U+0039 DIGIT NINE and U+002E FULL STOP ("."). Ignore any collected characters.
8. Skip whitespace (page 32).
9. Let *url* be the address of the current page.
10. If the character in *input* pointed to by *position* is a U+003B SEMICOLON (";"), then advance *position* to the next character. Otherwise, jump to the last step.
11. Skip whitespace (page 32).
12. If the character in *input* pointed to by *position* is one of U+0055 LATIN CAPITAL LETTER U or U+0075 LATIN SMALL LETTER U, then advance *position* to the next character. Otherwise, jump to the last step.
13. If the character in *input* pointed to by *position* is one of U+0052 LATIN CAPITAL LETTER R or U+0072 LATIN SMALL LETTER R, then advance *position* to the next character. Otherwise, jump to the last step.
14. If the character in *input* pointed to by *position* is one of U+004C LATIN CAPITAL LETTER L or U+006C LATIN SMALL LETTER L, then advance *position* to the next character. Otherwise, jump to the last step.
15. Skip whitespace (page 32).
16. If the character in *input* pointed to by *position* is a U+003D EQUALS SIGN ("="), then advance *position* to the next character. Otherwise, jump to the last step.
17. Skip whitespace (page 32).
18. Let *url* be equal to the substring of *input* from the character at *position* to the end of the string.
19. Strip any trailing space characters (page 32) from the end of *url*.
20. Strip any U+0009 CHARACTER TABULATION, U+000A LINE FEED (LF), and U+000D CARRIAGE RETURN (CR) characters from *url*.
21. Resolve (page 55) the *url* value to an absolute URL (page 56). (For the purposes of determining the base URL (page 54), the *url* value comes from the value of a content attribute of the meta element.) If this fails, abort these steps.
22. Perform one or more of the following steps:

- Set a timer so that in *time* seconds, adjusted to take into account user or user agent preferences, if the user has not canceled the redirect, the user agent navigates (page 473) the document's browsing context to *url*, with replacement enabled (page 476), and with the document's browsing context as the source browsing context (page 473).
- Provide the user with an interface that, when selected, navigates a browsing context (page 414) to *url*, with the document's browsing context as the source browsing context (page 473).
- Do nothing.

In addition, the user agent may, as with anything, inform the user of any and all aspects of its operation, including the state of any timers, the destinations of any timed redirects, and so forth.

For meta elements in the Refresh state (page 117), the content attribute must have a value consisting either of:

- just a valid non-negative integer (page 33), or
- a valid non-negative integer (page 33), followed by a U+003B SEMICOLON (;), followed by one or more space characters (page 32), followed by either a U+0055 LATIN CAPITAL LETTER U or a U+0075 LATIN SMALL LETTER U, a U+0052 LATIN CAPITAL LETTER R or a U+0072 LATIN SMALL LETTER R, a U+004C LATIN CAPITAL LETTER L or a U+006C LATIN SMALL LETTER L, a U+003D EQUALS SIGN (=), and then a valid URL (page 52).

In the former case, the integer represents a number of seconds before the page is to be reloaded; in the latter case the integer represents a number of seconds before the page is to be replaced by the page at the given URL (page 52).

There must not be more than one meta element with any particular state in the document at a time.

4.2.5.4 Specifying the document's character encoding

A **character encoding declaration** is a mechanism by which the character encoding used to store or transmit a document is specified.

The following restrictions apply to character encoding declarations:

- The character encoding name given must be the name of the character encoding used to serialize the file.
- The value must be a valid character encoding name, and must be the preferred name for that encoding. [IANACHARSET]
- The character encoding declaration must be serialized without the use of character references (page 581) or character escapes of any kind.

If the document does not start with a BOM, and if its encoding is not explicitly given by Content-Type metadata (page 60), then the character encoding used must be an ASCII-compatible character encoding (page 25), and, in addition, if that encoding isn't

US-ASCII itself, then the encoding must be specified using a meta element with a charset attribute or a meta element in the Encoding declaration state (page 117).

If the document contains a meta element with a charset attribute or a meta element in the Encoding declaration state (page 117), then the character encoding used must be an ASCII-compatible character encoding (page 25).

Authors should not use JIS_X0212-1990, x-JIS0208, and encodings based on EBCDIC. Authors should not use UTF-32. Authors must not use the CESU-8, UTF-7, BOCU-1 and SCSU encodings. [CESU8] [UTF7] [BOCU1] [SCSU]

Authors are encouraged to use UTF-8. Conformance checkers may advise against authors using legacy encodings.

In XHTML, the XML declaration should be used for inline character encoding information, if necessary.

4.2.6 The style element

Categories

Metadata content (page 93).

If the scoped attribute is present: flow content (page 93).

Contexts in which this element may be used:

If the scoped attribute is absent: where metadata content (page 93) is expected.

If the scoped attribute is absent: in a noscript element that is a child of a head element.

If the scoped attribute is present: where flow content (page 93) is expected, but before any other flow content (page 93) other than other style elements and inter-element whitespace (page 92).

Content model:

Depends on the value of the type attribute.

Element-specific attributes:

media

type

scoped

Also, the title attribute has special semantics on this element.

DOM interface:

```
interface HTMLStyleElement : HTMLElement {  
    attribute boolean disabled;  
    attribute DOMString media;  
    attribute DOMString type;  
    attribute boolean scoped;  
};
```

The LinkStyle interface must also be implemented by this element, the styling processing model (page 122) defines how. [CSSOM]

The `style` element allows authors to embed style information in their documents. The `style` element is one of several inputs to the styling processing model (page 122).

If the `type` attribute is given, it must contain a valid MIME type, optionally with parameters, that designates a styling language. [RFC2046] If the attribute is absent, the type defaults to `text/css`. [RFC2138]

When examining types to determine if they support the language, user agents must not ignore unknown MIME parameters — types with unknown parameters must be assumed to be unsupported.

The `media` attribute says which media the styles apply to. The value must be a valid media query (page 30). [MQ] User agents must apply the styles to views while their state match the listed media, and must not apply them otherwise. [DOM3VIEWS]

The default, if the `media` attribute is omitted, is `all`, meaning that by default styles apply to all media.

The `scoped` attribute is a boolean attribute (page 32). If the attribute is present, then the user agent must apply the specified style information only to the `style` element's parent element (if any), and that element's child nodes. Otherwise, the specified styles must, if applied, be applied to the entire document.

The `title` attribute on `style` elements defines alternative style sheet sets (page 122). If the `style` element has no `title` attribute, then it has no title; the `title` attribute of ancestors does not apply to the `style` element.

Note: *The `title` attribute on `style` elements, like the `title` attribute on link elements, differs from the global `title` attribute in that a style block without a `title` does not inherit the `title` of the parent element: it merely has no title.*

All descendant elements must be processed, according to their semantics, before the `style` element itself is evaluated. For styling languages that consist of pure text, user agents must evaluate `style` elements by passing the concatenation of the contents of all the text nodes (page 24) that are direct children of the `style` element (not any other nodes such as comments or elements), in tree order (page 24), to the style system. For XML-based styling languages, user agents must pass all the children nodes of the `style` element to the style system.

Note: *This specification does not specify a style system, but CSS is expected to be supported by most Web browsers. [CSS21]*

The `media`, `type` and `scoped` DOM attributes must reflect (page 67) the respective content attributes of the same name.

The DOM `disabled` attribute behaves as defined for the alternative style sheets DOM (page 122).

4.2.7 Styling

The link and style elements can provide styling information for the user agent to use when rendering the document. The DOM Styling specification specifies what styling information is to be used by the user agent and how it is to be used. [CSSOM]

The style and link elements implement the LinkStyle interface. [CSSOM]

For style elements, if the user agent does not support the specified styling language, then the sheet attribute of the element's LinkStyle interface must return null. Similarly, link elements that do not represent external resource links that contribute to the styling processing model (page 508) (i.e. that do not have a `stylesheet` keyword in their `rel` attribute), and link elements whose specified resource has not yet been fetched, or is not in a supported styling language, must have their LinkStyle interface's sheet attribute return null.

Otherwise, the LinkStyle interface's sheet attribute must return a StyleSheet object with the attributes implemented as follows: [CSSOM]

The content type (type DOM attribute)

The content type must be the same as the style's specified type. For style elements, this is the same as the type content attribute's value, or `text/css` if that is omitted. For link elements, this is the Content-Type metadata of the specified resource (page 60).

The location (href DOM attribute)

For link elements, the location must be the result of resolving (page 55) the URL (page 52) given by the element's href content attribute, or the empty string if that fails. For style elements, there is no location.

The intended destination media for style information (media DOM attribute)

The media must be the same as the value of the element's media content attribute.

The style sheet title (title DOM attribute)

The title must be the same as the value of the element's title content attribute. If the attribute is absent, then the style sheet does not have a title. The title is used for defining **alternative style sheet sets**.

The **disabled** DOM attribute on link and style elements must return false and do nothing on setting, if the sheet attribute of their LinkStyle interface is null. Otherwise, it must return the value of the StyleSheet interface's disabled attribute on getting, and forward the new value to that same attribute on setting.

4.3 Scripting

Scripts allow authors to add interactivity to their documents.

Authors are encouraged to use declarative alternatives to scripting where possible, as declarative mechanisms are often more maintainable, and many users disable scripting.

For example, instead of using script to show or hide a section to show more details, the details element could be used.

Authors are also encouraged to make their applications degrade gracefully in the absence of scripting support.

For example, if an author provides a link in a table header to dynamically resort the table, the link could also be made to function without scripts by requesting the sorted table from the server.

4.3.1 The script element

Categories

Metadata content (page 93).
Phrasing content (page 94).

Contexts in which this element may be used:

Where metadata content (page 93) is expected.
Where phrasing content (page 94) is expected.

Content model:

If there is no `src` attribute, depends on the value of the `type` attribute.
If there *is* a `src` attribute, the element must be empty.

Element-specific attributes:

`src`
`async`
`defer`
`type`
`charset`

DOM interface:

```
interface HTMLScriptElement : HTMLElement {  
    attribute DOMString src;  
    attribute boolean async;  
    attribute boolean defer;  
    attribute DOMString type;  
    attribute DOMString charset;  
    attribute DOMString text;  
};
```

The `script` element allows authors to include dynamic script and data blocks in their documents.

When used to include dynamic scripts, the scripts may either be embedded inline or may be imported from an external file using the `src` attribute. If the language is not that described by "text/javascript", then the `type` attribute must be present. If the `type` attribute is present, its value must be the type of the script's language.

When used to include data blocks, the data must be embedded inline, the format of the data must be given using the `type` attribute, and the `src` attribute must not be specified.

The `type` attribute gives the language of the script or format of the data. If the attribute is present, its value must be a valid MIME type, optionally with parameters. The `charset`

parameter must not be specified. (The default, which is used if the attribute is absent, is "text/javascript".) [RFC2046]

The **src** attribute, if specified, gives the address of the external script resource to use. The value of the attribute must be a valid URL (page 52) identifying a script resource of the type given by the **type** attribute, if the attribute is present, or of the type "text/javascript", if the attribute is absent.

The **charset** attribute gives the character encoding of the external script resource. The attribute must not be specified if the **src** attribute is not present. If the attribute is set, its value must be a valid character encoding name, and must be the preferred name for that encoding. [IANACHARSET]

The encoding specified must be the encoding used by the script resource. If the **charset** attribute is omitted, the character encoding of the document will be used. If the script resource uses a different encoding than the document, then the attribute must be specified.

The **async** and **defer** attributes are boolean attributes (page 32) that indicate how the script should be executed.

There are three possible modes that can be selected using these attributes. If the **async** attribute is present, then the script will be executed asynchronously, as soon as it is available. If the **async** attribute is not present but the **defer** attribute is present, then the script is executed when the page has finished parsing. If neither attribute is present, then the script is fetched and executed immediately, before the user agent continues parsing the page. The exact processing details for these attributes is described below.

The **defer** attribute may be specified even if the **async** attribute is specified, to cause legacy Web browsers that only support **defer** (and not **async**) to fall back to the **defer** behavior instead of the synchronous blocking behavior that is the default.

Changing the **src**, **type**, **charset**, **async**, and **defer** attributes dynamically has no direct effect; these attribute are only used at specific times described below (namely, when the element is inserted into the document (page 24)).

script elements have four associated pieces of metadata. The first is a flag indicating whether or not the script block has been "**already executed**". Initially, script elements must have this flag unset (script blocks, when created, are not "already executed"). When a script element is cloned, the "already executed" flag, if set, must be propagated to the clone when it is created. The second is a flag indicating whether the element was "**parser-inserted**". This flag is set by the HTML parser (page 582) and is used to handle `document.write()` calls. The third and fourth pieces of metadata are **the script's type** and **the script's character encoding**. They are determined when the script is run, based on the attributes on the element at that time.

When an XML parser creates a script element, it must be marked as being "parser-inserted" (page 124). When the element's end tag is parsed, the user agent must run (page 125) the script element.

Note: Equivalent requirements exist for the HTML parser (page 582), but they are detailed in that section instead.

When a script element that is marked as neither having "already executed" (page 124) nor being "parser-inserted" (page 124) is inserted into a document (page 24), the user agent must run (page 125) the script element.

Running a script: When a script element is to be run, the user agent must act as follows:

1. If either:

- the script element has a type attribute and its value is the empty string, or
- the script element has no type attribute but it has a language attribute and that attribute's value is the empty string, or
- the script element has neither a type attribute nor a language attribute, then

...let *the script's type* for this script element be "text/javascript".

Otherwise, if the script element has a type attribute, let *the script's type* for this script element be the value of that attribute.

Otherwise, the element has a language attribute; let *the script's type* for this script element be the concatenation of the string "text/" followed by the value of the language attribute.

2. If the script element has a charset attribute, then let *the script's character encoding* for this script element be the encoding given by the charset attribute.

Otherwise, let *the script's character encoding* for this script element be the same as the encoding of the document itself (page 79).

3. If the script element is without script (page 428), or if the script element was created by an XML parser that itself was created as part of the processing of the innerHTML attribute's setter, or if the user agent does not support the scripting language (page 129) given by *the script's type* for this script element, then the user agent must abort these steps at this point. The script is not executed.

4. The user agent must set the element's "already executed" (page 124) flag.

5. If the element has a src attribute, then the specified resource must be fetched (page 59).

For historical reasons, if the URL (page 52) is a javascript: URL (page 430), then the user agent must not, despite the requirements in the definition of the fetching (page 59) algorithm, actually execute the given script; instead the user agent must act as if it had received an empty HTTP 400 response.

Once the fetching process has completed, and the script has **completed loading**, the user agent will have to complete the steps described below (page 126). (If the parser is still active at that time, those steps defer to the parser to handle the execution of pending scripts.)

For performance reasons, user agents may start fetching the script as soon as the attribute is set, instead, in the hope that the element will be inserted into the document. Either way, once the element is inserted into the document (page 24), the load must have started. If the UA performs such prefetching, but the element is never inserted in the document, or the src attribute is dynamically changed, then the user

agent will not execute the script, and the fetching process will have been effectively wasted.

6. Then, the first of the following options that describes the situation must be followed:

- ↪ **If the document is still being parsed, and the element has a defer attribute, and the element does not have an async attribute**

**

The element must be added to the end of the list of scripts that will execute when the document has finished parsing (page 126).

This isn't compatible

**

with IE for inline deferred scripts but then what IE does is pretty hard to pin down exactly. Do we want to keep this like it is? Be more compatible?

- ↪ **If the element has an async attribute and a src attribute**

The element must be added to the end of the list of scripts that will execute asynchronously (page 127).

- ↪ **If the element has an async attribute but no src attribute, and the list of scripts that will execute asynchronously (page 127) is not empty**

The element must be added to the end of the list of scripts that will execute asynchronously (page 127).

- ↪ **If the element has a src attribute and has been flagged as "parser-inserted" (page 124)**

The element is the pending external script. (There can only be one such script at a time.)

- ↪ **If the element has a src attribute**

The element must be added to the end of the list of scripts that will execute as soon as possible (page 127).

- ↪ **Otherwise**

The user agent must immediately execute the script (page 127), even if other scripts are already executing.

When a script completes loading: If the script's element was added to one of the lists mentioned above and the document is still being parsed, then the parser handles it. Otherwise, the UA must run the following steps as the task (page 429) that the networking task source (page 430) places on the task queue (page 429):

- ↪ **If the script's element was added to the *list of scripts that will execute when the document has finished parsing*:**

1. If the script's element is not the first element in the list, then do nothing yet. Stop going through these steps.
2. Otherwise, execute the script (page 127) (that is, the script associated with the first element in the list).
3. Remove the script's element from the list (i.e. shift out the first entry in the list).
4. If there are any more entries in the list, and if the script associated with the element that is now the first in the list is already loaded, then jump back to step two to execute it.

↪ If the script's element was added to the *list of scripts that will execute asynchronously*:

1. If the script is not the first element in the list, then do nothing yet. Stop going through these steps.
2. Execute the script (page 127) (the script associated with the first element in the list).
3. Remove the script's element from the list (i.e. shift out the first entry in the list).
4. If there are any more scripts in the list, and the element now at the head of the list had no `src` attribute when it was added to the list, or had one, but its associated script has finished loading, then jump back to step two to execute the script associated with this element.

↪ If the script's element was added to the *list of scripts that will execute as soon as possible*:

1. Execute the script (page 127).
2. Remove the script's element from the list.

Fetching an external script must delay the `load` event (page 657).

Executing a script block: When the steps above require that the script be executed, the user agent must act as follows:

↪ If the load resulted in an error (for example a **DNS error**, or an **HTTP 404 error**)

Executing the script must just consist of firing an `error` event (page 436) at the element.

↪ If the load was successful

1. If the script element's `Document` is the active document (page 414) in its browsing context (page 414), the user agent must execute the script:

↪ If the script is from an external file

That file must be used as the file to execute.

The file must be interpreted using the character encoding given by *the script's character encoding*, regardless of any metadata given by the file's `Content-Type` metadata (page 60).

**

This means that a **UTF-16** document will always assume external scripts are **UTF-16**...? This applies, e.g., to `document`'s created using `createDocument()`... It also means changing `document.charset` will affect the character encoding used to interpret scripts, is that really what happens?

**

↪ If the script is inline

For scripting languages that consist of pure text, user agents must use the value of the `DOM text` attribute (defined below) as the

script to execute, and for XML-based scripting languages, user agents must use all the child nodes of the `script` element as the script to execute.

In any case, the user agent must execute the script according to the semantics defined by the language associated with *the script's type* (see the scripting languages (page 129) section below).

The script execution context (page 428) of the script must be the `Window` object of that browsing context (page 414).

The script document context (page 428) of the script must be the `Document` object that owns the `script` element.

Note: The element's attributes' values might have changed between when the element was inserted into the document and when the script has finished loading, as may its other attributes; similarly, the element itself might have been taken back out of the DOM, or had other changes made. These changes do not in any way affect the above steps; only the values of the attributes at the time the script element is first inserted into the document matter.

2. Then, the user agent must fire a `load` event (page 436) at the `script` element.

The DOM attributes `src`, `type`, `charset`, `async`, and `defer`, each must reflect (page 67) the respective content attributes of the same name.

The DOM attribute `text` must return a concatenation of the contents of all the text nodes (page 24) that are direct children of the `script` element (ignoring any other nodes such as comments or elements), in tree order. On setting, it must act the same way as the `textContent` DOM attribute.

In this example, two `script` elements are used. One embeds an external script, and the other includes some data.

```
<script src="game-engine.js"></script>
<script type="text/x-game-map">
    .....U.....e
    o.....A....e
    ....A.....AAA....e
    .A..AAA...AAAAA...e
</script>
```

The data in this case might be used by the script to generate the map of a video game. The data doesn't have to be used that way, though; maybe the map data is actually embedded in other parts of the page's markup, and the data block here is just used by the site's search engine to help users who are looking for particular features in their game maps.

4.3.1.1 Scripting languages

A user agent is said to **support the scripting language** if *the script's type* matches the MIME type of a scripting language that the user agent implements.

The following lists some MIME types and the languages to which they refer:

text/javascript
text/javascript1.1
text/javascript1.2
text/javascript1.3

ECMAScript. [ECMA262]

text/javascript;e4x=1

ECMAScript with ECMAScript for XML. [ECMA357]

User agents may support other MIME types and other languages.

When examining types to determine if they support the language, user agents must not ignore unknown MIME parameters — types with unknown parameters must be assumed to be unsupported.

4.3.2 The noscript element

Categories

Metadata content (page 93).

Phrasing content (page 94).

Contexts in which this element may be used:

In a head element of an HTML document (page 76), if there are no ancestor noscript elements.

Where phrasing content (page 94) is expected in HTML documents (page 76), if there are no ancestor noscript elements.

Content model:

Without script (page 428), in a head element: in any order, zero or more link elements, zero or more style elements, and zero or more meta elements.

Without script (page 428), not in a head element: transparent (page 96), but there must be no noscript element descendants.

With script (page 428): text that conforms to the requirements given in the prose.

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The noscript element does not represent anything. It is used to present different markup to user agents that support scripting and those that don't support scripting, by affecting how the document is parsed.

The noscript element must not be used in XML documents (page 76).

Note: The noscript element is only effective in the HTML serialization, it has no effect in the XML serialization.

When used in HTML documents (page 76), the allowed content model is as follows:

In a head element, if the noscript element is without script (page 428), then the content model of a noscript element must contain only link, style, and meta elements. If the noscript element is with script (page 428), then the content model of a noscript element is text, except that invoking the HTML fragment parsing algorithm (page 661) with the noscript element as the *context* element and the text contents as the *input* must result in a list of nodes that consists only of link, style, and meta elements.

Outside of head elements, if the noscript element is without script (page 428), then the content model of a noscript element is transparent (page 96), with the additional restriction that a noscript element must not have a noscript element as an ancestor (that is, noscript can't be nested).

Outside of head elements, if the noscript element is with script (page 428), then the content model of a noscript element is text, except that the text must be such that running the following algorithm results in a conforming document with no noscript elements and no script elements, and such that no step in the algorithm causes an HTML parser (page 582) to flag a parse error (page 583):

1. Remove every script element from the document.
2. Make a list of every noscript element in the document. For every noscript element in that list, perform the following steps:
 1. Let the *parent element* be the parent element of the noscript element.
 2. Take all the children of the *parent element* that come before the noscript element, and call these elements *the before children*.
 3. Take all the children of the *parent element* that come after the noscript element, and call these elements *the after children*.
 4. Let *s* be the concatenation of all the text node (page 24) children of the noscript element.
 5. Set the innerHTML attribute of the *parent element* to the value of *s*. (This, as a side-effect, causes the noscript element to be removed from the document.)
 6. Insert *the before children* at the start of the *parent element*, preserving their original relative order.
 7. Insert *the after children* at the end of the *parent element*, preserving their original relative order.

The noscript element has no other requirements. In particular, children of the noscript element are not exempt from form submission, scripting, and so forth, even when the element is with script (page 428).

Note: All these contortions are required because, for historical reasons, the noscript element is handled differently by the HTML parser (page 582) based on whether scripting was enabled or not (page 597) when the parser was invoked. The element is not allowed in XML, because in XML the parser is not affected by such state, and thus the element would not have the desired effect.

Note: The `noscript` element interacts poorly with the `designMode` feature. Authors are encouraged to not use `noscript` elements on pages that will have `designMode` enabled.

4.3.3 The `eventsource` element

Categories

Metadata content (page 93).
Phrasing content (page 94).

Contexts in which this element may be used:

Where metadata content (page 93) is expected.
Where phrasing content (page 94) is expected.

Content model:

Empty.

Element-specific attributes:

`src`

DOM interface:

```
interface HTMLEventSourceElement : HTMLElement {  
    attribute DOMString src;  
};
```

The `eventsource` element represents a target for events generated by a remote server.

The `src` attribute, if specified, must give a valid URL (page 52) identifying a resource that uses the text/event-stream format.

When an `eventsource` element with a `src` attribute specified is inserted into the document (page 24), and when an `eventsource` element that is already in the document has a `src` attribute added, the user agent must run the `add declared event source` (page 131) algorithm.

While an `eventsource` element is in a document, if its `src` attribute is mutated, the user agent must run the `remove declared event source` (page 132) algorithm followed by the `add declared event source` (page 131) algorithm.

When an `eventsource` element with a `src` attribute specified is removed from a document, and when an `eventsource` element that is in a document with a `src` attribute specified has its `src` attribute removed, the user agent must run the `remove declared event source` (page 132) algorithm.

When it is created, an `eventsource` element must have its *current declared event source* set to "undefined".

The `add declared event source` algorithm is as follows:

1. Resolve (page 55) the URL (page 52) specified by the `eventsource` element's `src` attribute.

2. If that fails, then set the element's *current declared event source* to "undefined" and abort these steps.
3. Otherwise, act as if the `addEventSource()` method on the `eventsources` element had been invoked with the resulting absolute URL (page 56).
4. Let the element's *current declared event source* be that absolute URL (page 56).

The **remove declared event source** algorithm is as follows:

1. If the element's *current declared event source* is "undefined", abort these steps.
2. Otherwise, act as if the `removeEventSource()` method on the `eventsources` element had been invoked with the element's *current declared event source*.
3. Let the element's *current declared event source* be "undefined".

There can be more than one `eventsources` element per document, but authors should take care to avoid opening multiple connections to the same server as HTTP recommends a limit to the number of simultaneous connections that a user agent can open per server.

The **src** DOM attribute must reflect (page 67) the content attribute of the same name.

4.4 Sections

Some elements, for example address elements, are scoped to their nearest ancestor sectioning content (page 94). For such elements `x`, the elements that apply to a sectioning content (page 94) element `e` are all the `x` elements whose nearest sectioning content (page 94) ancestor is `e`.

4.4.1 The body element

Categories

Sectioning content (page 94).

Contexts in which this element may be used:

As the second element in an `html` element.

Content model:

Flow content (page 93).

Element-specific attributes:

None.

DOM interface:

```
interface HTMLBodyElement : HTMLElement {};
```

The `body` element represents the main content of the document.

In conforming documents, there is only one `body` element. The `document.body` DOM attribute provides scripts with easy access to a document's `body` element.

Note: Some DOM operations (for example, parts of the drag and drop (page 525) model) are defined in terms of "the body element (page 81)". This refers to a particular element in the DOM, as per the definition of the term, and not any arbitrary body element.

4.4.2 The section element

Categories

Flow content (page 93).
Sectioning content (page 94).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

Flow content (page 93).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The section element represents a generic document or application section. A section, in this context, is a thematic grouping of content, typically with a header, possibly with a footer.

Examples of sections would be chapters, the various tabbed pages in a tabbed dialog box, or the numbered sections of a thesis. A Web site's home page could be split into sections for an introduction, news items, contact information.

4.4.3 The nav element

Categories

Flow content (page 93).
Sectioning content (page 94).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

Flow content (page 93).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The nav element represents a section of a page that links to other pages or to parts within the page: a section with navigation links. Not all groups of links on a page need to be in a nav element — only sections that consist of primary navigation blocks are appropriate for the nav

element. In particular, it is common for footers to have a list of links to various key parts of a site, but the footer element is more appropriate in such cases.

In the following example, the page has several places where links are present, but only one of those places is considered a navigation section.

```
<body>
  <header>
    <h1>Wake up sheeple!</h1>
    <p><a href="news.html">News</a> -
      <a href="blog.html">Blog</a> -
      <a href="forums.html">Forums</a></p>
  </header>
  <nav>
    <h1>Navigation</h1>
    <ul>
      <li><a href="articles.html">Index of all articles</a><li>
        <li><a href="today.html">Things sheeple need to wake up for
          today</a><li>
        <li><a href="successes.html">Sheeple we have managed to wake</a><li>
    </ul>
  </nav>
  <article>
    <p>...page content would be here...</p>
  </article>
  <footer>
    <p>Copyright © 2006 The Example Company</p>
    <p><a href="about.html">About</a> -
      <a href="policy.html">Privacy Policy</a> -
      <a href="contact.html">Contact Us</a></p>
  </footer>
</body>
```

4.4.4 The article element

Categories

Flow content (page 93).
Sectioning content (page 94).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

Flow content (page 93).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The article element represents a section of a page that consists of a composition that forms an independent part of a document, page, or site. This could be a forum post, a magazine or newspaper article, a Web log entry, a user-submitted comment, or any other independent item of content.

Note: An article element is "independent" in that its contents could stand alone, for example in syndication. However, the element is still associated with its ancestors; for instance, contact information that applies (page 132) to a parent body element still covers the article as well.

When article elements are nested, the inner article elements represent articles that are in principle related to the contents of the outer article. For instance, a Web log entry on a site that accepts user-submitted comments could represent the comments as article elements nested within the article element for the Web log entry.

Author information associated with an article element (q.v. the address element) does not apply to nested article elements.

4.4.5 The aside element

Categories

Flow content (page 93).
Sectioning content (page 94).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

Flow content (page 93).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The aside element represents a section of a page that consists of content that is tangentially related to the content around the aside element, and which could be considered separate from that content. Such sections are often represented as sidebars in printed typography.

The following example shows how an aside is used to mark up background material on Switzerland in a much longer news story on Europe.

```
<aside>
  <h1>Switzerland</h1>
  <p>Switzerland, a land-locked country in the middle of geographic Europe, has not joined the geopolitical European Union, though it is a signatory to a number of European treaties.</p>
</aside>
```

The following example shows how an aside is used to mark up a pull quote in a longer article.

```
...  
<p>He later joined a large company, continuing on the same work.  
<q>I love my job. People ask me what I do for fun when I'm not at  
work. But I'm paid to do my hobby, so I never know what to  
answer. Some people wonder what they would do if they didn't have to  
work... but I know what I would do, because I was unemployed for a  
year, and I filled that time doing exactly what I do  
now.</q></p>  
  
<aside>  
  <q> People ask me what I do for fun when I'm not at work. But I'm  
  paid to do my hobby, so I never know what to answer. </q>  
</aside>  
  
<p>Of course his work – or should that be hobby? –  
isn't his only passion. He also enjoys other pleasures.</p>  
  
...
```

4.4.6 The h1, h2, h3, h4, h5, and h6 elements

Categories

Flow content (page 93).
Heading content (page 94).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

These elements define headers for their sections.

The semantics and meaning of these elements are defined in the section on headings and sections (page 140).

These elements have a **rank** given by the number in their name. The h1 element is said to have the highest rank, the h6 element has the lowest rank, and two elements with the same name have equal rank.

4.4.7 The header element

Categories

Flow content (page 93).

Heading content (page 94).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

Flow content (page 93), including at least one descendant that is heading content (page 94), but no sectioning content (page 94) descendants, no header element descendants, and no footer element descendants.

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The header element represents the header of a section. The element is typically used to group a set of h1-h6 elements to mark up a page's title with its subtitle or tagline. However, header elements may contain more than just the section's headings and subheadings — for example it would be reasonable for the header to include version history information.

For the purposes of document summaries, outlines, and the like, the text of header elements is defined to be the text of the highest ranked (page 136) h1-h6 element descendant of the header element, if there are any such elements, and the first such element if there are multiple elements with that rank (page 136). If there are no such elements, then the text of the header element is the empty string.

Other heading elements in the header element indicate subheadings or subtitles.

The rank (page 136) of a header element is the same as for an h1 element (the highest rank).

The section on headings and sections (page 140) defines how header elements are assigned to individual sections.

Here are some examples of valid headers. In each case, the emphasised text represents the text that would be used as the header in an application extracting header data and ignoring subheadings.

```
<header>
  <h1>The reality dysfunction</h1>
  <h2>Space is not the only void</h2>
</header>
<header>
  <h1>Dr. Strangelove</h1>
  <h2>Or: How I Learned to Stop Worrying and Love the Bomb</h2>
</header>
<header>
  <p>Welcome to...</p>
  <h1>Voidwars!</h1>
</header>
<header>
  <h1>Scalable Vector Graphics (SVG) 1.2</h1>
```

```

<h2>W3C Working Draft 27 October 2004</h2>
<dl>
  <dt>This version:</dt>
  <dd><a href="http://www.w3.org/TR/2004/WD-SVG12-20041027/">http://www.w3.org/TR/2004/WD-SVG12-20041027/</a></dd>
  <dt>Previous version:</dt>
  <dd><a href="http://www.w3.org/TR/2004/WD-SVG12-20040510/">http://www.w3.org/TR/2004/WD-SVG12-20040510/</a></dd>
  <dt>Latest version of SVG 1.2:</dt>
  <dd><a href="http://www.w3.org/TR/SVG12/">http://www.w3.org/TR/SVG12/</a></dd>
  <dt>Latest SVG Recommendation:</dt>
  <dd><a href="http://www.w3.org/TR/SVG/">http://www.w3.org/TR/SVG/</a></dd>
  <dt>Editor:</dt>
  <dd>Dean Jackson, W3C, <a href="mailto:dean@w3.org">dean@w3.org</a></dd>
  <dt>Authors:</dt>
  <dd>See <a href="#authors">Author List</a></dd>
</dl>
<p class="copyright"><a href="http://www.w3.org/Consortium/Legal/ipr-notic ...">http://www.w3.org/Consortium/Legal/ipr-notic ...</a></p>

```

4.4.8 The footer element

Categories

Flow content (page 93).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

Flow content (page 93), but with no heading content (page 94) descendants, no sectioning content (page 94) descendants, and no footer element descendants.

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The footer element represents the footer for the section it applies (page 132) to. A footer typically contains information about its section such as who wrote it, links to related documents, copyright data, and the like.

Contact information for the section given in a footer should be marked up using the address element.

Footers don't necessarily have to appear at the end of a section, though they usually do.

Here is a page with two footers, one at the top and one at the bottom, with the same content:

```
<body>
<footer><a href="..">Back to index...</a></footer>
<h1>Lorem ipsum</h1>
<p>A dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.</p>
<footer><a href="..">Back to index...</a></footer>
</body>
```

4.4.9 The address element

Categories

Flow content (page 93).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

Flow content (page 93), but with no heading content (page 94) descendants, no sectioning content (page 94) descendants, no footer element descendants, and no address element descendants.

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The address element represents the contact information for the section it applies (page 132) to. If it applies to the body element (page 81), then it instead applies to the document as a whole.

For example, a page at the W3C Web site related to HTML might include the following contact information:

```
<ADDRESS>
<A href="..../People/Raggett/">Dave Raggett</A>,
<A href="..../People/Arnaud/">Arnaud Le Hors</A>,
contact persons for the <A href="Activity">W3C HTML Activity</A>
</ADDRESS>
```

The address element must not be used to represent arbitrary addresses (e.g. postal addresses), unless those addresses are contact information for the section. (The p element is the appropriate element for marking up such addresses.)

The address element must not contain information other than contact information.

For example, the following is non-conforming use of the address element:

```
<ADDRESS>Last Modified: 1999/12/24 23:37:50</ADDRESS>
```

Typically, the address element would be included with other information in a footer element.

To determine the contact information for a sectioning content (page 94) element (such as a document's body element, which would give the contact information for the page), UAs must collect all the address elements that apply (page 132) to that sectioning content (page 94) element and its ancestor sectioning content (page 94) elements. The contact information is the collection of all the information given by those elements.

Note: Contact information for one sectioning content (page 94) element, e.g. an aside element, does not apply to its ancestor elements, e.g. the page's body.

4.4.10 Headings and sections

The h1-h6 elements and the header element are headings.

The first element of heading content (page 94) in an element of sectioning content (page 94) gives the header for that section. Subsequent headers of equal or higher rank (page 136) start new (implied) sections, headers of lower rank (page 136) start subsections that are part of the previous one.

Sectioning content (page 94) elements are always considered subsections of their nearest ancestor element of sectioning content (page 94), regardless of what implied sections other headings may have created.

Certain elements are said to be **sectioning roots**, including blockquote and td elements. These elements can have their own outlines, but the sections and headers inside these elements do not contribute to the outlines of their ancestors.

For the following fragment:

```
<body>
  <h1>Foo</h1>
  <h2>Bar</h2>
  <blockquote>
    <h3>Bla</h3>
  </blockquote>
  <p>Baz</p>
  <h2>Quux</h2>
  <section>
    <h3>Thud</h3>
  </section>
  <p>Grunt</p>
</body>
```

...the structure would be:

1. Foo (heading of explicit body section, containing the "Grunt" paragraph)

1. Bar (heading starting implied section, containing a block quote and the "Baz" paragraph)
2. Quux (heading starting implied section)
3. Thud (heading of explicit section section)

Notice how the section ends the earlier implicit section so that a later paragraph ("Grunt") is back at the top level.

Sections may contain headers of any rank (page 136), but authors are strongly encouraged to either use only h1 elements, or to use elements of the appropriate rank (page 136) for the section's nesting level.

Authors are also encouraged to explicitly wrap sections in elements of sectioning content (page 94), instead of relying on the implicit sections generated by having multiple heading in one element of sectioning content (page 94).

For example, the following is correct:

```
<body>
  <h4>Apples</h4>
  <p>Apples are fruit.</p>
  <section>
    <h2>Taste</h2>
    <p>They taste lovely.</p>
    <h6>Sweet</h6>
    <p>Red apples are sweeter than green ones.</p>
    <h1>Color</h1>
    <p>Apples come in various colors.</p>
  </section>
</body>
```

However, the same document would be more clearly expressed as:

```
<body>
  <h1>Apples</h1>
  <p>Apples are fruit.</p>
  <section>
    <h2>Taste</h2>
    <p>They taste lovely.</p>
    <section>
      <h3>Sweet</h3>
      <p>Red apples are sweeter than green ones.</p>
    </section>
  </section>
  <section>
    <h2>Color</h2>
    <p>Apples come in various colors.</p>
  </section>
</body>
```

Both of the documents above are semantically identical and would produce the same outline in compliant user agents.

4.4.10.1 Creating an outline

This section defines an algorithm for creating an outline for a sectioning content (page 94) element or a sectioning root (page 140) element. It is defined in terms of a walk over the nodes of a DOM tree, in tree order, with each node being visited when it is *entered* and when it is *exited* during the walk.

The **outline** for a sectioning content (page 94) element or a sectioning root (page 140) element consists of a list of one or more potentially nested sections (page 142). A **section** is a container that corresponds to some nodes in the original DOM tree. Each section can have one heading associated with it, and can contain any number of further nested sections. The algorithm for the outline also associates each node in the DOM tree with a particular section and potentially a heading. (The sections in the outline aren't section elements, though some may correspond to such elements — they are merely conceptual sections.)

The following markup fragment:

```
<body>
  <h1>A</h1>
  <p>B</p>
  <h2>C</h2>
  <p>D</p>
  <h2>E</h2>
  <p>F</p>
</body>
```

...results in the following outline being created for the body node (and thus the entire document):

1. Section created for body node.
Associated with heading "A".
Also associated with paragraph "B".
Nested sections:
 1. Section implied for first h2 element.
Associated with heading "C".
Also associated with paragraph "D".
No nested sections.
 2. Section implied for second h2 element.
Associated with heading "E".
Also associated with paragraph "F".
No nested sections.

The algorithm that must be followed during a walk of a DOM subtree rooted at a sectioning content (page 94) element or a sectioning root (page 140) element to determine that element's outline (page 142) is as follows:

1. Let *current outlinee* be null. (It holds the element whose outline (page 142) is being created.)
2. Let *current section* be null. (It holds a pointer to a section (page 142), so that elements in the DOM can all be associated with a section.)
3. Create a stack to hold elements, which is used to handle nesting. Initialize this stack to empty.

4. As you walk over the DOM in tree order (page 24), trigger the first relevant step below for each element as you enter and exit it.

↪ **If the top of the stack is an element, and you are exiting that element**

Note: *The element being exited is a heading content (page 94) element.*

Pop that element from the stack.

↪ **If the top of the stack is a heading content (page 94) element**

Do nothing.

↪ **When entering a sectioning content (page 94) element or a sectioning root (page 140) element**

If *current outlinee* is not null, push *current outlinee* onto the stack.

Let *current outlinee* be the element that is being entered.

Let *current section* be a newly created section (page 142) for the *current outlinee* element.

Let there be a new outline (page 142) for the new *current outlinee*, initialized with just the new *current section* as the only section (page 142) in the outline.

↪ **When exiting a sectioning content (page 94) element, if the stack is not empty**

Pop the top element from the stack, and let the *current outlinee* be that element.

Let *current section* be the last section in the outline (page 142) of the *current outlinee* element.

Append the outline (page 142) of the sectioning content (page 94) element being exited to the *current section*. (This does not change which section is the last section in the outline (page 142).)

↪ **When exiting a sectioning root (page 140) element, if the stack is not empty**

Run these steps:

1. Pop the top element from the stack, and let the *current outlinee* be that element.
2. Let *current section* be the last section in the outline (page 142) of the *current outlinee* element.
3. *Finding the deepest child*: If *current section* has no child sections, stop these steps.
4. Let *current section* be the last child section (page 142) of the current *current section*.
5. Go back to the substep labeled *finding the deepest child*.

↪ When exiting a sectioning content (page 94) element or a sectioning root (page 140) element

Note: The current outlinee is the element being exited.

Let *current section* be the first section (page 142) in the outline (page 142) of the *current outlinee* element.

Skip to the next step in the overall set of steps. (The walk is over.)

↪ If the *current outlinee* is null.

Do nothing.

↪ When entering a heading content (page 94) element

If the *current section* has no heading, let the element being entered be the heading for the *current section*.

Otherwise, if the element being entered has a rank (page 136) equal to or greater than the heading of the last section of the outline (page 142) of the *current outlinee*, then create a new section (page 142) and append it to the outline (page 142) of the *current outlinee* element, so that this new section is the new last section of that outline. Let *current section* be that new section. Let the element being entered be the new heading for the *current section*.

Otherwise, run these substeps:

1. Let *candidate section* be *current section*.
2. If the element being entered has a rank (page 136) lower than the rank (page 136) of the heading of the *candidate section*, then create a new section (page 142), and append it to *candidate section*. (This does not change which section is the last section in the outline.) Let *current section* be this new section. Let the element being entered be the new heading for the *current section*. Abort these substeps.
3. Let *new candidate section* be the section (page 142) that contains *candidate section* in the outline (page 142) of *current outlinee*.
4. Let *candidate section* be *new candidate section*.
5. Return to step 2.

Push the element being entered onto the stack. (This causes the algorithm to skip any descendants of the element.)

Note: Recall that h1 has the highest rank, and h6 has the lowest rank.

↪ Otherwise

Do nothing.

In addition, whenever you exit a node, after doing the steps above, if *current section* is not null, associate the node with the section (page 142) *current section*.

5. If the *current outlinee* is null, then there was no sectioning content (page 94) element or sectioning root (page 140) element in the DOM. There is no outline (page 142). Abort these steps.
6. Associate any nodes that were not associated with a section (page 142) in the steps above with *current outlinee* as their section.
7. Associate all nodes with the heading of the section (page 142) with which they are associated, if any.
8. If *current outlinee* is the body element (page 81), then the outline created for that element is the outline (page 142) of the entire document.

The tree of sections created by the algorithm above, or a proper subset thereof, must be used when generating document outlines, for example when generating tables of contents.

When creating an interactive table of contents, entries should jump the user to the relevant sectioning content (page 94) element, if the section (page 142) was created for a real element in the original document, or to the relevant heading content (page 94) element, if the section (page 142) in the tree was generated for a heading in the above process.

Note: Selecting the first section (page 142) of the document therefore always takes the user to the top of the document, regardless of where the first header in the body is to be found.

The following JavaScript function shows how the tree walk could be implemented. The root argument is the root of the tree to walk, and the enter and exit arguments are callbacks that are called with the nodes as they are entered and exited. [ECMA262]

```
function (root, enter, exit) {
  var node = root;
  start: while (node) {
    enter(node);
    if (node.firstChild) {
      node = node.firstChild;
      continue start;
    }
    while (node) {
      exit(node);
      if (node.nextSibling) {
        node = node.nextSibling;
        continue start;
      }
      if (node == root)
        node = null;
      else
        node = node.parentNode;
    }
  }
}
```

4.4.10.2 Distinguishing site-wide headings from page headings

Given the outline (page 142) of a document, but ignoring any sections created for nav and aside elements, and any of their descendants, if the only root of the tree is the body element (page 81)'s section (page 142), and it has only a single subsection which is created by an article element, then the heading of the body element (page 81) should be assumed to be a site-wide heading, and the heading of the article element should be assumed to be the page's heading.

If a page starts with a heading that is common to the whole site, the document must be authored such that, in the document's outline (page 142), ignoring any sections created for nav and aside elements and any of their descendants, the tree has only one root section (page 142), the body element (page 81)'s section, its heading is the site-wide heading, the body element (page 81) has just one subsection, that subsection is created by an article element, and that article's heading is the page heading.

If a page does not contain a site-wide heading, then the page must be authored such that, in the document's outline (page 142), ignoring any sections created for nav and aside elements and any of their descendants, either the body element (page 81) has no subsections, or it has more than one subsection, or it has a single subsection but that subsection is not created by an article element, or there is more than one section (page 142) at the root of the outline.

Note: Conceptually, a site is thus a document with many articles — when those articles are split into many pages, the heading of the original single page becomes the heading of the site, repeated on every page.

4.5 Grouping content

4.5.1 The p element

Categories

Flow content (page 93).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The p element represents a paragraph (page 96).

The following examples are conforming HTML fragments:

```
<p>The little kitten gently seated himself on a piece of carpet. Later in his life, this would be referred to as the time the cat sat on the mat.</p>
```

```

<fieldset>
  <legend>Personal information</legend>
  <p>
    <label>Name: <input name="n"></label>
    <label><input name="anon" type="checkbox"> Hide from other
users</label>
  </p>
  <p><label>Address: <textarea name="a"></textarea></label></p>
</fieldset>
<p>There was once an example from Femley,<br>
Whose markup was of dubious quality.<br>
The validator complained,<br>
So the author was pained,<br>
To move the error from the markup to the rhyming.</p>

```

The `p` element should not be used when a more specific element is more appropriate.

The following example is technically correct:

```

<section>
  <!-- ... -->
  <p>Last modified: 2001-04-23</p>
  <p>Author: fred@example.com</p>
</section>

```

However, it would be better marked-up as:

```

<section>
  <!-- ... -->
  <footer>Last modified: 2001-04-23</footer>
  <address>Author: fred@example.com</address>
</section>

```

Or:

```

<section>
  <!-- ... -->
  <footer>
    <p>Last modified: 2001-04-23</p>
    <address>Author: fred@example.com</address>
  </footer>
</section>

```

4.5.2 The `hr` element

Categories

Flow content (page 93).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

Empty.

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The hr element represents a paragraph (page 96)-level thematic break, e.g. a scene change in a story, or a transition to another topic within a section of a reference book.

4.5.3 The br element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Empty.

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The br element represents a line break.

br elements must be empty. Any content inside br elements must not be considered part of the surrounding text.

br elements must be used only for line breaks that are actually part of the content, as in poems or addresses.

The following example is correct usage of the br element:

```
<p>P. Sherman<br>
42 Wallaby Way<br>
Sydney</p>
```

br elements must not be used for separating thematic groups in a paragraph.

The following examples are non-conforming, as they abuse the br element:

```
<p><a ...>34 comments.</a><br>
<a ...>Add a comment.<a></p>
<p>Name: <input name="name"><br>
Address: <input name="address"></p>
```

Here are alternatives to the above, which are correct:

```
<p><a ...>34 comments.</a></p>
<p><a ...>Add a comment.<a></p>
<p>Name: <input name="name"></p>
<p>Address: <input name="address"></p>
```

If a paragraph (page 96) consists of nothing but a single br element, it represents a placeholder blank line (e.g. as in a template). Such blank lines must not be used for presentation purposes.

4.5.4 The pre element

Categories

Flow content (page 93).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The pre element represents a block of preformatted text, in which structure is represented by typographic conventions rather than by elements.

Note: In the HTML serialization, a leading newline character immediately following the pre element start tag is stripped.

Some examples of cases where the pre element could be used:

- Including an e-mail, with paragraphs indicated by blank lines, lists indicated by lines prefixed with a bullet, and so on.
- Including fragments of computer code, with structure indicated according to the conventions of that language.
- Displaying ASCII art.

To represent a block of computer code, the pre element can be used with a code element; to represent a block of computer output the pre element can be used with a samp element. Similarly, the kbd element can be used within a pre element to indicate text that the user is to enter.

In the following snippet, a sample of computer code is presented.

```
<p>This is the <code>Panel</code> constructor:</p>
<pre><code>function Panel(element, canClose, closeHandler) {
    this.element = element;
```

```
this.canClose = canClose;
this.closeHandler = function () { if (closeHandler) closeHandler() };
};</code></pre>
```

In the following snippet, samp and kbd elements are mixed in the contents of a pre element to show a session of Zork I.

```
<pre><samp>You are in an open field west of a big white house with a
boarded
front door.
There is a small mailbox here.

></samp> <kbd>open mailbox</kbd>

<samp>Opening the mailbox reveals:
A leaflet.

></samp></pre>
```

The following shows a contemporary poem that uses the pre element to preserve its unusual formatting, which forms an intrinsic part of the poem itself.

```
<pre>
maxling

it is with a      heart
                  heavy

that i admit loss of a feline
                  so          loved

a friend lost to the
                  unknown
                           (night)

~cdr 11dec07</pre>
```

4.5.5 The dialog element

Categories

Flow content (page 93).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

Zero or more pairs of one dt element followed by one dd element.

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The dialog element represents a conversation.

Each part of the conversation must have an explicit talker (or speaker) given by a dt element, and a discourse (or quote) given by a dd element.

This example demonstrates this using an extract from Abbot and Costello's famous sketch, *Who's on first*:

```
<dialog>
  <dt> Costello
  <dd> Look, you gotta first baseman?
  <dt> Abbott
  <dd> Certainly.
  <dt> Costello
  <dd> Who's playing first?
  <dt> Abbott
  <dd> That's right.
  <dt> Costello
  <dd> When you pay off the first baseman every month, who gets the
      money?
  <dt> Abbott
  <dd> Every dollar of it.
</dialog>
```

Note: *Text in a dt element in a dialog element is implicitly the source of the text given in the following dd element, and the contents of the dd element are implicitly a quote from that speaker. There is thus no need to include cite, q, or blockquote elements in this markup. Indeed, a q element inside a dd element in a conversation would actually imply the people talking were themselves quoting another work. See the cite, q, and blockquote elements for other ways to cite or quote.*

4.5.6 The blockquote element

Categories

Flow content (page 93).
Sectioning root (page 140).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

Flow content (page 93).

Element-specific attributes:

cite

DOM interface:

```
interface HTMLQuoteElement : HTMLElement {
  attribute DOMString cite;
};
```

Note: The `HTMLQuoteElement` interface is also used by the `q` element.

The `blockquote` element represents a section that is quoted from another source.

Content inside a `blockquote` must be quoted from another source, whose address, if it has one, should be cited in the `cite` attribute.

If the `cite` attribute is present, it must be a valid URL (page 52). User agents should allow users to follow such citation links.

If a `blockquote` element is preceded or followed (page 92) by a single paragraph (page 96) that contains a single `cite` element and that is itself not preceded or followed (page 92) by another `blockquote` element and does not itself have a `q` element descendant, then, the title of the work given by that `cite` element gives the source of the quotation contained in the `blockquote` element.

The `cite` DOM attribute must reflect (page 67) the element's `cite` content attribute.

Note: The best way to represent a conversation is not with the `cite` and `blockquote` elements, but with the `dialog` element.

4.5.7 The `ol` element

Categories

Flow content (page 93).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

Zero or more `li` elements.

Element-specific attributes:

`reversed`
`start`

DOM interface:

```
interface HTMLListElement : HTMLElement {  
    attribute boolean reversed;  
    attribute long start;  
};
```

The `ol` element represents a list of items, where the items have been intentionally ordered, such that changing the order would change the meaning of the document.

The items of the list are the `li` element child nodes of the `ol` element, in tree order (page 24).

The **reversed** attribute is a boolean attribute (page 32). If present, it indicates that the list is a descending list (... , 3, 2, 1). If the attribute is omitted, the list is an ascending list (1, 2, 3, ...).

The **start** attribute, if present, must be a valid integer (page 33) giving the ordinal value of the first list item.

If the start attribute is present, user agents must parse it as an integer (page 33), in order to determine the attribute's value. The default value, used if the attribute is missing or if the value cannot be converted to a number according to the referenced algorithm, is 1 if the element has no reversed attribute, and is the number of child li elements otherwise.

The first item in the list has the ordinal value given by the ol element's start attribute, unless that li element has a value attribute with a value that can be successfully parsed, in which case it has the ordinal value given by that value attribute.

Each subsequent item in the list has the ordinal value given by its value attribute, if it has one, or, if it doesn't, the ordinal value of the previous item, plus one if the reversed is absent, or minus one if it is present.

The **reversed** DOM attribute must reflect (page 67) the value of the reversed content attribute.

The **start** DOM attribute must reflect (page 67) the value of the start content attribute.

The following markup shows a list where the order matters, and where the ol element is therefore appropriate. Compare this list to the equivalent list in the ul section to see an example of the same items using the ul element.

```
<p>I have lived in the following countries (given in the order of when  
I first lived there):</p>  
<ol>  
  <li>Switzerland  
  <li>United Kingdom  
  <li>United States  
  <li>Norway  
</ol>
```

Note how changing the order of the list changes the meaning of the document. In the following example, changing the relative order of the first two items has changed the birthplace of the author:

```
<p>I have lived in the following countries (given in the order of when  
I first lived there):</p>  
<ol>  
  <li>United Kingdom  
  <li>Switzerland  
  <li>United States  
  <li>Norway  
</ol>
```

4.5.8 The ul element

Categories

Flow content (page 93).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

Zero or more li elements.

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The ul element represents a list of items, where the order of the items is not important — that is, where changing the order would not materially change the meaning of the document.

The items of the list are the li element child nodes of the ul element.

The following markup shows a list where the order does not matter, and where the ul element is therefore appropriate. Compare this list to the equivalent list in the ol section to see an example of the same items using the ol element.

```
<p>I have lived in the following countries:</p>
<ul>
  <li>Norway
  <li>Switzerland
  <li>United Kingdom
  <li>United States
</ul>
```

Note that changing the order of the list does not change the meaning of the document. The items in the snippet above are given in alphabetical order, but in the snippet below they are given in order of the size of their current account balance in 2007, without changing the meaning of the document whatsoever:

```
<p>I have lived in the following countries:</p>
<ul>
  <li>Switzerland
  <li>Norway
  <li>United Kingdom
  <li>United States
</ul>
```

4.5.9 The li element

Categories

None.

Contexts in which this element may be used:

- Inside ol elements.
- Inside ul elements.
- Inside menu elements.

Content model:

When the element is a child of a menu element: phrasing content (page 94).
Otherwise: flow content (page 93).

Element-specific attributes:

- If the element is a child of an ol element: value
- If the element is not the child of an ol element: None.

DOM interface:

```
interface HTMLLIElement : HTMLElement {  
    attribute long value;  
};
```

The li element represents a list item. If its parent element is an ol, ul, or menu element, then the element is an item of the parent element's list, as defined for those elements. Otherwise, the list item has no defined list-related relationship to any other li element.

The **value** attribute, if present, must be a valid integer (page 33) giving the ordinal value of the list item.

If the value attribute is present, user agents must parse it as an integer (page 33), in order to determine the attribute's value. If the attribute's value cannot be converted to a number, the attribute must be treated as if it was absent. The attribute has no default value.

The value attribute is processed relative to the element's parent ol element (q.v.), if there is one. If there is not, the attribute has no effect.

The **value** DOM attribute must reflect (page 67) the value of the value content attribute.

The following example, the top ten movies are listed (in reverse order). Note the way the list is given a title by using a figure element and its legend.

```
<figure>  
  <legend>The top 10 movies of all time</legend>  
  <ol>  
    <li value="10"><cite>Josie and the Pussycats</cite>, 2001</li>  
    <li value="9"><cite lang="sh">Црна мачка, бели мачор</cite>, 1998</li>  
    <li value="8"><cite>A Bug's Life</cite>, 1998</li>  
    <li value="7"><cite>Toy Story</cite>, 1995</li>  
    <li value="6"><cite>Monsters, Inc</cite>, 2001</li>  
    <li value="5"><cite>Cars</cite>, 2006</li>  
    <li value="4"><cite>Toy Story 2</cite>, 1999</li>  
    <li value="3"><cite>Finding Nemo</cite>, 2003</li>  
    <li value="2"><cite>The Incredibles</cite>, 2004</li>  
    <li value="1"><cite>Ratatouille</cite>, 2007</li>
```

```
</ol>
</figure>
```

The markup could also be written as follows, using the reversed attribute on the ol element:

```
<figure>
  <legend>The top 10 movies of all time</legend>
  <ol reversed>
    <li><cite>Josie and the Pussycats</cite>, 2001</li>
    <li><cite lang="sh">Црна мачка, бели мачор</cite>, 1998</li>
    <li><cite>A Bug's Life</cite>, 1998</li>
    <li><cite>Toy Story</cite>, 1995</li>
    <li><cite>Monsters, Inc</cite>, 2001</li>
    <li><cite>Cars</cite>, 2006</li>
    <li><cite>Toy Story 2</cite>, 1999</li>
    <li><cite>Finding Nemo</cite>, 2003</li>
    <li><cite>The Incredibles</cite>, 2004</li>
    <li><cite>Ratatouille</cite>, 2007</li>
  </ol>
</figure>
```

If the li element is the child of a menu element and itself has a child that defines a command (page 406), then the li element must match the :enabled and :disabled pseudo-classes in the same way as the first such child element does.

4.5.10 The dl element

Categories

Flow content (page 93).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

Zero or more groups each consisting of one or more dt elements followed by one or more dd elements.

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The dl element introduces an association list consisting of zero or more name-value groups (a description list). Each group must consist of one or more names (dt elements) followed by one or more values (dd elements).

Name-value groups may be terms and definitions, metadata topics and values, or any other groups of name-value data.

The values within a group are alternatives; multiple paragraphs forming part of the same value must all be given within the same dd element.

The order of the list of groups, and of the names and values within each group, may be significant.

If a `dl` element is empty, it contains no groups.

If a `dl` element contains non-whitespace (page 92) text nodes (page 24), or elements other than `dt` and `dd`, then those elements or text nodes (page 24) do not form part of any groups in that `dl`.

If a `dl` element contains only `dt` elements, then it consists of one group with names but no values.

If a `dl` element contains only `dd` elements, then it consists of one group with values but no names.

If a `dl` element starts with one or more `dd` elements, then the first group has no associated name.

If a `dl` element ends with one or more `dt` elements, then the last group has no associated value.

Note: When a `dl` element doesn't match its content model, it is often due to accidentally using `dd` elements in the place of `dt` elements and vice versa. Conformance checkers can spot such mistakes and might be able to advise authors how to correctly use the markup.

In the following example, one entry ("Authors") is linked to two values ("John" and "Luke").

```
<dl>
  <dt> Authors
  <dd> John
  <dd> Luke
  <dt> Editor
  <dd> Frank
</dl>
```

In the following example, one definition is linked to two terms.

```
<dl>
  <dt lang="en-US"> <dfn>color</dfn> </dt>
  <dt lang="en-GB"> <dfn>colour</dfn> </dt>
  <dd> A sensation which (in humans) derives from the ability of
       the fine structure of the eye to distinguish three differently
       filtered analyses of a view. </dd>
</dl>
```

The following example illustrates the use of the `dl` element to mark up metadata of sorts. At the end of the example, one group has two metadata labels ("Authors" and "Editors") and two values ("Robert Rothman" and "Daniel Jackson").

```
<dl>
  <dt> Last modified time </dt>
  <dd> 2004-12-23T23:33Z </dd>
  <dt> Recommended update interval </dt>
```

```

<dd> 60s </dd>
<dt> Authors </dt>
<dt> Editors </dt>
<dd> Robert Rothman </dd>
<dd> Daniel Jackson </dd>
</dl>

```

The following example shows the dl element used to give a set of instructions. The order of the instructions here is important (in the other examples, the order of the blocks was not important).

```

<p>Determine the victory points as follows (use the
first matching case):</p>
<dl>
  <dt> If you have exactly five gold coins </dt>
  <dd> You get five victory points </dd>
  <dt> If you have one or more gold coins, and you have one or more
silver coins </dt>
  <dd> You get two victory points </dd>
  <dt> If you have one or more silver coins </dt>
  <dd> You get one victory point </dd>
  <dt> Otherwise </dt>
  <dd> You get no victory points </dd>
</dl>

```

The following snippet shows a dl element being used as a glossary. Note the use of dfn to indicate the word being defined.

```

<dl>
  <dt><dfn>Apartment</dfn>, n.</dt>
  <dd>An execution context grouping one or more threads with one or
more COM objects.</dd>
  <dt><dfn>Flat</dfn>, n.</dt>
  <dd>A deflated tire.</dd>
  <dt><dfn>Home</dfn>, n.</dt>
  <dd>The user's login directory.</dd>
</dl>

```

Note: The *dl* element is inappropriate for marking up dialogue, since dialogue is ordered (each speaker/line pair comes after the next). For an example of how to mark up dialogue, see the *dialog* element.

4.5.11 The dt element

Categories

None.

Contexts in which this element may be used:

Before dd or dt elements inside dl elements.
Before a dd element inside a dialog element.

Content model:

Phrasing content (page 94).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The dt element represents the term, or name, part of a term-description group in a description list (dl element), and the talker, or speaker, part of a talker-discourse pair in a conversation (dialog element).

Note: *The dt element itself, when used in a dl element, does not indicate that its contents are a term being defined, but this can be indicated using the dfn element.*

If the dt element is the child of a dialog element, and it further contains a time element, then that time element represents a timestamp for when the associated discourse (dd element) was said, and is not part of the name of the talker.

The following extract shows how an IM conversation log could be marked up.

```
<dialog>
  <dt> <time>14:22</time> egof
  <dd> I'm not that nerdy, I've only seen 30% of the star trek episodes
  <dt> <time>14:23</time> kaj
  <dd> if you know what percentage of the star trek episodes you have
  seen, you are inarguably nerdy
  <dt> <time>14:23</time> egof
  <dd> it's unarguably
  <dt> <time>14:24</time> kaj
  <dd> you are not helping your case
</dialog>
```

4.5.12 The dd element

Categories

None.

Contexts in which this element may be used:

After dt or dd elements inside dl elements.

After a dt element inside a dialog element.

Content model:

Flow content (page 93).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The dd element represents the description, definition, or value, part of a term-description group in a description list (dl element), and the discourse, or quote, part in a conversation (dialog element).

4.6 Text-level semantics

4.6.1 The a element

Categories

Interactive content (page 95).

When the element only contains phrasing content (page 94): phrasing content (page 94).

Otherwise: flow content (page 93).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Transparent (page 96), but there must be no interactive content (page 95) descendant.

Element-specific attributes:

href
target
ping
rel
media
hreflang
type

DOM interface:

```
[Stringifies=href] interface HTMLAnchorElement : HTMLElement {  
    attribute DOMString href;  
    attribute DOMString target;  
    attribute DOMString ping;  
    attribute DOMString rel;  
    readonly attribute DOMTokenList relList;  
    attribute DOMString media;  
    attribute DOMString hreflang;  
    attribute DOMString type;  
};
```

The Command interface must also be implemented by this element.

If the a element has an href attribute, then it represents a hyperlink (page 497).

If the `a` element has no `href` attribute, then the element is a placeholder for where a link might otherwise have been placed, if it had been relevant.

The `target`, `ping`, `rel`, `media`, `hreflang`, and `type` attributes must be omitted if the `href` attribute is not present.

If a site uses a consistent navigation toolbar on every page, then the link that would normally link to the page itself could be marked up using an `a` element:

```
<nav>
  <ul>
    <li> <a href="/">Home</a> </li>
    <li> <a href="/news">News</a> </li>
    <li> <a>Examples</a> </li>
    <li> <a href="/legal">Legal</a> </li>
  </ul>
</nav>
```

Interactive user agents should allow users to follow hyperlinks (page 498) created using the `a` element. The `href`, `target` and `ping` attributes decide how the link is followed. The `rel`, `media`, `hreflang`, and `type` attributes may be used to indicate to the user the likely nature of the target resource before the user follows the link.

The activation behavior (page 96) of `a` elements that represent hyperlinks is to run the following steps:

1. If the `DOMActivate` event in question is not trusted (i.e. a `click()` method call was the reason for the event being dispatched), and the `a` element's `target` attribute is ****** ... then raise an `INVALID_ACCESS_ERR` exception and abort these steps.
2. If the target of the `click` event is an `img` element with an `ismap` attribute specified, then server-side image map processing must be performed, as follows:
 1. If the `DOMActivate` event was dispatched as the result of a real pointing-device-triggered `click` event on the `img` element, then let `x` be the distance in CSS pixels from the left edge of the image to the location of the click, and let `y` be the distance in CSS pixels from the top edge of the image to the location of the click. Otherwise, let `x` and `y` be zero.
 2. Let the ***hyperlink suffix*** be a U+003F QUESTION MARK character, the value of `x` expressed as a base-ten integer using ASCII digits (U+0030 DIGIT ZERO to U+0039 DIGIT NINE), a U+002C COMMA character, and the value of `y` expressed as a base-ten integer using ASCII digits.
 3. Finally, the user agent must follow the hyperlink (page 498) defined by the `a` element. If the steps above defined a *hyperlink suffix*, then take that into account when following the hyperlink.

The DOM attributes `href`, `ping`, `target`, `rel`, `media`, `hreflang`, and `type`, must reflect (page 67) the respective content attributes of the same name.

The DOM attribute `relList` must reflect (page 67) the `rel` content attribute.

The `a` element may be wrapped around entire paragraphs, lists, tables, and so forth, even entire sections, so long as there is no interactive content within (e.g. buttons or other links). This example shows how this can be used to make an entire advertising block into a link:

```
<aside class="advertising">
  <h1>Advertising</h1>
  <a href="http://ad.example.com/?adid=1929&pubid=1422">
    <section>
      <h1>Mellblomatic 9000!</h1>
      <p>Turn all your widgets into mellbloms!</p>
      <p>Only $9.99 plus shipping and handling.</p>
    </section>
  </a>
  <a href="http://ad.example.com/?adid=375&pubid=1422">
    <section>
      <h1>The Mellblom Browser</h1>
      <p>Web browsing at the speed of light.</p>
      <p>No other browser goes faster!</p>
    </section>
  </a>
</aside>
```

4.6.2 The `q` element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

`cite`

DOM interface:

The `q` element uses the `HTMLQuoteElement` interface.

The `q` element represents some phrasing content (page 94) quoted from another source.

Quotation punctuation (such as quotation marks), if any, must be placed inside the `q` element.

Content inside a `q` element must be quoted from another source, whose address, if it has one, should be cited in the `cite` attribute.

If the `cite` attribute is present, it must be a valid URL (page 52). User agents should allow users to follow such citation links.

If a q element is contained (directly or indirectly) in a paragraph (page 96) that contains a single cite element and has no other q element descendants, then, the title of the work given by that cite element gives the source of the quotation contained in the q element.

Here is a simple example of the use of the q element:

```
<p>The man said <q>"Things that are impossible just take longer"</q>. I disagreed with him.</p>
```

Here is an example with both an explicit citation link in the q element, and an explicit citation outside:

```
<p>The W3C page <cite>About W3C</cite> says the W3C's mission is <q cite="http://www.w3.org/Consortium/">"To lead the World Wide Web to its full potential by developing protocols and guidelines that ensure long-term growth for the Web"</q>. I disagree with this mission.</p>
```

In the following example, the quotation itself contains a quotation:

```
<p>In <cite>Example One</cite>, he writes <q>"The man said <q>'Things that are impossible just take longer'</q>. I disagreed with him"</q>. Well, I disagree even more!</p>
```

In the following example, there are no quotation marks:

```
<p>His best argument: <q>I disagree!</q></p>
```

4.6.3 The cite element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The cite element represents the title of a work (e.g. a book, a paper, an essay, a poem, a score, a song, a script, a film, a TV show, a game, a sculpture, a painting, a theatre production, a play, an opera, a musical, an exhibition, etc). This can be a work that is being quoted or referenced in detail (i.e. a citation), or it can just be a work that is mentioned in passing.

A person's name is not the title of a work — even if people call that person a piece of work — and the element must therefore not be used to mark up people's names. (In some cases, the b element might be appropriate for names; e.g. in a gossip article where the names of

famous people are keywords rendered with a different style to draw attention to them. In other cases, if an element is *really* needed, the `span` element can be used.)

A ship is similarly not a work, and the `element` must not be used to mark up ship names (the `i` element can be used for that purpose).

This next example shows a typical use of the `cite` element:

```
<p>My favourite book is <code><cite>The Reality Dysfunction</cite></code> by Peter F. Hamilton. My favourite comic is <code><cite>Pearls Before Swine</cite></code> by Stephan Pastis. My favourite track is <code><cite>Jive Samba</cite></code> by the Cannonball Adderley Sextet.</p>
```

This is correct usage:

```
<p>According to the Wikipedia article <code><cite>HTML</cite></code>, as it stood in mid-February 2008, leaving attribute values unquoted is unsafe. This is obviously an over-simplification.</p>
```

The following, however, is incorrect usage, as the `cite` element here is containing far more than the title of the work:

```
<!-- do not copy this example, it is an example of bad usage! -->
<p>According to <code><cite>the Wikipedia article on HTML</cite></code>, as it stood in mid-February 2008, leaving attribute values unquoted is unsafe. This is obviously an over-simplification.</p>
```

The `cite` element is obviously a key part of any citation in a bibliography, but it is only used to mark the title:

```
<p><code><cite>Universal Declaration of Human Rights</cite></code>, United Nations, December 1948. Adopted by General Assembly resolution 217 A (III).</p>
```

Note: A citation is not a quote (for which the `q` element is appropriate).

This is incorrect usage, because `cite` is not for quotes:

```
<p><code><cite>This is wrong!</cite></code>, said Ian.</p>
```

This is also incorrect usage, because a person is not a work:

```
<p><code><q>This is still wrong!</q>, said <code><cite>Ian</cite></code>.</p>
```

The correct usage does not use a `cite` element:

```
<p><code><q>This is correct</q>, said Ian.</code></p>
```

As mentioned above, the `b` element might be relevant for marking names as being keywords in certain kinds of documents:

```
<p>And then <code><b>Ian</b></code> said <code><q>this might be right, in a gossip column, maybe!</q></code>.</p>
```

Note: The `cite` element can apply to `blockquote` and `q` elements in certain cases described in the definitions of those elements.

This next example shows the use of cite alongside blockquote:

```
<p>His next piece was the aptly named <cite>Sonnet 130</cite>:</p>
<blockquote>
  <p>My mistress' eyes are nothing like the sun,<br>
  Coral is far more red, than her lips red,
  ...

```

4.6.4 The em element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The em element represents stress emphasis of its contents.

The level of emphasis that a particular piece of content has is given by its number of ancestor em elements.

The placement of emphasis changes the meaning of the sentence. The element thus forms an integral part of the content. The precise way in which emphasis is used in this way depends on the language.

These examples show how changing the emphasis changes the meaning. First, a general statement of fact, with no emphasis:

```
<p>Cats are cute animals.</p>
```

By emphasizing the first word, the statement implies that the kind of animal under discussion is in question (maybe someone is asserting that dogs are cute):

```
<p><em>Cats</em> are cute animals.</p>
```

Moving the emphasis to the verb, one highlights that the truth of the entire sentence is in question (maybe someone is saying cats are not cute):

```
<p>Cats <em>are</em> cute animals.</p>
```

By moving it to the adjective, the exact nature of the cats is reasserted (maybe someone suggested cats were *mean* animals):

```
<p>Cats are <em>cute</em> animals.</p>
```

Similarly, if someone asserted that cats were vegetables, someone correcting this might emphasize the last word:

```
<p>Cats are cute <em>animals</em>.</p>
```

By emphasizing the entire sentence, it becomes clear that the speaker is fighting hard to get the point across. This kind of emphasis also typically affects the punctuation, hence the exclamation mark here.

```
<p><em>Cats are cute animals!</em></p>
```

Anger mixed with emphasizing the cuteness could lead to markup such as:

```
<p><em>Cats are <em>cute</em> animals!</em></p>
```

4.6.5 The strong element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The strong element represents strong importance for its contents.

The relative level of importance of a piece of content is given by its number of ancestor strong elements; each strong element increases the importance of its contents.

Changing the importance of a piece of text with the strong element does not change the meaning of the sentence.

Here is an example of a warning notice in a game, with the various parts marked up according to how important they are:

```
<p><strong>Warning.</strong> This dungeon is dangerous.  
<strong>Avoid the ducks.</strong> Take any gold you find.  
<strong><strong>Do not take any of the diamonds</strong>,</strong>  
they are explosive and <strong>will destroy anything within  
ten meters.</strong></strong> You have been warned.</p>
```

4.6.6 The small element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The small element represents small print (part of a document often describing legal restrictions, such as copyrights or other disadvantages), or other side comments.

Note: The small element does not "de-emphasize" or lower the importance of text emphasised by the em element or marked as important with the strong element.

In this example the footer contains contact information and a copyright.

```
<footer>
  <address>
    For more details, contact
    <a href="mailto:js@example.com">John Smith</a>.
  </address>
  <p><small>© copyright 2038 Example Corp.</small></p>
</footer>
```

In this second example, the small element is used for a side comment.

```
<p>Example Corp today announced record profits for the
second quarter <small>(Full Disclosure: Foo News is a subsidiary of
Example Corp)</small>, leading to speculation about a third quarter
merger with Demo Group.</p>
```

In this last example, the small element is marked as being *important* small print.

```
<p><strong><small>Continued use of this service will result in a
kiss.</small></strong></p>
```

4.6.7 The mark element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The mark element represents a run of text in one document marked or highlighted for reference purposes, due to its relevance in another context. When used in a quotation or other block of text referred to from the prose, it indicates a highlight that was not originally present but which has been added to bring the reader's attention to a part of the text that might not have been considered important by the original author when the block was originally written, but which is now under previously unexpected scrutiny. When used in the main prose of a document, it indicates a part of the document that has been highlighted due to its likely relevance to the user's current activity.

- ** The rendering section will eventually suggest that user agents provide a way to let users jump between mark elements. Suggested rendering is a neon yellow background highlight, though UAs maybe should allow this to be toggled.

This example shows how the mark example can be used to bring attention to a particular part of a quotation:

```
<p lang="en-US">Consider the following quote:</p>
<blockquote lang="en-GB">
  <p>Look around and you will find, no-one's really
    <mark>colour</mark> blind.</p>
</blockquote>
<p lang="en-US">As we can tell from the <em>spelling</em> of the word,
  the person writing this quote is clearly not American.</p>
```

Another example of the mark element is highlighting parts of a document that are matching some search string. If someone looked at a document, and the server knew that the user was searching for the word "kitten", then the server might return the document with one paragraph modified as follows:

```
<p>I also have some <mark>kitten</mark>s who are visiting me
these days. They're really cute. I think they like my garden! Maybe I
should adopt a <mark>kitten</mark>.</p>
```

In the following snippet, a paragraph of text refers to a specific part of a code fragment.

```
<p>The highlighted part below is where the error lies:</p>
<pre><code>var i: Integer;
begin
  i := <mark>1.1</mark>;
end.</code></pre>
```

This is another example showing the use of mark to highlight a part of quoted text that was originally not emphasised. In this example, common typographic conventions have led the author to explicitly style mark elements in quotes to render in italics.

```
<article>
  <style>
    blockquote mark, q mark {
      font: inherit; font-style: italic;
      text-decoration: none;
      background: transparent; color: inherit;
    }
    .bubble em {
      font: inherit; font-size: larger;
      text-decoration: underline;
    }
  </style>
  <h1>She knew</h1>
  <p>Did you notice the subtle joke in the joke on panel 4?</p>
  <blockquote>
    <p class="bubble">I didn't <em>want</em> to believe. <mark>Of course on some level I realized it was a known-plaintext attack.</mark> But I
      couldn't admit it until I saw for myself.</p>
  </blockquote>
  <p>(Emphasis mine.) I thought that was great. It's so pedantic, yet it explains everything neatly.</p>
</article>
```

Note, incidentally, the distinction between the em element in this example, which is part of the original text being quoted, and the mark element, which is highlighting a part for comment.

The following example shows the difference between denoting the *importance* of a span of text (strong) as opposed to denoting the *relevance* of a span of text (mark). It is an extract from a textbook, where the extract has had the parts relevant to the exam highlighted. The safety warnings, important though they may be, are apparently not relevant to the exam.

```
<h3>Wormhole Physics Introduction</h3>

<p><mark>A wormhole in normal conditions can be held open for a maximum of just under 39 minutes.</mark> Conditions that can increase the time include a powerful energy source coupled to one or both of the gates connecting the wormhole, and a large gravity well (such as a black hole).</p>

<p><mark>Momentum is preserved across the wormhole. Electromagnetic radiation can travel in both directions through a wormhole, but matter cannot.</mark></p>

<p>When a wormhole is created, a vortex normally forms. <strong>Warning: The vortex caused by the wormhole opening will annihilate anything in its path.</strong> Vortexes can be avoided when
```

```
    using sufficiently advanced dialing technology.</p>
<p><mark>An obstruction in a gate will prevent it from accepting a wormhole connection.</mark></p>
```

4.6.8 The dfn element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94), but there must be no descendant dfn elements.

Element-specific attributes:

None, but the title attribute has special semantics on this element.

DOM interface:

Uses HTMLElement.

The dfn element represents the defining instance of a term. The paragraph (page 96), description list group (page 156), or section (page 94) that is the nearest ancestor of the dfn element must also contain the definition(s) for the term (page 170) given by the dfn element.

Defining term: If the dfn element has a **title** attribute, then the exact value of that attribute is the term being defined. Otherwise, if it contains exactly one element child node and no child text nodes (page 24), and that child element is an abbr element with a title attribute, then the exact value of *that* attribute is the term being defined. Otherwise, it is the exact textContent of the dfn element that gives the term being defined.

If the title attribute of the dfn element is present, then it must contain only the term being defined.

Note: The title attribute of ancestor elements does not affect dfn elements.

An a element that links to a dfn element represents an instance of the term defined by the dfn element.

In the following fragment, the term "GDO" is first defined in the first paragraph, then used in the second.

```
<p>The <dfn><abbr title="Garage Door Opener">GDO</abbr></dfn>
is a device that allows off-world teams to open the iris.</p>
<!-- ... later in the document: -->
<p>Teal'c activated his <abbr title="Garage Door Opener">GDO</abbr>
and so Hammond ordered the iris to be opened.</p>
```

With the addition of an a element, the reference can be made explicit:

```
<p>The <dfn id=gdo><abbr title="Garage Door Opener">GDO</abbr></dfn>  
is a device that allows off-world teams to open the iris.</p>  
<!-- ... later in the document: -->  
<p>Teal'c activated his <a href=#gdo><abbr title="Garage Door  
Opener">GDO</abbr></a>  
and so Hammond ordered the iris to be opened.</p>
```

4.6.9 The abbr element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

None, but the `title` attribute has special semantics on this element.

DOM interface:

Uses `HTMLElement`.

The `abbr` element represents an abbreviation or acronym, optionally with its expansion. The `title` attribute may be used to provide an expansion of the abbreviation. The attribute, if specified, must contain an expansion of the abbreviation, and nothing else.

The paragraph below contains an abbreviation marked up with the `abbr` element. This paragraph defines the term (page 170) "Web Hypertext Application Technology Working Group".

```
<p>The <dfn id=whatwg><abbr title="Web Hypertext Application  
Technology Working Group">WHATWG</abbr></dfn> is a loose  
unofficial collaboration of Web browser manufacturers and interested  
parties who wish to develop new technologies designed to allow authors  
to write and deploy Applications over the World Wide Web.</p>
```

This paragraph has two abbreviations. Notice how only one is defined; the other, with no expansion associated with it, does not use the `abbr` element.

```
<p>The <abbr title="Web Hypertext Application Technology Working  
Group">WHATWG</abbr> started working on HTML5 in 2004.</p>
```

This paragraph links an abbreviation to its definition.

```
<p>The <a href="#whatwg"><abbr title="Web Hypertext Application  
Technology Working Group">WHATWG</abbr></a> community does not  
have much representation from Asia.</p>
```

This paragraph marks up an abbreviation without giving an expansion, possibly as a hook to apply styles for abbreviations (e.g. `smallcaps`).

<p>Philip` and Dashiva both denied that they were going to get the issue counts from past revisions of the specification to backfill the <abbr>WHATWG</abbr> issue graph.</p>

If an abbreviation is pluralized, the expansion's grammatical number (plural vs singular) must match the grammatical number of the contents of the element.

Here the plural is outside the element, so the expansion is in the singular:

<p>Two <abbr title="Working Group">WG</abbr>s worked on this specification: the <abbr>WHATWG</abbr> and the <abbr>HTMLWG</abbr>.</p>

Here the plural is inside the element, so the expansion is in the plural:

<p>Two <abbr title="Working Groups">WGs</abbr> worked on this specification: the <abbr>WHATWG</abbr> and the <abbr>HTMLWG</abbr>.</p>

4.6.10 The time element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

datetime

DOM interface:

```
interface HTMLTimeElement : HTMLElement {  
    attribute DOMString dateTime;  
    readonly attribute DOMTimeStamp date;  
    readonly attribute DOMTimeStamp time;  
    readonly attribute DOMTimeStamp timezone;  
};
```

The time element represents a date and/or a time.

The **datetime** attribute, if present, must contain a date or time string (page 44) that identifies the date or time being specified.

If the **datetime** attribute is not present, then the date or time must be specified in the content of the element, such that parsing the element's **textContent** according to the rules for parsing date or time strings in content (page 44) successfully extracts a date or time.

The **dateTime** DOM attribute must reflect (page 67) the **datetime** content attribute.

User agents, to obtain the **date**, **time**, and **timezone** represented by a time element, must follow these steps:

1. If the `datetime` attribute is present, then parse it according to the rules for parsing date or time strings in attributes (page 44), and let the result be *result*.
2. Otherwise, parse the element's `textContent` according to the rules for parsing date or time strings in content (page 44), and let the result be *result*.
3. If *result* is empty (because the parsing failed), then the date (page 173) is unknown, the time (page 173) is unknown, and the timezone (page 173) is unknown.
4. Otherwise: if *result* contains a date, then that is the date (page 173); if *result* contains a time, then that is the time (page 173); and if *result* contains a timezone, then the timezone is the element's timezone (page 173). (A timezone can only be present if both a date and a time are also present.)

The **date** DOM attribute must return null if the date (page 173) is unknown, and otherwise must return the time corresponding to midnight UTC (i.e. the first second) of the given date (page 173).

The **time** DOM attribute must return null if the time (page 173) is unknown, and otherwise must return the time corresponding to the given time (page 173) of 1970-01-01, with the timezone UTC.

The **timezone** DOM attribute must return null if the timezone (page 173) is unknown, and otherwise must return the time corresponding to 1970-01-01 00:00 UTC in the given timezone (page 173), with the timezone set to UTC (i.e. the time corresponding to 1970-01-01 at 00:00 UTC plus the offset corresponding to the timezone).

In the following snippet:

```
<p>Our first date was <time datetime="2006-09-23">a  
Saturday</time>.</p>
```

...the time element's date attribute would have the value 1,158,969,600,000ms, and the time and timezone attributes would return null.

In the following snippet:

```
<p>We stopped talking at <time datetime="2006-09-24 05:00 -7">5am the  
next morning</time>.</p>
```

...the time element's date attribute would have the value 1,159,056,000,000ms, the time attribute would have the value 18,000,000ms, and the timezone attribute would return -25,200,000ms. To obtain the actual time, the three attributes can be added together, obtaining 1,159,048,800,000, which is the specified date and time in UTC.

Finally, in the following snippet:

```
<p>Many people get up at <time>08:00</time>.</p>
```

...the time element's date attribute would have the value null, the time attribute would have the value 28,800,000ms, and the timezone attribute would return null.

- ** These APIs may be suboptimal. Comments on making them more useful to JS authors are welcome. The primary use cases for these elements are for marking up publication dates e.g. in blog entries, and for marking event dates in hCalendar markup. Thus the DOM APIs are likely to be used as ways to generate interactive calendar widgets or some such.

4.6.11 The progress element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

`value`
`max`

DOM interface:

```
interface HTMLProgressElement : HTMLElement {
    attribute float value;
    attribute float max;
    readonly attribute float position;
};
```

The progress element represents the completion progress of a task. The progress is either indeterminate, indicating that progress is being made but that it is not clear how much more work remains to be done before the task is complete (e.g. because the task is waiting for a remote host to respond), or the progress is a number in the range zero to a maximum, giving the fraction of work that has so far been completed.

There are two attributes that determine the current task completion represented by the element.

The **value** attribute specifies how much of the task has been completed, and the **max** attribute specifies how much work the task requires in total. The units are arbitrary and not specified.

Instead of using the attributes, authors are recommended to include the current value and the maximum value inline as text inside the element.

Here is a snippet of a Web application that shows the progress of some automated task:

```
<section>
  <h2>Task Progress</h2>
  <p>Progress: <progress><span id="p">0</span>%</progress></p>
  <script>
    var progressBar = document.getElementById('p');
    function updateProgress(newValue) {
      progressBar.textContent = newValue;
```

```
    }
  </script>
</section>
```

(The updateProgress() method in this example would be called by some other code on the page to update the actual progress bar as the task progressed.)

Author requirements: The max and value attributes, when present, must have values that are valid floating point numbers (page 34). The max attribute, if present, must have a value greater than zero. The value attribute, if present, must have a value equal to or greater than zero, and less than or equal to the value of the max attribute, if present.

Note: *The progress element is the wrong element to use for something that is just a gauge, as opposed to task progress. For instance, indicating disk space usage using progress would be inappropriate. Instead, the meter element is available for such use cases.*

User agent requirements: User agents must parse the max and value attributes' values according to the rules for parsing floating point number values (page 34).

If the value attribute is omitted, then user agents must also parse the textContent of the progress element in question using the steps for finding one or two numbers of a ratio in a string (page 36). These steps will return nothing, one number, one number with a denominator punctuation character, or two numbers.

Using the results of this processing, user agents must determine whether the progress bar is an indeterminate progress bar, or whether it is a determinate progress bar, and in the latter case, what its current and maximum values are, all as follows:

1. If the max attribute is omitted, and the value is omitted, and the results of parsing the textContent was nothing, then the progress bar is an indeterminate progress bar. Abort these steps.
2. Otherwise, it is a determinate progress bar.
3. If the max attribute is included, then, if a value could be parsed out of it, then the maximum value is that value.
4. Otherwise, if the max attribute is absent but the value attribute is present, or, if the max attribute is present but no value could be parsed from it, then the maximum is 1.
5. Otherwise, if neither attribute is included, then, if the textContent contained one number with an associated denominator punctuation character, then the maximum value is the value associated with that denominator punctuation character; otherwise, if the textContent contained two numbers, the maximum value is the higher of the two values; otherwise, the maximum value is 1.
6. If the value attribute is present on the element and a value could be parsed out of it, that value is the current value of the progress bar. Otherwise, if the attribute is present but no value could be parsed from it, the current value is zero.
7. Otherwise if the value attribute is absent and the max attribute is present, then, if the textContent was parsed and found to contain just one number, with no associated denominator punctuation character, then the current value is that number.

Otherwise, if the value attribute is absent and the max attribute is present then the current value is zero.

8. Otherwise, if neither attribute is present, then the current value is the lower of the one or two numbers that were found in the textContent of the element.
9. If the maximum value is less than or equal to zero, then it is reset to 1.
10. If the current value is less than zero, then it is reset to zero.
11. Finally, if the current value is greater than the maximum value, then the current value is reset to the maximum value.

UA requirements for showing the progress bar: When representing a progress element to the user, the UA should indicate whether it is a determinate or indeterminate progress bar, and in the former case, should indicate the relative position of the current value relative to the maximum value.

The **max** and **value** DOM attributes must reflect (page 67) the respective content attributes of the same name. When the relevant content attributes are absent, the DOM attributes must return zero. The value parsed from the textContent never affects the DOM values.

** Would be cool to have the value DOM attribute update the textContent in-line...

If the progress bar is an indeterminate progress bar, then the **position** DOM attribute must return -1. Otherwise, it must return the result of dividing the current value by the maximum value.

4.6.12 The meter element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

value
min
low
high
max
optimum

DOM interface:

```
interface HTMLMeterElement : HTMLElement {  
    attribute float value;  
    attribute float min;
```

```
attribute float max;  
attribute float low;  
attribute float high;  
attribute float optimum;  
};
```

The meter element represents a scalar measurement within a known range, or a fractional value; for example disk usage, the relevance of a query result, or the fraction of a voting population to have selected a particular candidate.

This is also known as a gauge.

Note: *The meter element should not be used to indicate progress (as in a progress bar). For that role, HTML provides a separate progress element.*

Note: *The meter element also does not represent a scalar value of arbitrary range — for example, it would be wrong to use this to report a weight, or height, unless there is a known maximum value.*

There are six attributes that determine the semantics of the gauge represented by the element.

The **min** attribute specifies the lower bound of the range, and the **max** attribute specifies the upper bound. The **value** attribute specifies the value to have the gauge indicate as the "measured" value.

The other three attributes can be used to segment the gauge's range into "low", "medium", and "high" parts, and to indicate which part of the gauge is the "optimum" part. The **low** attribute specifies the range that is considered to be the "low" part, and the **high** attribute specifies the range that is considered to be the "high" part. The **optimum** attribute gives the position that is "optimum"; if that is higher than the "high" value then this indicates that the higher the value, the better; if it's lower than the "low" mark then it indicates that lower values are better, and naturally if it is in between then it indicates that neither high nor low values are good.

Authoring requirements: The recommended way of giving the value is to include it as contents of the element, either as two numbers (the higher number represents the maximum, the other number the current value, and the minimum is assumed to be zero), or as a percentage or similar (using one of the characters such as "%"), or as a fraction.

The value, min, low, high, max, and optimum attributes are all optional. When present, they must have values that are valid floating point numbers (page 34), and their values must satisfy the following inequalities:

- $\text{min} \leq \text{value} \leq \text{max}$
- $\text{min} \leq \text{low} \leq \text{high} \leq \text{max}$
- $\text{min} \leq \text{optimum} \leq \text{max}$

All meter elements must have a value specified somehow, either using the value attribute or by including a number in the contents of the element.

Note: If no minimum or maximum is specified, then the range is assumed to be 0..1, and the value thus has to be within that range.

The following examples all represent a measurement of three quarters (of the maximum of whatever is being measured):

```
<meter>75%</meter>
<meter>750%</meter>
<meter>3/4</meter>
<meter>6 blocks used (out of 8 total)</meter>
<meter>max: 100; current: 75</meter>
<meter><object data="graph75.png">0.75</object></meter>
<meter min="0" max="100" value="75"></meter>
```

The following example is incorrect use of the element, because it doesn't give a range (and since the default maximum is 1, both of the gauges would end up looking maxed out):

```
<p>The grapefruit pie had a radius of <meter>12cm</meter>
and a height of <meter>2cm</meter>. <!-- BAD! -->
```

Instead, one would either not include the meter element, or use the meter element with a defined range to give the dimensions in context compared to other pies:

```
<p>The grapefruit pie had a radius of 12cm and a height of
2cm.</p>
<dl>
  <dt>Radius: <dd> <meter min=0 max=20 value=12>12cm</meter>
  <dt>Height: <dd> <meter min=0 max=10 value=2>2cm</meter>
</dl>
```

There is no explicit way to specify units in the meter element, but the units may be specified in the title attribute in free-form text.

The example above could be extended to mention the units:

```
<dl>
  <dt>Radius: <dd> <meter min=0 max=20 value=12
  title="centimeters">12cm</meter>
  <dt>Height: <dd> <meter min=0 max=10 value=2
  title="centimeters">2cm</meter>
</dl>
```

User agent requirements: User agents must parse the min, max, value, low, high, and optimum attributes using the rules for parsing floating point number values (page 34).

If the value attribute has been omitted, the user agent must also process the textContent of the element according to the steps for finding one or two numbers of a ratio in a string (page 36). These steps will return nothing, one number, one number with a denominator punctuation character, or two numbers.

User agents must then use all these numbers to obtain values for six points on the gauge, as follows. (The order in which these are evaluated is important, as some of the values refer to earlier ones.)

The minimum value

If the `min` attribute is specified and a value could be parsed out of it, then the minimum value is that value. Otherwise, the minimum value is zero.

The maximum value

If the `max` attribute is specified and a value could be parsed out of it, the maximum value is that value.

Otherwise, if the `max` attribute is specified but no value could be parsed out of it, or if it was not specified, but either or both of the `min` or `value` attributes were specified, then the maximum value is 1.

Otherwise, none of the `max`, `min`, and `value` attributes were specified. If the result of processing the `textContent` of the element was either nothing or just one number with no denominator punctuation character, then the maximum value is 1; if the result was one number but it had an associated denominator punctuation character, then the maximum value is the value associated with that denominator punctuation character (page 36); and finally, if there were two numbers parsed out of the `textContent`, then the maximum is the higher of those two numbers.

If the above machinations result in a maximum value less than the minimum value, then the maximum value is actually the same as the minimum value.

The actual value

If the `value` attribute is specified and a value could be parsed out of it, then that value is the actual value.

If the `value` attribute is not specified but the `max` attribute is specified and the result of processing the `textContent` of the element was one number with no associated denominator punctuation character, then that number is the actual value.

If neither of the `value` and `max` attributes are specified, then, if the result of processing the `textContent` of the element was one number (with or without an associated denominator punctuation character), then that is the actual value, and if the result of processing the `textContent` of the element was two numbers, then the actual value is the lower of the two numbers found.

Otherwise, if none of the above apply, the actual value is zero.

If the above procedure results in an actual value less than the minimum value, then the actual value is actually the same as the minimum value.

If, on the other hand, the result is an actual value greater than the maximum value, then the actual value is the maximum value.

The low boundary

If the `low` attribute is specified and a value could be parsed out of it, then the low boundary is that value. Otherwise, the low boundary is the same as the minimum value.

If the above results in a low boundary that is less than the minimum value, the low boundary is the minimum value.

The high boundary

If the high attribute is specified and a value could be parsed out of it, then the high boundary is that value. Otherwise, the high boundary is the same as the maximum value.

If the above results in a high boundary that is higher than the maximum value, the high boundary is the maximum value.

The optimum point

If the optimum attribute is specified and a value could be parsed out of it, then the optimum point is that value. Otherwise, the optimum point is the midpoint between the minimum value and the maximum value.

If the optimum point is then less than the minimum value, then the optimum point is actually the same as the minimum value. Similarly, if the optimum point is greater than the maximum value, then it is actually the maximum value instead.

All of which should result in the following inequalities all being true:

- minimum value ≤ actual value ≤ maximum value
- minimum value ≤ low boundary ≤ high boundary ≤ maximum value
- minimum value ≤ optimum point ≤ maximum value

UA requirements for regions of the gauge: If the optimum point is equal to the low boundary or the high boundary, or anywhere in between them, then the region between the low and high boundaries of the gauge must be treated as the optimum region, and the low and high parts, if any, must be treated as suboptimal. Otherwise, if the optimum point is less than the low boundary, then the region between the minimum value and the low boundary must be treated as the optimum region, the region between the low boundary and the high boundary must be treated as a suboptimal region, and the region between the high boundary and the maximum value must be treated as an even less good region. Finally, if the optimum point is higher than the high boundary, then the situation is reversed; the region between the high boundary and the maximum value must be treated as the optimum region, the region between the high boundary and the low boundary must be treated as a suboptimal region, and the remaining region between the low boundary and the minimum value must be treated as an even less good region.

UA requirements for showing the gauge: When representing a meter element to the user, the UA should indicate the relative position of the actual value to the minimum and maximum values, and the relationship between the actual value and the three regions of the gauge.

The following markup:

```
<h3>Suggested groups</h3>
<menu type="toolbar">
  <a href="?cmd=hsg" onclick="hideSuggestedGroups()>Hide suggested
  groups</a>
</menu>
<ul>
  <li>
    <p><a href="/group/comp.infosystems.www.authoring.stylesheets/
    view">comp.infosystems.www.authoring.stylesheets</a> -
      <a href="/group/comp.infosystems.www.authoring.stylesheets/
      subscribe">join</a></p>
```

```

<p>Group description: <strong>Layout/presentation on the
WWW.</strong></p>
<p><b><meter value="0.5">Moderate activity,</meter></b> Usenet, 618
subscribers</p>
</li>
<li>
  <p><a href="/group/netscape.public.mozilla.xpininstall/
view">netscape.public.mozilla.xpininstall</a> -
    <a href="/group/netscape.public.mozilla.xpininstall/
subscribe">join</a></p>
  <p>Group description: <strong>Mozilla XPIInstall
discussion.</strong></p>
  <p><b><meter value="0.25">Low activity,</meter></b> Usenet, 22
subscribers</p>
</li>
<li>
  <p><a href="/group/mozilla.dev.general/view">mozilla.dev.general</a>
  ->
    <a href="/group/mozilla.dev.general/subscribe">join</a></p>
  <p><b><meter value="0.25">Low activity,</meter></b> Usenet, 66
subscribers</p>
</li>
</ul>
```

Might be rendered as follows:

Suggested groups - [Hide suggested groups](#)

[comp.infosystems.www.authoring.stylesheets](#) - [join](#)
 Group description: Layout/presentation on the WWW.
 [Usenet, 618 subscribers](#)

[netscape.public.mozilla.xpininstall](#) - [join](#)
 Group description: Mozilla XPIInstall discussion.
 [Usenet, 22 subscribers](#)

[mozilla.dev.general](#) - [join](#)
 [Usenet, 66 subscribers](#)

User agents may combine the value of the title attribute and the other attributes to provide context-sensitive help or inline text detailing the actual values.

For example, the following snippet:

```
<meter min=0 max=60 value=23.2 title="seconds"></meter>
```

...might cause the user agent to display a gauge with a tooltip saying "Value: 23.2 out of 60." on one line and "seconds" on a second line.

The **min**, **max**, **value**, **low**, **high**, and **optimum** DOM attributes must reflect (page 67) the respective content attributes of the same name. When the relevant content attributes are absent, the DOM attributes must return zero. The value parsed from the `textContent` never affects the DOM values.

** Would be cool to have the value DOM attribute update the textContent in-line...

4.6.13 The code element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The code element represents a fragment of computer code. This could be an XML element name, a filename, a computer program, or any other string that a computer would recognize.

Although there is no formal way to indicate the language of computer code being marked up, authors who wish to mark code elements with the language used, e.g. so that syntax highlighting scripts can use the right rules, may do so by adding a class prefixed with "language-" to the element.

The following example shows how the element can be used in a paragraph to mark up element names and computer code, including punctuation.

```
<p>The <code>code</code> element represents a fragment of computer code.</p>
```

```
<p>When you call the <code>activate()</code> method on the <code>robotSnowman</code> object, the eyes glow.</p>
```

```
<p>The example below uses the <code>begin</code> keyword to indicate the start of a statement block. It is paired with an <code>end</code> keyword, which is followed by the <code>.</code> punctuation character (full stop) to indicate the end of the program.</p>
```

The following example shows how a block of code could be marked up using the pre and code elements.

```
<pre><code class="language-pascal">var i: Integer;
begin
  i := 1;
end.</code></pre>
```

A class is used in that example to indicate the language used.

Note: See the `pre` element for more details.

4.6.14 The `var` element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

None.

DOM interface:

Uses `HTMLElement`.

The `var` element represents a variable. This could be an actual variable in a mathematical expression or programming context, or it could just be a term used as a placeholder in prose.

In the paragraph below, the letter "n" is being used as a variable in prose:

```
<p>If there are <var>n</var> pipes leading to the ice  
cream factory then I expect at least</em> <var>n</var>  
flavours of ice cream to be available for purchase!</p>
```

4.6.15 The `samp` element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

None.

DOM interface:

Uses `HTMLElement`.

The `samp` element represents (sample) output from a program or computing system.

Note: See the `pre` and `kbd` elements for more details.

This example shows the `samp` element being used inline:

```
<p>The computer said <samp>Too much cheese in tray  
two</samp> but I didn't know what that meant.</p>
```

This second example shows a block of sample output. Nested samp and kbd elements allow for the styling of specific elements of the sample output using a style sheet.

```
<pre><samp><samp class="prompt">jdoe@mowmow:~$</samp> <kbd>ssh  
demo.example.com</kbd>  
Last login: Tue Apr 12 09:10:17 2005 from mowmow.example.com on pts/1  
Linux demo  
2.6.10-grsec+gg3+e+fhs6b+nfs+gr0501+++p3+c4a+gr2b-reslog-v6.189 #1 SMP  
Tue Feb 1 11:22:36 PST 2005 i686 unknown  
  
<samp class="prompt">jdoe@demo:~$</samp> <samp  
class="cursor">_</samp></samp></pre>
```

4.6.16 The kbd element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The kbd element represents user input (typically keyboard input, although it may also be used to represent other input, such as voice commands).

When the kbd element is nested inside a samp element, it represents the input as it was echoed by the system.

When the kbd element *contains* a samp element, it represents input based on system output, for example invoking a menu item.

When the kbd element is nested inside another kbd element, it represents an actual key or other single unit of input as appropriate for the input mechanism.

Here the kbd element is used to indicate keys to press:

```
<p>To make George eat an apple, press  
<kbd><kbd>Shift</kbd>+<kbd>F3</kbd></kbd></p>
```

In this second example, the user is told to pick a particular menu item. The outer kbd element marks up a block of input, with the inner kbd elements representing each

individual step of the input, and the samp elements inside them indicating that the steps are input based on something being displayed by the system, in this case menu labels:

```
<p>To make George eat an apple, select  
    <kbd><kbd><samp>File</samp></kbd>|<kbd><samp>Eat  
Apple...</samp></kbd></kbd>  
</p>
```

4.6.17 The sub and sup elements

Categories

Phrasing content (page 94).

Contexts in which these elements may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The sup element represents a superscript and the sub element represents a subscript.

These elements must be used only to mark up typographical conventions with specific meanings, not for typographical presentation for presentation's sake. For example, it would be inappropriate for the sub and sup elements to be used in the name of the LaTeX document preparation system. In general, authors should use these elements only if the absence of those elements would change the meaning of the content.

When the sub element is used inside a var element, it represents the subscript that identifies the variable in a family of variables.

```
<p>The coordinate of the <var>i</var>th point is  
(<var>x<sub>i</sub></var>,  
<var>y<sub>i</sub></var>).  
For example, the 10th point has coordinate  
(<var>x<sub>10</sub></var>, <var>y<sub>10</sub></var>).</p>
```

In certain languages, superscripts are part of the typographical conventions for some abbreviations.

```
<p>The most beautiful women are  
<span lang="fr"><abbr>Mlle</sup></abbr> Gwendoline</span> and  
<span lang="fr"><abbr>Mme</sup></abbr> Denise</span>.</p>
```

Mathematical expressions often use subscripts and superscripts. Authors are encouraged to use MathML for marking up mathematics, but authors may opt to use sub and sup if detailed mathematical markup is not desired. [MathML]

```
<var>E</var>=<var>m</var><var>c</var><sup>2</sup>
```

```
f(<var>x</var>, <var>n</var>) =  
log<sub>4</sub><var>x</var><sup><var>n</var></sup>
```

4.6.18 The span element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The span element doesn't mean anything on its own, but can be useful when used together with other attributes, e.g. class, lang, or dir.

4.6.19 The i element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The i element represents a span of text in an alternate voice or mood, or otherwise offset from the normal prose, such as a taxonomic designation, a technical term, an idiomatic phrase from another language, a thought, a ship name, or some other prose whose typical typographic presentation is italicized.

Terms in languages different from the main text should be annotated with lang attributes (xml:lang in XML).

The examples below show uses of the i element:

```
<p>The <i class="taxonomy">Felis silvestris catus</i> is cute.</p>  
<p>The term <i>prose content</i> is defined above.</p>  
<p>There is a certain <i lang="fr">je ne sais quoi</i> in the air.</p>
```

In the following example, a dream sequence is marked up using i elements.

```
<p>Raymond tried to sleep.</p>
<p><i>The ship sailed away on Thursday</i>, he
dreamt. <i>The ship had many people aboard, including a beautiful
princess called Carey. He watched her, day-in, day-out, hoping she
would notice him, but she never did.</i></p>
<p><i>Finally one night he picked up the courage to speak with
her</i></p>
<p>Raymond woke with a start as the fire alarm rang out.</p>
```

The i element should be used as a last resort when no other element is more appropriate. In particular, citations should use the cite element, defining instances of terms should use the dfn element, stress emphasis should use the em element, importance should be denoted with the strong element, quotes should be marked up with the q element, and small print should use the small element.

Authors are encouraged to use the class attribute on the i element to identify why the element is being used, so that if the style of a particular use (e.g. dream sequences as opposed to taxonomic terms) is to be changed at a later date, the author doesn't have to go through the entire document (or series of related documents) annotating each use.

Note: Style sheets can be used to format i elements, just like any other element can be restyled. Thus, it is not the case that content in i elements will necessarily be italicized.

4.6.20 The b element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The b element represents a span of text to be stylistically offset from the normal prose without conveying any extra importance, such as key words in a document abstract, product names in a review, or other spans of text whose typical typographic presentation is boldened.

The following example shows a use of the b element to highlight key words without marking them up as important:

```
<p>The <b>frobonitor</b> and <b>barbinator</b> components are
fried.</p>
```

In the following example, objects in a text adventure are highlighted as being special by use of the b element.

```
<p>You enter a small room. Your <b>sword</b> glows  
brighter. A <b>rat</b> scurries past the corner wall.</p>
```

Another case where the b element is appropriate is in marking up the lede (or lead) sentence or paragraph. The following example shows how a BBC article about kittens adopting a rabbit as their own could be marked up using HTML5 elements:

```
<article>  
  <h2>Kittens 'adopted' by pet rabbit</h2>  
  <p><b>Six abandoned kittens have found an unexpected new  
  mother figure – a pet rabbit.</b></p>  
  <p>Veterinary nurse Melanie Humble took the three-week-old  
  kittens to her Aberdeen home.</p>  
  [...]
```

The b element should be used as a last resort when no other element is more appropriate. In particular, headers should use the h1 to h6 elements, stress emphasis should use the em element, importance should be denoted with the strong element, and text marked or highlighted should use the mark element.

The following would be *incorrect* usage:

```
<p><b>WARNING!</b> Do not frob the barbinator!</p>
```

In the previous example, the correct element to use would have been strong, not b.

Note: Style sheets can be used to format b elements, just like any other element can be restyled. Thus, it is not the case that content in b elements will necessarily be boldened.

4.6.21 The bdo element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

None, but the dir global attribute has special requirements on this element.

DOM interface:

Uses HTMLElement.

The bdo element allows authors to override the Unicode bidi algorithm by explicitly specifying a direction override. [BIDI]

Authors must specify the `dir` attribute on this element, with the value `ltr` to specify a left-to-right override and with the value `rtl` to specify a right-to-left override.

If the element has the `dir` attribute set to the exact value `ltr`, then for the purposes of the bidi algorithm, the user agent must act as if there was a U+202D LEFT-TO-RIGHT OVERRIDE character at the start of the element, and a U+202C POP DIRECTIONAL FORMATTING at the end of the element.

If the element has the `dir` attribute set to the exact value `rtl`, then for the purposes of the bidi algorithm, the user agent must act as if there was a U+202E RIGHT-TO-LEFT OVERRIDE character at the start of the element, and a U+202C POP DIRECTIONAL FORMATTING at the end of the element.

The requirements on handling the `bdo` element for the bidi algorithm may be implemented indirectly through the style layer. For example, an HTML+CSS user agent should implement these requirements by implementing the CSS `unicode-bidi` property. [CSS21]

4.6.22 The ruby element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

One or more groups of: phrasing content (page 94) followed either by a single `rt` element, or an `rp` element, an `rt` element, and another `rp` element.

Element-specific attributes:

None.

DOM interface:

Uses `HTMLElement`.

The `ruby` element allows one or more spans of phrasing content to be marked with ruby annotations.

A `ruby` element represents the spans of phrasing content it contains, ignoring all the child `rt` and `rp` elements and their descendants. Those spans of phrasing content have associated annotations created using the `rt` element.

In this example, each ideograph in the text ??? is annotated with its reading.

```
... <ruby>
? <rt> ?? </rt>
? <rt> ?? </rt>
? <rt> ?? </rt>
? <rt> ?       </rt>
</ruby> ...
```

This might be rendered as:

さいとうのぶお
... 斎藤信男 ...

4.6.23 The `rt` element

Categories

None.

Contexts in which this element may be used:

As a child of a ruby element.

Content model:

Phrasing content (page 94).

Element-specific attributes:

None.

DOM interface:

Uses `HTMLElement`.

The `rt` element marks the ruby text component of a ruby annotation.

An `rt` element that is a child of a `ruby` element represents an annotation (given by its children) for the zero or more nodes of phrasing content that immediately precedes it in the `ruby` element, ignoring `rp` elements.

An `rt` element that is not a child of a `ruby` element represents the same thing as its children.

4.6.24 The `rp` element

Categories

None.

Contexts in which this element may be used:

As a child of a ruby element, either immediately before or immediately after an `rt` element.

Content model:

If the `rp` element is immediately after an `rt` element that is immediately preceded by another `rp` element: a single character from Unicode character class Pe.
Otherwise: a single character from Unicode character class Ps.

Element-specific attributes:

None.

DOM interface:

Uses `HTMLElement`.

The `rp` element can be used to provide parentheses around a ruby text component of a ruby annotation, to be shown by user agents that don't support ruby annotations.

An `rp` element that is a child of a `ruby` element represents nothing and it and its contents must be ignored. An `rp` element whose parent element is not a `ruby` element represents the same thing as its children.

The example above, in which each ideograph in the text ???? is annotated with its reading, could be expanded to use `rp` so that in legacy user agent the readings are in parentheses:

```
... <ruby>
? <rp>(</rp><rt>??</rt><rp>) </rp>
? <rp>(</rp><rt>??</rt><rp>) </rp>
? <rp>(</rp><rt>??</rt><rp>) </rp>
? <rp>(</rp><rt>??</rt><rp>) </rp>
</ruby> ...
```

In conforming user agents the rendering would be as above, but in user agents that do not support ruby, the rendering would be:

```
... ? (??) ? (??) ? (??) ? (?) ...
```

4.6.25 Usage summary

** We need to summarize the various elements, in particular to distinguish b/i/em/strong/var/q/mark/cite.

4.6.26 Footnotes

HTML does not have a dedicated mechanism for marking up footnotes. Here are the recommended alternatives.

For short inline annotations, the `title` attribute should be used.

In this example, two parts of a dialog are annotated.

```
<dialog>
<dt>Customer
<dd>Hello! I wish to register a complaint. Hello. Miss?
<dt>Shopkeeper
<dd><span title="Colloquial pronunciation of 'What do you'">Watcha</span> mean, miss?
<dt>Customer
<dd>Uh, I'm sorry, I have a cold. I wish to make a complaint.
<dt>Shopkeeper
<dd>Sorry, <span title="This is, of course, a lie.">we're closing for lunch</span>.
```

For longer annotations, the `a` element should be used, pointing to an element later in the document. The convention is that the contents of the link be a number in square brackets.

In this example, a footnote in the dialog links to a paragraph below the dialog. The paragraph then reciprocally links back to the dialog, allowing the user to return to the location of the footnote.

```
<dialog>
  <dt>Announcer
  <dd>Number 16: The <i>hand</i>.
  <dt>Interviewer
  <dd>Good evening. I have with me in the studio tonight Mr Norman St John Polevaulter, who for the past few years has been contradicting people. Mr Polevaulter, why <em>do</em> you contradict people?
  <dt>Norman
  <dd>I don't. <a href="#fn1" id="r1">[1]</a>
  <dt>Interviewer
  <dd>You told me you did!
</dialog>
<section>
  <p id="fn1"><a href="#r1">[1]</a> This is, naturally, a lie, but paradoxically if it were true he could not say so without contradicting the interviewer and thus making it false.</p>
</section>
```

For side notes, longer annotations that apply to entire sections of the text rather than just specific words or sentences, the aside element should be used.

In this example, a sidebar is given after a dialog, giving some context to the dialog.

```
<dialog>
  <dt>Customer
  <dd>I will not buy this record, it is scratched.
  <dt>Shopkeeper
  <dd>I'm sorry?
  <dt>Customer
  <dd>I will not buy this record, it is scratched.
  <dt>Shopkeeper
  <dd>No no no, this's'a tobacconist's.
</dialog>
<aside>
  <p>In 1970, the British Empire lay in ruins, and foreign nationalists frequented the streets – many of them Hungarians (not the streets – the foreign nationals). Sadly, Alexander Yalt has been publishing incompetently-written phrase books.
</aside>
```

4.7 Edits

The ins and del elements represent edits to the document.

4.7.1 The ins element

Categories

When the element only contains phrasing content (page 94): phrasing content (page 94).

Otherwise: flow content (page 93).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Transparent (page 96).

Element-specific attributes:

`cite`
`datetime`

DOM interface:

Uses the `HTMLModElement` interface.

The `ins` element represents an addition to the document.

The following represents the addition of a single paragraph:

```
<aside>
  <ins>
    <p> I like fruit. </p>
  </ins>
</aside>
```

As does this, because everything in the `aside` element here counts as phrasing content (page 94) and therefore there is just one paragraph (page 96):

```
<aside>
  <ins>
    Apples are <em>tasty</em>.
  </ins>
  <ins>
    So are pears.
  </ins>
</aside>
```

`ins` elements should not cross implied paragraph (page 96) boundaries.

The following example represents the addition of two paragraphs, the second of which was inserted in two parts. The first `ins` element in this example thus crosses a paragraph boundary, which is considered poor form.

```
<aside>
  <ins datetime="2005-03-16T00:00Z">
    <p> I like fruit. </p>
    Apples are <em>tasty</em>.
  </ins>
  <ins datetime="2007-12-19T00:00Z">
    So are pears.
  </ins>
```

```
</ins>  
</aside>
```

Here is a better way of marking this up. It uses more elements, but none of the elements cross implied paragraph boundaries.

```
<aside>  
  <ins datetime="2005-03-16T00:00Z">  
    <p> I like fruit. </p>  
  </ins>  
  <ins datetime="2005-03-16T00:00Z">  
    Apples are <em>tasty</em>.  
  </ins>  
  <ins datetime="2007-12-19T00:00Z">  
    So are pears.  
  </ins>  
</aside>
```

4.7.2 The **del** element

Categories

When the element only contains phrasing content (page 94): phrasing content (page 94).

Otherwise: flow content (page 93).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Transparent (page 96).

Element-specific attributes:

`cite`
`datetime`

DOM interface:

Uses the `HTMLModElement` interface.

The `del` element represents a removal from the document.

`del` elements should not cross implied paragraph (page 96) boundaries.

4.7.3 Attributes common to `ins` and `del` elements

The `cite` attribute may be used to specify the address of a document that explains the change. When that document is long, for instance the minutes of a meeting, authors are encouraged to include a fragment identifier pointing to the specific part of that document that discusses the change.

If the `cite` attribute is present, it must be a valid URL (page 52) that explains the change. User agents should allow users to follow such citation links.

The **datetime** attribute may be used to specify the time and date of the change.

If present, the **datetime** attribute must be a valid **datetime** (page 40) value.

User agents must parse the **datetime** attribute according to the **parse a string as a datetime value** (page 41) algorithm. If that doesn't return a time, then the modification has no associated timestamp (the value is non-conforming; it is not a valid **datetime** (page 40)). Otherwise, the modification is marked as having been made at the given **datetime**. User agents should use the associated **timezone** information to determine which **timezone** to present the given **datetime** in.

The **ins** and **del** elements must implement the **HTMLModElement** interface:

```
interface HTMLModElement : HTMLElement {  
    attribute DOMString cite;  
    attribute DOMString dateTime;  
};
```

The **cite** DOM attribute must reflect (page 67) the element's **cite** content attribute. The **dateTime** DOM attribute must reflect (page 67) the element's **datetime** content attribute.

4.7.4 Edits and paragraphs

Since the **ins** and **del** elements do not affect **paragraphing** (page 96), it is possible, in some cases where paragraphs are implied (page 96) (without explicit **p** elements), for an **ins** or **del** element to span both an entire paragraph or other non-phrasing content (page 94) elements and part of another paragraph.

For example:

```
<section>  
<ins>  
  <p>  
    This is a paragraph that was inserted.  
  </p>  
  This is another paragraph whose first sentence was inserted  
  at the same time as the paragraph above.  
</ins>  
  This is a second sentence, which was there all along.  
</section>
```

By only wrapping some paragraphs in **p** elements, one can even get the end of one paragraph, a whole second paragraph, and the start of a third paragraph to be covered by the same **ins** or **del** element (though this is very confusing, and not considered good practice):

```
<section>  
  This is the first paragraph. <ins>This sentence was  
  inserted.  
  <p>This second paragraph was inserted.</p>  
  This sentence was inserted too.</ins> This is the  
  third paragraph in this example.  
</section>
```

However, due to the way implied paragraphs (page 96) are defined, it is not possible to mark up the end of one paragraph and the start of the very next one using the same ins or del element. You instead have to use one (or two) p element(s) and two ins or del elements:

For example:

```
<section>
  <p>This is the first paragraph. <del>This sentence was
  deleted.</del></p>
  <p><del>This sentence was deleted too.</del> That
  sentence needed a separate &lt;del&gt; element.</p>
</section>
```

Partly because of the confusion described above, authors are strongly recommended to always mark up all paragraphs with the p element, and to not have any ins or del elements that cross across any implied paragraphs (page 96).

4.7.5 Edits and lists

The content models of the ol and ul elements do not allow ins and del elements as children. Lists always represent all their items, including items that would otherwise have been marked as deleted.

To indicate that an item is inserted or deleted, an ins or del element can be wrapped around the contents of the li element. To indicate that an item has been replaced by another, a single li element can have one or more del elements followed by one or more ins elements.

In the following example, a list that started empty had items added and removed from it over time. The bits in the example that have been emphasised show the parts that are the "current" state of the list. The list item numbers don't take into account the edits, though.

```
<h1>Stop-ship bugs</h1>
<ol>
  <li><ins datetime="2008-02-12 15:20 Z">Bug 225: Rain detector
    doesn't work in snow</ins></li>
  <li><del datetime="2008-03-01 20:22 Z"><ins datetime="2008-02-14
  12:02 Z">Bug 228: Water buffer overflows in April</ins></del></li>
  <li><ins datetime="2008-02-16 13:50 Z">Bug 230: Water heater
    doesn't use renewable fuels</ins></li>
  <li><del datetime="2008-02-20 21:15 Z"><ins datetime="2008-02-16
  14:25 Z">Bug 232: Carbon dioxide emissions detected after
  startup</ins></del></li>
</ol>
```

In the following example, a list that started with just fruit was replaced by a list with just colors.

```
<h1>List of <del>fruits</del><ins>colors</ins></h1>
<ul>
  <li><del>Lime</del><ins>Green</ins></li>
  <li><del>Apple</del></li>
```

```
<li>Orange</li>
<li><del>Pear</del></li>
<li><ins>Teal</ins></li>
<li><del>Lemon</del><ins>Yellow</ins></li>
<li>Olive</li>
<li><ins>Purple</ins>
</ul>
```

4.8 Embedded content

4.8.1 The figure element

Categories

Flow content (page 93).
Sectioning root (page 140).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

Either: one legend element followed by flow content (page 93).
Or: Flow content (page 93) followed by one legend element.
Or: Flow content (page 93).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The figure element represents some flow content (page 93), optionally with a caption, which can be moved away from the main flow of the document without affecting the document's meaning.

The element can thus be used to annotate illustrations, diagrams, photos, code listings, etc, that are referred to from the main content of the document, but that could, without affecting the flow of the document, be moved away from that primary content, e.g. to the side of the page, to dedicated pages, or to an appendix.

The first legend element child of the element, if any, represents the caption of the figure element's contents. If there is no child legend element, then there is no caption.

The remainder of the element's contents, if any, represents the content.

This example shows the figure element to mark up a code listing.

```
<p>In <a href="#l4">listing 4</a> we see the primary core interface API declaration.</p>
<figure id="l4">
  <legend>Listing 4. The primary core interface API declaration.</legend>
  <pre><code>interface PrimaryCore {
    boolean verifyDataLine();</code></pre>
</figure>
```

```

    void sendData(in sequence<byte> data);
    void initSelfDestruct();
}></code></pre>
</figure>
<p>The API is designed to use UTF-8.</p>

```

Here we see a figure element to mark up a photo.

```

<figure>

<legend>Bubbles at work</legend>
</figure>

```

In this example, we see an image that is *not* a figure, as well as an image and a video that are.

<h2>Malinko's comics</h2>

<p>This case centered on some sort of "intellectual property" infringement related to a comic (see Exhibit A). The suit started after a trailer ending with these words:</p>

<p>...was aired. A lawyer, armed with a Bigger Notebook, launched a preemptive strike using snowballs. A complete copy of the trailer is included with Exhibit B.</p>

```

<figure>

<legend>Exhibit A. The alleged <cite>rough copy</cite> comic.</legend>
</figure>

```

```

<figure>
<video src="ex-b.mov"></video>
<legend>Exhibit A. The alleged <cite>rough copy</cite> comic.</legend>
</figure>

```

<p>The case was resolved out of court.</p>

Here, a part of a poem is marked up using figure.

```

<figure>
<p>'Twas brillig, and the slithy toves<br>
Did gyre and gimble in the wabe;<br>
All mimsy were the borogoves,<br>
And the mome raths outgrabe.</p>
<legend><cite>Jabberwocky</cite> (first verse). Lewis Carroll,
1832-98</legend>
</figure>

```

In this example, which could be part of a much larger work discussing a castle, the figure has three images in it.

```
<figure>
  
  
  
  <legend>The castle through the ages: 1423, 1858, and 1999
respectively.</legend>
</figure>
```

4.8.2 The `img` element

Categories

Embedded content (page 94).

Contexts in which this element may be used:

Where embedded content (page 94) is expected.

Content model:

Empty.

Element-specific attributes:

alt
src
usemap
ismap
width
height

DOM interface:

```
[NamedConstructor=Image(),
 NamedConstructor=Image(in unsigned long width),
 NamedConstructor=Image(in unsigned long width, in unsigned long
height)]
interface HTMLImageElement : HTMLElement {
    attribute DOMString alt;
    attribute DOMString src;
    attribute DOMString useMap;
    attribute boolean isMap;
    attribute unsigned long width;
    attribute unsigned long height;
    readonly attribute boolean complete;
};
```

An `img` element represents an image.

The image given by the `src` attribute is the embedded content, and the value of the `alt` attribute is the `img` element's fallback content (page 95).

The `src` attribute must be present, and must contain a valid URL (page 52) referencing a non-interactive, optionally animated, image resource that is neither paged nor scripted.

Note: *Images can thus be static bitmaps (e.g. PNGs, GIFs, JPEGs), single-page vector documents (single-page PDFs, XML files with an SVG root element), animated bitmaps (APNGs, animated GIFs), animated vector graphics (XML files with an SVG root element that use declarative SMIL animation), and so forth. However, this also precludes SVG files with script, multipage PDF files, interactive MNG files, HTML documents, plain text documents, and so forth.*

The requirements on the `alt` attribute's value are described in the next section (page 205).

- ** There has been some suggestion that the `longdesc` attribute from HTML4, or some other mechanism that is more powerful than `alt=""`, should be included. This has not yet been considered.

The `img` must not be used as a layout tool. In particular, `img` elements should not be used to display fully transparent images, as they rarely convey meaning and rarely add anything useful to the document.

When an `img` is created with a `src` attribute, and whenever the `src` attribute is set subsequently, the user agent must fetch (page 59) the resource specified by the `src` attribute's value, unless the user agent cannot support images, or its support for images has been disabled, or the user agent only fetches elements on demand.

Fetching the image must delay the `load` event (page 657).

⚠Warning! *This, unfortunately, can be used to perform a rudimentary port scan of the user's local network (especially in conjunction with scripting, though scripting isn't actually necessary to carry out such an attack). User agents may implement cross-origin (page 423) access control policies that mitigate this attack.*

If the image's type is a supported image type, and the image is a valid image of that type, then the image is said to be *available* (this affects exactly what the element represents, as defined below). This can be true even before the image is completely downloaded, if the user agent supports incremental rendering of images; in such cases, each task (page 429) that is queued (page 429) by the networking task source (page 430) while the image is being fetched (page 59) must update the presentation of the image appropriately.

Whether the image is fetched successfully or not (e.g. whether the response code was a 2xx code or equivalent) must be ignored when determining the image's type and whether it is a valid image.

Note: *This allows servers to return images with error responses, and have them displayed.*

The user agents should apply the image sniffing rules (page 65) to determine the type of the image, with the image's associated Content-Type headers (page 60) giving the *official type*. If these rules are not applied, then the type of the image must be the type given by the image's associated Content-Type headers (page 60).

User agents must not support non-image resources with the `img` element (e.g. XML files whose root element is an HTML element). User agents must not run executable code (e.g. scripts) embedded in the image resource. User agents must only display the first page of a multipage resource (e.g. a PDF file). User agents must not allow the resource to act in an interactive fashion, but should honour any animation in the resource.

This specification does not specify which image types are to be supported.

The task (page 429) that is queued (page 429) by the networking task source (page 430) once the resource has been fetched (page 59), must, if the download was successful and the image is *available*, queue a task (page 429) to fire a `load` event (page 436) on the `img` element (this happens after `complete` starts returning true); and otherwise, if the fetching process fails without a response from the remote server, or completes but the image is not a valid or supported image, queue a task (page 429) to fire an `error` event (page 436) on the `img` element.

What an `img` element represents depends on the `src` attribute and the `alt` attribute.

↪ **If the `src` attribute is set and the `alt` attribute is set to the empty string**

The image is either decorative or supplemental to the rest of the content, redundant with some other information in the document.

If the image is *available* and the user agent is configured to display that image, then the element represents the image specified by the `src` attribute.

Otherwise, the element represents nothing, and may be omitted completely from the rendering. User agents may provide the user with a notification that an image is present but has been omitted from the rendering.

↪ **If the `src` attribute is set and the `alt` attribute is set to a value that isn't empty**

The image is a key part of the content; the `alt` attribute gives a textual equivalent or replacement for the image.

If the image is *available* and the user agent is configured to display that image, then the element represents the image specified by the `src` attribute.

Otherwise, the element represents the text given by the `alt` attribute. User agents may provide the user with a notification that an image is present but has been omitted from the rendering.

↪ **If the `src` attribute is set and the `alt` attribute is not**

The image might be a key part of the content, and there is no textual equivalent of the image available.

If the image is *available*, the element represents the image specified by the `src` attribute.

If the image is not *available* or if the user agent is not configured to display the image, then the user agent should display some sort of indicator that there is an image that is not being rendered, and may, if requested by the user, or if so configured, or when required to provide contextual information in response to navigation, provide caption information for the image, derived as follows:

1. If the image has a `title` attribute whose value is not the empty string, then the value of that attribute is the caption information; abort these steps.
2. If the image is the child of a `figure` element that has a child `legend` element, then the contents of the first such `legend` element are the caption information; abort these steps.
3. Run the algorithm to create the outline (page 142) for the document.
4. If the `img` element did not end up associated with a heading in the outline, or if there are any other images that are lacking an `alt` attribute and that are associated with the same heading in the outline as the `img` element in question, then there is no caption information; abort these steps.
5. The caption information is the heading with which the image is associated according to the outline.

↪ **If the `src` attribute is not set and either the `alt` attribute is set to the empty string or the `alt` attribute is not set at all**

The element represents nothing.

↪ **Otherwise**

The element represents the text given by the `alt` attribute.

The `alt` attribute does not represent advisory information. User agents must not present the contents of the `alt` attribute in the same way as content of the `title` attribute.

User agents may always provide the user with the option to display any image, or to prevent any image from being displayed. User agents may also apply image analysis heuristics to help the user make sense of the image when the user is unable to make direct use of the image, e.g. due to a visual disability or because they are using a text terminal with no graphics capabilities.

The *contents* of `img` elements, if any, are ignored for the purposes of rendering.

The `usemap` attribute, if present, can indicate that the image has an associated image map (page 291).

The `ismap` attribute, when used on an element that is a descendant of an `a` element with an `href` attribute, indicates by its presence that the element provides access to a server-side image map. This affects how events are handled on the corresponding `a` element.

The `ismap` attribute is a boolean attribute (page 32). The attribute must not be specified on an element that does not have an ancestor `a` element with an `href` attribute.

The `img` element supports dimension attributes (page 295).

The DOM attributes **alt**, **src**, **useMap**, and **isMap** each must reflect (page 67) the respective content attributes of the same name.

The DOM attributes **width** and **height** must return the rendered width and height of the image, in CSS pixels, if the image is being rendered, and is being rendered to a visual medium; or else the intrinsic width and height of the image, in CSS pixels, if the image is *available* but not being rendered to a visual medium; or else 0, if the image is not *available* or its dimensions are not known. [CSS21]

The DOM attribute **complete** must return true if the user agent has fetched the image specified in the **src** attribute, and it is a valid image, even if the final task (page 429) queued by the networking task source (page 430) for the fetching (page 59) of the image resource has not yet been processed. Otherwise, the attribute must return false.

Note: *The value of complete can thus change while a script is executing.*

Three constructors are provided for creating `HTMLImageElement` objects (in addition to the factory methods from DOM Core such as `createElement()`): `Image()`, `Image(width)`, and `Image(width, height)`. When invoked as constructors, these must return a new `HTMLImageElement` object (a new `img` element). If the `width` argument is present, the new object's `width` content attribute must be set to `width`. If the `height` argument is also present, the new object's `height` content attribute must be set to `height`.

A single image can have different appropriate alternative text depending on the context.

In each of the following cases, the same image is used, yet the `alt` text is different each time. The image is the coat of arms of the Canton Geneva in Switzerland.

Here it is used as a supplementary icon:

```
<p>I lived in  Carouge.</p>
```

Here it is used as an icon representing the town:

```
<p>Home town: </p>
```

Here it is used as part of a text on the town:

```
<p>Carouge has a coat of arms.</p>
<p></p>
<p>It is used as decoration all over the town.</p>
```

Here it is used as a way to support a similar text where the description is given as well as, instead of as an alternative to, the image:

```
<p>Carouge has a coat of arms.</p>
<p></p>
<p>The coat of arms depicts a lion, sitting in front of a tree.
It is used as decoration all over the town.</p>
```

Here it is used as part of a story:

```
<p>He picked up the folder and a piece of paper fell out.</p>
<p></p>
<p>He stared at the folder. S! The answer he had been looking for all
this time was simply the letter S! How had he not seen that before? It
all
came together now. The phone call where Hector had referred to a
lion's tail,
the time Marco had stuck his tongue out...</p>
```

Here it is not known at the time of publication what the image will be, only that it will be a coat of arms of some kind, and thus no replacement text can be provided, and instead only a brief caption for the image is provided, in the title attribute:

```

<p>The last user to have uploaded a coat of arms uploaded this one:</p>
<p></p>
```

Ideally, the author would find a way to provide real replacement text even in this case, e.g. by asking the previous user. Not providing replacement text makes the document more difficult to use for people who are unable to view images, e.g. blind users, or users or very low-bandwidth connections or who pay by the byte, or users who are forced to use a text-only Web browser.

Here are some more examples showing the same picture used in different contexts, with different appropriate alternate texts each time.

```

<article>
  <h1>My cats</h1>
  <h2>Fluffy</h2>
  <p>Fluffy is my favourite.</p>
  
  <p>She's just too cute.</p>
  <h2>Miles</h2>
  <p>My other cat, Miles just eats and sleeps.</p>
</article>
<article>
  <h1>Photography</h1>
  <h2>Shooting moving targets indoors</h2>
  <p>The trick here is to know how to anticipate; to know at what speed
and
what distance the subject will pass by.</p>
  
  <h2>Nature by night</h2>
  <p>To achieve this, you'll need either an extremely sensitive film, or
immense flash lights.</p>
</article>
<article>
  <h1>About me</h1>
  <h2>My pets</h2>
  <p>I've got a cat named Fluffy and a dog named Miles.</p>
  
<p>My dog Miles and I like go on long walks together.</p>
<h2>music</h2>
<p>After our walks, having emptied my mind, I like listening to
Bach.</p>
</article>
<article>
<h1>Fluffy and the Yarn</h1>
<p>Fluffy was a cat who liked to play with yarn. He also liked to
jump.</p>
<aside></aside>
<p>He would play in the morning, he would play in the evening.</p>
</article>

```

4.8.2.1 Requirements for providing text to act as an alternative for images

The requirements for the alt attribute depend on what the image is intended to represent, as described in the following sections.

4.8.2.1.1 A link or button containing nothing but the image

When an a (page 160) element that is a hyperlink (page 497), or a button element, has no textual content but contains one or more images, the alt attributes must contain text that together convey the purpose of the link or button.

In this example, a user is asked to pick his preferred color from a list of three. Each color is given by an image, but for users who have configured their user agent not to display images, the color names are used instead:

```

<h1>Pick your color</h1>
<ul>
  <li><a href="green.html"></a></li>
  <li><a href="blue.html"></a></li>
  <li><a href="red.html"></a></li>
</ul>

```

In this example, each button has a set of images to indicate the kind of color output desired by the user. The first image is used in each case to give the alternative text.

```

<button name="rgb"></button>
<button name="cmyk"></button>

```

Since each image represents one part of the text, it could also be written like this:

```

<button name="rgb"></button>
<button name="cmyk"></button>

```

However, with other alternative text, this might not work, and putting all the alternative text into one image in each case might make more sense:

```
<button name="rgb"></button>
<button name="cmyk"></button>
```

4.8.2.1.2 A phrase or paragraph with an alternative graphical representation: charts, diagrams, graphs, maps, illustrations

Sometimes something can be more clearly stated in graphical form, for example as a flowchart, a diagram, a graph, or a simple map showing directions. In such cases, an image can be given using the `img` element, but the lesser textual version must still be given, so that users who are unable to view the image (e.g. because they have a very slow connection, or because they are using a text-only browser, or because they are listening to the page being read out by a hands-free automobile voice Web browser, or simply because they are blind) are still able to understand the message being conveyed.

The text must be given in the `alt` attribute, and must convey the same message as the image specified in the `src` attribute.

It is important to realize that the alternative text is a *replacement* for the image, not a description of the image.

In the following example we have a flowchart in image form, with text in the `alt` attribute rephrasing the flowchart in prose form:

```
<p>In the common case, the data handled by the tokenization stage comes from the network, but it can also come from script.</p>
<p></p>
```

Here's another example, showing a good solution and a bad solution to the problem of including an image in a description.

First, here's the good solution. This sample shows how the alternative text should just be what you would have put in the prose if the image had never existed.

```
<!-- This is the correct way to do things. -->
<p>
  You are standing in an open field west of a house.
  
  There is a small mailbox here.
</p>
```

Second, here's the bad solution. In this incorrect way of doing things, the alternative text is simply a description of the image, instead of a textual replacement for the image. It's

bad because when the image isn't shown, the text doesn't flow as well as in the first example.

```
<!-- This is the wrong way to do things. -->
<p>
    You are standing in an open field west of a house.
    
    There is a small mailbox here.
</p>
```

4.8.2.1.3 A short phrase or label with an alternative graphical representation: icons, logos

A document can contain information in iconic form. The icon is intended to help users of visual browsers to recognize features at a glance.

In some cases, the icon is supplemental to a text label conveying the same meaning. In those cases, the alt attribute must be present but must be empty.

Here the icons are next to text that conveys the same meaning, so they have an empty alt attribute:

```
<nav>
    <p><a href="/help/"> Help</a></p>
    <p><a href="/configure/">
        Configuration Tools</a></p>
</nav>
```

In other cases, the icon has no text next to it describing what it means; the icon is supposed to be self-explanatory. In those cases, an equivalent textual label must be given in the alt attribute.

Here, posts on a news site are labeled with an icon indicating their topic.

```
<body>
    <article>
        <header>
            <h1>Ratatouille wins <i>Best Movie of the Year</i> award</h1>
            <p></p>
        </header>
        <p>Pixar has won yet another <i>Best Movie of the Year</i> award,
            making this its 8th win in the last 12 years.</p>
    </article>
    <article>
        <header>
            <h1>Latest TWiT episode is online</h1>
            <p></p>
        </header>
        <p>The latest TWiT episode has been posted, in which we hear
            several tech news stories as well as learning much more about the
            iPhone. This week, the panelists compare how reflective their
            iPhones' Apple logos are.</p>
```

```
</article>  
</body>
```

Many pages include logos, insignia, flags, or emblems, which stand for a particular entity such as a company, organization, project, band, software package, country, or some such.

If the logo is being used to represent the entity, e.g. as a page header, the alt attribute must contain the name of the entity being represented by the logo. The alt attribute must *not* contain text like the word "logo", as it is not the fact that it is a logo that is being conveyed, it's the entity itself.

If the logo is being used next to the name of the entity that it represents, then the logo is supplemental, and its alt attribute must instead be empty.

If the logo is merely used as decorative material (as branding, or, for example, as a side image in an article that mentions the entity to which the logo belongs), then the entry below on purely decorative images applies. If the logo is actually being discussed, then it is being used as a phrase or paragraph (the description of the logo) with an alternative graphical representation (the logo itself), and the first entry above applies.

In the following snippets, all four of the above cases are present. First, we see a logo used to represent a company:

```
<h1></h1>
```

Next, we see a paragraph which uses a logo right next to the company name, and so doesn't have any alternative text:

```
<article>  
  <h2>News</h2>  
  <p>We have recently been looking at buying the  ABΓ company, a small Greek company  
    specializing in our type of product.</p>
```

In this third snippet, we have a logo being used in an aside, as part of the larger article discussing the acquisition:

```
<aside><p></p></aside>  
<p>The ABΓ company has had a good quarter, and our  
  pie chart studies of their accounts suggest a much bigger blue slice  
  than its green and orange slices, which is always a good sign.</p>  
</article>
```

Finally, we have an opinion piece talking about a logo, and the logo is therefore described in detail in the alternative text.

```
<p>Consider for a moment their logo:</p>
```

```
<p></p>
```

```
<p>How unoriginal can you get? I mean, oooooh, a question mark, how  
<em>revolutionary</em>, how utterly <em>ground-breaking</em>, I'm  
  sure everyone will rush to adopt those specifications now! They could  
  at least have tried for some sort of, I don't know, sequence of
```

rounded squares with varying shades of green and bold white outlines, at least that would look good on the cover of a blue book.</p>

This example shows how the alternative text should be written such that if the image isn't available, and the text is used instead, the text flows seamlessly into the surrounding text, as if the image had never been there in the first place.

4.8.2.1.4 Text that has been rendered to a graphic for typographical effect

Sometimes, an image just consists of text, and the purpose of the image is not to highlight the actual typographic effects used to render the text, but just to convey the text itself.

In such cases, the alt attribute must be present but must consist of the same text as written in the image itself.

Consider a graphic containing the text "Earth Day", but with the letters all decorated with flowers and plants. If the text is merely being used as a header, to spice up the page for graphical users, then the correct alternative text is just the same text "Earth Day", and no mention need be made of the decorations:

```
<h1></h1>
```

4.8.2.1.5 A graphical representation of some of the surrounding text

In many cases, the image is actually just supplementary, and its presence merely reinforces the surrounding text. In these cases, the alt attribute must be present but its value must be the empty string.

In general, an image falls into this category if removing the image doesn't make the page any less useful, but including the image makes it a lot easier for users of visual browsers to understand the concept.

A flowchart that repeats the previous paragraph in graphical form:

```
<p>The network passes data to the Tokeniser stage, which  
passes data to the Tree Construction stage. From there, data goes  
to both the DOM and to Script Execution. Script Execution is  
linked to the DOM, and, using document.write(), passes data to  
the Tokeniser.</p>
```

```
<p></p>
```

In these cases, it would be wrong to include alternative text that consists of just a caption. If a caption is to be included, then either the title attribute can be used, or the figure and legend elements can be used. In the latter case, the image would in fact be a phrase or paragraph with an alternative graphical representation, and would thus require alternative text.

```
<!-- Using the title="" attribute -->  
<p>The network passes data to the Tokeniser stage, which  
passes data to the Tree Construction stage. From there, data goes  
to both the DOM and to Script Execution. Script Execution is  
linked to the DOM, and, using document.write(), passes data to  
the Tokeniser.</p>
```

```

<p></p>  

<!-- Using <figure> and <legend> -->  

<p>The network passes data to the Tokeniser stage, which  

    passes data to the Tree Construction stage. From there, data goes  

    to both the DOM and to Script Execution. Script Execution is  

    linked to the DOM, and, using document.write(), passes data to  

    the Tokeniser.</p>  

<figure>  

    <legend>Flowchart representation of the parsing model.</legend>  

</figure>  

<!-- This is WRONG. Do not do this. Instead, do what the above  

examples do. -->  

<p>The network passes data to the Tokeniser stage, which  

    passes data to the Tree Construction stage. From there, data goes  

    to both the DOM and to Script Execution. Script Execution is  

    linked to the DOM, and, using document.write(), passes data to  

    the Tokeniser.</p>  

<p></p>  

<!-- Never put the image's caption in the alt="" attribute! -->

```

A graph that repeats the previous paragraph in graphical form:

```

<p>According to a study covering several billion pages,  

    about 62% of documents on the Web in 2007 triggered the Quirks  

    rendering mode of Web browsers, about 30% triggered the Almost  

    Standards mode, and about 9% triggered the Standards mode.</p>  

<p></p>

```

4.8.2.1.6 A purely decorative image that doesn't add any information but is still specific to the surrounding content

In general, if an image is decorative but isn't especially page-specific, for example an image that forms part of a site-wide design scheme, the image should be specified in the site's CSS, not in the markup of the document.

However, a decorative image that isn't discussed by the surrounding text still has some relevance can be included in a page using the `img` element. Such images are decorative, but still form part of the content. In these cases, the `alt` attribute must be present but its value must be the empty string.

Examples where the image is purely decorative despite being relevant would include things like a photo of the Black Rock City landscape in a blog post about an event at Burning Man, or an image of a painting inspired by a poem, on a page reciting that

poem. The following snippet shows an example of the latter case (only the first verse is included in this snippet):

```
<h1>The Lady of Shalott</h1>
<p></p>
<p>On either side the river lie<br>
Long fields of barley and of rye,<br>
That clothe the wold and meet the sky;<br>
And through the field the road run by<br>
To many-tower'd Camelot;<br>
And up and down the people go,<br>
Gazing where the lilies blow<br>
Round an island there below,<br>
The island of Shalott.</p>
```

4.8.2.1.7 A group of images that form a single larger picture with no links

When a picture has been sliced into smaller image files that are then displayed together to form the complete picture again, one of the images must have its alt attribute set as per the relevant rules that would be appropriate for the picture as a whole, and then all the remaining images must have their alt attribute set to the empty string.

In the following example, a picture representing a company logo for XYZ Corp has been split into two pieces, the first containing the letters "XYZ" and the second with the word "Corp". The alternative text ("XYZ Corp") is all in the first image.

```
<h1></h1>
```

In the following example, a rating is shown as three filled stars and two empty stars. While the alternative text could have been "★★★☆☆", the author has instead decided to more helpfully give the rating in the form "3/5". That is the alternative text of the first image, and the rest have blank alternative text.

```
<p>Rating: <meter max=5 value=3></meter></p>
```

4.8.2.1.8 A group of images that form a single larger picture with links

Generally, image maps (page 291) should be used instead of slicing an image for links.

However, if an image is indeed sliced and any of the components of the sliced picture are the sole contents of links, then one image per link must have alternative text in its alt attribute representing the purpose of the link.

In the following example, a picture representing the flying spaghetti monster emblem, with each of the left noodly appendages and the right noodly appendages in different images, so that the user can pick the left side or the right side in an adventure.

```
<h1>The Church</h1>
<p>You come across a flying spaghetti monster. Which side of His
Noodliness do you wish to reach out for?</p>
```

```

<p><a href="?go=left" ></a>
    
    <a href="?go=right"></a></p>

```

4.8.2.1.9 A key part of the content

In some cases, the image is a critical part of the content. This could be the case, for instance, on a page that is part of a photo gallery. The image is the whole *point* of the page containing it.

How to provide alternative text for an image that is a key part of the content depends on the image's provenance.

The general case

When it is possible for detailed alternative text to be provided, for example if the image is part of a series of screenshots in a magazine review, or part of a comic strip, or is a photograph in a blog entry about that photograph, text that conveys can serve as a substitute for the image must be given as the contents of the alt attribute.

A screenshot in a gallery of screenshots for a new OS, with some alternative text:

```

<figure>
    
    <legend>Screenshot of a KDE desktop.</legend>
</figure>

```

A graph in a financial report:

```



```

Note that "sales graph" would be inadequate alternative text for a sales graph. Text that would be a good *caption* is not generally suitable as replacement text.

Images that defy a complete description

In certain cases, the nature of the image might be such that providing thorough alternative text is impractical. For example, the image could be indistinct, or could be a complex fractal, or could be a detailed topographical map.

In these cases, the alt attribute must contain some suitable alternative text, but it may be somewhat brief.

Sometimes there simply is no text that can do justice to an image. For example, there is little that can be said to usefully describe a Rorschach inkblot test. However, a description, even if brief, is still better than nothing:

```
<figure>
  
  <legend>A black outline of the first of the ten cards
  in the Rorschach inkblot test.</legend>
</figure>
```

Note that the following would be a very bad use of alternative text:

```
<!-- This example is wrong. Do not copy it. -->
<figure>
  
  <legend>A black outline of the first of the ten cards
  in the Rorschach inkblot test.</legend>
</figure>
```

Including the caption in the alternative text like this isn't useful because it effectively duplicates the caption for users who don't have images, taunting them twice yet not helping them any more than if they had only read or heard the caption once.

Another example of an image that defies full description is a fractal, which, by definition, is infinite in complexity.

The following example shows one possible way of providing alternative text for the full view of an image of the Mandelbrot set.

```

```

Images whose contents are not known

In some unfortunate cases, there might be no alternative text available at all, either because the image is obtained in some automated fashion without any associated alternative text (e.g. a Webcam), or because the page is being generated by a script using user-provided images where the user did not provide suitable or usable alternative text (e.g. photograph sharing sites), or because the author does not himself know what the images represent (e.g. a blind photographer sharing an image on his blog).

In such cases, the alt attribute's value may be omitted, but one of the following conditions must be met as well:

- The title attribute is present and has a non-empty value.
- The img element is in a figure element that contains a legend element that contains content other than inter-element whitespace (page 92).
- The img element is the only img element without an alt attribute in its section (page 142), and its section (page 142) has an associated heading.

Note: Such cases are to be kept to an absolute minimum. If there is even the slightest possibility of the author having the ability to provide real alternative text, then it would not be acceptable to omit the alt attribute.

A photo on a photo-sharing site, if the site received the image with no metadata other than the caption:

```
<figure>

<legend>Bubbles traveled everywhere with us.</legend>
</figure>
```

It could also be marked up like this:

```
<h1>Bubbles traveled everywhere with us.</h1>


```

In either case, though, it would be better if a detailed description of the important parts of the image obtained from the user and included on the page.

A blind user's blog in which a photo taken by the user is shown. Initially, the user might not have any idea what the photo he took shows:

```
<article>
<h1>I took a photo</h1>
<p>I went out today and took a photo!</p>
<figure>

<legend>A photograph taken blindly from my front porch.</legend>
</figure>
</article>
```

Eventually though, the user might obtain a description of the image from his friends and could then include alternative text:

```
<article>
<h1>I took a photo</h1>
```

```

<p>I went out today and took a photo!</p>
<figure>
  
  <legend>A photograph taken blindly from my front porch.</legend>
</figure>
</article>
```

Sometimes the entire point of the image is that a textual description is not available, and the user is to provide the description. For instance, the point of a CAPTCHA image is to see if the user can literally read the graphic. Here is one way to mark up a CAPTCHA (note the title attribute):

```

<p><label>What does this image say?

<input type=text name=captcha></label>
(If you cannot see the image, you can use an <a
href="?audio">audio</a> test instead.)</p>
```

Another example would be software that displays images and asks for alternative text precisely for the purpose of then writing a page with correct alternative text. Such a page could have a table of images, like this:

Images	Descriptions
	<input name="alt2421">
	<input name="alt2422">

Notice that even in this example, as much useful information as possible is still included in the title attribute.

Note: Since some users cannot use images at all (e.g. because they have a very slow connection, or because they are using a text-only browser, or because they are listening to the page being read out by a hands-free automobile voice Web browser, or simply because they are blind), the alt attribute is only allowed to be omitted rather than being provided with replacement text when no alternative text is available and none can be made available, as in the above examples. Lack of

effort from the part of the author is not an acceptable reason for omitting the alt attribute.

4.8.2.1.10 An image not intended for the user

Generally authors should avoid using img elements for purposes other than showing images.

If an img element is being used for purposes other than showing an image, e.g. as part of a service to count page views, then the alt attribute must be the empty string.

4.8.2.1.11 An image in an e-mail or document intended for a specific person who is known to be able to view images

When an image is included in a communication (such as an HTML e-mail) aimed at someone who is known to be able to view images, the alt attribute may be omitted. However, even in such cases it is strongly recommended that alternative text be included (as appropriate according to the kind of image involved, as described in the above entries), so that the e-mail is still usable should the user use a mail client that does not support images, or should the e-mail be forwarded on to other users whose abilities might not include easily seeing images.

4.8.2.1.12 General guidelines

The most general rule for writing alternative text is that the intent is that replacing every image with the text of its alt attribute not change the meaning of the page.

So, in general, alternative text can be written by considering what one would have written had one not been able to include the image.

A corollary to this is that the alt attribute's value should never contain text that could be considered the image's *caption*, *title*, or *legend*. It is supposed to contain replacement text that could be used by users *instead* of the image; it is not meant to supplement the image. The title attribute can be used for supplemental information.

Note: One way to think of alternative text is to think about what how you would read the page containing the image to someone over the phone, without mentioning that there is an image present. Whatever you say instead of the image is typically a good start for writing the alternative text.

4.8.3 The iframe element

Categories

Embedded content (page 94).

Contexts in which this element may be used:

Where embedded content (page 94) is expected.

Content model:

Text that conforms to the requirements given in the prose.

Element-specific attributes:

```
src  
name  
sandbox  
seamless  
width  
height
```

DOM interface:

```
interface HTMLIFrameElement : HTMLElement {  
    attribute DOMString src;  
    attribute DOMString name;  
    attribute DOMString sandbox;  
    attribute boolean seamless;  
    attribute DOMString width;  
    attribute DOMString height;  
};
```

Objects implementing the `HTMLIFrameElement` interface must also implement the `EmbeddingElement` interface defined in the Window Object specification. [WINDOW]

The `iframe` element introduces a new nested browsing context (page 414).

The `src` attribute gives the address of a page that the nested browsing context (page 414) is to contain. The attribute, if present, must be a valid URL (page 52). When the browsing context is created, if the attribute is present, the user agent must navigate (page 473) the element's browsing context to the given URL (page 52), with replacement enabled (page 476), and with the `iframe` element's document's browsing context (page 414) as the source browsing context (page 473). If the user navigates (page 473) away from this page, the `iframe`'s corresponding Window object will reference new Document objects, but the `src` attribute will not change.

Whenever the `src` attribute is set, the nested browsing context (page 414) must be navigated (page 473) to the URL (page 52) given by that attribute's value, with the `iframe` element's document's browsing context (page 414) as the source browsing context (page 473).

If the `src` attribute is not set when the element is created, the browsing context will remain at the initial `about:blank` page.

The `name` attribute, if present, must be a valid browsing context name (page 417). When the browsing context is created, if the attribute is present, the browsing context name (page 417) must be set to the value of this attribute; otherwise, the browsing context name (page 417) must be set to the empty string.

Whenever the `name` attribute is set, the nested browsing context (page 414)'s name (page 417) must be changed to the new value. If the attribute is removed, the browsing context name (page 417) must be set to the empty string.

When content loads in an `iframe`, after any load events are fired within the content itself, the user agent must fire a load event (page 436) at the `iframe` element. When content fails to load (e.g. due to a network error), then the user agent must fire an error event (page 436) at the element instead.

When there is an active parser in the `iframe`, and when anything in the `iframe` that is delaying the load event (page 657) in the `iframe`'s browsing context (page 414), the `iframe` must delay the load event (page 657).

Note: If, during the handling of the load event, the browsing context (page 414) in the `iframe` is again navigated (page 473), that will further delay the load event (page 657).

The `sandbox` attribute, when specified, enables a set of extra restrictions on any content hosted by the `iframe`. Its value must be an unordered set of unique space-separated tokens (page 49). The allowed values are `allow-same-origin`, `allow-forms`, and `allow-scripts`.

While the `sandbox` attribute is specified, the `iframe` element's nested browsing context (page 415), and all the browsing contexts nested (page 415) within it (either directly or indirectly through other nested browsing contexts) must have the following flags set:

The `sandboxed navigation` browsing context flag

This flag prevents content from navigating browsing contexts other than the sandboxed browsing context itself (page 473) (or browsing contexts further nested inside it).

This flag also prevents content from creating new auxiliary browsing contexts (page 418), e.g. using the `target` attribute or the `window.open()` method.

The `sandboxed plugins` browsing context flag

This flag prevents content from instantiating plugins (page 25), whether using the `embed` element (page 222), the `object` element (page 227), the `applet` element (page 674), or through navigation (page 479) of a nested browsing context (page 415).

The `sandboxed annoyances` browsing context flag

This flag prevents content from showing notifications (page 441) outside of the nested browsing context (page 415).

The `sandboxed origin` browsing context flag, unless the `sandbox` attribute's value, when split on spaces (page 49), is found to have the `allow-same-origin` keyword set

This flag forces content into a unique origin (page 424) for the purposes of the same-origin policy (page 423).

This flag also prevents script from reading the `document.cookies` DOM attribute (page 78).

The `allow-same-origin` attribute is intended for two cases.

First, it can be used to allow content from the same site to be sandboxed to disable scripting, while still allowing access to the DOM of the sandboxed content.

Second, it can be used to embed content from a third-party site, sandboxed to prevent that site from opening popup windows, etc, without preventing the embedded page from communicating back to its originating site, using the database APIs to store data, etc.

The `sandboxed forms` browsing context flag, unless the `sandbox` attribute's value, when split on spaces (page 49), is found to have the `allow-forms` keyword set

This flag blocks form submission (page 368).

The `sandboxed scripts` browsing context flag, unless the `sandbox` attribute's value, when split on spaces (page 49), is found to have the `allow-scripts` keyword set

This flag blocks script execution (page 428).

These flags must not be set unless the conditions listed above define them as being set.

In this example, some completely-unknown, potentially hostile, user-provided HTML content is embedded in a page. Because it is sandboxed, it is treated by the user agent as being from a unique origin, despite the content being served from the same site. Thus it is affected by all the normal cross-site restrictions. In addition, the embedded page has scripting disabled, plugins disabled, forms disabled, and it cannot navigate any frames or windows other than itself (or any frames or windows it itself embeds).

```
<p>We're not scared of you! Here is your content, unedited:</p>
<iframe sandbox src="getusercontent.cgi?id=12193"></iframe>
```

Note that cookies are still sent to the server in the `getusercontent.cgi` request, though they are not visible in the `document.cookies` DOM attribute.

In this example, a gadget from another site is embedded. The gadget has scripting and forms enabled, and the origin sandbox restrictions are lifted, allowing the gadget to communicate with its originating server. The sandbox is still useful, however, as it disables plugins and popups, thus reducing the risk of the user being exposed to malware and other annoyances.

```
<iframe sandbox="allow-same-origin allow-forms allow-scripts"
src="http://maps.example.com/embedded.html"></iframe>
```

The `seamless` attribute is a boolean attribute. When specified, it indicates that the `iframe` element's browsing context (page 414) is to be rendered in a manner that makes it appear to be part of the containing document (seamlessly included in the parent document). Specifically, when the attribute is set on an element and while the browsing context (page 414)'s active document (page 414) has the same origin (page 426) as the `iframe` element's document, or the browsing context (page 414)'s active document (page 414)'s `address` has the same origin (page 426) as the `iframe` element's document, the following requirements apply:

- The user agent must set the **seamless browsing context flag** to true for that browsing context (page 414). This will cause links to open in the parent browsing context (page 473).
- In a CSS-supporting user agent: the user agent must add all the style sheets that apply to the `iframe` element to the cascade of the active document (page 414) of the

`iframe` element's nested browsing context (page 415), at the appropriate cascade levels, before any style sheets specified by the document itself.

- In a CSS-supporting user agent: the user agent must, for the purpose of CSS property inheritance only, treat the root element of the active document (page 414) of the `iframe` element's nested browsing context (page 415) as being a child of the `iframe` element. (Thus inherited properties on the root element of the document in the `iframe` will inherit the computed values of those properties on the `iframe` element instead of taking their initial values.)
- In visual media, in a CSS-supporting user agent: the user agent should set the intrinsic width of the `iframe` to the width that the element would have if it was a non-replaced block-level element with 'width: auto'.
- In visual media, in a CSS-supporting user agent: the user agent should set the intrinsic height of the `iframe` to the height of the bounding box around the content rendered in the `iframe` at its current width.
- In visual media, in a CSS-supporting user agent: the user agent must force the height of the initial containing block of the active document (page 414) of the nested browsing context (page 415) of the `iframe` to zero.

Note: This is intended to get around the otherwise circular dependency of percentage dimensions that depend on the height of the containing block, thus affecting the height of the document's bounding box, thus affecting the height of the viewport, thus affecting the size of the initial containing block.

- In speech media, the user agent should render the nested browsing context (page 415) without announcing that it is a separate document.
- User agents should, in general, act as if the active document (page 414) of the `iframe`'s nested browsing context (page 415) was part of the document that the `iframe` is in.

For example if the user agent supports listing all the links in a document, links in "seamlessly" nested documents would be included in that list without being significantly distinguished from links in the document itself.

** Parts of the above might get moved into the rendering section at some point.

If the attribute is not specified, or if the origin (page 423) conditions listed above are not met, then the user agent should render the nested browsing context (page 415) in a manner that is clearly distinguishable as a separate browsing context (page 414), and the seamless browsing context flag (page 219) must be set to false for that browsing context (page 414).

⚠Warning! It is important that user agents recheck the above conditions whenever the active document (page 414) of the nested browsing context (page 415) of the `iframe` changes, such that the seamless browsing context flag (page 219) gets unset if the nested browsing context (page 415) is navigated (page 473) to another origin.

In this example, the site's navigation is embedded using a client-side include using an `iframe`. Any links in the `iframe` will, in new user agents, be automatically opened in the `iframe`'s parent browsing context; for legacy user agents, the site could also include a `base` element with a `target` attribute with the value `_parent`. Similarly, in new user agents the styles of the parent page will be automatically applied to the contents of the frame, but to support legacy user agents authors might wish to include the styles explicitly.

```
<nav><iframe seamless src="nav.include.html"></iframe></nav>
```

The `iframe` element supports dimension attributes (page 295) for cases where the embedded content has specific dimensions (e.g. ad units have well-defined dimensions).

An `iframe` element never has fallback content (page 95), as it will always create a nested browsing context (page 414), regardless of whether the specified initial contents are successfully used.

Descendants of `iframe` elements represent nothing. (In legacy user agents that do not support `iframe` elements, the contents would be parsed as markup that could act as fallback content.)

- ** The content model of `iframe` elements is text, except that the text must be such that ...
- ** anyone have any bright ideas?

Note: *The HTML parser (page 582) treats markup inside `iframe` elements as text.*

The DOM attributes `src`, `name`, `sandbox`, and `seamless` must reflect (page 67) the respective content attributes of the same name.

4.8.4 The `embed` element

Categories

Embedded content (page 94).

Contexts in which this element may be used:

Where embedded content (page 94) is expected.

Content model:

Empty.

Element-specific attributes:

`src`
`type`
`width`
`height`

Any other attribute that has no namespace (see prose).

DOM interface:

```
interface HTMLEmbedElement : HTMLElement {  
    attribute DOMString src;  
    attribute DOMString type;  
    attribute DOMString width;  
    attribute DOMString height;  
};
```

Depending on the type of content instantiated by the embed element, the node may also support other interfaces.

The embed element represents an integration point for an external (typically non-HTML) application or interactive content.

The **src** attribute gives the address of the resource being embedded. The attribute, if present, must contain a valid URL (page 52).

The **type** attribute, if present, gives the MIME type of the plugin to instantiate. The value must be a valid MIME type, optionally with parameters. If both the type attribute and the src attribute are present, then the type attribute must specify the same type as the explicit Content-Type metadata (page 60) of the resource given by the src attribute. [RFC2046]

When the element is created with neither a src attribute nor a type attribute, and when attributes are removed such that neither attribute is present on the element anymore, any plugins instantiated for the element must be removed, and the embed element represents nothing.

When the sandboxed plugins browsing context flag (page 218) is set on the browsing context (page 414) for which the embed element's document is the active document (page 414), then the user agent must render the embed element in a manner that conveys that the plugin (page 25) was disabled. The user agent may offer the user the option to override the sandbox and instantiate the plugin (page 25) anyway; if the user invokes such an option, the user agent must act as if the sandboxed plugins browsing context flag (page 218) was not set for the purposes of this element.

⚠Warning! Plugins are disabled in sandboxed browsing contexts because they might not honor the restrictions imposed by the sandbox (e.g. they might allow scripting even when scripting in the sandbox is disabled). User agents should convey the danger of overriding the sandbox to the user if an option to do so is provided.

When the element is created with a src attribute, and whenever the src attribute is subsequently set, and whenever the type attribute is set or removed while the element has a src attribute, if the element is not in a sandboxed browsing context, user agents should fetch (page 59) the specified resource, find and instantiate an appropriate plugin (page 25) based on the content's type (page 222), and hand that plugin (page 25) the content of the resource, replacing any previously instantiated plugin for the element.

Fetching the resource must delay the load event (page 657).

The **type of the content** being embedded is defined as follows:

1. If the element has a type attribute, then the value of the type attribute is the content's type.
2. Otherwise, if the <path> (page 54) component of the URL (page 52) of the specified resource matches a pattern that a plugin (page 25) supports, then the content's type is the type that that plugin can handle.

For example, a plugin might say that it can handle resources with <path> (page 54) components that end with the four character string ".swf".

** It would be better if browsers didn't do extension sniffing like this, and only based their decision on the actual contents of the resource. Couldn't we just apply the sniffed type of a resource steps?

3. Otherwise, if the specified resource has explicit Content-Type metadata (page 60), then that is the content's type.
4. Otherwise, the content has no type and there can be no appropriate plugin (page 25) for it.

Whether the resource is fetched successfully or not (e.g. whether the response code was a 2xx code or equivalent) must be ignored when determining the resource's type and when handing the resource to the plugin.

Note: This allows servers to return data for plugins even with error responses (e.g. HTTP 500 Internal Server Error codes can still contain plugin data).

When the element is created with a type attribute and no src attribute, and whenever the type attribute is subsequently set, so long as no src attribute is set, and whenever the src attribute is removed when the element has a type attribute, if the element is not in a sandboxed browsing context, user agents should find and instantiate an appropriate plugin (page 25) based on the value of the type attribute.

Any (namespace-less) attribute may be specified on the embed element, so long as its name is XML-compatible (page 24) and contains no characters in the range U+0041 .. U+005A (LATIN CAPITAL LETTER A LATIN CAPITAL LETTER Z).

Note: All attributes in HTML documents (page 76) get lowercased automatically, so the restriction on uppercase letters doesn't affect such documents.

The user agent should pass the names and values of all the attributes of the embed element that have no namespace to the plugin (page 25) used, when it is instantiated.

If the plugin (page 25) instantiated for the embed element supports a scriptable interface, the HTMLEmbedElement object representing the element should expose that interface while the element is instantiated.

The embed element has no fallback content (page 95). If the user agent can't find a suitable plugin, then the user agent must use a default plugin. (This default could be as simple as saying "Unsupported Format".)

The embed element supports dimension attributes (page 295).

The DOM attributes **src** and **type** each must reflect (page 67) the respective content attributes of the same name.

4.8.5 The object element

Categories

Embedded content (page 94).

Contexts in which this element may be used:

Where embedded content (page 94) is expected.

Content model:

Zero or more param elements, then, transparent (page 96).

Element-specific attributes:

data
type
name
usemap
width
height

DOM interface:

```
interface HTMLObjectElement : HTMLElement {  
    attribute DOMString data;  
    attribute DOMString type;  
    attribute DOMString name;  
    attribute DOMString useMap;  
    attribute DOMString width;  
    attribute DOMString height;  
};
```

Objects implementing the `HTMLObjectElement` interface must also implement the `EmbeddingElement` interface defined in the Window Object specification. [WINDOW]

Depending on the type of content instantiated by the `object` element, the node may also support other interfaces.

The `object` element can represent an external resource, which, depending on the type of the resource, will either be treated as an image, as a nested browsing context (page 414), or as an external resource to be processed by a plugin (page 25).

The **data** attribute, if present, specifies the address of the resource. If present, the attribute must be a valid URL (page 52).

The **type** attribute, if present, specifies the type of the resource. If present, the attribute must be a valid MIME type, optionally with parameters. [RFC2046]

One or both of the `data` and `type` attributes must be present.

The **name** attribute, if present, must be a valid browsing context name (page 417).

When the element is created, and subsequently whenever the **classid** attribute changes or is removed, or, if the **classid** attribute is not present, whenever the **data** attribute changes or is removed, or, if neither **classid** attribute nor the **data** attribute are present, whenever the **type** attribute changes or is removed, the user agent must run the following steps to determine what the object element represents:

1. If the **classid** attribute is present, and has a value that isn't the empty string, then: if the user agent can find a plugin (page 25) suitable according to the value of the **classid** attribute, and plugins aren't being sandboxed (page 227), then that plugin (page 25) should be used (page 227), and the value of the **data** attribute, if any, should be passed to the plugin (page 25). If no suitable plugin (page 25) can be found, or if the plugin (page 25) reports an error, jump to the last step in the overall set of steps (fallback).
2. If the **data** attribute is present, then:
 1. If the **type** attribute is present and its value is not a type that the user agent supports, and is not a type that the user agent can find a plugin (page 25) for, then the user agent may jump to the last step in the overall set of steps (fallback) without fetching the content to examine its real type.
 2. Fetch (page 59) the resource specified by the **data** attribute.

The fetching of the resource must delay the **load** event (page 657).

3. If the resource is not yet available (e.g. because the resource was not available in the cache, so that loading the resource required making a request over the network), then jump to the last step in the overall set of steps (fallback). When the resource becomes available, or if the load fails, restart this algorithm from this step. Resources can load incrementally; user agents may opt to consider a resource "available" whenever enough data has been obtained to begin processing the resource.
4. If the load failed (e.g. an HTTP 404 error, a DNS error), fire an **error** event (page 436) at the element, then jump to the last step in the overall set of steps (fallback).
5. Determine the *resource type*, as follows:
 1. Let the *resource type* be unknown.
 2. If the resource has associated Content-Type metadata (page 60), then let the *resource type* be the type specified in the resource's Content-Type metadata (page 60).
 3. If the *resource type* is unknown or "application/octet-stream" and there is a **type** attribute present on the object element, then change the *resource type* to instead be the type specified in that **type** attribute.

Otherwise, if the *resource type* is "application/octet-stream" but there is no **type** attribute on the object element, then change the

resource type to be unknown, so that the sniffing rules in the next step are invoked.

4. If the *resource type* is still unknown, then change the *resource type* to instead be the sniffed type of the resource (page 61).
6. Handle the content as given by the first of the following cases that matches:

↳ **If the *resource type* can be handled by a plugin (page 25) and plugins aren't being sandboxed (page 227)**

The user agent should use that plugin (page 227) and pass the content of the resource to that plugin (page 25). If the plugin (page 25) reports an error, then jump to the last step in the overall set of steps (fallback).

↳ **If the *resource type* is an XML MIME type**

↳ **If the *resource type* is HTML**

↳ **If the *resource type* does not start with "image/"**

The object element must be associated with a nested browsing context (page 414), if it does not already have one. The element's nested browsing context (page 414) must then be navigated (page 473) to the given resource, with replacement enabled (page 476), and with the object element's document's browsing context (page 414) as the source browsing context (page 473). (The data attribute of the object element doesn't get updated if the browsing context gets further navigated to other locations.)

If the name attribute is present, the browsing context name (page 417) must be set to the value of this attribute; otherwise, the browsing context name (page 417) must be set to the empty string.

**
**

navigation might end up treating it as something else, because it can do sniffing. how should we handle that?

↳ **If the *resource type* starts with "image/", and support for images has not been disabled**

Apply the image sniffing (page 65) rules to determine the type of the image.

The object element represents the specified image. The image is not a nested browsing context (page 414).

If the image cannot be rendered, e.g. because it is malformed or in an unsupported format, jump to the last step in the overall set of steps (fallback).

↳ **Otherwise**

The given *resource type* is not supported. Jump to the last step in the overall set of steps (fallback).

7. The element's contents are not part of what the object element represents.

8. Once the resource is completely loaded, fire a `load` event (page 436) at the element.
3. If the `data` attribute is absent but the `type` attribute is present, plugins aren't being sandboxed (page 227), and the user agent can find a plugin (page 25) suitable according to the value of the `type` attribute, then that plugin (page 25) should be used (page 227). If no suitable plugin (page 25) can be found, or if the plugin (page 25) reports an error, jump to the next step (fallback).
4. (Fallback.) The `object` element represents what the element's contents represent, ignoring any leading `param` element children. This is the element's fallback content (page 95).

When the algorithm above instantiates a plugin (page 25), the user agent should pass the names and values of all the parameters given by `param` elements that are children of the `object` element to the plugin (page 25) used. If the plugin (page 25) supports a scriptable interface, the `HTMLObjectElement` object representing the element should expose that interface. The plugin (page 25) is not a nested browsing context (page 414).

If the sandboxed plugins browsing context flag (page 218) is set on the browsing context (page 414) for which the `object` element's document is the active document (page 414), then the steps above must always act as if they had failed to find a plugin (page 25), even if one would otherwise have been used.

Due to the algorithm above, the contents of `object` elements act as fallback content (page 95), used only when referenced resources can't be shown (e.g. because it returned a 404 error). This allows multiple `object` elements to be nested inside each other, targeting multiple user agents with different capabilities, with the user agent picking the first one it supports.

Whenever the `name` attribute is set, if the `object` element has a nested browsing context (page 414), its `name` (page 417) must be changed to the new value. If the attribute is removed, if the `object` element has a browsing context (page 414), the browsing context name (page 417) must be set to the empty string.

The `usemap` attribute, if present while the `object` element represents an image, can indicate that the object has an associated image map (page 291). The attribute must be ignored if the `object` element doesn't represent an image.

The `object` element supports dimension attributes (page 295).

The DOM attributes `data`, `type`, `name`, and `useMap` each must reflect (page 67) the respective content attributes of the same name.

In the following example, a Java applet is embedded in a page using the `object` element. (Generally speaking, it is better to avoid using applets like these and instead use native JavaScript and HTML to provide the functionality, since that way the application will work on all Web browsers without requiring a third-party plugin. Many devices, especially embedded devices, do not support third-party technologies like Java.)

```
<figure>
<object type="application/x-java-applet">
  <param name="code" value="MyJavaClass">
  <p>You do not have Java available, or it is disabled.</p>
</object>
```

```
<legend>My Java Clock</legend>
</figure>
```

In this example, an HTML page is embedded in another using the object element.

```
<figure>
<object data="clock.html"></object>
<legend>My HTML Clock</legend>
</figure>
```

4.8.6 The param element

Categories

None.

Contexts in which this element may be used:

As a child of an object element, before any flow content (page 93).

Content model:

Empty.

Element-specific attributes:

name
value

DOM interface:

```
interface HTMLParamElement : HTMLElement {
    attribute DOMString name;
    attribute DOMString value;
};
```

The param element defines parameters for plugins invoked by object elements.

The **name** attribute gives the name of the parameter.

The **value** attribute gives the value of the parameter.

Both attributes must be present. They may have any value.

If both attributes are present, and if the parent element of the param is an object element, then the element defines a **parameter** with the given name/value pair.

The DOM attributes **name** and **value** must both reflect (page 67) the respective content attributes of the same name.

4.8.7 The video element

Categories

Embedded content (page 94).

If the element has a **controls** attribute: Interactive content (page 95).

Contexts in which this element may be used:

Where embedded content (page 94) is expected.

Content model:

If the element has a `src` attribute: transparent (page 96).

If the element does not have a `src` attribute: one or more source elements, then, transparent (page 96).

Element-specific attributes:

`src`
`poster`
`autoplay`
`start`
`loopstart`
`loopend`
`end`
`playcount`
`controls`
`width`
`height`

DOM interface:

```
interface HTMLVideoElement : HTMLMediaElement {  
    attribute DOMString width;  
    attribute DOMString height;  
    readonly attribute unsigned long videoWidth;  
    readonly attribute unsigned long videoHeight;  
    attribute DOMString poster;  
};
```

A video element represents a video or movie.

Content may be provided inside the video element. User agents should not show this content to the user; it is intended for older Web browsers which do not support video, so that legacy video plugins can be tried, or to show text to the users of these older browser informing them of how to access the video contents.

Note: In particular, this content is not fallback content (page 95) intended to address accessibility concerns. To make video content accessible to the blind, deaf, and those with other physical or cognitive disabilities, authors are expected to provide alternative media streams and/or to embed accessibility aids (such as caption or subtitle tracks) into their media streams.

The video element is a media element (page 236) whose media data (page 238) is ostensibly video data, possibly with associated audio data.

The `src`, `autoplay`, `start`, `loopstart`, `loopend`, `end`, `playcount`, and `controls` attributes are the attributes common to all media elements (page 238).

The **poster** attribute gives the address of an image file that the user agent can show while no video data is available. The attribute, if present, must contain a valid URL (page 52). If the specified resource is to be used, it must be fetched (page 59) when the element is created or when the poster attribute is set. The **poster frame** is then the image obtained from that resource, if any.

Note: *The image given by the poster attribute, the poster frame (page 230), is intended to be a representative frame of the video (typically one of the first non-blank frames) that gives the user an idea of what the video is like.*

The **poster** DOM attribute must reflect (page 67) the poster content attribute.

When no video data is available (the element's readyState attribute is either HAVE NOTHING or HAVE_METADATA), video elements represent either the poster frame (page 230), or nothing.

When a video element is paused (page 250) and the current playback position (page 246) is the first frame of video, the element represents either the frame of video corresponding to the current playback position (page 246) or the poster frame (page 230), at the discretion of the user agent.

Notwithstanding the above, the poster frame (page 230) should be preferred over nothing, but the poster frame (page 230) should not be shown again after a frame of video has been shown.

When a video element is paused (page 250) at any other position, the element represents the frame of video corresponding to the current playback position (page 246), or, if that is not yet available (e.g. because the video is seeking or buffering), the last frame of the video to have been rendered.

When a video element is potentially playing (page 250), it represents the frame of video at the continuously increasing "current" position (page 246). When the current playback position (page 246) changes such that the last frame rendered is no longer the frame corresponding to the current playback position (page 246) in the video, the new frame must be rendered. Similarly, any audio associated with the video must, if played, be played synchronized with the current playback position (page 246), at the specified volume (page 256) with the specified mute state (page 256).

When a video element is neither potentially playing (page 250) nor paused (page 250) (e.g. when seeking or stalled), the element represents the last frame of the video to have been rendered.

Note: *Which frame in a video stream corresponds to a particular playback position is defined by the video stream's format.*

In addition to the above, the user agent may provide messages to the user (such as "buffering", "no video loaded", "error", or more detailed information) by overlaying text or icons on the video or other areas of the element's playback area, or in another appropriate manner.

User agents that cannot render the video may instead make the element represent a link to an external video playback utility or to the video data itself.

The intrinsic width and height of the video are the aspect-ratio corrected dimensions given by the video data itself: the **intrinsic width** is the number of pixels per line of the video data multiplied by the pixel ratio given by the resource, multiplied by the resolution of the resource; the **intrinsic height** is the number of pixels per column of the video data multiplied by the resolution of the resource. The **resolution of the resource** is the physical distance intended for each pixel of video data, and assumes square pixels, with the resource's pixel ratio then adjusting the width of the pixels to the actual aspect-ratio-corrected width. In the absence of resolution information defining the mapping of pixels in the video to physical dimensions, user agents should assume that one pixel in the video corresponds to one CSS pixel. The **pixel ratio of the resource** is the corrected aspect ratio of the video divided by the ratio of the number of pixels per line to the number of pixels per column. In the absence of pixel ratio information in the resource, user agents should assume a default of 1.0 (square pixels).

The **videoWidth** DOM attribute must return the intrinsic width (page 231) of the video in CSS pixels. The **videoHeight** DOM attribute must return the intrinsic height (page 231) of the video in CSS pixels. If no video data is available, then the attributes must return 0.

If the video's pixel ratio override (page 240)'s is *none*, then the video's **adjusted width** is the same as the video's intrinsic width (page 231). If the video has a pixel ratio override (page 240) other than *none*, then the adjusted width of the video is the number of pixels per line of the video data multiplied by the video's pixel ratio override (page 240), multiplied by the resolution of the resource (page 231); the pixel ratio of the resource (page 231) is thus ignored.

The video's **adjusted height** is the same as the video's intrinsic height (page 231).

The **adjusted aspect ratio** of a video is the ratio of its adjusted width (page 231) to its adjusted height (page 231).

User agents may adjust the adjusted width (page 231) and height (page 231) of the video to ensure that each pixel of video data corresponds to at least one device pixel, so long as this doesn't affect the adjusted aspect ratio (page 231) (this is especially relevant for pixel ratios that are less than 1.0).

The **video** element supports dimension attributes (page 295).

Video content should be rendered inside the element's playback area such that the video content is shown centered in the playback area at the largest possible size that fits completely within it, with the video content's adjusted aspect ratio (page 231) being preserved. Thus, if the aspect ratio of the playback area does not match the adjusted aspect ratio (page 231) of the video, the video will be shown letterboxed. Areas of the element's playback area that do not contain the video represent nothing.

The intrinsic width of a video element's playback area is the adjusted width (page 231) of the video resource, if that is available; otherwise it is the intrinsic width of the poster frame (page 230), if that is available; otherwise it is 300 CSS pixels.

The intrinsic height of a video element's playback area is the intrinsic height (page 231) of the video resource, if that is available; otherwise it is the intrinsic height of the poster frame (page 230), if that is available; otherwise it is 150 CSS pixels.

Note: The poster frame (page 230) is not affected by the pixel ratio conversions.

User agents should provide controls to enable or disable the display of closed captions associated with the video stream, though such features should, again, not interfere with the page's normal rendering.

User agents may allow users to view the video content in manners more suitable to the user (e.g. full-screen or in an independent resizable window). As for the other user interface features, controls to enable this should not interfere with the page's normal rendering unless the user agent is exposing a user interface (page 256). In such an independent context, however, user agents may make full user interfaces visible, with, e.g., play, pause, seeking, and volume controls, even if the controls attribute is absent.

User agents may allow video playback to affect system features that could interfere with the user's experience; for example, user agents could disable screensavers while video playback is in progress.

⚠ Warning! User agents should not provide a public API to cause videos to be shown full-screen. A script, combined with a carefully crafted video file, could trick the user into thinking a system-modal dialog had been shown, and prompt the user for a password. There is also the danger of "mere" annoyance, with pages launching full-screen videos when links are clicked or pages navigated. Instead, user-agent specific interface features may be provided to easily allow the user to obtain a full-screen playback mode.

- ** The spec does not currently define the interaction of the "controls" attribute with the "height" and "width" attributes. This will likely be defined in the rendering section based on implementation experience. So far, browsers seem to be making the controls overlay-only, thus somewhat sidestepping the issue.

4.8.7.1 Video and audio codecs for video elements

User agents may support any video and audio codecs and container formats.

- ** It would be helpful for interoperability if all browsers could support the same codecs. However, there are no known codecs that satisfy all the current players: we need a codec that is known to not require per-unit or per-distributor licensing, that is compatible with the open source development model, that is of sufficient quality as to be usable, and that is not an additional submarine patent risk for large companies. This is an ongoing issue and this section will be updated once more information is available.

Note: Certain user agents might support no codecs at all, e.g. text browsers running over SSH connections.

4.8.8 The audio element

Categories

Embedded content (page 94).

If the element has a controls attribute: Interactive content (page 95).

Contexts in which this element may be used:

Where embedded content (page 94) is expected.

Content model:

If the element has a src attribute: transparent (page 96).

If the element does not have a src attribute: one or more source elements, then, transparent (page 96).

Element-specific attributes:

src
autoplay
start
loopstart
loopend
end
playcount
controls

DOM interface:

```
[NamedConstructor=Audio(),
 NamedConstructor=Audio(in DOMString src)]
interface HTMLAudioElement : HTMLMediaElement {
    // no members
};
```

An audio element represents a sound or audio stream.

Content may be provided inside the audio element. User agents should not show this content to the user; it is intended for older Web browsers which do not support audio, so that legacy audio plugins can be tried, or to show text to the users of these older browser informing them of how to access the audio contents.

Note: In particular, this content is not fallback content (page 95) intended to address accessibility concerns. To make audio content accessible to the deaf or to those with other physical or cognitive disabilities, authors are expected to provide alternative media streams and/or to embed accessibility aids (such as transcriptions) into their media streams.

The audio element is a media element (page 236) whose media data (page 238) is ostensibly audio data.

The src, autoplay, start, loopstart, loopend, end, playcount, and controls attributes are the attributes common to all media elements (page 238).

When an audio element is potentially playing (page 250), it must have its audio data played synchronized with the current playback position (page 246), at the specified volume (page 256) with the specified mute state (page 256).

When an audio element is not potentially playing (page 250), audio must not play for the element.

Two constructors are provided for creating `HTMLAudioElement` objects (in addition to the factory methods from DOM Core such as `createElement()`): `Audio()` and `Audio(src)`. When invoked as constructors, these must return a new `HTMLAudioElement` object (a new audio element). If the `src` argument is present, the object created must have its `src` content attribute set to the provided value, and the user agent must invoke the `load()` method on the object before returning.

4.8.8.1 Audio codecs for audio elements

User agents may support any audio codecs and container formats.

User agents must support the WAVE container format with audio encoded using the PCM format.

4.8.9 The source element

Categories

None.

Contexts in which this element may be used:

As a child of a media element (page 236), before any flow content (page 93).

Content model:

Empty.

Element-specific attributes:

`src`
`type`
`media`
`pixelratio`

DOM interface:

```
interface HTMLSourceElement : HTMLElement {
    attribute DOMString src;
    attribute DOMString type;
    attribute DOMString media;
    attribute float pixelRatio;
};
```

The `source` element allows authors to specify multiple media resources (page 238) for media elements (page 236).

The `src` attribute gives the address of the media resource (page 238). The value must be a valid URL (page 52). This attribute must be present.

The **type** attribute gives the type of the media resource (page 238), to help the user agent determine if it can play this media resource (page 238) before fetching it. Its value must be a MIME type. The **codecs** parameter may be specified and might be necessary to specify exactly how the resource is encoded. [RFC2046] [RFC4281]

The following list shows some examples of how to use the **codecs=MIME** parameter in the **type** attribute.

H.264 Simple baseline profile video (main and extended video compatible) level 3 and Low-Complexity AAC audio in MP4 container

```
<source src="video.mp4" type="video/mp4; codecs="avc1.42E01E, mp4a.40.2" ;>
```

H.264 Extended profile video (baseline-compatible) level 3 and Low-Complexity AAC audio in MP4 container

```
<source src="video.mp4" type="video/mp4; codecs="avc1.58A01E, mp4a.40.2" ;>
```

H.264 Main profile video level 3 and Low-Complexity AAC audio in MP4 container

```
<source src="video.mp4" type="video/mp4; codecs="avc1.4D401E, mp4a.40.2" ;>
```

H.264 "High" profile video (incompatible with main, baseline, or extended profiles) level 3 and Low-Complexity AAC audio in MP4 container

```
<source src="video.mp4" type="video/mp4; codecs="avc1.64001E, mp4a.40.2" ;>
```

MPEG-4 Visual Simple Profile Level 0 video and Low-Complexity AAC audio in MP4 container

```
<source src="video.mp4" type="video/mp4; codecs="mp4v.20.8, mp4a.40.2" ;>
```

MPEG-4 Advanced Simple Profile Level 0 video and Low-Complexity AAC audio in MP4 container

```
<source src="video.mp4" type="video/mp4; codecs="mp4v.20.240, mp4a.40.2" ;>
```

MPEG-4 Visual Simple Profile Level 0 video and AMR audio in 3GPP container

```
<source src="video.3gp" type="video/3gpp; codecs="mp4v.20.8, samr" ;>
```

Theora video and Vorbis audio in Ogg container

```
<source src="video.ogv" type="video/ogg; codecs="theora, vorbis" ;>
```

Theora video and Speex audio in Ogg container

```
<source src="video.ogv" type="video/ogg; codecs="theora, speex" ;>
```

Vorbis audio alone in Ogg container

```
<source src="audio.ogg" type="audio/ogg; codecs=vorbis">
```

Speex audio alone in Ogg container

```
<source src="audio.spx" type="audio/ogg; codecs=speex">
```

FLAC audio alone in Ogg container

```
<source src="audio.oga" type="audio/ogg; codecs=flac">
```

Dirac video and Vorbis audio in Ogg container

```
<source src="video.ogv" type="video/ogg; codecs="dirac, vorbis"">
```

Theora video and Vorbis audio in Matroska container

```
<source src="video.mkv" type="video/x-matroska; codecs="theora, vorbis"">
```

The **media** attribute gives the intended media type of the media resource (page 238), to help the user agent determine if this media resource (page 238) is useful to the user before fetching it. Its value must be a valid media query (page 30). [MQ]

Either the **type** attribute, the **media** attribute or both, must be specified, unless this is the last source element child of the parent element.

The **pixelratio** attribute allows the author to specify the pixel ratio of anamorphic media resources (page 238) that do not self-describe their pixel ratio (page 231). The attribute value, if specified, must be a valid floating point number (page 34) giving the ratio of the correct rendered width of each pixel to the actual height of each pixel in the image. The default value, if the attribute is omitted or cannot be parsed, is 1.0.

Note: *The only way this default is used is in deciding what number the pixelRatio DOM attribute will return if the content attribute is omitted or cannot be parsed. If the content attribute is omitted or cannot be parsed, then the user agent doesn't adjust the intrinsic width (page 231) of the video at all; the intrinsic dimensions and the pixel ratio (page 231) of the video are honoured.*

If a source element is inserted into a media element (page 236) that is already in a document and whose **networkState** is in the **NETWORK_EMPTY** state, the user agent must queue a task (page 429) that implicitly invokes the **load()** method on the media element (page 236), and ignores any resulting exceptions. The task source (page 429) for this task is the media element (page 236)'s own media element new resource task source (page 238).

The DOM attributes **src**, **type**, and **media** must reflect (page 67) the respective content attributes of the same name.

The DOM attribute **pixelRatio** must reflect (page 67) the **pixelratio** content attribute.

4.8.10 Media elements

Media elements implement the following interface:

```
interface HTMLMediaElement : HTMLElement {  
  
    // error state  
    readonly attribute MediaError error;  
  
    // network state
```

```

        attribute DOMString src;
readonly attribute DOMString currentSrc;
const unsigned short NETWORK_EMPTY = 0;
const unsigned short NETWORK_IDLE = 1;
const unsigned short NETWORK_LOADING = 2;
const unsigned short NETWORK_LOADED = 3;
readonly attribute unsigned short networkState;
readonly attribute float bufferingRate;
readonly attribute boolean bufferingThrottled;
readonly attribute TimeRanges buffered;
readonly attribute ByteRanges bufferedBytes;
readonly attribute unsigned long totalBytes;
void load();

// ready state
const unsigned short HAVE NOTHING = 0;
const unsigned short HAVE_METADATA = 1;
const unsigned short HAVE SOME DATA = 2;
const unsigned short HAVE CURRENT DATA = 3;
const unsigned short HAVE FUTURE DATA = 4;
const unsigned short HAVE ENOUGH DATA = 5;
readonly attribute unsigned short readyState;
readonly attribute boolean seeking;

// playback state
        attribute float currentTime;
readonly attribute float duration;
readonly attribute boolean paused;
        attribute float defaultPlaybackRate;
        attribute float playbackRate;
readonly attribute TimeRanges played;
readonly attribute TimeRanges seekable;
readonly attribute boolean ended;
        attribute boolean autoplay;
void play();
void pause();

// looping
        attribute float start;
        attribute float end;
        attribute float loopStart;
        attribute float loopEnd;
        attribute unsigned long playCount;
        attribute unsigned long currentLoop;

// cue ranges
void addCueRange(in DOMString className, in DOMString id, in float
start, in float end, in boolean pauseOnExit, in CueRangeCallback
enterCallback, in CueRangeCallback exitCallback);
void removeCueRanges(in DOMString className);

```

```

    // controls
        attribute boolean controls;
        attribute float volume;
        attribute boolean muted;
    };

** // CueRangeCallback waiting on WebIDL

```

The **media element attributes**, `src`, `autoplay`, `start`, `loopstart`, `lopend`, `end`, `playcount`, and `controls`, apply to all media elements (page 236). They are defined in this section.

Media elements (page 236) are used to present audio data, or video and audio data, to the user. This is referred to as **media data** in this section, since this section applies equally to media elements (page 236) for audio or for video. The term **media resource** is used to refer to the complete set of media data, e.g. the complete video file, or complete audio file.

Media elements (page 236) use two task queues (page 429), the **media element event task source** for asynchronous events and callbacks, and the **media element new resource task source** for handling implicit loads. Unless otherwise specified, the task source (page 429) for all the tasks queued (page 429) in this section and its subsections is the media element event task source (page 238).

The `canPlayType()` method can be used to probe the user agent to determine what types are supported.

4.8.10.1 Error codes

All media elements (page 236) have an associated error status, which records the last error the element encountered since the `load()` method was last invoked. The **error** attribute, on getting, must return the `MediaError` object created for this last error, or null if there has not been an error.

```

interface MediaError {
    const unsigned short MEDIA_ERR_ABORTED = 1;
    const unsigned short MEDIA_ERR_NETWORK = 2;
    const unsigned short MEDIA_ERR_DECODE = 3;
    readonly attribute unsigned short code;
};

```

The **code** attribute of a `MediaError` object must return the code for the error, which must be one of the following:

MEDIA_ERR_ABORTED (numeric value 1)

The fetching process for the media resource (page 238) was aborted by the user agent at the user's request.

MEDIA_ERR_NETWORK (numeric value 2)

A network error of some description caused the user agent to stop fetching the media resource (page 238).

MEDIA_ERR_DECODE (numeric value 3)

An error of some description occurred while decoding the media resource (page 238).

4.8.10.2 Location of the media resource

The **src** content attribute on media elements (page 236) gives the address of the media resource (video, audio) to show. The attribute, if present, must contain a valid URL (page 52).

If the **src** attribute of a media element (page 236) that is already in a document and whose **networkState** is in the **NETWORK_EMPTY** state is added, changed, or removed, the user agent must queue a task (page 429) that implicitly invokes the **load()** method on the media element (page 236), and ignores any resulting exceptions. The task source (page 429) for this task is the media element (page 236)'s own media element new resource task source (page 238).

Note: *If a **src** attribute is specified, the resource it specifies is the media resource (page 238) that will be used. Otherwise, the resource specified by the first suitable source element child of the media element (page 236) is the one used.*

The **src** DOM attribute on media elements (page 236) must reflect (page 67) the respective content attribute of the same name.

To **pick a media resource** for a media element (page 236), a user agent must use the following steps:

1. Let the *chosen resource's pixel ratio override* be *none*.
2. If the media element (page 236) has a **src** attribute, then resolve (page 55) the URL (page 52) given in that attribute. If that is successful, then the resulting absolute URL (page 56) is the address of the media resource (page 238); jump to the last step.
3. Otherwise, let *candidate* be the first source element child in the media element (page 236), or null if there is no such child.
4. *Loop:* this is the start of the loop that looks at the source elements.
5. If *candidate* is not null and it has a **pixelratio** attribute, and the result of applying the rules for parsing floating point number values (page 34) to the value of that attribute is not an error, then let the *chosen resource's pixel ratio override* be that result; otherwise, reset it back to *none*.
6. If either:
 - *candidate* is null, or
 - the *candidate* element has no **src** attribute, or
 - resolving (page 55) the URL (page 52) given by the *candidate* element's **src** attribute fails, or
 - the *candidate* element has a **type** attribute and that attribute's value, when parsed as a MIME type, does not represent a type that the user agent can

render (including any codecs described by the codec parameter), or [RFC2046] [RFC4281]

- the *candidate* element has a media attribute and that attribute's value, when processed according to the rules for media queries (page 30), does not match the current environment, [MQ]

...then the *candidate* is not suitable; go to the next step.

Otherwise, the result of resolving (page 55) the URL (page 52) given in that *candidate* element's src attribute is the address of the media resource (page 238); jump to the last step.

7. Let *candidate* be the next source element child in the media element (page 236), or null if there are no more such children.
8. If *candidate* is not null, return to the step labeled *loop*.
9. There is no media resource (page 238). Abort these steps.
10. Let the address of the **chosen media resource** be the absolute URL (page 56) that was found before jumping to this step, and let its **pixel ratio override** be the value of the *chosen resource*'s *pixel ratio override*.

The **currentSrc** DOM attribute must return the empty string if the media element (page 236)'s networkState has the value NETWORK_EMPTY, and the absolute URL (page 56) that is the address of the chosen media resource (page 240) otherwise.

4.8.10.3 Network states

As media elements (page 236) interact with the network, their current network activity is represented by the **networkState** attribute. On getting, it must return the current network state of the element, which must be one of the following values:

NETWORK_EMPTY (numeric value 0)

The element has not yet been initialized. All attributes are in their initial states.

NETWORK_IDLE (numeric value 1)

The element has a chosen media resource (page 240), but the user agent is not using the network to obtain any more of the resource than is already obtained at this time.

NETWORK_LOADING (numeric value 2)

The user agent is actively trying to download data for the chosen media resource (page 240).

NETWORK_LOADED (numeric value 3)

The entire media resource (page 238) has been obtained and is available to the user agent locally. Network connectivity could be lost without affecting the media playback.

The algorithm for the load() method defined below describes exactly when the networkState attribute changes value and what events fire to indicate changes in this state.

Note: Some resources, e.g. streaming Web radio, can never reach the NETWORK_LOADED state.

4.8.10.4 Loading the media resource

All media elements (page 236) have a **begun flag**, which must begin in the false state, and an **autoplaying flag**, which must begin in the true state.

When the **load()** method on a media element (page 236) is invoked, the user agent must run the following steps. Note that this algorithm might get aborted, e.g. if the **load()** method itself is invoked again.

1. If there are any tasks (page 429) from the media element (page 236)'s media element new resource task source (page 238) or its media element event task source (page 238) in one of the task queues (page 429), then remove those tasks.

Note: Basically, pending events, callbacks, and loads for the media element are discarded when the media element starts loading a new resource.

2. Any already-running instance of this algorithm for this element must be aborted. If those method calls have not yet returned, they must finish the step they are on, and then immediately return. This is not blocking; this algorithm must not wait for the earlier instances to abort before continuing.
3. If the element's begun flag (page 241) is true, then the begun flag (page 241) must be set to false, the error attribute must be set to a new `MediaError` object whose code attribute is set to `MEDIA_ERR_ABORTED`, and the user agent must fire a progress event (page 436) called `abort` at the media element (page 236).
4. The error attribute must be set to null and the autoplaying flag (page 241) must be set to true.
5. The playbackRate attribute must be set to the value of the defaultPlaybackRate attribute.
6. If the media element (page 236)'s networkState is not set to `NETWORK_EMPTY`, then the following substeps must be followed:
 1. The networkState attribute must be set to `NETWORK_EMPTY` (page 240).
 2. If readyState is not set to `HAVE NOTHING`, it must be set to that state.
 3. If the paused attribute is false, it must be set to true.
 4. If seeking is true, it must be set to false.
 5. The current playback position (page 246) must be set to 0.
 6. The currentLoop DOM attribute must be set to 0.
 7. The user agent must fire a simple event (page 436) called `emptied` at the media element (page 236).

7. The user agent must pick a media resource (page 239) for the media element (page 236). If that fails, the method must raise an INVALID_STATE_ERR exception, and abort these steps.
8. The networkState attribute must be set to NETWORK_IDLE (page 240).

Note: The currentSrc attribute starts returning the new value.

9. The user agent must then set the begun flag (page 241) to true and fire a progress event (page 436) called loadstart at the media element (page 236).
10. The method must return, but these steps must continue.
11. **Note: Playback of any previously playing media resource (page 238) for this element stops.**
12. If a fetching process is in progress for the media element (page 236), the user agent should stop it.
13. The user agent must then begin to fetch (page 59) the chosen media resource (page 240). The rate of the download may be throttled, however, in response to user preferences (including throttling it to zero until the user indicates that the download can start), or to balance the download with other connections sharing the same bandwidth.
14. While the fetching process is progressing, the user agent must set the networkState to NETWORK_LOADING and queue a task (page 429) to fire a progress event (page 436) called progress at the element every 350ms ($\pm 200\text{ms}$) or for every byte received, whichever is least frequent.

If at any point the user agent has received no data for more than about three seconds, the user agent must queue a task (page 429) to fire a progress event (page 436) called stalled at the element.

User agents may allow users to selectively block or slow media data (page 238) downloads. When a media element (page 236)'s download has been blocked, the user agent must act as if it was stalled (as opposed to acting as if the connection was closed).

User agents may decide to not download more content at any time, e.g. after buffering five minutes of a one hour media resource, while waiting for the user to decide whether to play the resource or not, or while waiting for user input in an interactive resource. When a media element (page 236)'s download has been suspended, the user agent must set the networkState to NETWORK_IDLE and queue a task (page 429) to fire a progress event (page 436) called suspend at the element.

The user agent may use whatever means necessary to fetch the resource (within the constraints put forward by this and other specifications); for example, reconnecting to the server in the face of network errors, using HTTP partial range requests, or switching to a streaming protocol. The user agent must consider a resource erroneous only if it has given up trying to fetch it.

The networking task source (page 430) tasks (page 429) to process the data as it is being fetched must, when appropriate, include the relevant substeps from the following list:

↪ **If the media data (page 238) cannot be fetched at all, due to network errors, causing the user agent to give up trying to fetch the resource**

DNS errors and HTTP 4xx and 5xx errors (and equivalents in other protocols) must cause the user agent to execute the following steps. User agents may also follow these steps in response to other network errors of similar severity.

1. The user agent should cancel the fetching process.
2. The error attribute must be set to a new MediaError object whose code attribute is set to MEDIA_ERR_NETWORK.
3. The begun flag (page 241) must be set to false and the user agent must queue a task (page 429) to fire a progress event (page 436) called error at the media element (page 236).
4. The element's networkState attribute must be switched to the NETWORK_EMPTY (page 240) value and the user agent must queue a task (page 429) to fire a simple event (page 436) called emptied at the element.
5. These steps must be aborted.

↪ **If the media data (page 238) can be fetched but is in an unsupported format, or can otherwise not be rendered at all**

The server returning a file of the wrong kind (e.g. one that turns out to not be pure audio when the media element (page 236) is an audio element), or the file using unsupported codecs for all the data, must cause the user agent to execute the following steps. User agents may also execute these steps in response to other codec-related fatal errors, such as the file requiring more resources to process than the user agent can provide in real time.

1. The user agent should cancel the fetching process.
2. The error attribute must be set to a new MediaError object whose code attribute is set to MEDIA_ERR_DECODE.
3. The begun flag (page 241) must be set to false and the user agent must queue a task (page 429) to fire a progress event (page 436) called error at the media element (page 236).
4. The element's networkState attribute must be switched to the NETWORK_EMPTY (page 240) value and the user agent must queue a task (page 429) to fire a simple event (page 436) called emptied at the element.
5. These steps must be aborted.

↪ **If the media data (page 238) fetching process is aborted by the user**

The fetching process is aborted by the user, e.g. because the user navigated the browsing context to another page, the user agent must execute the following steps. These steps are not followed if the `load()` method itself is reinvoked, as the steps above handle that particular kind of abort.

1. The user agent should cancel the fetching process.
2. The `error` attribute must be set to a new `MediaError` object whose `code` attribute is set to `MEDIA_ERR_ABORT`.
3. The `begun` flag (page 241) must be set to false and the user agent must queue a task (page 429) to fire a progress event (page 436) called `abort` at the media element (page 236).
4. If the media element (page 236)'s `readyState` attribute has a value equal to `HAVE_NOTHING`, the element's `networkState` attribute must be switched to the `NETWORK_EMPTY` (page 240) value and the user agent must queue a task (page 429) to fire a simple event (page 436) called `emptied` at the element. (If the `readyState` attribute has a value greater than `HAVE_NOTHING`, then this doesn't happen; the available data, if any, will be playable.)
5. These steps must be aborted.

↪ **If the media data (page 238) can be fetched but has non-fatal errors or uses, in part, codecs that are unsupported, preventing the user agent from rendering the content completely correctly but not preventing playback altogether**

The server returning data that is partially usable but cannot be optimally rendered must cause the user agent to execute the following steps.

**
**

1. Should we fire a 'warning' event? Set the 'error' flag to '`MEDIA_ERR_SUBOPTIMAL`' or something?

↪ **Once enough of the media data (page 238) has been fetched to determine the duration of the media resource (page 238), its dimensions, and other metadata**

The user agent must follow these substeps:

1. The current playback position (page 246) must be set to the *effective start*.
2. The `readyState` attribute must be set to `HAVE_METADATA`.
3. **Note: A number of attributes, including `duration`, `buffered`, and `played`, become available.**
4. **Note: The user agent will (page 246) queue a task (page 429) to fire a simple event (page 436) called `durationchange` at the element at this point.**

5. The user agent must queue a task (page 429) to fire a simple event (page 436) called loadedmetadata at the element.

↳ **Once enough of the media data (page 238) has been fetched to enable the user agent to display the frame at the effective start (page 246) of the media resource (page 238)**

The user agent must follow these substeps:

1. The readyState attribute must change to HAVE_CURRENT_DATA.
2. The user agent must queue a task (page 429) to fire a simple event (page 436) called loadeddata at the element.

When the user agent has completely fetched of the entire media resource (page 238), it must move on to the next step. This might never happen, e.g. when streaming an infinite resource such as Web radio.

15. If the fetching process completes without errors, the begun flag (page 241) must be set to false, the networkState attribute must be set to NETWORK_LOADED, and the user agent must queue a task (page 429) to fire a progress event (page 436) called load at the element.

If a media element (page 236) whose networkState has the value NETWORK_EMPTY is inserted into a document (page 24), the user agent must queue a task (page 429) that implicitly invokes the load() method on the media element (page 236), and ignores any resulting exceptions. The task source (page 429) for this task is the media element (page 236)'s own media element new resource task source (page 238).

The **bufferingRate** attribute must return the average number of bits received per second for the current download over the past few seconds. If there is no download in progress, the attribute must return 0.

The **bufferingThrottled** attribute must return true if the user agent is intentionally throttling the bandwidth used by the download (including when throttling to zero to pause the download altogether), and false otherwise.

The **buffered** attribute must return a static normalized TimeRanges object (page 257) that represents the ranges of the media resource (page 238), if any, that the user agent has buffered, at the time the attribute is evaluated.

Note: *Typically this will be a single range anchored at the zero point, but if, e.g. the user agent uses HTTP range requests in response to seeking, then there could be multiple ranges.*

The **bufferedBytes** attribute must return a static normalized ByteRanges object (page 258) that represents the ranges of the media resource (page 238), if any, that the user agent has buffered, at the time the attribute is evaluated.

The **totalBytes** attribute must return the length of the media resource (page 238), in bytes, if it is known and finite. If it is not known, is infinite (e.g. streaming radio), or if no media data (page 238) is available, the attribute must return 0.

User agents may discard previously buffered data.

Note: Thus, a time or byte position included within a range of the objects returned by the buffered or bufferedBytes attributes at one time can end up being not included in the range(s) of objects returned by the same attributes at a later time.

4.8.10.5 Offsets into the media resource

The **duration** attribute must return the length of the media resource (page 238), in seconds. If no media data (page 238) is available, then the attributes must return 0. If media data (page 238) is available but the length is not known, the attribute must return the Not-a-Number (NaN) value. If the media resource (page 238) is known to be unbounded (e.g. a streaming radio), then the attribute must return the positive Infinity value.

When the length of the media resource (page 238) changes (e.g. from being unknown to known, or from indeterminate to known, or from a previously established length to a new length) the user agent must queue a task (page 429) to fire a simple event (page 436) called durationchange at the media element (page 236).

Media elements (page 236) have a **current playback position**, which must initially be zero. The current position is a time.

The **currentTime** attribute must, on getting, return the current playback position (page 246), expressed in seconds. On setting, the user agent must seek (page 254) to the new value (which might raise an exception).

If the media resource (page 238) is a streaming resource, then the user agent might be unable to obtain certain parts of the resource after it has expired from its buffer. The **earliest possible position** is the earliest position in the stream that the user agent can ever obtain again.

The **start** content attribute gives the offset into the media resource (page 238) at which playback is to begin. The default value is the default start position of the media resource (page 238), or 0 if not enough media data (page 238) has been obtained yet to determine the default start position or if the resource doesn't specify a default start position.

The **effective start** is the smaller of the start DOM attribute, the *earliest possible position*, and the end of the media resource (page 238).

The **loopstart** content attribute gives the offset into the media resource (page 238) at which playback is to begin when looping a clip. The default value of the loopstart content attribute is the value of the start DOM attribute.

The **effective loop start** is the smaller of the loopStart DOM attribute, the *earliest possible position*, and the end of the media resource (page 238).

The **loopend** content attribute gives an offset into the media resource (page 238) at which playback is to jump back to the loopstart, when looping the clip. The default value of the loopend content attribute is the value of the end DOM attribute.

The **effective loop end** is the greater of the start, loopStart, and loopEnd DOM attributes and the *earliest possible position*, except if that is greater than the end of the media resource (page 238), in which case that's its value.

The **end** content attribute gives an offset into the media resource (page 238) at which playback is to end. The default value is infinity.

The **effective end** is the greater of the start, loopStart, and end DOM attributes and the *earliest possible position*, except if that is greater than the end of the media resource (page 238), in which case that's its value.

The start, loopstart, loopend, and end attributes must, if specified, contain value time offsets. To get the time values they represent, user agents must use the rules for parsing time offsets (page 49).

The **start**, **loopStart**, **loopEnd**, and **end** DOM attributes must reflect (page 67) the start, loopstart, loopend, and end content attributes on the media element (page 236) respectively.

The **playcount** content attribute gives the number of times to play the clip. The default value is 1.

The **playCount** DOM attribute must reflect (page 67) the playcount content attribute on the media element (page 236). The value must be limited to only positive non-zero numbers (page 68).

The **currentLoop** attribute must initially have the value 0. It gives the index of the current loop. It is changed during playback as described below.

When any of the start, loopStart, loopEnd, end, playCount, and currentLoop DOM attributes change value (either through content attribute mutations reflecting into the DOM attribute, if applicable, or through direct mutations of the DOM attribute), or if the *earliest possible position* changes, the user agent must apply the following steps:

1. If the playCount DOM attribute's value is less than or equal to the currentLoop DOM attribute's value, then the currentLoop DOM attribute's value must be set to playCount-1 (which will make the current loop the last loop).
2. If the media element (page 236)'s readyState is in the HAVE NOTHING state, then the user agent must at this point abort these steps.
3. If the currentLoop is zero, and the current playback position (page 246) is before the *effective start*, the user agent must seek (page 254) to the *effective start*.
4. If the currentLoop is greater than zero, and the current playback position (page 246) is before the *effective loop start*, the user agent must seek (page 254) to the *effective loop start*.
5. If the currentLoop is less than playCount-1, and the current playback position (page 246) is after the *effective loop end*, the user agent must seek (page 254) to the *effective loop start*, and increase currentLoop by 1.
6. If the currentLoop is equal to playCount-1, and the current playback position (page 246) is after the *effective end*, the user agent must seek (page 254) to the *effective end* and then the looping will end.

4.8.10.6 The ready states

Media elements (page 236) have a *ready state*, which describes to what degree they are ready to be rendered at the current playback position (page 246). The possible values are as follows; the ready state of a media element at any particular time is the greatest value describing the state of the element:

HAVE NOTHING (numeric value 0)

No information regarding the media resource (page 238) is available. No data for the current playback position (page 246) is available. Media elements (page 236) whose networkState attribute is NETWORK_EMPTY are always in the HAVE NOTHING state.

HAVE_METADATA (numeric value 1)

Enough of the resource has been obtained that the metadata attributes are initialized (e.g. the length is known). The API will no longer raise an exception when seeking.

HAVE_SOME_DATA (numeric value 2)

Data for the immediate current playback position (page 246) is not available, but there is at least one playback position for which the *ready state* would have a value of HAVE_CURRENT_DATA or greater.

HAVE_CURRENT_DATA (numeric value 3)

Data for the immediate current playback position (page 246) is available, but not enough data is available that the user agent could successfully advance the current playback position (page 246) at all without immediately reverting to the HAVE_SOME_DATA state. For example, in video this corresponds to the user agent having data from the current frame, but not the next frame.

HAVE_FUTURE_DATA (numeric value 4)

Data for the immediate current playback position (page 246) is available, as well as enough data for the user agent to advance the current playback position (page 246) at least a little without immediately reverting to the HAVE_SOME_DATA state. For example, In video this corresponds to the user agent having data for at least the current frame and the next frame.

HAVE_ENOUGH_DATA (numeric value 5)

Data for the immediate current playback position (page 246) is available, as well as enough data for the user agent to advance the current playback position (page 246) at least a little without immediately reverting to the HAVE_SOME_DATA state, and, in addition, the user agent estimates that data is being fetched at a rate where the current playback position (page 246), if it were to advance at the rate given by the defaultPlaybackRate attribute, would not overtake the available data before playback reaches the effective end (page 247) of the media resource (page 238) on the last loop (page 247).

When the ready state of a media element (page 236) whose networkState is not NETWORK_EMPTY changes, the user agent must follow the steps given below:

- ↪ If the previous ready state was HAVE NOTHING, and the new ready state is HAVE_METADATA

Note: A loadedmetadata DOM event will be fired (page 245) as part of the load() algorithm.

- ↪ If the previous ready state was HAVE_METADATA, and the new ready state is HAVE_CURRENT_DATA

Note: A loadeddata DOM event will be fired (page 245) as part of the load() algorithm.
- ↪ If the previous ready state was HAVE_FUTURE_DATA or more, and the new ready state is HAVE_CURRENT_DATA or less

Note: A waiting DOM event can be fired (page 251), depending on the current state of playback.
- ↪ If the previous ready state was HAVE_CURRENT_DATA or less, and the new ready state is HAVE_FUTURE_DATA

The user agent must queue a task (page 429) to fire a simple event (page 436) called canplay.

- ↪ If the new ready state is HAVE_ENOUGH_DATA

The user agent must queue a task (page 429) to fire a simple event (page 436) called canplay, and then queue a task (page 429) to fire a simple event (page 436) called canplaythrough.

If the autoplaying flag (page 241) is true, and the paused attribute is true, and the media element (page 236) has an autoplay attribute specified, then the user agent may also set the paused attribute to false and queue a task (page 429) to fire a simple event (page 436) called play.

Note: User agents are not required to autoplay, and it is suggested that user agents honor user preferences on the matter. Authors are urged to use the autoplay attribute rather than using script to force the video to play, so as to allow the user to override the behavior if so desired.

Note: It is possible for the ready state of a media element to jump between these states discontinuously. For example, the state of a media element can jump straight from HAVE_SOME_DATA to HAVE_ENOUGH_DATA without passing through the HAVE_CURRENT_DATA and HAVE_FUTURE_DATA states.

The **readyState** DOM attribute must, on getting, return the value described above that describes the current ready state of the media element (page 236).

The **autoplay** attribute is a boolean attribute (page 32). When present, the algorithm described herein will cause the user agent to automatically begin playback of the media resource (page 238) as soon as it can do so without stopping.

The **autoplay** DOM attribute must reflect (page 67) the content attribute of the same name.

4.8.10.7 Cue ranges

Media elements (page 236) have a set of **cue ranges**. Each cue range is made up of the following information:

A class name

A group of related ranges can be given the same class name so that they can all be removed at the same time.

An identifier

A string can be assigned to each cue range for identification by script. The string need not be unique and can contain any value.

A start time

An end time

The actual time range, using the same timeline as the media resource (page 238) itself.

A "pause" boolean

A flag indicating whether to pause playback on exit.

An "enter" callback

A callback that is called when the current playback position (page 246) enters the range.

An "exit" callback

A callback that is called when the current playback position (page 246) exits the range.

An "active" boolean

A flag indicating whether the range is active or not.

The **addCueRange(*className*, *id*, *start*, *end*, *pauseOnExit*, *enterCallback*, *exitCallback*)** method must, when called, add a cue range (page 249) to the media element (page 236), that cue range having the class name *className*, the identifier *id*, the start time *start* (in seconds), the end time *end* (in seconds), the "pause" boolean with the same value as *pauseOnExit*, the "enter" callback *enterCallback*, the "exit" callback *exitCallback*, and an "active" boolean that is true if the current playback position (page 246) is equal to or greater than the start time and less than the end time, and false otherwise.

The **removeCueRanges(*className*)** method must, when called, remove all the cue ranges (page 249) of the media element (page 236) which have the class name *className*.

4.8.10.8 Playing the media resource

The **paused** attribute represents whether the media element (page 236) is paused or not. The attribute must initially be true.

A media element (page 236) is said to be **potentially playing** when its paused attribute is false, the readyState attribute is either HAVE_FUTURE_DATA or HAVE_FUTURE_DATA, the element has not ended playback (page 250), playback has not stopped due to errors (page 251), and the element has not paused for user interaction (page 251).

A media element (page 236) is said to have **ended playback** when the element's readyState attribute is HAD_METADATA or greater, the current playback position (page 246) is equal to the **effective end** of the media resource (page 238), and the currentLoop attribute is equal to playCount-1.

The **ended** attribute must return true if the media element (page 236) has ended playback (page 250), and false otherwise.

A media element (page 236) is said to have **stopped due to errors** when the element's readyState attribute is HAVE_METADATA or greater, and the user agent encounters a non-fatal error (page 244) during the processing of the media data (page 238), and due to that error, is not able to play the content at the current playback position (page 246).

A media element (page 236) is said to have **paused for user interaction** when its paused attribute is false, the readyState attribute is either HAVE_FUTURE_DATA or HAVE_ENOUGH_DATA and the user agent has reached a point in the media resource (page 238) where the user has to make a selection for the resource to continue.

It is possible for a media element (page 236) to have both ended playback (page 250) and paused for user interaction (page 251) at the same time.

When a media element (page 236) that is potentially playing (page 250) stops playing because it has paused for user interaction (page 251), the user agent must queue a task (page 429) to fire a simple event (page 436) called timeupdate at the element.

When a media element (page 236) that is potentially playing (page 250) stops playing because its readyState attribute changes to a value lower than HAVE_FUTURE_DATA, without the element having ended playback (page 250), or playback having stopped due to errors (page 251), or playback having paused for user interaction (page 251), or the seeking algorithm (page 254) being invoked, the user agent must queue a task (page 429) to fire a simple event (page 436) called timeupdate at the element, and queue a task (page 429) to fire a simple event (page 436) called waiting at the element.

When currentLoop is less than playCount-1 and the current playback position (page 246) reaches the *effective loop end*, then the user agent must increase currentLoop by 1 and seek (page 254) to the *effective loop start*.

When currentLoop is equal to the playCount-1 and the current playback position (page 246) reaches the *effective end*, then the user agent must follow these steps:

1. The user agent must stop playback.
2. **Note: The ended attribute becomes true.**
3. The user agent must queue a task (page 429) to fire a simple event (page 436) called timeupdate at the element.
4. The user agent must queue a task (page 429) to fire a simple event (page 436) called ended at the element.

The **defaultPlaybackRate** attribute gives the desired speed at which the media resource (page 238) is to play, as a multiple of its intrinsic speed. The attribute is mutable, but on setting, if the new value is 0.0, a NOT_SUPPORTED_ERR exception must be raised instead of the value being changed. It must initially have the value 1.0.

The **playbackRate** attribute gives the speed at which the media resource (page 238) plays, as a multiple of its intrinsic speed. If it is not equal to the defaultPlaybackRate, then the implication is that the user is using a feature such as fast forward or slow motion playback. The attribute is mutable, but on setting, if the new value is 0.0, a NOT_SUPPORTED_ERR exception must be raised instead of the value being changed. Otherwise, the playback must change speed (if the element is potentially playing (page 250)). It must initially have the value 1.0.

When the defaultPlaybackRate or playbackRate attributes change value (either by being set by script or by being changed directly by the user agent, e.g. in response to user control) the user agent must queue a task (page 429) to fire a simple event (page 436) called ratechange at the media element (page 236).

The **played** attribute must return a static normalized TimeRanges object (page 257) that represents the ranges of the media resource (page 238), if any, that the user agent has so far rendered, at the time the attribute is evaluated.

When the **play()** method on a media element (page 236) is invoked, the user agent must run the following steps.

1. If the media element (page 236)'s networkState attribute has the value NETWORK_EMPTY, then the user agent must invoke the `load()` method and wait for it to return. If that raises an exception, that exception must be reraised by the `play()` method.
2. If the playback has ended (page 250), then the user agent must set currentLoop to zero and seek (page 254) to the *effective start*.

Note: *If this involved a seek, the user agent will (page 255) queue a task (page 429) to fire a simple event (page 436) called timeupdate at the media element (page 236).*

3. The playbackRate attribute must be set to the value of the defaultPlaybackRate attribute.

Note: *If this caused the playbackRate attribute to change value, the user agent will (page 252) queue a task (page 429) to fire a simple event (page 436) called ratechange at the media element (page 236).*

4. If the media element (page 236)'s paused attribute is true, it must be set to false.
5. The media element (page 236)'s autoplaying flag (page 241) must be set to false.
6. The method must then return.
7. If the fourth step above changed the value of paused, the user agent must queue a task (page 429) to fire a simple event (page 436) called play at the element.

When the **pause()** method is invoked, the user agent must run the following steps:

1. If the media element (page 236)'s networkState attribute has the value NETWORK_EMPTY, then the user agent must invoke the `load()` method and wait for it to return. If that raises an exception, that exception must be reraised by the `pause()` method.
2. If the media element (page 236)'s paused attribute is false, it must be set to true.
3. The media element (page 236)'s autoplaying flag (page 241) must be set to false.
4. The method must then return.

- If the second step above changed the value of paused, then the user agent must queue a task (page 429) to fire a simple event (page 436) called timeupdate at the element, and queue a task (page 429) to fire a simple event (page 436) called pause at the element.

When a media element (page 236) is potentially playing (page 250) and its Document is an active document (page 414), its current playback position (page 246) must increase monotonically at playbackRate units of media time per unit time of wall clock time.

Note: *This specification doesn't define how the user agent achieves the appropriate playback rate — depending on the protocol and media available, it is plausible that the user agent could negotiate with the server to have the server provide the media data at the appropriate rate, so that (except for the period between when the rate is changed and when the server updates the stream's playback rate) the client doesn't actually have to drop or interpolate any frames.*

When the playbackRate is negative (playback is backwards), any corresponding audio must be muted. When the playbackRate is so low or so high that the user agent cannot play audio usefully, the corresponding audio must also be muted. If the playbackRate is not 1.0, the user agent may apply pitch adjustments to the audio as necessary to render it faithfully.

Media elements (page 236) that are potentially playing (page 250) while not in a Document must not play any video, but should play any audio component. Media elements must not stop playing just because all references to them have been removed; only once a media element to which no references exist has reached a point where no further audio remains to be played for that element (e.g. because the element is paused or because the end of the clip has been reached) may the element be garbage collected.

When the current playback position (page 246) of a media element (page 236) changes (e.g. due to playback or seeking), the user agent must run the following steps. If the current playback position (page 246) changes while the steps are running, then the user agent must wait for the steps to complete, and then must immediately rerun the steps. (These steps are thus run as often as possible or needed — if one iteration takes a long time, this can cause certain ranges to be skipped over as the user agent rushes ahead to "catch up".)

- Let *current ranges* be an ordered list of cue ranges (page 249), initialized to contain all the cue ranges (page 249) of the media element (page 236) whose start times are less than or equal to the current playback position (page 246) and whose end times are greater than the current playback position (page 246), in the order they were added to the element.
- Let *other ranges* be an ordered list of cue ranges (page 249), initialized to contain all the cue ranges (page 249) of the media element (page 236) that are not present in *current ranges*, in the order they were added to the element.
- If none of the cue ranges (page 249) in *current ranges* have their "active" boolean set to "false" (inactive) and none of the cue ranges (page 249) in *other ranges* have their "active" boolean set to "true" (active), then abort these steps.

4. If the time was reached through the usual monotonic increase of the current playback position during normal playback, the user agent must then queue a task (page 429) to fire a simple event (page 436) called `timeupdate` at the element. (In the other cases, such as explicit seeks, relevant events get fired as part of the overall process of changing the current playback position.)
5. If the time was reached through the usual monotonic increase of the current playback position during normal playback, and there are cue ranges (page 249) in *other ranges* that have both their "active" boolean and their "pause" boolean set to "true", then immediately act as if the element's `pause()` method had been invoked. (In the other cases, such as explicit seeks, playback is not paused by exiting a cue range, even if that cue range has its "pause" boolean set to "true".)
6. For each non-null "exit" callback of the cue ranges (page 249) in *other ranges* that have their "active" boolean set to "true" (active), in list order, queue a task (page 429) that invokes the callback, passing the cue range's identifier as the callback's only argument.
7. For each non-null "enter" callback of the cue ranges (page 249) in *current ranges* that have their "active" boolean set to "false" (inactive), in list order, queue a task (page 429) that invokes the callback, passing the cue range's identifier as the callback's only argument.
8. Set the "active" boolean of all the cue ranges (page 249) in the *current ranges* list to "true" (active), and the "active" boolean of all the cue ranges (page 249) in the *other ranges* list to "false" (inactive).

When a media element (page 236) is removed from a Document, if the media element (page 236)'s `networkState` attribute has a value other than `NETWORK_EMPTY` then the user agent must act as if the `pause()` method had been invoked.

Note: If the media element (page 236)'s Document stops being an active document, then the playback will stop (page 253) until the document is active again.

4.8.10.9 Seeking

The `seeking` attribute must initially have the value `false`.

When the user agent is required to `seek` to a particular *new playback position* in the media resource (page 238), it means that the user agent must run the following steps:

1. If the media element (page 236)'s `readyState` is `HAVE NOTHING`, then the user agent must raise an `INVALID_STATE_ERR` exception (if the seek was in response to a DOM method call or setting of a DOM attribute), and abort these steps.
2. If `currentLoop` is 0, let `min` be the *effective start*. Otherwise, let it be the *effective loop start*.
3. If `currentLoop` is equal to `playCount-1`, let `max` be the *effective end*. Otherwise, let it be the *effective loop end*.
4. If the *new playback position* is more than `max`, let it be `max`.

5. If the *new playback position* is less than *min*, let it be *min*.
6. If the (possibly now changed) *new playback position* is not in one of the ranges given in the **seekable** attribute, then the user agent must raise an INDEX_SIZE_ERR exception (if the seek was in response to a DOM method call or setting of a DOM attribute), and abort these steps.
7. The current playback position (page 246) must be set to the given *new playback position*.
8. The seeking DOM attribute must be set to true.
9. The user agent must queue a task (page 429) to fire a simple event (page 436) called `timeupdate` at the element.
10. If the media element (page 236) was potentially playing (page 250) immediately before it started seeking, but seeking caused its `readyState` attribute to change to a value lower than HAVE_FUTURE_FRAME, the user agent must queue a task (page 429) to fire a simple event (page 436) called `waiting` at the element.
11. If, when it reaches this step, the user agent has still not established whether or not the media data (page 238) for the *new playback position* is available, and, if it is, decoded enough data to play back that position, the user agent must queue a task (page 429) to fire a simple event (page 436) called `seeking` at the element.
12. If the seek was in response to a DOM method call or setting of a DOM attribute, then continue the script. The remainder of these steps must be run asynchronously.
13. The user agent must wait until it has established whether or not the media data (page 238) for the *new playback position* is available, and, if it is, until it has decoded enough data to play back that position.
14. The seeking DOM attribute must be set to false.
15. The user agent must queue a task (page 429) to fire a simple event (page 436) called `seeked` at the element.

The **seekable** attribute must return a static normalized TimeRanges object (page 257) that represents the ranges of the media resource (page 238), if any, that the user agent is able to seek to, at the time the attribute is evaluated, notwithstanding the looping attributes (i.e. the `effective start` and `effective end`, etc, don't affect the `seekable` attribute).

Note: *If the user agent can seek to anywhere in the media resource (page 238), e.g. because it a simple movie file and the user agent and the server support HTTP Range requests, then the attribute would return an object with one range, whose start is the time of the first frame (typically zero), and whose end is the same as the time of the first frame plus the duration attribute's value (which would equal the time of the last frame).*

Note: *The range might be continuously changing, e.g. if the user agent is buffering a sliding window on an infinite stream. This is the behavior seen with DVRs viewing live TV, for instance.*

Media resources (page 238) might be internally scripted or interactive. Thus, a media element (page 236) could play in a non-linear fashion. If this happens, the user agent must act as if the algorithm for seeking (page 254) was used whenever the current playback position (page 246) changes in a discontinuous fashion (so that the relevant events fire).

4.8.10.10 User interface

The **controls** attribute is a boolean attribute (page 32). If the attribute is present, or if the media element (page 236) is without script (page 428), then the user agent should **expose a user interface to the user**. This user interface should include features to begin playback, pause playback, seek to an arbitrary position in the content (if the content supports arbitrary seeking), change the volume, and show the media content in manners more suitable to the user (e.g. full-screen video or in an independent resizable window). Other controls may also be made available.

If the attribute is absent, then the user agent should avoid making a user interface available that could conflict with an author-provided user interface. User agents may make the following features available, however, even when the attribute is absent:

User agents may provide controls to affect playback of the media resource (e.g. play, pause, seeking, and volume controls), but such features should not interfere with the page's normal rendering. For example, such features could be exposed in the media element (page 236)'s context menu.

Where possible (specifically, for starting, stopping, pausing, and unpauseing playback, for muting or changing the volume of the audio, and for seeking), user interface features exposed by the user agent must be implemented in terms of the DOM API described above, so that, e.g., all the same events fire.

The **controls** DOM attribute must reflect (page 67) the content attribute of the same name.

The **volume** attribute must return the playback volume of any audio portions of the media element (page 236), in the range 0.0 (silent) to 1.0 (loudest). Initially, the volume must be 1.0, but user agents may remember the last set value across sessions, on a per-site basis or otherwise, so the volume may start at other values. On setting, if the new value is in the range 0.0 to 1.0 inclusive, the attribute must be set to the new value and the playback volume must be correspondingly adjusted as soon as possible after setting the attribute, with 0.0 being silent, and 1.0 being the loudest setting, values in between increasing in loudness. The range need not be linear. The loudest setting may be lower than the system's loudest possible setting; for example the user could have set a maximum volume. If the new value is outside the range 0.0 to 1.0 inclusive, then, on setting, an **INDEX_SIZE_ERR** exception must be raised instead.

The **muted** attribute must return true if the audio channels are muted and false otherwise. Initially, the audio channels should not be muted (false), but user agents may remember the last set value across sessions, on a per-site basis or otherwise, so the muted state may start as muted (true). On setting, the attribute must be set to the new value; if the new value is true, audio playback for this media resource (page 238) must then be muted, and if false, audio playback must then be enabled.

Whenever either the muted or volume attributes are changed, the user agent must queue a task (page 429) to fire a simple event (page 436) called **volumechange** at the media element (page 236).

4.8.10.11 Time ranges

Objects implementing the TimeRanges interface represent a list of ranges (periods) of time.

```
interface TimeRanges {
    readonly attribute unsigned long length;
    float start(in unsigned long index);
    float end(in unsigned long index);
};
```

The **length** DOM attribute must return the number of ranges represented by the object.

The **start(index)** method must return the position of the start of the *index*th range represented by the object, in seconds measured from the start of the timeline that the object covers.

The **end(index)** method must return the position of the end of the *index*th range represented by the object, in seconds measured from the start of the timeline that the object covers.

These methods must raise INDEX_SIZE_ERR exceptions if called with an *index* argument greater than or equal to the number of ranges represented by the object.

When a TimeRanges object is said to be a **normalized TimeRanges object**, the ranges it represents must obey the following criteria:

- The start of a range must be greater than the end of all earlier ranges.
- The start of a range must be less than the end of that same range.

In other words, the ranges in such an object are ordered, don't overlap, aren't empty, and don't touch (adjacent ranges are folded into one bigger range).

The timelines used by the objects returned by the buffered, seekable and played DOM attributes of media elements (page 236) must be the same as that element's media resource (page 238)'s timeline.

4.8.10.12 Byte ranges

Objects implementing the ByteRanges interface represent a list of ranges of bytes.

```
interface ByteRanges {
    readonly attribute unsigned long length;
    unsigned long start(in unsigned long index);
    unsigned long end(in unsigned long index);
};
```

The **length** DOM attribute must return the number of ranges represented by the object.

The **start(index)** method must return the position of the first byte of the *index*th range represented by the object.

The `end(index)` method must return the position of the byte immediately after the last byte of the `index`th range represented by the object. (The byte position returned by this method is not in the range itself. If the first byte of the range is the byte at position 0, and the entire stream of bytes is in the range, then the value of the position of the byte returned by this method for that range will be the same as the number of bytes in the stream.)

These methods must raise `INDEX_SIZE_ERR` exceptions if called with an `index` argument greater than or equal to the number of ranges represented by the object.

When a `ByteRanges` object is said to be a **normalized ByteRanges object**, the ranges it represents must obey the following criteria:

- The start of a range must be greater than the end of all earlier ranges.
- The start of a range must be less than the end of that same range.

In other words, the ranges in such an object are ordered, don't overlap, aren't empty, and don't touch (adjacent ranges are folded into one bigger range).

4.8.10.13 Event summary

The following events fire on media elements (page 236) as part of the processing model described above:

Event name	Interface	Dispatched when...	Preconditions
<code>loadstart</code>	<code>ProgressEvent [PROGRESS]</code>	The user agent begins fetching the media data (page 238), synchronously during the <code>load()</code> method call.	<code>networkState</code> equals <code>NETWORK_LOADING</code>
<code>progress</code>	<code>ProgressEvent [PROGRESS]</code>	The user agent is fetching media data (page 238).	<code>networkState</code> equals <code>NETWORK_LOADING</code>
<code>suspend</code>	<code>ProgressEvent [PROGRESS]</code>	The user agent is intentionally not currently fetching media data (page 238), but does not have the entire media resource (page 238) downloaded.	<code>networkState</code> equals <code>NETWORK_IDLE</code>
<code>load</code>	<code>ProgressEvent [PROGRESS]</code>	The user agent finishes fetching the entire media resource (page 238).	<code>networkState</code> equals <code>NETWORK_LOADED</code>
<code>abort</code>	<code>ProgressEvent [PROGRESS]</code>	The user agent stops fetching the media data (page 238) before it is completely downloaded. This can be fired synchronously during the <code>load()</code> method call.	<code>error</code> is an object with the code <code>MEDIA_ERR_ABORTED</code> . <code>networkState</code> equals either <code>NETWORK_EMPTY</code> or <code>NETWORK_LOADED</code> , depending on when the download was aborted.
<code>error</code>	<code>ProgressEvent [PROGRESS]</code>	An error occurs while fetching the media data (page 238).	<code>error</code> is an object with the code <code>MEDIA_ERR_NETWORK_ERROR</code> or higher. <code>networkState</code> equals either <code>NETWORK_EMPTY</code> or <code>NETWORK_LOADED</code> , depending on when the download was aborted.
<code>emptied</code>	<code>Event</code>	A media element (page 236) whose <code>networkState</code> was previously not in the <code>NETWORK_EMPTY</code> state has just switched to that state (either	<code>networkState</code> is <code>NETWORK_EMPTY</code> ; all the DOM attributes are in their initial states.

Event name	Interface	Dispatched when...	Preconditions
		because of a fatal error during load that's about to be reported, or because the <code>load()</code> method was reinvoked, in which case it is fired synchronously during the <code>load()</code> method call).	
stalled	ProgressEvent	The user agent is trying to fetch media data (page 238), but data is unexpectedly not forthcoming.	networkState is <code>NETWORK_LOADING</code> .
play	Event	Playback has begun. Fired after the <code>play</code> method has returned.	<code>paused</code> is newly false.
pause	Event	Playback has been paused. Fired after the <code>pause</code> method has returned.	<code>paused</code> is newly true.
loadedmetadata	Event	The user agent has just received the metadata, such as duration or dimensions, for the media resource (page 238).	<code>readyState</code> is newly equal to <code>HAVE_METADATA</code> or greater for the first time.
loadeddata	Event	The user agent can render the media data (page 238) at the current playback position (page 246) for the first time.	<code>readyState</code> newly increased to <code>HAVE_CURRENT_DATA</code> or greater for the first time.
waiting	Event	Playback has stopped because the next frame is not available, but the user agent expects that frame to become available in due course.	<code>readyState</code> is newly equal to or less than <code>HAVE_CURRENT_DATA</code> , and <code>paused</code> is false. Either seeking is true, or the current playback position (page 246) is not contained in any of the ranges in <code>buffered</code> . It is possible for playback to stop for two other reasons without <code>paused</code> being false, but those two reasons do not fire this event: maybe playback ended (page 250), or playback stopped due to errors (page 251).
canplay	Event	The user agent can resume playback of the media data (page 238), but estimates that if playback were to be started now, the media resource (page 238) could not be rendered at the current playback rate up to its end without having to stop for further buffering of content.	<code>readyState</code> newly increased to <code>HAVE_FUTURE_DATA</code> or greater.
canplaythrough	Event	The user agent estimates that if playback were to be started now, the media resource (page 238) could be rendered at the current playback rate all the way to its end without having to stop for further buffering.	<code>readyState</code> is newly equal to <code>HAVE_ENOUGH_DATA</code> .
seeking	Event	The seeking DOM attribute changed to true and the seek operation is taking long enough that the user agent has time to fire the event.	
seeked	Event	The seeking DOM attribute changed to false.	

Event name	Interface	Dispatched when...	Preconditions
<code>timeupdate</code>	Event	The current playback position (page 246) changed as part of normal playback or in an especially interesting way, for example discontinuously.	
<code>ended</code>	Event	Playback has stopped because the end of the media resource (page 238) was reached.	currentTime equals the <i>effective end</i> ; ended is true.
<code>ratechange</code>	Event	Either the defaultPlaybackRate or the playbackRate attribute has just been updated.	
<code>durationchange</code>	Event	The duration attribute has just been updated.	
<code>volumechange</code>	Event	Either the volume attribute or the muted attribute has changed. Fired after the relevant attribute's setter has returned.	

4.8.10.14 Security and privacy considerations

- ** Talk about making sure interactive media files (e.g. SVG) don't have access to the container DOM (XSS potential); talk about not exposing any sensitive data like metadata from tracks in the media files (intranet snooping risk)

4.8.11 The canvas element

Categories

Embedded content (page 94).

Contexts in which this element may be used:

Where embedded content (page 94) is expected.

Content model:

Transparent (page 96).

Element-specific attributes:

width
height

DOM interface:

```
interface HTMLCanvasElement : HTMLElement {
    attribute unsigned long width;
    attribute unsigned long height;

    DOMString toDataURL();
    DOMString toDataURL(in DOMString type, [Variadic] in any args);
```

```
DOMObject getContext(in DOMString contextId);  
};
```

The canvas element represents a resolution-dependent bitmap canvas, which can be used for rendering graphs, game graphics, or other visual images on the fly.

Authors should not use the canvas element in a document when a more suitable element is available. For example, it is inappropriate to use a canvas element to render a page heading: if the desired presentation of the heading is graphically intense, it should be marked up using appropriate elements (typically h1) and then styled using CSS and supporting technologies such as XBL.

When authors use the canvas element, they should also provide content that, when presented to the user, conveys essentially the same function or purpose as the bitmap canvas. This content may be placed as content of the canvas element. The contents of the canvas element, if any, are the element's fallback content (page 95).

In interactive visual media, if the canvas element is with script (page 428), the canvas element represents an embedded element with a dynamically created image.

In non-interactive, static, visual media, if the canvas element has been previously painted on (e.g. if the page was viewed in an interactive visual medium and is now being printed, or if some script that ran during the page layout process painted on the element), then the canvas element represents embedded content (page 94) with the current image and size. Otherwise, the element represents its fallback content (page 95) instead.

In non-visual media, and in visual media if the canvas element is without script (page 428), the canvas element represents its fallback content (page 95) instead.

The canvas element has two attributes to control the size of the coordinate space: **width** and **height**. These attributes, when specified, must have values that are valid non-negative integers (page 33). The rules for parsing non-negative integers (page 33) must be used to obtain their numeric values. If an attribute is missing, or if parsing its value returns an error, then the default value must be used instead. The width attribute defaults to 300, and the height attribute defaults to 150.

The intrinsic dimensions of the canvas element equal the size of the coordinate space, with the numbers interpreted in CSS pixels. However, the element can be sized arbitrarily by a style sheet. During rendering, the image is scaled to fit this layout size.

The size of the coordinate space does not necessarily represent the size of the actual bitmap that the user agent will use internally or during rendering. On high-definition displays, for instance, the user agent may internally use a bitmap with two device pixels per unit in the coordinate space, so that the rendering remains at high quality throughout.

Whenever the width and height attributes are set (whether to a new value or to the previous value), the bitmap and any associated contexts must be cleared back to their initial state and reinitialized with the newly specified coordinate space dimensions.

The **width** and **height** DOM attributes must reflect (page 67) the respective content attributes of the same name.

Only one square appears to be drawn in the following example:

```
// canvas is a reference to a <canvas> element
var context = canvas.getContext('2d');
context.fillRect(0,0,50,50);
canvas.setAttribute('width', '300'); // clears the canvas
context.fillRect(0,100,50,50);
canvas.width = canvas.width; // clears the canvas
context.fillRect(100,0,50,50); // only this square remains
```

When the canvas is initialized it must be set to fully transparent black.

To draw on the canvas, authors must first obtain a reference to a **context** using the **getContext(contextId)** method of the canvas element.

This specification only defines one context, with the name "2d". If `getContext()` is called with that exact string for its `contextId` argument, then the UA must return a reference to an object implementing `CanvasRenderingContext2D`. Other specifications may define their own contexts, which would return different objects.

Vendors may also define experimental contexts using the syntax `vendorname-context`, for example, `moz-3d`.

When the UA is passed an empty string or a string specifying a context that it does not support, then it must return null. String comparisons must be case-sensitive (page 31).

Arguments other than the `contextId` must be ignored.

Note: A future version of this specification will probably define a 3d context (probably based on the OpenGL ES API).

The **toDataURL()** method must, when called with no arguments, return a data: URL containing a representation of the image as a PNG file. [PNG].

If the canvas has no pixels (i.e. either its horizontal dimension or its vertical dimension is zero) then the method must return the string "data: , ". (This is the shortest data: URL; it represents the empty string in a text/plain resource.)

The **toDataURL(type)** method (when called with one or more arguments) must return a data: URL containing a representation of the image in the format given by `type`. The possible values are MIME types with no parameters, for example `image/png`, `image/jpeg`, or even maybe `image/svg+xml` if the implementation actually keeps enough information to reliably render an SVG image from the canvas.

For image types that do not support an alpha channel, the image must be composited onto a solid black background using the source-over operator, and the resulting image must be the one used to create the data: URL.

Only support for `image/png` is required. User agents may support other types. If the user agent does not support the requested type, it must return the image using the PNG format.

User agents must convert the provided type to lower case before establishing if they support that type and before creating the data: URL.

Note: When trying to use types other than `image/png`, authors can check if the image was really returned in the requested format by checking to see if the returned string starts with one the exact strings "`data:image/png,`" or "`data:image/png;`". If it does, the image is PNG, and thus the requested type was not supported. (The one exception to this is if the canvas has either no height or no width, in which case the result might simply be "`data: ,.`".)

If the method is invoked with the first argument giving a type corresponding to one of the types given in the first column of the following table, and the user agent supports that type, then the subsequent arguments, if any, must be treated as described in the second cell of that row.

Type	Other arguments
<code>image/</code> <code>jpeg</code>	The second argument, if it is a number between 0.0 and 1.0, must be treated as the desired quality level. If it is not a number or is outside that range, the user agent must use its default value, as if the argument had been omitted.

Other arguments must be ignored and must not cause the user agent to raise an exception. A future version of this specification will probably define other parameters to be passed to `toDataURL()` to allow authors to more carefully control compression settings, image metadata, etc.

4.8.11.1 The 2D context

When the `getContext()` method of a canvas element is invoked with `2d` as the argument, a `CanvasRenderingContext2D` object is returned.

There is only one `CanvasRenderingContext2D` object per canvas, so calling the `getContext()` method with the `2d` argument a second time must return the same object.

The 2D context represents a flat Cartesian surface whose origin (0,0) is at the top left corner, with the coordinate space having x values increasing when going right, and y values increasing when going down.

```
interface CanvasRenderingContext2D {

    // back-reference to the canvas
    readonly attribute HTMLCanvasElement canvas;

    // state
    void save(); // push state on state stack
    void restore(); // pop state stack and restore state

    // transformations (default transform is the identity matrix)
    void scale(in float x, in float y);
    void rotate(in float angle);
    void translate(in float x, in float y);
    void transform(in float m11, in float m12, in float m21, in float m22,
    in float dx, in float dy);
    void setTransform(in float m11, in float m12, in float m21, in float
    m22, in float dx, in float dy);
}
```

```

// compositing
    attribute float globalAlpha; // (default 1.0)
    attribute DOMString globalCompositeOperation; // (default
source-over)

// colors and styles
    attribute DOMObject strokeStyle; // (default black)
    attribute DOMObject fillStyle; // (default black)
    CanvasGradient createLinearGradient(in float x0, in float y0, in float
x1, in float y1);
    CanvasGradient createRadialGradient(in float x0, in float y0, in float
r0, in float x1, in float y1, in float r1);
    CanvasPattern createPattern(in HTMLImageElement image, in DOMString
repetition);
    CanvasPattern createPattern(in HTMLCanvasElement image, in DOMString
repetition);

// line caps/joins
    attribute float lineWidth; // (default 1)
    attribute DOMString lineCap; // "butt", "round", "square"
(default "butt")
    attribute DOMString lineJoin; // "round", "bevel", "miter"
(default "miter")
    attribute float miterLimit; // (default 10)

// shadows
    attribute float shadowOffsetX; // (default 0)
    attribute float shadowOffsetY; // (default 0)
    attribute float shadowBlur; // (default 0)
    attribute DOMString shadowColor; // (default transparent
black)

// rects
void clearRect(in float x, in float y, in float w, in float h);
void fillRect(in float x, in float y, in float w, in float h);
void strokeRect(in float x, in float y, in float w, in float h);

// path API
void beginPath();
void closePath();
void moveTo(in float x, in float y);
void lineTo(in float x, in float y);
void quadraticCurveTo(in float cpx, in float cpy, in float x, in float
y);
void bezierCurveTo(in float cp1x, in float cp1y, in float cp2x, in
float cp2y, in float x, in float y);
void arcTo(in float x1, in float y1, in float x2, in float y2, in
float radius);
void rect(in float x, in float y, in float w, in float h);
void arc(in float x, in float y, in float radius, in float startAngle,

```

```

    in float endAngle, in boolean anticlockwise);
    void fill();
    void stroke();
    void clip();
    boolean isPointInPath(in float x, in float y);

    // text
        attribute DOMString font; // (default 10px sans-serif)
        attribute DOMString textAlign; // "start", "end", "left",
"right", "center" (default: "start")
        attribute DOMString textBaseline; // "top", "hanging",
"middle", "alphabetic", "ideographic", "bottom" (default: "alphabetic")
    void fillText(in DOMString text, in float x, in float y);
    void fillText(in DOMString text, in float x, in float y, in float
maxWidth);
    void strokeText(in DOMString text, in float x, in float y);
    void strokeText(in DOMString text, in float x, in float y, in float
maxWidth);
    TextMetrics measureText(in DOMString text);

    // drawing images
    void drawImage(in HTMLImageElement image, in float dx, in float dy);
    void drawImage(in HTMLImageElement image, in float dx, in float dy, in
float dw, in float dh);
    void drawImage(in HTMLImageElement image, in float sx, in float sy, in
float sw, in float sh, in float dx, in float dy, in float dw, in float
dh);
    void drawImage(in HTMLCanvasElement image, in float dx, in float dy);
    void drawImage(in HTMLCanvasElement image, in float dx, in float dy,
in float dw, in float dh);
    void drawImage(in HTMLCanvasElement image, in float sx, in float sy, in
float sw, in float sh, in float dx, in float dy, in float dw, in float
dh);

    // pixel manipulation
    ImageData createImageData(in float sw, in float sh);
    ImageData getImageData(in float sx, in float sy, in float sw, in float
sh);
    void putImageData(in ImageData imagedata, in float dx, in float dy);
    void putImageData(in ImageData imagedata, in float dx, in float dy, in
float dirtyX, in float dirtyY, in float dirtyWidth, in float
dirtyHeight);
};

interface CanvasGradient {
    // opaque object
    void addColorStop(in float offset, in DOMString color);
};

interface CanvasPattern {
    // opaque object

```

```

};

interface TextMetrics {
    readonly attribute float width;
};

interface ImageData {
    readonly attribute unsigned long width;
    readonly attribute unsigned long height;
    readonly attribute CanvasPixelArray data;
};

interface CanvasPixelArray {
    readonly attribute unsigned long length;
    [IndexGetter] octet XXX5(in unsigned long index);
    [IndexSetter] void XXX6(in unsigned long index, in octet value);
};

```

The **canvas** attribute must return the canvas element that the context paints on.

Unless otherwise stated, for the 2D context interface, any method call with a numeric argument whose value is infinite or a NaN value must be ignored.

Whenever the CSS value `currentColor` is used as a color in this API, the "computed value of the 'color' property" for the purposes of determining the computed value of the `currentColor` keyword is the computed value of the 'color' property on the element in question at the time that the color is specified (e.g. when the appropriate attribute is set, or when the method is called; not when the color is rendered or otherwise used). If the computed value of the 'color' property is undefined for a particular case (e.g. because the element is not in a document), then the "computed value of the 'color' property" for the purposes of determining the computed value of the `currentColor` keyword is fully opaque black. [CSS3COLOR]

4.8.11.1.1 The canvas state

Each context maintains a stack of drawing states. **Drawing states** consist of:

- The current transformation matrix (page 267).
- The current clipping region (page 277).
- The current values of the following attributes: `strokeStyle`, `fillStyle`, `globalAlpha`, `lineWidth`, `lineCap`, `lineJoin`, `miterLimit`, `shadowOffsetX`, `shadowOffsetY`, `shadowBlur`, `shadowColor`, `globalCompositeOperation`, `font`, `textAlign`, `textBaseline`.

Note: *The current path and the current bitmap are not part of the drawing state. The current path is persistent, and can only be reset using the `beginPath()` method. The current bitmap is a property of the canvas, not the context.*

The `save()` method must push a copy of the current drawing state onto the drawing state stack.

The **restore()** method must pop the top entry in the drawing state stack, and reset the drawing state it describes. If there is no saved state, the method must do nothing.

4.8.11.1.2 Transformations

The transformation matrix is applied to coordinates when creating shapes and paths.

When the context is created, the transformation matrix must initially be the identity transform. It may then be adjusted using the transformation methods.

The transformations must be performed in reverse order. For instance, if a scale transformation that doubles the width is applied, followed by a rotation transformation that rotates drawing operations by a quarter turn, and a rectangle twice as wide as it is tall is then drawn on the canvas, the actual result will be a square.

The **scale(x, y)** method must add the scaling transformation described by the arguments to the transformation matrix. The *x* argument represents the scale factor in the horizontal direction and the *y* argument represents the scale factor in the vertical direction. The factors are multiples.

The **rotate(angle)** method must add the rotation transformation described by the argument to the transformation matrix. The *angle* argument represents a clockwise rotation angle expressed in radians.

The **translate(x, y)** method must add the translation transformation described by the arguments to the transformation matrix. The *x* argument represents the translation distance in the horizontal direction and the *y* argument represents the translation distance in the vertical direction. The arguments are in coordinate space units.

The **transform(m11, m12, m21, m22, dx, dy)** method must multiply the current transformation matrix with the matrix described by:

$$\begin{matrix} m11 & m21 & dx \\ m12 & m22 & dy \\ 0 & 0 & 1 \end{matrix}$$

The **setTransform(m11, m12, m21, m22, dx, dy)** method must reset the current transform to the identity matrix, and then invoke the **transform(m11, m12, m21, m22, dx, dy)** method with the same arguments.

4.8.11.1.3 Compositing

All drawing operations are affected by the global compositing attributes, **globalAlpha** and **globalCompositeOperation**.

The **globalAlpha** attribute gives an alpha value that is applied to shapes and images before they are composited onto the canvas. The value must be in the range from 0.0 (fully transparent) to 1.0 (no additional transparency). If an attempt is made to set the attribute to a value outside this range, the attribute must retain its previous value. When the context is created, the **globalAlpha** attribute must initially have the value 1.0.

The **globalCompositeOperation** attribute sets how shapes and images are drawn onto the existing bitmap, once they have had **globalAlpha** and the current transformation matrix

applied. It must be set to a value from the following list. In the descriptions below, the source image, *A*, is the shape or image being rendered, and the destination image, *B*, is the current state of the bitmap.

source-atop

A atop *B*. Display the source image wherever both images are opaque. Display the destination image wherever the destination image is opaque but the source image is transparent. Display transparency elsewhere.

source-in

A in *B*. Display the source image wherever both the source image and destination image are opaque. Display transparency elsewhere.

source-out

A out *B*. Display the source image wherever the source image is opaque and the destination image is transparent. Display transparency elsewhere.

source-over (default)

A over *B*. Display the source image wherever the source image is opaque. Display the destination image elsewhere.

destination-atop

B atop *A*. Same as *source-atop* but using the destination image instead of the source image and vice versa.

destination-in

B in *A*. Same as *source-in* but using the destination image instead of the source image and vice versa.

destination-out

B out *A*. Same as *source-out* but using the destination image instead of the source image and vice versa.

destination-over

B over *A*. Same as *source-over* but using the destination image instead of the source image and vice versa.

lighter

A plus *B*. Display the sum of the source image and destination image, with color values approaching 1 as a limit.

copy

A (*B* is ignored). Display the source image instead of the destination image.

xor

A xor *B*. Exclusive OR of the source image and destination image.

vendorName-operationName

Vendor-specific extensions to the list of composition operators should use this syntax.

These values are all case-sensitive — they must be used exactly as shown. User agents must not recognize values that are not a case-sensitive (page 31) match for one of the values given above.

The operators in the above list must be treated as described by the Porter-Duff operator given at the start of their description (e.g. A over B). [PORTERDUFF]

On setting, if the user agent does not recognize the specified value, it must be ignored, leaving the value of `globalCompositeOperation` unaffected.

When the context is created, the `globalCompositeOperation` attribute must initially have the value `source-over`.

4.8.11.1.4 Colors and styles

The `strokeStyle` attribute represents the color or style to use for the lines around shapes, and the `fillStyle` attribute represents the color or style to use inside the shapes.

Both attributes can be either strings, `CanvasGradients`, or `CanvasPatterns`. On setting, strings must be parsed as CSS `<color>` values and the color assigned, and `CanvasGradient` and `CanvasPattern` objects must be assigned themselves. [CSS3COLOR] If the value is a string but is not a valid color, or is neither a string, a `CanvasGradient`, nor a `CanvasPattern`, then it must be ignored, and the attribute must retain its previous value.

On getting, if the value is a color, then the serialization of the color (page 269) must be returned. Otherwise, if it is not a color but a `CanvasGradient` or `CanvasPattern`, then the respective object must be returned. (Such objects are opaque and therefore only useful for assigning to other attributes or for comparison to other gradients or patterns.)

The **serialization of a color** for a color value is a string, computed as follows: if it has alpha equal to 1.0, then the string is a lowercase six-digit hex value, prefixed with a "#" character (U+0023 NUMBER SIGN), with the first two digits representing the red component, the next two digits representing the green component, and the last two digits representing the blue component, the digits being in the range 0-9 a-f (U+0030 to U+0039 and U+0061 to U+0066). Otherwise, the color value has alpha less than 1.0, and the string is the color value in the CSS `rgba()` functional-notation format: the literal string `rgba` (U+0072 U+0067 U+0062 U+0061) followed by a U+0028 LEFT PARENTHESIS, a base-ten integer in the range 0-255 representing the red component (using digits 0-9, U+0030 to U+0039, in the shortest form possible), a literal U+002C COMMA and U+0020 SPACE, an integer for the green component, a comma and a space, an integer for the blue component, another comma and space, a U+0030 DIGIT ZERO, a U+002E FULL STOP (representing the decimal point), one or more digits in the range 0-9 (U+0030 to U+0039) representing the fractional part of the alpha value, and finally a U+0029 RIGHT PARENTHESIS.

When the context is created, the `strokeStyle` and `fillStyle` attributes must initially have the string value `#000000`.

There are two types of gradients, linear gradients and radial gradients, both represented by objects implementing the opaque `CanvasGradient` interface.

Once a gradient has been created (see below), stops are placed along it to define how the colors are distributed along the gradient. The color of the gradient at each stop is the color specified for that stop. Between each such stop, the colors and the alpha component must be linearly interpolated over the RGBA space without premultiplying the alpha value to find the color to use at that offset. Before the first stop, the color must be the color of the first stop. After the last stop, the color must be the color of the last stop. When there are no stops, the gradient is transparent black.

The **addColorStop(offset, color)** method on the CanvasGradient interface adds a new stop to a gradient. If the *offset* is less than 0, greater than 1, infinite, or NaN, then an INDEX_SIZE_ERR exception must be raised. If the *color* cannot be parsed as a CSS color, then a SYNTAX_ERR exception must be raised. Otherwise, the gradient must have a new stop placed, at offset *offset* relative to the whole gradient, and with the color obtained by parsing *color* as a CSS <color> value. If multiple stops are added at the same offset on a gradient, they must be placed in the order added, with the first one closest to the start of the gradient, and each subsequent one infinitesimally further along towards the end point (in effect causing all but the first and last stop added at each point to be ignored).

The **createLinearGradient(x0, y0, x1, y1)** method takes four arguments that represent the start point (x_0, y_0) and end point (x_1, y_1) of the gradient. If any of the arguments to `createLinearGradient()` are infinite or NaN, the method must raise a NOT_SUPPORTED_ERR exception. Otherwise, the method must return a linear CanvasGradient initialized with the specified line.

Linear gradients must be rendered such that all points on a line perpendicular to the line that crosses the start and end points have the color at the point where those two lines cross (with the colors coming from the interpolation and extrapolation (page 269) described above). The points in the linear gradient must be transformed as described by the current transformation matrix (page 267) when rendering.

If $x_0 = x_1$ and $y_0 = y_1$, then the linear gradient must paint nothing.

The **createRadialGradient(x0, y0, r0, x1, y1, r1)** method takes six arguments, the first three representing the start circle with origin (x_0, y_0) and radius r_0 , and the last three representing the end circle with origin (x_1, y_1) and radius r_1 . The values are in coordinate space units. If any of the arguments are infinite or NaN, a NOT_SUPPORTED_ERR exception must be raised. If either of r_0 or r_1 are negative, an INDEX_SIZE_ERR exception must be raised. Otherwise, the method must return a radial CanvasGradient initialized with the two specified circles.

Radial gradients must be rendered by following these steps:

1. If $x_0 = x_1$ and $y_0 = y_1$ and $r_0 = r_1$, then the radial gradient must paint nothing. Abort these steps.
2. Let $x(\omega) = (x_1 - x_0)\omega + x_0$

$$\text{Let } y(\omega) = (y_1 - y_0)\omega + y_0$$

$$\text{Let } r(\omega) = (r_1 - r_0)\omega + r_0$$

Let the color at ω be the color at that position on the gradient (with the colors coming from the interpolation and extrapolation (page 269) described above).

3. For all values of ω where $r(\omega) > 0$, starting with the value of ω nearest to positive infinity and ending with the value of ω nearest to negative infinity, draw the circumference of the circle with radius $r(\omega)$ at position $(x(\omega), y(\omega))$, with the color at ω , but only painting on the parts of the canvas that have not yet been painted on by earlier circles in this step for this rendering of the gradient.

Note: This effectively creates a cone, touched by the two circles defined in the creation of the gradient, with the part of the cone before the start

circle (0.0) using the color of the first offset, the part of the cone after the end circle (1.0) using the color of the last offset, and areas outside the cone untouched by the gradient (transparent black).

Gradients must be painted only where the relevant stroking or filling effects require that they be drawn.

The points in the radial gradient must be transformed as described by the current transformation matrix (page 267) when rendering.

Patterns are represented by objects implementing the opaque `CanvasPattern` interface.

To create objects of this type, the `createPattern(image, repetition)` method is used. The first argument gives the image to use as the pattern (either an `HTMLImageElement` or an `HTMLCanvasElement`). Modifying this image after calling the `createPattern()` method must not affect the pattern. The second argument must be a string with one of the following values: `repeat`, `repeat-x`, `repeat-y`, `no-repeat`. If the empty string or null is specified, `repeat` must be assumed. If an unrecognized value is given, then the user agent must raise a `SYNTAX_ERR` exception. User agents must recognize the four values described above exactly (e.g. they must not do case folding). The method must return a `CanvasPattern` object suitably initialized.

The *image* argument must be an instance of an `HTMLImageElement` or `HTMLCanvasElement`. If the *image* is of the wrong type or null, the implementation must raise a `TYPE_MISMATCH_ERR` exception.

If the *image* argument is an `HTMLImageElement` object whose `complete` attribute is false, then the implementation must raise an `INVALID_STATE_ERR` exception.

If the *image* argument is an `HTMLCanvasElement` object with either a horizontal dimension or a vertical dimension equal to zero, then the implementation must raise an `INVALID_STATE_ERR` exception.

Patterns must be painted so that the top left of the first image is anchored at the origin of the coordinate space, and images are then repeated horizontally to the left and right (if the `repeat-x` string was specified) or vertically up and down (if the `repeat-y` string was specified) or in all four directions all over the canvas (if the `repeat` string was specified). The images are not scaled by this process; one CSS pixel of the image must be painted on one coordinate space unit. Of course, patterns must actually be painted only where the stroking or filling effect requires that they be drawn, and are affected by the current transformation matrix.

When the `createPattern()` method is passed, as its *image* argument, an animated image, the poster frame of the animation, or the first frame of the animation if there is no poster frame, must be used.

4.8.11.1.5 Line styles

The `lineWidth` attribute gives the width of lines, in coordinate space units. On setting, zero, negative, infinite, and NaN values must be ignored, leaving the value unchanged.

When the context is created, the `lineWidth` attribute must initially have the value 1.0.

The **lineCap** attribute defines the type of endings that UAs will place on the end of lines. The three valid values are butt, round, and square. The butt value means that the end of each line has a flat edge perpendicular to the direction of the line (and that no additional line cap is added). The round value means that a semi-circle with the diameter equal to the width of the line must then be added on to the end of the line. The square value means that a rectangle with the length of the line width and the width of half the line width, placed flat against the edge perpendicular to the direction of the line, must be added at the end of each line. On setting, any other value than the literal strings butt, round, and square must be ignored, leaving the value unchanged.

When the context is created, the **lineCap** attribute must initially have the value butt.

The **lineJoin** attribute defines the type of corners that UAs will place where two lines meet. The three valid values are bevel, round, and miter.

On setting, any other value than the literal strings bevel, round, and miter must be ignored, leaving the value unchanged.

When the context is created, the **lineJoin** attribute must initially have the value miter.

A join exists at any point in a subpath shared by two consecutive lines. When a subpath is closed, then a join also exists at its first point (equivalent to its last point) connecting the first and last lines in the subpath.

In addition to the point where the join occurs, two additional points are relevant to each join, one for each line: the two corners found half the line width away from the join point, one perpendicular to each line, each on the side furthest from the other line.

A filled triangle connecting these two opposite corners with a straight line, with the third point of the triangle being the join point, must be rendered at all joins. The **lineJoin** attribute controls whether anything else is rendered. The three aforementioned values have the following meanings:

The bevel value means that this is all that is rendered at joins.

The round value means that a filled arc connecting the two aforementioned corners of the join, abutting (and not overlapping) the aforementioned triangle, with the diameter equal to the line width and the origin at the point of the join, must be rendered at joins.

The miter value means that a second filled triangle must (if it can given the miter length) be rendered at the join, with one line being the line between the two aforementioned corners, abutting the first triangle, and the other two being continuations of the outside edges of the two joining lines, as long as required to intersect without going over the miter length.

The miter length is the distance from the point where the lines touch on the inside of the join to the intersection of the line edges on the outside of the join. The miter limit ratio is the maximum allowed ratio of the miter length to half the line width. If the miter length would cause the miter limit ratio to be exceeded, this second triangle must not be rendered.

The miter limit ratio can be explicitly set using the **miterLimit** attribute. On setting, zero, negative, infinite, and NaN values must be ignored, leaving the value unchanged.

When the context is created, the **miterLimit** attribute must initially have the value 10.0.

4.8.11.1.6 Shadows

All drawing operations are affected by the four global shadow attributes.

The **shadowColor** attribute sets the color of the shadow.

When the context is created, the shadowColor attribute initially must be fully-transparent black.

On getting, the serialization of the color (page 269) must be returned.

On setting, the new value must be parsed as a CSS `<color>` value and the color assigned. If the value is not a valid color, then it must be ignored, and the attribute must retain its previous value. [CSS3COLOR]

The **shadowOffsetX** and **shadowOffsetY** attributes specify the distance that the shadow will be offset in the positive horizontal and positive vertical distance respectively. Their values are in coordinate space units. They are not affected by the current transformation matrix.

When the context is created, the shadow offset attributes must initially have the value 0.

On getting, they must return their current value. On setting, the attribute being set must be set to the new value, except if the value is infinite or NaN, in which case the new value must be ignored.

The **shadowBlur** attribute specifies the size of the blurring effect. (The units do not map to coordinate space units, and are not affected by the current transformation matrix.)

When the context is created, the shadowBlur attribute must initially have the value 0.

On getting, the attribute must return its current value. On setting the attribute must be set to the new value, except if the value is negative, infinite or NaN, in which case the new value must be ignored.

When shadows are drawn, they must be rendered as follows:

1. Let A be the source image for which a shadow is being created.
2. Let B be an infinite transparent black bitmap, with a coordinate space and an origin identical to A .
3. Copy the alpha channel of A to B , offset by `shadowOffsetX` in the positive x direction, and `shadowOffsetY` in the positive y direction.
4. If `shadowBlur` is greater than 0:
 1. If `shadowBlur` is less than 8, let σ be half the value of `shadowBlur`; otherwise, let σ be the square root of multiplying the value of `shadowBlur` by 2.
 2. Perform a 2D Gaussian Blur on B , using σ as the standard deviation.

User agents may limit values of σ to an implementation-specific maximum value to avoid exceeding hardware limitations during the Gaussian blur operation.

5. Set the red, green, and blue components of every pixel in B to the red, green, and blue components (respectively) of the color of `shadowColor`.

6. Multiply the alpha component of every pixel in B by the alpha component of the color of shadowColor.
7. The shadow is in the bitmap B , and is rendered as part of the drawing model described below.

4.8.11.1.7 Simple shapes (rectangles)

There are three methods that immediately draw rectangles to the bitmap. They each take four arguments; the first two give the x and y coordinates of the top left of the rectangle, and the second two give the width w and height h of the rectangle, respectively.

The current transformation matrix (page 267) must be applied to the following four coordinates, which form the path that must then be closed to get the specified rectangle: (x, y) , $(x+w, y)$, $(x+w, y+h)$, $(x, y+h)$.

Shapes are painted without affecting the current path, and are subject to the clipping region (page 277), and, with the exception of `clearRect()`, also shadow effects (page 273), global alpha (page 267), and global composition operators (page 267).

The `clearRect(x, y, w, h)` method must clear the pixels in the specified rectangle that also intersect the current clipping region to a fully transparent black, erasing any previous image. If either height or width are zero, this method has no effect.

The `fillRect(x, y, w, h)` method must paint the specified rectangular area using the `fillStyle`. If either height or width are zero, this method has no effect.

The `strokeRect(x, y, w, h)` method must stroke the specified rectangle's path using the `strokeStyle`, `lineWidth`, `lineJoin`, and (if appropriate) `miterLimit` attributes. If both height and width are zero, this method has no effect, since there is no path to stroke (it's a point). If only one of the two is zero, then the method will draw a line instead (the path for the outline is just a straight line along the non-zero dimension).

4.8.11.1.8 Complex shapes (paths)

The context always has a current path. There is only one current path, it is not part of the drawing state.

A **path** has a list of zero or more subpaths. Each subpath consists of a list of one or more points, connected by straight or curved lines, and a flag indicating whether the subpath is closed or not. A closed subpath is one where the last point of the subpath is connected to the first point of the subpath by a straight line. Subpaths with fewer than two points are ignored when painting the path.

Initially, the context's path must have zero subpaths.

The points and lines added to the path by these methods must be transformed according to the current transformation matrix (page 267) as they are added.

The `beginPath()` method must empty the list of subpaths so that the context once again has zero subpaths.

The **moveTo(x, y)** method must create a new subpath with the specified point as its first (and only) point.

The **closePath()** method must do nothing if the context has no subpaths. Otherwise, it must mark the last subpath as closed, create a new subpath whose first point is the same as the previous subpath's first point, and finally add this new subpath to the path. (If the last subpath had more than one point in its list of points, then this is equivalent to adding a straight line connecting the last point back to the first point, thus "closing" the shape, and then repeating the last `moveTo()` call.)

New points and the lines connecting them are added to subpaths using the methods described below. In all cases, the methods only modify the last subpath in the context's paths.

The **lineTo(x, y)** method must do nothing if the context has no subpaths. Otherwise, it must connect the last point in the subpath to the given point (x, y) using a straight line, and must then add the given point (x, y) to the subpath.

The **quadraticCurveTo(cpx, cpy, x, y)** method must do nothing if the context has no subpaths. Otherwise it must connect the last point in the subpath to the given point (x, y) using a quadratic Bézier curve with control point (cpx, cpy), and must then add the given point (x, y) to the subpath. [BEZIER]

The **bezierCurveTo(cp1x, cp1y, cp2x, cp2y, x, y)** method must do nothing if the context has no subpaths. Otherwise, it must connect the last point in the subpath to the given point (x, y) using a cubic Bézier curve with control points (cp1x, cp1y) and (cp2x, cp2y). Then, it must add the point (x, y) to the subpath. [BEZIER]

The **arcTo(x1, y1, x2, y2, radius)** method must do nothing if the context has no subpaths. If the context *does* have a subpath, then the behavior depends on the arguments and the last point in the subpath.

Negative values for *radius* must cause the implementation to raise an INDEX_SIZE_ERR exception.

Let the point (x0, y0) be the last point in the subpath.

If the point (x0, y0) is equal to the point (x1, y1), or if the point (x1, y1) is equal to the point (x2, y2), or if the radius *radius* is zero, then the method must add the point (x1, y1) to the subpath, and connect that point to the previous point (x0, y0) by a straight line.

Otherwise, if the points (x0, y0), (x1, y1), and (x2, y2) all lie on a single straight line, then: if the direction from (x0, y0) to (x1, y1) is the same as the direction from (x1, y1) to (x2, y2), then the method must add the point (x1, y1) to the subpath, and connect that point to the previous point (x0, y0) by a straight line; otherwise, the direction from (x0, y0) to (x1, y1) is the opposite of the direction from (x1, y1) to (x2, y2), and the method must add a point (x_∞, y_∞) to the subpath, and connect that point to the previous point (x0, y0) by a straight line, where (x_∞, y_∞) is the point that is infinitely far away from (x1, y1), that lies on the same line as (x0, y0), (x1, y1), and (x2, y2), and that is on the same side of (x1, y1) on that line as (x2, y2).

Otherwise, let *The Arc* be the shortest arc given by circumference of the circle that has radius *radius*, and that has one point tangent to the half-infinite line that crosses the point (x0, y0) and ends at the point (x1, y1), and that has a different point tangent to the half-infinite line

that ends at the point (x_1, y_1) and crosses the point (x_2, y_2) . The points at which this circle touches these two lines are called the start and end tangent points respectively.

The method must connect the point (x_0, y_0) to the start tangent point by a straight line, adding the start tangent point to the subpath, and then must connect the start tangent point to the end tangent point by *The Arc*, adding the end tangent point to the subpath.

The **arc(*x, y, radius, startAngle, endAngle, anticlockwise*)** method draws an arc. If the context has any subpaths, then the method must add a straight line from the last point in the subpath to the start point of the arc. In any case, it must draw the arc between the start point of the arc and the end point of the arc, and add the start and end points of the arc to the subpath. The arc and its start and end points are defined as follows:

Consider a circle that has its origin at (x, y) and that has radius *radius*. The points at *startAngle* and *endAngle* along this circle's circumference, measured in radians clockwise from the positive x-axis, are the start and end points respectively.

If the *anticlockwise* argument is false and *endAngle-startAngle* is equal to or greater than 2π , or, if the *anticlockwise* argument is *true* and *startAngle-endAngle* is equal to or greater than 2π , then the arc is the whole circumference of this circle.

Otherwise, the arc is the path along the circumference of this circle from the start point to the end point, going anti-clockwise if the *anticlockwise* argument is *true*, and clockwise otherwise. Since the points are on the circle, as opposed to being simply angles from zero, the arc can never cover an angle greater than 2π radians. If the two points are the same, or if the radius is zero, then the arc is defined as being of zero length in both directions.

Negative values for *radius* must cause the implementation to raise an INDEX_SIZE_ERR exception.

The **rect(*x, y, w, h*)** method must create a new subpath containing just the four points (x, y) , $(x+w, y)$, $(x+w, y+h)$, $(x, y+h)$, with those four points connected by straight lines, and must then mark the subpath as closed. It must then create a new subpath with the point (x, y) as the only point in the subpath.

The **fill()** method must fill all the subpaths of the current path, using *fillStyle*, and using the non-zero winding number rule. Open subpaths must be implicitly closed when being filled (without affecting the actual subpaths).

Note: Thus, if two overlapping but otherwise independent subpaths have opposite windings, they cancel out and result in no fill. If they have the same winding, that area just gets painted once.

The **stroke()** method must calculate the strokes of all the subpaths of the current path, using the *lineWidth*, *lineCap*, *lineJoin*, and (if appropriate) *miterLimit* attributes, and then fill the combined stroke area using the *strokeStyle* attribute.

Note: Since the subpaths are all stroked as one, overlapping parts of the paths in one stroke operation are treated as if their union was what was painted.

Paths, when filled or stroked, must be painted without affecting the current path, and must be subject to shadow effects (page 273), global alpha (page 267), the clipping region (page

277), and global composition operators (page 267). (Transformations affect the path when the path is created, not when it is painted, though the stroke *style* is still affected by the transformation during painting.)

Zero-length line segments must be pruned before stroking a path. Empty subpaths must be ignored.

The `clip()` method must create a new **clipping region** by calculating the intersection of the current clipping region and the area described by the current path, using the non-zero winding number rule. Open subpaths must be implicitly closed when computing the clipping region, without affecting the actual subpaths. The new clipping region replaces the current clipping region.

When the context is initialized, the clipping region must be set to the rectangle with the top left corner at (0,0) and the width and height of the coordinate space.

The `isPointInPath(x, y)` method must return true if the point given by the *x* and *y* coordinates passed to the method, when treated as coordinates in the canvas coordinate space unaffected by the current transformation, is inside the current path as determined by the non-zero winding number rule; and must return false otherwise. Points on the path itself are considered to be inside the path. If either of the arguments is infinite or NaN, then the method must return false.

4.8.11.1.9 Text

The `font` DOM attribute, on setting, must be parsed the same way as the 'font' property of CSS (but without supporting property-independent stylesheet syntax like 'inherit'), and the resulting font must be assigned to the context, with the 'line-height' component forced to 'normal'. If the new value is syntactically incorrect, then it must be ignored, without assigning a new font value. [CSS]

Font names must be interpreted in the context of the canvas element's stylesheets; any fonts embedded using @font-face must therefore be available. [CSSWEBFONTS]

Only vector fonts should be used by the user agent; if a user agent were to use bitmap fonts then transformations would likely make the font look very ugly.

On getting, the `font` attribute must return the serialized form of the current font of the context. [CSSOM]

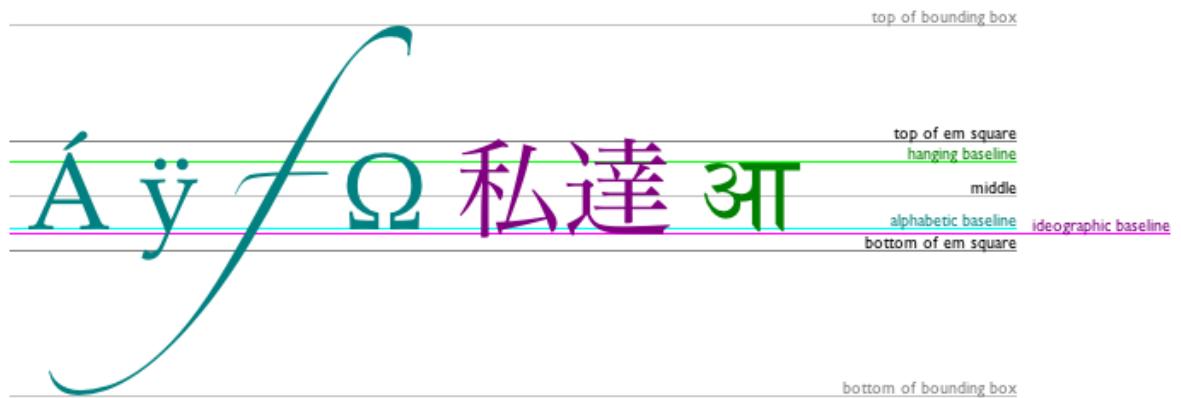
When the context is created, the font of the context must be set to 10px sans-serif. When the 'font-size' component is set to lengths using percentages, 'em' or 'ex' units, or the 'larger' or 'smaller' keywords, these must be interpreted relative to the computed value of the 'font-size' property of the corresponding canvas element at the time that the attribute is set. When the 'font-weight' component is set to the relative values 'bolder' and 'lighter', these must be interpreted relative to the computed value of the 'font-weight' property of the corresponding canvas element at the time that the attribute is set. If the computed values are undefined for a particular case (e.g. because the canvas element is not in a document), then the relative keywords must be interpreted relative to the normal-weight 10px sans-serif default.

The `textAlign` DOM attribute, on getting, must return the current value. On setting, if the value is one of start, end, left, right, or center, then the value must be changed to the

new value. Otherwise, the new value must be ignored. When the context is created, the `textAlign` attribute must initially have the value `start`.

The `textBaseline` DOM attribute, on getting, must return the current value. On setting, if the value is one of `top`, `hanging`, `middle`, `alphabetic`, `ideographic`, or `bottom`, then the value must be changed to the new value. Otherwise, the new value must be ignored. When the context is created, the `textBaseline` attribute must initially have the value `alphabetic`.

The `textBaseline` attribute's allowed keywords correspond to alignment points in the font:



The keywords map to these alignment points as follows:

top

The top of the em square

hanging

The hanging baseline

middle

The middle of the em square

alphabetic

The alphabetic baseline

ideographic

The ideographic baseline

bottom

The bottom of the em square

The `fillText()` and `strokeText()` methods take three or four arguments, `text`, `x`, `y`, and optionally `maxWidth`, and render the given `text` at the given `(x, y)` coordinates ensuring that the text isn't wider than `maxWidth` if specified, using the current font, `textAlign`, and `textBaseline` values. Specifically, when the methods are called, the user agent must run the following steps:

1. Let `font` be the current font of the browsing context, as given by the `font` attribute.
2. Replace all the space characters (page 32) in `text` with U+0020 SPACE characters.

3. Form a hypothetical infinitely wide CSS line box containing a single inline box containing the text *text*, with all the properties at their initial values except the 'font' property of the inline box set to *font* and the 'direction' property of the inline box set to the directionality (page 89) of the canvas element. [CSS]
4. If the *maxWidth* argument was specified and the hypothetical width of the inline box in the hypothetical line box is greater than *maxWidth* CSS pixels, then change *font* to have a more condensed font (if one is available or if a reasonably readable one can be synthesized by applying a horizontal scale factor to the font) or a smaller font, and return to the previous step.
5. Let the *anchor point* be a point on the inline box, determined by the *textAlign* and *textBaseline* values, as follows:

Horizontal position:

If *textAlign* is left

If *textAlign* is start and the directionality (page 89) of the canvas element is 'ltr'

If *textAlign* is end and the directionality (page 89) of the canvas element is 'rtl'

Let the *anchor point*'s horizontal position be the left edge of the inline box.

If *textAlign* is right

If *textAlign* is end and the directionality (page 89) of the canvas element is 'ltr'

If *textAlign* is start and the directionality (page 89) of the canvas element is 'rtl'

Let the *anchor point*'s horizontal position be the right edge of the inline box.

If *textAlign* is center

Let the *anchor point*'s horizontal position be half way between the left and right edges of the inline box.

Vertical position:

If *textBaseline* is top

Let the *anchor point*'s vertical position be the top of the em box of the first available font of the inline box.

If *textBaseline* is hanging

Let the *anchor point*'s vertical position be the hanging baseline of the first available font of the inline box.

If *textBaseline* is middle

Let the *anchor point*'s vertical position be half way between the bottom and the top of the em box of the first available font of the inline box.

If *textBaseline* is alphabetic

Let the *anchor point*'s vertical position be the alphabetic baseline of the first available font of the inline box.

If `textBaseline` is `ideographic`

Let the *anchor point*'s vertical position be the ideographic baseline of the first available font of the inline box.

If `textBaseline` is `bottom`

Let the *anchor point*'s vertical position be the bottom of the em box of the first available font of the inline box.

6. Paint the hypothetical inline box as the shape given by the text's glyphs, as transformed by the current transformation matrix (page 267), and anchored and sized so that before applying the current transformation matrix (page 267), the *anchor point* is at (x, y) and each CSS pixel is mapped to one coordinate space unit.

For `fillText()` `fillStyle` must be applied to the glyphs and `strokeStyle` must be ignored. For `strokeText()` the reverse holds and `strokeStyle` must be applied to the glyph outlines and `fillStyle` must be ignored.

Text is painted without affecting the current path, and is subject to shadow effects (page 273), global alpha (page 267), the clipping region (page 277), and global composition operators (page 267).

The `measureText()` method takes one argument, `text`. When the method is invoked, the user agent must replace all the space characters (page 32) in `text` with U+0020 SPACE characters, and then must form a hypothetical infinitely wide CSS line box containing a single inline box containing the text `text`, with all the properties at their initial values except the 'font' property of the inline element set to the current font of the browsing context, as given by the `font` attribute, and must then return a new `TextMetrics` object with its `width` attribute set to the width of that inline box, in CSS pixels. [CSS]

The `TextMetrics` interface is used for the objects returned from `measureText()`. It has one attribute, `width`, which is set by the `measureText()` method.

Note: *Glyphs rendered using `fillText()` and `strokeText()` can spill out of the box given by the font size (the em square size) and the width returned by `measureText()` (the text width). This version of the specification does not provide a way to obtain the bounding box dimensions of the text. If the text is to be rendered and removed, care needs to be taken to replace the entire area of the canvas that the clipping region covers, not just the box given by the em square height and measured text width.*

Note: *A future version of the 2D context API may provide a way to render fragments of documents, rendered using CSS, straight to the canvas. This would be provided in preference to a dedicated way of doing multiline layout.*

4.8.11.10 Images

To draw images onto the canvas, the `drawImage` method can be used.

This method is overloaded with three variants: `drawImage(image, dx, dy)`, `drawImage(image, dx, dy, dw, dh)`, and `drawImage(image, sx, sy, sw, sh, dx, dy)`,

dw, dh). (Actually it is overloaded with six; each of those three can take either an `HTMLImageElement` or an `HTMLCanvasElement` for the *image* argument.) If not specified, the *dw* and *dh* arguments must default to the values of *sw* and *sh*, interpreted such that one CSS pixel in the image is treated as one unit in the canvas coordinate space. If the *sx*, *sy*, *sw*, and *sh* arguments are omitted, they must default to 0, 0, the image's intrinsic width in image pixels, and the image's intrinsic height in image pixels, respectively.

The *image* argument must be an instance of an `HTMLImageElement` or `HTMLCanvasElement`. If the *image* is of the wrong type or null, the implementation must raise a `TYPE_MISMATCH_ERR` exception.

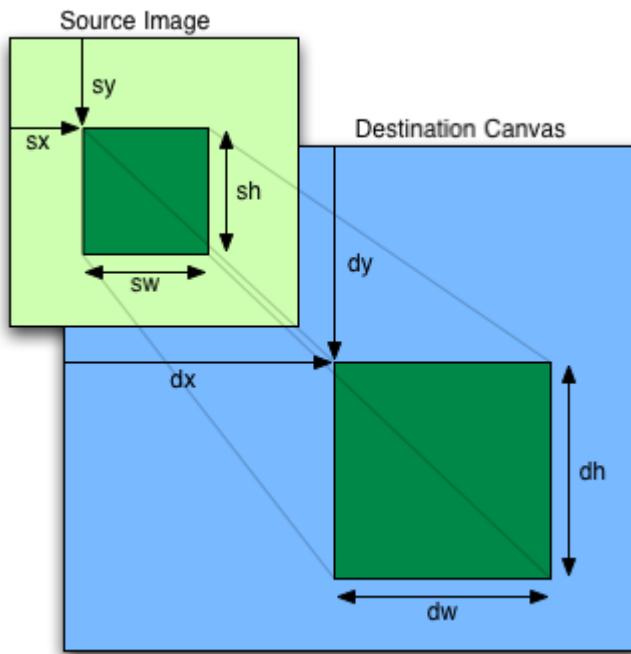
If the *image* argument is an `HTMLImageElement` object whose `complete` attribute is false, then the implementation must raise an `INVALID_STATE_ERR` exception.

The source rectangle is the rectangle whose corners are the four points (sx, sy) , $(sx+sw, sy)$, $(sx+sw, sy+sh)$, $(sx, sy+sh)$.

If the source rectangle is not entirely within the source image, or if one of the *sw* or *sh* arguments is zero, the implementation must raise an `INDEX_SIZE_ERR` exception.

The destination rectangle is the rectangle whose corners are the four points (dx, dy) , $(dx+dw, dy)$, $(dx+dw, dy+dh)$, $(dx, dy+dh)$.

When `drawImage()` is invoked, the region of the image specified by the source rectangle must be painted on the region of the canvas specified by the destination rectangle, after applying the current transformation matrix (page 267) to the points of the destination rectangle.



Note: When a canvas is drawn onto itself, the drawing model requires the source to be copied before the image is drawn back onto the canvas, so it is possible to copy parts of a canvas onto overlapping parts of itself.

When the `drawImage()` method is passed, as its `image` argument, an animated image, the poster frame of the animation, or the first frame of the animation if there is no poster frame, must be used.

Images are painted without affecting the current path, and are subject to shadow effects (page 273), global alpha (page 267), the clipping region (page 277), and global composition operators (page 267).

4.8.11.1.11 Pixel manipulation

The `createImageData(sw, sh)` method must return an `ImageData` object representing a rectangle with a width in CSS pixels equal to the absolute magnitude of `sw` and a height in CSS pixels equal to the absolute magnitude of `sh`, filled with transparent black.

The `getImageData(sx, sy, sw, sh)` method must return an `ImageData` object representing the underlying pixel data for the area of the canvas denoted by the rectangle whose corners are the four points (sx, sy) , $(sx+sw, sy)$, $(sx+sw, sy+sh)$, $(sx, sy+sh)$, in canvas coordinate space units. Pixels outside the canvas must be returned as transparent black. Pixels must be returned as non-premultiplied alpha values.

If any of the arguments to `createImageData()` or `getImageData()` are infinite or `Nan`, the method must instead raise a `NOT_SUPPORTED_ERR` exception. If either the `sw` or `sh` arguments are zero, the method must instead raise an `INDEX_SIZE_ERR` exception.

`ImageData` objects must be initialized so that their `width` attribute is set to `w`, the number of physical device pixels per row in the image data, their `height` attribute is set to `h`, the number of rows in the image data, and their `data` attribute is initialized to a `CanvasPixelArray` object holding the image data. At least one pixel's worth of image data must be returned.

The `CanvasPixelArray` object provides ordered, indexed access to the color components of each pixel of the image data. The data must be represented in left-to-right order, row by row top to bottom, starting with the top left, with each pixel's red, green, blue, and alpha components being given in that order for each pixel. Each component of each device pixel represented in this array must be in the range 0..255, representing the 8 bit value for that component. The components must be assigned consecutive indices starting with 0 for the top left pixel's red component.

The `CanvasPixelArray` object thus represents $h \times w \times 4$ integers. The `length` attribute of a `CanvasPixelArray` object must return this number.

The `XXX5(index)` method must return the value of the `index`th component in the array.

The `XXX6(index, value)` method must set the value of the `index`th component in the array to `value`. JS `undefined` values must be converted to zero. Other values must first be converted to numbers using JavaScript's `ToNumber` algorithm, and if the result is a `NaN` value, then the value must be converted to zero. If the result is less than 0, it must be clamped to zero. If the result is more than 255, it must be clamped to 255. If the number is not an integer, it should be rounded to the nearest integer using the IEEE 754r `convertToIntegerTiesToEven` rounding mode. [ECMA262] [IEEE754R]

- ** The above is not intended to cause these methods to get any unusual behaviour, it's just supposed to be the normal behaviour for passing values to a method expecting an octet type.

Note: The width and height (w and h) might be different from the sw and sh arguments to the above methods, e.g. if the canvas is backed by a high-resolution bitmap, or if the sw and sh arguments are negative.

The `putImageData(imagedata, dx, dy)` and `putImageData(imagedata, dx, dy, dirtyX, dirtyY, dirtyWidth, dirtyHeight)` methods write data from `ImageData` structures back to the canvas.

If any of the arguments to the method are infinite or NaN, the method must raise a `NOT_SUPPORTED_ERR` exception.

If the first argument to the method is null or not an `ImageData` object then the `putImageData()` method must raise a `TYPE_MISMATCH_ERR` exception.

When the last four arguments are omitted, they must be assumed to have the values 0, 0, the `width` member of the `imagedata` structure, and the `height` member of the `imagedata` structure, respectively.

When invoked with arguments that do not, per the last few paragraphs, cause an exception to be raised, the `putImageData()` method must act as follows:

1. Let dx_{device} be the x-coordinate of the device pixel in the underlying pixel data of the canvas corresponding to the `dx` coordinate in the canvas coordinate space.
Let dy_{device} be the y-coordinate of the device pixel in the underlying pixel data of the canvas corresponding to the `dy` coordinate in the canvas coordinate space.
2. If `dirtyWidth` is negative, let `dirtyX` be `dirtyX+dirtyWidth`, and let `dirtyWidth` be equal to the absolute magnitude of `dirtyWidth`.
If `dirtyHeight` is negative, let `dirtyY` be `dirtyY+dirtyHeight`, and let `dirtyHeight` be equal to the absolute magnitude of `dirtyHeight`.
3. If `dirtyX` is negative, let `dirtyWidth` be `dirtyWidth+dirtyX`, and let `dirtyX` be zero.
If `dirtyY` is negative, let `dirtyHeight` be `dirtyHeight+dirtyY`, and let `dirtyY` be zero.
4. If `dirtyX+dirtyWidth` is greater than the `width` attribute of the `imagedata` argument, let `dirtyWidth` be the value of that `width` attribute, minus the value of `dirtyX`.
If `dirtyY+dirtyHeight` is greater than the `height` attribute of the `imagedata` argument, let `dirtyHeight` be the value of that `height` attribute, minus the value of `dirtyY`.
5. If, after those changes, either `dirtyWidth` or `dirtyHeight` is negative or zero, stop these steps without affecting the canvas.
6. Otherwise, for all integer values of x and y where $dirtyX \leq x < dirtyX+dirtyWidth$ and $dirtyY \leq y < dirtyY+dirtyHeight$, copy the four channels of the pixel with coordinate (x, y) in the `imagedata` data structure to the pixel with coordinate $(dx_{device}+x, dy_{device}+y)$ in the underlying pixel data of the canvas.

The handling of pixel rounding when the specified coordinates do not exactly map to the device coordinate space is not defined by this specification, except that the following must result in no visible changes to the rendering:

```
context.putImageData(context.getImageData(x, y, w, h), x, y);
```

...for any value of *x*, *y*, *w*, and *h*, and the following two calls:

```
context.createImageData(w, h);
context.getImageData(0, 0, w, h);
```

...must return `ImageData` objects with the same dimensions, for any value of *w* and *h*. In other words, while user agents may round the arguments of these methods so that they map to device pixel boundaries, any rounding performed must be performed consistently for all of the `createImageData()`, `getImageData()` and `putImageData()` operations.

The current path, transformation matrix (page 267), shadow attributes (page 273), global alpha (page 267), the clipping region (page 277), and global composition operator (page 267) must not affect the `getImageData()` and `putImageData()` methods.

The data returned by `getImageData()` is at the resolution of the canvas backing store, which is likely to not be one device pixel to each CSS pixel if the display used is a high resolution display.

In the following example, the script generates an `ImageData` object so that it can draw onto it.

```
// canvas is a reference to a <canvas> element
var context = canvas.getContext('2d');

// create a blank slate
var data = context.createImageData(canvas.width, canvas.height);

// create some plasma
FillPlasma(data, 'green'); // green plasma

// add a cloud to the plasma
AddCloud(data, data.width/2, data.height/2); // put a cloud in the middle

// paint the plasma+cloud on the canvas
context.putImageData(data, 0, 0);

// support methods
function FillPlasma(data, color) { ... }
function AddCloud(data, x, y) { ... }
```

Here is an example of using `getImageData()` and `putImageData()` to implement an edge detection filter.

```
<!DOCTYPE HTML>
<html>
  <head>
    <title>Edge detection demo</title>
    <script>
```

```

var image = new Image();
function init() {
    image.onload = demo;
    image.src = "image.jpeg";
}
function demo() {
    var canvas = document.getElementsByTagName('canvas')[0];
    var context = canvas.getContext('2d');

    // draw the image onto the canvas
    context.drawImage(image, 0, 0);

    // get the image data to manipulate
    var input = context.getImageData(0, 0, canvas.width,
    canvas.height);

    // get an empty slate to put the data into
    var output = context.createImageData(canvas.width, canvas.height);

    // alias some variables for convenience
    // notice that we are using input.width and input.height here
    // as they might not be the same as canvas.width and canvas.height
    // (in particular, they might be different on high-res displays)
    var w = input.width, h = input.height;
    var inputData = input.data;
    var outputData = output.data;

    // edge detection
    for (var y = 1; y < h-1; y += 1) {
        for (var x = 1; x < w-1; x += 1) {
            for (var c = 0; c < 3; c += 1) {
                var i = (y*w + x)*4 + c;
                outputData[i] = 127 + -inputData[i - w*4 - 4] -
inputData[i - w*4] - inputData[i - w*4 + 4] +
                    -inputData[i - 4]      +
8*inputData[i]      - inputData[i + 4] +
                    -inputData[i + w*4 - 4] -
inputData[i + w*4] - inputData[i + w*4 + 4];
            }
            outputData[(y*w + x)*4 + 3] = 255; // alpha
        }
    }

    // put the image data back after manipulation
    context.putImageData(output, 0, 0);
}

</script>
</head>
<body onload="init()">
    <canvas></canvas>
</body>
</html>

```

4.8.11.1.12 Drawing model

When a shape or image is painted, user agents must follow these steps, in the order given (or act as if they do):

1. Render the shape or image, creating image A, as described in the previous sections. For shapes, the current fill, stroke, and line styles must be honored, and the stroke must itself also be subjected to the current transformation matrix.
2. Render the shadow from image A, using the current shadow styles, creating image B.
3. Multiply the alpha component of every pixel in B by globalAlpha.
4. Within the clipping region, composite B over the current canvas bitmap using the current composition operator.
5. Multiply the alpha component of every pixel in A by globalAlpha.
6. Within the clipping region, composite A over the current canvas bitmap using the current composition operator.

4.8.11.2 Color spaces and color correction

The canvas APIs must perform color correction at only two points: when rendering images with their own gamma correction and color space information onto the canvas, to convert the image to the color space used by the canvas (e.g. using the `drawImage()` method with an `HTMLImageElement` object), and when rendering the actual canvas bitmap to the output device.

Note: *Thus, in the 2D context, colors used to draw shapes onto the canvas will exactly match colors obtained through the `getImageData()` method.*

The `toDataURL()` method must not include color space information in the resource returned. Where the output format allows it, the color of pixels in resources created by `toDataURL()` must match those returned by the `getImageData()` method.

In user agents that support CSS, the color space used by a canvas element must match the color space used for processing any colors for that element in CSS.

The gamma correction and color space information of images must be handled in such a way that an image rendered directly using an `img` element would use the same colors as one painted on a canvas element that is then itself rendered. Furthermore, the rendering of images that have no color correction information (such as those returned by the `toDataURL()` method) must be rendered with no color correction.

Note: *Thus, in the 2D context, calling the `drawImage()` method to render the output of the `toDataURL()` method to the canvas, given the appropriate dimensions, has no visible effect.*

4.8.11.3 Security with canvas elements

Information leakage can occur if scripts from one origin (page 423) can access information (e.g. read pixels) from images from another origin (one that isn't the same (page 426)).

To mitigate this, canvas elements are defined to have a flag indicating whether they are *origin-clean*. All canvas elements must start with their *origin-clean* set to true. The flag must be set to false if any of the following actions occur:

- The element's 2D context's `drawImage()` method is called with an `HTMLImageElement` whose origin (page 423) is not the same (page 426) as that of the `Document` object that owns the canvas element.
- The element's 2D context's `drawImage()` method is called with an `HTMLCanvasElement` whose *origin-clean* flag is false.
- The element's 2D context's `fillStyle` attribute is set to a `CanvasPattern` object that was created from an `HTMLImageElement` whose origin (page 423) was not the same (page 426) as that of the `Document` object that owns the canvas element when the pattern was created.
- The element's 2D context's `fillStyle` attribute is set to a `CanvasPattern` object that was created from an `HTMLCanvasElement` whose *origin-clean* flag was false when the pattern was created.
- The element's 2D context's `strokeStyle` attribute is set to a `CanvasPattern` object that was created from an `HTMLImageElement` whose origin (page 423) was not the same (page 426) as that of the `Document` object that owns the canvas element when the pattern was created.
- The element's 2D context's `strokeStyle` attribute is set to a `CanvasPattern` object that was created from an `HTMLCanvasElement` whose *origin-clean* flag was false when the pattern was created.

Whenever the `toDataURL()` method of a canvas element whose *origin-clean* flag is set to false is called, the method must raise a security exception (page 430).

Whenever the `getImageData()` method of the 2D context of a canvas element whose *origin-clean* flag is set to false is called with otherwise correct arguments, the method must raise a security exception (page 430).

Note: Even resetting the canvas state by changing its width or height attributes doesn't reset the *origin-clean* flag.

4.8.12 The map element

Categories

Flow content (page 93).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

Flow content (page 93).

Element-specific attributes:

name

DOM interface:

```
interface HTMLMapElement : HTMLElement {  
    attribute DOMString name;  
    readonly attribute HTMLCollection areas;  
    readonly attribute HTMLCollection images;  
};
```

The map element, in conjunction with any area element descendants, defines an image map (page 291).

The **name** attribute gives the map a name so that it can be referenced. The attribute must be present and must have a non-empty value. Whitespace is significant in this attribute's value. If the **id** attribute is also specified, both attributes must have the same value.

The **areas** attribute must return an HTMLCollection rooted at the map element, whose filter matches only area elements.

The **images** attribute must return an HTMLCollection rooted at the Document node, whose filter matches only img and object elements that are associated with this map element according to the image map (page 291) processing model.

The DOM attribute **name** must reflect (page 67) the content attribute of the same name.

4.8.13 The area element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected, but only if there is a map element ancestor.

Content model:

Empty.

Element-specific attributes:

alt
coords
shape
href
target
ping
rel

media
hreflang
type

DOM interface:

```
interface HTMLAreaElement : HTMLElement {  
    attribute DOMString alt;  
    attribute DOMString coords;  
    attribute DOMString shape;  
    attribute DOMString href;  
    attribute DOMString target;  
    attribute DOMString ping;  
    attribute DOMString rel;  
    readonly attribute DOMTokenList relList;  
    attribute DOMString media;  
    attribute DOMString hreflang;  
    attribute DOMString type;  
};
```

The area element represents either a hyperlink with some text and a corresponding area on an image map (page 291), or a dead area on an image map.

If the area element has an href attribute, then the area element represents a hyperlink (page 497). In this case, the alt attribute must be present. It specifies the text of the hyperlink. Its value must be text that, when presented with the texts specified for the other hyperlinks of the image map (page 291), and with the alternative text of the image, but without the image itself, provides the user with the same kind of choice as the hyperlink would when used without its text but with its shape applied to the image. The alt attribute may be left blank if there is another area element in the same image map (page 291) that points to the same resource and has a non-blank alt attribute.

If the area element has no href attribute, then the area represented by the element cannot be selected, and the alt attribute must be omitted.

In both cases, the shape and coords attributes specify the area.

The shape attribute is an enumerated attribute (page 51). The following table lists the keywords defined for this attribute. The states given in the first cell of the rows with keywords give the states to which those keywords map. Some of the keywords are non-conforming, as noted in the last column.

State	Keywords	Notes
Circle state (page 290)	circ	Non-conforming
	circle	
Default state (page 290)	default	
Polygon state (page 290)	poly	
	polygon	Non-conforming
Rectangle state (page 290)	rect	
	rectangle	Non-conforming

The attribute may be omitted. The *missing value default* is the rectangle (page 290) state.

The **coords** attribute must, if specified, contain a valid list of integers (page 38). This attribute gives the coordinates for the shape described by the **shape** attribute. The processing for this attribute is described as part of the image map (page 291) processing model.

In the **circle state**, area elements must have a **coords** attribute present, with three integers, the last of which must be non-negative. The first integer must be the distance in CSS pixels from the left edge of the image to the center of the circle, the second integer must be the distance in CSS pixels from the top edge of the image to the center of the circle, and the third integer must be the radius of the circle, again in CSS pixels.

In the **default state** state, area elements must not have a **coords** attribute. (The area is the whole image.)

In the **polygon state**, area elements must have a **coords** attribute with at least six integers, and the number of integers must be even. Each pair of integers must represent a coordinate given as the distances from the left and the top of the image in CSS pixels respectively, and all the coordinates together must represent the points of the polygon, in order.

In the **rectangle state**, area elements must have a **coords** attribute with exactly four integers, the first of which must be less than the third, and the second of which must be less than the fourth. The four points must represent, respectively, the distance from the left edge of the image to the left side of the rectangle, the distance from the top edge to the top side, the distance from the left edge to the right side, and the distance from the top edge to the bottom side, all in CSS pixels.

When user agents allow users to follow hyperlinks (page 498) created using the **area** element, as described in the next section, the **href**, **target** and **ping** attributes decide how the link is followed. The **rel**, **media**, **hreflang**, and **type** attributes may be used to indicate to the user the likely nature of the target resource before the user follows the link.

The **target**, **ping**, **rel**, **media**, **hreflang**, and **type** attributes must be omitted if the **href** attribute is not present.

The activation behavior (page 96) of area elements is to run the following steps:

1. If the **DOMActivate** event in question is not trusted (i.e. a `click()` method call was the reason for the event being dispatched), and the **area** element's **target** attribute is **...** then raise an **INVALID_ACCESS_ERR** exception.
**
2. Otherwise, the user agent must follow the hyperlink (page 498) defined by the **area** element, if any.

The DOM attributes **alt**, **coords**, **href**, **target**, **ping**, **rel**, **media**, **hreflang**, and **type**, each must reflect (page 67) the respective content attributes of the same name.

The DOM attribute **shape** must reflect (page 67) the shape content attribute, limited to only known values (page 67).

The DOM attribute **relList** must reflect (page 67) the **rel** content attribute.

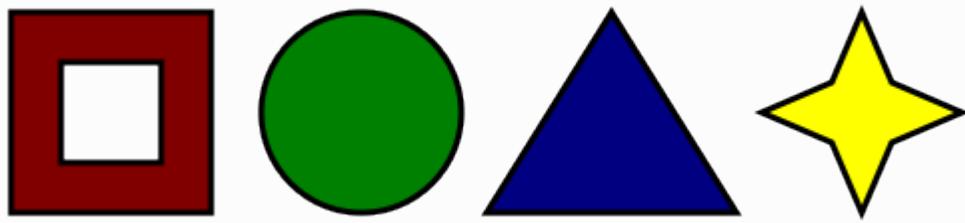
4.8.14 Image maps

4.8.14.1 Authoring

An **image map** allows geometric areas on an image to be associated with hyperlinks (page 497).

An image, in the form of an `img` element or an `object` element representing an image, may be associated with an image map (in the form of a `map` element) by specifying a `usemap` attribute on the `img` or `object` element. The `usemap` attribute, if specified, must be a valid hash-name reference (page 52) to a `map` element.

Consider an image that looks as follows:



If we wanted just the coloured areas to be clickable, we could do it as follows:

```
<p>
  Please select a shape:
  
  <map name="shapes">
    <area shape=rect coords="50,50,100,100"> <!-- the hole in the red
    box -->
    <area shape=rect coords="25,25,125,125" href="red.html" alt="Red
    box.">
    <area shape=circle coords="200,75,50" href="green.html" alt="Green
    circle.">
    <area shape=poly coords="325,25,262,125,388,125" href="blue.html"
    alt="Blue triangle.">
    <area shape=poly
    coords="450,25,435,60,400,75,435,90,450,125,465,90,500,75,465,60"
    href="yellow.html" alt="Yellow star.">
  </map>
</p>
```

4.8.14.2 Processing model

If an `img` element or an `object` element representing an image has a `usemap` attribute specified, user agents must process it as follows:

1. First, rules for parsing a hash-name reference (page 52) to a `map` element must be followed. This will return either an element (the *map*) or null.

2. If that returned null, then abort these steps. The image is not associated with an image map after all.
3. Otherwise, the user agent must collect all the area elements that are descendants of the *map*. Let those be the *areas*.

Having obtained the list of area elements that form the image map (the *areas*), interactive user agents must process the list in one of two ways.

If the user agent intends to show the text that the *img* element represents, then it must use the following steps.

Note: In user agents that do not support images, or that have images disabled, object elements cannot represent images, and thus this section never applies (the fallback content (page 95) is shown instead). The following steps therefore only apply to *img* elements.

1. Remove all the area elements in *areas* that have no *href* attribute.
2. Remove all the area elements in *areas* that have no *alt* attribute, or whose *alt* attribute's value is the empty string, *if* there is another area element in *areas* with the same value in the *href* attribute and with a non-empty *alt* attribute.
3. Each remaining area element in *areas* represents a hyperlink (page 497). Those hyperlinks should all be made available to the user in a manner associated with the text of the *img*.

In this context, user agents may represent area and *img* elements with no specified *alt* attributes, or whose *alt* attributes are the empty string or some other non-visible text, in a user-agent-defined fashion intended to indicate the lack of suitable author-provided text.

If the user agent intends to show the image and allow interaction with the image to select hyperlinks, then the image must be associated with a set of layered shapes, taken from the area elements in *areas*, in reverse tree order (so the last specified area element in the *map* is the bottom-most shape, and the first element in the *map*, in tree order, is the top-most shape).

Each area element in *areas* must be processed as follows to obtain a shape to layer onto the image:

1. Find the state that the element's *shape* attribute represents.
2. Use the rules for parsing a list of integers (page 38) to parse the element's *coords* attribute, if it is present, and let the result be the *coords* list. If the attribute is absent, let the *coords* list be the empty list.
3. If the number of items in the *coords* list is less than the minimum number given for the area element's current state, as per the following table, then the shape is empty; abort these steps.

State	Minimum number of items
Circle state (page 290)	3
Default state (page 290)	0

State	Minimum number of items
Polygon state (page 290)	6
Rectangle state (page 290)	4

4. Check for excess items in the *coords* list as per the entry in the following list corresponding to the shape attribute's state:

↪ **Circle state (page 290)**

Drop any items in the list beyond the third.

↪ **Default state (page 290)**

Drop all items in the list.

↪ **Polygon state (page 290)**

Drop the last item if there's an odd number of items.

↪ **Rectangle state (page 290)**

Drop any items in the list beyond the fourth.

5. If the shape attribute represents the rectangle state (page 290), and the first number in the list is numerically less than the third number in the list, then swap those two numbers around.
6. If the shape attribute represents the rectangle state (page 290), and the second number in the list is numerically less than the fourth number in the list, then swap those two numbers around.
7. If the shape attribute represents the circle state (page 290), and the third number in the list is less than or equal to zero, then the shape is empty; abort these steps.
8. Now, the shape represented by the element is the one described for the entry in the list below corresponding to the state of the shape attribute:

↪ **Circle state (page 290)**

Let x be the first number in *coords*, y be the second number, and r be the third number.

The shape is a circle whose center is x CSS pixels from the left edge of the image and y CSS pixels from the top edge of the image, and whose radius is r pixels.

↪ **Default state (page 290)**

The shape is a rectangle that exactly covers the entire image.

↪ **Polygon state (page 290)**

Let x_i be the $(2i)$ th entry in *coords*, and y_i be the $(2i+1)$ th entry in *coords* (the first entry in *coords* being the one with index 0).

Let *the coordinates* be (x_i, y_i) , interpreted in CSS pixels measured from the top left of the image, for all integer values of i from 0 to $(N/2)-1$, where N is the number of items in *coords*.

The shape is a polygon whose vertices are given by *the coordinates*, and whose interior is established using the even-odd rule. [GRAPHICS]

↪ Rectangle state (page 290)

Let x_1 be the first number in *coords*, y_1 be the second number, x_2 be the third number, and y_2 be the fourth number.

The shape is a rectangle whose top-left corner is given by the coordinate (x_1, y_1) and whose bottom right corner is given by the coordinate (x_2, y_2) , those coordinates being interpreted as CSS pixels from the top left corner of the image.

For historical reasons, the coordinates must be interpreted relative to the *displayed* image, even if it stretched using CSS or the `image` element's `width` and `height` attributes.

Mouse clicks on an image associated with a set of layered shapes per the above algorithm must be dispatched to the top-most shape covering the point that the pointing device indicated (if any), and then, must be dispatched again (with a new Event object) to the `image` element itself. User agents may also allow individual area elements representing hyperlinks (page 497) to be selected and activated (e.g. using a keyboard); events from this are not also propagated to the image.

Note: Because a `map` element (and its area elements) can be associated with multiple `img` and `object` elements, it is possible for an area element to correspond to multiple focusable areas of the document.

Image maps are *live*; if the DOM is mutated, then the user agent must act as if it had rerun the algorithms for image maps.

4.8.15 MathML

The `math` element from the MathML namespace (page 658) falls into the embedded content (page 94) category for the purposes of the content models in this specification.

User agents must handle text other than inter-element whitespace (page 92) found in MathML elements whose content models do not allow raw text by pretending for the purposes of MathML content models, layout, and rendering that that text is actually wrapped in an `mtext` element in the MathML namespace (page 658). (Such text is not, however, conforming.)

User agents must act as if any MathML element whose contents does not match the element's content model was replaced, for the purposes of MathML layout and rendering, by an `merror` element in the MathML namespace (page 658) containing some appropriate error message.

To enable authors to use MathML tools that only accept MathML in its XML form, interactive HTML user agents are encouraged to provide a way to export any MathML fragment as a namespace-well-formed XML fragment.

4.8.16 SVG

The `svg` element from the SVG namespace (page 658) falls into the embedded content (page 94) category for the purposes of the content models in this specification.

To enable authors to use SVG tools that only accept SVG in its XML form, interactive HTML user agents are encouraged to provide a way to export any SVG fragment as a namespace-well-formed XML fragment.

4.8.17 Dimension attributes

The **width** and **height** attributes on `img`, `iframe`, `embed`, `object`, and `video` elements may be specified to give the dimensions of the visual content of the element (the width and height respectively, relative to the nominal direction of the output medium), in CSS pixels. The attributes, if specified, must have values that are valid positive non-zero integers (page 38).

The specified dimensions given may differ from the dimensions specified in the resource itself, since the resource may have a resolution that differs from the CSS pixel resolution. (On screens, CSS pixels have a resolution of 96ppi, but in general the CSS pixel resolution depends on the reading distance.) If both attributes are specified, then the ratio of the specified width to the specified height must be the same as the ratio of the intrinsic width to the intrinsic height in the resource, or alternatively, in the case of the `video` element, the same as the adjusted ratio (page 231). The two attributes must be omitted if the resource in question does not have both an intrinsic width and an intrinsic height.

To parse the attributes, user agents must use the rules for parsing dimension values (page 38). This will return either an integer length, a percentage value, or nothing. The user agent requirements for processing the values obtained from parsing these attributes are described in the rendering section (page 672). If one of these attributes, when parsing, returns no value, it must be treated, for the purposes of those requirements, as if it was not specified.

The **width** and **height** DOM attributes on the `iframe`, `embed`, `object`, and `video` elements must reflect (page 67) the respective content attributes of the same name.

4.9 Tabular data

4.9.1 Introduction

This section is non-normative.

** ...examples, how to write tables accessibly, a brief mention of the table model, etc...

4.9.2 The `table` element

Categories

Flow content (page 93).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

In this order: optionally a `caption` element, followed by either zero or more `colgroup` elements, followed optionally by a `thead` element, followed optionally by a `tfoot`

element, followed by either zero or more tbody elements or one or more tr elements, followed optionally by a tfoot element (but there can only be one tfoot element child in total).

Element-specific attributes:

None.

DOM interface:

```
interface HTMLTableElement : HTMLElement {  
    attribute HTMLTableCaptionElement caption;  
    HTMLElement createCaption();  
    void deleteCaption();  
    attribute HTMLTableSectionElement tHead;  
    HTMLElement createTHead();  
    void deleteTHead();  
    attribute HTMLTableSectionElement tFoot;  
    HTMLElement createTFoot();  
    void deleteTFoot();  
    readonly attribute HTMLCollection tBodies;  
    HTMLElement createTBody();  
    readonly attribute HTMLCollection rows;  
    HTMLElement insertRow(in long index);  
    void deleteRow(in long index);  
};
```

The table element represents data with more than one dimension (a table (page 306)).

** we need some editorial text on how layout tables are bad practice and non-conforming

The children of a table element must be, in order:

1. Zero or one caption elements.
2. Zero or more colgroup elements.
3. Zero or one thead elements.
4. Zero or one tfoot elements, if the last element in the table is not a tfoot element.
5. Either:
 - Zero or more tbody elements, or
 - One or more tr elements. (**Only expressible in the XML serialization.**)
6. Zero or one tfoot element, if there are no other tfoot elements in the table.

The table element takes part in the table model (page 306).

The **caption** DOM attribute must return, on getting, the first caption element child of the table element, if any, or null otherwise. On setting, if the new value is a caption element, the first caption element child of the table element, if any, must be removed, and the new

value must be inserted as the first node of the table element. If the new value is not a caption element, then a HIERARCHY_REQUEST_ERR DOM exception must be raised instead.

The **createCaption()** method must return the first caption element child of the table element, if any; otherwise a new caption element must be created, inserted as the first node of the table element, and then returned.

The **deleteCaption()** method must remove the first caption element child of the table element, if any.

The **tHead** DOM attribute must return, on getting, the first thead element child of the table element, if any, or null otherwise. On setting, if the new value is a thead element, the first thead element child of the table element, if any, must be removed, and the new value must be inserted immediately before the first element in the table element that is neither a caption element nor a colgroup element, if any, or at the end of the table otherwise. If the new value is not a thead element, then a HIERARCHY_REQUEST_ERR DOM exception must be raised instead.

The **createTHead()** method must return the first thead element child of the table element, if any; otherwise a new thead element must be created and inserted immediately before the first element in the table element that is neither a caption element nor a colgroup element, if any, or at the end of the table otherwise, and then that new element must be returned.

The **deleteTHead()** method must remove the first thead element child of the table element, if any.

The **tFoot** DOM attribute must return, on getting, the first tfoot element child of the table element, if any, or null otherwise. On setting, if the new value is a tfoot element, the first tfoot element child of the table element, if any, must be removed, and the new value must be inserted immediately before the first element in the table element that is neither a caption element, a colgroup element, nor a thead element, if any, or at the end of the table if there are no such elements. If the new value is not a tfoot element, then a HIERARCHY_REQUEST_ERR DOM exception must be raised instead.

The **createTFoot()** method must return the first tfoot element child of the table element, if any; otherwise a new tfoot element must be created and inserted immediately before the first element in the table element that is neither a caption element, a colgroup element, nor a thead element, if any, or at the end of the table if there are no such elements, and then that new element must be returned.

The **deleteTFoot()** method must remove the first tfoot element child of the table element, if any.

The **tBodies** attribute must return an HTMLCollection rooted at the table node, whose filter matches only tbody elements that are children of the table element.

The **createTBody()** method must create a new tbody element, insert it immediately after the last tbody element in the table element, if any, or at the end of the table element if the table element has no tbody element children, and then must return the new tbody element.

The **rows** attribute must return an HTMLCollection rooted at the table node, whose filter matches only tr elements that are either children of the table element, or children of thead, tbody, or tfoot elements that are themselves children of the table element. The elements

in the collection must be ordered such that those elements whose parent is a `thead` are included first, in tree order, followed by those elements whose parent is either a `tbody` or `tfoot` element, again in tree order, followed finally by those elements whose parent is a `tfoot` element, still in tree order.

The behavior of the `insertRow(index)` method depends on the state of the table. When it is called, the method must act as required by the first item in the following list of conditions that describes the state of the table and the `index` argument:

↪ **If `index` is less than `-1` or greater than the number of elements in `rows` collection:**

The method must raise an `INDEX_SIZE_ERR` exception.

↪ **If the `rows` collection has zero elements in it, and the table has no `tbody` elements in it:**

The method must create a `tbody` element, then create a `tr` element, then append the `tr` element to the `tbody` element, then append the `tbody` element to the table element, and finally return the `tr` element.

↪ **If the `rows` collection has zero elements in it:**

The method must create a `tr` element, append it to the last `tbody` element in the table, and return the `tr` element.

↪ **If `index` is equal to `-1` or equal to the number of items in `rows` collection:**

The method must create a `tr` element, and append it to the parent of the last `tr` element in the `rows` collection. Then, the newly created `tr` element must be returned.

↪ **Otherwise:**

The method must create a `tr` element, insert it immediately before the `index`th `tr` element in the `rows` collection, in the same parent, and finally must return the newly created `tr` element.

When the `deleteRow(index)` method is called, the user agent must run the following steps:

1. If `index` is equal to `-1`, then `index` must be set to the number of items in the `rows` collection, minus one.
2. Now, if `index` is less than zero, or greater than or equal to the number of elements in the `rows` collection, the method must instead raise an `INDEX_SIZE_ERR` exception, and these steps must be aborted.
3. Otherwise, the method must remove the `index`th element in the `rows` collection from its parent.

4.9.3 The `caption` element

Categories

None.

Contexts in which this element may be used:

As the first element child of a `table` element.

Content model:

Phrasing content (page 94).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The caption element represents the title of the table that is its parent, if it has a parent and that is a table element.

The caption element takes part in the table model (page 306).

4.9.4 The colgroup element

Categories

None.

Contexts in which this element may be used:

As a child of a table element, after any caption elements and before any thead, tbody, tfoot, and tr elements.

Content model:

Zero or more col elements.

Element-specific attributes:

span

DOM interface:

```
interface HTMLTableColElement : HTMLElement {  
    attribute unsigned long span;  
};
```

The colgroup element represents a group (page 306) of one or more columns (page 306) in the table that is its parent, if it has a parent and that is a table element.

If the colgroup element contains no col elements, then the element may have a **span** content attribute specified, whose value must be a valid non-negative integer (page 33) greater than zero.

The colgroup element and its span attribute take part in the table model (page 306).

The **span** DOM attribute must reflect (page 67) the respective content attribute of the same name. The value must be limited to only positive non-zero numbers (page 68).

4.9.5 The col element

Categories

None.

Contexts in which this element may be used:

As a child of a colgroup element that doesn't have a span attribute.

Content model:

Empty.

Element-specific attributes:

span

DOM interface:

HTMLTableColElement, same as for colgroup elements. This interface defines one member, span.

If a col element has a parent and that is a colgroup element that itself has a parent that is a table element, then the col element represents one or more columns (page 306) in the column group (page 306) represented by that colgroup.

The element may have a **span** content attribute specified, whose value must be a valid non-negative integer (page 33) greater than zero.

The col element and its span attribute take part in the table model (page 306).

The **span** DOM attribute must reflect (page 67) the content attribute of the same name. The value must be limited to only positive non-zero numbers (page 68).

4.9.6 The tbody element

Categories

None.

Contexts in which this element may be used:

As a child of a table element, after any caption, colgroup, and thead elements, but only if there are no tr elements that are children of the table element.

Content model:

Zero or more tr elements

Element-specific attributes:

None.

DOM interface:

```
interface HTMLTableSectionElement : HTMLElement {  
    readonly attribute HTMLCollection rows;  
    HTMLElement insertRow(in long index);  
    void deleteRow(in long index);  
};
```

The `HTMLTableSectionElement` interface is also used for `thead` and `tfoot` elements.

The `tbody` element represents a block (page 306) of rows (page 306) that consist of a body of data for the parent table element, if the `tbody` element has a parent and it is a table.

The `tbody` element takes part in the table model (page 306).

The `rows` attribute must return an `HTMLCollection` rooted at the element, whose filter matches only `tr` elements that are children of the element.

The `insertRow(index)` method must, when invoked on an element *table section*, act as follows:

If `index` is less than `-1` or greater than the number of elements in the `rows` collection, the method must raise an `INDEX_SIZE_ERR` exception.

If `index` is equal to `-1` or equal to the number of items in the `rows` collection, the method must create a `tr` element, append it to the element *table section*, and return the newly created `tr` element.

Otherwise, the method must create a `tr` element, insert it as a child of the *table section* element, immediately before the `index`th `tr` element in the `rows` collection, and finally must return the newly created `tr` element.

The `deleteRow(index)` method must remove the `index`th element in the `rows` collection from its parent. If `index` is less than zero or greater than or equal to the number of elements in the `rows` collection, the method must instead raise an `INDEX_SIZE_ERR` exception.

4.9.7 The `thead` element

Categories

None.

Contexts in which this element may be used:

As a child of a `table` element, after any `caption`, and `colgroup` elements and before any `tbody`, `tfoot`, and `tr` elements, but only if there are no other `thead` elements that are children of the `table` element.

Content model:

Zero or more `tr` elements

Element-specific attributes:

None.

DOM interface:

`HTMLTableSectionElement`, as defined for `tbody` elements.

The `thead` element represents the block (page 306) of rows (page 306) that consist of the column labels (headers) for the parent `table` element, if the `thead` element has a parent and it is a `table`.

The thead element takes part in the table model (page 306).

4.9.8 The tfoot element

Categories

None.

Contexts in which this element may be used:

As a child of a table element, after any caption, colgroup, and thead elements and before any tbody and tr elements, but only if there are no other tfoot elements that are children of the table element.

As a child of a table element, after any caption, colgroup, thead, tbody, and tr elements, but only if there are no other tfoot elements that are children of the table element.

Content model:

Zero or more tr elements

Element-specific attributes:

None.

DOM interface:

HTMLTableSectionElement, as defined for tbody elements.

The tfoot element represents the block (page 306) of rows (page 306) that consist of the column summaries (footers) for the parent table element, if the tfoot element has a parent and it is a table.

The tfoot element takes part in the table model (page 306).

4.9.9 The tr element

Categories

None.

Contexts in which this element may be used:

As a child of a thead element.

As a child of a tbody element.

As a child of a tfoot element.

As a child of a table element, after any caption, colgroup, and thead elements, but only if there are no tbody elements that are children of the table element.

Content model:

Zero or more td or th elements

Element-specific attributes:

None.

DOM interface:

```
interface HTMLTableRowElement : HTMLElement {  
    readonly attribute long rowIndex;
```

```
readonly attribute long sectionRowIndex;  
readonly attribute HTMLCollection cells;  
HTMLElement insertCell(in long index);  
void deleteCell(in long index);  
};
```

The `tr` element represents a row (page 306) of cells (page 306) in a table (page 306).

The `tr` element takes part in the table model (page 306).

The `rowIndex` attribute must, if the element has a parent table element, or a parent `tbody`, `thead`, or `tfoot` element and a *grandparent* table element, return the index of the `tr` element in that table element's rows collection. If there is no such table element, then the attribute must return `-1`.

The `sectionRowIndex` attribute must, if the element has a parent table, `tbody`, `thead`, or `tfoot` element, return the index of the `tr` element in the parent element's rows collection (for tables, that's the rows collection; for table sections, that's the rows collection). If there is no such parent element, then the attribute must return `-1`.

The `cells` attribute must return an `HTMLCollection` rooted at the `tr` element, whose filter matches only `td` and `th` elements that are children of the `tr` element.

The `insertCell(index)` method must act as follows:

If `index` is less than `-1` or greater than the number of elements in the `cells` collection, the method must raise an `INDEX_SIZE_ERR` exception.

If `index` is equal to `-1` or equal to the number of items in `cells` collection, the method must create a `td` element, append it to the `tr` element, and return the newly created `td` element.

Otherwise, the method must create a `td` element, insert it as a child of the `tr` element, immediately before the `index`th `td` or `th` element in the `cells` collection, and finally must return the newly created `td` element.

The `deleteCell(index)` method must remove the `index`th element in the `cells` collection from its parent. If `index` is less than zero or greater than or equal to the number of elements in the `cells` collection, the method must instead raise an `INDEX_SIZE_ERR` exception.

4.9.10 The `td` element

Categories

Sectioning root (page 140).

Contexts in which this element may be used:

As a child of a `tr` element.

Content model:

Flow content (page 93).

Element-specific attributes:

colspan
rowspan
headers

DOM interface:

```
interface HTMLTableDataCellElement : HTMLTableCellElement {  
    attribute DOMString headers;  
};
```

The td element represents a data cell (page 306) in a table.

The td element may have a **headers** content attribute specified. The headers attribute, if specified, must contain a string consisting of an unordered set of unique space-separated tokens (page 49), each of which must have the value of an ID of a th element taking part in the same table (page 306) as the td element (as defined by the table model (page 306)).

The exact effect of the attribute is described in detail in the algorithm for assigning header cells to data cells (page 311), which user agents must apply to determine the relationships between data cells and header cells.

The td element and its colspan and rowspan attributes take part in the table model (page 306).

The **headers** DOM attribute must reflect (page 67) the content attribute of the same name.

4.9.11 The th element

Categories

None.

Contexts in which this element may be used:

As a child of a tr element.

Content model:

Phrasing content (page 94).

Element-specific attributes:

colspan
rowspan
scope

DOM interface:

```
interface HTMLTableHeaderCellElement : HTMLTableCellElement {  
    attribute DOMString scope;  
};
```

The th element represents a header cell (page 306) in a table.

The `th` element may have a **scope** content attribute specified. The `scope` attribute is an enumerated attribute (page 51) with five states, four of which have explicit keywords:

The `row` keyword, which maps to the `row` state

The `row` state means the header cell applies to all the remaining cells in the row.

The `col` keyword, which maps to the `column` state

The `column` state means the header cell applies to all the remaining cells in the column.

The `rowgroup` keyword, which maps to the `row group` state

The `row group` state means the header cell applies to all the remaining cells in the row group.

The `colgroup` keyword, which maps to the `column group` state

The `column group` state means the header cell applies to all the remaining cells in the column group.

The `auto` state

The `auto` state makes the header cell apply to a set of cells selected based on context.

The `scope` attribute's *missing value default* is the `auto` state.

The exact effect of these values is described in detail in the algorithm for assigning header cells to data cells (page 311), which user agents must apply to determine the relationships between data cells and header cells.

The `th` element and its `colspan` and `rowspan` attributes take part in the table model (page 306).

The `scope` DOM attribute must reflect (page 67) the content attribute of the same name.

4.9.12 Attributes common to `td` and `th` elements

The `td` and `th` elements may have a **colspan** content attribute specified, whose value must be a valid non-negative integer (page 33) greater than zero.

The `td` and `th` elements may also have a **rowspan** content attribute specified, whose value must be a valid non-negative integer (page 33).

The `td` and `th` elements implement interfaces that inherit from the `HTMLTableCellElement` interface:

```
interface HTMLTableCellElement : HTMLElement {  
    attribute long colSpan;  
    attribute long rowSpan;  
    readonly attribute long cellIndex;  
};
```

The `colSpan` DOM attribute must reflect (page 67) the content attribute of the same name. The value must be limited to only positive non-zero numbers (page 68).

The **rowSpan** DOM attribute must reflect (page 67) the content attribute of the same name. Its default value, which must be used if parsing the attribute as a non-negative integer (page 33) returns an error, is also 1.

The **cellIndex** DOM attribute must, if the element has a parent **tr** element, return the index of the cell's element in the parent element's **cells** collection. If there is no such parent element, then the attribute must return 0.

4.9.13 Processing model

The various table elements and their content attributes together define the **table model**.

A **table** consists of cells aligned on a two-dimensional grid of **slots** with coordinates (x, y) . The grid is finite, and is either empty or has one or more slots. If the grid has one or more slots, then the x coordinates are always in the range $0 \leq x < x_{width}$, and the y coordinates are always in the range $0 \leq y < y_{height}$. If one or both of x_{width} and y_{height} are zero, then the table is empty (has no slots). Tables correspond to **table** elements.

A **cell** is a set of slots anchored at a slot $(cell_x, cell_y)$, and with a particular *width* and *height* such that the cell covers all the slots with coordinates (x, y) where $cell_x \leq x < cell_x + width$ and $cell_y \leq y < cell_y + height$. Cells can either be *data cells* or *header cells*. Data cells correspond to **td** elements, and have zero or more associated header cells. Header cells correspond to **th** elements.

A **row** is a complete set of slots from $x=0$ to $x=x_{width}-1$, for a particular value of y . Rows correspond to **tr** elements.

A **column** is a complete set of slots from $y=0$ to $y=y_{height}-1$, for a particular value of x . Columns can correspond to **col** elements, but in the absence of **col** elements are implied.

A **row group** is a set of rows (page 306) anchored at a slot $(0, group_y)$ with a particular *height* such that the row group covers all the slots with coordinates (x, y) where $0 \leq x < x_{width}$ and $group_y \leq y < group_y + height$. Row groups correspond to **tbody**, **thead**, and **tfoot** elements. Not every row is necessarily in a row group.

A **column group** is a set of columns (page 306) anchored at a slot $(group_x, 0)$ with a particular *width* such that the column group covers all the slots with coordinates (x, y) where $group_x \leq x < group_x + width$ and $0 \leq y < y_{height}$. Column groups correspond to **colgroup** elements. Not every column is necessarily in a column group.

Row groups (page 306) cannot overlap each other. Similarly, column groups (page 306) cannot overlap each other.

A cell (page 306) cannot cover slots that are from two or more row groups (page 306). It is, however, possible for a cell to be in multiple column groups (page 306). All the slots that form part of one cell are part of zero or one row groups (page 306) and zero or more column groups (page 306).

In addition to cells (page 306), columns (page 306), rows (page 306), row groups (page 306), and column groups (page 306), tables (page 306) can have a **caption** element associated with them. This gives the table a heading, or legend.

A **table model error** is an error with the data represented by table elements and their descendants. Documents must not have table model errors.

4.9.13.1 Forming a table

To determine which elements correspond to which slots in a table (page 306) associated with a table element, to determine the dimensions of the table (x_{width} and y_{height}), and to determine if there are any table model errors (page 307), user agents must use the following algorithm:

1. Let x_{width} be zero.
2. Let y_{height} be zero.
3. Let *pending tfoot elements* be a list of tfoot elements, initially empty.
4. Let *the table* be the table (page 306) represented by the table element. The x_{width} and y_{height} variables give *the table's* dimensions. *The table* is initially empty.
5. If the table element has no children elements, then return *the table* (which will be empty), and abort these steps.
6. Associate the first caption element child of the table element with *the table*. If there are no such children, then it has no associated caption element.
7. Let the *current element* be the first element child of the table element.

If a step in this algorithm ever requires the *current element* to be **advanced to the next child of the table** when there is no such next child, then the user agent must jump to the step labeled *end*, near the end of this algorithm.

8. While the *current element* is not one of the following elements, advance (page 307) the *current element* to the next child of the table:
 - colgroup
 - thead
 - tbody
 - tfoot
 - tr
9. If the *current element* is a colgroup, follow these substeps:

1. *Column groups*: Process the *current element* according to the appropriate case below:

↪ **If the current element has any col element children**

Follow these steps:

1. Let x_{start} have the value of x_{width} .
2. Let the *current column* be the first col element child of the colgroup element.
3. *Columns*: If the *current column* col element has a span attribute, then parse its value using the rules for parsing non-negative integers (page 33).

If the result of parsing the value is not an error or zero, then let *span* be that value.

Otherwise, if the *col* element has no *span* attribute, or if trying to parse the attribute's value resulted in an error, then let *span* be 1.

4. Increase x_{width} by *span*.
5. Let the last *span* columns (page 306) in *the table* correspond to the *current column* *col* element.
6. If *current column* is not the last *col* element child of the *colgroup* element, then let the *current column* be the next *col* element child of the *colgroup* element, and return to the step labeled *columns*.
7. Let all the last columns (page 306) in *the table* from $x=x_{start}$ to $x=x_{width}-1$ form a new column group (page 306), anchored at the slot $(x_{start}, 0)$, with width $x_{width}-x_{start}$, corresponding to the *colgroup* element.

↪ **If the *current element* has no *col* element children**

1. If the *colgroup* element has a *span* attribute, then parse its value using the rules for parsing non-negative integers (page 33).

If the result of parsing the value is not an error or zero, then let *span* be that value.

Otherwise, if the *colgroup* element has no *span* attribute, or if trying to parse the attribute's value resulted in an error, then let *span* be 1.

2. Increase x_{width} by *span*.
3. Let the last *span* columns (page 306) in *the table* form a new column group (page 306), anchored at the slot $(x_{width}-span, 0)$, with width *span*, corresponding to the *colgroup* element.
2. Advance (page 307) the *current element* to the next child of the table.
3. While the *current element* is not one of the following elements, advance (page 307) the *current element* to the next child of the table:
 - *colgroup*
 - *thead*
 - *tbody*
 - *tfoot*
 - *tr*
4. If the *current element* is a *colgroup* element, jump to the step labeled *column groups* above.
10. Let $y_{current}$ be zero.

11. Let the *list of downward-growing cells* be an empty list.
12. **Rows:** While the *current element* is not one of the following elements, advance (page 307) the *current element* to the next child of the table:
 - `thead`
 - `tbody`
 - `tfoot`
 - `tr`
13. If the *current element* is a `tr`, then run the algorithm for processing rows (page 310), advance (page 307) the *current element* to the next child of the table, and return to the step labeled *rows*.
14. Run the algorithm for ending a row group (page 309).
15. If the *current element* is a `tfoot`, then add that element to the list of *pending tfoot elements*, advance (page 307) the *current element* to the next child of the table, and return to the step labeled *rows*.
16. The *current element* is either a `thead` or a `tbody`.
Run the algorithm for processing row groups (page 309).
17. Advance (page 307) the *current element* to the next child of the table.
18. Return to the step labeled *rows*.
19. *End:* For each `tfoot` element in the list of *pending tfoot elements*, in tree order, run the algorithm for processing row groups (page 309).
20. If there exists a row (page 306) or column (page 306) in the table (page 306) *the table* containing only slots that do not have a cell (page 306) anchored to them, then this is a table model error (page 307).
21. Return *the table*.

The **algorithm for processing row groups**, which is invoked by the set of steps above for processing `thead`, `tbody`, and `tfoot` elements, is:

1. Let y_{start} have the value of y_{height} .
2. For each `tr` element that is a child of the element being processed, in tree order, run the algorithm for processing rows (page 310).
3. If $y_{height} > y_{start}$, then let all the last rows (page 306) in *the table* from $y=y_{start}$ to $y=y_{height}-1$ form a new row group (page 306), anchored at the slot with coordinate $(0, y_{start})$, with height $y_{height}-y_{start}$, corresponding to the element being processed.
4. Run the algorithm for ending a row group (page 309).

The **algorithm for ending a row group**, which is invoked by the set of steps above when starting and ending a block of rows, is:

1. While $y_{current}$ is less than y_{height} , follow these steps:
 1. Run the algorithm for growing downward-growing cells (page 311).

2. Increase $y_{current}$ by 1.
2. Empty the *list of downward-growing cells*.

The **algorithm for processing rows**, which is invoked by the set of steps above for processing `tr` elements, is:

1. If y_{height} is equal to $y_{current}$, then increase y_{height} by 1. ($y_{current}$ is never greater than y_{height} .)
2. Let $x_{current}$ be 0.
3. Run the algorithm for growing downward-growing cells (page 311).
4. If the `tr` element being processed has no `td` or `th` element children, then increase $y_{current}$ by 1, abort this set of steps, and return to the algorithm above.
5. Let *current cell* be the first `td` or `th` element in the `tr` element being processed.
6. *Cells*: While $x_{current}$ is less than x_{width} and the slot with coordinate $(x_{current}, y_{current})$ already has a cell assigned to it, increase $x_{current}$ by 1.
7. If $x_{current}$ is equal to x_{width} , increase x_{width} by 1. ($x_{current}$ is never greater than x_{width} .)
8. If the *current cell* has a `colspan` attribute, then parse that attribute's value (page 33), and let `colspan` be the result.
If parsing that value failed, or returned zero, or if the attribute is absent, then let `colspan` be 1, instead.
9. If the *current cell* has a `rowspan` attribute, then parse that attribute's value (page 33), and let `rowspan` be the result.
If parsing that value failed or if the attribute is absent, then let `rowspan` be 1, instead.
10. If `rowspan` is zero, then let *cell grows downward* be true, and set `rowspan` to 1. Otherwise, let *cell grows downward* be false.
11. If $x_{width} < x_{current} + colspan$, then let x_{width} be $x_{current} + colspan$.
12. If $y_{height} < y_{current} + rowspan$, then let y_{height} be $y_{current} + rowspan$.
13. Let the slots with coordinates (x, y) such that $x_{current} \leq x < x_{current} + colspan$ and $y_{current} \leq y < y_{current} + rowspan$ be covered by a new cell (page 306) *c*, anchored at $(x_{current}, y_{current})$, which has width `colspan` and height `rowspan`, corresponding to the *current cell* element.
If the *current cell* element is a `th` element, let this new cell *c* be a header cell; otherwise, let it be a data cell. To establish what header cells apply to a data cell, use the algorithm for assigning header cells to data cells (page 311) described in the next section.
If any of the slots involved already had a cell (page 306) covering them, then this is a table model error (page 307). Those slots now have two cells overlapping.

14. If *cell grows downward* is true, then add the tuple $\{c, x_{current}, colspan\}$ to the *list of downward-growing cells*.
15. Increase $x_{current}$ by *colspan*.
16. If *current cell* is the last *td* or *th* element in the *tr* element being processed, then increase $y_{current}$ by 1, abort this set of steps, and return to the algorithm above.
17. Let *current cell* be the next *td* or *th* element in the *tr* element being processed.
18. Return to the step labelled *cells*.

When the algorithms above require the user agent to run the **algorithm for growing downward-growing cells**, the user agent must, for each $\{cell, cell_x, width\}$ tuple in the *list of downward-growing cells*, if any, extend the cell (page 306) *cell* so that it also covers the slots with coordinates $(x, y_{current})$, where $cell_x \leq x < cell_x + width$.

4.9.13.2 Forming relationships between data cells and header cells

Each data cell can be assigned zero or more header cells. The **algorithm for assigning header cells to data cells** is as follows.

1. For each header cell in the table, in tree order (page 24), run these substeps:
 1. Let $(header_x, header_y)$ be the coordinate of the slot to which the header cell is anchored.
 2. Let $header_{width}$ be the width of the header cell.
 3. Let $header_{height}$ be the height of the header cell.
 4. Let *data cells* be a list of data cells, initially empty.
 5. Examine the scope attribute of the *th* element corresponding to the header cell, and, based on its state, apply the appropriate substep:

↪ **If it is in the row (page 305) state**

Add all the data cells that cover slots with coordinates $(slot_x, slot_y)$, where $header_x + header_{width} \leq slot_x < x_{width}$ and $header_y \leq slot_y < header_y + header_{height}$, to the *data cells* list.

↪ **If it is in the column (page 305) state**

Add all the data cells that cover slots with coordinates $(slot_x, slot_y)$, where $header_x \leq slot_x < header_x + header_{width}$ and $header_y + header_{height} \leq slot_y < y_{height}$, to the *data cells* list.

↪ **If it is in the row group (page 305) state**

If the header cell is not in a row group (page 306), then do nothing.

Otherwise, let $(0, group_y)$ be the slot at which the row group is anchored, let *height* be the number of rows in the row group, and add all the data cells that cover slots with coordinates $(slot_x, slot_y)$, where $header_x \leq slot_x < x_{width}$ and $header_y \leq slot_y < group_y + height$, to the *data cells* list.

↪ If it is in the **column group** (page 305) state

If the header cell is not anchored in a column group (page 306), then do nothing.

Otherwise, let $(group_x, 0)$ be the slot at which that column group is anchored, let $width$ be the number of columns in the column group, and add all the data cells that cover slots with coordinates $(slot_x, slot_y)$, where $header_x \leq slot_x < group_x + width$ and $header_y \leq slot_y < y_{height}$, to the *data cells* list.

↪ Otherwise, it is in the **auto** state

Run these steps:

1. If the header cell is equivalent to a wide cell (page 313), let $header_{width}$ equal $x_{width} - header_x$.
2. Let x equal $header_x + header_{width}$.
3. Let y equal $header_y + header_{height}$.
4. *Horizontal*: If x is equal to x_{width} , then jump down to the step below labeled *vertical*.
5. If there is a header cell anchored at $(x, header_y)$ with height $header_{height}$, then jump down to the step below labeled *vertical*.
6. Add all the data cells that cover slots with coordinates $(slot_x, slot_y)$, where $slot_x = x$ and $header_y \leq slot_y < header_y + header_{height}$, to the *data cells* list.
7. Increase x by 1.
8. Jump up to the step above labeled *horizontal*.
9. *Vertical*: If y is equal to y_{height} , then jump to the step below labeled *end*.
10. If there is a header cell *cell* anchored at $(header_x, y)$, then follow these substeps:
 1. If the header cell *cell* is equivalent to a wide cell (page 313), then let $width$ be $x_{width} - header_x$. Otherwise, let $width$ be the width of the header cell *cell*.
 2. If $width$ is equal to $header_{width}$, then jump to the step below labeled *end*.
11. Add all the data cells that cover slots with coordinates $(slot_x, slot_y)$, where $header_x \leq slot_x < header_x + header_{width}$ and $slot_y = y$, to the *data cells* list.
12. Increase y by 1.

13. Jump up to the step above labeled *vertical*.
14. *End*: Coalesce all the duplicate entries in the *data cells* list, so that each data cell is only present once, in tree order.
6. Assign the header cell to all the data cells in the *data cells* list that correspond to *td* elements that do not have a *headers* attribute specified.
2. For each data cell in the table, in tree order (page 24), run these substeps:
 1. If the data cell corresponds to a *td* element that does not have a *headers* attribute specified, then skip these substeps and move on to the next data cell (if any).
 2. Otherwise, take the value of the *headers* attribute and split it on spaces (page 49), letting *id list* be the list of tokens obtained.
 3. For each token in the *id list*, run the following steps:
 1. Let *id* be the token.
 2. If there is a header cell in the table (page 306) whose corresponding *th* element has an ID that is equal to the value of *id*, then assign the first such header cell in tree order to the data cell.

A header cell anchored at $(header_x, header_y)$ with width $header_{width}$ and height $header_{height}$ is said to be **equivalent to a wide cell** if all the slots with coordinates $(slot_x, slot_y)$, where $header_x + header_{width} \leq slot_x < x_{width}$ and $header_y \leq slot_y < header_y + header_{height}$, are all either empty or covered by empty data cells (page 313).

A data cell is said to be an **empty data cell** if it contains no elements and its text content, if any, consists only of White_Space (page 32) characters.

User agents may remove empty data cells (page 313) when analyzing data in a table (page 306).

4.10 Forms

Forms allow unscripted client-server interaction: given a form, a user can provide data, submit it to the server, and have the server act on it accordingly (e.g. returning the results of a search or calculation). The elements used in forms can also be used for user interaction with no associated submission mechanism, in conjunction with scripts.

Mostly for historical reasons, elements in this section fall into several overlapping (but subtly different) categories in addition to the usual ones like flow content (page 93), phrasing content (page 94), and interactive content (page 95).

A number of the elements are **form-associated elements**, which means they can have a form owner (page 361) and, to expose this, have a *form content* attribute with a matching *form DOM* attribute.

The form-associated elements (page 313) fall into several subcategories:

Submittable elements

Denotes elements that can be used for constructing the form data set (page 368) when a `form` element is submit (page 368).

Resettable elements

Denotes elements that can be affected when a `form` element is reset (page 376).

Listed

Denotes elements that are listed in the `form.elements` and `fieldset.elements` APIs.

Labelable

Denotes elements that can be associated with `label` elements.

In addition, some submittable (page 314) can be, depending on their attributes, **buttons**. The prose below defines when an element is a button.

4.10.1 The `form` element

Categories

Flow content (page 93).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

Flow content (page 93), but with no `form` element descendants.

Element-specific attributes:

`accept-charset`
`action`
`enctype`
`method`
`name`
`target`

DOM interface:

```
interface HTMLFormElement : HTMLElement {  
    attribute DOMString accept-charset;  
    attribute DOMString action;  
    attribute DOMString enctype;  
    attribute DOMString method;  
    attribute DOMString name;  
    attribute DOMString target;  
  
    readonly attribute HTMLFormControlsCollection elements;  
    readonly attribute long length;  
    [IndexGetter] HTMLElement XXX7(in unsigned long index);  
    [NameGetter] Object XXX8(in DOMString name);  
  
    void submit();  
    void reset();
```

```
boolean checkValidity();

void dispatchFormInput();
void dispatchFormChange();
};
```

The `form` element represents a collection of form-associated elements (page 313), some of which can represent editable values that can be submitted to a server for processing.

The **accept-charset** attribute gives the character encodings that are to be used for the submission. If specified, the value must be an ordered set of space-separated tokens, and each token must be the preferred name of an ASCII-compatible character encoding (page 25). [IANACHARSET]

The **name** attribute represents the form's name within the `forms` collection. The value must not be the empty string, and the value must be unique amongst the `form` elements in the `forms` collection that it is in, if any.

The `action`, `enctype`, `method`, and `target` attributes are attributes for form submission (page 363).

The **accept-charset** and **name** DOM attributes must reflect (page 67) the respective content attributes of the same name.

The **elements** DOM attribute must return an `HTMLFormControlsCollection` rooted at the Document node, whose filter matches listed (page 314) elements whose form owner (page 361) is the `form` element, with the exception of `input` elements whose `type` attribute is in the Image Button (page 338) state, which must, for historical reasons, be excluded from this particular collection.

The **length** DOM attribute must return the number of nodes represented (page 69) by the `elements` collection.

The **XXX7()** method must return the value that would be returned by the `item()` method of the `elements` collection if it was invoked with the same arguments.

The **XXX8()** method must return the value that would be returned by the `namedItem()` method of the `elements` collection if it was invoked with the same arguments.

The **submit()** method, when invoked, must submit (page 368) the `form` element from the `form` element itself.

The **reset()** method, when invoked, must reset (page 376) the `form` element.

If the **checkValidity()** method is invoked, the user agent must statically validate the constraints (page 365) of the `form` element, and return true if the constraint validation return a *positive* result, and false if it returned a *negative* result.

If the **dispatchFormInput()** method is invoked, the user agent must broadcast `forminput` events (page 376) from the `form` element.

If the `dispatchFormChange()` method is invoked, the user agent must broadcast `formchange` events (page 376) from the `form` element.

4.10.2 The `fieldset` element

Categories

Flow content (page 93).

Listed (page 314) form-associated element (page 313).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

One legend element followed by flow content (page 93).

Element-specific attributes:

`disabled`

`form`

`name`

DOM interface:

```
interface HTMLFieldSetElement : HTMLElement {  
    attribute boolean disabled;  
    readonly attribute HTMLFormElement form;  
    attribute DOMString name;  
  
    readonly attribute DOMString type;  
  
    readonly attribute HTMLFormControlsCollection elements;  
  
    readonly attribute boolean willValidate;  
    readonly attribute ValidityState validity;  
    readonly attribute DOMString validationMessage;  
    boolean checkValidity();  
    void setCustomValidity(in DOMString error);  
};
```

The `fieldset` element represents a set of form controls grouped under a common name.

The name of the group is given by the first legend element that is a child of the `fieldset` element. The remainder of the descendants form the group.

The `disabled` attribute, when specified, causes all the form control descendants of the `fieldset` element to be disabled (page 362).

The `form` attribute is used to explicitly associate the `fieldset` element with its form owner (page 361). The `name` attribute represents the element's name.

The `disabled` DOM attribute must reflect (page 67) the content attribute of the same name.

The `type` DOM attribute must return the string "fieldset".

The **elements** DOM attribute must return an `HTMLFormControlsCollection` rooted at the `fieldset` element, whose filter matches listed (page 314) elements.

The `willValidate`, `validity`, and `validationMessage` attributes, and the `checkValidity()` and `setCustomValidity()` methods, are part of the constraint validation API (page 366).

Constraint validation: `fieldset` elements are always barred from constraint validation (page 364).

4.10.3 The `label` element

Categories

- Phrasing content (page 94).
- Interactive content (page 95).
- Form-associated element (page 313).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

If the element has a `for` attribute: Phrasing content (page 94), but with no descendant labelable form-associated elements (page 314) or `label` elements.
Otherwise: Phrasing content (page 94), but with at most one descendant labelable form-associated element (page 314), and with no descendant `label` elements.

Element-specific attributes:

- `form`
- `for`

DOM interface:

```
interface HTMLabelElement : HTMLElement {  
    readonly attribute HTMLFormElement form;  
        attribute DOMString htmlFor;  
    readonly attribute HTMLElement control;  
};
```

The `label` represents a caption in a user interface. The caption can be associated with a specific form control, known as the `label` element's **labeled control**.

Unless otherwise specified by the following rules, a `label` element has no labeled control (page 317).

The `for` attribute may be specified to indicate a form control with which the caption is to be associated. If the attribute is specified, the attribute's value must be the ID of a labelable form-associated element (page 314) in the same Document as the `label` element. If the attribute is specified and there is an element in the Document whose ID is equal to the value of the `for` attribute, and the first such element is a labelable form-associated element (page 314), then that element is the `label` element's labeled control (page 317).

If the `for` attribute is not specified, but the `label` element has a labelable form-associated element descendant, then the first such descendant in tree order (page 24) is the `label` element's labeled control (page 317).

The `label` element's exact default presentation and behavior, in particular what its activation behavior (page 96) might be, if anything, should match the platform's label behavior.

For example, on platforms where clicking a checkbox label checks the checkbox, clicking the `label` in the following snippet could trigger the user agent to run synthetic click activation steps (page 95) on the `input` element, as if the element itself had been triggered by the user:

```
<label><input type=checkbox name=lost> Lost</label>
```

On other platforms, the behavior might be just to focus the control, or do nothing.

The `form` attribute is used to explicitly associate the `label` element with its form owner (page 361).

The `htmlFor` DOM attribute must reflect (page 67) the `for` content attribute.

The `control` DOM attribute must return the `label` element's labeled control (page 317), if any, or null if there isn't one.

Labelable form-associated elements have a `NodeList` object associated with them that represents the list of `label` elements, in tree order (page 24), whose labeled control (page 317) is the element in question. The `labels` DOM attribute of labelable form-associated elements, on getting, must return that `NodeList` object.

4.10.4 The `input` element

Categories

Phrasing content (page 94).

Interactive content (page 95).

Listed (page 314), labelable (page 314), submittable (page 314), and resettable (page 314) form-associated element (page 313).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Empty.

Element-specific attributes:

- accept
- action
- alt
- autocomplete
- autofocus
- checked
- disabled
- enctype
- form
- list
- max
- maxlength

method
min
name
pattern
readonly
required
size
src
step
target
type
value

DOM interface:

```
interface HTMLInputElement : HTMLElement {
    attribute DOMString accept;
    attribute DOMString action;
    attribute DOMString alt;
    attribute boolean autocomplete;
    attribute boolean autofocus;
    attribute boolean defaultChecked;
    attribute boolean checked;
    attribute boolean disabled;
    attribute DOMString enctype;
    readonly attribute HTMLFormElement form;
    readonly attribute HTMLElement list;
    attribute DOMString max;
    attribute long maxLength;
    attribute DOMString method;
    attribute DOMString min;
    attribute DOMString name;
    attribute DOMString pattern;
    attribute boolean readOnly;
    attribute boolean required;
    attribute unsigned long size;
    attribute DOMString src;
    attribute DOMString step;
    attribute DOMString target;
    attribute DOMString type;
    attribute DOMString defaultValue;
    attribute DOMString value;
    attribute DOMTimeStamp valueAsDate;
    attribute float valueAsNumber;
    readonly attribute HTMLOptionElement selectedOption;

    void stepUp(in long n);
    void stepDown(in long n);

    readonly attribute boolean willValidate;
```

```

readonly attribute ValidityState validity;
readonly attribute DOMString validationMessage;
boolean checkValidity();
void setCustomValidity(in DOMString error);

readonly attribute NodeList labels;
};

```

The `input` element represents a typed data field, usually with a form control to allow the user to edit the data.

The **`type`** attribute controls the data type (and associated control) of the element. It is an enumerated attribute (page 51). The following table lists the keywords and states for the attribute — the keywords in the left column map to the states in the cell in the second column on the same row as the keyword.

Keyword	State	Data type	Control type
hidden	Hidden (page 324)	An arbitrary string	n/a
text	Text (page 324)	Text with no line breaks	Text field
email	E-mail (page 325)	An e-mail address	A text field
url	URL (page 325)	An IRI	A text field
password	Password (page 326)	Text with no line breaks (sensitive information)	Text field that obscures data entry
datetime	Date and Time (page 326)	A date and time (year, month, day, hour, minute, second, fraction of a second) with the time zone set to UTC	A date and time control
date	Date (page 328)	A date (year, month, day) with no time zone	A date control
month	Month (page 329)	A date consisting of a year and a month with no time zone	A month control
week	Week (page 330)	A date consisting of a year and a week number with no time zone	A week control
time	Time (page 331)	A time (hour, minute, seconds, fractional seconds) with no time zone	A time control
datetime-local	Local Date and Time (page 332)	A date and time (year, month, day, hour, minute, second, fraction of a second) with no time zone	A date and time control
number	Number (page 333)	A numerical value	A text field or spinner control
range	Range (page 334)	A numerical value, with the extra semantic that the exact value is not important	A slider control or similar
checkbox	Checkbox (page 335)	A set of zero or more values from a predefined list	A checkbox
radio	Radio Button (page 336)	An enumerated value	A radio button

Keyword	State	Data type	Control type
file	File Upload (page 337)	Zero or more files each with a MIME type and optionally a file name	A label and a button
submit	Submit Button (page 338)	An enumerated value, with the extra semantic that it must be the last value selected and initiates form submission	A button
image	Image Button (page 338)	A coordinate, relative to a particular image's size, with the extra semantic that it must be the last value selected and initiates form submission	Either a clickable image, or a button
reset	Reset Button (page 340)	n/a	A button
button	Button (page 341)	n/a	A button

The *missing value default* is the Text (page 324) state.

Which of the accept, action, alt, autocomplete, checked, enctype, and list, max, maxlength, method, min, pattern, readonly, required, size, src, step, and target attributes apply to an input element depends on the state of its type attribute. Similarly, the checked, valueAsDate, valueAsNumber, list, and selectedOption DOM attributes, and the stepUp() and stepDown() methods, are specific to certain states. The following table is non-normative and summarises which content attributes, DOM attributes, and methods apply to each state:

	Hidden (page 324)	Text (page 324), E-mail (page 325), URL (page 52)	Password (page 326)	Date and Time (page 326), Date (page 328), Month (page 329), Week (page 330), Time (page 172)	Local Date and Time (page 332), Number (page 333)	Range (page 334)	Checkbox (page 335), Radio Button (page 336)	File Upload (page 337)	Submit Button (page 338)	Image Button (page 338)	Reset Button (page 340), Button (page 348)
accept	Yes	.	.
action	Yes	Yes	.
alt	Yes
autocomplete	.	Yes	Yes	Yes	Yes	Yes
checked	Yes
enctype	Yes	Yes	.
list	.	Yes	.	Yes	Yes	Yes
max	.	.	.	Yes	Yes	Yes
maxlength	.	Yes	Yes
method	Yes	Yes	.
min	.	.	.	Yes	Yes	Yes
pattern	.	Yes	Yes

	Hidden (page 324)	Text (page 324), E-mail (page 325), URL (page 52)	Password (page 326)	Date and Time (page 326), Date (page 328), Month (page 329), Week (page 330), Time (page 172)	Local Date and Time (page 332), Number (page 333)	Range (page 334)	Checkbox (page 335), Radio Button (page 336)	File Upload (page 337)	Submit Button (page 338)	Image Button (page 338)	Res
readonly	.	Yes	Yes	Yes	Yes
required	.	Yes	Yes	Yes	Yes	.	Yes	Yes	.	.	.
size	.	Yes	Yes
src	Yes
step	.	.	.	Yes	Yes	Yes
target	Yes	Yes	.
checked	Yes
value	value (page 345)	value (page 345)	value (page 345)	value (page 345)	value (page 345)	value (page 345)	value (page 345)	default/on (page 346)	.	default (page 346)	default (page 346)
valueAsDate	.	.	.	Yes
valueAsNumber	.	.	.	Yes	Yes	Yes
list	.	Yes	.	Yes	Yes	Yes
selectedOption	.	Yes	.	Yes	Yes	Yes
stepUp()	.	.	.	Yes	Yes	Yes
stepDown()	.	.	.	Yes	Yes	Yes
input event	.	Yes	Yes	Yes	Yes	Yes	Yes
change event	.	Yes	Yes	Yes	Yes	Yes	Yes	Yes	.	.	.

When an input element's type attribute changes state, and when the element is first created, the element's rendering and behaviour must change to the new state's accordingly and the **value sanitization algorithm**, if one is defined for the type attribute's new state, must be invoked.

Each input element has a value (page 362), which is exposed by the value DOM attribute. Some states define an **algorithm to convert a string to a number**, an **algorithm to convert a number to a string**, an **algorithm to convert a string to a Date object**, and an **algorithm to convert a Date object to a string**, which are used by max, min, step, valueAsDate, valueAsNumber, stepUp(), and stepDown().

Each input element has a boolean **dirty value flag**. When it is true, the element is said to have a **dirty value**. The dirty value flag (page 322) must be initially set to false when the element is created, and must be set to true whenever the user interacts with the control in a way that changes the value (page 362).

The **value** content attribute gives the default value (page 362) of the input element. When the value content attribute is added, set, or removed, if the control does not have a *dirty value* (page 322), the user agent must set the value (page 362) of the element to the value of the value content attribute, if there is one, or the empty string otherwise, and then run the current value sanitization algorithm (page 322), if one is defined.

Each input element has a checkedness (page 362), which is exposed by the checked DOM attribute.

Each input element has a boolean **dirty checkedness flag**. When it is true, the element is said to have a **dirty checkedness**. The dirty checkedness flag (page 323) must be initially set to false when the element is created, and must be set to true whenever the user interacts with the control in a way that changes the checkedness (page 362).

The **checked** content attribute gives the default checkedness (page 362) of the input element. When the checked content attribute is added, if the control does not have *dirty checkedness* (page 322), the user agent must set the checkedness (page 362) of the element to true; when the checked content attribute is removed, if the control does not have *dirty checkedness* (page 322), the user agent must set the checkedness (page 362) of the element to false.

The reset algorithm (page 376) for input elements is to set the dirty value flag (page 322) and dirty checkedness flag (page 323) back to false, set the value (page 362) of the element to the value of the value content attribute, if there is one, or the empty string otherwise, set the checkedness (page 362) of the element to true if the element has a checked content attribute and false if it does not, and then invoke the value sanitization algorithm (page 322), if the type attribute's current state defines one.

Each input element has a boolean **mutability flag**. When it is true, the element is said to be **mutable**, and when it is false the element is **immutable**. Unless otherwise specified, an input element is always *mutable* (page 323). Unless otherwise specified, the user agent should not allow the user to modify the element's value (page 362) or checkedness (page 362).

When an input element is disabled (page 362), it is *immutable* (page 323).

When an input element does not have a Document node as one of its ancestors (i.e. when it is not in the document), it is *immutable* (page 323).

Note: The **readonly** attribute can also in some cases make an input element immutable (page 323).

The **form** attribute is used to explicitly associate the input element with its form owner (page 361). The **name** attribute represents the element's name. The **disabled** attribute is used to make the control non-interactive and to prevent its value from being submitted. The **autofocus** attribute controls focus.

The **accept**, **alt**, **autocomplete**, **max**, **min**, **pattern**, **required**, **size**, **src**, **step**, and **type** DOM attributes must reflect (page 67) the respective content attributes of the same name. The **maxLength** DOM attribute must reflect (page 67) the **maxlength** content attribute. The **readOnly** DOM attribute must reflect (page 67) the **readonly** content attribute. The **defaultChecked** DOM attribute must reflect (page 67) the **checked** content attribute. The **defaultValue** DOM attribute must reflect (page 67) the **value** content attribute.

The `willValidate`, `validity`, and `validationMessage` attributes, and the `checkValidity()` and `setCustomValidity()` methods, are part of the constraint validation API (page 366). The `labels` attribute provides a list of the element's labels.

4.10.4.1 States of the type attribute

4.10.4.1.1 Hidden state

When an `input` element's `type` attribute is in the Hidden (page 324) state, the rules in this section apply.

The `input` element represents a value that is not intended to be examined or manipulated by the user.

Constraint validation: If an `input` element's `type` attribute is in the Hidden (page 324) state, it is barred from constraint validation (page 364).

If the `name` attribute is present and has a value that is a case-sensitive (page 31) match for the string `"_charset_"`, then the element's `value` attribute must be omitted.

Bookkeeping details

- The `value` DOM attribute applies to this element and is in mode value (page 345).
- The following content attributes must not be specified and do not apply to the element: `accept`, `action`, `alt`, `autocomplete`, `checked`, `enctype`, `list`, `max`, `maxlength`, `method`, `min`, `pattern`, `readonly`, `required`, `size`, `src`, `step`, and `target`.
- The following DOM attributes and methods do not apply to the element: `checked`, `list`, `selectedOption`, `valueAsDate`, and `valueAsNumber` DOM attributes `stepUp()`, and `stepDown()` methods.
- The `input` and `change` events do not fire.

4.10.4.1.2 Text state

When an `input` element's `type` attribute is in the Text (page 324) state, the rules in this section apply.

The `input` element represents a one line plain text edit control for the element's `value` (page 362).

If the element is *mutable* (page 323), its `value` (page 362) should be editable by the user. User agents must not allow users to insert U+000A LINE FEED (LF) or U+000D CARRIAGE RETURN (CR) characters into the element's `value` (page 362).

The value sanitization algorithm (page 322) is as follows: Strip line breaks (page 32) from the `value` (page 362).

Bookkeeping details

- The following common `input` element content attributes and DOM attributes apply to the element: `autocomplete`, `list`, `maxlength`, `pattern`, `readonly`, `required`, and `size` content attributes; `list`, `selectedOption`, and `value` DOM attributes.
- The `value` DOM attribute is in mode value (page 345).
- The `input` and `change` events apply.
- The following content attributes must not be specified and do not apply to the element: `accept`, `action`, `alt`, `checked`, `enctype`, `max`, `method`, `min`, `src`, `step`, and `target`.

- The following DOM attributes and methods do not apply to the element: checked, valueAsDate, and valueAsNumber DOM attributes stepUp(), and stepDown() methods.

4.10.4.1.3 E-mail state

When an input element's type attribute is in the E-mail (page 325) state, the rules in this section apply.

The input element represents a control for editing a single e-mail address given in the element's value (page 362).

If the element is *mutable* (page 323), the user agent should allow the user to change the e-mail address represented by its value (page 362). User agents may allow the user to set the value (page 362) to a string that is not an e-mail address. User agents should allow the user to set the value (page 362) to the empty string. User agents must not allow users to insert U+000A LINE FEED (LF) or U+000D CARRIAGE RETURN (CR) characters into the value (page 362).

The value sanitization algorithm (page 322) is as follows: Strip line breaks (page 32) from the value (page 362).

Constraint validation: While the value (page 362) of the element does not match the addr-spec token defined in RFC 2822 section 3.4.1, excluding the CFWS subtoken everywhere, and excluding the FWS subtoken everywhere except in the quoted-string subtoken, the element is suffering from a type mismatch (page 365). [RFC2822]

Bookkeeping details

- The following common input element content attributes and DOM attributes apply to the element: autocomplete, list, maxlength, pattern, readonly, required, and size content attributes; list, selectedOption, and value DOM attributes.
- The value DOM attribute is in mode value (page 345).
- The input and change events apply.
- The following content attributes must not be specified and do not apply to the element: accept, action, alt, checked, enctype, max, method, min, src, step, and target.
- The following DOM attributes and methods do not apply to the element: checked, valueAsDate, and valueAsNumber DOM attributes; stepUp(), and stepDown() methods.

4.10.4.1.4 URL state

When an input element's type attribute is in the URL (page 325) state, the rules in this section apply.

The input element represents a control for editing a single URL (page 52) given in the element's value (page 362).

If the is *mutable* (page 323), the user agent should allow the user to change the URL represented by its value (page 362). User agents may allow the user to set the value (page 362) to a string that is not a valid URL (page 52). User agents should allow the user to set the value (page 362) to the empty string. User agents must not allow users to insert U+000A LINE FEED (LF) or U+000D CARRIAGE RETURN (CR) characters into the value (page 362).

The value sanitization algorithm (page 322) is as follows: Strip line breaks (page 32) from the value (page 362).

Constraint validation: While the value (page 362) of the element is not a valid URL (page 52), the element is suffering from a type mismatch (page 365).

Bookkeeping details

- The following common input element content attributes and DOM attributes apply to the element: autocomplete, list, maxlength, pattern, readonly, required, and size content attributes; list, selectedOption, and value DOM attributes.
- The value DOM attribute is in mode value (page 345).
- The input and change events apply.
- The following content attributes must not be specified and do not apply to the element: accept, action, alt, checked, enctype, max, method, min, src, step, and target.
- The following DOM attributes and methods do not apply to the element: checked, valueAsDate, and valueAsNumber DOM attributes; stepUp(), and stepDown() methods.

4.10.4.1.5 Password state

When an input element's type attribute is in the Password (page 326) state, the rules in this section apply.

The input element represents a one line plain text edit control for the element's value (page 362). The user agent should obscure the value so that people other than the user cannot see it.

If the element is *mutable* (page 323), its value (page 362) should be editable by the user. User agents must not allow users to insert U+000A LINE FEED (LF) or U+000D CARRIAGE RETURN (CR) characters into the value (page 362).

The value sanitization algorithm (page 322) is as follows: Strip line breaks (page 32) from the value (page 362).

Bookkeeping details

- The following common input element content attributes and DOM attributes apply to the element: autocomplete, maxlength, pattern, readonly, required, and size content attributes; value DOM attribute.
- The input and change events apply.
- The following content attributes must not be specified and do not apply to the element: accept, action, alt, checked, enctype, list, max, method, min, src, step, and target.
- The following DOM attributes and methods do not apply to the element: checked, list, selectedOption, valueAsDate, and valueAsNumber DOM attributes; stepUp(), and stepDown() methods.

4.10.4.1.6 Date and Time state

When an input element's type attribute is in the Date and Time (page 326) state, the rules in this section apply.

The input element represents a control for setting the element's value (page 362) to a string representing a specific UTC date and time (page 47). User agents may display the date and time in whatever timezone is appropriate for the user.

If the element is *mutable* (page 323), the user agent should allow the user to change the UTC date and time (page 47) represented by its value (page 362), as obtained by parsing a UTC date and time (page 48) from it. User agents must not allow the user to set the value (page

362) to a string that is not a valid UTC date and time string (page 47). If the user agent provides a user interface for selecting a UTC date and time (page 47), then the value (page 362) must be set to a valid UTC date and time string (page 47) representing the user's selection. User agents should allow the user to set the value (page 362) to the empty string.

The value sanitization algorithm (page 322) is as follows: If the value (page 362) of the element is not a valid UTC date and time string (page 47), then set it to the empty string instead.

The min attribute, if specified, must have a value that is a valid UTC date and time string (page 47). The max attribute, if specified, must have a value that is a valid UTC date and time string (page 47).

The step attribute is expressed in seconds. The step scale factor (page 345) is 1000 (which converts the seconds to milliseconds, as used in the other algorithms). The default step (page 345) is 60 seconds.

When the element is suffering from a step mismatch (page 365), the user agent may round the element's value (page 362) to the nearest UTC date and time (page 47) for which the element would not suffer from a step mismatch (page 365).

The algorithm to convert a string to a number (page 322), given a string *input*, is as follows: If parsing a UTC date and time (page 48) from *input* results in an error, then return an error; otherwise, return the number of milliseconds elapsed from midnight UTC on the morning of 1970-01-01 (the time represented by the value "1970-01-01T00:00:00.0Z") to the parsed UTC date and time (page 47), ignoring leap seconds.

The algorithm to convert a number to a string (page 322), given a number *input*, is as follows: Return a valid UTC date and time string (page 47) that represents the date and time in UTC (page 47) that is *input* milliseconds after midnight UTC on the morning of 1970-01-01 (the time represented by the value "1970-01-01T00:00:00.0Z").

The algorithm to convert a string to a Date object (page 322), given a string *input*, is as follows: If parsing a UTC date and time (page 48) from *input* results in an error, then return an error; otherwise, return a Date object representing the parsed UTC date and time (page 47).

The algorithm to convert a Date object to a string (page 322), given a Date object *input*, is as follows: Return a valid UTC date and time string (page 47) that represents the date and time in UTC (page 47) that is represented by *input*.

Bookkeeping details

- The following common input element content attributes, DOM attributes, and methods apply to the element: autocomplete, list, max, min, readonly, required, and step content attributes; list, value, valueAsDate, valueAsNumber, and selectedOption DOM attributes; stepUp(), and stepDown() methods.
- The value DOM attribute is in mode value (page 345).
- The input and change events apply.
- The following content attributes must not be specified and do not apply to the element: accept, alt, action, checked, enctype, maxlength, method, pattern, size, src, and target.
- The checked DOM attribute does not apply to the element.

4.10.4.1.7 Date state

When an input element's type attribute is in the Date (page 328) state, the rules in this section apply.

The input element represents a control for setting the element's value (page 362) to a string representing a specific date (page 48).

If the element is *mutable* (page 323), the user agent should allow the user to change the date (page 48) represented by its value (page 362), as obtained by parsing a date (page 48) from it. User agents must not allow the user to set the value (page 362) to a string that is not a valid date string (page 48). If the user agent provides a user interface for selecting a date (page 48), then the value (page 362) must be set to a valid date string (page 48) representing the user's selection. User agents should allow the user to set the value (page 362) to the empty string.

The value sanitization algorithm (page 322) is as follows: If the value (page 362) of the element is not a valid date string (page 48), then set it to the empty string instead.

The min attribute, if specified, must have a value that is a valid date string (page 48). The max attribute, if specified, must have a value that is a valid date string (page 48).

The step attribute is expressed in days. The step scale factor (page 345) is 86,400,000 (which converts the days to milliseconds, as used in the other algorithms). The default step (page 345) is 1 day.

When the element is suffering from a step mismatch (page 365), the user agent may round the element's value (page 362) to the nearest date (page 48) for which the element would not suffer from a step mismatch (page 365).

The algorithm to convert a string to a number (page 322), given a string *input*, is as follows: If parsing a date (page 48) from *input* results in an error, then return an error; otherwise, return the number of milliseconds elapsed from midnight UTC on the morning of 1970-01-01 (the time represented by the value "1970-01-01T00:00:00.0Z") to midnight UTC on the morning of the parsed date (page 48), ignoring leap seconds.

The algorithm to convert a number to a string (page 322), given a number *input*, is as follows: Return a valid date string (page 48) that represents the date (page 48) that, in UTC, is current *input* milliseconds after midnight UTC on the morning of 1970-01-01 (the time represented by the value "1970-01-01T00:00:00.0Z").

The algorithm to convert a string to a Date object (page 322), given a string *input*, is as follows: If parsing a date (page 48) from *input* results in an error, then return an error; otherwise, return a Date object representing midnight UTC on the morning of the parsed date (page 48).

The algorithm to convert a Date object to a string (page 322), given a Date object *input*, is as follows: Return a valid date string (page 48) that represents the date (page 48) current at the time represented by *input* in the UTC timezone.

Bookkeeping details

- The following common input element content attributes, DOM attributes, and methods apply to the element: autocomplete, list, max, min, readonly, required, and step content attributes; list, value, valueAsDate, valueAsNumber, and selectedOption DOM attributes; stepUp(), and stepDown() methods.

- The value DOM attribute is in mode value (page 345).
- The input and change events apply.
- The following content attributes must not be specified and do not apply to the element: accept, action, alt, checked, enctype, maxlength, method, pattern, size, src, and target.
- The checked DOM attribute does not apply to the element.

4.10.4.1.8 Month state

When an input element's type attribute is in the Month (page 329) state, the rules in this section apply.

The input element represents a control for setting the element's value (page 362) to a string representing a specific month (page 48).

If the element is *mutable* (page 323), the user agent should allow the user to change the month (page 48) represented by its value (page 362), as obtained by parsing a month (page 48) from it. User agents must not allow the user to set the value (page 362) to a string that is not a valid month string (page 48). If the user agent provides a user interface for selecting a month (page 48), then the value (page 362) must be set to a valid month string (page 48) representing the user's selection. User agents should allow the user to set the value (page 362) to the empty string.

The value sanitization algorithm (page 322) is as follows: If the value (page 362) of the element is not a valid month string (page 48), then set it to the empty string instead.

The min attribute, if specified, must have a value that is a valid month string (page 48). The max attribute, if specified, must have a value that is a valid month string (page 48).

The step attribute is expressed in months. The step scale factor (page 345) is 1 (there is no conversion needed as the algorithms use months). The default step (page 345) is 1 month.

When the element is suffering from a step mismatch (page 365), the user agent may round the element's value (page 362) to the nearest month (page 48) for which the element would not suffer from a step mismatch (page 365).

The algorithm to convert a string to a number (page 322), given a string *input*, is as follows: If parsing a month time (page 48) from *input* results in an error, then return an error; otherwise, return the number of months between January 1970 and the parsed month (page 48).

The algorithm to convert a number to a string (page 322), given a number *input*, is as follows: Return a valid month string (page 48) that represents the month that has *input* months between it and January 1970.

The algorithm to convert a string to a Date object (page 322), given a string *input*, is as follows: If parsing a month (page 48) from *input* results in an error, then return an error; otherwise, return a Date object representing midnight UTC on the morning of the first day of the parsed month (page 48).

The algorithm to convert a Date object to a string (page 322), given a Date object *input*, is as follows: Return a valid month string (page 48) that represents the month (page 48) current at the time represented by *input* in the UTC timezone.

Bookkeeping details

- The following common input element content attributes, DOM attributes, and methods apply to the element: autocomplete, list, max, min, readonly, required, and step content attributes; list, value, valueAsDate, valueAsNumber, and selectedOption DOM attributes; stepUp(), and stepDown() methods.
- The value DOM attribute is in mode value (page 345).
- The input and change events apply.
- The following content attributes must not be specified and do not apply to the element: accept, action, alt, checked, enctype, maxlength, method, pattern, size, src, and target.
- The checked DOM attribute does not apply to the element.

4.10.4.1.9 Week state

When an input element's type attribute is in the Week (page 330) state, the rules in this section apply.

The input element represents a control for setting the element's value (page 362) to a string representing a specific week (page 48).

If the element is *mutable* (page 323), the user agent should allow the user to change the week (page 48) represented by its value (page 362), as obtained by parsing a week (page 48) from it. User agents must not allow the user to set the value (page 362) to a string that is not a valid week string (page 48). If the user agent provides a user interface for selecting a week (page 48), then the value (page 362) must be set to a valid week string (page 48) representing the user's selection. User agents should allow the user to set the value (page 362) to the empty string.

The value sanitization algorithm (page 322) is as follows: If the value (page 362) of the element is not a valid week string (page 48), then set it to the empty string instead.

The min attribute, if specified, must have a value that is a valid week string (page 48). The max attribute, if specified, must have a value that is a valid week string (page 48).

The step attribute is expressed in weeks. The step scale factor (page 345) is 604,800,000 (which converts the weeks to milliseconds, as used in the other algorithms). The default step (page 345) is 1 week.

When the element is suffering from a step mismatch (page 365), the user agent may round the element's value (page 362) to the nearest week (page 48) for which the element would not suffer from a step mismatch (page 365).

The algorithm to convert a string to a number (page 322), given a string *input*, is as follows: If parsing a week string (page 48) from *input* results in an error, then return an error; otherwise, return the number of milliseconds elapsed from midnight UTC on the morning of 1970-01-01 (the time represented by the value "1970-01-01T00:00:00.0Z") to midnight UTC on the morning of the Monday of the parsed week (page 48), ignoring leap seconds.

The algorithm to convert a number to a string (page 322), given a number *input*, is as follows: Return a valid week string (page 48) that represents the week (page 48) that, in UTC, is current *input* milliseconds after midnight UTC on the morning of 1970-01-01 (the time represented by the value "1970-01-01T00:00:00.0Z").

The algorithm to convert a string to a Date object (page 322), given a string *input*, is as follows: If parsing a week (page 48) from *input* results in an error, then return an error; otherwise, return a Date object representing midnight UTC on the morning of the Monday of the parsed week (page 48).

The algorithm to convert a Date object to a string (page 322), given a Date object *input*, is as follows: Return a valid week string (page 48) that represents the week (page 48) current at the time represented by *input* in the UTC timezone.

Bookkeeping details

- The following common input element content attributes, DOM attributes, and methods apply to the element: autocomplete, list, max, min, readonly, required, and step content attributes; list, value, valueAsDate, valueAsNumber, and selectedOption DOM attributes; stepUp(), and stepDown() methods.
- The value DOM attribute is in mode value (page 345).
- The input and change events apply.
- The following content attributes must not be specified and do not apply to the element: accept, action, alt, checked, enctype, maxlength, method, pattern, size, src, and target.
- The checked DOM attribute does not apply to the element.

4.10.4.1.10 Time state

When an input element's type attribute is in the Time (page 331) state, the rules in this section apply.

The input element represents a control for setting the element's value (page 362) to a string representing a specific time (page 49).

If the element is *mutable* (page 323), the user agent should allow the user to change the time (page 49) represented by its value (page 362), as obtained by parsing a time (page 49) from it. User agents must not allow the user to set the value (page 362) to a string that is not a valid time string (page 49). If the user agent provides a user interface for selecting a time (page 49), then the value (page 362) must be set to a valid time string (page 49) representing the user's selection. User agents should allow the user to set the value (page 362) to the empty string.

The value sanitization algorithm (page 322) is as follows: If the value (page 362) of the element is not a valid time string (page 49), then set it to the empty string instead.

The min attribute, if specified, must have a value that is a valid time string (page 49). The max attribute, if specified, must have a value that is a valid time string (page 49).

The step attribute is expressed in seconds. The step scale factor (page 345) is 1000 (which converts the seconds to milliseconds, as used in the other algorithms). The default step (page 345) is 60 seconds.

When the element is suffering from a step mismatch (page 365), the user agent may round the element's value (page 362) to the nearest time (page 49) for which the element would not suffer from a step mismatch (page 365).

The algorithm to convert a string to a number (page 322), given a string *input*, is as follows: If parsing a time (page 49) from *input* results in an error, then return an error; otherwise, return the number of milliseconds elapsed from midnight to the parsed time (page 49) on a day with no time changes.

The algorithm to convert a number to a string (page 322), given a number *input*, is as follows: Return a valid time string (page 49) that represents the time (page 49) that is *input* milliseconds after midnight on a day with no time changes.

The algorithm to convert a string to a Date object (page 322), given a string *input*, is as follows: If parsing a time (page 49) from *input* results in an error, then return an error; otherwise, return a Date object representing the parsed time (page 49) in UTC on 1970-01-01.

The algorithm to convert a Date object to a string (page 322), given a Date object *input*, is as follows: Return a valid time string (page 49) that represents the UTC time (page 49) component that is represented by *input*.

Bookkeeping details

- The following common input element content attributes, DOM attributes, and methods apply to the element: autocomplete, list, max, min, readonly, required, and step content attributes; list, value, valueAsDate, valueAsNumber, and selectedOption DOM attributes; stepUp(), and stepDown() methods.
- The value DOM attribute is in mode value (page 345).
- The input and change events apply.
- The following content attributes must not be specified and do not apply to the element: accept, action, alt, checked, enctype, maxlength, method, pattern, size, src, and target.
- The checked DOM attribute does not apply to the element.

4.10.4.1.11 Local Date and Time state

When an input element's type attribute is in the Local Date and Time (page 332) state, the rules in this section apply.

The input element represents a control for setting the element's value (page 362) to a string representing a local date and time (page 48), with no time zone information.

If the element is *mutable* (page 323), the user agent should allow the user to change the date and time (page 48) represented by its value (page 362), as obtained by parsing a date and time (page 48) from it. User agents must not allow the user to set the value (page 362) to a string that is not a valid local date and time string (page 48). If the user agent provides a user interface for selecting a local date and time (page 48), then the value (page 362) must be set to a valid local date and time string (page 48) representing the user's selection. User agents should allow the user to set the value (page 362) to the empty string.

The value sanitization algorithm (page 322) is as follows: If the value (page 362) of the element is not a valid local date and time string (page 48), then set it to the empty string instead.

The min attribute, if specified, must have a value that is a valid local date and time string (page 48). The max attribute, if specified, must have a value that is a valid local date and time string (page 48).

The step attribute is expressed in seconds. The step scale factor (page 345) is 1000 (which converts the seconds to milliseconds, as used in the other algorithms). The default step (page 345) is 60 seconds.

When the element is suffering from a step mismatch (page 365), the user agent may round the element's value (page 362) to the nearest local date and time (page 48) for which the element would not suffer from a step mismatch (page 365).

The algorithm to convert a string to a number (page 322), given a string *input*, is as follows: If parsing a date and time (page 48) from *input* results in an error, then return an error; otherwise, return the number of milliseconds elapsed from midnight on the morning of 1970-01-01 (the time represented by the value "1970-01-01T00:00:00.0") to the parsed local date and time (page 48), ignoring leap seconds.

The algorithm to convert a number to a string (page 322), given a number *input*, is as follows: Return a valid local date and time string (page 48) that represents the date and time that is *input* milliseconds after midnight on the morning of 1970-01-01 (the time represented by the value "1970-01-01T00:00:00.0").

Bookkeeping details

- The following common input element content attributes, DOM attributes, and methods apply to the element: autocomplete, list, max, min, readonly, required, and step content attributes; list, value, valueAsNumber, and selectedOption DOM attributes; stepUp(), and stepDown() methods.
- The value DOM attribute is in mode value (page 345).
- The input and change events apply.
- The following content attributes must not be specified and do not apply to the element: accept, action, alt, checked, enctype, maxlength, method, pattern, size, src, and target.
- The following DOM attributes do not apply to the element: valueAsDate and checked.

4.10.4.1.12 Number state

When an `input` element's `type` attribute is in the Number (page 333) state, the rules in this section apply.

The `input` element represents a control for setting the element's value (page 362) to a string representing a number.

If the element is *mutable* (page 323), the user agent should allow the user to change the number represented by its value (page 362), as obtained from applying the rules for parsing floating point number values (page 34) to it. User agents must not allow the user to set the value (page 362) to a string that is not a valid floating point number (page 34). If the user agent provides a user interface for selecting a number, then the value (page 362) must be set to a valid floating point number (page 34) representing the user's selection. User agents should allow the user to set the value (page 362) to the empty string.

The value sanitization algorithm (page 322) is as follows: If the value (page 362) of the element is not a valid floating point number (page 34), then set it to the empty string instead.

The `min` attribute, if specified, must have a value that is a valid floating point number (page 34). The `max` attribute, if specified, must have a value that is a valid floating point number (page 34).

The step scale factor (page 345) is 1. The default step (page 345) is 1 (allowing only integers, unless the `min` attribute has a non-integer value).

When the element is suffering from a step mismatch (page 365), the user agent may round the element's value (page 362) to the nearest number for which the element would not suffer from a step mismatch (page 365).

The algorithm to convert a string to a number (page 322), given a string *input*, is as follows: If applying the rules for parsing floating point number values (page 34) to *input* results in an error, then return an error; otherwise, return the resulting number.

The algorithm to convert a number to a string (page 322), given a number *input*, is as follows: Return a valid floating point number (page 34) that represents *input*.

Bookkeeping details

- The following common input element content attributes, DOM attributes, and methods apply to the element: autocomplete, list, max, min, readonly, required, and step content attributes; list, value, valueAsNumber, and selectedOption DOM attributes; stepUp(), and stepDown() methods.
- The value DOM attribute is in mode value (page 345).
- The input and change events apply.
- The following content attributes must not be specified and do not apply to the element: accept, action, alt, checked, enctype, maxlength, method, pattern, size, src, and target.
- The following DOM attributes do not apply to the element: valueAsDate and checked.

4.10.4.1.13 Range state

When an input element's type attribute is in the Range (page 334) state, the rules in this section apply.

The input element represents a control for setting the element's value (page 362) to a string representing a number, but with the caveat that the exact value is not important, letting UAs provide a simpler interface than they do for the Number (page 333) state.

Note: In this state, the range and step constraints are enforced even during user input, and there is no way to set the value to the empty string.

If the element is *mutable* (page 323), the user agent should allow the user to change the number represented by its value (page 362), as obtained from applying the rules for parsing floating point number values (page 34) to it. User agents must not allow the user to set the value (page 362) to a string that is not a valid floating point number (page 34). If the user agent provides a user interface for selecting a number, then the value (page 362) must be set to a valid floating point number (page 34) representing the user's selection. User agents must not allow the user to set the value (page 362) to the empty string.

The value sanitization algorithm (page 322) is as follows: If the value (page 362) of the element is not a valid floating point number (page 34), then set it to a valid floating point number (page 34) that represents the default value (page 334).

The min attribute, if specified, must have a value that is a valid floating point number (page 34). The default minimum (page 344) is 0. The max attribute, if specified, must have a value that is a valid floating point number (page 34). The default minimum (page 344) is 100.

The **default value** is the minimum (page 344) plus half the difference between the minimum (page 344) and the maximum (page 344), unless the maximum (page 344) is less than the minimum (page 344), in which case the default value (page 334) is the minimum (page 344).

When the element is suffering from a range underflow, the user agent must set the element's value (page 362) to a valid floating point number (page 34) that represents the minimum (page 344).

When the element is suffering from a range overflow, if the maximum (page 344) is not less than the minimum (page 344), the user agent must set the element's value (page 362) to a valid floating point number (page 34) that represents the maximum (page 344).

The step scale factor (page 345) is 1. The default step (page 345) is 1 (allowing only integers, unless the `min` attribute has a non-integer value).

When the element is suffering from a step mismatch (page 365), the user agent must round the element's value (page 362) to the nearest number for which the element would not suffer from a step mismatch (page 365), and which is greater than or equal to the minimum (page 344), and, if the maximum (page 344) is not less than the minimum (page 344), which is less than or equal to the maximum (page 344).

The algorithm to convert a string to a number (page 322), given a string *input*, is as follows: If applying the rules for parsing floating point number values (page 34) to *input* results in an error, then return an error; otherwise, return the resulting number.

The algorithm to convert a number to a string (page 322), given a number *input*, is as follows: Return a valid floating point number (page 34) that represents *input*.

Bookkeeping details

- The following common input element content attributes, DOM attributes, and methods apply to the element: `autocomplete`, `list`, `max`, `min`, and `step` content attributes; `list`, `value`, `valueAsNumber`, and `selectedOption` DOM attributes; `stepUp()`, and `stepDown()` methods.
- The `value` DOM attribute is in mode value (page 345).
- The `input` and `change` events apply.
- The following content attributes must not be specified and do not apply to the element: `accept`, `action`, `alt`, `checked`, `enctype`, `maxlength`, `method`, `pattern`, `readonly`, `required`, `size`, `src`, and `target`.
- The following DOM attributes do not apply to the element: `valueAsDate` and `checked`.

4.10.4.1.14 Checkbox state

When an `input` element's `type` attribute is in the Checkbox (page 335) state, the rules in this section apply.

The `input` element represents a two-state control that represents the element's checkedness (page 362) state. If the element's checkedness (page 362) state is true, the control represents a positive selection, and if it is false, a negative selection.

If the element is *mutable* (page 323), then: The pre-click activation steps (page 96) consist of setting the element's checkedness (page 362) to its opposite value (i.e. true if it is false, false if it is true). The canceled activation steps (page 96) consist of setting the checkedness (page 362) back to the value it had before the pre-click activation steps (page 96) were run. The activation behavior (page 96) is to fire a simple event (page 436) called `change` at the element.

Constraint validation: If the element is *required* (page 343) and its checkedness is false, then the element is suffering from being missing (page 365).

Bookkeeping details

- The following common input element content attributes and DOM attributes apply to the element: checked, and required content attributes; checked and value DOM attributes.
- The value DOM attribute is in mode default/on (page 346).
- The change event applies.
- The following content attributes must not be specified and do not apply to the element: accept, action, alt, autocomplete, enctype, list, max, maxlength, method, min, pattern, readonly, size, src, step, and target.
- The following DOM attributes and methods do not apply to the element: list, selectedOption, valueAsDate and valueAsNumber DOM attributes; stepDown() and stepUp() methods.
- The input event does not apply.

4.10.4.1.15 Radio Button state

When an input element's type attribute is in the Radio Button (page 336) state, the rules in this section apply.

The input element represents a control that, when used in conjunction with other input elements, forms a *radio button group* (page 336) in which only one control can have its checkedness (page 362) state set to true. If the element's checkedness (page 362) state is true, the control represents the selected control in the group, and if it is false, it indicates a control in the group that is not selected.

The **radio button group** that contains an input element *a* also contains all the other input elements *b* that fulfill all of the following conditions:

- The input element *b*'s type attribute is in the Radio Button (page 336) state.
- Either neither *a* nor *b* have a form owner (page 361), or they both have one and it is the same for both.
- They both have a name attribute, and the value of *a*'s name attribute is a compatibility caseless (page 31) match for the value of *b*'s name attribute.

A document must not contain an input element whose *radio button group* (page 336) contains only that element.

When any of the following events occur, if the element's checkedness (page 362) state is true after the event, the checkedness (page 362) state of all the other elements in the same *radio button group* (page 336) must be set to false:

- The element's checkedness (page 362) state is set to true (for whatever reason).
- The element's name attribute is added, removed, or changes value.
- The element's form owner (page 361) changes.

If the element is *mutable* (page 323), then: The pre-click activation steps (page 96) consist of setting the element's checkedness (page 362) to true. The canceled activation steps (page 96) consist of setting the element's checkedness (page 362) to false. The activation behavior (page 96) is to fire a simple event (page 436) called change at the element.

Constraint validation: If the element is *required* (page 343) and all of the input elements in the *radio button group* (page 336) have a checkedness that is false, then the element is suffering from being missing (page 365).

Bookkeeping details

- The following common input element content attributes and DOM attributes apply to the element: checked and required content attributes; checked and value DOM attributes.
- The value DOM attribute is in mode default/on (page 346).
- The change event applies.
- The following content attributes must not be specified and do not apply to the element: accept, action, alt, autocomplete, enctype, list, max, maxlength, method, min, pattern, readonly, size, src, step, and target.
- The following DOM attributes and methods do not apply to the element: list, selectedOption, valueAsDate and valueAsNumber DOM attributes; stepDown() and stepUp() methods.
- The input event does not apply.

4.10.4.1.16 File Upload state

When an input element's type attribute is in the File Upload (page 337) state, the rules in this section apply.

The input element represents a list of **selected files**, each file consisting of a file name, a file type, and a file body (the contents of the file).

If the element is *mutable* (page 323), the user agent should allow the user to change the files on the list, e.g. adding or removing files. Files can be from the filesystem or created on the fly, e.g. a picture taken from a camera connected to the user's device.

Constraint validation: If the element is *required* (page 343) and the list of selected files (page 337) is empty, then the element is suffering from being missing (page 365).

There must be no more than one file in the list of selected files (page 337).

The **accept** attribute may be specified to provide user agents with a hint of what file types the server will be able to accept.

If specified, the attribute must consist of a set of comma-separated tokens (page 51), each of which must be an ASCII case-insensitive (page 31) match for one of the following:

The string audio/*

Indicates that sound files are accepted.

The string video/*

Indicates that video files are accepted.

The string image/*

Indicates that image files are accepted.

A valid MIME type, with no parameters

Indicates that files of the specified type are accepted. RFC[2046]

The tokens must not be ASCII case-insensitive (page 31) matches for any of the other tokens (i.e. duplicates are not allowed).

User agents should prevent the user from selecting files that are not accepted by one (or more) of these tokens.

Bookkeeping details

- The following common input element content attributes apply to the element: accept and required.
- The change event applies.
- The following content attributes must not be specified and do not apply to the element: action, alt, autocomplete, checked, enctype, list, max, maxlength, method, min, pattern, readonly, size, src, step, and target.
- The element's value attribute must be omitted.
- The following DOM attributes and methods do not apply to the element: checked, list, selectedOption, value, valueAsDate and valueAsNumber DOM attributes; stepDown() and stepUp() methods.
- The input event does not apply.

4.10.4.1.17 Submit Button state

When an input element's type attribute is in the Submit Button (page 338) state, the rules in this section apply.

The input element represents a button that, when activated, submits the form. If the element has a value attribute, the button's label must be the value of that attribute; otherwise, it must be an implementation-defined string that means "Submit" or some such.

If the element is *mutable* (page 323), the user agent should allow the user to activate the element.

The element's activation behavior (page 96), if the element has a form owner (page 361), is to submit (page 368) the form owner (page 361) from the input element; otherwise, it is to do nothing.

The action, enctype, method, and target attributes are attributes for form submission (page 363).

Bookkeeping details

- The following common input element content attributes and DOM attributes apply to the element: action, enctype, method, and target content attributes; value DOM attribute.
- The value DOM attribute is in mode default (page 346).
- The following content attributes must not be specified and do not apply to the element: accept, alt, autocomplete, checked, list, max, maxlength, min, pattern, readonly, required size, src, and step.
- The following DOM attributes and methods do not apply to the element: checked, list, selectedOption, valueAsDate and valueAsNumber DOM attributes; stepDown() and stepUp() methods.
- The input and change events do not fire.

4.10.4.1.18 Image Button state

When an input element's type attribute is in the Image Button (page 338) state, the rules in this section apply.

The input element represents either an image from which a user can select a coordinate and submit the form, or alternatively a button from which the user can submit the form.

The image is given by the **src** attribute. The **src** attribute must be present, and must contain a valid URL (page 52) referencing a non-interactive, optionally animated, image resource that is neither paged nor scripted.

When any of the following events occur, the user agent must fetch (page 59) the resource specified by the **src** attribute's value, unless the user agent cannot support images, or its support for images has been disabled, or the user agent only fetches elements on demand:

- The input element's type attribute is first set to the Image Button (page 338) state (possibly when the element is first created), and the **src** attribute is present.
- The input element's type attribute is changed back to the Image Button (page 338) state, and the **src** attribute is present, and its value has changed since the last time the type attribute was in the Image Button (page 338) state.
- The input element's type attribute is in the Image Button (page 338) state, and the **src** attribute is set or changed.

Fetching the image must delay the load event (page 657).

If the image was successfully obtained, with no network errors, and the image's type is a supported image type, and the image is a valid image of that type, then the image is said to be *available*. If this is true before the image is completely downloaded, each task (page 429) that is queued (page 429) by the networking task source (page 430) while the image is being fetched (page 59) must update the presentation of the image appropriately.

The user agents should apply the image sniffing rules (page 65) to determine the type of the image, with the image's associated Content-Type headers (page 60) giving the *official type*. If these rules are not applied, then the type of the image must be the type given by the image's associated Content-Type headers (page 60).

User agents must not support non-image resources with the **input** element. User agents must not run executable code embedded in the image resource. User agents must only display the first page of a multipage resource. User agents must not allow the resource to act in an interactive fashion, but should honour any animation in the resource.

The task (page 429) that is queued (page 429) by the networking task source (page 430) once the resource has been fetched (page 59), must, if the download was successful and the image is *available*, queue a task (page 429) to fire a load event (page 436) on the **input** element; and otherwise, if the fetching process fails without a response from the remote server, or completes but the image is not a valid or supported image, queue a task (page 429) to fire an error event (page 436) on the **input** element.

The **alt** attribute provides the textual label for the alternative button for users and user agents who cannot use the image. The **alt** attribute must also be present, and must contain a non-empty string.

If the **src** attribute is set, and the image is *available* and the user agent is configured to display that image, then: The element represents a control for selecting a coordinate (page 340) from the image specified by the **src** attribute; if the element is *mutable* (page 323), the user agent should allow the user to select this coordinate (page 340). The activation behavior (page 96) in this case consists of taking the user's selected coordinate (page 340), and then,

if the element has a form owner (page 361), submitting (page 368) the input element's form owner (page 361) from the input element.

Otherwise, the element represents a submit button whose label is given by the value of the alt attribute; if the element is *mutable* (page 323), the user agent should allow the user to activate the button. The activation behavior (page 96) in this case consists of setting the selected coordinate (page 340) to (0,0), and then, if the element has a form owner (page 361), submitting (page 368) the input element's form owner (page 361) from the input element.

The **selected coordinate** must consist of an x-component and a y-component. The x-component must be greater than or equal to zero, and less than or equal to the rendered width, in CSS pixels, of the image. The y-component must be greater than or equal to zero, and less than or equal to the rendered height, in CSS pixels, of the image.

The action, enctype, method, and target attributes are attributes for form submission (page 363).

Bookkeeping details

- The following common input element content attributes and DOM attributes apply to the element: action, alt, enctype, method, src, and target content attributes; value DOM attribute.
- The value DOM attribute is in mode default (page 346).
- The following content attributes must not be specified and do not apply to the element: accept, autocomplete, checked, list, max, maxlength, min, pattern, readonly, required size, and step.
- The element's value attribute must be omitted.
- The following DOM attributes and methods do not apply to the element: checked, list, selectedOption, valueAsDate and valueAsNumber DOM attributes; stepDown() and stepUp() methods.
- The input and change events do not fire.

Note: Many aspects of this state's behavior are similar to the behavior of the `img` element. Readers are encouraged to read that section, where many of the same requirements are described in more detail.

4.10.4.1.19 Reset Button state

When an input element's type attribute is in the Reset Button (page 340) state, the rules in this section apply.

The input element represents a button that, when activated, resets the form. If the element has a value attribute, the button's label must be the value of that attribute; otherwise, it must be an implementation-defined string that means "Reset" or some such.

If the element is *mutable* (page 323), the user agent should allow the user to activate the element.

The element's activation behavior (page 96), if the element has a form owner (page 361), is to reset (page 376) the form owner (page 361); otherwise, it is to do nothing.

Constraint validation: The element is barred from constraint validation (page 364).

The value DOM attribute applies to this element and is in mode default (page 346).

Bookkeeping details

- The following content attributes must not be specified and do not apply to the element: accept, action, alt, autocomplete, checked, enctype, list, max, maxlength, method, min, pattern, readonly, required size, src, step, and target.
- The following DOM attributes and methods do not apply to the element: checked, list, selectedOption, valueAsDate and valueAsNumber DOM attributes; stepDown() and stepUp() methods.
- The input and change events do not fire.

4.10.4.1.20 Button state

When an input element's type attribute is in the Button (page 341) state, the rules in this section apply.

The input element represents a button with no default behavior. If the element has a value attribute, the button's label must be the value of that attribute; otherwise, it must be the empty string.

If the element is *mutable* (page 323), the user agent should allow the user to activate the element. The element's activation behavior (page 96) is to do nothing.

Constraint validation: The element is barred from constraint validation (page 364).

The value DOM attribute applies to this element and is in mode default (page 346).

Bookkeeping details

- The following content attributes must not be specified and do not apply to the element: accept, action, alt, autocomplete, checked, enctype, list, max, maxlength, method, min, pattern, readonly, required size, src, step, and target.
- The following DOM attributes and methods do not apply to the element: checked, list, selectedOption, valueAsDate and valueAsNumber DOM attributes; stepDown() and stepUp() methods.
- The input and change events do not fire.

4.10.4.2 Common input element attributes

These attributes only apply to an input element if its type attribute is in a state whose definition declares that the attribute applies. When an attribute doesn't apply to an input element, user agents must ignore (page 24) the attribute.

4.10.4.2.1 The autocomplete attribute

The **autocomplete** attribute is an enumerated attribute (page 51). The attribute has two states. The **on** keyword maps to the **on** state, and the **off** keyword maps to the **off** state. The attribute may also be omitted. The *missing value default* is the **on** (page 341) state.

The **off** (page 341) state indicates that the control's input data is either particularly sensitive (for example the activation code for a nuclear weapon) or is a value that will never be reused (for example a one-time-key for a bank login) and the user will therefore have to explicitly enter the data each time, instead of being able to rely on the UA to prefill the value for him.

Conversely, the **on** (page 341) state indicates that the value is not particularly sensitive and the user can expect to be able to rely on his user agent to remember values he has entered for that control.

When an `input` element's `autocomplete` attribute is in the `on` (page 341) state, the user agent may store the value entered by the user so that if the user returns to the page, the UA can prefill the form. When an `input` element's `autocomplete` attribute is in the `off` (page 341) state, the user agent should not remember the control's value (page 362).

The autocompletion mechanism must be implemented by the user agent acting as if the user had modified the element's value (page 362), and must be done at a time where the element is *mutable* (page 323) (e.g. just after the element has been inserted into the document, or when the user agent stops parsing (page 656)).

Banks frequently do not want UAs to prefill login information:

```
<p>Account: <input type="text" name="ac" autocomplete="off"></p>
<p>PIN: <input type="text" name="pin" autocomplete="off"></p>
```

A user agent may allow the user to disable support for this attribute's `off` (page 341) state (causing the attribute to always be in the `on` (page 341) state and always allowing values to be remembered and prefilled). Support for the `off` (page 341) state should be enabled by default, and the ability to disable support should not be trivially accessible, as there are significant security implications for the user if support for this attribute is disabled.

4.10.4.2.2 The `list` attribute

The `list` attribute is used to identify an element that lists predefined options suggested to the user.

If present, its value must be the ID of a `select` or `datagrid` element in the same document.

The **suggestions source element** is the first element in the document in tree order (page 24) to have an ID equal to the value of the `list` attribute, if that element is either a `select` or `datalist` element. If there is no `list` attribute, or if there is no element with that ID, or if the first element with that ID is not either a `select` or `datalist` element, then there is no suggestions source element (page 342).

If there is a suggestions source element (page 342), then each `option` element that is a descendant of the suggestions source element (page 342), that is not disabled (page 355), and whose value (page 355) is a string that isn't the empty string and that the user would be allowed to enter as the `input` element's value (page 362), represents a suggestion. Each suggestion has a value (page 355) and a label (page 355).

When the user agent is allowing the user to edit the `input` element's value (page 362), the user agent should offer the suggestions to the user in a manner suitable for the type of control used. The user agent may use the suggestion's label (page 355) to identify the suggestion if appropriate. If the user selects a suggestion, then the `input` element's value (page 362) must be set to the selected suggestion's value (page 355), as if the user had written that value himself.

User agents should filter the suggestions to hide suggestions that would cause the element to not satisfy its constraints (page 365).

If the `list` attribute does not apply, there is no suggestions source element (page 342).

4.10.4.2.3 The `readonly` attribute

The **readonly** attribute is a boolean attribute (page 32). When specified, the element is *immutable* (page 323).

Constraint validation: If the `readonly` attribute is specified on an input element, the element is barred from constraint validation (page 364).

4.10.4.2.4 The `size` attribute

The **size** attribute gives the number of characters that, in a visual rendering, the user agent is to allow the user to see while editing the element's value (page 362).

The `size` attribute, if specified, must have a value that is a valid non-negative integer (page 33) greater than zero.

If the attribute is present, then its value must be parsed using the rules for parsing non-negative integers (page 33), and if the result is a number greater than zero, then the user agent should ensure that at least that many characters are visible.

** The rendering section will define this in more detail.

The `size` DOM attribute limited to only positive non-zero numbers (page 68).

4.10.4.2.5 The `required` attribute

The **required** attribute is a boolean attribute (page 32). When specified, the element is *required*.

Constraint validation: If the element is *required* (page 343), and its value DOM attribute applies and is in the mode value (page 345), and the element is *mutable* (page 323), and the element's value (page 362) is the empty string, then the element is suffering from being missing (page 365).

4.10.4.2.6 The `maxlength` attribute

The **maxlength** attribute, when it applies, is a form control `maxlength` attribute (page 363) controlled by the `input` element's dirty value flag (page 322).

If the `input` element has a maximum allowed value length (page 363), then the codepoint length (page 32) of the value of the element's `value` attribute must be equal to or less than the element's maximum allowed value length (page 363).

4.10.4.2.7 The `pattern` attribute

The **pattern** attribute specifies a regular expression against which the control's value (page 362) is to be checked.

If specified, the attribute's value must match the *Pattern* production of ECMA 262's grammar. [ECMA262]

Constraint validation: If the element's value (page 362) is not the empty string, and the element's pattern attribute is specified and the attribute's value, when compiled as an ECMA 262 regular expression with the global, ignoreCase, and multiline flags *disabled* (see ECMA 262, sections 15.10.7.2 through 15.10.7.4), compiles successfully but the resulting regular expression does not match the entirety of the element's value (page 362), then the element is suffering from a pattern mismatch (page 365). [ECMA262]

Note: This implies that the regular expression language used for this attribute is the same as that defined in ECMA 262, except that the pattern attribute must match the entire value, not just any subset (somewhat as if it implied a `^(: at the start of the pattern and a)$ at the end)`.

4.10.4.2.8 The min and max attributes

The **min** and **max** attributes indicate the allowed range of values for the element.

Their syntax is defined by the section that defines the type attribute's current state.

If the element has a **min** attribute, and the result of applying the algorithm to convert a string to a number (page 322) to the value of the **min** attribute is a a number, then that number is the element's **minimum**; otherwise, if the the type attribute's current state defines a **default minimum**, then that is the minimum (page 344); otherwise, the element has no minimum (page 344).)

Constraint validation: When the element has a minimum (page 344), and the result of applying the algorithm to convert a string to a number (page 322) to the string given by the element's value (page 362) is a number, and the number obtained from that algorithm is less than the minimum (page 344), the element is suffering from an underflow (page 365).

The **min** attribute also defines the step base (page 345).

If the element has a **max** attribute, and the result of applying the algorithm to convert a string to a number (page 322) to the value of the **max** attribute is a a number, then that number is the element's **maximum**; otherwise, if the the type attribute's current state defines a **default maximum**, then that is the maximum (page 344); otherwise, the element has no maximum (page 344).)

Constraint validation: When the element has a maximum (page 344), and the result of applying the algorithm to convert a string to a number (page 322) to the string given by the element's value (page 362) is a number, and the number obtained from that algorithm is more than the maximum (page 344), the element is suffering from an overflow (page 365).

The **max** attribute's value (the maximum (page 344)) must not be less than the **min** attribute's value (its minimum (page 344)).

Note: If an element has a maximum (page 344) that is less than its minimum (page 344), then so long as the element has a value (page 362), it will either be suffering from an underflow (page 365) or suffering from an overflow (page 365).

4.10.4.2.9 The step attribute

The **step** attribute indicates the granularity that is expected (and required) of the value (page 362), by limiting the allowed values. The section that defines the type attribute's current state also defines the **default step** and the **step scale factor**, which are used in processing the attribute as described below.

The step attribute, if specified, must either have a value that is a valid floating point number (page 34) that parses (page 34) to a number that is greater than zero, or must have a value that is an ASCII case-insensitive (page 31) match for the string "any".

The attribute provides the **allowed value step** for the element, as follows:

1. If the attribute is absent, then the allowed value step (page 345) is the default step (page 345) multiplied by the step scale factor (page 345).
2. Otherwise, if the attribute's value is an ASCII case-insensitive (page 31) match for the string "any", then there is no allowed value step (page 345).
3. Otherwise, if the rules for parsing floating point number values (page 34), when they are applied to the attribute's value, return an error, zero, or a number less than zero, then the allowed value step (page 345) is the default step (page 345) multiplied by the step scale factor (page 345).
4. Otherwise, the allowed value step (page 345) is the number returned by the rules for parsing floating point number values (page 34) when they are applied to the attribute's value, multiplied by the step scale factor (page 345).

The **step base** is the result of applying the algorithm to convert a string to a number (page 322) to the value of the **min** attribute, unless the element does not have a **min** attribute specified or the result of applying that algorithm is an error, in which case the step base (page 345) is zero.

Constraint validation: When the element has an allowed value step (page 345), and the result of applying the algorithm to convert a string to a number (page 322) to the string given by the element's value (page 362) is a number, and that number subtracted from the step base (page 345) is not an integral multiple of the allowed value step (page 345), the element is suffering from a step mismatch (page 365).

4.10.4.3 Common input element APIs

The **value** DOM attribute allows scripts to manipulate the value (page 362) of an input element. If the attribute applies, then it is in one of the following modes, which define its behavior:

value

On getting, it must return the current value (page 362) of the element. On setting, it must set the element's value (page 362) to the new value, set the element's dirty value flag (page 322) to true, and then invoke the value sanitization algorithm (page 322), if the element's type attribute's current state defines one.

default

On getting, if the element has a value attribute, it must return that attribute's value; otherwise, it must return the empty string. On setting, it must set the element's value attribute to the new value.

default/on

On getting, if the element has a value attribute, it must return that attribute's value; otherwise, it must return the string "on". On setting, it must set the element's value attribute to the new value.

If the attribute does not apply, then on getting and setting it must throw an INVALID_ACCESS_ERR exception.

The **checked** DOM attribute allows scripts to manipulate the checkedness (page 362) of an input element. On getting, it must return the current checkedness (page 362) of the element; and on setting, it must set the element's checkedness (page 362) to the new value and set the element's dirty checkedness flag (page 323) to true.

The **valueAsDate** DOM attribute represents the value (page 362) of the element, interpreted as a date.

On getting, if the valueAsDate attribute does not apply, as defined for the input element's type attribute's current state, then return null. Otherwise, run the algorithm to convert a string to a Date object (page 322) defined for that state; if the algorithm returned a Date object, then return it, otherwise, return null.

On setting, if the valueAsDate attribute does not apply, as defined for the input element's type attribute's current state, then throw an INVALID_ACCESS_ERR exception; otherwise, if the new value is null, then set the value (page 362) of the element to the empty string; otherwise, run the algorithm to convert a Date object to a string (page 322), as defined for that state, on the new value, and set the value (page 362) of the element to resulting string.

The **valueAsNumber** DOM attribute represents the value (page 362) of the element, interpreted as a number.

On getting, if the valueAsNumber attribute does not apply, as defined for the input element's type attribute's current state, then return a Not-a-Number (NaN) value. Otherwise, if the valueAsDate attribute applies, run the algorithm to convert a string to a Date object (page 322) defined for that state; if the algorithm returned a Date object, then return the *time value* of the object (the number of milliseconds from midnight UTC the morning of 1970-01-01 to the time represented by the Date object), otherwise, return a Not-a-Number (NaN) value. Otherwise, run the algorithm to convert a string to a number (page 322) defined for that state; if the algorithm returned a number, then return it, otherwise, return a Not-a-Number (NaN) value.

On setting, if the valueAsNumber attribute does not apply, as defined for the input element's type attribute's current state, then throw an INVALID_ACCESS_ERR exception. Otherwise, if the valueAsDate attribute applies, run the algorithm to convert a Date object to a string (page 322) defined for that state, passing it a Date object whose *time value* is the new value, and set the value (page 362) of the element to resulting string. Otherwise, run the algorithm

to convert a number to a string (page 322), as defined for that state, on the new value, and set the value (page 362) of the element to resulting string.

The **stepUp()** and **stepDown()** methods, when invoked, must run the following algorithm:

1. If the **stepUp()** and **stepDown()** methods do not apply, as defined for the input element's type attribute's current state, then throw an **INVALID_ACCESS_ERR** exception, and abort these steps.
2. If the element has no allowed value step (page 345), then throw an **INVALID_ACCESS_ERR** exception, and abort these steps.
3. If applying the algorithm to convert a string to a number (page 322) to the string given by the element's value (page 362) results in an error, then throw an **INVALID_ACCESS_ERR** exception, and abort these steps; otherwise, let *value* be the result of that algorithm.
4. If the method invoked was the **stepDown()** method, negate *value*.
5. Let *value* be the result of adding the allowed value step (page 345) to *value*.
6. Let *value as string* be the result of running the algorithm to convert a number to a string (page 322), as defined for the input element's type attribute's current state, on *value*.
7. If the element has a minimum (page 344), and the *value* is less than that minimum (page 344), then throw a **INVALID_ACCESS_ERR** exception.
8. If the element has a maximum (page 344), and the *value* is greater than that maximum (page 344), then throw a **INVALID_ACCESS_ERR** exception.
9. Set the value (page 362) of the element to *value as string*.

The **list** DOM attribute must return the current suggestions source element (page 342), if any, or null otherwise.

The **selectedOption** DOM attribute must return the first option element, in tree order (page 24), to be a child of the suggestions source element (page 342) and whose value (page 355) matches the input element's value (page 362), if any. If there is no suggestions source element (page 342), or if it contains no matching option element, then the **selectedOption** attribute must return null.

4.10.4.4 Common event behaviors

When the **input** event applies, the user agent must queue a task (page 429) to fire a simple event (page 436) called **input** at the input element any time the user causes the element's value (page 362) to change. User agents may wait for a suitable break in the user's interaction before queuing the task; for example, a user agent could wait for the user to have not hit a key for 100ms, so as to only fire the event when the user pauses, instead of continuously for each keystroke.

When the **change** event applies, if the element does not have an activation behavior (page 96) defined but uses a user interface that involves an explicit commit action, then the user

agent must queue a task (page 429) to fire a simple event (page 436) called change at the input element any time the user commits a change to the element's value (page 362) or list of selected files (page 337).

An example of a user interface with a commit action would be a File Upload (page 337) control that consists of a single button that brings up a file selection dialog: when the dialog is closed, if that the file selection (page 337) changed as a result, then the user has committed a new file selection (page 337).

Another example of a user interface with a commit action would be a Date (page 328) control that allows both text-based user input and user selection from a drop-down calendar: while text input might not have an explicit commit step, selecting a date from the drop down calendar and then dismissing the drop down would be a commit action.

Note: *In addition, when the change event applies, change events can also be fired as part of the element's action behavior and as part of the unfocusing steps (page 516).*

The task source (page 429) for these task is the user interaction task source (page 430).

4.10.5 The button element

Categories

Phrasing content (page 94).
Interactive content (page 95).
Listed (page 314), labelable (page 314), and submittable (page 314) form-associated element (page 313).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

action
autofocus
disabled
enctype
form
method
name
target
type
value

DOM interface:

```
interface HTMLButtonElement : HTMLElement {  
    attribute DOMString action;  
    attribute boolean autofocus;  
    attribute boolean disabled;
```

```

        attribute DOMString enctype;
readonly attribute HTMLFormElement form;
        attribute DOMString method;
        attribute DOMString name;
        attribute DOMString target;
        attribute DOMString type;
        attribute DOMString value;

readonly attribute boolean willValidate;
readonly attribute ValidityState validity;
readonly attribute DOMString validationMessage;
boolean checkValidity();
void setCustomValidity(in DOMString error);

readonly attribute NodeList labels;
};

```

The button element represents a button. If the element is not disabled (page 362), then the user agent should allow the user to activate the button.

The **type** attribute controls the behavior of the button when it is activated. It is an enumerated attribute (page 51). The following table lists the keywords and states for the attribute — the keywords in the left column map to the states in the cell in the second column on the same row as the keyword.

Keyword	State	Brief description
submit	Submit Button (page 349)	Submits the form.
reset	Reset Button (page 349)	Resets the form.
button	Button (page 349)	Does nothing.

The *missing value default* is the Submit Button (page 349) state.

If the element is not disabled (page 362), the activation behavior (page 96) of the button element is to run the steps defined in the following list for the current state of the element's type attribute.

Submit Button

If the element has a form owner (page 361), the element must submit (page 368) the form owner (page 361) from the button element.

Reset Button

If the element has a form owner (page 361), the element must reset (page 376) the form owner (page 361).

Button

Do nothing.

The form attribute is used to explicitly associate the button element with its form owner (page 361). The name attribute represents the element's name. The disabled attribute is used to make the control non-interactive and to prevent its value from being submitted. The

`autofocus` attribute controls focus. The `action`, `enctype`, `method`, and `target` attributes are attributes for form submission (page 363).

The `value` attribute gives the element's value for the purposes of form submission. The `value` attribute must not be present unless the `form` attribute is present. The element's value (page 362) is the value of the element's `value` attribute, if there is one, or the empty string otherwise.

The `value` and `type` DOM attributes must reflect (page 67) the respective content attributes of the same name.

The `willValidate`, `validity`, and `validationMessage` attributes, and the `checkValidity()` and `setCustomValidity()` methods, are part of the constraint validation API (page 366). The `labels` attribute provides a list of the element's labels.

4.10.6 The select element

Categories

Phrasing content (page 94).

Interactive content (page 95).

Listed (page 314), labelable (page 314), submittable (page 314), and resettable (page 314) form-associated element (page 313).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Zero or more `option` or `optgroup` elements.

Element-specific attributes:

`autofocus`
`disabled`
`form`
`multiple`
`name`
`size`

DOM interface:

```
interface HTMLSelectElement : HTMLElement {  
    attribute boolean autofocus;  
    attribute boolean disabled;  
    readonly attribute HTMLFormElement form;  
    attribute boolean multiple;  
    attribute DOMString name;  
    attribute boolean size;  
  
    readonly attribute DOMString type;  
  
    readonly attribute HTMLOptionsCollection options;  
    attribute unsigned long length;  
    [IndexGetter] HTMLElement XXX9(in unsigned long index);
```

```
void add(in HTMLElement element, in HTMLElement before);
void add(in HTMLElement element, in long before);
void remove(in long index);

readonly attribute HTMLCollection selectedOptions;
attribute long selectedIndex;
attribute DOMString value;

readonly attribute boolean willValidate;
readonly attribute ValidityState validity;
readonly attribute DOMString validationMessage;
boolean checkValidity();
void setCustomValidity(in DOMString error);

readonly attribute NodeList labels;
};
```

The `select` element represents a control for selecting amongst a set of options.

The **multiple** attribute is a boolean attribute (page 32). If the attribute is present, then the `select` element represents a control for selecting zero or more options from the list of options (page 351). If the attribute is absent, then the `select` element represents a control for selecting a single option from the list of options (page 351).

The **list of options** for a `select` element consists of all the `option` element children of the `select` element, and all the `option` element children of all the `optgroup` element children of the `select` element, in tree order (page 24).

The **size** attribute gives the number of options to show to the user. The `size` attribute, if specified, must have a value that is a valid non-negative integer (page 33) greater than zero. If the `multiple` attribute is present, then the `size` attribute's default value is 4. If the `multiple` attribute is absent, then the `size` attribute's default value is 1.

If the `multiple` attribute is absent, and the element is not disabled (page 362), then the user agent should allow the user to pick an `option` element in its list of options (page 351) that is itself not disabled (page 355). Upon this `option` element being picked (either through a click, or through unfocusing the element after changing its value, or through any other mechanism), and before the relevant user interaction event is queued (e.g. before the `click` event), the user agent must set the selectedness (page 355) of the picked `option` element to true and then queue a task (page 429) to fire a simple event (page 436) called `change` at the `select` element, using the user interaction task source (page 430) as the task source.

If the `multiple` attribute is absent, whenever an `option` element in the `select` element's list of options (page 351) has its selectedness (page 355) set to true, and whenever an `option` element with its selectedness (page 355) set to true is added to the `select` element's list of options (page 351), the user agent must set the selectedness (page 355) of all the other `option` element in its list of options (page 351) to false.

If the `multiple` attribute is absent, whenever there are no `option` elements in the `select` element's list of options (page 351) that have their selectedness (page 355) set to true, the

user agent must set the selectedness (page 355) of the first option element in the list of options (page 351) in tree order (page 24) that is not disabled (page 355), if any, to true.

If the `multiple` attribute is present, and the element is not disabled (page 362), then the user agent should allow the user to toggle the selectedness (page 355) of the option elements in its list of options (page 351) that are themselves not disabled (page 355). Upon the selectedness (page 355) of one or more option elements being changed by the user, and before the relevant user interaction event is queued (e.g. before a related click event), the user agent must queue a task (page 429) to fire a simple event (page 436) called `change` at the `select` element, using the user interaction task source (page 430) as the task source.

The reset algorithm (page 376) for `select` elements is to go through all the option elements in the element's list of options (page 351), and set their selectedness (page 355) to true if the option element has a `selected` attribute, and false otherwise.

The `form` attribute is used to explicitly associate the `select` element with its form owner (page 361). The `name` attribute represents the element's name. The `disabled` attribute is used to make the control non-interactive and to prevent its value from being submitted. The `autofocus` attribute controls focus.

The `type` DOM attribute, on getting, must return the string "select-one" if the `multiple` attribute is absent, and the string "select-multiple" if the `multiple` attribute is present.

The `options` DOM attribute must return an `HTMLOptionsCollection` rooted at the `select` node, whose filter matches the elements in the list of options (page 351).

The `length` DOM attribute, on getting and setting, must act like the `length` attribute on the `HTMLOptionsCollection` object returned by the `options` attribute. Similarly, the `add()` and `remove()` methods must act like their namesake methods on that same `HTMLOptionsCollection` object.

The `selectedOptions` DOM attribute must return an `HTMLCollection` rooted at the `select` node, whose filter matches the elements in the list of options (page 351) that have their selectedness (page 355) set to true.

The `selectedIndex` DOM attribute, on getting, must return the index (page 355) of the first option element in the list of options (page 351) in tree order (page 24) that has its selectedness (page 355) set to true, if any. If there isn't one, then it must return -1.

On setting, the `selectedIndex` attribute must set the selectedness (page 355) of all the option elements in the list of options (page 351) to false, and then the option element in the list of options (page 351) whose index (page 355) is the given new value, if any, must have its selectedness (page 355) set to true.

The `value` DOM attribute, on getting, must return the value (page 355) of the first option element in the list of options (page 351) in tree order (page 24) that has its selectedness (page 355) set to true, if any. If there isn't one, then it must return the empty string.

On setting, the `value` attribute must set the selectedness (page 355) of all the option elements in the list of options (page 351) to false, and then first the option element in the list of options (page 351), in tree order (page 24), whose value (page 355) is equal to the given new value, if any, must have its selectedness (page 355) set to true.

The **multiple** and **size** DOM attributes must reflect (page 67) the respective content attributes of the same name. The **size** DOM attribute limited to only positive non-zero numbers (page 68).

The **willValidate**, **validity**, and **validationMessage** attributes, and the **checkValidity()** and **setCustomValidity()** methods, are part of the constraint validation API (page 366). The **labels** attribute provides a list of the element's labels.

4.10.7 The **datalist** element

Categories

Phrasing content (page 94).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Either: phrasing content (page 94).
Or: Zero or more option elements.

Element-specific attributes:

None.

DOM interface:

```
interface HTMLDataListElement : HTMLElement {  
    readonly attribute HTMLCollection options;  
};
```

The **datalist** element represents a set of option elements that represent predefined options for other controls. The contents of the element represents fallback content for legacy user agents, intermixed with option elements that represent the predefined options. In the rendering, the **datalist** element represents nothing and it, along with its children, should be hidden.

The **datalist** element is hooked up to an **input** element using the **list** attribute on the **input** element.

The **options** DOM attribute must return an **HTMLCollection** rooted at the **datalist** node, whose filter matches option elements.

Constraint validation: If an element has a **datalist** element ancestor, it is barred from constraint validation (page 364).

4.10.8 The **optgroup** element

Categories

None.

Contexts in which this element may be used:

As a child of a **select** element.

Content model:

Zero or more option elements.

Element-specific attributes:

disabled
label

DOM interface:

```
interface HTMLOptGroupElement : HTMLElement {  
    attribute boolean disabled;  
    attribute DOMString label;  
};
```

The optgroup element represents a group of option elements with a common label.

The element's group of option elements consists of the option elements that are children of the optgroup element.

When showing option elements in select elements, user agents should show the option elements of such groups as being related to each other and separate from other option elements.

The **disabled** attribute is a boolean attribute (page 32) and can be used to disable (page 355) a group of option elements together.

The **label** attribute must be specified. Its value gives the name of the group, for the purposes of the user interface. User agents should use this attribute's value when labelling the group of option elements in a select element.

The **disabled** and **label** attributes must reflect (page 67) the respective content attributes of the same name.

4.10.9 The option element

Categories

None.

Contexts in which this element may be used:

- As a child of a select element.
- As a child of a datalist element.
- As a child of an optgroup element.

Content model:

Text.

Element-specific attributes:

disabled
label
selected
value

DOM interface:

```
[Constructor(),
 Constructor(in DOMString text),
 Constructor(in DOMString text, in DOMString value),
 Constructor(in DOMString text, in DOMString value, in boolean
 defaultSelected)]
interface HTMLOptionElement : HTMLElement {
    attribute boolean disabled;
    readonly attribute HTMLFormElement form;
    attribute DOMString label;
    attribute boolean defaultSelected;
    attribute boolean selected;
    attribute DOMString value;

    readonly attribute DOMString text;
    readonly attribute long index;
};
```

The option element represents an option in a select element or as part of a list of suggestions in a datalist element.

The **disabled** attribute is a boolean attribute (page 32). An option element is **disabled** if its disabled is present or if it is a child of an optgroup element whose disabled attribute is present.

The **label** attribute provides a label for element. The **label** of an option element is the value of the label attribute, if there is one, or the textContent of the element, if there isn't.

The **value** attribute provides a value for element. The **value** of an option element is the value of the value attribute, if there is one, or the textContent of the element, if there isn't.

The **selected** attribute represents the default selectedness (page 355) of the element.

The **selectedness** of an option element is a boolean state, initially false. If the element is disabled (page 355), then the element's selectedness (page 355) is always false and cannot be set to true. When the element is created, its selectedness (page 355) must be set to true if the element has a selected attribute. Whenever an option element's selected attribute is added, its selectedness (page 355) must be set to true.

An option element's **index** is the number of option element that are in the same list of options (page 351) but that come before it in tree order (page 24). If the option element is not in a list of options (page 351), then the option element's index (page 355) is zero.

The **disabled**, **label**, and **value** DOM attributes must reflect (page 67) the respective content attributes of the same name. The **defaultSelected** DOM attribute must reflect (page 67) the selected content attribute.

The **selected** DOM attribute must return true if the element's selectedness (page 355) is true, and false otherwise.

The **index** DOM attribute must return the element's index (page 355).

The **text** DOM attribute must return the same value as the `textContent` DOM attribute on the element.

The **form** DOM attribute's behavior depends on whether the option element is in a `select` element or not. If the option has a `select` element as its parent, or has a `colgroup` element as its parent and that `colgroup` element has a `select` element as its parent, then the `form` DOM attribute must return the same value as the `form` DOM attribute on that `select` element. Otherwise, it must return null.

Four constructors are provided for creating `HTMLOptionElement` objects (in addition to the factory methods from DOM Core such as `createElement()`): `Option()`, `Option(text)`, `Option(text, value)`, and `Option(text, value, defaultSelected)`. When invoked as constructors, these must return a new `HTMLOptionElement` object (a new option element). If the `text` argument is present, the new object must have a text node with the value of that argument as its data as its only child. If the `value` argument is present, the new object must have a `value` attribute set with the value of the argument as its value. If the `defaultSelected` argument is present and true, the new object must have a `selected` attribute set with no value.

4.10.10 The `textarea` element

Categories

Phrasing content (page 94).
Interactive content (page 95).
Listed (page 314), labelable (page 314), submittable (page 314), and resettable (page 314) form-associated element (page 313).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Text.

Element-specific attributes:

`autofocus`
`cols`
`disabled`
`form`
`maxlength`
`name`
`readonly`
`required`
`rows`
`wrap`

DOM interface:

```
interface HTMLTextAreaElement : HTMLElement {  
    attribute boolean autofocus;  
    attribute unsigned long cols;  
    attribute boolean disabled;  
    readonly attribute HTMLFormElement form;
```

```

        attribute long maxLength;
        attribute DOMString name;
        attribute boolean readOnly;
        attribute boolean required;
        attribute unsigned long rows;
        attribute DOMString wrap;

        readonly attribute DOMString type;
        attribute DOMString defaultValue;
        attribute DOMString value;

        readonly attribute boolean willValidate;
        readonly attribute ValidityState validity;
        readonly attribute DOMString validationMessage;
        boolean checkValidity();
        void setCustomValidity(in DOMString error);

        readonly attribute NodeList labels;
    };

```

The `textarea` element represents a multiline plain text edit control for the element's **raw value**. The contents of the control represent the control's default value.

The raw value (page 357) of a `textarea` control must be initially the empty string.

The **readonly** attribute is a boolean attribute (page 32) used to control whether the text can be edited by the user or not.

Constraint validation: If the `readonly` attribute is specified on a `textarea` element, the element is barred from constraint validation (page 364).

A `textarea` element is **mutable** if it is neither disabled (page 362) nor has a `readonly` attribute specified.

When a `textarea` is mutable (page 357), its raw value (page 357) should be editable by the user.

A `textarea` element has a **dirty value flag**, which must be initially set to false, and must be set to true whenever the user interacts with the control in a way that changes the raw value (page 357).

When the `textarea` element's `textContent` DOM attribute changes value, if the element's dirty value flag (page 357) is false, then the element's raw value (page 357) must be set to the value of the element's `textContent` DOM attribute.

The reset algorithm (page 376) for `textarea` elements is to set the element's value (page 357) to the value of the element's `textContent` DOM attribute.

The **cols** attribute specifies the expected maximum number of characters per line. If the `cols` attribute is specified, its value must be a valid non-negative integer (page 33) greater than zero. If applying the rules for parsing non-negative integers (page 33) to the attribute's

value results in a number greater than zero, then the element's **character width** is that value; otherwise, it is 20.

The user agent may use the `textarea` element's character width (page 358) as a hint to the user as to how many characters the server prefers per line (e.g. for visual user agents by making the width of the control be that many characters). In visual renderings, the user agent should wrap the user's input in the rendering so that each line is no wider than this number of characters.

The **rows** attribute specifies the number of lines to show. If the `rows` attribute is specified, its value must be a valid non-negative integer (page 33) greater than zero. If applying the rules for parsing non-negative integers (page 33) to the attribute's value results in a number greater than zero, then the element's **character height** is that value; otherwise, it is 2.

Visual user agents should set the height of the control to the number of lines given by character height (page 358).

The **wrap** attribute is an enumerated attribute (page 51) with two keywords and states: the **soft** keyword which maps to the **Soft** state, and the **hard** keyword which maps to the **Hard** state. The *missing value default* is the Soft (page 358) state.

If the element's `wrap` attribute is in the Hard (page 358) state, the `cols` attribute must be specified.

The element's value (page 362) is defined to be the element's raw value (page 357) with the following transformation applied:

1. Replace every occurrence of a U+000D CARRIAGE RETURN (CR) character not followed by a U+000A LINE FEED (LF) character, and every occurrence of a U+000A LINE FEED (LF) character not proceeded by a U+000D CARRIAGE RETURN (CR) character, by a two-character string consisting of a U+000D CARRIAGE RETURN - U+000A LINE FEED (CRLF) character pair.
2. If the element's `wrap` attribute is in the Hard (page 358) state, insert U+000D CARRIAGE RETURN - U+000A LINE FEED (CRLF) character pairs into the string using a UA-defined algorithm so that each line so that each line has no more than character width (page 358) characters. The the purposes of this requirement, lines are delimited by the start of the string, the end of the string, and U+000D CARRIAGE RETURN - U+000A LINE FEED (CRLF) character pairs.

The **maxlength** attribute is a form control `maxlength` attribute (page 363) controlled by the `textarea` element's dirty value flag (page 357).

If the `textarea` element has a maximum allowed value length (page 363), then the element's children must be such that the codepoint length (page 32) of the value of the element's `textContent` DOM attribute is equal to or less than the element's maximum allowed value length (page 363).

The **required** attribute is a boolean attribute (page 32).

Constraint validation: If the element has its `required` attribute specified, and the element is mutable (page 357), and the element's value (page 362) is the empty string, then the element is suffering from being missing (page 365).

The `form` attribute is used to explicitly associate the `textarea` element with its form owner (page 361). The `name` attribute represents the element's name. The `disabled` attribute is used to make the control non-interactive and to prevent its value from being submitted. The `autofocus` attribute controls focus.

The `cols`, `required`, `rows`, and `wrap` attributes must reflect (page 67) the respective content attributes of the same name. The `cols` and `rows` attributes are limited to only positive non-zero numbers (page 68). The `maxLength` DOM attribute must reflect (page 67) the `maxlength` content attribute. The `readonly` DOM attribute must reflect (page 67) the `readonly` content attribute.

The `type` DOM attribute must return the value "textarea".

The `defaultValue` DOM attribute must act like the element's `textContent` DOM attribute.

The `value` attribute must, on getting, return the element's value (page 362); on setting, it must set the element's value (page 362) to the new value.

The `willValidate`, `validity`, and `validationMessage` attributes, and the `checkValidity()` and `setCustomValidity()` methods, are part of the constraint validation API (page 366). The `labels` attribute provides a list of the element's labels.

4.10.11 The output element

Categories

Phrasing content (page 94).

Listed (page 314) and resettable (page 314) form-associated element (page 313).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94).

Element-specific attributes:

`for`
`form`
`name`

DOM interface:

```
interface HTMLOutputElement : HTMLElement {  
    attribute DOMString htmlFor;  
    readonly attribute HTMLFormElement form;  
    attribute DOMString name;  
  
    readonly attribute DOMString type;  
    attribute DOMString defaultValue;  
    attribute DOMString value;  
  
    readonly attribute boolean willValidate;  
    readonly attribute ValidityState validity;  
    readonly attribute DOMString validationMessage;
```

```
boolean checkValidity();
void setCustomValidity(in DOMString error);
};
```

The output element represents the result of a calculation.

The **for** content attribute allows an explicit relationship to be made between the result of a calculation and the elements that represent the values that went into the calculation or that otherwise influenced the calculation. The **for** attribute, if specified, must contain a string consisting of an unordered set of unique space-separated tokens (page 49), each of which must have the value of an ID of an element in the same Document.

The **form** attribute is used to explicitly associate the output element with its form owner (page 361). The **name** attribute represents the element's name.

The element has a **value mode flag** which is either *value* or *default*. Initially the value mode flag (page 360) must be set to *default*.

When the to value mode flag (page 360) is in mode *default*, the contents of the element represent both the value of the element and its default value. When the value mode flag (page 360) is in mode *value*, the contents of the element represent the value of the element only, and the default value is only accessible using the **defaultValue** DOM attribute.

The element also has a **defaultValue**. Initially, the default value (page 360) must be the empty string.

Whenever the element's descendants are changed in any way, if the value mode flag (page 360) is in mode *default*, the element's default value (page 360) must be set to the value of the element's **textContent** DOM attribute.

The reset algorithm (page 376) for textarea elements is to set the element's **textContent** DOM attribute to the value of the element's **defaultValue** DOM attribute (thus replacing the element's child nodes), and then to set the element's value mode flag (page 360) to *default*.

The **value** DOM attribute must act like the element's **textContent** DOM attribute, except that on setting, in addition, before the child nodes are changed, the element's value mode flag (page 360) must be set to *value*.

The **defaultValue** DOM attribute, on getting, must return the element's default value (page 360). On setting, the attribute must set the element's default value (page 360), and, if the element's value mode flag (page 360) is in the mode *default*, set the element's **textContent** DOM attribute as well.

The **type** attribute must return the string "output".

The **htmlFor** DOM attribute must reflect (page 67) the **for** content attribute.

The **willValidate**, **validity**, and **validationMessage** attributes, and the **checkValidity()** and **setCustomValidity()** methods, are part of the constraint validation API (page 366).

Constraint validation: output elements are always barred from constraint validation (page 364).

4.10.12 Association of controls and forms

A form-associated element (page 313) can have a relationship with a form element, which is called the element's **form owner**. If a form-associated element (page 313) is not associated with a form element, its form owner (page 361) is said to be null.

A form-associated element (page 313) is, by default, associated with its nearest ancestor form element (as described below), but may have a **form** attribute specified to override this.

If a form-associated element (page 313) has a **form** attribute specified, then its value must be the ID of a form element in the element's owner Document.

When a form-associated element (page 313) is created, its form owner (page 361) must be initialized to null (no owner).

When a form-associated element (page 313) is to be **associated** with a form, its form owner (page 361) must be set to that form.

When a form-associated element (page 313)'s ancestor chain changes, e.g. because it or one of its ancestors was inserted (page 24) or removed from a Document, then the user agent must reset the form owner (page 361) of that element.

When a form-associated element (page 313)'s **form** attribute is added, removed, or has its value changed, then the user agent must reset the form owner (page 361) of that element.

When a form-associated element (page 313) has a **form** attribute and the ID of any of the form elements in the Document changes, then the user agent must reset the form owner (page 361) of that form-associated element (page 313).

When the user agent is to **reset the form owner** of a form-associated element (page 313), it must run the following steps:

1. If the element's form owner (page 361) is not null, and the element's **form** content attribute is not present, and the element's form owner (page 361) is one of the ancestors of the element after the change to the ancestor chain, then do nothing, and abort these steps.
2. Let the element's form owner (page 361) be null.
3. If the element has a **form** content attribute, then run these substeps:
 1. If the first element in the Document to have an ID that is equal to the element's **form** content attribute's value is a form element, then associate (page 361) the form-associated element (page 313) with that form element.
 2. Abort the "reset the form owner" steps.
4. Otherwise, if the form-associated element (page 313) in question has an ancestor form element, then associate (page 361) the form-associated element (page 313) with the nearest such ancestor form element.
5. Otherwise, the element is left unassociated.

Form-associated elements (page 313) have a **form** DOM attribute, which, on getting, must return the element's form owner (page 361), or null if there isn't one.

Constraint validation: If an element has no form owner (page 361), it is barred from constraint validation (page 364).

4.10.13 Attributes common to form controls

4.10.13.1 Naming form controls

The **name** content attribute gives the name of the form control, as used in form submission and in the `form` element's `elements` object. If the attribute is specified, its value must not be the empty string.

Constraint validation: If an element does not have a `name` attribute specified, or its `name` attribute's value is the empty string, then it is barred from constraint validation (page 364).

The `name` DOM attribute must reflect (page 67) the `name` content attribute.

4.10.13.2 Enabling and disabling form controls

The **disabled** content attribute is a boolean attribute (page 32).

A form control is **disabled** if its `disabled` attribute is set, or if it is a descendant of a `fieldset` element whose `disabled` attribute is set.

Constraint validation: If an element is disabled (page 362), it is barred from constraint validation (page 364).

The `disabled` DOM attribute must reflect (page 67) the `disabled` content attribute.

4.10.13.3 A form control's value

Form controls have a **value** and a **checkedness**. (The latter is only used by `input` elements.) These are used to describe how the user interacts with the control.

4.10.13.4 Autofocusing a form control

The **autofocus** content attribute allows the user to indicate that a control is to be focused as soon as the page is loaded, allowing the user to just start typing without having to manually focus the main control.

The `autofocus` attribute is a boolean attribute (page 32).

There must not be more than one element in the document with the `autofocus` attribute specified.

Whenever an element with the `autofocus` attribute specified is inserted into a document (page 24), the user agent should queue a task (page 429) that checks to see if the element is focusable (page 516), and if so, runs the focusing steps (page 516) for that element. User agents may also change the scrolling position of the document, or perform some other action that brings the element to the user's attention. The task source (page 429) for this task is the DOM manipulation task source (page 430).

User agents may ignore this attribute if the user has indicated (for example, by starting to type in a form control) that he does not wish focus to be changed.

Note: Focusing the control does not imply that the user agent must focus the browser window if it has lost focus.

The **autofocus** DOM attribute must reflect (page 67) the content attribute of the same name.

In the following snippet, the text control would be focused when the document was loaded.

```
<input maxlength="256" name="q" value="" autofocus>
<input type="submit" value="Search">
```

4.10.13.5 Limiting user input length

A **form control maxlen attribute**, controlled by a *dirty value flag* declares a limit on the number of characters a user can input.

If an element has its form control **maxlength attribute** (page 363) specified, the attribute's value must be a valid non-negative integer (page 33). If the attribute is specified and applying the rules for parsing non-negative integers (page 33) to its value results in a number, then that number is the element's **maximum allowed value length**. If the attribute is omitted or parsing its value results in an error, then there is no maximum allowed value length (page 363).

Constraint validation: If an element has a maximum allowed value length (page 363), and its *dirty value flag* is false, and the codepoint length (page 32) of the element's value (page 362) is greater than the element's maximum allowed value length (page 363), then the element is suffering from being too long (page 365).

User agents may prevent the user from causing the element's value (page 362) to be set to a value whose codepoint length (page 32) is greater than the element's maximum allowed value length (page 363).

4.10.13.6 Form submission

Attributes for form submission can be specified both on **form** elements and on elements that represent buttons that submit forms, e.g. an **input** element whose **type** attribute is in the Submit Button (page 338) state. The attributes on the buttons, when omitted, default to the values given on the corresponding the **form** element.

The **action** content attribute, if specified, must have a value that is a valid URL (page 52).

The **action** of an element is the value of the element's **action** attribute, if it has one, or the value of its form owner (page 361)'s **action** attribute, if it has one, or else the empty string.

The **method** content attribute is an enumerated attribute (page 51) with the following keywords and states:

- The keyword **GET**, mapping to the state **GET**, indicating the HTTP GET method.
- The keyword **POST**, mapping to the state **POST**, indicating the HTTP POST method.

- The keyword **PUT**, mapping to the state **PUT**, indicating the HTTP PUT method.
- The keyword **DELETE**, mapping to the state **DELETE**, indicating the HTTP DELETE method.

The *missing value default* is the GET (page 363) state.

The **method** of an element is one of those four states. If the element has a `method` attribute, then the element's method (page 364) is that attribute's state; otherwise, it is the form owner (page 361)'s `method` attribute's state.

The **enctype** content attribute is an enumerated attribute (page 51) with the following keywords and states:

- The "**application/x-www-form-urlencoded**" keyword and corresponding state.
- The "**multipart/form-data**" keyword and corresponding state.
- The "**text/plain**" keyword and corresponding state.

The *missing value default* is the `application/x-www-form-urlencoded` state.

The **enctype** of an element is one of those four states. If the element has a `enctype` attribute, then the element's `enctype` (page 364) is that attribute's state; otherwise, it is the form owner (page 361)'s `enctype` attribute's state.

The **target** content attribute, if present, must be a valid browsing context name or keyword (page 418).

The **target** of an element is the value of the element's `target` attribute, if it has one; or the value of its form owner (page 361)'s `target` attribute, if it has one; or, if one of the child nodes of the head element (page 80) is a base element with a `target` attribute, then the the value of the `target` attribute of the first such base element; or, if there is no such element, the empty string.

The **action**, **method**, **enctype**, and **target** DOM attributes must reflect (page 67) the respective content attributes of the same name.

4.10.14 Constraints

4.10.14.1 Definitions

A listed form-associated element (page 314) is a **candidate for constraint validation** unless a condition has **barred the element from constraint validation**. (For example, an element is barred from constraint validation (page 364) if it is an output or fieldset element.)

An element can have a **custom validity error message** defined. Initially, an element must have its custom validity error message (page 364) set to the empty string. When its value is not the empty string, the element is suffering from a custom error (page 365). It can be set using the `setCustomValidity()` method. The user agent should use the custom validity error message (page 364) when alerting the user to the problem with the control.

An element can be constrained in various ways. The following is the list of **validity states** that a form control can be in, making the control invalid for the purposes of constraint validation:

Suffering from being missing

When a control has no value (page 362) but has a required attribute (`input required`, `textarea required`).

Suffering from a type mismatch

When a control that allows arbitrary user input has a value (page 362) that is not in the correct syntax (E-mail (page 325), URL (page 325)).

Suffering from a pattern mismatch

When a control has a value (page 362) that doesn't satisfy the `pattern` attribute.

Suffering from being too long

When a control has a value (page 362) that is too long for the form control `maxlength` attribute (page 363) (`input maxlength`, `textarea maxlength`).

Suffering from an underflow

When a control has a value (page 362) that is too low for the `min` attribute.

Suffering from an overflow

When a control has a value (page 362) that is too high for the `max` attribute.

Suffering from a step mismatch

When a control has a value (page 362) that doesn't fit the rules given by the `step` attribute.

Suffering from a custom error

When a control's custom validity error message (page 364) (as set by the element's `setCustomValidity()` method) is not the empty string.

An element **satisfies its constraints** if it is not suffering from any of the above validity states (page 365).

4.10.14.2 Constraint validation

When the user agent is required to **statically validate the constraints** of form element `form`, it must run the following steps, which return either a *positive* result (all the controls in the form are valid) or a *negative* result (there are invalid controls) along with a (possibly empty) list of elements that are invalid and for which no script has claimed responsibility:

1. Let `controls` be a list of all the submittable (page 314) elements whose form owner (page 361) is `form`, in tree order (page 24).
2. Let `invalid controls` be an initially empty list of elements.
3. For each element `field` in `controls`, in tree order (page 24), run the following substeps:
 1. If `field` is not a candidate for constraint validation (page 364), then move on to the next element.

2. Otherwise, if *field* satisfies its constraints (page 365), then move on to the next element.
3. Otherwise, add *field* to *invalid controls*.
4. If *invalid controls* is empty, then return a *positive* result and abort these steps.
5. Let *unhandled invalid controls* be an initially empty list of elements.
6. For each element *field* in *invalid controls*, if any, in tree order (page 24), run the following substeps:
 1. Fire a simple event (page 436) named *invalid* at *field*.
 2. If the event was not canceled, then add *field* to *unhandled invalid controls*.
7. Return a *negative* result with the list of elements in the *unhandled invalid controls* list.

If a user agent is to **interactively validate the constraints** of form element *form*, then the user agent must run the following steps:

1. Statically validate the constraints (page 365) of *form*, and let *unhandled invalid controls* be the list of elements returned if the result was *negative*.
2. If the result was *positive*, then return that result and abort these steps.
3. Report the problems with the constraints of at least one of the elements given in *unhandled invalid controls* to the user. User agents may focus one of those elements in the process, by running the focusing steps (page 516) for that element, and may change the scrolling position of the document, or perform some other action that brings the element to the user's attention. User agents may report more than one constraint violation. User agents may coalesce related constraint violation reports if appropriate (e.g. if multiple radio buttons in a set are marked as required, only one error need be reported). If one of the controls is not visible to the user (e.g. it has the `hidden` attribute set) then user agents may report a script error.
4. Return a *negative* result.

4.10.14.3 The constraint validation API

The `willValidate` attribute must return true if an element is a candidate for constraint validation (page 364), and false otherwise (i.e. false if any conditions are barring it from constraint validation (page 364)).

The `setCustomValidity(message)`, when invoked, must set the custom validity error message (page 364) to the value of the given *message* argument.

The `validity` attribute must return a `ValidityState` object that represents the validity states (page 365) of the element. This object is live, and the same object must be returned each time the element's `validity` attribute is retrieved.

```
interface ValidityState {
  readonly attribute boolean valueMissing;
  readonly attribute boolean typeMismatch;
```

```

readonly attribute boolean patternMismatch;
readonly attribute boolean tooLong;
readonly attribute boolean rangeUnderflow;
readonly attribute boolean rangeOverflow;
readonly attribute boolean stepMismatch;
readonly attribute boolean customError;
readonly attribute boolean valid;
};

```

A `ValidityState` object has the following attributes. On getting, they must return true if the corresponding condition given in the following list is true, and false otherwise.

valueMissing

The control is suffering from being missing (page 365).

typeMismatch

The control is suffering from a type mismatch (page 365).

patternMismatch

The control is suffering from a pattern mismatch (page 365).

tooLong

The control is suffering from being too long (page 365).

rangeUnderflow

The control is suffering from an underflow (page 365).

rangeOverflow

The control is suffering from an overflow (page 365).

stepMismatch

The control is suffering from a step mismatch (page 365).

customError

The control is suffering from a custom error (page 365).

valid

None of the other conditions are true.

When the `checkValidity()` method is invoked, if the element is a candidate for constraint validation (page 364) and does not satisfy its constraints (page 365), the user agent must fire a simple event (page 436) named `invalid` at the element and return false. Otherwise, it must only return true without doing anything else.

The `validationMessage` attribute must return the empty string if the element is not a candidate for constraint validation (page 364) or if it is one but it satisfies its constraints (page 365); otherwise, it must return a suitably localised message that the user agent would show the user if this were the only form with a validity constraint problem. If the element is suffering from a custom error (page 365), then the custom validity error message (page 364) should be present in the return value.

4.10.15 Form submission

When a form *form* is **submitted** from an element *submitter* (typically a button), the user agent must run the following steps:

1. If *form* is in a Document that has no associated browsing context (page 414) or whose browsing context (page 414) has its sandboxed forms browsing context flag (page 219) set, then abort these steps without doing anything.
2. If the *submitter* is anything but a form element, then interactively validate the constraints (page 366) of *form* and examine the result: if the result is negative (the constraint validation concluded that there were invalid fields and probably informed the user of this) then abort these steps.
3. If the *submitter* is anything but a form element, then fire a simple event (page 436) that bubbles, named *submit*, at *form*. If the event's default action is prevented (i.e. if the event is canceled) then abort these steps. Otherwise, continue (effectively the default action is to perform the submission).
4. Let *controls* be a list of all the submittable (page 314) elements whose form owner (page 361) is *form*, in tree order (page 24).
5. Let the *form data set* be a list of name-value-type tuples, initially empty.
6. **Constructing the form data set.** For each element *field* in *controls*, in tree order (page 24), run the following substeps:
 1. If any of the following conditions are met, then skip these substeps for this element:
 - The *field* element has a *datalist* element ancestor.
 - The *field* element is disabled (page 362).
 - The *field* element is a button (page 314) but it is not *submitter*.
 - The *field* element is an *input* element whose *type* attribute is in the Checkbox (page 335) state and whose checkedness (page 362) is false.
 - The *field* element is an *input* element whose *type* attribute is in the Radio Button (page 336) state and whose checkedness (page 362) is false.
 - The *field* element is an *input* element whose *type* attribute is in the File Upload (page 337) state but the control does not have any files selected.
 - Otherwise, process *field* as follows:
 2. Let *type* be the value of the *type* DOM attribute of *field*.
 3. If the *field* element is an *input* element whose *type* attribute is in the Image Button (page 338) state, then run these further nested substeps:

1. If the *field* element has an `name` attribute specified and its value is not the empty string, let *name* be that value followed by a single U+002E FULL STOP (.) character. Otherwise, let *name* be the empty string.
 2. Let *name_x* be the string consisting of the concatenation of *name* and a single U+0078 LATIN SMALL LETTER X (x) character.
 3. Let *name_y* be the string consisting of the concatenation of *name* and a single U+0079 LATIN SMALL LETTER Y (y) character.
 4. The *field* element is *submitter*, and before this algorithm was invoked the user indicated a coordinate (page 340). Let *x* be the *x*-component of the coordinate selected by the user, and let *y* be the *y*-component of the coordinate selected by the user.
 5. Append an entry in the *form data set* with the name *name_x*, the value *x*, and the type *type*.
 6. Append an entry in the *form data set* with the name *name_y* and the value *y*, and the type *type*.
 7. Skip the remaining substeps for this element: if there are any more elements in *controls*, return to the top of the constructing the form data set (page 368) step, otherwise, jump to the next step in the overall form submission algorithm.
4. If the *field* element does not have a `name` attribute specified, or its `name` attribute's value is the empty string, skip these substeps for this element: if there are any more elements in *controls*, return to the top of the constructing the form data set (page 368) step, otherwise, jump to the next step in the overall form submission algorithm.
 5. Let *name* be the value of the *field* element's `name` attribute.
 6. If the *field* element is a *select* element, then for each *option* element in the *select* element whose `selectedness` (page 355) is true, append an entry in the *form data set* with the *name* as the name, the `value` (page 355) of the *option* element as the value, and `type` as the type.
 7. Otherwise, if the *field* element is an *input* element whose `type` attribute is in the Checkbox (page 335) state or the Radio Button (page 336) state, then then run these further nested substeps:
 1. If the *field* element has a `value` attribute specified, then let *value* be the value of that attribute; otherwise, let *value* be the string "on".
 2. Append an entry in the *form data set* with *name* as the name, *value* as the value, and `type` as the type.
 8. Otherwise, if the *field* element is an *input* element whose `type` attribute is in the File Upload (page 337) state, then for each file selected (page 337) in the *input* element, append an entry in the *form data set* with the *name* as the name, the file (consisting of the name, the type, and the body) as the value, and `type` as the type.

9. Otherwise, append an entry in the *form data set* with *name* as the name, the value (page 362) of the *field* element as the value, and *type* as the type.
7. Let *action* be the *submitter* element's action (page 363).
8. If *action* is the empty string, let *action* be the document's address.
9. Resolve (page 55) the URL (page 52) *action*. If this fails, abort these steps.
10. Let *scheme* be the <scheme> (page 53) of the resulting absolute URL (page 56).
11. Let *enctype* be the *submitter* element's *enctype* (page 364).
12. Let *method* be the *submitter* element's *method* (page 364).
13. Let *target* be the *submitter* element's *target* (page 364).
14. Select the appropriate row in the table below based on the value of *scheme* as given by the first cell of each row. Then, select the appropriate cell on that row based on the value of *method* as given in the first cell of each column. Then, jump to the steps named in that cell and defined below the table.

	GET (page 363)	POST (page 363)	PUT (page 364)	DELETE (page 364)
http	Mutate action (page 370)	Submit as entity body (page 371)	Submit as entity body (page 371)	Delete action (page 371)
https	Mutate action (page 370)	Submit as entity body (page 371)	Submit as entity body (page 371)	Delete action (page 371)
ftp	Get action (page 371)	Get action (page 371)	Get action (page 371)	Get action (page 371)
javascript	Get action (page 371)	Get action (page 371)	Get action (page 371)	Get action (page 371)
data	Get action (page 371)	Post to data: (page 371)	Put to data: (page 372)	Get action (page 371)
mailto	Mail with headers (page 372)	Mail as body (page 373)	Mail with headers (page 372)	Mail with headers (page 372)

If *scheme* is not one of those listed in this table, then the behavior is not defined by this specification. User agents should, in the absence of another specification defining this, act in a manner analogous to that defined in this specification for similar schemes.

The behaviors are as follows:

Mutate action

Let *query* be the result of encoding the *form data set* using the application/x-www-form-urlencoded encoding algorithm (page 373), interpreted as a US-ASCII string.

Let *destination* be a new URL (page 52) that is equal to the *action* except that its <query> (page 54) component is replaced by *query* (adding a U+003F QUESTION MARK (?) character if appropriate).

Let *target browsing context* be the form submission target browsing context (page 373).

Navigate (page 473) *target browsing context* to *destination*. If *target browsing context* was newly created for this purpose by the steps above, then it must be navigated with replacement enabled (page 476).

Submit as entity body

Let *entity body* be the result of encoding the *form data set* using the appropriate form encoding algorithm (page 373).

Let *target browsing context* be the form submission target browsing context (page 373).

Let *MIME type* be determined as follows:

If *enctype* is application/x-www-form-urlencoded

Let *MIME type* be "application/x-www-form-urlencoded".

If *enctype* is multipart/form-data

Let *MIME type* be "multipart/form-data".

If *enctype* is text/plain

Let *MIME type* be "text/plain".

Navigate (page 473) *target browsing context* to *action* using the HTTP method given by *method* and with *entity body* as the entity body, of type *MIME type*. If *target browsing context* was newly created for this purpose by the steps above, then it must be navigated with replacement enabled (page 476).

Delete action

Let *target browsing context* be the form submission target browsing context (page 373).

Navigate (page 473) *target browsing context* to *action* using the DELETE method. If *target browsing context* was newly created for this purpose by the steps above, then it must be navigated with replacement enabled (page 476).

Get action

Let *target browsing context* be the form submission target browsing context (page 373).

Navigate (page 473) *target browsing context* to *action*. If *target browsing context* was newly created for this purpose by the steps above, then it must be navigated with replacement enabled (page 476).

Post to data:

Let *data* be the result of encoding the *form data set* using the appropriate form encoding algorithm (page 373).

If *action* contains the string "%%%%" (four U+0025 PERCENT SIGN characters), then %-escape all bytes in *data* that, if interpreted as US-ASCII, do not match the unreserved production in the URI Generic Syntax, and then, treating the result as a US-ASCII string, further %-escape all the U+0025 PERCENT SIGN characters in the resulting string and replace the first occurrence of "%%%%" in *action* with the resulting double-escaped string. [RFC3986]

Otherwise, if *action* contains the string "%%" (two U+0025 PERCENT SIGN characters in a row, but not four), then %-escape all characters in *data* that, if interpreted as US-ASCII, do not match the unreserved production in the URI Generic Syntax, and then, treating the result as a US-ASCII string, replace the first occurrence of "%%" in *action* with the resulting escaped string. [RFC3986]

Let *target browsing context* be the form submission target browsing context (page 373).

Navigate (page 473) *target browsing context* to the potentially modified *action*. If *target browsing context* was newly created for this purpose by the steps above, then it must be navigated with replacement enabled (page 476).

Put to data:

Let *data* be the result of encoding the *form data set* using the appropriate form encoding algorithm (page 373).

Let *MIME type* be determined as follows:

If enctype is application/x-www-form-urlencoded

Let *MIME type* be "application/x-www-form-urlencoded".

If enctype is multipart/form-data

Let *MIME type* be "multipart/form-data".

If enctype is text/plain

Let *MIME type* be "text/plain".

Let *destination* be the result of concatenating the following:

1. The string "data:".
2. The value of *MIME type*.
3. The string ";base64,".
4. A base-64 encoded representation of *data*. [RFC2045]

Let *target browsing context* be the form submission target browsing context (page 373).

Navigate (page 473) *target browsing context* to *destination*. If *target browsing context* was newly created for this purpose by the steps above, then it must be navigated with replacement enabled (page 476).

Mail with headers

Let *headers* be the resulting encoding the *form data set* using the application/x-www-form-urlencoded encoding algorithm (page 373), interpreted as a US-ASCII string.

Replace occurrences of U+002B PLUS SIGN characters (+) in *headers* with the string "%20".

Let *destination* consist of all the characters from the first character in *action* to the character immediately before the first U+003F QUESTION MARK character (?), if any, or the end of the string if there are none.

Append a single U+003F QUESTION MARK character (?) to *destination*.

Append *headers* to *destination*.

Let *target browsing context* be the form submission target browsing context (page 373).

Navigate (page 473) *target browsing context* to *destination*. If *target browsing context* was newly created for this purpose by the steps above, then it must be navigated with replacement enabled (page 476).

Mail as body

Let *body* be the resulting encoding the *form data set* using the appropriate form encoding algorithm (page 373) and then %-escaping all the bytes in the resulting byte string that, when interpreted as US-ASCII, do not match the unreserved production in the URI Generic Syntax. [RFC3986]

Let *destination* have the same value as *action*.

If *destination* does not contain a U+003F QUESTION MARK character (?), append a single U+003F QUESTION MARK character (?) to *destination*. Otherwise, append a single U+0026 AMPERSAND character (&).

Append the string "body=" to *destination*.

Append *body*, interpreted as a US-ASCII string, to *destination*.

Let *target browsing context* be the form submission target browsing context (page 373).

Navigate (page 473) *target browsing context* to *destination*. If *target browsing context* was newly created for this purpose by the steps above, then it must be navigated with replacement enabled (page 476).

The form submission target browsing context is obtained, when needed by the behaviors described above, as follows: If the user indicated a specific browsing context (page 414) to use when submitting the form, then that is the target browsing context. Otherwise, apply the rules for choosing a browsing context given a browsing context name (page 418) using *target* as the name and the browsing context (page 414) of *form* as the context in which the algorithm is executed; the resulting browsing context (page 414) is the target browsing context.

The **appropriate form encoding algorithm** is determined as follows:

If *enctype* is application/x-www-form-urlencoded

Use the application/x-www-form-urlencoded encoding algorithm (page 373).

If *enctype* is multipart/form-data

Use the multipart/form-data encoding algorithm (page 375).

If *enctype* is text/plain

Use the text/plain encoding algorithm (page 375).

4.10.15.1 URL-encoded form data

The **application/x-www-form-urlencoded encoding algorithm** is as follows:

1. Let *result* be the empty string.

2. If the `form` element has an `accept-charset` attribute, then, taking into account the characters found in the `form data set`'s names and values, and the character encodings supported by the user agent, select a character encoding from the list given in the form's `accept-charset` attribute that is an ASCII-compatible character encoding (page 25). If none of the encodings are supported, then let the selected character encoding be UTF-8.

Otherwise, if the document's character encoding (page 79) is an ASCII-compatible character encoding (page 25), then that is the selected character encoding.

Otherwise, let the selected character encoding be UTF-8.

3. Let `charset` be the preferred MIME name of the selected character encoding.
4. If the entry's name is "`_charset_`" and its type is "hidden", replace its value with `charset`.
5. If the entry's type is "file", replace its value with the file's filename only.
6. For each entry in the `form data set`, perform these substeps:
 1. For each character in the entry's name and value that cannot be expressed using the selected character encoding, replace the character by a string consisting of a U+0026 AMPERSAND character (&), one or more characters in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9) representing the Unicode codepoint of the character in base ten, and finally a U+003B SEMICOLON character (;).
 2. For each character in the entry's name and value, apply the following subsubsteps:
 1. If the character isn't in the range U+0020, U+002A, U+002D, U+002E, U+0030 .. U+0039, U+0041 .. U+005A, U+005F, U+0061 .. U+007A then replace the character with a string formed as follows: Start with the empty string, and then, taking each byte of the character when expressed in the selected character encoding in turn, append to the string a U+0025 PERCENT SIGN character (%) followed by two characters in the ranges U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9) and U+0041 LATIN CAPITAL LETTER A to U+005A LATIN CAPITAL LETTER Z representing the hexadecimal value of the byte (zero-padded if necessary).
 2. If the character is a U+0020 SPACE character, replace it with a single U+002B PLUS SIGN character (+).
 3. If the entry's name is "isindex", its type is "text", and this is the first entry in the `form data set`, then append the value to `result` and skip the rest of the substeps for this entry, moving on to the next entry, if any, or the next step in the overall algorithm otherwise.
 4. If this is not the first entry, append a single U+0026 AMPERSAND character (&) to `result`.
 5. Append the entry's name to `result`.

6. Append a single U+003D EQUALS SIGN character (=) to *result*.
7. Append the entry's value to *result*.
7. Encode *result* as US-ASCII and return the resulting byte stream.

4.10.15.2 Multipart form data

The **multipart/form-data encoding algorithm** is to encode the *form data set* using the rules described by RFC2388, *Returning Values from Forms: multipart/form-data*, and return the resulting byte stream. [RFC2388]

Each entry in the *form data set* is a *field*, the name of the entry is the *field name* and the value of the entry is the *field value*.

The order of parts must be the same as the order of fields in the *form data set*. Multiple entries with the same name must be treated as distinct fields.

4.10.15.3 Plain text form data

The **text/plain encoding algorithm** is as follows:

1. Let *result* be the empty string.
2. If the *form* element has an *accept-charset* attribute, then, taking into account the characters found in the *form data set*'s names and values, and the character encodings supported by the user agent, select a character encoding from the list given in the *form*'s *accept-charset* attribute. If none of the encodings are supported, then let the selected character encoding be UTF-8.
Otherwise, the selected character encoding is the document's character encoding (page 79).
3. Let *charset* be the preferred MIME name of the selected character encoding.
4. If the entry's name is "*_charset_*" and its type is "hidden", replace its value with *charset*.
5. If the entry's type is "file", replace its value with the file's filename only.
6. For each entry in the *form data set*, perform these substeps:
 1. Append the entry's name to *result*.
 2. Append a single U+003D EQUALS SIGN character (=) to *result*.
 3. Append the entry's value to *result*.
 4. Append a U+000D CARRIAGE RETURN (CR) U+000A LINE FEED (LF) character pair to *result*.
7. Encode *result* using the selected character encoding and return the resulting byte stream.

4.10.16 Resetting a form

When a form *form* is **reset**, the user agent must invoke the reset algorithm (page 376) of each resettable (page 314) elements whose form owner (page 361) is *form*, and must then broadcast **formchange** events (page 376) from *form*.

Each resettable (page 314) element defines its own **reset algorithm**.

4.10.17 Event dispatch

When the user agent is to **broadcast forminput events** or **broadcast formchange events** from a form element *form*, it must run the following steps:

1. Let *controls* be a list of all the resettable (page 314) elements whose form owner (page 361) is *form*.
2. If the user agent was to broadcast forminput events (page 376), let *event name* be **forminput**. Otherwise the user agent was to broadcast formchange events (page 376); let *event name* be **formchange**.
3. For each element in *controls*, in tree order (page 24), fire a simple event (page 436) named *event name* at the element.

** Still need to define when **formchange** and **forminput** fire naturally.

4.11 Interactive elements

4.11.1 The details element

Categories

Flow content (page 93).
Interactive content (page 95).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

One legend element followed by flow content (page 93).

Element-specific attributes:

open

DOM interface:

```
interface HTMLDetailsElement : HTMLElement {  
    attribute boolean open;  
};
```

The **details** element represents additional information or controls which the user can obtain on demand.

Note: The details element is not appropriate for footnotes. Please see the section on footnotes (page 191) for details on how to mark up footnotes.

The first element child of a details element, if it is a legend element, represents the summary of the details.

If the first element is not a legend element, the UA should provide its own legend (e.g. "Details").

The **open** content attribute is a boolean attribute (page 32). If present, it indicates that the details should be shown to the user. If the attribute is absent, the details should not be shown.

If the attribute is removed, then the details should be hidden. If the attribute is added, the details should be shown.

The user agent should allow the user to request that the details be shown or hidden. To honor a request for the details to be shown, the user agent must set the open attribute on the element to the value open. To honour a request for the details to be hidden, the user agent must remove the open attribute from the element.

The **open** attribute must reflect (page 67) the open content attribute.

- ** Rendering will be described in the Rendering section in due course. Basically CSS :open and :closed match the element, it's a block-level element by default, and when it matches :closed it renders as if it had an XBL binding attached to it whose template was just <template>><content>
includes="legend:first-child">Details</content></template>, and when it's :open it acts as if it had an XBL binding attached to it whose template was just <template>><content>
includes="legend:first-child">Details</content><content/></template> or some such.
- **

4.11.2 The datagrid element

Categories

- Flow content (page 93).
- Interactive content (page 95).
- Sectioning root (page 140).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

- Either: Nothing.
- Or: Flow content (page 93), but where the first element child node, if any, is not a table, select, or datalist element.
- Or: A single table element.
- Or: A single select element.
- Or: A single datalist element.

Element-specific attributes:

`multiple`
`disabled`

DOM interface:

```
interface HTMLDataGridElement : HTMLElement {  
    attribute DataGridDataProvider data;  
    readonly attribute DataGridSelection selection;  
    attribute boolean multiple;  
    attribute boolean disabled;  
    void updateEverything();  
    void updateRowsChanged(in RowSpecification row, in unsigned long  
    count);  
    void updateRowsInserted(in RowSpecification row, in unsigned long  
    count);  
    void updateRowsRemoved(in RowSpecification row, in unsigned long  
    count);  
    void updateRowChanged(in RowSpecification row);  
    void updateColumnChanged(in unsigned long column);  
    void updateCellChanged(in RowSpecification row, in unsigned long  
    column);  
};
```

- ** One possible thing to be added is a way to detect when a row/selection has been deleted,
- ** activated, etc, by the user (delete key, enter key, etc).

The datagrid element represents an interactive representation of tree, list, or tabular data.

The data being presented can come either from the content, as elements given as children of the datagrid element, or from a scripted data provider given by the data DOM attribute.

The `multiple` and `disabled` attributes are boolean attributes (page 32). Their effects are described in the processing model sections below.

The `multiple` and `disabled` DOM attributes must reflect (page 67) the `multiple` and `disabled` content attributes respectively.

4.11.2.1 The datagrid data model

This section is non-normative.

In the datagrid data model, data is structured as a set of rows representing a tree, each row being split into a number of columns. The columns are always present in the data model, although individual columns may be hidden in the presentation.

Each row can have child rows. Child rows may be hidden or shown, by closing or opening (respectively) the parent row.

Rows are referred to by the path along the tree that one would take to reach the row, using zero-based indices. Thus, the first row of a list is row "0", the second row is row "1"; the first child row of the first row is row "0,0", the second child row of the first row is row "0,1"; the fourth child of the seventh child of the third child of the tenth row is "9,2,6,3", etc.

The columns can have captions. Those captions are not considered a row in their own right, they are obtained separately.

Selection of data in a datagrid operates at the row level. If the `multiple` attribute is present, multiple rows can be selected at once, otherwise the user can only select one row at a time.

The datagrid element can be disabled entirely by setting the `disabled` attribute.

Columns, rows, and cells can each have specific flags, known as classes, applied to them by the data provider. These classes affect the functionality (page 382) of the datagrid element, and are also passed to the style system (page 672). They are similar in concept to the `class` attribute, except that they are not specified on elements but are given by scripted data providers.

4.11.2.2 How rows are identified

The chains of numbers that give a row's path, or identifier, are represented by objects that implement the `RowSpecification` (page 379) interface.

```
[NoInterfaceObject] interface RowSpecification {
    // binding-specific interface
};
```

In ECMAScript, two classes of objects are said to implement this interface: Numbers representing non-negative integers, and homogeneous arrays of Numbers representing non-negative integers. Thus, `[1,0,9]` is a `RowSpecification` (page 379), as is `1` on its own. However, `[1,0.2,9]` is not a `RowSpecification` (page 379) object, since its second value is not an integer.

User agents must always represent `RowSpecifications` in ECMAScript by using arrays, even if the path only has one number.

The root of the tree is represented by the empty path; in ECMAScript, this is the empty array `([])`. Only the `getRowCount()` and `GetChildAtPosition()` methods ever get called with the empty path.

4.11.2.3 The data provider interface

The conformance criteria in this section apply to any implementation of the `DataGridDataProvider`, including (and most commonly) the content author's implementation(s).

```
// To be implemented by Web authors as a JS object
[NoInterfaceObject] interface DataGridDataProvider {
    void initialize(in HTMLDataGridElement datagrid);
    unsigned long getCount(in RowSpecification row);
```

```

    unsigned long getChildAtPosition(in RowSpecification parentRow, in
unsigned long position);
    unsigned long getColumnCount();
    DOMString getCaptionText(in unsigned long column);
    void getCaptionClasses(in unsigned long column, in DOMTokenList
classes);
    DOMString getRowImage(in RowSpecification row);
    HTMLMenuItemElement getRowMenu(in RowSpecification row);
    void getRowClasses(in RowSpecification row, in DOMTokenList classes);
    DOMString getCellData(in RowSpecification row, in unsigned long
column);
    void getCellClasses(in RowSpecification row, in unsigned long column,
in DOMTokenList classes);
    void toggleColumnSortState(in unsigned long column);
    void setCellCheckedState(in RowSpecification row, in unsigned long
column, in long state);
    void cycleCell(in RowSpecification row, in unsigned long column);
    void editCell(in RowSpecification row, in unsigned long column, in
DOMString data);
};


```

The DataGridDataProvider interface represents the interface that objects must implement to be used as custom data views for datagrid elements.

Not all the methods are required. The minimum number of methods that must be implemented in a useful view is two: the `getRowCount()` and `getCellData()` methods.

Once the object is written, it must be hooked up to the datagrid using the **data** DOM attribute.

The following methods may be usefully implemented:

initialize(datagrid)

Called by the datagrid element (the one given by the *datagrid* argument) after it has first populated itself. This would typically be used to set the initial selection of the datagrid element when it is first loaded. The data provider could also use this method call to register a select event handler on the datagrid in order to monitor selection changes.

getRowCount(row)

Must return the number of rows that are children of the specified *row*, including rows that are off-screen. If *row* is empty, then the number of rows at the top level must be returned. If the value that this method would return for a given *row* changes, the relevant update methods on the datagrid must be called first. Otherwise, this method must always return the same number. For a list (as opposed to a tree), this method must return 0 whenever it is called with a *row* identifier that is not empty.

getChildAtPosition(parentRow, position)

Must return the index of the row that is a child of *parentRow* and that is to be positioned as the *position*th row under *parentRow* when rendering the children of *parentRow*. If *parentRow* is empty, then *position* refers to the *position*th row at the top level of the data grid. May be omitted if the rows are always to be sorted in the natural order. (The natural order is the one where the method always returns *position*.) For a given

parentRow, this method must never return the same value for different values of *position*. The returned value *x* must be in the range $0 \leq x < n$, where *n* is the value returned by `getRowCount(parentRow)`.

`getColumnCount()`

Must return the number of columns currently in the data model (including columns that might be hidden). May be omitted if there is only one column. If the value that this method would return changes, the datagrid's `updateEverything()` method must be called.

`getCaptionText(column)`

Must return the caption, or label, for column *column*. May be omitted if the columns have no captions. If the value that this method would return changes, the datagrid's `updateColumnChanged()` method must be called with the appropriate column index.

`getCaptionClasses(column, classes)`

Must add the classes that apply to column *column* to the *classes* object. May be omitted if the columns have no special classes. If the classes that this method would add changes, the datagrid's `updateColumnChanged()` method must be called with the appropriate column index. Some classes have predefined meanings (page 382).

`getRowImage(row)`

Must return a URL (page 52) giving the address of an image that represents row *row*, or the empty string if there is no applicable image. May be omitted if no rows have associated images. If the value that this method would return changes, the datagrid's update methods must be called to update the row in question.

`getRowMenu(row)`

Must return an `HTMLElement` object that is to be used as a context menu for row *row*, or null if there is no particular context menu. May be omitted if none of the rows have a special context menu. As this method is called immediately before showing the menu in question, no precautions need to be taken if the return value of this method changes.

`getRowClasses(row, classes)`

Must add the classes that apply to row *row* to the *classes* object. May be omitted if the rows have no special classes. If the classes that this method would add changes, the datagrid's update methods must be called to update the row in question. Some classes have predefined meanings (page 382).

`getCellData(row, column)`

Must return the value of the cell on row *row* in column *column*. For text cells, this must be the text to show for that cell. For progress bar cells (page 383), this must be either a floating point number in the range 0.0 to 1.0 (converted to a string representation), indicating the fraction of the progress bar to show as full (1.0 meaning complete), or the empty string, indicating an indeterminate progress bar. If the value that this method would return changes, the datagrid's update methods must be called to update the rows that changed. If only one cell changed, the `updateCellChanged()` method may be used.

`getCellClasses(row, column, classes)`

Must add the classes that apply to the cell on row *row* in column *column* to the *classes* object. May be omitted if the cells have no special classes. If the classes that this

method would add changes, the datagrid's update methods must be called to update the rows or cells in question. Some classes have predefined meanings (page 382).

`toggleColumnSortState(column)`

Called by the datagrid when the user tries to sort the data using a particular column *column*. The data provider must update its state so that the `GetChildAtPosition()` method returns the new order, and the classes of the columns returned by `getCaptionClasses()` represent the new sort status. There is no need to tell the datagrid that it the data has changed, as the datagrid automatically assumes that the entire data model will need updating.

`setCellCheckedState(row, column, state)`

Called by the datagrid when the user changes the state of a checkbox cell on row *row*, column *column*. The checkbox should be toggled to the state given by *state*, which is a positive integer (1) if the checkbox is to be checked, zero (0) if it is to be unchecked, and a negative number (-1) if it is to be set to the indeterminate state. There is no need to tell the datagrid that the cell has changed, as the datagrid automatically assumes that the given cell will need updating.

`cycleCell(row, column)`

Called by the datagrid when the user changes the state of a cyclable cell on row *row*, column *column*. The data provider should change the state of the cell to the new state, as appropriate. There is no need to tell the datagrid that the cell has changed, as the datagrid automatically assumes that the given cell will need updating.

`editCell(row, column, data)`

Called by the datagrid when the user edits the cell on row *row*, column *column*. The new value of the cell is given by *data*. The data provider should update the cell accordingly. There is no need to tell the datagrid that the cell has changed, as the datagrid automatically assumes that the given cell will need updating.

The following classes (for rows, columns, and cells) may be usefully used in conjunction with this interface:

Class name	Applies to	Description
checked	Cells	The cell has a checkbox and it is checked. (The cyclable and progress classes override this, though.)
cyclable	Cells	The cell can be cycled through multiple values. (The progress class overrides this, though.)
editable	Cells	The cell can be edited. (The cyclable, progress, checked, unchecked and indeterminate classes override this, though.)
header	Rows	The row is a heading, not a data row.
indeterminate	Cells	The cell has a checkbox, and it can be set to an indeterminate state. If neither the checked nor unchecked classes are present, then the checkbox is in that state, too. (The cyclable and progress classes override this, though.)
initially-hidden	Columns	The column will not be shown when the datagrid is initially rendered. If this class is not present on the column when the datagrid is initially rendered, the column will be visible if space allows.
initially-closed	Rows	The row will be closed when the datagrid is initially rendered. If neither this class nor the initially-open class is present on the row when the datagrid is initially rendered, the initial state will depend on platform conventions.

initially-open	Rows	The row will be opened when the datagrid is initially rendered. If neither this class nor the <code>initially-closed</code> class is present on the row when the datagrid is initially rendered, the initial state will depend on platform conventions.
progress	Cells	The cell is a progress bar.
reversed	Columns	If the cell is sorted, the sort direction is descending, instead of ascending.
selectable-separator	Rows	The row is a normal, selectable, data row, except that instead of having data, it only has a separator. (The header and separator classes override this, though.)
separator	Rows	The row is a separator row, not a data row. (The header class overrides this, though.)
sortable	Columns	The data can be sorted by this column.
sorted	Columns	The data is sorted by this column. Unless the <code>reversed</code> class is also present, the sort direction is ascending.
unchecked	Cells	The cell has a checkbox and, unless the <code>checked</code> class is present as well, it is unchecked. (The <code>cyclable</code> and <code>progress</code> classes override this, though.)

4.11.2.4 The default data provider

The user agent must supply a default data provider for the case where the datagrid's `data` attribute is null. It must act as described in this section.

The behavior of the default data provider depends on the nature of the first element child of the datagrid.

↪ While the first element child is a table element

`getRowCount(row)`: The number of rows returned by the default data provider for the root of the tree (when `row` is empty) must be the total number of `tr` elements that are children of `tbody` elements that are children of the table, if there are any such child `tbody` elements. If there are no such `tbody` elements then the number of rows returned for the root must be the number of `tr` elements that are children of the table.

When `row` is not empty, the number of rows returned must be zero.

Note: *The table-based default data provider cannot represent a tree.*

Note: *Rows in thead elements do not contribute to the number of rows returned, although they do affect the columns and column captions. Rows in tfoot elements are ignored (page 24) completely by this algorithm.*

`getChildAtPosition(row, i)`: The default data provider must return the mapping appropriate to the current sort order (page 385).

`getColumnCount()`: The number of columns returned must be the number of `td` element children in the first `tr` element child of the first `tbody` element child of the table, if there are any such `tbody` elements. If there are no such `tbody` elements, then it must be the number of `td` element children in the first `tr` element child of the

table, if any, or otherwise 1. If the number that would be returned by these rules is 0, then 1 must be returned instead.

getCaptionText(*i*): If the table has no thead element child, or if its first thead element child has no tr element child, the default data provider must return the empty string for all captions. Otherwise, the value of the textContent attribute of the *i*th th element child of the first tr element child of the first thead element child of the table element must be returned. If there is no such th element, the empty string must be returned.

getCaptionClasses(*i, classes*): If the table has no thead element child, or if its first thead element child has no tr element child, the default data provider must not add any classes for any of the captions. Otherwise, each class in the class attribute of the *i*th th element child of the first tr element child of the first thead element child of the table element must be added to the *classes*. If there is no such th element, no classes must be added. The user agent must then:

1. Remove the sorted and reversed classes.
2. If the table element has a class attribute that includes the sortable class, add the sortable class.
3. If the column is the one currently being used to sort the data, add the sorted class.
4. If the column is the one currently being used to sort the data, and it is sorted in descending order, add the reversed class as well.

The various row- and cell- related methods operate relative to a particular element, the element of the row or cell specified by their arguments.

For rows: Since the default data provider for a table always returns 0 as the number of children for any row other than the root, the path to the row passed to these methods will always consist of a single number. In the prose below, this number is referred to as *i*.

If the table has tbody element children, the element for the *i*th row is the *i*th tr element that is a child of a tbody element that is a child of the table element. If the table does not have tbody element children, then the element for the *i*th real row is the *i*th tr element that is a child of the table element.

For cells: Given a row and its element, the row's *i*th cell's element is the *i*th td element child of the row element.

Note: *The colspan and rowspan attributes are ignored (page 24) by this algorithm.*

getRowImage(*i*): The URL (page 52) of the row's image is the absolute URL (page 56) obtained by resolving (page 55) the value of the src attribute of the first img element child of the row's first cell's element, if there is one and resolving its attribute is successful. Otherwise, the URL (page 52) of the row's image is the empty string.

getRowMenu(*i*): If the row's first cell's element has a menu element child, then the row's menu is the first menu element child of the row's first cell's element. Otherwise, the row has no menu.

getRowClasses(*i, classes*): The default data provider must never add a class to the row's classes.

toggleColumnSortState(*i*): If the data is already being sorted on the given column, then the user agent must change the current sort mapping to be the inverse of the current sort mapping; if the sort order was ascending before, it is now descending, otherwise it is now ascending. Otherwise, if the current sort column is another column, or the data model is currently not sorted, the user agent must create a new mapping, which maps rows in the data model to rows in the DOM so that the rows in the data model are sorted by the specified column, in ascending order. (Which sort comparison operator to use is left up to the UA to decide.)

When the sort mapping is changed, the values returned by the `getChildAtPosition()` method for the default data provider will change appropriately (page 383).

getCellData(*i, j*), getCellClasses(*i, j, classes*), getCellCheckedState(*i, j, state*), cycleCell(*i, j*), and editCell(*i, j, data*): See the common definitions below (page 388).

The data provider must call the datagrid's update methods appropriately whenever the descendants of the datagrid mutate. For example, if a `tr` is removed, then the `updateRowsRemoved()` methods would probably need to be invoked, and any change to a cell or its descendants must cause the cell to be updated. If the `table` element stops being the first child of the datagrid, then the data provider must call the `updateEverything()` method on the datagrid. Any change to a cell that is in the column that the data provider is currently using as its sort column must also cause the sort to be reperformed, with a call to `updateEverything()` if the change did affect the sort order.

↪ While the first element child is a select or datalist element

The default data provider must return 1 for the column count, the empty string for the column's caption, and must not add any classes to the column's classes.

For the rows, assume the existence of a node filter view of the descendants of the first element child of the datagrid element (the `select` or `datalist` element), that skips all nodes other than `optgroup` and `option` elements, as well as any descendants of any `option` elements.

Given a path `row`, the corresponding element is the one obtained by drilling into the view, taking the child given by the path each time.

Given the following XML markup:

```
<datagrid>
  <select>
    <!-- the options and optgroups have had their labels and
        values removed
        to make the underlying structure clearer -->
    <optgroup>
```

```

<option/>
<option/>
</optgroup>
<optgroup>
<option/>
<optgroup id="a">
<option/>
<option/>
<bogus/>
<option id="b"/>
</optgroup>
<option/>
</optgroup>
</select>
</datagrid>
```

The path "1,1,2" would select the element with ID "b". In the filtered view, the text nodes, comment nodes, and bogus elements are ignored; so for instance, the element with ID "a" (path "1,1") has only 3 child nodes in the view.

`getRowCount(row)` must drill through the view to find the element corresponding to the method's argument, and return the number of child nodes in the filtered view that the corresponding element has. (If the `row` is empty, the corresponding element is the `select` element at the root of the filtered view.)

`getChildAtPosition(row, position)` must return `position`. (The `select/datalist` default data provider does not support sorting the data grid.)

`getRowImage(i)` must return the empty string, `getRowMenu(i)` must return null.

`getRowClasses(row, classes)` must add the classes from the following list to `classes` when their condition is met:

- If the `row`'s corresponding element is an `optgroup` element: `header`
- If the `row`'s corresponding element contains other elements that are also in the view, and the element's `class` attribute contains the `closed` class: `initially-closed`
- If the `row`'s corresponding element contains other elements that are also in the view, and the element's `class` attribute contains the `open` class: `initially-open`

The `getCellData(row, cell)` method must return the value of the `label` attribute if the `row`'s corresponding element is an `optgroup` element, otherwise, if the `row`'s corresponding element is an `option` element, its `label` attribute if it has one, otherwise the value of its `textContent` DOM attribute.

The `getCellClasses(row, cell, classes)` method must add no classes.

**
**

autoselect some rows when initialized, reflect the selection in the select, reflect the multiple attribute somehow.

The data provider must call the datagrid's update methods appropriately whenever the descendants of the datagrid mutate.

↳ While the first element child is another element

The default data provider must return 1 for the column count, the empty string for the column's caption, and must not add any classes to the column's classes.

For the rows, assume the existence of a node filter view of the descendants of the datagrid that skips all nodes other than li, h1-h6, and hr elements, and skips any descendants of menu elements.

Given this view, each element in the view represents a row in the data model. The element corresponding to a path *row* is the one obtained by drilling into the view, taking the child given by the path each time. The element of the row of a particular method call is the element given by drilling into the view along the path given by the method's arguments.

`getRowCount(row)` must return the number of child elements in this view for the given row, or the number of elements at the root of the view if the *row* is empty.

In the following example, the elements are identified by the paths given by their child text nodes:

```
<datagrid>
  <ol>
    <li> row 0 </li>
    <li> row 1
      <ol>
        <li> row 1,0 </li>
      </ol>
    </li>
    <li> row 2 </li>
  </ol>
</datagrid>
```

In this example, only the li elements actually appear in the data grid; the ol element does not affect the data grid's processing model.

`getChildAtPosition(row, position)` must return *position*. (The generic default data provider does not support sorting the data grid.)

`getRowImage(i)` must return the absolute URL (page 56) obtained from resolving (page 55) the value of the src attribute of the first img element descendant (in the real DOM) of the row's element, that is not also a descendant of another element in the filtered view that is a descendant of the row's element, if such an element exists and resolving its attribute is successful. Otherwise, it must return the empty string.

In the following example, the row with path "1,0" returns "http://example.com/a" as its image URL, and the other rows (including the row with path "1") return the empty string:

```
<datagrid>
  <ol>
    <li> row 0 </li>
    <li> row 1
```

```

<ol>
  <li> row 1,0  </li>
</ol>
</li>
<li> row 2 </li>
</ol>
</datagrid>

```

`getRowMenu(i)` must return the first menu element descendant (in the real DOM) of the row's element, that is not also a descendant of another element in the filtered view that is a descendant of the row's element. (This is analogous to the image case above.)

`getRowClasses(i, classes)` must add the classes from the following list to `classes` when their condition is met:

- If the row's element contains other elements that are also in the view, and the element's `class` attribute contains the `closed` class: `initially-closed`
- If the row's element contains other elements that are also in the view, and the element's `class` attribute contains the `open` class: `initially-open`
- If the row's element is an `h1-h6` element: `header`
- If the row's element is an `hr` element: `separator`

The `getCellData(i, j)`, `getCellClasses(i, j, classes)`, `getCellCheckedState(i, j, state)`, `cycleCell(i, j)`, and `editCell(i, j, data)` methods must act as described in the common definitions below (page 388), treating the row's element as being the cell's element.

**

selection handling?

The data provider must call the datagrid's update methods appropriately whenever the descendants of the datagrid mutate.

→ Otherwise, while there is no element child

The data provider must return 0 for the number of rows, 1 for the number of columns, the empty string for the first column's caption, and must add no classes when asked for that column's classes. If the datagrid's child list changes such that there is a first element child, then the data provider must call the `updateEverything()` method on the datagrid.

4.11.2.4.1 Common default data provider method definitions for cells

These definitions are used for the cell-specific methods of the default data providers (other than in the select/datalist case). How they behave is based on the contents of an element that represents the cell given by their first two arguments. Which element that is is defined in the previous section.

Cyclable cells

If the first element child of a cell's element is a select element that has a no multiple attribute and has at least one option element descendent, then the cell acts as a cyclable cell.

The "current" option element is the selected option element, or the first option element if none is selected.

The `getCellData()` method must return the `textContent` of the current option element (the `label` attribute is ignored (page 24) in this context as the optgroups are not displayed).

The `getCellClasses()` method must add the `cyclable` class and then all the classes of the current option element.

The `cycleCell()` method must change the selection of the select element such that the next option element after the current option element is the only one that is selected (in tree order (page 24)). If the current option element is the last option element descendent of the select, then the first option element descendent must be selected instead.

The `setCellCheckedState()` and `editCell()` methods must do nothing.

Progress bar cells

If the first element child of a cell's element is a progress element, then the cell acts as a progress bar cell.

The `getCellData()` method must return the value returned by the progress element's position DOM attribute.

The `getCellClasses()` method must add the `progress` class.

The `setCellCheckedState()`, `cycleCell()`, and `editCell()` methods must do nothing.

Checkbox cells

If the first element child of a cell's element is an input element that has a type attribute with the value checkbox, then the cell acts as a check box cell.

The `getCellData()` method must return the `textContent` of the cell element.

The `getCellClasses()` method must add the `checked` class if the input element's checkedness (page 362) is true, and the `unchecked` class otherwise.

The `setCellCheckedState()` method must set the input element's checkbox checkedness (page 362) to true if the method's third argument is 1, and to false otherwise.

The `cycleCell()` and `editCell()` methods must do nothing.

Editable cells

If the first element child of a cell's element is an input element that has a type attribute with the value text or that has no type attribute at all, then the cell acts as an editable cell.

The `getCellData()` method must return the value of the input element.

The `getCellClasses()` method must add the `editable` class.

The `editCell()` method must set the `input` element's `value` DOM attribute to the value of the third argument to the method.

The `setCellCheckedState()` and `cycleCell()` methods must do nothing.

4.11.2.5 Populating the datagrid element

A datagrid must be disabled until its end tag has been parsed (in the case of a datagrid element in the original document markup) or until it has been inserted into the document (page 24) (in the case of a dynamically created element). After that point, the element must fire a single load event at itself, which doesn't bubble and cannot be canceled.

** The end-tag parsing thing should be moved to the parsing section.

The datagrid must then populate itself using the data provided by the data provider assigned to the `data` DOM attribute. After the view is populated (using the methods described below), the datagrid must invoke the `initialize()` method on the data provider specified by the `data` attribute, passing itself (the `HTMLDataGridElement` object) as the only argument.

When the `data` attribute is null, the datagrid must use the default data provider described in the previous section.

To obtain data from the data provider, the element must invoke methods on the data provider object in the following ways:

To determine the total number of columns

Invoke the `getRowCount()` method with no arguments. The return value is the number of columns. If the return value is zero or negative, not an integer, or simply not a numeric type, or if the method is not defined, then 1 must be used instead.

To get the captions to use for the columns

Invoke the `getCaptionText()` method with the index of the column in question. The index i must be in the range $0 \leq i < N$, where N is the total number of columns. The return value is the string to use when referring to that column. If the method returns null or the empty string, the column has no caption. If the method is not defined, then none of the columns have any captions.

To establish what classes apply to a column

Invoke the `getCaptionClasses()` method with the index of the column in question, and an object implementing the `DOMTokenList` interface, associated with an anonymous empty string. The index i must be in the range $0 \leq i < N$, where N is the total number of columns. The tokens contained in the string underlying `DOMTokenList` object when the method returns represent the classes that apply to the given column. If the method is not defined, no classes apply to the column.

To establish whether a column should be initially included in the visible columns

Check whether the `initially-hidden` class applies to the column. If it does, then the column should not be initially included; if it does not, then the column should be initially included.

To establish whether the data can be sorted relative to a particular column

Check whether the sortable class applies to the column. If it does, then the user agent should offer the user the option to have the data displayed sorted by that column; if it does not, then the user agent must not allow the user to ask for the data to be sorted by that column.

To establish if a column is a sorted column

If the user agent can handle multiple columns being marked as sorted simultaneously: Check whether the sorted class applies to the column. If it does, then that column is the sorted column, otherwise it is not.

If the user agent can only handle one column being marked as sorted at a time: Check each column in turn, starting with the first one, to see whether the sorted class applies to that column. The first column that has that class, if any, is the sorted column. If none of the columns have that class, there is no sorted column.

To establish the sort direction of a sorted column

Check whether the reversed class applies to the column. If it does, then the sort direction is descending (down; first rows have the highest values), otherwise it is ascending (up; first rows have the lowest values).

To determine the total number of rows

Determine the number of rows for the root of the data grid, and determine the number of child rows for each open row. The total number of rows is the sum of all these numbers.

To determine the number of rows for the root of the data grid

Invoke the getRowCount() method with a RowSpecification object representing the empty path as its only argument. The return value is the number of rows at the top level of the data grid. If the return value of the method is negative, not an integer, or simply not a numeric type, or if the method is not defined, then zero must be used instead.

To determine the number of child rows for a row

Invoke the getRowCount() method with a RowSpecification object representing the path to the row in question. The return value is the number of child rows for the given row. If the return value of the method is negative, not an integer, or simply not a numeric type, or if the method is not defined, then zero must be used instead.

To determine what order to render rows in

Invoke the getChildAtPosition() method with a RowSpecification object representing the path to the parent of the rows that are being rendered as the first argument, and the position that is being rendered as the second argument. The return value is the index of the row to render in that position.

If the rows are:

1. Row "0"
 1. Row "0,0"
 2. Row "0,1"
2. Row "1"
 1. Row "1,0"

2. Row "1,1"

...and the getChildAtPosition() method is implemented as follows:

```
function getChildAtPosition(parent, child) {  
    // always return the reverse order  
    return getRowCount(parent)-child-1;  
}
```

...then the rendering would actually be:

1. Row "1"

 1. Row "1,1"

 2. Row "1,0"

2. Row "0"

 1. Row "0,1"

 2. Row "0,0"

If the return value of the method is negative, larger than the number of rows that the getRowCount() method reported for that parent, not an integer, or simply not a numeric type, then the entire data grid should be disabled. Similarly, if the method returns the same value for two or more different values for the second argument (with the same first argument, and assuming that the data grid hasn't had relevant update methods invoked in the meantime), then the data grid should be disabled. Instead of disabling the data grid, the user agent may act as if the getChildAtPosition() method was not defined on the data provider (thus disabling sorting for that data grid, but still letting the user interact with the data). If the method is not defined, then the return value must be assumed to be the same as the second argument (an identity transform; the data is rendered in its natural order).

To establish what classes apply to a row

Invoke the getRowClasses() method with a RowSpecification object representing the row in question, and a DOMTokenList associated with an empty string. The tokens contained in the DOMTokenList object's underlying string when the method returns represent the classes that apply to the row in question. If the method is not defined, no classes apply to the row.

To establish whether a row is a data row or a special row

Examine the classes that apply to the row. If the header class applies to the row, then it is not a data row, it is a subheading. The data from the first cell of the row is the text of the subheading, the rest of the cells must be ignored. Otherwise, if the separator class applies to the row, then in the place of the row, a separator should be shown. Otherwise, if the selectable-separator class applies to the row, then the row should be a data row, but represented as a separator. (The difference between a separator and a selectable-separator is that the former is not an item that can be actually selected, whereas the second can be selected and thus has a context menu that applies to it, and so forth.) For both kinds of separator rows, the data of the rows' cells must all be ignored. If none of those three classes apply then the row is a simple data row.

To establish whether a row is openable

Determine the number of child rows for that row. If there are one or more child rows, then the row is openable.

To establish whether a row should be initially open or closed

If the row is openable (page 393), examine the classes that apply to the row. If the `initially-open` class applies to the row, then it should be initially open. Otherwise, if the `initially-closed` class applies to the row, then it must be initially closed.

Otherwise, if neither class applies to the row, or if the row is not openable, then the initial state of the row should be based on platform conventions.

To obtain a URL (page 52) identifying an image representing a row

Invoke the `getRowImage()` method with a `RowSpecification` object representing the row in question. The return value is a URL (page 52). Immediately resolve (page 55) that URL as if it came from an attribute of the `datagrid` element to obtain an absolute URL (page 56) identifying the image that represents the row. If the method returns the empty string, null, or if the method is not defined, then the row has no associated image.

To obtain a context menu appropriate for a particular row

Invoke the `getRowMenu()` method with a `RowSpecification` object representing the row in question. The return value is a reference to an object implementing the `HTMLMenuItemElement` interface, i.e. a menu element DOM node. (This element must then be interpreted as described in the section on context menus to obtain the actual context menu to use.) If the method returns something that is not an `HTMLMenuItemElement`, or if the method is not defined, then the row has no associated context menu. User agents may provide their own default context menu, and may add items to the author-provided context menu. For example, such a menu could allow the user to change the presentation of the `datagrid` element.

To establish the value of a particular cell

Invoke the `getCellData()` method with the first argument being a `RowSpecification` object representing the row of the cell in question and the second argument being the index of the cell's column. The second argument must be a non-negative integer less than the total number of columns. The return value is the value of the cell. If the return value is null or the empty string, or if the method is not defined, then the cell has no data. (For progress bar cells, the cell's value must be further interpreted, as described below.)

To establish what classes apply to a cell

Invoke the `getCellClasses()` method with the first argument being a `RowSpecification` object representing the row of the cell in question, the second argument being the index of the cell's column, and the third being an object implementing the `DOMTokenList` interface, associated with an empty string. The second argument must be a non-negative integer less than the total number of columns. The tokens contained in the `DOMTokenList` object's underlying string when the method returns represent the classes that apply to that cell. If the method is not defined, no classes apply to the cell.

To establish the type of a cell

Examine the classes that apply to the cell. If the `progress` class applies to the cell, it is a progress bar. Otherwise, if the `cyclable` class applies to the cell, it is a cycling cell whose value can be cycled between multiple states. Otherwise, none of these classes apply, and the cell is a simple text cell.

To establish the value of a progress bar cell

If the value x of the cell is a string that can be converted to a floating-point number (page 34) in the range $0.0 \leq x \leq 1.0$, then the progress bar has that value (0.0 means no progress, 1.0 means complete). Otherwise, the progress bar is an indeterminate progress bar.

To establish how a simple text cell should be presented

Check whether one of the checked, unchecked, or indeterminate classes applies to the cell. If any of these are present, then the cell has a checkbox, otherwise none are present and the cell does not have a checkbox. If the cell has no checkbox, check whether the editable class applies to the cell. If it does, then the cell value is editable, otherwise the cell value is static.

To establish the state of a cell's checkbox, if it has one

Check whether the checked class applies to the cell. If it does, the cell is checked. Otherwise, check whether the unchecked class applies to the cell. If it does, the cell is unchecked. Otherwise, the indeterminate class applies to the cell and the cell's checkbox is in an indeterminate state. When the indeterminate class applies to the cell, the checkbox is a tristate checkbox, and the user can set it to the indeterminate state. Otherwise, only the checked and/or unchecked classes apply to the cell, and the cell can only be toggled between those two states.

If the data provider ever raises an exception while the datagrid is invoking one of its methods, the datagrid must act, for the purposes of that particular method call, as if the relevant method had not been defined.

A RowSpecification object p with n path components passed to a method of the data provider must fulfill the constraint $0 \leq p_i < m-1$ for all integer values of i in the range $0 \leq i < n-1$, where m is the value that was last returned by the `getRowCount()` method when it was passed the RowSpecification object q with $i-1$ items, where $p_i = q_i$ for all integer values of i in the range $0 \leq i < n-1$, with any changes implied by the update methods taken into account.

The data model is considered stable: user agents may assume that subsequent calls to the data provider methods will return the same data, until one of the update methods is called on the datagrid element. If a user agent is returned inconsistent data, for example if the number of rows returned by `getRowCount()` varies in ways that do not match the calls made to the update methods, the user agent may disable the datagrid. User agents that do not disable the datagrid in inconsistent cases must honor the most recently returned values.

User agents may cache returned values so that the data provider is never asked for data that could contradict earlier data. User agents must not cache the return value of the `getRowMenu` method.

The exact algorithm used to populate the data grid is not defined here, since it will differ based on the presentation used. However, the behavior of user agents must be consistent with the descriptions above. For example, it would be non-conformant for a user agent to make cells have both a checkbox and be editable, as the descriptions above state that cells that have a checkbox cannot be edited.

4.11.2.6 Updating the datagrid

Whenever the `data` attribute is set to a new value, the datagrid must clear the current selection, remove all the displayed rows, and plan to repopulate itself using the information from the new data provider at the earliest opportunity.

There are a number of update methods that can be invoked on the `datagrid` element to cause it to refresh itself in slightly less drastic ways:

When the `updateEverything()` method is called, the user agent must repopulate the entire datagrid. If the number of rows decreased, the selection must be updated appropriately. If the number of rows increased, the new rows should be left unselected.

When the `updateRowsChanged(row, count)` method is called, the user agent must refresh the rendering of the rows starting from the row specified by `row`, and including the `count` next siblings of the row (or as many next siblings as it has, if that is less than `count`), including all descendant rows.

When the `updateRowsInserted(row, count)` method is called, the user agent must assume that `count` new rows have been inserted, such that the first new row is identified by `row`. The user agent must update its rendering and the selection accordingly. The new rows should not be selected.

When the `updateRowsRemoved(row, count)` method is called, the user agent must assume that `count` rows have been removed starting from the row that used to be identifier by `row`. The user agent must update its rendering and the selection accordingly.

The `updateRowChanged(row)` method must be exactly equivalent to calling `updateRowsChanged(row, 1)`.

When the `updateColumnChanged(column)` method is called, the user agent must refresh the rendering of the specified column `column`, for all rows.

When the `updateCellChanged(row, column)` method is called, the user agent must refresh the rendering of the cell on row `row`, in column `column`.

Any effects the update methods have on the datagrid's selection is not considered a change to the selection, and must therefore not fire the `select` event.

These update methods should be called only by the data provider, or code acting on behalf of the data provider. In particular, calling the `updateRowsInserted()` and `updateRowsRemoved()` methods without actually inserting or removing rows from the data provider is likely to result in inconsistent renderings (page 394), and the user agent is likely to disable the data grid.

4.11.2.7 Requirements for interactive user agents

This section only applies to interactive user agents.

If the `datagrid` element has a `disabled` attribute, then the user agent must disable the `datagrid`, preventing the user from interacting with it. The `datagrid` element should still continue to update itself when the data provider signals changes to the data, though. Obviously, conformance requirements stating that `datagrid` elements must react to users in particular ways do not apply when one is disabled.

If a row is openable (page 393), then the user agent should offer to the user the option of toggling the row's open/closed state. When a row's open/closed state changes, the user agent must update the rendering to match the new state.

If a cell is a cell whose value can be cycled between multiple states (page 393), then the user agent should allow the user to activate the cell to cycle its value. When the user activates this "cycling" behavior of a cell, then the datagrid must invoke the data provider's `cycleCell()` method, with a `RowSpecification` object representing the cell's row as the first argument and the cell's column index as the second. The datagrid must then act as if the datagrid's `updateCellChanged()` method had been invoked with those same arguments.

When a cell has a checkbox (page 394), the user agent should allow the user to set the checkbox's state. When the user changes the state of a checkbox in such a cell, the datagrid must invoke the data provider's `setCellCheckedState()` method, with a `RowSpecification` object representing the cell's row as the first argument, the cell's column index as the second, and the checkbox's new state as the third. The state should be represented by the number 1 if the new state is checked, 0 if the new state is unchecked, and -1 if the new state is indeterminate (which must be possible only if the cell has the `indeterminate` class set). The datagrid must then act as if the datagrid's `updateCellChanged()` method had been invoked, specifying the same cell.

If a cell is editable (page 394), the user agent should allow the user to edit the data for that cell, and doing so must cause the user agent to invoke the `editCell()` method of the data provider with three arguments: a `RowSpecification` object representing the cell's row, the cell's column's index, and the new text entered by the user. The user agent must then act as if the `updateCellChanged()` method had been invoked, with the same row and column specified.

4.11.2.8 The selection

This section only applies to interactive user agents. For other user agents, the selection attribute must return null.

```
interface DataGridSelection {
    readonly attribute unsigned long length;
    [IndexGetter] RowSpecification item(in unsigned long index);
    boolean isSelected(in RowSpecification row);
    void setSelected(in RowSpecification row, in boolean selected);

    void selectAll();
    void invert();
    void clear();
};
```

Each datagrid element must keep track of which rows are currently selected. Initially no rows are selected, but this can be changed via the methods described in this section.

The selection of a datagrid is represented by its **selection** DOM attribute, which must be a `DataGridSelection` object.

`DataGridSelection` objects represent the rows in the selection. In the selection the rows must be ordered in the natural order of the data provider (and not, e.g., the rendered order).

Rows that are not rendered because one of their ancestors is closed must share the same selection state as their nearest rendered ancestor. Such rows are not considered part of the selection for the purposes of iterating over the selection.

Note: This selection API doesn't allow for hidden rows to be selected because it is trivial to create a data provider that has infinite depth, which would then require the selection to be infinite if every row, including every hidden row, was selected.

The **length** attribute must return the number of rows currently present in the selection. The **item(index)** method must return the *index*th row in the selection. If the argument is out of range (less than zero or greater than the number of selected rows minus one), then it must raise an INDEX_SIZE_ERR exception. [DOM3CORE]

The **isSelected()** method must return the selected state of the row specified by its argument. If the specified row exists and is selected, it must return true, otherwise it must return false.

The **setSelected()** method takes two arguments, *row* and *selected*. When invoked, it must set the selection state of row *row* to selected if *selected* is true, and unselected if it is false. If *row* is not a row in the data grid, the method must raise an INDEX_SIZE_ERR exception. If the specified row is not rendered because one of its ancestors is closed, the method must do nothing.

The **selectAll()** method must mark all the rows in the data grid as selected. After a call to selectAll(), the length attribute will return the number of rows in the data grid, not counting children of closed rows.

The **invert()** method must cause all the rows in the selection that were marked as selected to now be marked as not selected, and vice versa.

The **clear()** method must mark all the rows in the data grid to be marked as not selected. After a call to clear(), the length attribute will return zero.

If the datagrid element has a **multiple** attribute, then the user agent should allow the user to select any number of rows (zero or more). If the attribute is not present, then the user agent must not allow the user to select more than a single row at a time, and selecting another one must unselect all the other rows.

Note: This only applies to the user. Scripts can select multiple rows even when the multiple attribute is absent.

Whenever the selection of a datagrid changes, whether due to the user interacting with the element, or as a result of calls to methods of the selection object, a **select** event that bubbles but is not cancelable must be fired on the datagrid element. If changes are made to the selection via calls to the object's methods during the execution of a script, then the select events must be coalesced into one, which must then be fired when the script execution has completed.

Note: The DataGridSelection interface has no relation to the Selection interface.

4.11.2.9 Columns and captions

This section only applies to interactive user agents.

Each datagrid element must keep track of which columns are currently being rendered. User agents should initially show all the columns except those with the `initially-hidden` class, but may allow users to hide or show columns. User agents should initially display the columns in the order given by the data provider, but may allow this order to be changed by the user.

If columns are not being used, as might be the case if the data grid is being presented in an icon view, or if an overview of data is being read in an aural context, then the text of the first column of each row should be used to represent the row.

If none of the columns have any captions (i.e. if the data provider does not provide a `getCaptionText()` method), then user agents may avoid showing the column headers at all. This may prevent the user from performing actions on the columns (such as reordering them, changing the sort column, and so on).

Note: *Whatever the order used for rendering, and irrespective of what columns are being shown or hidden, the "first column" as referred to in this specification is always the column with index zero, and the "last column" is always the column with the index one less than the value returned by the `getColumnCount()` method of the data provider.*

If a column is sortable (page 391), then the user agent should allow the user to request that the data be sorted using that column. When the user does so, then the datagrid must invoke the data provider's `toggleColumnSortState()` method, with the column's index as the only argument. The datagrid must *then* act as if the datagrid's `updateEverything()` method had been invoked.

4.11.3 The command element

Categories

Metadata content (page 93).
Phrasing content (page 94).

Contexts in which this element may be used:

Where metadata content (page 93) is expected.
Where phrasing content (page 94) is expected.

Content model:

Empty.

Element-specific attributes:

`type`
`label`
`icon`
`disabled`
`checked`
`radiogroup`
`default`

Also, the `title` attribute has special semantics on this element.

DOM interface:

```
interface HTMLCommandElement : HTMLElement {  
    attribute DOMString type;  
    attribute DOMString label;  
    attribute DOMString icon;  
    attribute boolean disabled;  
    attribute boolean checked;  
    attribute DOMString radiogroup;  
    attribute boolean default;  
    void click(); // shadows HTMLElement.click()  
};
```

The Command interface must also be implemented by this element.

The command element represents a command that the user can invoke.

The **type** attribute indicates the kind of command: either a normal command with an associated action, or a state or option that can be toggled, or a selection of one item from a list of items.

The attribute's value must be either "command", "checkbox", or "radio", denoting each of these three types of commands respectively. The attribute may also be omitted if the element is to represent the first of these types, a simple command.

The **label** attribute gives the name of the command, as shown to the user.

The **title** attribute gives a hint describing the command, which might be shown to the user to help him.

The **icon** attribute gives a picture that represents the command. If the attribute is specified, the attribute's value must contain a valid URL (page 52).

The **disabled** attribute is a boolean attribute (page 32) that, if present, indicates that the command is not available in the current state.

Note: *The distinction between disabled and hidden is subtle. A command should be disabled if, in the same context, it could be enabled if only certain aspects of the situation were changed. A command should be marked as hidden if, in that situation, the command will never be enabled. For example, in the context menu for a water faucet, the command "open" might be disabled if the faucet is already open, but the command "eat" would be marked hidden since the faucet could never be eaten.*

The **checked** attribute is a boolean attribute (page 32) that, if present, indicates that the command is selected.

The **radiogroup** attribute gives the name of the group of commands that will be toggled when the command itself is toggled, for commands whose type attribute has the value "radio". The scope of the name is the child list of the parent element.

If the command element is used when generating a context menu, then the **default** attribute indicates, if present, that the command is the one that would have been invoked if the user had directly activated the menu's subject instead of using its context menu. The default attribute is a boolean attribute (page 32).

- ** Need an example that shows an element that, if double-clicked, invokes an action, but that also has a context menu, showing the various command attributes off, and that has a default command.

The **type**, **label**, **icon**, **disabled**, **checked**, **radiogroup**, and **default** DOM attributes must reflect (page 67) the respective content attributes of the same name.

The **click()** method's behavior depends on the value of the type attribute of the element, as follows:

→ **If the type attribute has the value checkbox**

If the element has a checked attribute, the UA must remove that attribute.

Otherwise, the UA must add a checked attribute, with the literal value checked. The UA must then fire a click event (page 436) at the element.

→ **If the type attribute has the value radio**

If the element has a parent, then the UA must walk the list of child nodes of that parent element, and for each node that is a command element, if that element has a radiogroup attribute whose value exactly matches the current element's (treating missing radiogroup attributes as if they were the empty string), and has a checked attribute, must remove that attribute and fire a click event (page 436) at the element.

Then, the element's checked attribute attribute must be set to the literal value checked and a click event must be fired at the element.

→ **Otherwise**

The UA must fire a click event (page 436) at the element.

Note: Firing a synthetic click event at the element does not cause any of the actions described above to happen.

- ** should change all the above so it actually is just triggered by a click event, then we could
** remove the shadowing click() method and rely on actual events.

- ** Need to define the command="" attribute

Note: command elements are not rendered unless they form part of a menu (page 403).

4.11.4 The bb element

Categories

Phrasing content (page 94).
Interactive content (page 95).

Contexts in which this element may be used:

Where phrasing content (page 94) is expected.

Content model:

Phrasing content (page 94), but there must be no interactive content (page 95) descendant.

Element-specific attributes:

type

DOM interface:

```
interface HTMLBrowserButtonElement : HTMLElement {  
    attribute DOMString type;  
    readonly attribute boolean supported;  
    readonly attribute boolean disabled;  
};
```

The Command interface must also be implemented by this element.

The bb element represents a user agent command that the user can invoke.

The **type** attribute indicates the kind of command. The **type** attribute is an enumerated attribute (page 51). The following table lists the keywords and states for the attribute — the keywords in the left column map to the states listed in the cell in the second column on the same row as the keyword.

Keyword	State
makeapp	<i>make application (page 402)</i>

The *missing value default* state is the *null (page 401)* state.

Each state has an *action* and a *relevance*, defined in the following sections.

When the attribute is in the **null** state, the *action* is to not do anything, and the *relevance* is unconditionally false.

A bb element whose type attribute is in a state whose *relevance* is true must be enabled. Conversely, a bb element whose type attribute is in a state whose *relevance* is false must be disabled.

If a bb element is enabled, it must match the `:enabled` pseudo-class; otherwise, it must match the `:disabled` pseudo-class.

User agents should allow users to invoke bb elements when they are enabled. When a user invokes a bb element, its type attribute's state's *action* must be invoked.

When the element has no descendant element children and has no descendant text node children of non-zero length, the element represents a browser button with a

user-agent-defined icon or text representing the type attribute's state's *action* and *relevance* (enabled vs disabled). Otherwise, the element represents its descendants.

The **type** DOM attribute must reflect (page 67) the content attribute of the same name.

The **supported** DOM attribute must return true if the type attribute is in a state other than the *null* (page 401) state and the user agent supports that state's *action* (i.e. when the attribute's value is one that the user agent recognises and supports), and false otherwise.

The **disabled** DOM attribute must return true if the element is disabled, and false otherwise (i.e. it returns the opposite of the type attribute's state's *relevance*).

4.11.4.1 Browser button types

4.11.4.1.1 The make application state

Some user agents support making sites accessible as independent applications, as if they were not Web sites at all. The *make application* (page 402) state exists to allow Web pages to offer themselves to the user as targets for this mode of operation.

The *action* of the *make application* (page 402) state is to confirm the user's intent to use the current site in a standalone fashion, and, provided the user's intent is confirmed, offer the user a way to make the resource identified by the document's address available in such a fashion.

⚠Warning! *The confirmation is needed because it is relatively easy to trick users into activating buttons. The confirmation could, e.g. take the form of asking the user where to "save" the application, or non-modal information panel that is clearly from the user agent and gives the user the opportunity to drag an icon to their system's application launcher.*

The *relevance* of the *make application* (page 402) state is false if the user agent is already handling the site in such a fashion, or if the user agent doesn't support making the site available in that fashion, and true otherwise.

In the following example, a few links are listed on an application's page, to allow the user perform certain actions, including making the application standalone:

```
<menu>
  <li><a href="settings.html"
    onclick="panels.show('settings')">Settings</a>
    <li><bb type="makeapp">Download standalone application</bb>
    <li><a href="help.html" onclick="panels.show('help')">Help</a>
    <li><a href="logout.html" onclick="panels.show('logout')">Sign out</a>
</menu>
```

With the following stylesheet, it could be made to look like a single line of text with vertical bars separating the options, with the "make app" option disappearing when it's not supported or relevant:

```
menu li { display: none; }
menu li:enabled { display: inline; }
menu li:not(:first-child)::before { content: ' | ';
```

This could look like this:

[Settings](#) | [Download standalone application](#) | [Help](#) | [Sign out](#)

The following example shows another way to do the same thing as the previous one, this time not relying on CSS support to hide the "make app" link if it doesn't apply:

```
<menu>
  <a href="settings.html"
  onclick="panels.show('settings')">Settings</a> |
  <bb type="makeapp" id="makeapp"> </bb>
  <a href="help.html" onclick="panels.show('help')">Help</a> |
  <a href="logout.html" onclick="panels.show('logout')">Sign out</a>
</menu>
<script>
  var bb = document.getElementById('makeapp');
  if (bb.supported && bb.enabled) {
    bb.parentNode.nextSibling.textContent = ' | ';
    bb.textContent = 'Download standalone application';
  } else {
    bb.parentNode.removeChild(bb);
  }
</script>
```

4.11.5 The menu element

Categories

Flow content (page 93).

If the element's type attribute is in the tool bar (page 404) state: Interactive content (page 95).

If there is a menu element ancestor: phrasing content (page 94).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

If there is a menu element ancestor: where phrasing content (page 94) is expected.

Content model:

Either: Zero or more li elements.

Or: Phrasing content (page 94).

Element-specific attributes:

type

label

DOM interface:

```
interface HTMLMenuElement : HTMLElement {
  attribute DOMString type;
  attribute DOMString label;
};
```

The menu element represents a list of commands.

The **type** attribute is an enumerated attribute (page 51) indicating the kind of menu being declared. The attribute has three states. The context keyword maps to the **context menu** state, in which the element is declaring a context menu. The toolbar keyword maps to the **tool bar** state, in which the element is declaring a tool bar. The attribute may also be omitted. The *missing value default* is the **list** state, which indicates that the element is merely a list of commands that is neither declaring a context menu nor defining a tool bar.

If a menu element's type attribute is in the context menu (page 404) state, then the element represents the commands of a context menu, and the user can only interact with the commands if that context menu is activated.

If a menu element's type attribute is in the tool bar (page 404) state, then the element represents a list of active commands that the user can immediately interact with.

If a menu element's type attribute is in the list (page 404) state, then the element either represents an unordered list of items (each represented by an **li** element), each of which represents a command that the user can perform or activate, or, if the element has no **li** element children, flow content (page 93) describing available commands.

The **label** attribute gives the label of the menu. It is used by user agents to display nested menus in the UI. For example, a context menu containing another menu would use the nested menu's label attribute for the submenu's menu label.

The **type** and **label** DOM attributes must reflect (page 67) the respective content attributes of the same name.

4.11.5.1 Introduction

This section is non-normative.

**

...

4.11.5.2 Building menus and tool bars

A menu (or tool bar) consists of a list of zero or more of the following components:

- Commands (page 406), which can be marked as default commands
- Separators
- Other menus (which allows the list to be nested)

The list corresponding to a particular menu element is built by iterating over its child nodes. For each child node in tree order (page 24), the required behavior depends on what the node is, as follows:

↳ An element that defines a command (page 406)

Append the command to the menu, respecting its facets (page 407). If the element is a command element with a default attribute, mark the command as being a default command.

- ↪ **An hr element**
- ↪ **An option element that has a value attribute set to the empty string, and has a disabled attribute, and whose textContent consists of a string of one or more hyphens (U+002D HYPHEN-MINUS)**
 - Append a separator to the menu.
- ↪ **An li element**
 - Iterate over the children of the li element.
- ↪ **A menu element with no label attribute**
- ↪ **A select element**
 - Append a separator to the menu, then iterate over the children of the menu or select element, then append another separator.
- ↪ **A menu element with a label attribute**
- ↪ **An optgroup element**
 - Append a submenu to the menu, using the value of the element's label attribute as the label of the menu. The submenu must be constructed by taking the element and creating a new menu for it using the complete process described in this section.
- ↪ **Any other node**
 - Ignore (page 24) the node.

** We should support label in the algorithm above -- just iterate through the contents like with li, to support input elements in label elements. Also, optgroup elements without labels should be ignored (maybe? or at least should say they have no label so that they are dropped below), and select elements inside label elements may need special processing.

Once all the nodes have been processed as described above, the user agent must post-process the menu as follows:

1. Any menu item with no label, or whose label is the empty string, must be removed.
2. Any sequence of two or more separators in a row must be collapsed to a single separator.
3. Any separator at the start or end of the menu must be removed.

4.11.5.3 Context menus

The **contextmenu** attribute gives the element's context menu (page 405). The value must be the ID of a menu element in the DOM. If the node that would be obtained by invoking the `getElementById()` method using the attribute's value as the only argument is null or not a menu element, then the element has no assigned context menu. Otherwise, the element's assigned context menu is the element so identified.

When an element's context menu is requested (e.g. by the user right-clicking the element, or pressing a context menu key), the UA must fire a `contextmenu` event (page 436) on the element for which the menu was requested.

Note: Typically, therefore, the firing of the contextmenu event will be the default action of a mouseup or keyup event. The exact sequence of events is UA-dependent, as it will vary based on platform conventions.

The default action of the contextmenu event depends on whether the element has a context menu assigned (using the contextmenu attribute) or not. If it does not, the default action must be for the user agent to show its default context menu, if it has one.

- ** Context menus should inherit (so clicking on a span in a paragraph with a context menu should show the menu).

If the element *does* have a context menu assigned, then the user agent must fire a show event (page 436) on the relevant menu element.

The default action of *this* event is that the user agent must show a context menu built (page 404) from the menu element.

The user agent may also provide access to its default context menu, if any, with the context menu shown. For example, it could merge the menu items from the two menus together, or provide the page's context menu as a submenu of the default menu.

If the user dismisses the menu without making a selection, nothing in particular happens.

If the user selects a menu item that represents a command, then the UA must invoke that command's Action (page 407).

Context menus must not, while being shown, reflect changes in the DOM; they are constructed as the default action of the show event and must remain like that until dismissed.

User agents may provide means for bypassing the context menu processing model, ensuring that the user can always access the UA's default context menus. For example, the user agent could handle right-clicks that have the Shift key depressed in such a way that it does not fire the contextmenu event and instead always shows the default context menu.

The **contextMenu** attribute must reflect (page 67) the contextmenu content attribute.

4.11.5.4 Toolbars

Toolbars are a kind of menu that is always visible.

When a menu element has a type attribute with the value toolbar, then the user agent must build (page 404) the menu for that menu element and render it in the document in a position appropriate for that menu element.

The user agent must reflect changes made to the menu's DOM immediately in the UI.

4.11.6 Commands

A **command** is the abstraction behind menu items, buttons, and links. Once a command is defined, other parts of the interface can refer to the same command, allowing many access points to a single feature to share aspects such as the disabled state.

Commands are defined to have the following **facets**:

Type

The kind of command: "command", meaning it is a normal command; "radio", meaning that triggering the command will, amongst other things, set the Checked State (page 407) to true (and probably uncheck some other commands); or "checkbox", meaning that triggering the command will, amongst other things, toggle the value of the Checked State (page 407).

ID

The name of the command, for referring to the command from the markup or from script. If a command has no ID, it is an **anonymous command**.

Label

The name of the command as seen by the user.

Hint

A helpful or descriptive string that can be shown to the user.

Icon

An absolute URL (page 56) identifying a graphical image that represents the action. A command might not have an Icon.

Hidden State

Whether the command is hidden or not (basically, whether it should be shown in menus).

Disabled State

Whether the command is relevant and can be triggered or not.

Checked State

Whether the command is checked or not.

Action

The actual effect that triggering the command will have. This could be a scripted event handler, a URL (page 52) to which to navigate (page 473), or a form submission.

Triggers

The list of elements that can trigger the command. The element defining a command is always in the list of elements that can trigger the command. For anonymous commands, only the element defining the command is on the list, since other elements have no way to refer to it.

Commands are represented by elements in the DOM. Any element that can define a command also implements the Command interface:

** Actually even better would be to just mix it straight into those interfaces somehow.

```
[NoInterfaceObject] interface Command {  
    readonly attribute DOMString commandType;  
    readonly attribute DOMString id;  
    readonly attribute DOMString label;  
    readonly attribute DOMString title;  
    readonly attribute DOMString icon;
```

```
readonly attribute boolean hidden;
readonly attribute boolean disabled;
readonly attribute boolean checked;
void click();
readonly attribute HTMLCollection triggers;
readonly attribute Command command;
};
```

The Command interface is implemented by any element capable of defining a command. (If an element can define a command, its definition will list this interface explicitly.) All the attributes of the Command interface are read-only. Elements implementing this interface may implement other interfaces that have attributes with identical names but that are mutable; in bindings that flatten all supported interfaces on the object, the mutable attributes must shadow the readonly attributes defined in the Command interface.

The **commandType** attribute must return a string whose value is either "command", "radio", or "checked", depending on whether the Type (page 407) of the command defined by the element is "command", "radio", or "checked" respectively. If the element does not define a command, it must return null.

The **id** attribute must return the command's ID (page 407), or null if the element does not define a command or defines an anonymous command (page 407). This attribute will be shadowed by the id DOM attribute on the HTMLElement interface.

The **label** attribute must return the command's Label (page 407), or null if the element does not define a command or does not specify a Label (page 407). This attribute will be shadowed by the label DOM attribute on option and command elements.

The **title** attribute must return the command's Hint (page 407), or null if the element does not define a command or does not specify a Hint (page 407). This attribute will be shadowed by the title DOM attribute on the HTMLElement interface.

The **icon** attribute must return the absolute URL (page 56) of the command's Icon (page 407). If the element does not specify an icon, or if the element does not define a command, then the attribute must return null. This attribute will be shadowed by the icon DOM attribute on command elements.

The **hidden** attribute must return true if the command's Hidden State (page 407) is that the command is hidden, and false if it is that the command is not hidden. If the element does not define a command, the attribute must return false. This attribute will be shadowed by the hidden DOM attribute on the HTMLElement interface.

The **disabled** attribute must return true if the command's Disabled State (page 407) is that the command is disabled, and false if the command is not disabled. This attribute is not affected by the command's Hidden State (page 407). If the element does not define a command, the attribute must return false. This attribute will be shadowed by the disabled attribute on button, input, option, and command elements.

The **checked** attribute must return true if the command's Checked State (page 407) is that the command is checked, and false if it is that the command is not checked. If the element does not define a command, the attribute must return false. This attribute will be shadowed by the checked attribute on input and command elements.

The `click()` method must trigger the Action (page 407) for the command. If the element does not define a command, this method must do nothing. This method will be shadowed by the `click()` method on HTML elements (page 23), and is included only for completeness.

The `triggers` attribute must return a list containing the elements that can trigger the command (the command's Triggers (page 407)). The list must be live (page 24). While the element does not define a command, the list must be empty.

The `commands` attribute of the document's `HTMLDocument` interface must return an `HTMLCollection` rooted at the `Document` node, whose filter matches only elements that define commands and have IDs.

The following elements can define commands: `a`, `button`, `input`, `option`, `command`, `bb`.

4.11.6.1 Using the `a` element to define a command

An `a` element with an `href` attribute defines a command (page 406).

The Type (page 407) of the command is "command".

The ID (page 407) of the command is the value of the `id` attribute of the element, if the attribute is present and not empty. Otherwise the command is an anonymous command (page 407).

The Label (page 407) of the command is the string given by the element's `textContent` DOM attribute.

The Hint (page 407) of the command is the value of the `title` attribute of the element. If the attribute is not present, the Hint (page 407) is the empty string.

The Icon (page 407) of the command is the absolute URL (page 56) obtained from resolving (page 55) the value of the `src` attribute of the first `img` element descendant of the element, if there is such an element and resolving its attribute is successful. Otherwise, there is no Icon (page 407) for the command.

The Hidden State (page 407) of the command is true (`hidden`) if the element has a `hidden` attribute, and false otherwise.

The Disabled State (page 407) facet of the command is always false. (The command is always enabled.)

The Checked State (page 407) of the command is always false. (The command is never checked.)

The Action (page 407) of the command is to fire a `click` event (page 436) at the element.

4.11.6.2 Using the `button` element to define a command

A `button` element always defines a command (page 406).

The Type (page 407), ID (page 407), Label (page 407), Hint (page 407), Icon (page 407), Hidden State (page 407), Checked State (page 407), and Action (page 407) facets of the command are determined as for `a` elements (page 409) (see the previous section).

The Disabled State (page 407) of the command mirrors the disabled (page 362) state of the button.

4.11.6.3 Using the input element to define a command

An input element whose type attribute is in one of the Submit Button (page 338), Reset Button (page 340), Button (page 341), Radio Button (page 336), or Checkbox (page 335) states defines a command (page 406).

The Type (page 407) of the command is "radio" if the type attribute is in the Radio Button state, "checkbox" if the type attribute is in the Checkbox state, and "command" otherwise.

The ID (page 407) of the command is the value of the id attribute of the element, if the attribute is present and not empty. Otherwise the command is an anonymous command (page 407).

The Label (page 407) of the command depends on the Type of the command:

If the Type (page 407) is "command", then it is the string given by the value attribute, if any, and a UA-dependent value that the UA uses to label the button itself if the attribute is absent.

Otherwise, the Type (page 407) is "radio" or "checkbox". If the element is a labeled control (page 317), the textContent of the first label element in tree order (page 24) whose labeled control (page 317) is the element in question is the Label (page 407) (in DOM terms, this the string given by `element.labels[0].textContent`). Otherwise, the value of the value attribute, if present, is the Label (page 407). Otherwise, the Label (page 407) is the empty string.

The Hint (page 407) of the command is the value of the title attribute of the input element. If the attribute is not present, the Hint (page 407) is the empty string.

There is no Icon (page 407) for the command.

The Hidden State (page 407) of the command is true (hidden) if the element has a hidden attribute, and false otherwise.

The Disabled State (page 407) of the command mirrors the disabled (page 362) state of the control.

The Checked State (page 407) of the command is true if the command is of Type (page 407) "radio" or "checkbox" and the element is checked (page 362) attribute, and false otherwise.

The Action (page 407) of the command is to fire a click event (page 436) at the element.

4.11.6.4 Using the option element to define a command

An option element with an ancestor select element and either no value attribute or a value attribute that is not the empty string defines a command (page 406).

The Type (page 407) of the command is "radio" if the option's nearest ancestor select element has no multiple attribute, and "checkbox" if it does.

The ID (page 407) of the command is the value of the `id` attribute of the element, if the attribute is present and not empty. Otherwise the command is an anonymous command (page 407).

The Label (page 407) of the command is the value of the `option` element's `label` attribute, if there is one, or the value of the `option` element's `textContent` DOM attribute if it doesn't.

The Hint (page 407) of the command is the string given by the element's `title` attribute, if any, and the empty string if the attribute is absent.

There is no Icon (page 407) for the command.

The Hidden State (page 407) of the command is true (hidden) if the element has a `hidden` attribute, and false otherwise.

The Disabled State (page 407) of the command is true (disabled) if the element has a `disabled` attribute, and false otherwise.

The Checked State (page 407) of the command is true (checked) if the element's `selected` DOM attribute is true, and false otherwise.

The Action (page 407) of the command depends on its Type (page 407). If the command is of Type (page 407) "radio" then this must set the `selected` DOM attribute of the `option` element to true, otherwise it must toggle the state of the `selected` DOM attribute (set it to true if it is false and vice versa). Then a change event must be fired (page 436) on the `option` element's nearest ancestor `select` element (if there is one), as if the selection had been changed directly.

4.11.6.5 Using the command element to define a command

A command element defines a command (page 406).

The Type (page 407) of the command is "radio" if the command's `type` attribute is "radio", "checkbox" if the attribute's value is "checkbox", and "command" otherwise.

The ID (page 407) of the command is the value of the `id` attribute of the element, if the attribute is present and not empty. Otherwise the command is an anonymous command (page 407).

The Label (page 407) of the command is the value of the element's `label` attribute, if there is one, or the empty string if it doesn't.

The Hint (page 407) of the command is the string given by the element's `title` attribute, if any, and the empty string if the attribute is absent.

The Icon (page 407) for the command is the absolute URL (page 56) obtained from resolving (page 55) the value of the element's `icon` attribute, if it has such an attribute and resolving it is successful. Otherwise, there is no Icon (page 407) for the command.

The Hidden State (page 407) of the command is true (hidden) if the element has a `hidden` attribute, and false otherwise.

The Disabled State (page 407) of the command is true (disabled) if the element has a `disabled` attribute, and false otherwise.

The Checked State (page 407) of the command is true (checked) if the element has a checked attribute, and false otherwise.

The Action (page 407) of the command is to invoke the behavior described in the definition of the `click()` method of the `HTMLCommandElement` interface.

4.11.6.6 Using the `bb` element to define a command

A `bb` element always defines a command (page 406).

The Type (page 407) of the command is "command".

The ID (page 407) of the command is the value of the `id` attribute of the element, if the attribute is present and not empty. Otherwise the command is an anonymous command (page 407).

The Label (page 407) of the command is the string given by the element's `textContent` DOM attribute, if that is not the empty string, or a user-agent-defined string appropriate for the `bb` element's type attribute's state.

The Hint (page 407) of the command is the value of the `title` attribute of the element. If the attribute is not present, the Hint (page 407) is a user-agent-defined string appropriate for the `bb` element's type attribute's state.

The Icon (page 407) of the command is the absolute URL (page 56) obtained from resolving (page 55) the value of the `src` attribute of the first `img` element descendant of the element, if there is such an element and resolving its attribute is successful. Otherwise, the Icon (page 407) is a user-agent-defined image appropriate for the `bb` element's type attribute's state.

The Hidden State (page 407) facet of the command is true (hidden) if the `bb` element's type attribute's state is `null` (page 401) or if the element has a `hidden` attribute, and false otherwise.

The Disabled State (page 407) facet of the command is true if the `bb` element's type attribute's state's `relevance` is false, and true otherwise.

The Checked State (page 407) of the command is always false. (The command is never checked.)

The Action (page 407) of the command is to perform the `action` of the `bb` element's type attribute's state.

4.12 Miscellaneous elements

4.12.1 The `legend` element

Categories

None.

Contexts in which this element may be used:

As the first child of a `fieldset` element.

As the first child of a `details` element.

As a child of a figure element, if there are no other legend element children of that element.

Content model:

Phrasing content (page 94).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The legend element represents a title or explanatory caption for the rest of the contents of the legend element's parent element.

4.12.2 The div element

Categories

Flow content (page 93).

Contexts in which this element may be used:

Where flow content (page 93) is expected.

Content model:

Flow content (page 93).

Element-specific attributes:

None.

DOM interface:

Uses HTMLElement.

The div element represents nothing at all. It can be used with the class, lang/xml:lang, and title attributes to mark up semantics common to a group of consecutive elements.

- ** Allowing div elements to contain phrasing content makes it easy for authors to abuse div, using it with the class="" attribute to the point of not having any other elements in the markup. This is a disaster from an accessibility point of view, and it would be nice if we could somehow make such pages non-compliant without preventing people from using divs as the extension mechanism that they are, to handle things the spec can't otherwise do (like making new widgets).

5 Web browsers

This section describes features that apply most directly to Web browsers. Having said that, unless specified elsewhere, the requirements defined in this section *do* apply to all user agents, whether they are Web browsers or not.

5.1 Browsing contexts

A **browsing context** is a collection of one or more Document objects, and one or more views (page 414).

At any one time, one of the Documents in a browsing context (page 414) is the **active document**. The collection of Documents is the browsing context (page 414)'s session history (page 467).

A **view** is a user agent interface tied to a particular media used for the presentation of Document objects in some media. A view may be interactive. Each view is represented by an AbstractView object. Each view belongs to a browsing context (page 414). [DOM2VIEWS]

Note: *The document attribute of an AbstractView object representing a view (page 414) gives the Document object of the view's browsing context (page 414)'s active document (page 414). [DOM2VIEWS]*

Note: *Events that use the UIEvent interface are related to a specific view (page 414) (the view in which the event happened); the AbstractView of that view is given in the event object's view attribute. [DOM3EVENTS]*

Note: *A typical Web browser has one obvious view (page 414) per browsing context (page 414): the browser's window (screen media). If a page is printed, however, a second view becomes evident, that of the print media. The two views always share the same underlying Document, but they have a different presentation of that document. A speech browser also establishes a browsing context, one with a view in the speech media.*

Note: *A Document does not necessarily have a browsing context (page 414) associated with it. In particular, data mining tools are likely to never instantiate browsing contexts.*

The main view (page 414) through which a user primarily interacts with a user agent is the **default view**.

Note: *The default view (page 414) of a Document is given by the defaultView attribute on the Document object's DocumentView interface. [DOM3VIEWS]*

When a browsing context (page 414) is first created, it must be created with a single Document in its session history, whose address is about:blank, which is marked as being an HTML document (page 76), and whose character encoding (page 79) is UTF-8. The Document must have a single child html node, which itself has a single child body node. If the browsing

context (page 414) is created specifically to be immediately navigated, then that initial navigation will have replacement enabled (page 476).

The origin (page 423) of the about:blank Document is set when the Document is created, in a manner dependent on whether the browsing context (page 414) created is a nested browsing context (page 415), as follows:

↪ **If the new browsing context (page 414) is a nested browsing context (page 415)**

The origin (page 423) of the about:blank Document is the origin (page 423) of the active document (page 414) of the new browsing context (page 414)'s parent browsing context (page 415) at the time of its creation.

↪ **If the new browsing context (page 414) is an auxiliary browsing context (page 416)**

The origin (page 423) of the about:blank Document is the origin (page 423) of the active document (page 414) of the new browsing context (page 414)'s opener browsing context (page 416) at the time of the new browsing context's creation.

↪ **Otherwise**

The origin (page 423) of the about:blank Document is a globally unique identifier assigned when the new browsing context (page 414) is created.

5.1.1 Nested browsing contexts

Certain elements (for example, iframe elements) can instantiate further browsing contexts (page 414). These are called **nested browsing contexts**. If a browsing context P has an element E in one of its Documents D that nests another browsing context C inside it, then P is said to be the **parent browsing context** of C , C is said to be a **child browsing context** of P , C is said to be **nested through** D , and E is said to be the **browsing context container** of C .

A browsing context A is said to be an ancestor of a browsing context B if there exists a browsing context A' that is a child browsing context (page 415) of A and that is itself an ancestor of B , or if there is a browsing context P that is a child browsing context (page 415) of A and that is the parent browsing context (page 415) of B .

The browsing context with no parent browsing context (page 415) is the **top-level browsing context** of all the browsing contexts nested (page 415) within it (either directly or indirectly through other nested browsing contexts).

The transitive closure of parent browsing contexts (page 415) for a nested browsing context (page 415) gives the list of **ancestor browsing contexts**.

A Document is said to be **fully active** when it is the active document (page 414) of its browsing context (page 414), and either its browsing context is a top-level browsing context (page 415), or the Document through which (page 415) that browsing context is nested (page 415) is itself fully active (page 415).

Because they are nested through an element, child browsing contexts (page 415) are always tied to a specific Document in their parent browsing context (page 415). User agents must not allow the user to interact with child browsing contexts (page 415) of elements that are in Documents that are not themselves fully active (page 415).

A nested browsing context (page 415) can have a seamless browsing context flag (page 219) set, if it is embedded through an `iframe` element with a `seamless` attribute.

5.1.1.1 Navigating nested browsing contexts in the DOM

The `top` DOM attribute on the `Window` object of a browsing context (page 414) b must return the `Window` object of its top-level browsing context (page 415) (which would be its own `Window` object if it was a top-level browsing context (page 415) itself).

The `parent` DOM attribute on the `Window` object of a browsing context (page 414) b must return the `Window` object of the parent browsing context (page 415), if there is one (i.e. if b is a child browsing context (page 415)), or the `Window` object of the browsing context (page 414) b itself, otherwise (i.e. if it is a top-level browsing context (page 415)).

The `frameElement` DOM attribute on the `Window` object of a browsing context (page 414) b , on getting, must run the following algorithm:

1. If b is not a child browsing context (page 415), return null and abort these steps.
2. If the parent browsing context (page 415)'s active document (page 414) does not have the same effective script origin (page 423) as the script that is accessing the `frameElement` attribute, then throw a security exception (page 430).
3. Otherwise, return the browsing context container (page 415) for b .

5.1.2 Auxiliary browsing contexts

It is possible to create new browsing contexts that are related to a top level browsing context without being nested through an element. Such browsing contexts are called **auxiliary browsing contexts**. Auxiliary browsing contexts are always top-level browsing contexts (page 415).

An auxiliary browsing context (page 416) has an **opener browsing context**, which is the browsing context (page 414) from which the auxiliary browsing context (page 416) was created, and it has a **furthest ancestor browsing context**, which is the top-level browsing context (page 415) of the opener browsing context (page 416) when the auxiliary browsing context (page 416) was created.

5.1.2.1 Navigating auxiliary browsing contexts in the DOM

The `opener` DOM attribute on the `Window` object must return the `Window` object of the browsing context (page 414) from which the current browsing context was created (its opener browsing context (page 416)), if there is one and it is still available.

5.1.3 Secondary browsing contexts

User agents may support **secondary browsing contexts**, which are browsing contexts (page 414) that form part of the user agent's interface, apart from the main content area.

5.1.4 Security

A browsing context (page 414) *A* is **allowed to navigate** a second browsing context (page 414) *B* if one of the following conditions is true:

- Either the origin (page 423) of the active document (page 414) of *A* is the same (page 426) as the origin (page 423) of the active document (page 414) of *B*, or
- The browsing context *A* is a nested browsing context (page 415) and its top-level browsing context (page 415) is *B*, or
- The browsing context *B* is an auxiliary browsing context (page 416) and *A* is allowed to navigate (page 417) *B*'s opener browsing context (page 416), or
- The browsing context *B* is not a top-level browsing context (page 415), but there exists an ancestor browsing context (page 415) of *B* whose active document (page 414) has the same (page 426) origin (page 423) as the active document (page 414) of *A* (possibly in fact being *A* itself).

5.1.5 Groupings of browsing contexts

Each browsing context (page 414) is defined as having a list of zero or more **directly reachable browsing contexts**. These are:

- All the browsing context (page 414)'s child browsing contexts (page 415).
- The browsing context (page 414)'s parent browsing context (page 415).
- All the browsing contexts (page 414) that have the browsing context (page 414) as their opener browsing context (page 416).
- The browsing context (page 414)'s opener browsing context (page 416).

The transitive closure of all the browsing contexts (page 414) that are directly reachable browsing contexts (page 417) forms a **unit of related browsing contexts**.

Each unit of related browsing contexts (page 417) is then further divided into the smallest number of groups such that every member of each group has an effective script origin (page 423) that, through appropriate manipulation of the `document.domain` attribute, could be made to be the same as other members of the group, but could not be made the same as members of any other group. Each such group is a **unit of related similar-origin browsing contexts**.

5.1.6 Browsing context names

Browsing contexts can have a **browsing context name**. By default, a browsing context has no name (its name is not set).

A **valid browsing context name** is any string with at least one character that does not start with a U+005F LOW LINE character. (Names starting with an underscore are reserved for special keywords.)

A **valid browsing context name or keyword** is any string that is either a valid browsing context name (page 417) or that is an ASCII case-insensitive (page 31) match for one of: _blank, _self, _parent, or _top.

The **rules for choosing a browsing context given a browsing context name** are as follows. The rules assume that they are being applied in the context of a browsing context (page 414).

1. If the given browsing context name is the empty string or _self, then the chosen browsing context must be the current one.
2. If the given browsing context name is _parent, then the chosen browsing context must be the *parent* browsing context (page 415) of the current one, unless there isn't one, in which case the chosen browsing context must be the current browsing context.
3. If the given browsing context name is _top, then the chosen browsing context must be the most top-level browsing context (page 415) of the current one.
4. If the given browsing context name is not _blank and there exists a browsing context whose name (page 417) is the same as the given browsing context name, and the current browsing context is allowed to navigate (page 417) that browsing context, and the user agent determines that the two browsing contexts are related enough that it is ok if they reach each other, then that browsing context must be the chosen one. If there are multiple matching browsing contexts, the user agent should select one in some arbitrary consistent manner, such as the most recently opened, most recently focused, or more closely related.
5. Otherwise, a new browsing context is being requested, and what happens depends on the user agent's configuration and/or abilities:

If the current browsing context has the sandboxed navigation browsing context flag (page 218) set.

The user agent may offer to create a new top-level browsing context (page 415) or reuse an existing top-level browsing context (page 415). If the user picks one of those options, then the designated browsing context must be the chosen one (the browsing context's name isn't set to the given browsing context name). Otherwise (if the user agent doesn't offer the option to the user, or if the user declines to allow a browsing context to be used) there must not be a chosen browsing context.

If the user agent has been configured such that in this instance it will create a new browsing context, and the browsing context is being requested as part of following a hyperlink (page 498) whose link types (page 500) include the noreferrer keyword

A new top-level browsing context (page 415) must be created. If the given browsing context name is not _blank, then the new top-level browsing context's name must be the given browsing context name (otherwise, it has no name). The chosen browsing context must be this new browsing context. If it is immediately navigated (page 473), then the navigation will be done with replacement enabled (page 476).

If the user agent has been configured such that in this instance it will create a new browsing context, and the noreferrer keyword doesn't apply

A new auxiliary browsing context (page 416) must be created, with the opener browsing context (page 416) being the current one. If the given browsing context name is not `_blank`, then the new auxiliary browsing context's name must be the given browsing context name (otherwise, it has no name). The chosen browsing context must be this new browsing context. If it is immediately navigated (page 473), then the navigation will be done with replacement enabled (page 476).

If the user agent has been configured such that in this instance it will reuse the current browsing context

The chosen browsing context is the current browsing context.

If the user agent has been configured such that in this instance it will not find a browsing context

There must not be a chosen browsing context.

User agent implementors are encouraged to provide a way for users to configure the user agent to always reuse the current browsing context.

5.2 The default view

The `AbstractView` object of default views (page 414) must also implement the `Window` and `EventTarget` interfaces.

```
[NoInterfaceObject] interface Window {
    // the current browsing context
    readonly attribute Window window;
    readonly attribute Window self;
        attribute DOMString name;
    [PutForwards=href] readonly attribute Location location;
    readonly attribute History history;
    readonly attribute UndoManager undoManager;
    Selection getSelection();

    // other browsing contexts
    readonly attribute Window frames;
    readonly attribute unsigned long length;
    [IndexGetter] Window XXX4(in unsigned long index);
    readonly attribute Window top;
    readonly attribute Window opener;
    readonly attribute Window parent;
    readonly attribute Element frameElement;
    Window open();
    Window open(in DOMString url);
    Window open(in DOMString url, in DOMString target);
    Window open(in DOMString url, in DOMString target, in DOMString
features);
    Window open(in DOMString url, in DOMString target, in DOMString
features, in DOMString replace);
```

```

// the user agent
readonly attribute Navigator navigator;
readonly attribute Storage localStorage;
readonly attribute Storage sessionStorage;
Database openDatabase(in DOMString name, in DOMString version, in
DOMString displayName, in unsigned long estimatedSize);

// user prompts
void alert(in DOMString message);
boolean confirm(in DOMString message);
DOMString prompt(in DOMString message);
DOMString prompt(in DOMString message, in DOMString default);
void print();
any showModalDialog(in DOMString url);
any showModalDialog(in DOMString url, in any arguments);
void showNotification(in DOMString title, in DOMString subtitle, in
DOMString description);
void showNotification(in DOMString title, in DOMString subtitle, in
DOMString description, in VoidCallback onclick);

// cross-document messaging
void postMessage(in DOMString message, in DOMString targetOrigin);
void postMessage(in DOMString message, in MessagePort messagePort, in
DOMString targetOrigin);

// event handler DOM attributes
attribute EventListener onabort;
attribute EventListener onbeforeunload;
attribute EventListener onblur;
attribute EventListener onchange;
attribute EventListener onclick;
attribute EventListener oncontextmenu;
attribute EventListener ondblclick;
attribute EventListener ondrag;
attribute EventListener ondragend;
attribute EventListener ondragenter;
attribute EventListener ondragleave;
attribute EventListener ondragover;
attribute EventListener ondragstart;
attribute EventListener ondrop;
attribute EventListener onerror;
attribute EventListener onfocus;
attribute EventListener onhashchange;
attribute EventListener onkeydown;
attribute EventListener onkeypress;
attribute EventListener onkeyup;
attribute EventListener onload;
attribute EventListener onmessage;
attribute EventListener onmousedown;
attribute EventListener onmousemove;

```

```

        attribute EventListener onmouseout;
        attribute EventListener onmouseover;
        attribute EventListener onmouseup;
        attribute EventListener onmousewheel;
        attribute EventListener onresize;
        attribute EventListener onscroll;
        attribute EventListener onselect;
        attribute EventListener onstorage;
        attribute EventListener onsubmit;
        attribute EventListener onunload;
    };
** // VoidCallback waiting on WebIDL

```

The **window**, **frames**, and **self** DOM attributes must all return the Window object itself.

The Window object also provides the scope for script execution. Each Document in a browsing context (page 414) has an associated **list of added properties** that, when a document is active (page 414), are available on the Document's default view (page 414)'s Window object. A Document object's list of added properties (page 421) must be empty when the Document object is created.

5.2.1 Security

User agents must raise a security exception (page 430) whenever any of the members of a Window object are accessed by scripts whose effective script origin (page 423) is not the same as the Window object's browsing context (page 414)'s active document (page 414)'s effective script origin (page 423), with the following exceptions:

- The location object
- The postMessage() method with two arguments
- The postMessage() method with three arguments
- The frames attribute
- The XXX4 method

User agents must not allow scripts to override the location object's setter.

5.2.2 APIs for creating and navigating browsing contexts by name

The **open()** method on Window objects provides a mechanism for navigating (page 473) an existing browsing context (page 414) or opening and navigating an auxiliary browsing context (page 416).

The method has four arguments, though they are all optional.

The first argument, *url*, must be a valid URL (page 52) for a page to load in the browsing context. If no arguments are provided, or if the first argument is the empty string, then the

url argument defaults to "about:blank". The argument must be resolved (page 55) to an absolute URL (page 56) (or an error) when the method is invoked.

The second argument, *target*, specifies the name (page 417) of the browsing context that is to be navigated. It must be a valid browsing context name or keyword (page 418). If fewer than two arguments are provided, then the *name* argument defaults to the value "_blank".

The third argument, *features*, has no effect and is supported for historical reasons only.

The fourth argument, *replace*, specifies whether or not the new page will replace (page 476) the page currently loaded in the browsing context, when *target* identifies an existing browsing context (as opposed to leaving the current page in the browsing context's session history (page 467)). When three or fewer arguments are provided, *replace* defaults to false.

When the method is invoked, the user agent must first select a browsing context (page 414) to navigate by applying the rules for choosing a browsing context given a browsing context name (page 418) using the *target* argument as the name and the browsing context (page 414) of the script as the context in which the algorithm is executed, unless the user has indicated a preference, in which case the browsing context to navigate may instead be the one indicated by the user.

For example, suppose there is a user agent that supports control-clicking a link to open it in a new tab. If a user clicks in that user agent on an element whose *onclick* handler uses the `window.open()` API to open a page in an iframe, but, while doing so, holds the control key down, the user agent could override the selection of the target browsing context to instead target a new tab.

Then, the user agent must navigate (page 473) the selected browsing context (page 414) to the absolute URL (page 56) (or error) obtained from resolving (page 55) *url*. If the *replace* is true, then replacement must be enabled (page 476); otherwise, it must not be enabled unless the browsing context (page 414) was just created as part of the rules for choosing a browsing context given a browsing context name (page 418). The navigation must be done with the script browsing context (page 428) of the script that invoked the method as the source browsing context (page 473).

The method must return the `Window` object of the default view of the browsing context (page 414) that was navigated, or null if no browsing context was navigated.

The `name` attribute of the `Window` object must, on getting, return the current name of the browsing context (page 414), and, on setting, set the name of the browsing context (page 414) to the new value.

Note: *The name gets reset (page 481) when the browsing context is navigated to another domain.*

5.2.3 Accessing other browsing contexts

The `length` DOM attribute on the `Window` interface must return the number of child browsing contexts (page 415) of the active (page 414) Document.

The `XXX4(index)` method must return the `index`th child browsing context (page 415) of the active (page 414) Document, sorted in document order of the elements nesting those browsing contexts.

5.3 Origin

The **origin** of a resource and the **effective script origin** of a resource are both either opaque identifiers or tuples consisting of a scheme component, a host component, a port component, and optionally extra data.

Note: *The extra data could include the certificate of the site when using encrypted connections, to ensure that if the site's secure certificate changes, the origin is considered to change as well.*

These characteristics are defined as follows:

For URLs

The origin (page 423) and effective script origin (page 423) of the URL (page 52) is whatever is returned by the following algorithm:

1. Let `url` be the URL (page 52) for which the origin (page 423) is being determined.
2. Parse (page 53) `url`.
3. If `url` does not use a server-based naming authority, or if parsing `url` failed, or if `url` is not an absolute URL (page 56), then return a new globally unique identifier.
4. Let `scheme` be the `<scheme>` (page 53) component of `url`, converted to lowercase (page 31).
5. If the UA doesn't support the protocol given by `scheme`, then return a new globally unique identifier.
6. If `scheme` is "file", then the user agent may return a UA-specific value.
7. Let `host` be the `<host>` (page 53) component of `url`.
8. Apply the IDNA ToASCII algorithm to `host`, with both the `AllowUnassigned` and `UseSTD3ASCIIRules` flags set. Let `host` be the result of the ToASCII algorithm.
If ToASCII fails to convert one of the components of the string, e.g. because it is too long or because it contains invalid characters, then return a new globally unique identifier. [RFC3490]
9. Let `host` be the result of converting `host` to lowercase (page 31).
10. If there is no `<port>` (page 53) component, then let `port` be the default port for the protocol given by `scheme`. Otherwise, let `port` be the `<port>` (page 53) component of `url`.
11. Return the tuple (`scheme, host, port`).

In addition, if the URL (page 52) is in fact associated with a Document object that was created by parsing the resource obtained from fetching URL (page 52), and this was done over a secure connection, then the server's secure certificate may be added to the origin as additional data.

For scripts

The origin (page 423) and effective script origin (page 423) of a script are determined from another resource, called the *owner*:

- ↪ **If a script is in a script element**
The owner is the Document to which the script element belongs.
- ↪ **If a script is in an event handler content attribute (page 432)**
The owner is the Document to which the attribute node belongs.
- ↪ **If a script is a function or other code reference created by another script**
The owner is the script that created it.
- ↪ **If a script is a javascript: URL (page 430) that was returned as the location of an HTTP redirect (or equivalent in other protocols)**
The owner is the URL (page 52) that redirected to the javascript: URL (page 430).
- ↪ **If a script is a javascript: URL (page 430) in an attribute**
The owner is the Document of the element on which the attribute is found.
- ↪ **If a script is a javascript: URL (page 430) in a style sheet**
The owner is the URL (page 52) of the style sheet.
- ↪ **If a script is a javascript: URL (page 430) to which a browsing context (page 414) is being navigated (page 473), the URL having been provided by the user (e.g. by using a bookmarklet)**
The owner is the Document of the browsing context (page 414)'s active document (page 414).
- ↪ **If a script is a javascript: URL (page 430) to which a browsing context (page 414) is being navigated (page 473), the URL having been declared in markup**
The owner is the Document of the element (e.g. an a or area element) that declared the URL.
- ↪ **If a script is a javascript: URL (page 430) to which a browsing context (page 414) is being navigated (page 473), the URL having been provided by script**
The owner is the script that provided the URL.

The origin (page 423) of the script is then equal to the origin (page 423) of the owner, and the effective script origin (page 423) of the script is equal to the effective script origin (page 423) of the owner.

For Document objects and images

- ↪ **If a Document is in a browsing context (page 414) whose sandboxed origin browsing context flag (page 218) is set**
The origin (page 423) is a globally unique identifier assigned when the Document is created.
- ↪ **If a Document or image was returned by the XMLHttpRequest API**
The origin (page 423) and effective script origin (page 423) are equal to the origin (page 423) and effective script origin (page 423) of the Document object that was the active document (page 414) of the browsing context of the Window

object from which the XMLHttpRequest constructor was invoked. (That is, they track the Document to which the XMLHttpRequest object's Document pointer pointed when it was created.) [XHR]

↪ **If a Document or image was generated from a javascript: URL (page 430)**

The origin (page 423) is equal to the origin (page 423) of the script of that javascript: URL (page 430).

↪ **If a Document or image was served over the network and has an address that uses a URL scheme with a server-based naming authority**

The origin (page 423) is the origin (page 423) of the address of the Document or image.

↪ **If a Document or image was generated from a data: URL that was returned as the location of an HTTP redirect (or equivalent in other protocols)**

The origin (page 423) is the origin (page 423) of the URL (page 52) that redirected to the data: URL.

↪ **If a Document or image was generated from a data: URL found in another Document or in a script**

The origin (page 423) is the origin (page 423) of the Document or script in which the data: URL was found.

↪ **If a Document has the address "about:blank"**

The origin (page 423) of the Document is the origin it was assigned when its browsing context was created (page 415).

↪ **If a Document or image was obtained in some other manner (e.g. a data: URL typed in by the user, a Document created using the createDocument() API, a data: URL returned as the location of an HTTP redirect, etc)**

The origin (page 423) is a globally unique identifier assigned when the Document or image is created.

When a Document is created, unless stated otherwise above, its effective script origin (page 423) is initialized to the origin (page 423) of the Document. However, the document.domain attribute can be used to change it.

The **Unicode serialization of an origin** is the string obtained by applying the following algorithm to the given origin (page 423):

1. If the origin (page 423) in question is not a scheme/host/port tuple, then return the empty string and abort these steps.
2. Otherwise, let *result* be the scheme part of the origin (page 423) tuple.
3. Append the string "://*" to *result*.*
4. Apply the IDNA ToUnicode algorithm to each component of the host part of the origin (page 423) tuple, and append the results — each component, in the same order, separated by U+002E FULL STOP characters (".") — to *result*.
5. If the port part of the origin (page 423) tuple gives a port that is different from the default port for the protocol given by the scheme part of the origin (page 423) tuple, then append a U+003A COLON character ":" and the given port, in base ten, to *result*.
6. Return *result*.

The **ASCII serialization of an origin** is the string obtained by applying the following algorithm to the given origin (page 423):

1. If the origin (page 423) in question is not a scheme/host/port tuple, then return the empty string and abort these steps.
2. Otherwise, let *result* be the scheme part of the origin (page 423) tuple.
3. Append the string "://*" to *result*.*
4. Apply the IDNA ToASCII algorithm the host part of the origin (page 423) tuple, with both the AllowUnassigned and UseSTD3ASCIIRules flags set, and append the results *result*.

If ToASCII fails to convert one of the components of the string, e.g. because it is too long or because it contains invalid characters, then return the empty string and abort these steps. [RFC3490]

5. If the port part of the origin (page 423) tuple gives a port that is different from the default port for the protocol given by the scheme part of the origin (page 423) tuple, then append a U+003A COLON character ":" and the given port, in base ten, to *result*.
6. Return *result*.

Two origins (page 423) are said to be the **same origin** if the following algorithm returns true:

1. Let *A* be the first origin (page 423) being compared, and *B* be the second origin (page 423) being compared.
2. If *A* and *B* are both opaque identifiers, and their value is equal, then return true.
3. Otherwise, if either *A* or *B* or both are opaque identifiers, return false.
4. If *A* and *B* have scheme components that are not identical, return false.
5. If *A* and *B* have host components that are not identical, return false.
6. If *A* and *B* have port components that are not identical, return false.
7. If either *A* or *B* have additional data, but that data is not identical for both, return false.
8. Return true.

5.3.1 Relaxing the same-origin restriction

The **domain** attribute on Document objects must be initialized to the document's domain (page 427), if it has one, and the empty string otherwise. On getting, the attribute must return its current value, unless the document was created by XMLHttpRequest, in which case it must throw an INVALID_ACCESS_ERR exception. On setting, the user agent must run the following algorithm:

1. If the document was created by XMLHttpRequest, throw an INVALID_ACCESS_ERR exception and abort these steps.

2. Apply the IDNA ToASCII algorithm to the new value, with both the AllowUnassigned and UseSTD3ASCIIRules flags set. Let *new value* be the result of the ToASCII algorithm.
If ToASCII fails to convert one of the components of the string, e.g. because it is too long or because it contains invalid characters, then throw a security exception (page 430) and abort these steps. [RFC3490]
3. If *new value* is not exactly equal to the current value of the document.domain attribute, then run these substeps:
 1. If the current value is an IP address, throw a security exception (page 430) and abort these steps.
 2. If *new value*, prefixed by a U+002E FULL STOP ("."), does not exactly match the end of the current value, throw a security exception (page 430) and abort these steps.
 4. Set the attribute's value to *new value*.
 5. Set the host part of the effective script origin (page 423) tuple of the Document to *new value*.
 6. Set the port part of the effective script origin (page 423) tuple of the Document to "manual override" (a value that, for the purposes of comparing origins (page 426), is identical to "manual override" but not identical to any other value).

The **domain** of a Document is the host part of the document's origin (page 423), if that is a scheme/host/port tuple. If it isn't, then the document does not have a domain.

Note: *The domain attribute is used to enable pages on different hosts of a domain to access each others' DOMs.*

5.4 Scripting

Various mechanisms can cause author-provided executable code to run in the context of a document. These mechanisms include, but are probably not limited to:

- Processing of script elements.
- Processing of inline javascript: URLs (e.g. the src attribute of img elements, or an @import rule in a CSS style element block).
- Event handlers, whether registered through the DOM using addEventListener(), by explicit event handler content attributes (page 432), by event handler DOM attributes (page 432), or otherwise.
- Processing of technologies like XBL or SVG that have their own scripting features.

When a script is created, it is associated with a script execution context (page 428), a script browsing context (page 428), and a script document context (page 428).

5.4.1 Script execution contexts

The **script execution context** of a script is defined when that script is created. In this specification, it is either a Window object or an empty object; other specifications might make the script execution context be some other object.

When the script execution context (page 428) of a script is an empty object, it can't do anything that interacts with the environment.

A script execution context (page 428) always has an associated browsing context (page 414), known as the **script browsing context**. If the script execution context (page 428) is a Window object, then that object's browsing context (page 414) is it. Otherwise, the script execution context (page 428) is associated explicitly with a browsing context (page 414) when it is created.

Every script whose script execution context (page 428) is a Window object is also associated with a Document object, known as its **script document context**. It is used to resolve (page 55) URLs. The document is assigned when the script is created, as with the script browsing context (page 428).

It is said that **scripting is disabled** in a script execution context (page 428) when any of the following conditions are true:

- The user agent does not support scripting.
- The user has disabled scripting for this script execution context (page 428). (User agents may provide users with the option to disable scripting globally, on a per-origin basis, or in other ways down to the granularity of individual script execution contexts (page 428).)
- The script execution context (page 428)'s associated browsing context (page 414)'s active document (page 414) has designMode enabled.
- The script execution context (page 428)'s associated browsing context (page 414) has the sandboxed scripts browsing context flag (page 219) set.

A node is said to be **without script** if either the Document object of the node (the node itself, if it is itself a Document object) does not have an associated browsing context (page 414), or scripting is disabled (page 428) in that browsing context (page 414).

A node is said to be **with script** if it is not without script (page 428).

** If you can find a better pair of terms than "with script" and "without script" let me know.
** The only things I can find that are less confusing are also way, way longer.

When a script is to be executed in a script execution context (page 428) in which scripting is disabled (page 428), the script must do nothing and return nothing (a void return value).

Note: Thus, for instance, enabling designMode will disable any event handler attributes, event listeners, timeouts, etc, that were set by scripts in the document.

5.4.2 Event loops

To coordinate events, user interaction, scripts, rendering, networking, and so forth, user agents must use **event loops** as described in this section.

There must be at least one event loop (page 429) per user agent, and at most one event loop (page 429) per unit of related similar-origin browsing contexts (page 417).

An event loop (page 429) always has at least one browsing context (page 414). If an event loop (page 429)'s browsing contexts (page 414) all go away, then the event loop (page 429) goes away as well. A browsing context (page 414) always has an event loop (page 429) coordinating its activities.

An event loop (page 429) has one or more **task queues**. A task queue (page 429) is an ordered list of **tasks**, which can be:

Events

Asynchronously dispatching an Event object at a particular EventTarget object is a task.

Note: *Not all events are dispatched using the task queue (page 429), many are dispatched synchronously during other tasks.*

Parsing

The HTML parser (page 582) tokenising a single byte, and then processing any resulting tokens, is a task.

Callbacks

Calling a callback asynchronously is a task.

Using a resource

When an algorithm fetches (page 59) a resource, if the fetching occurs asynchronously then the processing of the resource once some or all of the resource is available is a task.

Reacting to DOM manipulation

Some elements have tasks that trigger in response to DOM manipulation, e.g. when that element is inserted into the document (page 24).

When a user agent is to **queue a task**, it must add the given task to one of the task queues (page 429) of the relevant event loop (page 429). All the tasks from one particular **task source** (e.g. the callbacks generated by timers, the events dispatched for mouse movements, the tasks queued for the parser) must always be added to the same task queue (page 429), but tasks from different task sources (page 429) may be placed in different task queues (page 429).

For example, a user agent could have one task queue (page 429) for mouse and key events (the user interaction task source (page 430)), and another for everything else. The user agent could then give keyboard and mouse events preference over other tasks three quarters of the time, keeping the interface responsive but not starving other task queues, and never processing events from any one task source (page 429) out of order.

An event loop (page 429) must continually run through the following steps for as long as it exists:

1. Run the oldest task on one of the event loop (page 429)'s task queues (page 429).
The user agent may pick any task queue (page 429).
2. Remove that task from its task queue (page 429).
3. If necessary, update the rendering or user interface of any Document or browsing context (page 414) to reflect the current state.
4. Return to the first step of the event loop (page 429).

5.4.2.1 Generic task sources

The following task sources (page 429) are used by a number of mostly unrelated features in this and other specifications.

The DOM manipulation task source

This task source (page 429) is used for features that react to DOM manipulations, such as things that happen asynchronously when an element is inserted into the document (page 24).

Asynchronous mutation events must be dispatched using tasks (page 429) queued (page 429) with the DOM manipulation task source (page 430). [DOMEVENTS]

The user interaction task source

This task source (page 429) is used for features that react to user interaction, for example keyboard or mouse input.

Asynchronous events sent in response to user input (e.g. click events) must be dispatched using tasks (page 429) queued (page 429) with the user interaction task source (page 430). [DOMEVENTS]

The networking task source

This task source (page 429) is used for features that trigger in response to network activity.

5.4.3 Security exceptions

** Define **security exception**.

5.4.4 The javascript: protocol

A URL using the javascript: protocol must, if and when **dereferenced**, be evaluated by executing the script obtained using the content retrieval operation defined for javascript: URLs. [JSURL]

When a browsing context (page 414) is navigated (page 473) to a javascript: URL, and the active document (page 414) of that browsing context has the same origin (page 426) as the script given by that URL, the script execution context (page 428) must be the Window object of the browsing context (page 414) being navigated, and the script document context (page 428) must be that active document (page 414).

When a browsing context (page 414) is navigated (page 473) to a javascript: URL, and the active document (page 414) of that browsing context has an origin (page 423) that is *not* the same (page 426) as that of the script given by the URL, the script execution context (page 428) must be an empty object, and the script browsing context (page 428) must be the browsing context (page 414) being navigated (page 473).

Otherwise, if the Document object of the element, attribute, or style sheet from which the javascript: URL was reached has an associated browsing context (page 414), the script execution context (page 428) must be an empty object, and the script execution context (page 428)'s associated browsing context (page 414) must be that browsing context (page 414).

Otherwise, the script is not executed and its return value is void.

If the result of executing the script is void (there is no return value), then the URL must be treated in a manner equivalent to an HTTP resource with an HTTP 204 No Content response.

Otherwise, the URL must be treated in a manner equivalent to an HTTP resource with a 200 OK response whose Content-Type metadata (page 60) is text/html and whose response body is the return value converted to a string value.

Note: Certain contexts, in particular img elements, ignore the Content-Type metadata (page 60).

So for example a javascript: URL for a src attribute of an img element would be evaluated in the context of an empty object as soon as the attribute is set; it would then be sniffed to determine the image type and decoded as an image.

A javascript: URL in an href attribute of an a element would only be evaluated when the link was followed (page 498).

The src attribute of an iframe element would be evaluated in the context of the iframe's own browsing context (page 414); once evaluated, its return value (if it was not void) would replace that browsing context (page 414)'s document, thus changing the variables visible in that browsing context (page 414).

Note: The rules for handling script execution in a script execution context (page 428) include making the script not execute (and just return void) in certain cases, e.g. in a sandbox or when the user has disabled scripting altogether.

5.4.5 Events

- ** We need to define how to handle events that are to be fired on a Document that is no longer the active document of its browsing context, and for Documents that have no browsing context. Do the events fire? Do the handlers in that document not fire? Do we just define scripting to be disabled when the document isn't active, with events still running as is? See also the script element section, which says scripts don't run when the document isn't active.

5.4.5.1 Event handler attributes

HTML elements (page 23) can have **event handler attributes** specified. These act as bubbling event listeners for the element on which they are specified.

Each event handler attribute has two parts, an event handler content attribute (page 432) and an event handler DOM attribute (page 432). Event handler attributes must initially be set to null. When their value changes (through the changing of their event handler content attribute or their event handler DOM attribute), they will either be null, or have an `EventListener` object assigned to them.

Objects other than `Element` objects, in particular `Window`, only have event handler DOM attribute (page 432) (since they have no content attributes).

Event handler content attributes, when specified, must contain valid ECMAScript code matching the ECMAScript FunctionBody production. [ECMA262]

When an event handler content attribute is set, if the element is owned by a Document that is in a browsing context (page 414), its new value must be interpreted as the body of an anonymous function with a single argument called `event`, with the new function's scope chain being linked from the activation object of the handler, to the element, to the element's form element if it is a form control, to the Document object, to the Window object of the browsing context (page 414) of that Document. The function's `this` parameter must be the `Element` object representing the element. The resulting function must then be set as the value of the corresponding event handler attribute, and the new value must be set as the value of the content attribute. If the given function body fails to compile, then the corresponding event handler attribute must be set to null instead (the content attribute must still be updated to the new value, though).

Note: See ECMA262 Edition 3, sections 10.1.6 and 10.2.3, for more details on activation objects. [ECMA262]

The script execution context (page 428) of the event handler must be the `Window` object at the end of the scope chain. The script document context (page 428) of the event handler must be the `Document` object that owns the event handler content attribute that was set.

When an event handler content attribute is set on an element owned by a Document that is not in a browsing context (page 414), the corresponding event handler attribute is not changed.

Note: Removing an event handler content attribute does not reset the corresponding event handler attribute either.

** How do we allow non-JS event handlers?

Event handler DOM attributes, on setting, must set the corresponding event handler attribute to their new value, and on getting, must return whatever the current value of the corresponding event handler attribute is (possibly null).

The following are the event handler attributes that must be supported by all HTML elements (page 23), as both content attributes and DOM attributes, and on `Window` objects, as DOM attributes:

onabort

Must be invoked whenever an abort event is targeted at or bubbles through the element.

onbeforeunload

Must be invoked whenever a beforeunload event is targeted at or bubbles through the element.

onblur

Must be invoked whenever a blur event is targeted at or bubbles through the element.

onchange

Must be invoked whenever a change event is targeted at or bubbles through the element.

onclick

Must be invoked whenever a click event is targeted at or bubbles through the element.

oncontextmenu

Must be invoked whenever a contextmenu event is targeted at or bubbles through the element.

ondblclick

Must be invoked whenever a dblclick event is targeted at or bubbles through the element.

ondrag

Must be invoked whenever a drag event is targeted at or bubbles through the element.

ondragend

Must be invoked whenever a dragend event is targeted at or bubbles through the element.

ondragenter

Must be invoked whenever a dragenter event is targeted at or bubbles through the element.

ondragleave

Must be invoked whenever a dragleave event is targeted at or bubbles through the element.

ondragover

Must be invoked whenever a dragover event is targeted at or bubbles through the element.

ondragstart

Must be invoked whenever a dragstart event is targeted at or bubbles through the element.

ondrop

Must be invoked whenever a drop event is targeted at or bubbles through the element.

onerror

Must be invoked whenever an error event is targeted at or bubbles through the element.

Note: *The onerror handler is also used for reporting script errors (page 437).*

onfocus

Must be invoked whenever a focus event is targeted at or bubbles through the element.

onhashchange

Must be invoked whenever a hashchange event is targeted at or bubbles through the element.

onkeydown

Must be invoked whenever a keydown event is targeted at or bubbles through the element.

onkeypress

Must be invoked whenever a keypress event is targeted at or bubbles through the element.

onkeyup

Must be invoked whenever a keyup event is targeted at or bubbles through the element.

onload

Must be invoked whenever a load event is targeted at or bubbles through the element.

onmessage

Must be invoked whenever a message event is targeted at or bubbles through the element.

onmousedown

Must be invoked whenever a mousedown event is targeted at or bubbles through the element.

onmousemove

Must be invoked whenever a mousemove event is targeted at or bubbles through the element.

onmouseout

Must be invoked whenever a mouseout event is targeted at or bubbles through the element.

onmouseover

Must be invoked whenever a mouseover event is targeted at or bubbles through the element.

onmouseup

Must be invoked whenever a mouseup event is targeted at or bubbles through the element.

onmousewheel

Must be invoked whenever a mousewheel event is targeted at or bubbles through the element.

onresize

Must be invoked whenever a resize event is targeted at or bubbles through the element.

onscroll

Must be invoked whenever a scroll event is targeted at or bubbles through the element.

onselect

Must be invoked whenever a select event is targeted at or bubbles through the element.

onstorage

Must be invoked whenever a storage event is targeted at or bubbles through the element.

onsubmit

Must be invoked whenever a submit event is targeted at or bubbles through the element.

onunload

Must be invoked whenever an unload event is targeted at or bubbles through the element.

When an event handler attribute is invoked, its argument must be set to the Event object of the event in question. If the function returns the exact boolean value false, the event's preventDefault() method must then be invoked. Exception: for historical reasons, for the HTML mouseover event, the preventDefault() method must be called when the function returns true instead.

All event handler attributes on an element, whether set to null or to a function, must be registered as event listeners on the element, as if the addEventListenerNS() method on the Element object's EventTarget interface had been invoked when the element was created, with the event type (*type* argument) equal to the type described for the event handler attribute in the list above, the namespace (*namespaceURI* argument) set to null, the listener set to be a target and bubbling phase listener (*useCapture* argument set to false), the event group set to the default group (*evtGroup* argument set to null), and the event listener itself (*listener* argument) set to do nothing while the event handler attribute is null, and set to invoke the function associated with the event handler attribute otherwise. (The *listener* argument is emphatically *not* the event handler attribute itself.)

5.4.5.2 Event firing

** maybe this should be moved higher up (terminology? conformance? DOM?) Also, the whole terminology thing should be changed so that we don't define any specific events here, we only define 'simple event', 'progress event', 'mouse event', 'key event', and the like, and have the actual dispatch use those generic terms when firing events.

Certain operations and methods are defined as firing events on elements. For example, the `click()` method on the `HTMLElement` interface is defined as firing a `click` event on the element. [DOM3EVENTS]

Firing a click event means that a `click` event with no namespace, which bubbles and is cancelable, and which uses the `MouseEvent` interface, must be dispatched at the given element. The event object must have its `screenX`, `screenY`, `clientX`, `clientY`, and `button` attributes set to 0, its `ctrlKey`, `shiftKey`, `altKey`, and `metaKey` attributes set according to the current state of the key input device, if any (false for any keys that are not available), its `detail` attribute set to 1, and its `relatedTarget` attribute set to null. The `getModifierState()` method on the object must return values appropriately describing the state of the key input device at the time the event is created.

Firing a change event means that a `change` event with no namespace, which bubbles but is not cancelable, and which uses the `Event` interface, must be dispatched at the given element.

Firing a contextmenu event means that a `contextmenu` event with no namespace, which bubbles and is cancelable, and which uses the `Event` interface, must be dispatched at the given element.

Firing a simple event called e means that an event with the name `e`, with no namespace, which does not bubble but is cancelable (unless otherwise stated), and which uses the `Event` interface, must be dispatched at the given element.

- ** **Firing a show event** means firing a simple event called `show` (page 436). Actually this
should fire an event that has modifier information (shift/ctrl etc), as well as having a pointer
to the node on which the menu was fired, and with which the menu was associated (which
could be an ancestor of the former).

Firing a load event means firing a simple event called `load` (page 436). **Firing an error event** means firing a simple event called `error` (page 436).

- ** **Firing a progress event called e** means something that hasn't yet been defined, in the [PROGRESS] spec.

The default action of these event is to do nothing unless otherwise stated.

- ** If you dispatch a custom "click" event at an element that would normally have default actions, should they get triggered? If so, we need to go through the entire spec and make sure that any default actions are defined in terms of *any* event of the right type on that element, not those that are dispatched in expected ways.

5.4.5.3 Events and the Window object

When an event is dispatched at a DOM node in a Document in a browsing context (page 414), if the event is not a load event, the user agent must also dispatch the event to the Window, as follows:

1. In the capture phase, the event must be dispatched to the Window object before being dispatched to any of the nodes.
2. In the bubble phase, the event must be dispatched to the Window object at the end of the phase, unless bubbling has been prevented.

5.4.5.4 Runtime script errors

This section only applies to user agents that support scripting in general and ECMAScript in particular.

Whenever a runtime script error occurs in one of the scripts associated with the document, the value of the onerror event handler DOM attribute of the Window object must be processed, as follows:

↪ If the value is a function

The function referenced by the onerror attribute must be invoked with three arguments, before notifying the user of the error.

The three arguments passed to the function are all DOMStrings; the first must give the message that the UA is considering reporting, the second must give the absolute URL (page 56) of the resource in which the error occurred, and the third must give the line number in that resource on which the error occurred.

If the function returns false, then the error should not be reported to the user. Otherwise, if the function returns another value (or does not return at all), the error should be reported to the user.

Any exceptions thrown or errors caused by this function must be reported to the user immediately after the error that the function was called for, without calling the function again.

↪ If the value is null

The error should not be reported to the user.

↪ If the value is anything else

The error should be reported to the user.

The initial value of onerror must be undefined.

5.5 User prompts

5.5.1 Simple dialogs

The `alert(message)` method, when invoked, must show the given *message* to the user. The user agent may make the method wait for the user to acknowledge the message before returning; if so, the user agent must pause (page 23) while the method is waiting.

The `confirm(message)` method, when invoked, must show the given *message* to the user, and ask the user to respond with a positive or negative response. The user agent must then pause (page 23) as the method waits for the user's response. If the user responds positively, the method must return true, and if the user responds negatively, the method must return false.

The `prompt(message, default)` method, when invoked, must show the given *message* to the user, and ask the user to either respond with a string value or abort. The user agent must then pause (page 23) as the method waits for the user's response. The second argument is optional. If the second argument (*default*) is present, then the response must be defaulted to the value given by *default*. If the user aborts, then the method must return null; otherwise, the method must return the string that the user responded with.

5.5.2 Printing

The `print()` method, when invoked, must run the printing steps (page 438).

User agents should also run the printing steps (page 438) whenever the user attempts to obtain a physical form (e.g. printed copy), or the representation of a physical form (e.g. PDF copy), of a document.

The **printing steps** are as follows:

1. The user agent may display a message to the user and/or may abort these steps.
 - For instance, a kiosk browser could silently ignore any invocations of the `print()` method.
 - For instance, a browser on a mobile device could detect that there are no printers in the vicinity and display a message saying so before continuing to offer a "save to PDF" option.
2. The user agent must fire a simple event (page 436) called `beforeprint` at the `Window` object of the browsing context of the `Document` that is being printed, as well as any nested browsing contexts (page 415) in it.
 - The `beforeprint` event can be used to annotate the printed copy, for instance adding the time at which the document was printed.
3. The user agent should offer the user the opportunity to obtain a physical form (page 672) (or the representation of a physical form) of the document. The user agent may wait for the user to either accept or decline before returning; if so, the user agent must pause (page 23) while the method is waiting. Even if the user agent doesn't wait at this point, the user agent must use the state of the relevant documents as

they are at this point in the algorithm if and when it eventually creates the alternate form.

4. The user agent must fire a simple event (page 436) called `afterprint` at the `Window` object of the browsing context of the `Document` that is being printed, as well as any nested browsing contexts (page 415) in it.

The `afterprint` event can be used to revert annotations added in the earlier event, as well as showing post-printing UI. For instance, if a page is walking the user through the steps of applying for a home loan, the script could automatically advance to the next step after having printed a form or other.

5.5.3 Dialogs implemented using separate documents

The `showModalDialog(url, arguments)` method, when invoked, must cause the user agent to run the following steps:

1. If the user agent is configured such that this invocation of `showModalDialog()` is somehow disabled, then the method returns the empty string; abort these steps.

Note: User agents are expected to disable this method in certain cases to avoid user annoyance. For instance, a user agent could require that a site be white-listed before enabling this method, or the user agent could be configured to only allow one modal dialog at a time.

2. Let *the list of background browsing contexts* be a list of all the browsing contexts that:
 - are part of the same unit of related browsing contexts (page 417) as the browsing context of the `Window` object on which the `showModalDialog()` method was called, and that
 - have an active document (page 414) whose origin (page 423) is the same (page 426) as the origin (page 423) of the script that called the `showModalDialog()` method at the time the method was called,

...as well as any browsing contexts that are nested inside any of the browsing contexts matching those conditions.
3. Disable the user interface for all the browsing contexts in *the list of background browsing contexts*. This should prevent the user from navigating those browsing contexts, causing events to be sent to those browsing context, or editing any content in those browsing contexts. However, it does not prevent those browsing contexts from receiving events from sources other than the user, from running scripts, from running animations, and so forth.
4. Create a new auxiliary browsing context (page 416), with the opener browsing context (page 416) being the browsing context of the `Window` object on which the `showModalDialog()` method was called. The new auxiliary browsing context has no name.

Note: This browsing context implements the `ModalWindow` interface.

5. Let the dialog arguments (page 440) of the new browsing context be set to the value of *arguments*.
6. Let the dialog arguments' origin (page 440) be the origin (page 423) of the script that called the `showModalDialog()` method.
7. Navigate (page 473) the new browsing context to *url*, with replacement enabled (page 476), and with the script browsing context (page 428) of the script that invoked the method as the source browsing context (page 473).
8. Wait for the browsing context to be closed. (The user agent must allow the user to indicate that the browsing context is to be closed.)
9. Reenable the user interface for all the browsing contexts in *the list of background browsing contexts*.
10. Return the auxiliary browsing context (page 416)'s return value (page 440).

Browsing contexts created by the above algorithm must implement the `ModalWindow` interface:

```
[XXX] interface ModalWindow {  
    readonly attribute any dialogArguments;  
    attribute DOMString returnValue;  
};
```

Such browsing contexts have associated **dialog arguments**, which are stored along with the **dialog arguments' origin**. These values are set by the `showModalDialog()` method in the algorithm above, when the browsing context is created, based on the arguments provided to the method.

The **dialogArguments** DOM attribute, on getting, must check whether its browsing context's active document (page 414)'s origin (page 423) is the same (page 426) as the dialog arguments' origin (page 440). If it is, then the browsing context's dialog arguments (page 440) must be returned unchanged. Otherwise, if the dialog arguments (page 440) are an object, then the empty string must be returned, and if the dialog arguments (page 440) are not an object, then the stringification of the dialog arguments (page 440) must be returned.

These browsing contexts also have an associated **return value**. The return value (page 440) of a browsing context must be initialized to the empty string when the browsing context is created.

The **returnValue** DOM attribute, on getting, must return the return value (page 440) of its browsing context, and on setting, must set the return value (page 440) to the given new value.

5.5.4 Notifications

Notifications are short, transient messages that bring the user's attention to new information, or remind the user of scheduled events.

Since notifications can be annoying if abused, this specification defines a mechanism that scopes notifications to a site's existing rendering area unless the user explicitly indicates that the site can be trusted.

To this end, each origin (page 423) can be flagged as being a **trusted notification source**. By default origins should not be flagged as such, but user agents may allow users to whitelist origins or groups of origins as being trusted notification sources (page 441). Only origins flagged as trusted in this way are allowed to show notification UI outside of their tab.

For example, a user agent could allow a user to mark all subdomains and ports of example.org as trusted notification sources. Then, mail.example.org and calendar.example.org would both be able to show notifications, without the user having to flag them individually.

The **showNotification(*title*, *subtitle*, *description*, *onclick*)** method, when invoked, must cause the user agent to show a notification.

If the method was invoked from a script whose script browsing context (page 428) has the sandboxed annoyances browsing context flag (page 218) set, then the notification must be shown within that browsing context (page 414). The notification is said to be a **sandboxed notification**.

Otherwise, if the origin (page 423) of the script browsing context (page 428) of the script that invoked the method is *not* flagged as being a trusted notification source (page 441), then the notification should be rendered within the top-level browsing context (page 415) of the script browsing context (page 428) of the script that invoked the method. The notification is said to be a **normal notification**. User agents should provide a way to set the origin's trusted notification source (page 441) flag from the notification, so that the user can benefit from notifications even when the user agent is not the active application.

Otherwise, the origin (page 423) is flagged as a trusted notification source (page 441), and the notification should be shown using the platform conventions for system-wide notifications. The notification is said to be a **trusted notification**. User agents may provide a way to unset the origin's trusted notification source (page 441) flag from within the notification, so as to allow users to easily disable notifications from sites that abuse the privilege.

For example, if a site contains a gadget of a mail application in a sandboxed iframe and that frame triggers a notification upon the receipt of a new e-mail message, that notification would be displayed on top of the gadget only.

However, if the user then goes to the main site of that mail application, the notification would be displayed over the entire rendering area of the tab for the site.

The notification, in this case, would have a button on it to let the user indicate that he trusts the site. If the user clicked this button, the next notification would use the system-wide notification system, appearing even if the tab for the mail application was buried deep inside a minimised window.

The style of notifications varies from platform to platform. On some, it is typically displayed as a "toast" window that slides in from the bottom right corner. In others, notifications are shown as semi-transparent white-on-grey overlays centered over the screen. Other schemes could include simulated ticker tapes, and speech-synthesis playback.

When a normal notification (page 441) (but not a sandboxed notification (page 441)) is shown, the user agent may bring the user's attention to the top-level browsing context (page 415) of the script browsing context (page 428) of the script that invoked the method, if that would be useful; but user agents should not use system-wide notification mechanisms to do so.

When a trusted notification (page 441) is shown, the user agent should bring the user's attention to the notification and the script browsing context (page 428) of the script that invoked the method, as per the platform conventions for attracting the user's attention to applications.

In the case of normal notifications (page 441), typically the only attention-grabbing device that would be employed would be something like flashing the tab's caption, or making it bold, or some such.

In addition, in the case of a trusted notification (page 441), the entire window could flash, or the browser's application icon could bounce or flash briefly, or a short sound effect could be played.

Notifications should include the following content:

- The *title*, *subtitle*, and *description* strings passed to the method. They may be truncated or abbreviated if necessary.
- The application name (page 115), if available, or else the document title (page 80), of the active document (page 414) of the script browsing context (page 428) of the script that invoked the method.
- An icon chosen from the external resource links (page 111) of type icon, if any are available.

If a new notification from one browsing context (page 414) has *title*, *subtitle*, and *description* strings that are identical to the *title*, *subtitle*, and *description* strings of an already-active notification from the same browsing context (page 414) or another browsing context (page 414) with the same origin (page 423), the user agent should not display the new notification, but should instead add an indicator to the already-active notification that another identical notification would otherwise have been shown.

For instance, if a user has his mail application open in three windows, and thus the same "New Mail" notification is fired three times each time a mail is received, instead of displaying three identical notifications each time, the user agent could just show one, with the title "New Mail x3".

Notifications should have a lifetime based on the platform conventions for notifications. However, the lifetime of a notification should not begin until the user has had the opportunity to see it, so if a notification is spawned for a browsing context (page 414) that is hidden, it should be shown for its complete lifetime once the user brings that browsing context (page 414) into view.

User agents should support multiple notifications at once.

User agents should support user interaction with notifications, if and as appropriate given the platform conventions. If a user activates a notification, and the *onclick* callback argument was present and is not null, then the script browsing context (page 428) of the function given

by *onclick* should be brought to the user's attention, and the *onclick* callback should then be invoked.

5.6 System state and capabilities

The **navigator** attribute of the Window interface must return an instance of the Navigator interface, which represents the identity and state of the user agent (the client), and allows Web pages to register themselves as potential protocol and content handlers:

```
interface Navigator {
    // client identification
    readonly attribute DOMString appName;
    readonly attribute DOMString appVersion;
    readonly attribute DOMString platform;
    readonly attribute DOMString userAgent;

    // system state
    readonly attribute boolean onLine;
    void registerProtocolHandler(in DOMString protocol, in DOMString url,
        in DOMString title);
    void registerContentHandler(in DOMString mimeType, in DOMString url,
        in DOMString title);

    // abilities
    short canPlayType(in DOMString type);
};
```

5.6.1 Client identification

In certain cases, despite the best efforts of the entire industry, Web browsers have bugs and limitations that Web authors are forced to work around.

This section defines a collection of attributes that can be used to determine, from script, the kind of user agent in use, in order to work around these issues.

Client detection should always be limited to detecting known current versions; future versions and unknown versions should always be assumed to be fully compliant.

appName

Must return either the string "Netscape" or the full name of the browser, e.g. "Mellblom Browernator".

appVersion

Must return either the string "4.0" or a string representing the version of the browser in detail, e.g. "1.0 (VMS; en-US) Mellblomenator/9000".

platform

Must return either the empty string or a string representing the platform on which the browser is executing, e.g. "MacIntel", "Win32", "FreeBSD i386", "WebTV OS".

userAgent

Must return the string used for the value of the "User-Agent" header in HTTP requests, or the empty string if no such header is ever sent.

5.6.2 Custom protocol and content handlers

The **registerProtocolHandler()** method allows Web sites to register themselves as possible handlers for particular protocols. For example, an online fax service could register itself as a handler of the fax: protocol ([RFC2806]), so that if the user clicks on such a link, he is given the opportunity to use that Web site. Analogously, the **registerContentHandler()** method allows Web sites to register themselves as possible handlers for content in a particular MIME type. For example, the same online fax service could register itself as a handler for image/g3fax files ([RFC1494]), so that if the user has no native application capable of handling G3 Facsimile byte streams, his Web browser can instead suggest he use that site to view the image.

User agents may, within the constraints described in this section, do whatever they like when the methods are called. A UA could, for instance, prompt the user and offer the user the opportunity to add the site to a shortlist of handlers, or make the handlers his default, or cancel the request. UAs could provide such a UI through modal UI or through a non-modal transient notification interface. UAs could also simply silently collect the information, providing it only when relevant to the user.

There is an example of how these methods could be presented to the user (page 447) below.

The arguments to the methods have the following meanings:

***protocol* (**registerProtocolHandler()** only)**

A scheme, such as ftp or fax. The scheme must be compared in an ASCII case-insensitive (page 31) manner by user agents for the purposes of comparing with the scheme part of URLs that they consider against the list of registered handlers.

The *protocol* value, if it contains a colon (as in "ftp:"), will never match anything, since schemes don't contain colons.

***mimeType* (**registerContentHandler()** only)**

A MIME type, such as model/vrml or text/richtext. The MIME type must be compared in an ASCII case-insensitive (page 31) manner by user agents for the purposes of comparing with MIME types of documents that they consider against the list of registered handlers.

User agents must compare the given values only to the MIME type/subtype parts of content types, not to the complete type including parameters. Thus, if *mimeType* values passed to this method include characters such as commas or whitespace, or include MIME parameters, then the handler being registered will never be used.

url

The URL (page 52) of the page that will handle the requests. When the user agent uses this URL, it must replace the first occurrence of the exact literal string "%s" with an escaped version of the URL of the content in question (as defined below), then resolve (page 55) the resulting URL (using the document base URL (page 54) of the script document context (page 428) of the script that originally invoked the

`registerContentHandler()` or `registerProtocolHandler()` method), and then fetch (page 59) the resulting URL using the GET method (or equivalent for non-HTTP URLs).

To get the escaped version of the URL of the content in question, the user agent must resolve (page 55) the URL, and then every character in the URL that doesn't match the `<query>` production defined in RFC 3986 must be replaced by the percent-encoded form of the character. [RFC3986]

If the user had visited a site at `http://example.com/` that made the following call:

```
navigator.registerContentHandler('application/x-soup',  
    'soup?url=%s', 'SoupWeb™')
```

...and then, much later, while visiting `http://www.example.net/`, clicked on a link such as:

```
<a href="chickenkiwi.soup">Download our Chicken Kiwi soup!</a>  
...then, assuming this chickenkiwi.soup file was served with the MIME type  
application/x-soup, the UA might navigate to the following URL:
```

```
http://example.com/soup?url=http://www.example.net/  
chickenk%C3%AFwi.soup
```

This site could then fetch the `chickenkiwi.soup` file and do whatever it is that it does with soup (synthesize it and ship it to the user, or whatever).

title

A descriptive title of the handler, which the UA might use to remind the user what the site in question is.

User agents should raise security exceptions (page 430) if the methods are called with *protocol* or *mimeType* values that the UA deems to be "privileged". For example, a site attempting to register a handler for http URLs or text/html content in a Web browser would likely cause an exception to be raised.

User agents must raise a `SYNTAX_ERR` exception if the *url* argument passed to one of these methods does not contain the exact literal string "%s".

User agents must not raise any other exceptions (other than binding-specific exceptions, such as for an incorrect number of arguments in an ECMAScript implementation).

This section does not define how the pages registered by these methods are used, beyond the requirements on how to process the *url* value (see above). To some extent, the processing model for navigating across documents (page 473) defines some cases where these methods are relevant, but in general UAs may use this information wherever they would otherwise consider handing content to native plugins or helper applications.

UAs must not use registered content handlers to handle content that was returned as part of a non-GET transaction (or rather, as part of any non-idempotent transaction), as the remote site would not be able to fetch the same data.

5.6.2.1 Security and privacy

These mechanisms can introduce a number of concerns, in particular privacy concerns.

Hijacking all Web usage. User agents should not allow protocols that are key to its normal operation, such as http or https, to be rerouted through third-party sites. This would allow a user's activities to be trivially tracked, and would allow user information, even in secure connections, to be collected.

Hijacking defaults. It is strongly recommended that user agents do not automatically change any defaults, as this could lead the user to send data to remote hosts that the user is not expecting. New handlers registering themselves should never automatically cause those sites to be used.

Registration spamming. User agents should consider the possibility that a site will attempt to register a large number of handlers, possibly from multiple domains (e.g. by redirecting through a series of pages each on a different domain, and each registering a handler for video/mpeg — analogous practices abusing other Web browser features have been used by pornography Web sites for many years). User agents should gracefully handle such hostile attempts, protecting the user.

Misleading titles. User agents should not rely wholly on the *title* argument to the methods when presenting the registered handlers to the user, since sites could easily lie. For example, a site hostile.example.net could claim that it was registering the "Cuddly Bear Happy Content Handler". User agents should therefore use the handler's domain in any UI along with any title.

Hostile handler metadata. User agents should protect against typical attacks against strings embedded in their interface, for example ensuring that markup or escape characters in such strings are not executed, that null bytes are properly handled, that over-long strings do not cause crashes or buffer overruns, and so forth.

Leaking Intranet URLs. The mechanism described in this section can result in secret Intranet URLs being leaked, in the following manner:

1. The user registers a third-party content handler as the default handler for a content type.
2. The user then browses his corporate Intranet site and accesses a document that uses that content type.
3. The user agent contacts the third party and hands the third party the URL to the Intranet content.

No actual confidential file data is leaked in this manner, but the URLs themselves could contain confidential information. For example, the URL could be <http://www.corp.example.com/upcoming-aquisitions/the-sample-company.egf>, which might tell the third party that Example Corporation is intending to merge with The Sample Company. Implementors might wish to consider allowing administrators to disable this feature for certain subdomains, content types, or protocols.

Leaking secure URLs. User agents should not send HTTPS URLs to third-party sites registered as content handlers, in the same way that user agents do not send Referer headers from secure sites to third-party sites.

Leaking credentials. User agents must never send username or password information in the URLs that are escaped and included sent to the handler sites. User agents may even avoid attempting to pass to Web-based handlers the URLs of resources that are known to

require authentication to access, as such sites would be unable to access the resources in question without prompting the user for credentials themselves (a practice that would require the user to know whether to trust the third-party handler, a decision many users are unable to make or even understand).

5.6.2.2 Sample user interface

This section is non-normative.

A simple implementation of this feature for a desktop Web browser might work as follows.

The `registerProtocolHandler()` method could display a modal dialog box:

```
||[ Protocol Handler Registration ]|||||||||||||||||||||||||||  
|  
| This Web page:  
|  
|   Kittens at work  
|   http://kittens.example.org/  
|  
| ...would like permission to handle the protocol "x-meow:"  
| using the following Web-based application:  
|  
|   Kittens-at-work displayer  
|   http://kittens.example.org/?show=%s  
|  
| Do you trust the administrators of the "kittens.example.  
| org" domain?  
|  
|   ( Trust kittens.example.org ) (( Cancel ))
```

...where "Kittens at work" is the title of the page that invoked the method, "http://kittens.example.org/" is the URL of that page, "x-meow" is the string that was passed to the `registerProtocolHandler()` method as its first argument (*protocol*), "http://kittens.example.org/?show=%s" was the second argument (*url*), and "Kittens-at-work displayer" was the third argument (*title*).

If the user clicks the Cancel button, then nothing further happens. If the user clicks the "Trust" button, then the handler is remembered.

When the user then attempts to fetch a URL that uses the "x-meow:" scheme, then it might display a dialog as follows:

```
||[ Unknown Protocol ]|||||||||||||||||||||||||||  
|  
| You have attempted to access:  
|  
|   x-meow:S2l0dGVucyBhcmUgdGhlIGN1dGVzdCE%3D  
|  
| How would you like FerretBrowser to handle this resource?  
|  
|   (o) Contact the FerretBrowser plugin registry to see if
```

<p>there is an official way to handle this resource.</p> <p>() Pass this URL to a local application: [/no application selected/] (Choose)</p> <p>() Pass this URL to the "Kittens-at-work displayer" application at "kittens.example.org".</p> <p>[] Always do this for resources using the "x-meow" protocol in future.</p>
<p>(Ok) ((Cancel))</p>

...where the third option is the one that was primed by the site registering itself earlier.

If the user does select that option, then the browser, in accordance with the requirements described in the previous two sections, will redirect the user to "http://kittens.example.org/?show=x-meow%3AS2I0dGVucyBhcmUgdGhlIGN1dGVzdCE%253D".

The `registerContentHandler()` method would work equivalently, but for unknown MIME types instead of unknown protocols.

5.6.3 Client abilities

The `canPlayType(type)` method must return 1 if `type` is a MIME type that the user agent is confident represents a media resource (page 238) that it can render if used in a audio or video element, 0 if it cannot determine whether it could do so, and -1 if it is confident that it would not be able to render resources of that type.

5.7 Offline Web applications

5.7.1 Introduction

This section is non-normative.

**

...

5.7.2 Application caches

An **application cache** is a collection of resources. An application cache is identified by the absolute URL (page 56) of a resource manifest which is used to populate the cache.

Application caches are versioned, and there can be different instances of caches for the same manifest URL, each having a different version. A cache is newer than another if it was created after the other (in other words, caches in a group have a chronological order).

Each group of application caches for the same manifest URL has a common **update status**, which is one of the following: *idle, checking, downloading*.

Each group of application caches for the same manifest URL also has a common **lifecycle status**, which is one of the following: *new*, *mature*, *obsolete*. A **relevant application cache** is an application cache (page 448) whose lifecycle status (page 449) is *mature*.

A browsing context (page 414) is associated with the application cache appropriate for its active document (page 414), if any. A Document initially has no appropriate cache, but steps in the parser (page 623) and in the navigation (page 473) sections cause cache selection (page 460) to occur early in the page load process. A browsing context's associated cache can also change (page 481) during session history traversal (page 480).

An application cache consists of:

- One or more resources (including their out-of-band metadata, such as HTTP headers, if any), identified by URLs, each falling into one (or more) of the following categories:

Master entries

Documents that were added to the cache because a browsing context (page 414) was navigated (page 473) to that document and the document indicated that this was its cache, using the `manifest` attribute.

The manifest

The resource corresponding to the URL that was given in an master entry's `html` element's `manifest` attribute. The manifest is fetched and processed during the application cache update process (page 455). All the master entries (page 449) have the same origin (page 426) as the manifest.

Explicit entries

Resources that were listed in the cache's manifest (page 449). Explicit entries can also be marked as **foreign**, which means that they have a `manifest` attribute but that it doesn't point at this cache's manifest (page 449).

Fallback entries

Resources that were listed in the cache's manifest (page 449) as fallback entries.

Dynamic entries

Resources that were added to the cache by the `add()` method.

Note: A URL in the list can be flagged with multiple different types, and thus an entry can end up being categorized as multiple entries. For example, an entry can be an explicit entry and a dynamic entry at the same time.

- Zero or more **fallback namespaces**: URLs, used as prefix match patterns (page 460), each of which is mapped to a fallback entry (page 449). Each namespace URL has the same origin (page 426) as the manifest (page 449).
- Zero or more URLs that form the **online whitelist namespaces**.

Multiple application caches can contain the same resource, e.g. if their manifests all reference that resource. If the user agent is to **select an application cache** from a list of relevant application caches (page 449) that contain a resource, that the user agent must use the application cache that the user most likely wants to see the resource from, taking into account the following:

- which application cache was most recently updated,
- which application cache was being used to display the resource from which the user decided to look at the new resource, and
- which application cache the user prefers.

5.7.3 The cache manifest syntax

5.7.3.1 A sample manifest

This section is non-normative.

This example manifest requires two images and a style sheet to be cached and whitelists a CGI script.

```
CACHE MANIFEST
# the above line is required

# this is a comment
# there can be as many of these anywhere in the file
# they are all ignored
# comments can have spaces before them
# but must be alone on the line

# blank lines are ignored too

# these are files that need to be cached they can either be listed
# first, or a "CACHE:" header could be put before them, as is done
# lower down.
images/sound-icon.png
images/background.png
# note that each file has to be put on its own line

# here is a file for the online whitelist -- it isn't cached, and
# references to this file will bypass the cache, always hitting the
# network (or trying to, if the user is offline).
NETWORK:
comm.cgi

# here is another set of files to cache, this time just the CSS file.
CACHE:
style/default.css
```

5.7.3.2 Writing cache manifests

Manifests must be served using the text/cache-manifest MIME type. All resources served using the text/cache-manifest MIME type must follow the syntax of application cache manifests, as described in this section.

An application cache manifest is a text file, whose text is encoded using UTF-8. Data in application cache manifests is line-based. Newlines must be represented by U+000A LINE

FEED (LF) characters, U+000D CARRIAGE RETURN (CR) characters, or U+000D CARRIAGE RETURN (CR) U+000A LINE FEED (LF) pairs.

Note: This is a willful double violation of RFC2046. [RFC2046]

The first line of an application cache manifest must consist of the string "CACHE", a single U+0020 SPACE character, the string "MANIFEST", and zero or more U+0020 SPACE and U+0009 CHARACTER TABULATION (tab) characters. The first line may optionally be preceded by a U+FEFF BYTE ORDER MARK (BOM) character. If any other text is found on the first line, the user agent will ignore the entire file.

Subsequent lines, if any, must all be one of the following:

A blank line

Blank lines must consist of zero or more U+0020 SPACE and U+0009 CHARACTER TABULATION (tab) characters only.

A comment

Comment lines must consist of zero or more U+0020 SPACE and U+0009 CHARACTER TABULATION (tab) characters, followed by a single U+0023 NUMBER SIGN (#) character, followed by zero or more characters other than U+000A LINE FEED (LF) and U+000D CARRIAGE RETURN (CR) characters.

Note: Comments must be on a line on their own. If they were to be included on a line with a URL, the "#" would be mistaken for part of a fragment identifier.

A section header

Section headers change the current section. There are three possible section headers:

CACHE:

Switches to the explicit section.

FALLBACK:

Switches to the fallback section.

NETWORK:

Switches to the online whitelist section.

Section header lines must consist of zero or more U+0020 SPACE and U+0009 CHARACTER TABULATION (tab) characters, followed by one of the names above (including the U+003A COLON (:) character) followed by zero or more U+0020 SPACE and U+0009 CHARACTER TABULATION (tab) characters.

Ironically, by default, the current section is the explicit section.

Data for the current section

The format that data lines must take depends on the current section.

When the current section is the explicit section or the online whitelist section, data lines must consist of zero or more U+0020 SPACE and U+0009 CHARACTER TABULATION (tab) characters, a valid URL (page 52) identifying a resource other than the manifest itself, and then zero or more U+0020 SPACE and U+0009 CHARACTER TABULATION (tab) characters.

When the current section is the fallback section, data lines must consist of zero or more U+0020 SPACE and U+0009 CHARACTER TABULATION (tab) characters, a valid URL (page 52) identifying a resource other than the manifest itself, one or more U+0020 SPACE and U+0009 CHARACTER TABULATION (tab) characters, another valid URL (page 52) identifying a resource other than the manifest itself, and then zero or more U+0020 SPACE and U+0009 CHARACTER TABULATION (tab) characters.

Note: The URLs in data lines can't be empty strings, since those would be relative URLs to the manifest itself. Such lines would be confused with blank or invalid lines, anyway.

Manifests may contain sections more than once. Sections may be empty.

URLs that are to be fallback pages associated with fallback namespaces (page 449), and those namespaces themselves, must be given in fallback sections, with the namespace being the first URL of the data line, and the corresponding fallback page being the second URL. All the other pages to be cached must be listed in explicit sections.

Fallback namespaces (page 449) and fallback entries (page 449) must have the same origin (page 426) as the manifest itself.

A fallback namespace (page 449) must not be listed more than once.

URLs that the user agent is to put into the online whitelist (page 449) must all be specified in online whitelist sections. (This is needed for any URL that the page is intending to use to communicate back to the server.)

Relative URLs must be given relative to the manifest's own URL.

URLs in manifests must not have fragment identifiers (i.e. the U+0023 NUMBER SIGN character isn't allowed in URLs in manifests).

5.7.3.3 Parsing cache manifests

When a user agent is to **parse a manifest**, it means that the user agent must run the following steps:

1. The user agent must decode the byte stream corresponding with the manifest to be parsed, treating it as UTF-8. Bytes or sequences of bytes that are not valid UTF-8 sequences must be interpreted as a U+FFFD REPLACEMENT CHARACTER.
2. Let *explicit URLs* be an initially empty list of explicit entries (page 449).
3. Let *fallback URLs* be an initially empty mapping of fallback namespaces (page 449) to fallback entries (page 449).
4. Let *online whitelist URLs* be an initially empty list of URLs for a online whitelist (page 449).
5. Let *input* be the decoded text of the manifest's byte stream.
6. Let *position* be a pointer into *input*, initially pointing at the first character.

7. If *position* is pointing at a U+FEFF BYTE ORDER MARK (BOM) character, then advance *position* to the next character.
8. If the characters starting from *position* are "CACHE", followed by a U+0020 SPACE character, followed by "MANIFEST", then advance *position* to the next character after those. Otherwise, this isn't a cache manifest; abort this algorithm with a failure while checking for the magic signature.
9. Collect a sequence of characters (page 32) that are U+0020 SPACE or U+0009 CHARACTER TABULATION (tab) characters.
10. If *position* is not past the end of *input* and the character at *position* is neither a U+000A LINE FEED (LF) characters nor a U+000D CARRIAGE RETURN (CR) character, then this isn't a cache manifest; abort this algorithm with a failure while checking for the magic signature.
11. This is a cache manifest. The algorithm cannot fail beyond this point (though bogus lines can get ignored).
12. Let *mode* be "explicit".
13. *Start of line*: If *position* is past the end of *input*, then jump to the last step. Otherwise, collect a sequence of characters (page 32) that are U+000A LINE FEED (LF), U+000D CARRIAGE RETURN (CR), U+0020 SPACE, or U+0009 CHARACTER TABULATION (tab) characters.
14. Now, collect a sequence of characters (page 32) that are *not* U+000A LINE FEED (LF) or U+000D CARRIAGE RETURN (CR) characters, and let the result be *line*.
15. Drop any trailing U+0020 SPACE and U+0009 CHARACTER TABULATION (tab) characters at the end of *line*.
16. If *line* is the empty string, then jump back to the step labeled "start of line".
17. If the first character in *line* is a U+0023 NUMBER SIGN (#) character, then jump back to the step labeled "start of line".
18. If *line* equals "CACHE:" (the word "CACHE" followed by a U+003A COLON (:)) character), then set *mode* to "explicit" and jump back to the step labeled "start of line".
19. If *line* equals "FALLBACK:" (the word "FALLBACK" followed by a U+003A COLON (:)) character), then set *mode* to "fallback" and jump back to the step labeled "start of line".
20. If *line* equals "NETWORK:" (the word "NETWORK" followed by a U+003A COLON (:)) character), then set *mode* to "online whitelist" and jump back to the step labeled "start of line".
21. If *line* ends with a U+003A COLON (:) character, then set *mode* to "unknown" and jump back to the step labeled "start of line".
22. This is either a data line or it is syntactically incorrect.
23. Let *position* be a pointer into *line*, initially pointing at the start of the string.

24. Let *tokens* be a list of strings, initially empty.
25. While *position* doesn't point past the end of *line*:
 1. Let *current token* be an empty string.
 2. While *position* doesn't point past the end of *line* and the character at *position* is neither a U+0020 SPACE nor a U+0009 CHARACTER TABULATION (tab) character, add the character at *position* to *current token* and advance *position* to the next character in *input*.
 3. Add *current token* to the *tokens* list.
 4. While *position* doesn't point past the end of *line* and the character at *position* is either a U+0020 SPACE or a U+0009 CHARACTER TABULATION (tab) character, advance *position* to the next character in *input*.
26. Process *tokens* as follows:
 - ↪ **If mode is "explicit"**

Resolve (page 55) the first item in *tokens*; ignore the rest.

If this fails, then jump back to the step labeled "start of line".

If the resulting absolute URL (page 56) has a different <scheme> (page 53) component than the manifest's URL (compared in an ASCII case-insensitive (page 31) manner), then jump back to the step labeled "start of line".

Drop the <fragment> (page 54) component of the resulting absolute URL (page 56), if it has one.

Add the resulting absolute URL (page 56) to the *explicit URLs*.
 - ↪ **If mode is "fallback"**

Let *part one* be the first token in *tokens*, and let *part two* be the second token in *tokens*.

Resolve (page 55) *part one* and *part two*.

If either fails, then jump back to the step labeled "start of line".

If the absolute URL (page 56) corresponding to either *part one* or *part two* does not have the same origin (page 426) as the manifest's URL, then jump back to the step labeled "start of line".

Drop any the <fragment> (page 54) components of the resulting absolute URLs (page 56).

If the absolute URL (page 56) corresponding to *part one* is already in the *fallback URLs* mapping as a fallback namespace (page 449), then jump back to the step labeled "start of line".

Otherwise, add the absolute URL (page 56) corresponding to *part one* to the *fallback URLs* mapping as a fallback namespace (page 449), mapped to the absolute URL (page 56) corresponding to *part two* as the fallback entry (page 449).

↪ **If mode is "online whitelist"**

Resolve (page 55) the first item in *tokens*; ignore the rest.

If this fails, then jump back to the step labeled "start of line".

If the resulting absolute URL (page 56) has a different <scheme> (page 53) component than the manifest's URL (compared in an ASCII case-insensitive (page 31) manner), then jump back to the step labeled "start of line".

Drop the <fragment> (page 54) component of the resulting absolute URL (page 56), if it has one.

Add the resulting absolute URL (page 56) to the *online whitelist URLs*.

↪ **If mode is "unknown"**

Do nothing. The line is ignored.

27. Jump back to the step labeled "start of line". (That step jumps to the next, and last, step when the end of the file is reached.)
28. Return the *explicit URLs* list, the *fallback URLs* mapping, and the *online whitelist URLs*.

Note: If a resource is listed in the explicit section and matches an entry in the online whitelist, or if a resource matches both an entry in the fallback section and the online whitelist, the resource will taken from the cache, and the online whitelist entry will be ignored.

5.7.4 Updating an application cache

When the user agent is required (by other parts of this specification) to start the **application cache update process** for a manifest (page 449) URL or for an application cache (page 448), potentially given a particular browsing context (page 414), and potentially given a new master (page 449) resource, the user agent must run the following steps:

- ** the event stuff needs to be more consistent -- something about showing every step of the ui or no steps or something; and we need to deal with showing ui for browsing contexts that open when an update is already in progress, and we may need to give applications control over the ui the first time they cache themselves (right now the original cache is done without notifications to the browsing contexts); also, we need to update this so all event firing uses queues

1. Atomically, so as to avoid race conditions, perform the following substeps:
 1. Let *manifest URL* be the URL of the manifest (page 449) to be updated, or of the manifest (page 449) of the application cache (page 448) to be updated, as appropriate.
 2. If these steps were invoked with a URL (as opposed to a specific cache), and there is no application cache (page 448) identified by *manifest URL* whose lifecycle status (page 449) is not *obsolete*, then create a new application

cache (page 448) identified with that URL and set the group's lifecycle status (page 449) to *new*.

3. If these steps were invoked with a new master (page 449) resource, then flag the resource's Document as a candidate for this manifest URL's caches, so that it will be associated with an application cache identified by this manifest URL (page 459) later, when such an application cache (page 448) is ready.
4. Let *cache group* be the group of application caches (page 448) identified by *manifest URL*.
5. Let *cache* be the most recently updated application cache (page 448) identified by *manifest URL* (that is, the newest version found in *cache group*).
6. If these steps were invoked with a browsing context (page 414), and the status (page 448) of the *cache group* is *checking* or *downloading*, then fire a simple event (page 436) called *checking* at the ApplicationCache singleton of that browsing context (page 414).
7. If these steps were invoked with a browsing context (page 414), and the status (page 448) of the *cache group* is *downloading*, then also fire a simple event (page 436) called *downloading* at the ApplicationCache singleton of that browsing context (page 414).
8. If the status (page 448) of the *cache group* is either *checking* or *downloading*, then abort this instance of the update process, as an update is already in progress for them.
9. Set the status (page 448) of this group of caches to *checking*.

The remainder of the steps run asynchronously.

2. If there is already a resource with the URL of *manifest URL* in *cache*, and that resource is categorized as a manifest (page 449), then this is an **upgrade attempt**. Otherwise, this is a **cache attempt**.

Note: *If this is a cache attempt (page 456), then cache is forcibly the only application cache in cache group, and it hasn't ever been populated from its manifest (i.e. this update is an attempt to download the application for the first time). It also can't have any browsing contexts associated with it.*

3. Fire a simple event (page 436) called *checking* at the ApplicationCache singleton of each browsing context (page 414) whose active document (page 414) is associated with a cache in *cache group*. The default action of this event should be the display of some sort of user interface indicating to the user that the user agent is checking for the availability of updates.

Note: *Again, if this is a cache attempt (page 456), then cache group has only one cache and it has no browsing contexts associated with it, so no events are dispatched due to this step or any of the other steps that fire events other than the final cached event.*

4. Fetch (page 59) the resource from *manifest URL*, and let *manifest* be that resource.

If the resource is labeled with the MIME type `text/cache-manifest`, parse *manifest* according to the rules for parsing manifests (page 452), obtaining a list of explicit entries (page 449), fallback entries (page 449) and the fallback namespaces (page 449) that map to them, and entries for the online whitelist (page 449).

5. If the previous step fails due to a 404 or 410 response or equivalent, then run the cache removal steps (page 460)

If the previous step fails in some other way (e.g. the server returns another 4xx or 5xx response or equivalent, or there is a DNS error, or the connection times out, or the user cancels the download, or the parser for manifests fails when checking the magic signature), or if the resource is labeled with a MIME type other than `text/cache-manifest`, then run the cache failure steps (page 460).

6. If this is an upgrade attempt (page 456) and the newly downloaded *manifest* is byte-for-byte identical to the manifest found in *cache*, or if the server reported it as "304 Not Modified" or equivalent, then run these substeps:

1. Fire a simple event (page 436) called `noupdate` at the `ApplicationCache` singleton of each browsing context (page 414) whose active document (page 414) is associated with a cache in *cache group*. The default action of this event should be the display of some sort of user interface indicating to the user that the application is up to date.
2. If there are any pending downloads of master entries (page 449) that are being stored in the cache, then wait for all of them to have completed. If any of these downloads fail (e.g. the server returns a 4xx or 5xx response or equivalent, or there is a DNS error, or the connection times out, or the user cancels the download), then run the cache failure steps (page 460).
3. Let the status (page 448) of the group of caches to which *cache* belongs be *idle*. If appropriate, remove any user interface indicating that an update for this cache is in progress.
4. Abort the update process.

7. Set the status (page 448) of *cache group* to *downloading*.

8. Fire a simple event (page 436) called `downloading` at the `ApplicationCache` singleton of each browsing context (page 414) whose active document (page 414) is associated with a cache in *cache group*. The default action of this event should be the display of some sort of user interface indicating to the user that a new version is being downloaded.

9. If this is an upgrade attempt (page 456), then let *new cache* be a newly created application cache (page 448) identified by manifest URL, being a new version in *cache group*. Otherwise, let *new cache* and *cache* be the same version of the application cache.

10. Let *file list* be an empty list of URLs with flags.

11. Add all the URLs in the list of explicit entries (page 449) obtained by parsing *manifest* to *file list*, each flagged with "explicit entry".

12. Add all the URLs in the list of fallback entries (page 449) obtained by parsing *manifest* to *file list*, each flagged with "fallback entry".
13. If this is an upgrade attempt (page 456), then add all the URLs of master entries (page 449) in *cache* to *file list*, each flagged with "master entry".
14. If this is an upgrade attempt (page 456), then add all the URLs of dynamic entries (page 449) in *cache* to *file list*, each flagged with "dynamic entry".
15. If any URL is in *file list* more than once, then merge the entries into one entry for that URL, that entry having all the flags that the original entries had.
16. For each URL in *file list*, run the following steps. These steps may be run in parallel for two or more of the URLs at a time.
 1. Fire a simple event (page 436) called `progress` at the `ApplicationCache` singleton of each browsing context (page 414) whose active document (page 414) is associated with a cache in *cache group*. The default action of this event should be the display of some sort of user interface indicating to the user that a file is being downloaded in preparation for updating the application.
 2. Fetch (page 59) the resource. If this is an upgrade attempt (page 456), then use *cache* as an HTTP cache, and honor HTTP caching semantics (such as expiration, ETags, and so forth) with respect to that cache. User agents may also have other caches in place that are also honored.

Note: If the resource in question is already being downloaded for other reasons then the existing download process can be used for the purposes of this step, as defined by the fetching (page 59) algorithm.

An example of a resource that might already be being downloaded is a large image on a Web page that is being seen for the first time. The image would get downloaded to satisfy the `img` element on the page, as well as being listed in the cache manifest. According to the rules for fetching (page 59) that image only need be downloaded once, and it can be used both for the cache and for the rendered Web page.

3. If the previous steps fails (e.g. the server returns a 4xx or 5xx response or equivalent, or there is a DNS error, or the connection times out, or the user cancels the download), or if the server returned a redirect, then run the cache failure steps (page 460).

Note: Redirects are fatal because they are either indicative of a network problem (e.g. a captive portal); or would allow resources to be added to the cache under URLs that differ from any URL that the networking model will allow access to, leaving orphan entries; or would allow resources to be stored under URLs different than their true URLs. All of these situations are bad.

4. Otherwise, the fetching succeeded. Store the resource in the *new cache*.

5. If the URL being processed was flagged as an "explicit entry" in *file list*, then categorize the entry as an explicit entry (page 449).
6. If the URL being processed was flagged as a "fallback entry" in *file list*, then categorize the entry as a fallback entry (page 449).
7. If the URL being processed was flagged as an "master entry" in *file list*, then categorize the entry as a master entry (page 449).
8. If the URL being processed was flagged as an "dynamic entry" in *file list*, then categorize the entry as a dynamic entry (page 449).
9. As an optimization, if the resource is an HTML or XML file whose root element is an `html` element with a `manifest` attribute whose value doesn't match the manifest URL of the application cache being processed, then the user agent should mark the entry as being foreign (page 449).
17. Store the list of fallback namespaces (page 449), and the URLs of the fallback entries (page 449) that they map to, in *new cache*.
18. Store the URLs that form the new online whitelist (page 449) in *new cache*.
19. Wait for all pending downloads of master entries (page 449) that are being stored in the cache to have completed.

For example, if the browsing context (page 414)'s active document (page 414) isn't itself listed in the cache manifest, then it might still be being downloaded.

If any of these downloads fail (e.g. the connection times out, or the user cancels the download), then run the cache failure steps (page 460).
20. Store *manifest* in *new cache*, if it's not there already, and categorize this entry (whether newly added or not) as the manifest (page 449).
21. If this is a cache attempt (page 456), then:

Associate any `Document` objects that were flagged as candidates (page 456) for this manifest URL's caches with *cache*.

Fire a simple event (page 436) called `cached` at the `ApplicationCache` singleton of each browsing context (page 414) whose active document (page 414) is associated with a cache in *cache group*. The default action of this event should be the display of some sort of user interface indicating to the user that the application has been cached and that they can now use it offline.

Set the lifecycle status (page 449) of *cache group* to *mature*.

Set the update status (page 448) of *cache group* to *idle*.
22. Otherwise, this is an upgrade attempt (page 456):

Fire a simple event (page 436) called `updateready` at the `ApplicationCache` singleton of each browsing context (page 414) whose active document (page 414) is associated with a cache in *cache group*. The default action of this event should be the display of some sort of user interface indicating to the user that a new version is available and that they can activate it by reloading the page.

Set the status (page 448) of *cache group* to *idle*.

The **cache removal steps** are as follows:

1. If this is a cache attempt (page 456), then discard *cache* and abort the update process.
2. Fire a simple event (page 436) called *obsolete* at the ApplicationCache singleton of each browsing context (page 414) whose active document (page 414) is associated with a cache in *cache group*. The default action of this event should be the display of some sort of user interface indicating to the user that the application is no longer available for offline use.
3. Set the lifecycle status (page 449) of *cache group* to *obsolete*.
4. Let the update status (page 448) of the group of caches to which *cache* belongs be *idle*. If appropriate, remove any user interface indicating that an update for this cache is in progress. Abort the update process.

The **cache failure steps** are as follows:

1. Fire a simple event (page 436) called *error* at the ApplicationCache singleton of each browsing context (page 414) whose active document (page 414) is associated with a cache in *cache group*. The default action of this event should be the display of some sort of user interface indicating to the user that the user agent failed to save the application for offline use.
2. If this is a cache attempt (page 456), then discard *cache* and abort the update process.
3. Otherwise, let the status (page 448) of the group of caches to which *cache* belongs be *idle*. If appropriate, remove any user interface indicating that an update for this cache is in progress. Abort the update process.

5.7.5 Processing model

The processing model of application caches for offline support in Web applications is part of the navigation (page 473) model, but references the algorithms defined in this section.

A URL **matches a fallback namespace** if there exists a relevant application cache (page 449) whose manifest (page 449)'s URL has the same origin (page 426) as the URL in question, and that has a fallback namespace (page 449) that is a prefix match (page 31) for the URL being examined. If multiple fallback namespaces match the same URL, the longest one is the one that matches. A URL looking for a fallback namespace can match more than one application cache at a time, but only matches one namespace in each cache.

If a manifest `http://example.com/app1/manifest` declares that `http://example.com/resources/images` is a fallback namespace, and the user navigates to `HTTP://EXAMPLE.COM:80/resources/images/cat.png`, then the user agent will decide that the application cache identified by `http://example.com/app1/manifest` contains a namespace with a match for that URL.

When the **application cache selection algorithm** algorithm is invoked with a manifest URL, the user agent must run the first applicable set of steps from the following list:

↪ **If the resource is not being loaded as part of navigation of a browsing context (page 414)**

Do nothing.

↪ **If the resource being loaded was loaded from an application cache and the URL of that application cache's manifest is the same as the manifest URL with which the algorithm was invoked**

Associate the Document with the cache from which it was loaded. Invoke the application cache update process (page 455) for that cache and with the browsing context (page 414) being navigated.

↪ **If the resource being loaded was loaded from an application cache and the URL of that application cache's manifest is *not* the same as the manifest URL with which the algorithm was invoked**

Mark the entry for this resource in the application cache from which it was loaded as foreign (page 449).

Restart the current navigation from the top of the navigation algorithm (page 473), undoing any changes that were made as part of the initial load (changes can be avoided by ensuring that the step to update the session history with the new page (page 476) is only ever completed *after* the application cache selection algorithm is run, though this is not required).

Note: *The navigation will not result in the same resource being loaded, because "foreign" entries are never picked during navigation.*

User agents may notify the user of the inconsistency between the cache manifest and the resource's own metadata, to aid in application development.

↪ **If the resource being loaded was not loaded from an application cache, but it was loaded using HTTP GET or equivalent**

1. If the manifest URL does not have the same origin (page 426) as the resource's own URL, then invoke the application cache selection algorithm (page 461) again, but without a manifest, and abort these steps.
2. Otherwise, invoke the application cache update process (page 455) for the given manifest URL, with the browsing context (page 414) being navigated, and with the resource's Document as the new master resource.

↪ **Otherwise**

Invoke the application cache selection algorithm (page 461) again, but without a manifest.

When the **application cache selection algorithm** is invoked *without* a manifest, then the user agent must run the first applicable set of steps from the following list:

If the resource is being loaded as part of navigation of a browsing context (page 414), and the resource was fetched from a particular application cache (page 448)

The user agent must associate the Document with that application cache and invoke the application cache update process (page 455) for that cache, with that browsing context (page 414).

If the resource is being loaded as part of navigation of a child browsing context (page 415)

The user agent must associate the Document with that application cache associated with the active document (page 414) of the parent browsing context (page 415).

Otherwise

Nothing special happens with respect to application caches.

5.7.5.1 Changes to the networking model

When a browsing context (page 414) is associated with an application cache (page 448), any and all resource loads must go through the following steps instead of immediately invoking the mechanisms appropriate to that resource's scheme:

1. If the resource is not to be fetched using the HTTP GET mechanism or equivalent, then fetch (page 59) the resource normally and abort these steps.
2. If the resource's URL is an master entry (page 449), the manifest (page 449), an explicit entry (page 449), a fallback entry (page 449), or a dynamic entry (page 449) in the application cache (page 448), then get the resource from the cache (instead of fetching it), and abort these steps.
3. If the resource's URL has the same origin (page 426) as the manifest's URL, and there is a fallback namespace (page 449) in the application cache (page 448) that is a prefix match (page 31) for the resource's URL, then:

Fetch (page 59) the resource normally. If this results in a redirect to a resource with another origin (page 423) (indicative of a captive portal), or a 4xx or 5xx status code or equivalent, or if there were network errors (but not if the user canceled the download), then instead get, from the cache, the resource of the fallback entry (page 449) corresponding to the namespace with the longest matching <path> (page 54) component. Abort these steps.

4. If there is an entry in the application cache (page 448)'s online whitelist (page 449) that has the same origin (page 426) as the resource's URL and that is a prefix match (page 31) for the resource's URL, then fetch (page 59) the resource normally and abort these steps.
5. Fail the resource load.

Note: The above algorithm ensures that resources that are not present in the manifest will always fail to load (at least, after the cache has been primed the first time), making the testing of offline applications simpler.

5.7.6 Application cache API

```
interface ApplicationCache {  
  
    // update status  
    const unsigned short UNCACHED = 0;  
    const unsigned short IDLE = 1;  
    const unsigned short CHECKING = 2;  
    const unsigned short DOWNLOADING = 3;  
    const unsigned short UPDATEREADY = 4;  
    const unsigned short OBSOLETE = 5;  
    readonly attribute unsigned short status;  
  
    // updates  
    void update();  
    void swapCache();  
  
    // dynamic entries  
    readonly attribute unsigned long length;  
    DOMString item(in unsigned long index);  
    void add(in DOMString url);  
    void remove(in DOMString url);  
  
    // events  
    attribute EventListener onchecking;  
    attribute EventListener onerror;  
    attribute EventListener onnoupdate;  
    attribute EventListener ondownloading;  
    attribute EventListener onprogress;  
    attribute EventListener onupdateready;  
    attribute EventListener oncached;  
    attribute EventListener onobsolete;  
  
};
```

Objects implementing the `ApplicationCache` interface must also implement the `EventTarget` interface.

There is a one-to-one mapping from `Document` objects to `ApplicationCache` objects. The `applicationCache` attribute on `Window` objects must return the `ApplicationCache` object associated with the active document (page 414) of the `Window`'s browsing context (page 414).

An `ApplicationCache` object might be associated with an application cache (page 448). When the `Document` object that the `ApplicationCache` object maps to is associated with an application cache, then that is the application cache with which the `ApplicationCache` object is associated. Otherwise, the `ApplicationCache` object is associated with the application cache that the `Document` object's browsing context (page 414) is associated with, if any.

The **status** attribute, on getting, must return the current state of the application cache (page 448) ApplicationCache object is associated with, if any. This must be the appropriate value from the following list:

UNCACHED (numeric value 0)

The ApplicationCache object is not associated with an application cache (page 448) at this time.

IDLE (numeric value 1)

The ApplicationCache object is associated with an application cache (page 448) whose group is in the *idle* update status (page 448) and the *mature* lifecycle status (page 449) and that application cache is the newest cache in its group.

CHECKING (numeric value 2)

The ApplicationCache object is associated with an application cache (page 448) whose group is in the *checking* update status (page 448).

DOWNLOADING (numeric value 3)

The ApplicationCache object is associated with an application cache (page 448) whose group is in the *downloading* update status (page 448).

UPDATEREADY (numeric value 4)

The ApplicationCache object is associated with an application cache (page 448) whose group is in the *idle* update status (page 448) and the *mature* lifecycle status (page 449) but that application cache is *not* the newest cache in its group.

OBSOLETE (numeric value 5)

The ApplicationCache object is associated with an application cache (page 448) whose group is in the *obsolete* lifecycle status (page 449).

The **length** attribute must return the number of dynamic entries (page 449) in the application cache (page 448) with which the ApplicationCache object is associated, if any, and zero if the object is not associated with any application cache.

The dynamic entries (page 449) in the application cache (page 448) are ordered in the same order as they were added to the cache by the `add()` method, with the oldest entry being the zeroth entry, and the most recently added entry having the index `length-1`.

The **item(*index*)** method must return the absolute URL (page 56) of the dynamic entry (page 449) with index *index* from the application cache (page 448), if one is associated with the ApplicationCache object. If the object is not associated with any application cache, or if the *index* argument is lower than zero or greater than `length-1`, the method must instead raise an `INDEX_SIZE_ERR` exception.

The **add(*url*)** method must run the following steps:

1. If the ApplicationCache object is not associated with any application cache, then raise an `INVALID_STATE_ERR` exception and abort these steps.
2. Resolve (page 55) the *url* argument. If this fails, raise a `SYNTAX_ERR` exception and abort these steps.

3. If there is already a resource in in the application cache (page 448) with which the ApplicationCache object is associated that has the address *url*, then ensure that entry is categorized as a dynamic entry (page 449) and return and abort these steps.
4. If *url* has a different <schema> (page 53) component than the manifest's URL, then raise a security exception (page 430).
5. Return, but do not abort these steps.
6. Fetch (page 59) the resource referenced by *url*.
7. If this results in a redirect, or a 4xx or 5xx status code or equivalent, or if there were network errors, or if the user canceled the download, then abort these steps.
8. Wait for there to be no running scripts, or at least no running scripts that can reach an ApplicationCache object associated with the application cache (page 448) with which this ApplicationCache object is associated.

Add the fetched resource to the application cache (page 448) and categorize it as a dynamic entry (page 449) before letting any such scripts resume.

** We can make the add() API more usable (i.e. make it possible to detect progress and distinguish success from errors without polling and timeouts) if we have the method return an object that is a target of Progress Events, much like the XMLHttpRequestEventTarget interface. This would also make this far more complex to spec and implement.

The **remove(*url*)** method must resolve (page 55) the *url* argument and, if that is successful, remove the dynamic entry (page 449) categorization of any entry whose address is the resulting absolute URL (page 56) in the application cache (page 448) with which the ApplicationCache object is associated. If this removes the last categorization of an entry in that cache, then the entry must be removed entirely (such that if it is re-added, it will be loaded from the network again). If the ApplicationCache object is not associated with any application cache, then the method must raise an INVALID_STATE_ERR exception instead.

If the **update()** method is invoked, the user agent must invoke the application cache update process (page 455), in the background, for the application cache (page 448) with which the ApplicationCache object is associated, but with no browsing context (page 414). If there is no such application cache, then the method must raise an INVALID_STATE_ERR exception instead.

If the **swapCache()** method is invoked, the user agent must run the following steps:

1. Let *document* be the Document with which the ApplicationCache object is associated.
2. Check that *document* is associated with an application cache (page 448). If it is not, then raise an INVALID_STATE_ERR exception and abort these steps.

Note: This is not the same thing as the ApplicationCache object being itself associated with an application cache (page 448)! In particular, the Document with which the ApplicationCache object is associated can only itself be associated with an application cache if it is in a top-level browsing context (page 415).

3. Let *cache* be the application cache (page 448) with which the ApplicationCache object is associated. (By definition, this is the same as the one that was found in the previous step.)
4. If the group of application caches (page 448) to which *cache* belongs has the lifecycle status (page 449) *obsolete*, unassociate *document* from *cache* and abort these steps.
5. Otherwise, check that there is an application cache in the same group as *cache* which has an entry categorized as a manifest (page 449) that is newer than *cache*. If there is not, then raise an INVALID_STATE_ERR exception and abort these steps.
6. Let *new cache* be the newest application cache (page 448) in the same group as *cache* which has an entry categorized as a manifest (page 449).
7. Unassociate *document* from *cache* and instead associate it with *new cache*.

The following are the event handler DOM attributes (page 432) that must be supported by objects implementing the ApplicationCache interface:

onchecking

Must be invoked whenever an checking event is targeted at or bubbles through the ApplicationCache object.

onerror

Must be invoked whenever an error event is targeted at or bubbles through the ApplicationCache object.

onnoupdate

Must be invoked whenever an noupdate event is targeted at or bubbles through the ApplicationCache object.

ondownloading

Must be invoked whenever an downloading event is targeted at or bubbles through the ApplicationCache object.

onprogress

Must be invoked whenever an progress event is targeted at or bubbles through the ApplicationCache object.

onupdateready

Must be invoked whenever an updateready event is targeted at or bubbles through the ApplicationCache object.

oncached

Must be invoked whenever a cached event is targeted at or bubbles through the ApplicationCache object.

onobsolete

Must be invoked whenever an obsolete event is targeted at or bubbles through the ApplicationCache object.

5.7.7 Browser state

The `navigator.onLine` attribute must return false if the user agent will not contact the network when the user follows links or when a script requests a remote page (or knows that such an attempt would fail), and must return true otherwise.

When the value that would be returned by the `navigator.onLine` attribute of the Window changes from true to false, the user agent must fire a simple event (page 436) called **offline** at the body element (page 81).

On the other hand, when the value that would be returned by the `navigator.onLine` attribute of the Window changes from false to true, the user agent must fire a simple event (page 436) called **online** at the body element (page 81).

5.8 Session history and navigation

5.8.1 The session history of browsing contexts

The sequence of Documents in a browsing context (page 414) is its **session history**.

History objects provide a representation of the pages in the session history of browsing contexts (page 414). Each browsing context has a distinct session history.

Each Document object in a browsing context's session history is associated with a unique instance of the History object, although they all must model the same underlying session history.

The `history` attribute of the Window interface must return the object implementing the History interface for that Window object's active document (page 414).

History objects represent their browsing context (page 414)'s session history as a flat list of session history entries (page 467). Each **session history entry** consists of either a URL (page 52) or a state object (page 467), or both, and may in addition have a title, a Document object, form data, a scroll position, and other information associated with it.

Note: This does not imply that the user interface need be linear. See the notes below (page 472).

URLs without associated state objects (page 467) are added to the session history as the user (or script) navigates from page to page.

A **state object** is an object representing a user interface state.

Pages can add (page 469) state objects (page 467) between their entry in the session history and the next ("forward") entry. These are then returned to the script (page 470) when the user (or script) goes back in the history, thus enabling authors to use the "navigation" metaphor even in one-page applications.

Every Document in the session history is defined to have a **last activated entry**, which is the state object (page 467) entry associated with that Document which was most recently activated. Initially, the last activated entry (page 467) of a Document must be the first entry for the Document, representing the fact that no state object (page 467) entry has yet been activated.

At any point, one of the entries in the session history is the **current entry**. This is the entry representing the active document (page 414) of the browsing context (page 414). The current entry (page 468) is usually an entry for the location (page 471) of the Document. However, it can also be one of the entries for state objects (page 467) added to the history by that document.

Entries that consist of state objects (page 467) share the same Document as the entry for the page that was active when they were added.

Contiguous entries that differ just by fragment identifier also share the same Document.

Note: All entries that share the same Document (and that are therefore merely different states of one particular document) are contiguous by definition.

User agents may discard (page 482) the Document objects of entries other than the current entry (page 468) that are not referenced from any script, reloading the pages afresh when the user or script navigates back to such pages. This specification does not specify when user agents should discard Document objects and when they should cache them.

Entries that have had their Document objects discarded must, for the purposes of the algorithms given below, act as if they had not. When the user or script navigates back or forwards to a page which has no in-memory DOM objects, any other entries that shared the same Document object with it must share the new object as well.

When state object entries are added, a URL can be provided. This URL is used to replace the state object entry if the Document is evicted.

When a user agent discards the Document object from an entry in the session history, it must also discard all the entries that share that Document but do not have an associated URL (i.e. entries that only have a state object (page 467)). Entries that shared that Document object but had a state object and have a different URL must then have their *state objects* removed. Removed entries are not recreated if the user or script navigates back to the page. If there are no state object entries for that Document object then no entries are removed.

5.8.2 The History interface

```
interface History {
    readonly attribute long length;
    void go(in long delta);
    void go();
    void back();
    void forward();
    void pushState(in DOMObject data, in DOMString title);
    void pushState(in DOMObject data, in DOMString title, in DOMString url);
    void clearState();
};
```

The **length** attribute of the History interface must return the number of entries in this session history (page 467).

The actual entries are not accessible from script.

The **go(delta)** method causes the UA to move the number of steps specified by *delta* in the session history.

If the index of the current entry (page 468) plus *delta* is less than zero or greater than or equal to the number of items in the session history (page 468), then the user agent must do nothing.

If the *delta* is zero, then the user agent must act as if the `location.reload()` method was called instead.

Otherwise, the user agent must cause the current browsing context (page 414) to traverse the history (page 480) to the specified entry. The *specified entry* is the one whose index equals the index of the current entry (page 468) plus *delta*.

When the user navigates through a browsing context (page 414), e.g. using a browser's back and forward buttons, the user agent must translate this action into the equivalent invocations of the `history.go(delta)` method on the various affected window objects.

Some of the other members of the History interface are defined in terms of the `go()` method, as follows:

Member	Definition
<code>go()</code>	Must do the same as <code>go(0)</code>
<code>back()</code>	Must do the same as <code>go(-1)</code>
<code>forward()</code>	Must do the same as <code>go(1)</code>

The **pushState(data, title, url)** method adds a state object to the history.

When this method is invoked, the user agent must run the following steps:

1. If a third argument is specified, run these substeps:
 1. Resolve (page 55) the value of the third argument.
 2. If that fails, raise a security exception (page 430) and abort the `pushState()` steps.
 3. Compare the resulting absolute URL (page 56) to the document's address. If any part of these two URLs (page 52) differ other than the `<path>` (page 54), `<query>` (page 54), and `<fragment>` (page 54) components, then raise a security exception (page 430) and abort the `pushState()` steps.

For the purposes of the comparison in the above substeps, the `<path>` (page 54) and `<query>` (page 54) components can only be the same if the URLs use a hierarchical `<scheme>` (page 53).

2. Remove from the session history (page 467) any entries for the Document from the entry after the current entry (page 468) up to the last entry in the session history that references the same Document object, if any. If the current entry (page 468) is the last entry in the session history, or if there are no entries after the current entry (page 468) that reference the same Document object, then no entries are removed.

3. Add a state object entry to the session history, after the current entry (page 468), with the specified *data* as the state object, the given *title* as the title, and, if the third argument is present, the absolute URL (page 56) that was found in the first step as the URL of the entry.
4. Set this new entry as being the last activated entry (page 467) for the Document.
5. Update the current entry (page 468) to be the this newly added entry.

Note: The title is purely advisory. User agents might use the title in the user interface.

User agents may limit the number of state objects added to the session history per page. If a page hits the UA-defined limit, user agents must remove the entry immediately after the first entry for that Document object in the session history after having added the new entry. (Thus the state history acts as a FIFO buffer for eviction, but as a LIFO buffer for navigation.)

The `clearState()` method removes all the state objects for the Document object from the session history.

When this method is invoked, the user agent must remove from the session history all the entries from the first state object entry for that Document object up to the last entry that references that same Document object, if any.

Then, if the current entry (page 468) was removed in the previous step, the current entry (page 468) must be set to the last entry for that Document object in the session history.

5.8.3 Activating state object entries

When an entry in the session history is activated (which happens during session traversal (page 480), as described above), the user agent must run the following steps:

1. First, the user agent must set this new entry as being the last activated entry (page 467) for the Document to which the entry belongs.
2. If the entry is a state object (page 467) entry, let *state* be that state object. Otherwise, the entry is the first entry for the Document; let *state* be null.
3. The user agent must then fire a `popstate` event in no namespace on the body element (page 81) using the PopStateEvent interface, with the *state* attribute set to the value of *state*. This event bubbles but is not cancelable and has no default action.

```
interface PopStateEvent : Event {
  readonly attribute DOMObject state;
  void initPopStateEvent(in DOMString typeArg, in boolean canBubbleArg,
  in boolean cancelableArg, in DOMObject stateArg);
  void initPopStateEventNS(in DOMString namespaceURIArg, in DOMString
  typeArg, in boolean canBubbleArg, in boolean cancelableArg, in DOMObject
  stateArg);
};
```

The **initPopStateEvent()** and **initPopStateEventNS()** methods must initialize the event in a manner analogous to the similarly-named methods in the DOM3 Events interfaces.
[DOM3EVENTS]

The **state** attribute represents the context information for the event, or null, if the state represented is the initial state of the Document.

5.8.4 The Location interface

Each Document object in a browsing context (page 414)'s session history is associated with a unique instance of a Location object.

The **location** attribute of the HTMLDocument interface must return the Location object for that Document object, if it is in a browsing context (page 414), and null otherwise.

The **location** attribute of the Window interface must return the Location object for that Window object's active document (page 414).

Location objects provide a representation of their document's address, and allow the current entry (page 468) of the browsing context (page 414)'s session history to be changed, by adding or replacing entries in the history object.

```
interface Location {
    readonly attribute DOMString href;
    void assign(in DOMString url);
    void replace(in DOMString url);
    void reload();

    // URL decomposition attributes
    attribute DOMString protocol;
    attribute DOMString host;
    attribute DOMString hostname;
    attribute DOMString port;
    attribute DOMString pathname;
    attribute DOMString search;
    attribute DOMString hash;
};
```

The **href** attribute must return the address of the page represented by the associated Document object, as an absolute URL (page 56).

On setting, the user agent must act as if the **assign()** method had been called with the new value as its argument.

When the **assign(*url*)** method is invoked, the UA must navigate (page 473) the browsing context (page 414) to the specified *url*.

When the **replace(*url*)** method is invoked, the UA must navigate (page 473) the browsing context (page 414) to the specified *url* with replacement enabled (page 476).

Navigation for the `assign()` and `replace()` methods must be done with the script browsing context (page 428) of the script that invoked the method as the source browsing context (page 473).

The Location interface also has the complement of URL decomposition attributes (page 57), **protocol**, **host**, **port**, **hostname**, **pathname**, **search**, and **hash**. These must follow the rules given for URL decomposition attributes, with the input (page 57) being the address of the page represented by the associated Document object, as an absolute URL (page 56) (same as the `href` attribute), and the common setter action (page 57) being the same as setting the `href` attribute to the new output value.

5.8.4.1 Security

User agents must raise a security exception (page 430) whenever any of the members of a Location object are accessed by scripts whose effective script origin (page 423) is not the same (page 426) as the Location object's associated Document's effective script origin (page 423), with the following exceptions:

- The `href` setter, if the script is running in a browsing context (page 414) that is allowed to navigate (page 417) the browsing context with which the Location object is associated

User agents must not allow scripts to override the `href` attribute's setter.

5.8.5 Implementation notes for session history

This section is non-normative.

The History interface is not meant to place restrictions on how implementations represent the session history to the user.

For example, session history could be implemented in a tree-like manner, with each page having multiple "forward" pages. This specification doesn't define how the linear list of pages in the history object are derived from the actual session history as seen from the user's perspective.

Similarly, a page containing two iframes has a history object distinct from the iframes' history objects, despite the fact that typical Web browsers present the user with just one "Back" button, with a session history that interleaves the navigation of the two inner frames and the outer page.

Security: It is suggested that to avoid letting a page "hijack" the history navigation facilities of a UA by abusing `pushState()`, the UA provide the user with a way to jump back to the previous page (rather than just going back to the previous state). For example, the back button could have a drop down showing just the pages in the session history, and not showing any of the states. Similarly, an aural browser could have two "back" commands, one that goes back to the previous state, and one that jumps straight back to the previous page.

In addition, a user agent could ignore calls to `pushState()` that are invoked on a timer, or from event handlers that do not represent a clear user action, or that are invoked in rapid succession.

5.9 Browsing the Web

5.9.1 Navigating across documents

Certain actions cause the browsing context (page 414) to **navigate** to a new resource. Navigation always involves **source browsing context**, which is the browsing context which was responsible for starting the navigation.

For example, following a hyperlink (page 498), form submission (page 368), and the `window.open()` and `location.assign()` methods can all cause a browsing context to navigate.

A user agent may provide various ways for the user to explicitly cause a browsing context to navigate, in addition to those defined in this specification.

When a browsing context is navigated to a new resource, the user agent must run the following steps:

1. If the source browsing context (page 473) is not the same as the browsing context (page 414) being navigated, and the source browsing context (page 473) is not one of the ancestor browsing contexts (page 415) of the browsing context (page 414) being navigated, and the source browsing context (page 473) has its sandboxed navigation browsing context flag (page 218) set, then abort these steps. The user agent may offer to open the new resource in a new top-level browsing context (page 415) or in the top-level browsing context (page 415) of the source browsing context (page 473), at the user's option, in which case the user agent must navigate (page 473) that designated top-level browsing context (page 415) to the new resource as if the user had requested it independently.
2. If the source browsing context (page 473) is the same as the browsing context (page 414) being navigated, and this browsing context has its seamless browsing context flag (page 219) set, then find the nearest ancestor browsing context (page 415) that does not have its seamless browsing context flag (page 219) set, and continue these steps as if *that* browsing context (page 414) was the one that was going to be navigated (page 473) instead.
3. Cancel any preexisting attempt to navigate the browsing context (page 414).
4. Resolve (page 55) the URL (page 52) of the new resource. If that fails, the user agent may abort these steps, or may treat the URL as identifying some sort of user-agent defined error resource, which could display some sort of inline content, or could be handled using a mechanism that does not affect the browsing context.
5. *Fragment identifiers:* If the absolute URL (page 56) of the new resource is the same as the address of the active document (page 414) of the browsing context (page 414) being navigated, ignoring any `<fragment>` (page 54) components of those URLs (page 52), and the new resource is to be fetched using HTTP GET or equivalent, then navigate to that fragment identifier (page 479) and abort these steps.
6. If the new resource is to be handled by displaying some sort of inline content, e.g. an error message because the specified scheme is not one of the supported protocols, or an inline prompt to allow the user to select a registered handler (page 444) for the given scheme, then display the inline content (page 479) and abort these steps.

7. If the new resource is to be handled using a mechanism that does not affect the browsing context, e.g. ignoring the navigation request altogether because the specified scheme is not one of the supported protocols, then abort these steps and proceed with that mechanism instead.
8. If the new resource is to be fetched using HTTP GET or equivalent, then check if there are any relevant application caches (page 449) that are identified by a URL with the same origin (page 426) as the URL in question, and that have this URL as one of their entries, excluding entries marked as foreign (page 449). If so, then the user agent must then get the resource from the most appropriate application cache (page 449) of those that match.

Otherwise, fetch (page 59) the new resource. If this results in a redirect, return to the step labeled "fragment identifiers" (page 473) with the new resource.

For example, imagine an HTML page with an associated application cache displaying an image and a form, where the image is also used by several other application caches. If the user right-clicks on the image and chooses "View Image", then the user agent could decide to show the image from any of those caches, but it is likely that the most useful cache for the user would be the one that was used for the aforementioned HTML page. On the other hand, if the user submits the form, and the form does a POST submission, then the user agent will not use an application cache at all; the submission will be made to the network.

9. If fetching the resource is synchronous (i.e. for javascript: URLs (page 430) and about:blank), then this must be synchronous, but if fetching the resource depends on external resources, as it usually does for URLs that use HTTP or other networking protocols, then at this point the user agents must yield to whatever script invoked the navigation steps, if they were invoked by script.
10. Wait for one or more bytes to be available or for the user agent to establish that the resource in question is empty. During this time, the user agent may allow the user to cancel this navigation attempt or start other navigation attempts.
11. If the resource was not fetched from an application cache (page 448), and was to be fetched using HTTP GET or equivalent, and its URL matches the fallback namespace (page 460) of one or more relevant application caches (page 449), and the user didn't cancel the navigation attempt during the previous step, and the navigation attempt failed (e.g. the server returned a 4xx or 5xx status code or equivalent, or there was a DNS error), then:

Let *candidate* be the fallback resource (page 449) specified for the fallback namespace (page 449) in question. If multiple application caches match, the user agent must use the fallback of the most appropriate application cache (page 449) of those that match.

If *candidate* is not marked as foreign (page 449), then the user agent must discard the failed load and instead continue along these steps using *candidate* as the resource.

For the purposes of session history (and features that depend on session history, e.g. bookmarking) the user agent must use the URL of the resource that was requested (the one that matched the fallback namespace (page 449)), not the fallback resource. However, the user agent may indicate to the user that the original page

load failed, that the page used was a fallback resource, and what the URL of the fallback resource actually is.

12. If the document's out-of-band metadata (e.g. HTTP headers), not counting any type information (page 60) (such as the Content-Type HTTP header), requires some sort of processing that will not affect the browsing context, then perform that processing and abort these steps.

Such processing might be triggered by, amongst other things, the following:

- ***HTTP status codes (e.g. 204 No Content or 205 Reset Content)***
- ***HTTP Content-Disposition headers***
- ***Network errors***

13. Let *type* be the sniffed type of the resource (page 61).
14. If the user agent has been configured to process resources of the given *type* using some mechanism other than rendering the content in a browsing context (page 414), then skip this step. Otherwise, if the *type* is one of the following types, jump to the appropriate entry in the following list, and process the resource as described there:
 - ↪ **"text/html"**
Follow the steps given in the HTML document (page 476) section, and abort these steps.
 - ↪ **Any type ending in "+xml"**
 - ↪ **"application/xml"**
 - ↪ **"text/xml"**
Follow the steps given in the XML document (page 477) section. If that section determines that the content is *not* to be displayed as a generic XML document, then proceed to the next step in this overall set of steps.
Otherwise, abort these steps.
 - ↪ **"text/plain"**
Follow the steps given in the plain text file (page 478) section, and abort these steps.
 - ↪ **A supported image type**
Follow the steps given in the image (page 478) section, and abort these steps.
 - ↪ **A type that will use an external application to render the content in the browsing context (page 414)**
Follow the steps given in the plugin (page 479) section, and abort these steps.
15. *Non-document content*: If, given *type*, the new resource is to be handled by displaying some sort of inline content, e.g. a native rendering of the content, an error message because the specified type is not supported, or an inline prompt to allow the user to select a registered handler (page 444) for the given type, then display the inline content (page 479) and abort these steps.

16. Otherwise, the document's type is such that the resource will not affect the browsing context, e.g. because the resource is to be handed to an external application.
Process the resource appropriately.

Some of the sections below, to which the above algorithm defers in certain cases, require the user agent to **update the session history with the new page**. When a user agent is required to do this, it must follow the set of steps given below that is appropriate for the situation at hand. From the point of view of any script, these steps must occur atomically.

- ** 1. pause for scripts -- but don't use the "pause" definition since that involves not running script!
- ** 2. onbeforeunload, and if present set flag that we will kill document
- ** 3. onunload, and if present set flag that we will kill document
- ** 4. if flag is set: discard the Document (page 482)

5. If the navigation was initiated for entry update of an entry

1. Replace the entry being updated with a new entry representing the new resource and its Document object and related state. The user agent may propagate state from the old entry to the new entry (e.g. scroll position).
2. Traverse the history (page 480) to the new entry.

Otherwise

1. Remove all the entries after the current entry (page 468) in the browsing context (page 414)'s Document object's History object.

Note: This doesn't necessarily have to affect (page 472) the user agent's user interface.

2. Append a new entry at the end of the History object representing the new resource and its Document object and related state.
3. Traverse the history (page 480) to the new entry.
4. If the navigation was initiated with **replacement enabled**, remove the entry immediately before the new current entry (page 468) in the session history.

5.9.2 Page load processing model for HTML files

When an HTML document is to be loaded in a browsing context (page 414), the user agent must create a Document object, mark it as being an HTML document (page 76), create an HTML parser (page 582), associate it with the document, and begin to use the bytes provided for the document as the input stream (page 584) for that parser.

Note: The input stream (page 584) converts bytes into characters for use in the tokeniser. This process relies, in part, on character encoding information found in the real Content-Type metadata (page 60) of the resource; the "sniffed type" is not used for this purpose.

When no more bytes are available, an EOF character is implied, which eventually causes a load event to be fired.

After creating the Document object, but potentially before the page has finished parsing, the user agent must update the session history with the new page (page 476).

Note: Application cache selection (page 460) happens in the HTML parser (page 623).

5.9.3 Page load processing model for XML files

When faced with displaying an XML file inline, user agents must first create a Document object, following the requirements of the XML and Namespaces in XML recommendations, RFC 3023, DOM3 Core, and other relevant specifications. [XML] [XMLNS] [RFC3023] [DOM3CORE]

The actual HTTP headers and other metadata, not the headers as mutated or implied by the algorithms given in this specification, are the ones that must be used when determining the character encoding according to the rules given in the above specifications. Once the character encoding is established, the document's character encoding (page 79) must be set to that character encoding.

If the root element, as parsed according to the XML specifications cited above, is found to be an `html` element with an attribute `manifest`, then, as soon as the element is inserted into the document (page 24), the user agent must resolve (page 55) the value of that attribute, and if that is successful, must run the application cache selection algorithm (page 460) with the resulting absolute URL (page 56) as the manifest URL. Otherwise, if the attribute is absent or resolving it fails, then as soon as the root element is inserted into the document (page 24), the user agent must run the application cache selection algorithm (page 461) with no manifest.

Note: Because the processing of the `manifest` attribute happens only once the root element is parsed, any URLs referenced by processing instructions before the root element (such as `<?xml-stylesheet?>` and `<?xbl?>` PIs) will be fetched from the network and cannot be cached.

User agents may examine the namespace of the root Element node of this Document object to perform namespace-based dispatch to alternative processing tools, e.g. determining that the content is actually a syndication feed and passing it to a feed handler. If such processing is to take place, abort the steps in this section, and jump to the next step (page 475) (labeled "non-document content") in the navigate (page 473) steps above.

Otherwise, then, with the newly created Document, the user agents must update the session history with the new page (page 476). User agents may do this before the complete document has been parsed (thus achieving *incremental rendering*).

Error messages from the parse process (e.g. XML namespace well-formedness errors) may be reported inline by mutating the Document.

5.9.4 Page load processing model for text files

When a plain text document is to be loaded in a browsing context (page 414), the user agent should create a Document object, mark it as being an HTML document (page 76), create an HTML parser (page 582), associate it with the document, act as if the tokeniser had emitted a start tag token with the tag name "pre", set the tokenization (page 597) stage's content model flag (page 597) to *PLAINTEXT*, and begin to pass the stream of characters in the plain text document to that tokeniser.

The rules for how to convert the bytes of the plain text document into actual characters are defined in RFC 2046, RFC 2646, and subsequent versions thereof. [RFC2046] [RFC2646]

The document's character encoding (page 79) must be set to the character encoding used to decode the document.

Upon creation of the Document object, the user agent must run the application cache selection algorithm (page 461) with no manifest.

When no more character are available, an EOF character is implied, which eventually causes a load event to be fired.

After creating the Document object, but potentially before the page has finished parsing, the user agent must update the session history with the new page (page 476).

User agents may add content to the head element of the Document, e.g. linking to stylesheet or an XBL binding, providing script, giving the document a title, etc.

5.9.5 Page load processing model for images

When an image resource is to be loaded in a browsing context (page 414), the user agent should create a Document object, mark it as being an HTML document (page 76), append an html element to the Document, append a head element and a body element to the html element, append an img to the body element, and set the src attribute of the img element to the address of the image.

Then, the user agent must act as if it had stopped parsing (page 656).

Upon creation of the Document object, the user agent must run the application cache selection algorithm (page 461) with no manifest.

After creating the Document object, but potentially before the page has finished fully loading, the user agent must update the session history with the new page (page 476).

User agents may add content to the head element of the Document, or attributes to the img element, e.g. to link to stylesheet or an XBL binding, to provide a script, to give the document a title, etc.

5.9.6 Page load processing model for content that uses plugins

When a resource that requires an external resource to be rendered is to be loaded in a browsing context (page 414), the user agent should create a Document object, mark it as being an HTML document (page 76), append an `html` element to the Document, append a `head` element and a `body` element to the `html` element, append an `embed` to the `body` element, and set the `src` attribute of the `embed` element to the address of the resource.

Then, the user agent must act as if it had stopped parsing (page 656).

Upon creation of the Document object, the user agent must run the application cache selection algorithm (page 461) with no manifest.

After creating the Document object, but potentially before the page has finished fully loading, the user agent must update the session history with the new page (page 476).

User agents may add content to the `head` element of the Document, or attributes to the `embed` element, e.g. to link to stylesheet or an XBL binding, or to give the document a title.

Note: If the *sandboxed plugins browsing context flag* (page 218) is set on the browsing context (page 414), the synthesized `embed` element will fail to render the content (page 222).

5.9.7 Page load processing model for inline content that doesn't have a DOM

When the user agent is to display a user agent page inline in a browsing context (page 414), the user agent should create a Document object, mark it as being an HTML document (page 76), and then either associate that Document with a custom rendering that is not rendered using the normal Document rendering rules, or mutate that Document until it represents the content the user agent wants to render.

Once the page has been set up, the user agent must act as if it had stopped parsing (page 656).

Upon creation of the Document object, the user agent must run the application cache selection algorithm (page 461) with no manifest.

After creating the Document object, but potentially before the page has been completely set up, the user agent must update the session history with the new page (page 476).

5.9.8 Navigating to a fragment identifier

When a user agent is supposed to navigate to a fragment identifier, then the user agent must update the session history with the new page (page 476), where "the new page" has the same Document as before but with the URL having the newly specified fragment identifier.

Part of that algorithm involves the user agent having to scroll to the fragment identifier (page 479), which is the important part for this step.

When the user agent is required to **scroll to the fragment identifier**, it must change the scrolling position of the document, or perform some other action, such that the indicated part

of the document (page 480) is brought to the user's attention. If there is no indicated part, then the user agent must not scroll anywhere.

The indicated part of the document is the one that the fragment identifier, if any, identifies. The semantics of the fragment identifier in terms of mapping it to a specific DOM Node is defined by the MIME type specification of the document's MIME Type (for example, the processing of fragment identifiers for XML MIME types is the responsibility of RFC3023).

For HTML documents (and the text/html MIME type), the following processing model must be followed to determine what the indicated part of the document (page 480) is.

1. Parse (page 53) the URL (page 52), and let *fragid* be the <fragment> (page 54) component of the URL.
2. If *fragid* is the empty string, then the indicated part of the document is the top of the document.
3. If there is an element in the DOM that has an ID exactly equal to *fragid*, then the first such element in tree order is the indicated part of the document (page 480); stop the algorithm here.
4. If there is an element in the DOM that has a name attribute whose value is exactly equal to *fragid*, then the first such element in tree order is the indicated part of the document (page 480); stop the algorithm here.
5. Otherwise, there is no indicated part of the document.

For the purposes of the interaction of HTML with Selectors' :target pseudo-class, the *target element* is the indicated part of the document (page 480), if that is an element; otherwise there is no *target element*. [SELECTORS]

5.9.9 History traversal

When a user agent is required to **traverse the history** to a *specified entry*, the user agent must act as follows:

1. If there is no longer a Document object for the entry in question, the user agent must navigate (page 473) the browsing context to the location for that entry to perform an entry update (page 476) of that entry, and abort these steps. The "navigate (page 473)" algorithm reinvokes this "traverse" algorithm to complete the traversal, at which point there *is* a Document object and so this step gets skipped. The navigation must be done using the same source browsing context (page 473) as was used the first time this entry was created.
2. If appropriate, update the current entry (page 468) in the browsing context (page 414)'s Document object's History object to reflect any state that the user agent wishes to persist.

For example, some user agents might want to persist the scroll position, or the values of form controls.
3. If the *specified entry* has a different Document object than the current entry (page 468) then the user agent must run the following substeps:

**

1. **freeze any timers, intervals, XMLHttpRequests, database transactions, etc**
2. The user agent must move any properties that have been added to the browsing context's default view's Window object to the active document (page 414)'s Document's list of added properties (page 421).
3. If the browsing context is a top-level browsing context (page 415) (and not an auxiliary browsing context (page 416)), and the origin (page 423) of the Document of the *specified entry* is not the same (page 426) as the origin (page 423) of the Document of the current entry (page 468), then the following sub-sub-steps must be run:
 1. The current browsing context name (page 417) must be stored with all the entries in the history that are associated with Document objects with the same origin (page 426) as the active document (page 414) and that are contiguous with the current entry (page 468).
 2. The browsing context's browsing context name (page 417) must be unset.
4. The user agent must make the *specified entry*'s Document object the active document (page 414) of the browsing context (page 414). (If it is a top-level browsing context (page 415), this might change (page 449) which application cache (page 448) it is associated with.)
5. If the *specified entry* has a browsing context name (page 417) stored with it, then the following sub-sub-steps must be run:
 1. The browsing context's browsing context name (page 417) must be set to the name stored with the specified entry.
 2. Any browsing context name (page 417) stored with the entries in the history that are associated with Document objects with the same origin (page 426) as the new active document (page 414), and that are contiguous with the specified entry, must be cleared.
6. The user agent must move any properties that have been added to the active document (page 414)'s Document's list of added properties (page 421) to browsing context's default view's Window object.

**

**

7. **unfreeze any timers, intervals, XMLHttpRequests, database transactions, etc**

4. If there are any entries with state objects between the last activated entry (page 467) for the Document of the *specified entry* and the *specified entry* itself (not inclusive), then the user agent must iterate through every entry between that last activated entry (page 467) and the *specified entry*, starting with the entry closest to the current entry (page 468), and ending with the one closest to the *specified entry*. For each entry, if the entry is a state object, the user agent must activate the state object (page 470).
5. If the *specified entry* is a state object or the first entry for a Document, the user agent must activate that entry (page 470).

6. If the *specified entry* has a URL that differs from the current entry (page 468)'s only by its fragment identifier, and the two share the same Document object, then fire a simple event (page 436) with the name hashchange at the body element (page 81), and, if the new URL has a fragment identifier, scroll to the fragment identifier (page 479).
7. User agents may also update other aspects of the document view when the location changes in this way, for instance the scroll position, values of form fields, etc.
8. The current entry (page 468) is now the *specified entry*.

** how does the changing of the global attributes affect .watch() when seen from other Windows?

5.9.10 Closing a browsing context

- ** Closing a browsing context and discarding it (vs closing it and keeping it around in memory).
- ** when a browsing context is closed, all session history entries' Document objects must be discarded.
- ** When a user agent is to **discard a Document**, any frozen timers, intervals, XMLHttpRequests, database transactions, etc, must be killed, and any MessagePorts owned by the Window object must be unentangled.
- ** Also, unload events should fire.

5.10 Structured client-side storage

5.10.1 Storing name/value pairs

5.10.1.1 Introduction

This section is non-normative.

This specification introduces two related mechanisms, similar to HTTP session cookies, for storing structured data on the client side. [RFC2109] [RFC2965]

The first is designed for scenarios where the user is carrying out a single transaction, but could be carrying out multiple transactions in different windows at the same time.

Cookies don't really handle this case well. For example, a user could be buying plane tickets in two different windows, using the same site. If the site used cookies to keep track of which ticket the user was buying, then as the user clicked from page to page in both windows, the ticket currently being purchased would "leak" from one window to the other, potentially causing the user to buy two tickets for the same flight without really noticing.

To address this, this specification introduces the `sessionStorage` DOM attribute. Sites can add data to the session storage, and it will be accessible to any page from the same site opened in that window.

For example, a page could have a checkbox that the user ticks to indicate that he wants insurance:

```
<label>
  <input type="checkbox" onchange="sessionStorage.insurance = checked">
    I want insurance on this trip.
</label>
```

A later page could then check, from script, whether the user had checked the checkbox or not:

```
if (sessionStorage.insurance) { ... }
```

If the user had multiple windows opened on the site, each one would have its own individual copy of the session storage object.

The second storage mechanism is designed for storage that spans multiple windows, and lasts beyond the current session. In particular, Web applications may wish to store megabytes of user data, such as entire user-authored documents or a user's mailbox, on the client side for performance reasons.

Again, cookies do not handle this case well, because they are transmitted with every request.

The `localStorage` DOM attribute is used to access a page's local storage area.

The site at `example.com` can display a count of how many times the user has loaded its page by putting the following at the bottom of its page:

```
<p>
  You have viewed this page
  <span id="count">an untold number of</span>
  time(s).
</p>
<script>
  if (!localStorage.pageLoadCount)
    localStorage.pageLoadCount = 0;
  localStorage.pageLoadCount = parseInt(localStorage.pageLoadCount,
  10) + 1;
  document.getElementById('count').textContent =
  localStorage.pageLoadCount;
</script>
```

Each site has its own separate storage area.

Storage areas (both session storage and local storage) store strings. To store structured data in a storage area, you must first convert it to a string.

5.10.1.2 The Storage interface

```
interface Storage {
    readonly attribute unsigned long length;
    [IndexGetter] DOMString key(in unsigned long index);
    [NameGetter] DOMString getItem(in DOMString key);
    [NameSetter] void setItem(in DOMString key, in DOMString data);
    [XXX] void removeItem(in DOMString key);
    void clear();
};
```

Each Storage object provides access to a list of key/value pairs, which are sometimes called items. Keys and values are strings. Any string (including the empty string) is a valid key.

Note: To store more structured data, authors may consider using the SQL interfaces (page 488) instead.

Each Storage object is associated with a list of key/value pairs when it is created, as defined in the sections on the `sessionStorage` and `localStorage` attributes. Multiple separate objects implementing the Storage interface can all be associated with the same list of key/value pairs simultaneously.

The `length` attribute must return the number of key/value pairs currently present in the list associated with the object.

The `key(n)` method must return the name of the *n*th key in the list. The order of keys is user-agent defined, but must be consistent within an object between changes to the number of keys. (Thus, adding (page 484) or removing (page 484) a key may change the order of the keys, but merely changing the value of an existing key must not.) If *n* is less than zero or greater than or equal to the number of key/value pairs in the object, then this method must raise an `INDEX_SIZE_ERR` exception.

The `getItem(key)` method must return the current value associated with the given *key*. If the given key does not exist in the list associated with the object then this method must return null.

The `setItem(key, value)` method must first check if a key/value pair with the given *key* already exists in the list associated with the object.

If it does not, then a new key/value pair must be added to the list, with the given *key* and *value*.

If the given *key* does exist in the list, then it must have its value updated to the value given in the *value* argument.

If it couldn't set the new value, the method must raise an `INVALID_ACCESS_ERR` exception. (Setting could fail if, e.g., the user has disabled storage for the domain, or if the quota has been exceeded.)

The `removeItem(key)` method must cause the key/value pair with the given *key* to be removed from the list associated with the object, if it exists. If no item with that key exists, the method must do nothing.

The `setItem()` and `removeItem()` methods must be atomic with respect to failure. That is, changes to the data storage area must either be successful, or the data storage area must not be changed at all.

The `clear()` method must atomically cause the list associated with the object to be emptied of all key/value pairs.

When the `setItem()`, `removeItem()`, and `clear()` methods are invoked, events are fired on other `HTMLDocument` objects that can access the newly stored or removed data, as defined in the sections on the `sessionStorage` and `localStorage` attributes.

5.10.1.3 The `sessionStorage` attribute

The `sessionStorage` attribute represents the set of storage areas specific to the current top-level browsing context (page 415).

Each top-level browsing context (page 415) has a unique set of session storage areas, one for each origin (page 423).

User agents should not expire data from a browsing context's session storage areas, but may do so when the user requests that such data be deleted, or when the UA detects that it has limited storage space, or for security reasons. User agents should always avoid deleting data while a script that could access that data is running. When a top-level browsing context is destroyed (and therefore permanently inaccessible to the user) the data stored in its session storage areas can be discarded with it, as the API described in this specification provides no way for that data to ever be subsequently retrieved.

Note: *The lifetime of a browsing context can be unrelated to the lifetime of the actual user agent process itself, as the user agent may support resuming sessions after a restart.*

When a new `HTMLDocument` is created, the user agent must check to see if the document's top-level browsing context (page 415) has allocated a session storage area for that document's origin (page 423). If it has not, a new storage area for that document's origin (page 423) must be created.

The `Storage` object for the document's associated `Window` object's `sessionStorage` attribute must then be associated with that origin (page 423)'s session storage area for that top-level browsing context (page 415).

When a new top-level browsing context (page 415) is created by cloning an existing browsing context (page 414), the new browsing context must start with the same session storage areas as the original, but the two sets must from that point on be considered separate, not affecting each other in any way.

When a new top-level browsing context (page 415) is created by a script in an existing browsing context (page 414), or by the user following a link in an existing browsing context, or in some other way related to a specific `HTMLDocument`, then the session storage area of the origin (page 423) of that `HTMLDocument` must be copied into the new browsing context when it is created. From that point on, however, the two session storage areas must be considered separate, not affecting each other in any way.

When the `setItem()`, `removeItem()`, and `clear()` methods are called on a Storage object `x` that is associated with a session storage area, then in every `HTMLDocument` object whose `Window` object's `sessionStorage` attribute's Storage object is associated with the same storage area, other than `x`, a storage event must be fired, as described below (page 486).

5.10.1.4 The `localStorage` attribute

The `localStorage` object provides a Storage object for an origin (page 423).

User agents must have a set of local storage areas, one for each origin (page 423).

User agents should expire data from the local storage areas only for security reasons or when requested to do so by the user. User agents should always avoid deleting data while a script that could access that data is running. Data stored in local storage areas should be considered potentially user-critical. It is expected that Web applications will use the local storage areas for storing user-written documents.

When the `localStorage` attribute is accessed, the user agent must check to see if it has allocated local storage area for the for the origin (page 423) of the active document (page 414) of the browsing context (page 414) of the `Window` object on which the method was invoked. If it has not, a new storage area for that origin (page 423) must be created.

The user agent must then create a Storage object associated with that origin's local storage area, and return it.

When the `setItem()`, `removeItem()`, and `clear()` methods are called on a Storage object `x` that is associated with a local storage area, then in every `HTMLDocument` object whose `Window` object's `localStorage` attribute's Storage object is associated with the same storage area, other than `x`, a storage event must be fired, as described below (page 486).

5.10.1.5 The `storage` event

The `storage` event is fired in an `HTMLDocument` when a storage area changes, as described in the previous two sections (for session storage (page 486), for local storage (page 486)).

When this happens, the user agent must dispatch an event with the name `storage`, with no namespace, which does not bubble but is cancelable, and which uses the `StorageEvent`, at the `body` element (page 81) of each active (page 414) `HTMLDocument` object affected.

If the event is being fired due to an invocation of the `setItem()` or `removeItem()` methods, the event must have its `key` attribute set to the name of the key in question, its `oldValue` attribute set to the old value of the key in question, or null if the key is newly added, and its `newValue` attribute set to the new value of the key in question, or null if the key was removed.

Otherwise, if the event is being fired due to an invocation of the `clear()` method, the event must have its `key`, `oldValue`, and `newValue` attributes set to null.

In addition, the event must have its `url` attribute set to the address of the page whose Storage object was affected, and its `source` attribute set to the `Window` object of the browsing context (page 414) that that document is in, if the two documents are in the same unit of related browsing contexts (page 417), or null otherwise.

5.10.1.5.1 Event definition

```
interface StorageEvent : Event {  
    readonly attribute DOMString key;  
    readonly attribute DOMString oldValue;  
    readonly attribute DOMString newValue;  
    readonly attribute DOMString url;  
    readonly attribute Window source;  
    void initStorageEvent(in DOMString typeArg, in boolean canBubbleArg,  
    in boolean cancelableArg, in DOMString keyArg, in DOMString oldValueArg,  
    in DOMString newValueArg, in DOMString urlArg, in Window sourceArg);  
    void initStorageEventNS(in DOMString namespaceURI, in DOMString  
    typeArg, in boolean canBubbleArg, in boolean cancelableArg, in DOMString  
    keyArg, in DOMString oldValueArg, in DOMString newValueArg, in DOMString  
    urlArg, in Window sourceArg);  
};
```

The **initStorageEvent()** and **initStorageEventNS()** methods must initialize the event in a manner analogous to the similarly-named methods in the DOM3 Events interfaces.

[DOM3EVENTS]

The **key** attribute represents the key being changed.

The **oldValue** attribute represents the old value of the key being changed.

The **newValue** attribute represents the new value of the key being changed.

The **url** attribute represents the address of the document that changed the key.

The **source** attribute represents the Window that changed the key.

5.10.1.6 Threads

Multiple browsing contexts must be able to access the local storage areas simultaneously in a predictable manner. Scripts must not be able to detect any concurrent script execution.

This is required to guarantee that the `length` attribute of a `Storage` object never changes while a script is executing, other than in a way that is predictable by the script itself.

Note: There are various ways of implementing this requirement. One is to just have one event loop (page 429) for all browsing contexts (page 414). Another is that if a script running in one browsing context accesses a storage area, the user agent blocks scripts in other browsing contexts when they try to access the same storage area until the event loop (page 429) running the first script has completed running the task that started that script. Another (potentially more efficient but certainly more complex) implementation strategy is to use optimistic transactional script execution. This specification does not require any particular implementation strategy, so long as the requirement above is met.

5.10.2 Database storage

5.10.2.1 Introduction

This section is non-normative.

**

...

5.10.2.2 Databases

Each *origin* (page 423) has an associated set of databases. Each database has a name and a current version. There is no way to enumerate or delete the databases available for a domain from this API.

Note: *Each database has one version at a time, a database can't exist in multiple versions at once. Versions are intended to allow authors to manage schema changes incrementally and non-destructively, and without running the risk of old code (e.g. in another browser window) trying to write to a database with incorrect assumptions.*

The `openDatabase()` method returns a `Database` object. The method takes four arguments: a database name, a database version, a display name, and an estimated size, in bytes, of the data that will be stored in the database.

The `openDatabase()` method must use and create databases from the *origin* (page 423) of the active document (page 414) of the browsing context (page 414) of the `Window` object on which the method was invoked.

If the database version provided is not the empty string, and the database already exists but has a different version, then the method must raise an `INVALID_STATE_ERR` exception.

The user agent may also raise a security exception (page 430) in case the request violates a policy decision (e.g. if the user agent is configured to not allow the page to open databases).

Otherwise, if the database version provided is the empty string, or if the database doesn't yet exist, or if the database exists and the version provided to the `openDatabase()` method is the same as the current version associated with the database, then the method must return a `Database` object representing the database that has the name that was given. If no such database exists, it must be created first.

All strings including the empty string are valid database names. Database names must be compared in a case-sensitive (page 31) manner.

Note: *Implementations can support this even in environments that only support a subset of all strings as database names by mapping database names (e.g. using a hashing algorithm) to the supported set of names.*

User agents are expected to use the display name and the estimated database size to optimize the user experience. For example, a user agent could use the estimated size to suggest an initial quota to the user. This allows a site that is aware that it will try to use hundreds of megabytes to declare this upfront, instead of the user agent prompting the user for permission to increase the quota every five megabytes.

```

interface Database {
    void transaction(in SQLTransactionCallback callback);
    void transaction(in SQLTransactionCallback callback, in
SQLTransactionErrorCallback errorCallback);
    void transaction(in SQLTransactionCallback callback, in
SQLTransactionErrorCallback errorCallback, in VoidCallback
successCallback);

    readonly attribute DOMString version;
    void changeVersion(in DOMString oldVersion, in DOMString newVersion,
in SQLTransactionCallback callback, in SQLTransactionErrorCallback
errorCallback, in VoidCallback successCallback);
};

interface SQLTransactionCallback {
    void handleEvent(in SQLTransaction transaction);
};

interface SQLTransactionErrorCallback {
    void handleEvent(in SQLError error);
};

```

The **transaction()** method takes one or two arguments. When called, the method must immediately return and then asynchronously run the transaction steps (page 493) with the *transaction callback* being the first argument, the *error callback* being the second argument, if any, the *success callback* being the third argument, if any, and with no *preflight operation* or *postflight operation*.

The version that the database was opened with is the **expected version** of this Database object. It can be the empty string, in which case there is no expected version — any version is fine.

On getting, the **version** attribute must return the current version of the database (as opposed to the expected version (page 489) of the Database object).

The **changeVersion()** method allows scripts to atomically verify the version number and change it at the same time as doing a schema update. When the method is invoked, it must immediately return, and then asynchronously run the transaction steps (page 493) with the *transaction callback* being the third argument, the *error callback* being the fourth argument, the *success callback* being the fifth argument, the *preflight operation* being the following:

1. Check that the value of the first argument to the `changeVersion()` method exactly matches the database's actual version. If it does not, then the *preflight operation* fails.

...and the *postflight operation* being the following:

1. Change the database's actual version to the value of the second argument to the `changeVersion()` method.
2. Change the Database object's expected version to the value of the second argument to the `changeVersion()` method.

5.10.2.3 Executing SQL statements

The `transaction()` and `changeVersion()` methods invoke callbacks with `SQLTransaction` objects.

```
typedef sequence<Object> ObjectArray;  
  
interface SQLTransaction {  
    void executeSql(in DOMString sqlStatement);  
    void executeSql(in DOMString sqlStatement, in ObjectArray arguments);  
    void executeSql(in DOMString sqlStatement, in ObjectArray arguments,  
in SQLStatementCallback callback);  
    void executeSql(in DOMString sqlStatement, in ObjectArray arguments,  
in SQLStatementCallback callback, in SQLStatementErrorCallback  
errorCallback);  
};  
  
interface SQLStatementCallback {  
    void handleEvent(in SQLTransaction transaction, in SQLResultSet  
resultSet);  
};  
  
interface SQLStatementErrorCallback {  
    boolean handleEvent(in SQLTransaction transaction, in SQLError error);  
};
```

When the `executeSql(sqlStatement, arguments, callback, errorCallback)` method is invoked, the user agent must run the following algorithm. (This algorithm is relatively simple and doesn't actually execute any SQL — the bulk of the work is actually done as part of the transaction steps (page 493).)

1. If the method was not invoked during the execution of a `SQLTransactionCallback`, `SQLStatementCallback`, or `SQLStatementErrorCallback` then raise an `INVALID_STATE_ERR` exception. (Calls from inside a `SQLTransactionErrorCallback` thus raise an exception. The `SQLTransactionErrorCallback` handler is only called once a transaction has failed, and no SQL statements can be added to a failed transaction.)
2. Parse the first argument to the method (`sqlStatement`) as an SQL statement, with the exception that ? characters can be used in place of literals in the statement. [SQL]
3. Replace each ? placeholder with the value of the argument in the `arguments` array with the same position. (So the first ? placeholder gets replaced by the first value in the `arguments` array, and generally the *n*th ? placeholder gets replaced by the *n*th value in the `arguments` array.)

If the second argument is omitted or null, then treat the `arguments` array as empty.

The result is *the statement*.

**

Implementation feedback is requested on what to do with arguments that are of types that are not supported by the underlying SQL backend. For example, SQLite

**

doesn't support booleans, so what should the UA do if passed a boolean? The Gears team suggests failing, not silently converting types.

4. If the syntax of *sqlStatement* is not valid (except for the use of ? characters in the place of literals), or the statement uses features that are not supported (e.g. due to security reasons), or the number of items in the *arguments* array is not equal to the number of ? placeholders in the statement, or the statement cannot be parsed for some other reason, then mark *the statement* as bogus.
5. If the Database object that the SQLTransaction object was created from has an expected version (page 489) that is neither the empty string nor the actual version of the database, then mark *the statement* as bogus. (Error code 2 (page 492).)
6. Queue up *the statement* in the transaction, along with the third argument (if any) as the statement's result set callback and the fourth argument (if any) as the error callback.

The user agent must act as if the database was hosted in an otherwise completely empty environment with no resources. For example, attempts to read from or write to the file system will fail.

SQL inherently supports multiple concurrent connections. Authors should make appropriate use of the transaction features to handle the case of multiple scripts interacting with the same database simultaneously (as could happen if the same page was opened in two different browsing contexts (page 414)).

User agents must consider statements that use the BEGIN, COMMIT, and ROLLBACK SQL features as being unsupported (and thus will mark them as bogus), so as to not let these statements interfere with the explicit transactions managed by the database API itself.

Note: A future version of this specification will probably define the exact SQL subset required in more detail.

5.10.2.4 Database query results

The executeSql() method invokes its callback with a SQLResultSet object as an argument.

```
interface SQLResultSet {  
    readonly attribute long insertId;  
    readonly attribute long rowsAffected;  
    readonly attribute SQLResultSetRowList rows;  
};
```

The **insertId** attribute must return the row ID of the row that the SQLResultSet object's SQL statement inserted into the database, if the statement inserted a row. If the statement inserted multiple rows, the ID of the last row must be the one returned. If the statement did not insert a row, then the attribute must instead raise an INVALID_ACCESS_ERR exception.

The **rowsAffected** attribute must return the number of rows that were affected by the SQL statement. If the statement did not affect any rows, then the attribute must return zero. For "SELECT" statements, this returns zero (querying the database doesn't affect any rows).

The **rows** attribute must return a `SQLResultSetRowList` representing the rows returned, in the order returned by the database. If no rows were returned, then the object will be empty (its length will be zero).

```
interface SQLResultSetRowList {
    readonly attribute unsigned long length;
    [IndexGetter] DOMObject item(in unsigned long index);
};
```

`SQLResultSetRowList` objects have a **length** attribute that must return the number of rows it represents (the number of rows returned by the database).

The **item(index)** attribute must return the row with the given index *index*. If there is no such row, then the method must raise an `INDEX_SIZE_ERR` exception.

Each row must be represented by a native ordered dictionary data type. In the ECMAScript binding, this must be `Object`. Each row object must have one property (or dictionary entry) per column, with those properties enumerating in the order that these columns were returned by the database. Each property must have the name of the column and the value of the cell, as they were returned by the database.

5.10.2.5 Errors

Errors in the database API are reported using callbacks that have a `SQLError` object as one of their arguments.

```
interface SQLError {
    readonly attribute unsigned long code;
    readonly attribute DOMString message;
};
```

The **code** DOM attribute must return the most appropriate code from the following table:

Code	Situation
0	The transaction failed for reasons unrelated to the database itself and not covered by any other error code.
1	The statement failed for database reasons not covered by any other error code.
2	The statement failed because the expected version (page 489) of the database didn't match the actual database version.
3	The statement failed because the data returned from the database was too large. The SQL "LIMIT" modifier might be useful to reduce the size of the result set.
4	The statement failed because there was not enough remaining storage space, or the storage quota was reached and the user declined to give more space to the database.
5	The statement failed because the transaction's first statement was a read-only statement, and a subsequent statement in the same transaction tried to modify the database, but the transaction failed to obtain a write lock before another transaction obtained a write lock and changed a part of the database that the former transaction was depending upon.
6	An <code>INSERT</code> , <code>UPDATE</code> , or <code>REPLACE</code> statement failed due to a constraint failure. For example, because a row was being inserted and the value given for the primary key column duplicated the value of an existing row.

- ** We should define a more thorough list of codes. Implementation feedback is requested to determine what codes are needed.

The **message** DOM attribute must return an error message describing the error encountered. The message should be localized to the user's language.

5.10.2.6 Processing model

The **transaction steps** are as follows. These steps must be run asynchronously. These steps are invoked with a *transaction callback*, optionally an *error callback*, optionally a *success callback*, optionally a *preflight operation*, and optionally a *postflight operation*.

1. Open a new SQL transaction to the database, and create a `SQLTransaction` object that represents that transaction.
2. If an error occurred in the opening of the transaction, jump to the last step.
3. If a *preflight operation* was defined for this instance of the transaction steps, run that. If it fails, then jump to the last step. (This is basically a hook for the `changeVersion()` method.)
4. Queue a task (page 429) to invoke the *transaction callback* with the aforementioned `SQLTransaction` object as its only argument, and wait for that task to be run.
5. If the callback couldn't be called (e.g. it was null), or if the callback was invoked and raised an exception, jump to the last step.
6. While there are any statements queued up in the transaction, perform the following steps for each queued up statement in the transaction, oldest first. Each statement has a statement, optionally a result set callback, and optionally an error callback.
 1. If the statement is marked as bogus, jump to the "in case of error" steps below.
 2. Execute the statement in the context of the transaction. [SQL]
 3. If the statement failed, jump to the "in case of error" steps below.
 4. Create a `SQLResultSet` object that represents the result of the statement.
 5. If the statement has a result set callback, queue a task (page 429) to invoke it with the `SQLTransaction` object as its first argument and the new `SQLResultSet` object as its second argument, and wait for that task to be run.
 6. If the callback was invoked and raised an exception, jump to the last step in the overall steps.
 7. Move on to the next statement, if any, or onto the next overall step otherwise.

In case of error (or more specifically, if the above substeps say to jump to the "in case of error" steps), run the following substeps:

1. If the statement had an associated error callback, then queue a task (page 429) to invoke that error callback with the `SQLTransaction` object and a newly constructed `SQLError` object that represents the error that caused these substeps to be run as the two arguments, respectively, and wait for the task to be run.
2. If the error callback returns false, then move on to the next statement, if any, or onto the next overall step otherwise.
3. Otherwise, the error callback did not return false, or there was no error callback. Jump to the last step in the overall steps.
7. If a *postflight operation* was defined for this instance of the transaction steps, run that. If it fails, then jump to the last step. (This is basically a hook for the `changeVersion()` method.)
8. Commit the transaction.
9. If an error occurred in the committing of the transaction, jump to the last step.
10. Queue a task (page 429) to invoke the *success callback*.
11. End these steps. The next step is only used when something goes wrong.
12. Queue a task (page 429) to invoke the *error callback* with a newly constructed `SQLError` object that represents the last error to have occurred in this transaction. Rollback the transaction. Any still-pending statements in the transaction are discarded.

5.10.3 Disk space

User agents should limit the total amount of space allowed for storage areas and databases.

User agents should guard against sites storing data in the storage areas or databases of subdomains, e.g. storing up to the limit in `a1.example.com`, `a2.example.com`, `a3.example.com`, etc, circumventing the main `example.com` storage limit.

User agents may prompt the user when quotas are reached, allowing the user to grant a site more space. This enables sites to store many user-created documents on the user's computer, for instance.

User agents should allow users to see how much space each domain is using.

A mostly arbitrary limit of five megabytes per domain is recommended. Implementation feedback is welcome and will be used to update this suggestion in future.

5.10.4 Privacy

5.10.4.1 User tracking

A third-party advertiser (or any entity capable of getting content distributed to multiple sites) could use a unique identifier stored in its local storage area or in its client-side database to track a user across multiple sessions, building a profile of the user's interests to allow for

highly targeted advertising. In conjunction with a site that is aware of the user's real identity (for example an e-commerce site that requires authenticated credentials), this could allow oppressive groups to target individuals with greater accuracy than in a world with purely anonymous Web usage.

There are a number of techniques that can be used to mitigate the risk of user tracking:

- Blocking third-party storage: user agents may restrict access to the `localStorage` and database objects to scripts originating at the domain of the top-level document of the browsing context (page 414), for instance denying access to the API for pages from other domains running in iframes.
- Expiring stored data: user agents may automatically delete stored data after a period of time.

For example, a user agent could treat third-party local storage areas as session-only storage, deleting the data once the user had closed all the browsing contexts that could access it.

This can restrict the ability of a site to track a user, as the site would then only be able to track the user across multiple sessions when he authenticates with the site itself (e.g. by making a purchase or logging in to a service).

However, this also puts the user's data at risk.

- Treating persistent storage as cookies: user agents should present the persistent storage and database features to the user in a way that does not distinguish them from HTTP session cookies. [RFC2109] [RFC2965]

This might encourage users to view persistent storage with healthy suspicion.

- Site-specific white-listing of access to local storage areas and databases: user agents may allow sites to access session storage areas in an unrestricted manner, but require the user to authorize access to local storage areas and databases.
- Origin (page 423)-tracking of persistent storage data: user agents may record the origins of sites that contained content from third-party origins that caused data to be stored.

If this information is then used to present the view of data currently in persistent storage, it would allow the user to make informed decisions about which parts of the persistent storage to prune. Combined with a blacklist ("delete this data and prevent this domain from ever storing data again"), the user can restrict the use of persistent storage to sites that he trusts.

- Shared blacklists: user agents may allow users to share their persistent storage domain blacklists.

This would allow communities to act together to protect their privacy.

While these suggestions prevent trivial use of these APIs for user tracking, they do not block it altogether. Within a single domain, a site can continue to track the user during a session, and can then pass all this information to the third party along with any identifying information (names, credit card numbers, addresses) obtained by the site. If a third party cooperates with multiple sites to obtain such information, a profile can still be created.

However, user tracking is to some extent possible even with no cooperation from the user agent whatsoever, for instance by using session identifiers in URLs, a technique already commonly used for innocuous purposes but easily repurposed for user tracking (even retroactively). This information can then be shared with other sites, using visitors' IP addresses and other user-specific data (e.g. user-agent headers and configuration settings) to combine separate sessions into coherent user profiles.

5.10.4.2 Cookie resurrection

If the user interface for persistent storage presents data in the persistent storage features separately from data in HTTP session cookies, then users are likely to delete data in one and not the other. This would allow sites to use the two features as redundant backup for each other, defeating a user's attempts to protect his privacy.

5.10.5 Security

5.10.5.1 DNS spoofing attacks

Because of the potential for DNS spoofing attacks, one cannot guarantee that a host claiming to be in a certain domain really is from that domain. To mitigate this, pages can use SSL. Pages using SSL can be sure that only pages using SSL that have certificates identifying them as being from the same domain can access their local storage areas and databases.

5.10.5.2 Cross-directory attacks

Different authors sharing one host name, for example users hosting content on `geocities.com`, all share one persistent storage object and one set of databases. There is no feature to restrict the access by pathname. Authors on shared hosts are therefore recommended to avoid using the persistent storage features, as it would be trivial for other authors to read from and write to the same storage area or database.

Note: Even if a path-restriction feature was made available, the usual DOM scripting security model would make it trivial to bypass this protection and access the data from any path.

5.10.5.3 Implementation risks

The two primary risks when implementing these persistent storage features are letting hostile sites read information from other domains, and letting hostile sites write information that is then read from other domains.

Letting third-party sites read data that is not supposed to be read from their domain causes *information leakage*. For example, a user's shopping wishlist on one domain could be used by another domain for targeted advertising; or a user's work-in-progress confidential documents stored by a word-processing site could be examined by the site of a competing company.

Letting third-party sites write data to the storage areas of other domains can result in *information spoofing*, which is equally dangerous. For example, a hostile site could add items to a user's wishlist; or a hostile site could set a user's session identifier to a known ID that the hostile site can then use to track the user's actions on the victim site.

Thus, strictly following the origin (page 423) model described in this specification is important for user security.

5.10.5.4 SQL and user agents

User agent implementors are strongly encouraged to audit all their supported SQL statements for security implications. For example, `LOAD DATA INFILE` is likely to pose security risks and there is little reason to support it.

In general, it is recommended that user agents not support features that control how databases are stored on disk. For example, there is little reason to allow Web authors to control the character encoding used in the disk representation of the data, as all data in ECMAScript is implicitly UTF-16.

5.10.5.5 SQL injection

Authors are strongly recommended to make use of the `?` placeholder feature of the `executeSql()` method, and to never construct SQL statements on the fly.

5.11 Links

5.11.1 Hyperlink elements

The `a`, `area`, and `link` elements can, in certain situations described in the definitions of those elements, represent **hyperlinks**.

The `href` attribute on a hyperlink element must have a value that is a valid URL (page 52). This URL is the *destination resource* of the hyperlink.

The href attribute on a and area elements is not required; when those elements do not have href attributes they do not represent hyperlinks.

The href attribute on the link element is required, but whether a link element represents a hyperlink or not depends on the value of the rel attribute of that element.

The `target` attribute, if present, must be a valid browsing context name or keyword (page 418). User agents use this name when following hyperlinks (page 498).

The `ping` attribute, if present, gives the URLs of the resources that are interested in being notified if the user follows the hyperlink. The value must be a space separated list of one or more valid URLs. The value is used by the user agent when following hyperlinks (page 498).

For `a` and `area` elements that represent hyperlinks, the relationship between the document containing the hyperlink and the destination resource indicated by the hyperlink is given by the value of the element's `rel` attribute, which must be a set of space-separated tokens (page 49). The allowed values and their meanings (page 500) are defined below. The `rel` attribute has no default value. If the attribute is omitted or if none of the values in the attribute are recognized by the UA, then the document has no particular relationship with the destination resource other than there being a hyperlink between the two.

The **media** attribute describes for which media the target document was designed. It is purely advisory. The value must be a valid media query (page 30). [MQ] The default, if the **media** attribute is omitted, is all.

The **hreflang** attribute on hyperlink elements, if present, gives the language of the linked resource. It is purely advisory. The value must be a valid RFC 3066 language code.

[RFC3066] User agents must not consider this attribute authoritative — upon fetching the resource, user agents must use only language information associated with the resource to determine its language, not metadata included in the link to the resource.

The **type** attribute, if present, gives the MIME type of the linked resource. It is purely advisory. The value must be a valid MIME type, optionally with parameters. [RFC2046] User agents must not consider the **type** attribute authoritative — upon fetching the resource, user agents must not use metadata included in the link to the resource to determine its type.

5.11.2 Following hyperlinks

When a user *follows a hyperlink*, the user agent must navigate (page 473) a browsing context (page 414) to the URL (page 52) given by the **href** attribute of that hyperlink. In the case of server-side image maps, the URL of the hyperlink must further have its *hyperlink suffix* appended to it.

If the user indicated a specific browsing context (page 414) when following the hyperlink, or if the user agent is configured to follow hyperlinks by navigating a particular browsing context, then that must be the browsing context (page 414) that is navigated.

Otherwise, if the hyperlink element is an **a** or **area** element that has a **target** attribute, then the browsing context (page 414) that is navigated must be chosen by applying the rules for choosing a browsing context given a browsing context name (page 418), using the value of the **target** attribute as the browsing context name. If these rules result in the creation of a new browsing context (page 414), it must be navigated with replacement enabled (page 476).

Otherwise, if the hyperlink element is a sidebar hyperlink (page 508) and the user agent implements a feature that can be considered a secondary browsing context, such a secondary browsing context may be selected as the browsing context to be navigated.

Otherwise, if the hyperlink element is an **a** or **area** element with no **target** attribute, but one of the child nodes of the **head** element (page 80) is a **base** element with a **target** attribute, then the browsing context that is navigated must be chosen by applying the rules for choosing a browsing context given a browsing context name (page 418), using the value of the **target** attribute of the first such **base** element as the browsing context name. If these rules result in the creation of a new browsing context (page 414), it must be navigated with replacement enabled (page 476).

Otherwise, the browsing context that must be navigated is the same browsing context as the one which the hyperlink element itself is in.

The navigation must be done with the browsing context (page 414) that contains the **Document** object with which the hyperlink's element in question is associated as the source browsing context (page 473).

5.11.2.1 Hyperlink auditing

If an a or area hyperlink element has a ping attribute and the user follows the hyperlink, the user agent must take the ping attribute's value, split that string on spaces, resolve (page 55) each resulting token, and then should send a request (as described below) to each of the resulting absolute URLs (page 56). (Tokens that fail to resolve are ignored.) This may be done in parallel with the primary request, and is independent of the result of that request.

User agents should allow the user to adjust this behavior, for example in conjunction with a setting that disables the sending of HTTP Referer headers. Based on the user's preferences, UAs may either ignore (page 24) the ping attribute altogether, or selectively ignore URLs in the list (e.g. ignoring any third-party URLs).

For URLs that are HTTP URLs, the requests must be performed by fetching (page 59) the specified URLs using the POST method, with an empty entity body in the request. All relevant cookie and HTTP authentication headers must be included in the request. Which other headers are required depends on the URLs involved.

↪ **If both the address of the Document object containing the hyperlink being audited and the ping URL have the same origin (page 426)**

The request must include a Ping-From HTTP header with, as its value, the address of the document containing the hyperlink, and a Ping-To HTTP header with, as its value, the address of the absolute URL (page 56) of the target of the hyperlink. The request must not include a Referer HTTP header.

↪ **Otherwise, if the origins are different, but the document containing the hyperlink being audited was not retrieved over an encrypted connection**

The request must include a Referer HTTP header [sic] with, as its value, the location of the document containing the hyperlink, a Ping-From HTTP header with the same value, and a Ping-To HTTP header with, as its value, the address of the target of the hyperlink.

↪ **Otherwise, the origins are different and the document containing the hyperlink being audited was retrieved over an encrypted connection**

The request must include a Ping-To HTTP header with, as its value, the address of the target of the hyperlink. The request must neither include a Referer HTTP header nor include a Ping-From HTTP header.

Note: To save bandwidth, implementors might also wish to consider omitting optional headers such as Accept from these requests.

User agents must, unless otherwise specified by the user, honor the HTTP headers (including, in particular, redirects and HTTP cookie headers), but must ignore any entity bodies returned in the responses. User agents may close the connection prematurely once they start receiving an entity body. [RFC2109] [RFC2965]

For URLs that are not HTTP URLs, the requests must be performed by fetching (page 59) the specified URL normally, and discarding the results.

When the ping attribute is present, user agents should clearly indicate to the user that following the hyperlink will also cause secondary requests to be sent in the background, possibly including listing the actual target URLs.

The `ping` attribute is redundant with pre-existing technologies like HTTP redirects and JavaScript in allowing Web pages to track which off-site links are most popular or allowing advertisers to track click-through rates.

However, the `ping` attribute provides these advantages to the user over those alternatives:

- **It allows the user to see the final target URL unobscured.**
- **It allows the UA to inform the user about the out-of-band notifications.**
- **It allows the paranoid user to disable the notifications without losing the underlying link functionality.**
- **It allows the UA to optimize the use of available network bandwidth so that the target page loads faster.**

Thus, while it is possible to track users without this feature, authors are encouraged to use the `ping` attribute so that the user agent can improve the user experience.

5.11.3 Link types

The following table summarizes the link types that are defined by this specification. This table is non-normative; the actual definitions for the link types are given in the next few sections.

In this section, the term *referenced document* refers to the resource identified by the element representing the link, and the term *current document* refers to the resource within which the element representing the link finds itself.

To determine which link types apply to a link, a, or area element, the element's rel attribute must be split on spaces (page 49). The resulting tokens are the link types that apply to that element.

Unless otherwise specified, a keyword must not be specified more than once per rel attribute.

Link type	Effect on...		Brief description
	link	a and area	
alternate	Hyperlink (page 111)	Hyperlink (page 497)	Gives alternate representations of the current document.
archives	Hyperlink (page 111)	Hyperlink (page 497)	Provides a link to a collection of records, documents, or other materials of historical interest.
author	Hyperlink (page 111)	Hyperlink (page 497)	Gives a link to the current document's author.
bookmark	<i>not allowed</i>	Hyperlink (page 497)	Gives the permalink for the nearest ancestor section.

Link type	Effect on...		Brief description
	link	a and area	
external	<i>not allowed</i>	Hyperlink (page 497)	Indicates that the referenced document is not part of the same site as the current document.
feed	Hyperlink (page 111)	Hyperlink (page 497)	Gives the address of a syndication feed for the current document.
first	Hyperlink (page 111)	Hyperlink (page 497)	Indicates that the current document is a part of a series, and that the first document in the series is the referenced document.
help	Hyperlink (page 111)	Hyperlink (page 497)	Provides a link to context-sensitive help.
icon	External Resource (page 111)	<i>not allowed</i>	Imports an icon to represent the current document.
index	Hyperlink (page 111)	Hyperlink (page 497)	Gives a link to the document that provides a table of contents or index listing the current document.
last	Hyperlink (page 111)	Hyperlink (page 497)	Indicates that the current document is a part of a series, and that the last document in the series is the referenced document.
license	Hyperlink (page 111)	Hyperlink (page 497)	Indicates that the current document is covered by the copyright license described by the referenced document.
next	Hyperlink (page 111)	Hyperlink (page 497)	Indicates that the current document is a part of a series, and that the next document in the series is the referenced document.
nofollow	<i>not allowed</i>	Hyperlink (page 497)	Indicates that the current document's original author or publisher does not endorse the referenced document.
noreferrer	<i>not allowed</i>	Hyperlink (page 497)	Requires that the user agent not send an HTTP Referer header if the user follows the hyperlink.
pingback	External Resource (page 111)	<i>not allowed</i>	Gives the address of the pingback server that handles pingbacks to the current document.
prefetch	External Resource (page 111)	<i>not allowed</i>	Specifies that the target resource should be preemptively cached.
prev	Hyperlink (page 111)	Hyperlink (page 497)	Indicates that the current document is a part of a series, and that the previous document in the series is the referenced document.
search	Hyperlink (page 111)	Hyperlink (page 497)	Gives a link to a resource that can be used to search through the current document and its related pages.
stylesheet	External Resource (page 111)	<i>not allowed</i>	Imports a stylesheet.
sidebar	Hyperlink (page 111)	Hyperlink (page 497)	Specifies that the referenced document, if retrieved, is intended to be shown in the browser's sidebar (if it has one).

Link type	Effect on...		Brief description
	link	a and area	
tag	Hyperlink (page 111)	Hyperlink (page 497)	Gives a tag (identified by the given address) that applies to the current document.
up	Hyperlink (page 111)	Hyperlink (page 497)	Provides a link to a document giving the context for the current document.

Some of the types described below list synonyms for these values. These are to be handled as specified by user agents, but must not be used in documents.

5.11.3.1 Link type "alternate"

The alternate keyword may be used with link, a, and area elements. For link elements, if the rel attribute does not also contain the keyword stylesheet, it creates a hyperlink (page 111); but if it *does* also contain the keyword stylesheet, the alternate keyword instead modifies the meaning of the stylesheet keyword in the way described for that keyword, and the rest of this subsection doesn't apply.

The alternate keyword indicates that the referenced document is an alternate representation of the current document.

The nature of the referenced document is given by the media, hreflang, and type attributes.

If the alternate keyword is used with the media attribute, it indicates that the referenced document is intended for use with the media specified.

If the alternate keyword is used with the hreflang attribute, and that attribute's value differs from the root element (page 24)'s language (page 88), it indicates that the referenced document is a translation.

If the alternate keyword is used with the type attribute, it indicates that the referenced document is a reformulation of the current document in the specified format.

The media, hreflang, and type attributes can be combined when specified with the alternate keyword.

For example, the following link is a French translation that uses the PDF format:

```
<link rel=alternate type=application/pdf hreflang=fr href=manual-fr>
```

If the alternate keyword is used with the type attribute set to the value application/rss+xml or the value application/atom+xml, then the user agent must treat the link as it would if it had the feed keyword specified as well.

The alternate link relationship is transitive — that is, if a document links to two other documents with the link type "alternate", then, in addition to implying that those documents are alternative representations of the first document, it is also implying that those two documents are alternative representations of each other.

5.11.3.2 Link type "archives"

The archives keyword may be used with link, a, and area elements. For link elements, it creates a hyperlink (page 111).

The archives keyword indicates that the referenced document describes a collection of records, documents, or other materials of historical interest.

|| A blog's index page could link to an index of the blog's past posts with rel="archives".

Synonyms: For historical reasons, user agents must also treat the keyword "archive" like the archives keyword.

5.11.3.3 Link type "author"

The author keyword may be used with link, a, and area elements. For link elements, it creates a hyperlink (page 111).

For a and area elements, the author keyword indicates that the referenced document provides further information about the author of the section that the element defining the hyperlink applies (page 132) to.

For link elements, the author keyword indicates that the referenced document provides further information about the author for the page as a whole.

Note: *The "referenced document" can be, and often is, a mailto: URL giving the e-mail address of the author. [MAILTO]*

Synonyms: For historical reasons, user agents must also treat link, a, and area elements that have a rev attribute with the value "made" as having the author keyword specified as a link relationship.

5.11.3.4 Link type "bookmark"

The bookmark keyword may be used with a and area elements.

The bookmark keyword gives a permalink for the nearest ancestor article element of the linking element in question, or of the section the linking element is most closely associated with (page 144), if there are no ancestor article elements.

|| The following snippet has three permalinks. A user agent could determine which permalink applies to which part of the spec by looking at where the permalinks are given.

```
...
<body>
  <h1>Example of permalinks</h1>
  <div id="a">
    <h2>First example</h2>
    <p><a href="a.html" rel="bookmark">This</a> permalink applies to
       only the content from the first H2 to the second H2. The DIV isn't
       exactly that section, but it roughly corresponds to it.</p>
  </div>
```

```

<h2>Second example</h2>
<article id="b">
  <p><a href="b.html" rel="bookmark">This</a> permalink applies to
  the outer ARTICLE element (which could be, e.g., a blog post).</p>
  <article id="c">
    <p><a href="c.html" rel="bookmark">This</a> permalink applies to
    the inner ARTICLE element (which could be, e.g., a blog
comment).</p>
  </article>
</article>
</body>
...

```

5.11.3.5 Link type "external"

The external keyword may be used with a and area elements.

The external keyword indicates that the link is leading to a document that is not part of the site that the current document forms a part of.

5.11.3.6 Link type "feed"

The feed keyword may be used with link, a, and area elements. For link elements, it creates a hyperlink (page 111).

The feed keyword indicates that the referenced document is a syndication feed. If the alternate link type is also specified, then the feed is specifically the feed for the current document; otherwise, the feed is just a syndication feed, not necessarily associated with a particular Web page.

The first link, a, or area element in the document (in tree order) that creates a hyperlink with the link type feed must be treated as the default syndication feed for the purposes of feed autodiscovery.

Note: *The feed keyword is implied by the alternate link type in certain cases (q.v.).*

The following two link elements are equivalent: both give the syndication feed for the current page:

```

<link rel="alternate" type="application/atom+xml" href="data.xml">
<link rel="feed alternate" href="data.xml">

```

The following extract offers various different syndication feeds:

```

<p>You can access the planets database using Atom feeds:</p>
<ul>
  <li><a href="recently-visited-planets.xml" rel="feed">Recently
  Visited Planets</a></li>
  <li><a href="known-bad-planets.xml" rel="feed">Known Bad
  Planets</a></li>
  <li><a href="unexplored-planets.xml" rel="feed">Unexplored

```

```
    Planets</a></li>
  </ul>
```

5.11.3.7 Link type "help"

The help keyword may be used with link, a, and area elements. For link elements, it creates a hyperlink (page 111).

For a and area elements, the help keyword indicates that the referenced document provides further help information for the parent of the element defining the hyperlink, and its children.

In the following example, the form control has associated context-sensitive help. The user agent could use this information, for example, displaying the referenced document if the user presses the "Help" or "F1" key.

```
<p><label> Topic: <input name=topic> <a href="help/topic.html"
rel="help">(Help)</a></label></p>
```

For link elements, the help keyword indicates that the referenced document provides help for the page as a whole.

5.11.3.8 Link type "icon"

The icon keyword may be used with link elements, for which it creates an external resource link (page 111).

The specified resource is an icon representing the page or site, and should be used by the user agent when representing the page in the user interface.

Icons could be auditory icons, visual icons, or other kinds of icons. If multiple icons are provided, the user agent must select the most appropriate icon according to the type, media, and sizes attributes. If there are multiple equally appropriate icons, user agents must use the last one declared in tree order (page 24). If the user agent tries to use an icon but that icon is determined, upon closer examination, to in fact be inappropriate (e.g. because it uses an unsupported format), then the user agent must try the next-most-appropriate icon as determined by the attributes.

There is no default type for resources given by the icon keyword. However, for the purposes of determining the type of the resource (page 112), user agents must expect the resource to be an image.

The **sizes** attribute gives the sizes of icons for visual media.

If specified, the attribute must have a value that is an unordered set of unique space-separated tokens (page 49). The values must all be either any or a value that consists of two valid non-negative integers (page 33) that do not have a leading U+0030 DIGIT ZERO (0) character and that are separated by a single U+0078 LATIN SMALL LETTER X character.

The keywords represent icon sizes.

To parse and process the attribute's value, the user agent must first split the attribute's value on spaces (page 49), and must then parse each resulting keyword to determine what it represents.

The **any** keyword represents that the resource contains a scalable icon, e.g. as provided by an SVG image.

Other keywords must be further parsed as follows to determine what they represent:

- If the keyword doesn't contain exactly one U+0078 LATIN SMALL LETTER X character, then this keyword doesn't represent anything. Abort these steps for that keyword.
- Let *width string* be the string before the "x".
- Let *height string* be the string after the "x".
- If either *width string* or *height string* start with a U+0030 DIGIT ZERO (0) character or contain any characters other than characters in the range U+0030 DIGIT ZERO (0) to U+0039 DIGIT NINE (9), then this keyword doesn't represent anything. Abort these steps for that keyword.
- Apply the rules for parsing non-negative integers (page 33) to *width string* to obtain *width*.
- Apply the rules for parsing non-negative integers (page 33) to *height string* to obtain *height*.
- The keyword represents that the resource contains a bitmap icon with a width of *width* device pixels and a height of *height* device pixels.

The keywords specified on the `sizes` attribute must not represent icon sizes that are not actually available in the linked resource.

If the attribute is not specified, then the user agent must assume that the given icon is appropriate, but less appropriate than an icon of a known and appropriate size.

The following snippet shows the top part of an application with several icons.

```
<!DOCTYPE HTML>
<html>
  <head>
    <title>lsForums – Inbox</title>
    <link rel=icon href=favicon.png sizes="16x16">
    <link rel=icon href=windows.ico sizes="32x32 48x48">
    <link rel=icon href=mac.icns sizes="128x128 512x512 8192x8192
32768x32768">
    <link rel=icon href=iphone.png sizes="59x60">
    <link rel=icon href=gnome.svg sizes="any">
    <link rel=stylesheet href=lsforums.css>
    <script src=lsforums.js></script>
    <meta name=application-name content="lsForums">
  </head>
  <body>
    ...
  </body>
```

5.11.3.9 Link type "license"

The license keyword may be used with link, a, and area elements. For link elements, it creates a hyperlink (page 111).

The license keyword indicates that the referenced document provides the copyright license terms under which the current document is provided.

Synonyms: For historical reasons, user agents must also treat the keyword "copyright" like the license keyword.

5.11.3.10 Link type "nofollow"

Thenofollow keyword may be used with a and area elements.

Thenofollow keyword indicates that the link is not endorsed by the original author or publisher of the page, or that the link to the referenced document was included primarily because of a commercial relationship between people affiliated with the two pages.

5.11.3.11 Link type "noreferrer"

The noreferrer keyword may be used with a and area elements.

If a user agent follows a link defined by an a or area element that has the noreferrer keyword, the user agent must not include a Referer HTTP header (or equivalent for other protocols) in the request.

This keyword also causes the opener attribute to remain null (page 418) if the hyperlink creates a new browsing context (page 414).

5.11.3.12 Link type "pingback"

The pingback keyword may be used with link elements, for which it creates an external resource link (page 111).

For the semantics of the pingback keyword, see the Pingback 1.0 specification. [PINGBACK]

5.11.3.13 Link type "prefetch"

The prefetch keyword may be used with link elements, for which it creates an external resource link (page 111).

The prefetch keyword indicates that preemptively fetching and caching the specified resource is likely to be beneficial, as it is highly likely that the user will require this resource.

There is no default type for resources given by the prefetch keyword.

5.11.3.14 Link type "search"

The search keyword may be used with link, a, and area elements. For link elements, it creates a hyperlink (page 111).

The search keyword indicates that the referenced document provides an interface specifically for searching the document and its related resources.

Note: OpenSearch description documents can be used with link elements and the search link type to enable user agents to autodiscover search interfaces. [OPENSEARCH]

5.11.3.15 Link type "stylesheet"

The stylesheet keyword may be used with link elements, for which it creates an external resource link (page 111) that contributes to the styling processing model (page 122).

The specified resource is a resource that describes how to present the document. Exactly how the resource is to be processed depends on the actual type of the resource.

If the alternate keyword is also specified on the link element, then the link is an alternative stylesheet; in this case, the title attribute must be specified on the link element, with a non-empty value.

The default type for resources given by the stylesheet keyword is text/css.

Quirk: If the document has been set to quirks mode (page 79) and the Content-Type metadata (page 60) of the external resource is not a supported style sheet type, the user agent must instead assume it to be text/css.

5.11.3.16 Link type "sidebar"

The sidebar keyword may be used with link, a, and area elements. For link elements, it creates a hyperlink (page 111).

The sidebar keyword indicates that the referenced document, if retrieved, is intended to be shown in a secondary browsing context (page 416) (if possible), instead of in the current browsing context (page 414).

A hyperlink element (page 497) with the sidebar keyword specified is a **sidebar hyperlink**.

5.11.3.17 Link type "tag"

The tag keyword may be used with link, a, and area elements. For link elements, it creates a hyperlink (page 111).

The tag keyword indicates that the tag that the referenced document represents applies to the current document.

5.11.3.18 Hierarchical link types

Some documents form part of a hierarchical structure of documents.

A hierarchical structure of documents is one where each document can have various subdocuments. The document of which a document is a subdocument is said to be the document's *parent*. A document with no parent forms the top of the hierarchy.

A document may be part of multiple hierarchies.

5.11.3.18.1 Link type "index"

The index keyword may be used with link, a, and area elements. For link elements, it creates a hyperlink (page 111).

The index keyword indicates that the document is part of a hierarchical structure, and that the link is leading to the document that is the top of the hierarchy. It conveys more information when used with the up keyword (q.v.).

Synonyms: For historical reasons, user agents must also treat the keywords "top", "contents", and "toc" like the index keyword.

5.11.3.18.2 Link type "up"

The up keyword may be used with link, a, and area elements. For link elements, it creates a hyperlink (page 111).

The up keyword indicates that the document is part of a hierarchical structure, and that the link is leading to the document that is the parent of the current document.

The up keyword may be repeated within a rel attribute to indicate the hierarchical distance from the current document to the referenced document. Each occurrence of the keyword represents one further level. If the index keyword is also present, then the number of up keywords is the depth of the current page relative to the top of the hierarchy. Only one link is created for the set of one or more up keywords and, if present, the index keyword.

If the page is part of multiple hierarchies, then they should be described in different paragraphs (page 96). User agents must scope any interpretation of the up and index keywords together indicating the depth of the hierarchy to the paragraph (page 96) in which the link finds itself, if any, or to the document otherwise.

When two links have both the up and index keywords specified together in the same scope and contradict each other by having a different number of up keywords, the link with the greater number of up keywords must be taken as giving the depth of the document.

This can be used to mark up a navigation style sometimes known as bread crumbs. In the following example, the current page can be reached via two paths.

```
<nav>
  <p>
    <a href="/" rel="index up up up">Main</a> >
    <a href="/products/" rel="up up">Products</a> >
    <a href="/products/dishwashers/" rel="up">Dishwashers</a> >
    <a>Second hand</a>
  </p>
  <p>
    <a href="/" rel="index up up">Main</a> >
```

```
<a href="/second-hand/" rel="up">Second hand</a> >
<a>Dishwashers</a>
</p>
</nav>
```

Note: The `relList` DOM attribute (e.g. on the `a` element) does not currently represent multiple `up` keywords (the interface hides duplicates).

5.11.3.19 Sequential link types

Some documents form part of a sequence of documents.

A sequence of documents is one where each document can have a *previous sibling* and a *next sibling*. A document with no previous sibling is the start of its sequence, a document with no next sibling is the end of its sequence.

A document may be part of multiple sequences.

5.11.3.19.1 Link type "first"

The `first` keyword may be used with `link`, `a`, and `area` elements. For `link` elements, it creates a hyperlink (page 111).

The `first` keyword indicates that the document is part of a sequence, and that the link is leading to the document that is the first logical document in the sequence.

Synonyms: For historical reasons, user agents must also treat the keywords "begin" and "start" like the `first` keyword.

5.11.3.19.2 Link type "last"

The `last` keyword may be used with `link`, `a`, and `area` elements. For `link` elements, it creates a hyperlink (page 111).

The `last` keyword indicates that the document is part of a sequence, and that the link is leading to the document that is the last logical document in the sequence.

Synonyms: For historical reasons, user agents must also treat the keyword "end" like the `last` keyword.

5.11.3.19.3 Link type "next"

The `next` keyword may be used with `link`, `a`, and `area` elements. For `link` elements, it creates a hyperlink (page 111).

The `next` keyword indicates that the document is part of a sequence, and that the link is leading to the document that is the next logical document in the sequence.

5.11.3.19.4 Link type "prev"

The prev keyword may be used with link, a, and area elements. For link elements, it creates a hyperlink (page 111).

The prev keyword indicates that the document is part of a sequence, and that the link is leading to the document that is the previous logical document in the sequence.

Synonyms: For historical reasons, user agents must also treat the keyword "previous" like the prev keyword.

5.11.3.20 Other link types

Other than the types defined above, only types defined as extensions in the WHATWG Wiki RelExtensions page may be used with the rel attribute on link, a, and area elements.
[WHATWGWiki]

Anyone is free to edit the WHATWG Wiki RelExtensions page at any time to add a type. Extension types must be specified with the following information:

Keyword

The actual value being defined. The value should not be confusingly similar to any other defined value (e.g. differing only in case).

Effect on... link

One of the following:

not allowed

The keyword is not allowed to be specified on link elements.

Hyperlink

The keyword may be specified on a link element; it creates a hyperlink link (page 111).

External Resource

The keyword may be specified on a link element; it creates a external resource link (page 111).

Effect on... a and area

One of the following:

not allowed

The keyword is not allowed to be specified on a and area elements.

Hyperlink

The keyword may be specified on a and area elements.

Brief description

A short description of what the keyword's meaning is.

Link to more details

A link to a more detailed description of the keyword's semantics and requirements. It could be another page on the Wiki, or a link to an external page.

Synonyms

A list of other keyword values that have exactly the same processing requirements. Authors must not use the values defined to be synonyms, they are only intended to allow user agents to support legacy content.

Status

One of the following:

Proposal

The keyword has not received wide peer review and approval. It is included for completeness because pages use the keyword. Pages should not use the keyword.

Accepted

The keyword has received wide peer review and approval. It has a specification that unambiguously defines how to handle pages that use the keyword, including when they use them in incorrect ways. Pages may use the keyword.

Rejected

The keyword has received wide peer review and it has been found to have significant problems. Pages must not use the keyword. When a keyword has this status, the "Effect on... link" and "Effect on... a and area" information should be set to "not allowed".

If a keyword is added with the "proposal" status and found to be redundant with existing values, it should be removed and listed as a synonym for the existing value. If a keyword is added with the "proposal" status and found to be harmful, then it should be changed to "rejected" status, and its "Effect on..." information should be changed accordingly.

Conformance checkers must use the information given on the WHATWG Wiki RelExtensions page to establish if a value not explicitly defined in this specification is allowed or not. When an author uses a new type not defined by either this specification or the Wiki page, conformance checkers should offer to add the value to the Wiki, with the details described above, with the "proposal" status.

This specification does not define how new values will get approved. It is expected that the Wiki will have a community that addresses this.

6 User Interaction

This section describes various features that allow authors to enable users to edit documents and parts of documents interactively.

6.1 Introduction

This section is non-normative.

** Would be nice to explain how these features work together.

6.2 The hidden attribute

All elements may have the hidden content attribute set. The hidden attribute is a boolean attribute (page 32). When specified on an element, it indicates that the element is not yet, or is no longer, relevant. User agents should not render elements that have the hidden attribute specified.

In the following skeletal example, the attribute is used to hide the Web game's main screen until the user logs in:

```
<h1>The Example Game</h1>
<section id="login">
  <h2>Login</h2>
  <form>
    ...
    <!-- calls login() once the user's credentials have been checked
    -->
  </form>
  <script>
    function login() {
      // switch screens
      document.getElementById('login').hidden = true;
      document.getElementById('game').hidden = false;
    }
  </script>
</section>
<section id="game" hidden>
  ...
</section>
```

The hidden attribute must not be used to hide content that could legitimately be shown in another presentation. For example, it is incorrect to use hidden to hide panels in a tabbed dialog, because the tabbed interface is merely a kind of overflow presentation — showing all the form controls in one big page with a scrollbar would be equivalent, and no less correct.

Elements in a section hidden by the hidden attribute are still active, e.g. scripts and form controls in such sections still render execute and submit respectively. Only their presentation to the user changes.

The **hidden** DOM attribute must reflect (page 67) the content attribute of the same name.

6.3 Activation

The **click()** method must, if the element has a defined activation behavior (page 96), run synthetic click activation steps (page 95) on the element. Otherwise, the user agent must fire a click event (page 436) at the element.

6.4 Scrolling elements into view

The **scrollIntoView([top])** method, when called, must cause the element on which the method was called to have the attention of the user called to it.

Note: *In a speech browser, this could happen by having the current playback position move to the start of the given element.*

In visual user agents, if the argument is present and has the value false, the user agent should scroll the element into view such that both the bottom and the top of the element are in the viewport, with the bottom of the element aligned with the bottom of the viewport. If it isn't possible to show the entire element in that way, or if the argument is omitted or is true, then the user agent should instead align the top of the element with the top of the viewport. If the entire scrollable part of the content is visible all at once (e.g. if a page is shorter than the viewport), then the user agent should not scroll anything. Visual user agents should further scroll horizontally as necessary to bring the element to the attention of the user.

Non-visual user agents may ignore the argument, or may treat it in some media-specific manner most useful to the user.

6.5 Focus

When an element is *focused*, key events received by the document must be targeted at that element. There may be no element focused; when no element is focused, key events received by the document must be targeted at the body element (page 81).

User agents may track focus for each browsing context (page 414) or Document individually, or may support only one focused element per top-level browsing context (page 415) — user agents should follow platform conventions in this regard.

Which element(s) within a top-level browsing context (page 415) currently has focus must be independent of whether or not the top-level browsing context (page 415) itself has the *system focus*.

Note: *When an element is focused, the element matches the CSS :focus pseudo-class.*

6.5.1 Sequential focus navigation

The **tabindex** content attribute specifies whether the element is focusable, whether it can be reached using sequential focus navigation, and the relative order of the element for the purposes of sequential focus navigation. The name "tab index" comes from the common use of the "tab" key to navigate through the focusable elements. The term "tabbing" refers to moving forward through the focusable elements that can be reached using sequential focus navigation.

The `tabindex` attribute, if specified, must have a value that is a valid integer (page 33).

If the attribute is specified, it must be parsed using the rules for parsing integers (page 33). The attribute's values have the following meanings:

If the attribute is omitted or parsing the value returns an error

The user agent should follow platform conventions to determine if the element is to be focusable and, if so, whether the element can be reached using sequential focus navigation, and if so, what its relative order should be.

If the value is a negative integer

The user agent must allow the element to be focused, but should not allow the element to be reached using sequential focus navigation.

If the value is a zero

The user agent must allow the element to be focused, should allow the element to be reached using sequential focus navigation, and should follow platform conventions to determine the element's relative order.

If the value is greater than zero

The user agent must allow the element to be focused, should allow the element to be reached using sequential focus navigation, and should place the element in the sequential focus navigation order so that it is:

- before any focusable element whose `tabindex` attribute has been omitted or whose value, when parsed, returns an error,
- before any focusable element whose `tabindex` attribute has a value equal to or less than zero,
- after any element whose `tabindex` attribute has a value greater than zero but less than the value of the `tabindex` attribute on the element,
- after any element whose `tabindex` attribute has a value equal to the value of the `tabindex` attribute on the element but that is earlier in the document in tree order (page 24) than the element,
- before any element whose `tabindex` attribute has a value equal to the value of the `tabindex` attribute on the element but that is later in the document in tree order (page 24) than the element, and
- before any element whose `tabindex` attribute has a value greater than the value of the `tabindex` attribute on the element.

An element is **focusable** if the tabindex attribute's definition above defines the element to be focusable *and* the element is being rendered.

The **tabIndex** DOM attribute must reflect (page 67) the value of the tabindex content attribute. If the attribute is not present, or parsing its value returns an error, then the DOM attribute must return 0 for elements that are focusable and –1 for elements that are not focusable.

6.5.2 Focus management

The **focusing steps** are as follows:

1. If focusing the element will remove the focus from another element, then run the unfocusing steps (page 516) for that element.
2. Make the element the currently focused element in its top-level browsing context (page 415).

Some elements, most notably area, can correspond to more than one distinct focusable area. If a particular area was indicated when the element was focused, then that is the area that must get focus; otherwise, e.g. when using the focus() method, the first such region in tree order is the one that must be focused.

3. Fire a simple event (page 436) that doesn't bubble called focus at the element.

User agents must run the focusing steps (page 516) for an element whenever the user moves the focus to a focusable (page 516) element.

The **unfocusing steps** are as follows:

1. If the element is an input element, and the change event applies to the element, and the element does not have a defined activation behavior (page 96), and the user has changed the element's value (page 362) or its list of selected files (page 337) while the control was focused without committing that change, then fire a simple event (page 436) called change at the element.
2. Unfocus the element.
3. Fire a simple event (page 436) that doesn't bubble called blur at the element.

User agents should run the unfocusing steps (page 516) for an element whenever the user moves the focus away from any focusable (page 516) element.

6.5.3 Document-level focus APIs

The **activeElement** attribute must return the element in the document that is focused. If no element in the Document is focused, this must return the body element (page 81).

The **hasFocus()** method must return true if the document, one of its nested browsing contexts (page 414), or any element in the document or its browsing contexts currently has the system focus.

6.5.4 Element-level focus APIs

The **focus()** method, when invoked, must run the following algorithm:

1. If the element is marked as *locked for focus* (page 517), then abort these steps.
2. If the element is not focusable (page 516), then abort these steps.
3. Mark the element as **locked for focus**.
4. If the element is not already focused, run the focusing steps (page 516) for the element.
5. Unmark the element as *locked for focus* (page 517).

The **blur()** method, when invoked, should run the unfocusing steps (page 516) for the element. User agents may selectively or uniformly ignore calls to this method for usability reasons.

6.6 The text selection APIs

Every browsing context (page 414) has **a selection**. The selection can be empty, and the selection can have more than one range (a disjointed selection). The user agent should allow the user to change the selection. User agents are not required to let the user select more than one range, and may collapse multiple ranges in the selection to a single range when the user interacts with the selection. (But, of course, the user agent may let the user create selections with multiple ranges.)

This one selection must be shared by all the content of the browsing context (though not by nested browsing contexts (page 414)), including any editing hosts in the document. (Editing hosts that are not inside a document cannot have a selection.)

If the selection is empty (collapsed, so that it has only one segment and that segment's start and end points are the same) then the selection's position should equal the caret position. When the selection is not empty, this specification does not define the caret position; user agents should follow platform conventions in deciding whether the caret is at the start of the selection, the end of the selection, or somewhere else.

On some platforms (such as those using Wordstar editing conventions), the caret position is totally independent of the start and end of the selection, even when the selection is empty. On such platforms, user agents may ignore the requirement that the cursor position be linked to the position of the selection altogether.

Mostly for historical reasons, in addition to the browsing context (page 414)'s selection (page 517), each `textarea` and `input` element has an independent selection. These are the **text field selections**.

User agents may selectively ignore attempts to use the API to adjust the selection made after the user has modified the selection. For example, if the user has just selected part of a word, the user agent could ignore attempts to use the API call to immediately unselect the selection altogether, but could allow attempts to change the selection to select the entire word.

User agents may also allow the user to create selections that are not exposed to the API.

The datagrid and select elements also have selections, indicating which items have been picked by the user. These are not discussed in this section.

Note: This specification does not specify how selections are presented to the user. The Selectors specification, in conjunction with CSS, can be used to style text selections using the ::selection pseudo-element. [SELECTORS] [CSS21]

6.6.1 APIs for the browsing context selection

The `getSelection()` method on the Window interface must return the Selection object representing the selection (page 517) of that Window object's browsing context (page 414).

For historical reasons, the `getSelection()` method on the HTMLDocument interface must return the same Selection object.

```
[Stringifies] interface Selection {
    readonly attribute Node anchorNode;
    readonly attribute long anchorOffset;
    readonly attribute Node focusNode;
    readonly attribute long focusOffset;
    readonly attribute boolean isCollapsed;
    void collapse(in Node parentNode, in long offset);
    void collapseToStart();
    void collapseToEnd();
    void selectAllChildren(in Node parentNode);
    void deleteFromDocument();
    readonly attribute long rangeCount;
    Range getRangeAt(in long index);
    void addRange(in Range range);
    void removeRange(in Range range);
    void removeAllRanges();
};
```

The Selection interface represents a list of Range objects. The first item in the list has index 0, and the last item has index *count*-1, where *count* is the number of ranges in the list. [DOM2RANGE]

All of the members of the Selection interface are defined in terms of operations on the Range objects represented by this object. These operations can raise exceptions, as defined for the Range interface; this can therefore result in the members of the Selection interface raising exceptions as well, in addition to any explicitly called out below.

The **anchorNode** attribute must return the value returned by the `startContainer` attribute of the last Range object in the list, or null if the list is empty.

The **anchorOffset** attribute must return the value returned by the `startOffset` attribute of the last Range object in the list, or 0 if the list is empty.

The **focusNode** attribute must return the value returned by the endContainer attribute of the last Range object in the list, or null if the list is empty.

The **focusOffset** attribute must return the value returned by the endOffset attribute of the last Range object in the list, or 0 if the list is empty.

The **isCollapsed** attribute must return true if there are zero ranges, or if there is exactly one range and its collapsed attribute is itself true. Otherwise it must return false.

The **collapse(*parentNode*, *offset*)** method must raise a WRONG_DOCUMENT_ERR DOM exception if *parentNode*'s Document is not the HTMLDocument object with which the Selection object is associated. Otherwise it is, and the method must remove all the ranges in the Selection list, then create a new Range object, add it to the list, and invoke its setStart() and setEnd() methods with the *parentNode* and *offset* values as their arguments.

The **collapseToStart()** method must raise an INVALID_STATE_ERR DOM exception if there are no ranges in the list. Otherwise, it must invoke the collapse() method with the startContainer and startOffset values of the first Range object in the list as the arguments.

The **collapseToEnd()** method must raise an INVALID_STATE_ERR DOM exception if there are no ranges in the list. Otherwise, it must invoke the collapse() method with the endContainer and endOffset values of the last Range object in the list as the arguments.

The **selectAllChildren(*parentNode*)** method must invoke the collapse() method with the *parentNode* value as the first argument and 0 as the second argument, and must then invoke the selectNodeContents() method on the first (and only) range in the list with the *parentNode* value as the argument.

The **deleteFromDocument()** method must invoke the deleteContents() method on each range in the list, if any, from first to last.

The **rangeCount** attribute must return the number of ranges in the list.

The **getRangeAt(*index*)** method must return the *index*th range in the list. If *index* is less than zero or greater or equal to the value returned by the rangeCount attribute, then the method must raise an INDEX_SIZE_ERR DOM exception.

The **addRange(*range*)** method must add the given *range* Range object to the list of selections, at the end (so the newly added range is the new last range). Duplicates are not prevented; a range may be added more than once in which case it appears in the list more than once, which (for example) will cause stringification (page 520) to return the range's text twice.

The **removeRange(*range*)** method must remove the first occurrence of *range* in the list of ranges, if it appears at all.

The **removeAllRanges()** method must remove all the ranges from the list of ranges, such that the rangeCount attribute returns 0 after the removeAllRanges() method is invoked (and until a new range is added to the list, either through this interface or via user interaction).

Objects implementing this interface must **stringify** to a concatenation of the results of invoking the `toString()` method of the Range object on each of the ranges of the selection, in the order they appear in the list (first to last).

In the following document fragment, the emphasised parts indicate the selection.

```
<p>The cute girl likes the <cite>Oxford English Dictionary</cite>.</p>
```

If a script invoked `window.getSelection().toString()`, the return value would be "the Oxford English".

Note: *The Selection interface has no relation to the DataGridSelection interface.*

6.6.2 APIs for the text field selections

- ** When we define `HTMLTextAreaElement` and `HTMLInputElement` we will have to add the IDL given below to both of their IDLs.

The `input` and `textarea` elements define four members in their DOM interfaces for handling their text selection:

```
void select();
    attribute unsigned long selectionStart;
    attribute unsigned long selectionEnd;
void setSelectionRange(in unsigned long start, in unsigned long end);
```

These methods and attributes expose and control the selection of `input` and `textarea` text fields.

The `select()` method must cause the contents of the text field to be fully selected.

The `selectionStart` attribute must, on getting, return the offset (in logical order) to the character that immediately follows the start of the selection. If there is no selection, then it must return the offset (in logical order) to the character that immediately follows the text entry cursor.

On setting, it must act as if the `setSelectionRange()` method had been called, with the new value as the first argument, and the current value of the `selectionEnd` attribute as the second argument, unless the current value of the `selectionEnd` is less than the new value, in which case the second argument must also be the new value.

The `selectionEnd` attribute must, on getting, return the offset (in logical order) to the character that immediately follows the end of the selection. If there is no selection, then it must return the offset (in logical order) to the character that immediately follows the text entry cursor.

On setting, it must act as if the `setSelectionRange()` method had been called, with the current value of the `selectionStart` attribute as the first argument, and new value as the second argument.

The **setSelectionRange(*start*, *end*)** method must set the selection of the text field to the sequence of characters starting with the character at the *start*th position (in logical order) and ending with the character at the (*end*-1)th position. Arguments greater than the length of the value in the text field must be treated as pointing at the end of the text field. If *end* is less than or equal to *start* then the start of the selection and the end of the selection must both be placed immediately before the character with offset *end*. In UAs where there is no concept of an empty selection, this must set the cursor to be just before the character with offset *end*.

To obtain the currently selected text, the following JavaScript suffices:

```
var selectionText = control.value.substring(control.selectionStart,  
control.selectionEnd);
```

...where *control* is the input or textarea element.

Characters with no visible rendering, such as U+200D ZERO WIDTH JOINER, still count as characters. Thus, for instance, the selection can include just an invisible character, and the text insertion cursor can be placed to one side or another of such a character.

When these methods and attributes are used with input elements that are not displaying simple text fields, they must raise an `INVALID_STATE_ERR` exception.

6.7 The **contenteditable** attribute

The **contenteditable** attribute is a common attribute. User agents must support this attribute on all HTML elements (page 23).

The **contenteditable** attribute is an enumerated attribute (page 51) whose keywords are the empty string, true, and false. The empty string and the true keyword map to the *true* state. The false keyword maps to the *false* state. In addition, there is a third state, the *inherit* state, which is the *missing value default* (and the *invalid value default*).

If an HTML element (page 23) has a **contenteditable** attribute set to the true state, or it has its **contenteditable** attribute set to the inherit state and if its nearest ancestor HTML element (page 23) with the **contenteditable** attribute set to a state other than the inherit state has its attribute set to the true state, or if it and its ancestors all have their **contenteditable** attribute set to the inherit state but the Document has `designMode` enabled, then the UA must treat the element as **editable** (as described below).

Otherwise, either the HTML element (page 23) has a **contenteditable** attribute set to the false state, or its **contenteditable** attribute is set to the inherit state and its nearest ancestor HTML element (page 23) with the **contenteditable** attribute set to a state other than the inherit state has its attribute set to the false state, or all its ancestors have their **contenteditable** attribute set to the inherit state and the Document itself has `designMode` disabled; either way, the element is not editable.

The **contentEditable** DOM attribute, on getting, must return the string "true" if the content attribute is set to the true state, "false" if the content attribute is set to the false state, and "inherit" otherwise. On setting, if the new value is an ASCII case-insensitive (page 31) match for the string "inherit" then the content attribute must be removed, if the new value is an ASCII case-insensitive (page 31) match for the string "true" then the content attribute

must be set to the string "true", if the new value is an ASCII case-insensitive (page 31) match for the string "false" then the content attribute must be set to the string "false", and otherwise the attribute setter must raise a SYNTAX_ERR exception.

The **isContentEditable** DOM attribute, on getting, must return true if the element is editable (page 521), and false otherwise.

If an element is editable (page 521) and its parent element is not, or if an element is editable (page 521) and it has no parent element, then the element is an **editing host**. Editable elements can be nested. User agents must make editing hosts focusable (which typically means they enter the tab order (page 515)). An editing host can contain non-editable sections, these are handled as described below. An editing host can contain non-editable sections that contain further editing hosts.

When an editing host has focus, it must have a **caret position** that specifies where the current editing position is. It may also have a selection (page 517).

Note: *How the caret and selection are represented depends entirely on the UA.*

6.7.1 User editing actions

There are several actions that the user agent should allow the user to perform while the user is interacting with an editing host. How exactly each action is triggered is not defined for every action, but when it is not defined, suggested key bindings are provided to guide implementors.

Move the caret

User agents must allow users to move the caret to any position within an editing host, even into nested editable elements. This could be triggered as the default action of keydown events with various key identifiers and as the default action of mousedown events.

Change the selection

User agents must allow users to change the selection (page 517) within an editing host, even into nested editable elements. User agents may prevent selections from being made in ways that cross from editable elements into non-editable elements (e.g. by making each non-editable descendant atomically selectable, but not allowing text selection within them). This could be triggered as the default action of keydown events with various key identifiers and as the default action of mousedown events.

Insert text

This action must be triggered as the default action of a `textInput` event, and may be triggered by other commands as well. It must cause the user agent to insert the specified text (given by the event object's `data` attribute in the case of the `textInput` event) at the caret.

If the caret is positioned somewhere where phrasing content (page 94) is not allowed (e.g. inside an empty `ol` element), then the user agent must not insert the text directly at the caret position. In such cases the behavior is UA-dependent, but user agents must

not, in response to a request to insert text, generate a DOM that is less conformant than the DOM prior to the request.

User agents should allow users to insert new paragraphs into elements that contains only content other than paragraphs.

For example, given the markup:

```
<section>
  <dl>
    <dt> Ben </dt>
    <dd> Goat </dd>
  </dl>
</section>
```

...the user agent should allow the user to insert p elements before and after the dl element, as children of the section element.

Break block

UAs should offer a way for the user to request that the current paragraph be broken at the caret, e.g. as the default action of a keydown event whose identifier is the "Enter" key and that has no modifiers set.

The exact behavior is UA-dependent, but user agents must not, in response to a request to break a paragraph, generate a DOM that is less conformant than the DOM prior to the request.

Insert a line separator

UAs should offer a way for the user to request an explicit line break at the caret position without breaking the paragraph, e.g. as the default action of a keydown event whose identifier is the "Enter" key and that has a shift modifier set. Line separators are typically found within a poem verse or an address. To insert a line break, the user agent must insert a br element.

If the caret is positioned somewhere where phrasing content (page 94) is not allowed (e.g. in an empty ol element), then the user agent must not insert the br element directly at the caret position. In such cases the behavior is UA-dependent, but user agents must not, in response to a request to insert a line separator, generate a DOM that is less conformant than the DOM prior to the request.

Delete

UAs should offer a way for the user to delete text and elements, including non-editable descendants, e.g. as the default action of keydown events whose identifiers are "U+0008" or "U+007F".

Five edge cases in particular need to be considered carefully when implementing this feature: backspacing at the start of an element, backspacing when the caret is immediately after an element, forward-deleting at the end of an element, forward-deleting when the caret is immediately before an element, and deleting a selection (page 517) whose start and end points do not share a common parent node.

In any case, the exact behavior is UA-dependent, but user agents must not, in response to a request to delete text or an element, generate a DOM that is less conformant than the DOM prior to the request.

Insert, and wrap text in, semantic elements

UAs should offer the user the ability to mark text and paragraphs with semantics that HTML can express.

UAs should similarly offer a way for the user to insert empty semantic elements to subsequently fill by entering text manually.

UAs should also offer a way to remove those semantics from marked up text, and to remove empty semantic element that have been inserted.

In response to a request from a user to mark text up in italics, user agents should use the **i** element to represent the semantic. The **em** element should be used only if the user agent is sure that the user means to indicate stress emphasis.

In response to a request from a user to mark text up in bold, user agents should use the **b** element to represent the semantic. The **strong** element should be used only if the user agent is sure that the user means to indicate importance.

The exact behavior is UA-dependent, but user agents must not, in response to a request to wrap semantics around some text or to insert or remove a semantic element, generate a DOM that is less conformant than the DOM prior to the request.

Select and move non-editable elements nested inside editing hosts

UAs should offer a way for the user to move images and other non-editable parts around the content within an editing host. This may be done using the drag and drop (page 525) mechanism. User agents must not, in response to a request to move non-editable elements nested inside editing hosts, generate a DOM that is less conformant than the DOM prior to the request.

Edit form controls nested inside editing hosts

When an editable (page 521) form control is edited, the changes must be reflected in both its current value *and* its default value. For input elements this means updating the `defaultValue` DOM attribute as well as the `value` DOM attribute; for select elements it means updating the option elements' `defaultSelected` DOM attribute as well as the `selected` DOM attribute; for textarea elements this means updating the `defaultValue` DOM attribute as well as the `value` DOM attribute. (Updating the `default*` DOM attributes causes content attributes to be updated as well.)

User agents may perform several commands per user request; for example if the user selects a block of text and hits Enter, the UA might interpret that as a request to delete the content of the selection (page 517) followed by a request to break the paragraph at that position.

6.7.2 Making entire documents editable

Documents have a **designMode**, which can be either enabled or disabled.

The `designMode` DOM attribute on the Document object takes two values, "on" and "off". When it is set, the new value must be compared in an ASCII case-insensitive (page 31) manner to these two values. If it matches the "on" value, then `designMode` must be enabled, and if it matches the "off" value, then `designMode` must be disabled. Other values must be ignored.

When designMode is enabled, the DOM attribute must return the value "on", and when it is disabled, it must return the value "off".

The last state set must persist until the document is destroyed or the state is changed. Initially, documents must have their designMode disabled.

Enabling designMode causes scripts in general to be disabled (page 428) and the document to become editable.

6.8 Drag and drop

This section defines an event-based drag-and-drop mechanism.

This specification does not define exactly what a *drag-and-drop operation* actually is.

On a visual medium with a pointing device, a drag operation could be the default action of a mousedown event that is followed by a series of mousemove events, and the drop could be triggered by the mouse being released.

On media without a pointing device, the user would probably have to explicitly indicate his intention to perform a drag-and-drop operation, stating what he wishes to drag and what he wishes to drop, respectively.

However it is implemented, drag-and-drop operations must have a starting point (e.g. where the mouse was clicked, or the start of the selection (page 517) or element that was selected for the drag), may have any number of intermediate steps (elements that the mouse moves over during a drag, or elements that the user picks as possible drop points as he cycles through possibilities), and must either have an end point (the element above which the mouse button was released, or the element that was finally selected), or be canceled. The end point must be the last element selected as a possible drop point before the drop occurs (so if the operation is not canceled, there must be at least one element in the middle step).

6.8.1 Introduction

This section is non-normative.

** It's also currently non-existent.

6.8.2 The DragEvent and DataTransfer interfaces

The drag-and-drop processing model involves several events. They all use the DragEvent interface.

```
interface DragEvent : UIEvent {  
    readonly attribute DataTransfer dataTransfer;  
    void initDragEvent(in DOMString typeArg, in boolean canBubbleArg, in  
        boolean cancelableArg, in AbstractView viewArg, in long detailArg, in  
        DataTransfer dataTransferArg);  
    void initDragEventNS(in DOMString namespaceURIArg, in DOMString
```

```
    typeArg, in boolean canBubbleArg, in boolean cancelableArg, in  
    AbstractView viewArg, in long detailArg, in DataTransfer  
    dataTransferArg);  
};
```

- ** We should have modifier key information in here too (shift/ctrl, etc), like with mouse events and like with the context menu event.

The **initDragEvent()** and **initDragEventNS()** methods must initialize the event in a manner analogous to the similarly-named methods in the DOM3 Events interfaces.
[DOM3EVENTS]

The **dataTransfer** attribute of the DragEvent interface represents the context information for the event.

```
interface DataTransfer {  
    attribute DOMString dropEffect;  
    attribute DOMString effectAllowed;  
    readonly attribute DOMStringList types;  
    void clearData(in DOMString format);  
    void setData(in DOMString format, in DOMString data);  
    DOMString getData(in DOMString format);  
    void setDragImage(in Element image, in long x, in long y);  
    void addElement(in Element element);  
};
```

DataTransfer objects can conceptually contain various kinds of data.

When a DataTransfer object is created, it must be initialized as follows:

- The DataTransfer object must initially contain no data, no elements, and have no associated image.
- The DataTransfer object's effectAllowed attribute must be set to "uninitialized".
- The dropEffect attribute must be set to "none".

The **dropEffect** attribute controls the drag-and-drop feedback that the user is given during a drag-and-drop operation.

The attribute must ignore any attempts to set it to a value other than none, copy, link, and move. On getting, the attribute must return the last of those four values that it was set to.

The **effectAllowed** attribute is used in the drag-and-drop processing model to initialize the dropEffect attribute during the dragenter and dragover events.

The attribute must ignore any attempts to set it to a value other than none, copy, copyLink, copyMove, link, linkMove, move, all, and uninitialized. On getting, the attribute must return the last of those values that it was set to.

DataTransfer objects can hold pieces of data, each associated with a unique format. Formats are generally given by MIME types, with some values special-cased for legacy reasons.

The **clearData(*format*)** method must clear the DataTransfer object of any data associated with the given *format*. If *format* is the value "Text", then it must be treated as "text/plain". If the *format* is "URL", then it must be treated as "text/uri-list".

The **setData(*format*, *data*)** method must add *data* to the data stored in the DataTransfer object, labeled as being of the type *format*. This must replace any previous data that had been set for that format. If *format* is the value "Text", then it must be treated as "text/plain". If the *format* is "URL", then it must be treated as "text/uri-list".

The **getData(*format*)** method must return the data that is associated with the type *format*, if any, and must return the empty string otherwise. If *format* is the value "Text", then it must be treated as "text/plain". If the *format* is "URL", then the data associated with the "text/uri-list" format must be parsed as appropriate for text/uri-list data, and the first URL from the list must be returned. If there is no data with that format, or if there is but it has no URLs, then the method must return the empty string. [RFC2483]

The **types** attribute must return a live DOMStringList that contains the list of formats that are stored in the DataTransfer object.

The **setDragImage(*element*, *x*, *y*)** method sets which element to use to generate the drag feedback (page 530). The *element* argument can be any Element; if it is an img element, then the user agent should use the element's image (at its intrinsic size) to generate the feedback, otherwise the user agent should base the feedback on the given element (but the exact mechanism for doing so is not specified).

The **addElement(*element*)** method is an alternative way of specifying how the user agent is to render the drag feedback (page 530). It adds an element to the DataTransfer object.

6.8.3 Events fired during a drag-and-drop action

The following events are involved in the drag-and-drop model. Whenever the processing model described below causes one of these events to be fired, the event fired must use the DragEvent interface defined above, must have the bubbling and cancelable behaviors given in the table below, and must have the context information set up as described after the table, with the view attribute set to the view with which the user interacted to trigger the drag-and-drop event, and the detail attribute set to zero.

Event Name	Target	Bubbles?	Cancelable?	dataTransfer	effectAllowed	dropEffect	Default Action
dragstart	Source node (page 529)	✓ Bubbles	✓ Cancelable	Contains source node (page 529) unless a selection is being dragged, in which case it is empty	uninitialized	none	Initiate the drag-and-drop operation

Event Name	Target	Bubbles?	Cancelable?	dataTransfer	effectAllowed	dropEffect	Default Action
drag	Source node (page 529)	✓ Bubbles	✓ Cancelable	Empty	Same as last event (page 528)	none	Continue the drag-and-drop operation
dragenter	Immediate user selection (page 530) or the body element (page 81)	✓ Bubbles	✓ Cancelable	Empty	Same as last event (page 528)	Based on effectAllowed value (page 528)	Reject immediate user selection (page 530) as potential target element (page 530)
dragleave	Previous target element (page 530)	✓ Bubbles	—	Empty	Same as last event (page 528)	none	None
dragover	Current target element (page 530)	✓ Bubbles	✓ Cancelable	Empty	Same as last event (page 528)	Based on effectAllowed value (page 528)	Reset the current drag operation (page 530) to "none"
drop	Current target element (page 530)	✓ Bubbles	✓ Cancelable	getData() returns data set in dragstart event	Same as last event (page 528)	Current drag operation (page 530)	Varies
dragend	Source node (page 529)	✓ Bubbles	—	Empty	Same as last event (page 528)	Current drag operation (page 530)	Varies

The dataTransfer object's contents are empty except for dragstart events and drop events, for which the contents are set as described in the processing model, below.

The effectAllowed attribute must be set to "uninitialized" for dragstart events, and to whatever value the field had after the last drag-and-drop event was fired for all other events (only counting events fired by the user agent for the purposes of the drag-and-drop model described below).

The dropEffect attribute must be set to "none" for dragstart, drag, and dragleave events (except when stated otherwise in the algorithms given in the sections below), to the value corresponding to the current drag operation (page 530) for drop and dragend events, and to a value based on the effectAllowed attribute's value and to the drag-and-drop source, as given by the following table, for the remaining events (dragenter and dragover):

effectAllowed	dropEffect
none	none
copy, copyLink, copyMove, all	copy
link, linkMove	link
move	move
uninitialized when what is being dragged is a selection from a text field	move
uninitialized when what is being dragged is a selection	copy

effectAllowed	dropEffect
uninitialized when what is being dragged is an a element with an href attribute	link
Any other case	copy

6.8.4 Drag-and-drop processing model

When the user attempts to begin a drag operation, the user agent must first determine what is being dragged. If the drag operation was invoked on a selection, then it is the selection that is being dragged. Otherwise, it is the first element, going up the ancestor chain, starting at the node that the user tried to drag, that has the DOM attribute draggable set to true. If there is no such element, then nothing is being dragged, the drag-and-drop operation is never started, and the user agent must not continue with this algorithm.

Note: *img elements and a elements with an href attribute have their draggable attribute set to true by default.*

If the user agent determines that something can be dragged, a dragstart event must then be fired.

If it is a selection that is being dragged, then this event must be fired on the node that the user started the drag on (typically the text node that the user originally clicked). If the user did not specify a particular node, for example if the user just told the user agent to begin a drag of "the selection", then the event must be fired on the deepest node that is a common ancestor of all parts of the selection.

- ** We should look into how browsers do other types (e.g. Firefox apparently also adds text/html for internal drag and drop of a selection).

If it is not a selection that is being dragged, then the event must be fired on the element that is being dragged.

The node on which the event is fired is the **source node**. Multiple events are fired on this node during the course of the drag-and-drop operation.

If it is a selection that is being dragged, the dataTransfer member of the event must be created with no nodes. Otherwise, it must be created containing just the source node (page 529). Script can use the addElement() method to add further elements to the list of what is being dragged.

If it is a selection that is being dragged, the dataTransfer member of the event must have the text of the selection added to it as the data associated with the text/plain format. Otherwise, if it is an img element being dragged, then the value of the element's src DOM attribute must be added, associated with the text/uri-list format. Otherwise, if it is an a element being dragged, then the value of the element's href DOM attribute must be added, associated with the text/uri-list format. Otherwise, no data is added to the object by the user agent.

If the event is canceled, then the drag-and-drop operation must not occur; the user agent must not continue with this algorithm.

If it is not canceled, then the drag-and-drop operation must be initiated.

Note: Since events with no event handlers registered are, almost by definition, never canceled, drag-and-drop is always available to the user if the author does not specifically prevent it.

The drag-and-drop feedback must be generated from the first of the following sources that is available:

1. The element specified in the last call to the `setDragImage()` method of the `dataTransfer` object of the `dragstart` event, if the method was called. In visual media, if this is used, the `x` and `y` arguments that were passed to that method should be used as hints for where to put the cursor relative to the resulting image. The values are expressed as distances in CSS pixels from the left side and from the top side of the image respectively. [CSS21]
2. The elements that were added to the `dataTransfer` object, both before the event was fired, and during the handling of the event using the `addElement()` method, if any such elements were indeed added.
3. The selection that the user is dragging.

The user agent must take a note of the data that was placed (page 527) in the `dataTransfer` object. This data will be made available again when the drop event is fired.

From this point until the end of the drag-and-drop operation, device input events (e.g. mouse and keyboard events) must be suppressed. In addition, the user agent must track all DOM changes made during the drag-and-drop operation, and add them to its undo history (page 536) as one atomic operation once the drag-and-drop operation has ended.

During the drag operation, the element directly indicated by the user as the drop target is called the **immediate user selection**. (Only elements can be selected by the user; other nodes must not be made available as drop targets.) However, the immediate user selection (page 530) is not necessarily the **current target element**, which is the element currently selected for the drop part of the drag-and-drop operation. The immediate user selection (page 530) changes as the user selects different elements (either by pointing at them with a pointing device, or by selecting them in some other way). The current target element (page 530) changes when the immediate user selection (page 530) changes, based on the results of event handlers in the document, as described below.

Both the current target element (page 530) and the immediate user selection (page 530) can be null, which means no target element is selected. They can also both be elements in other (DOM-based) documents, or other (non-Web) programs altogether. (For example, a user could drag text to a word-processor.) The current target element (page 530) is initially null.

In addition, there is also a **current drag operation**, which can take on the values "none", "copy", "link", and "move". Initially it has the value "none". It is updated by the user agent as described in the steps below.

User agents must, every 350ms ($\pm 200\text{ms}$), perform the following steps in sequence. (If the user agent is still performing the previous iteration of the sequence when the next iteration becomes due, the user agent must not execute the overdue iteration, effectively "skipping missed frames" of the drag-and-drop operation.)

1. First, the user agent must fire a drag event at the source node (page 529). If this event is canceled, the user agent must set the current drag operation (page 530) to none (no drag operation).
2. Next, if the drag event was not canceled and the user has not ended the drag-and-drop operation, the user agent must check the state of the drag-and-drop operation, as follows:
 1. First, if the user is indicating a different immediate user selection (page 530) than during the last iteration (or if this is the first iteration), and if this immediate user selection (page 530) is not the same as the current target element (page 530), then the current target element (page 530) must be updated, as follows:
 1. If the new immediate user selection (page 530) is null, or is in a non-DOM document or application, then set the current target element (page 530) to the same value.
 2. Otherwise, the user agent must fire a dragenter event at the immediate user selection (page 530).
 3. If the event is canceled, then the current target element (page 530) must be set to the immediate user selection (page 530).
 4. Otherwise, if the current target element (page 530) is not the body element (page 81), the user agent must fire a dragenter event at the body element (page 81), and the current target element (page 530) must be set to the body element (page 81), regardless of whether that event was canceled or not. (If the body element (page 81) is null, then the current target element (page 530) would be set to null too in this case, it wouldn't be set to the Document object.)
 2. If the previous step caused the current target element (page 530) to change, and if the previous target element was not null or a part of a non-DOM document, the user agent must fire a dragleave event at the previous target element.
 3. If the current target element (page 530) is a DOM element, the user agent must fire a dragover event at this current target element (page 530).

If the dragover event is not canceled, the current drag operation (page 530) must be reset to "none".

Otherwise, the current drag operation (page 530) must be set based on the values the effectAllowed and dropEffect attributes of the dataTransfer object had after the event was handled, as per the following table:

effectAllowed	dropEffect	Drag operation
uninitialized, copy, copyLink, copyMove, or all	copy	"copy"
uninitialized, link, copyLink, linkMove, or all	link	"link"
uninitialized, move, copyMove, linkMove, or all	move	"move"
Any other case		"none"

Then, regardless of whether the dragover event was canceled or not, the drag feedback (e.g. the mouse cursor) must be updated to match the current drag operation (page 530), as follows:

Drag operation	Feedback
"copy"	Data will be copied if dropped here.
"link"	Data will be linked if dropped here.
"move"	Data will be moved if dropped here.
"none"	No operation allowed, dropping here will cancel the drag-and-drop operation.

4. Otherwise, if the current target element (page 530) is not a DOM element, the user agent must use platform-specific mechanisms to determine what drag operation is being performed (none, copy, link, or move). This sets the *current drag operation*.
3. Otherwise, if the user ended the drag-and-drop operation (e.g. by releasing the mouse button in a mouse-driven drag-and-drop interface), or if the drag event was canceled, then this will be the last iteration. The user agent must execute the following steps, then stop looping.
 1. If the current drag operation (page 530) is none (no drag operation), or, if the user ended the drag-and-drop operation by canceling it (e.g. by hitting the Escape key), or if the current target element (page 530) is null, then the drag operation failed. If the current target element (page 530) is a DOM element, the user agent must fire a dragleave event at it; otherwise, if it is not null, it must use platform-specific conventions for drag cancellation.
 2. Otherwise, the drag operation was as success. If the current target element (page 530) is a DOM element, the user agent must fire a drop event at it; otherwise, it must use platform-specific conventions for indicating a drop.

When the target is a DOM element, the dropEffect attribute of the event's dataTransfer object must be given the value representing the current drag operation (page 530) (copy, link, or move), and the object must be set up so that the getData() method will return the data that was added during the dragstart event.

If the event is canceled, the current drag operation (page 530) must be set to the value of the dropEffect attribute of the event's dataTransfer object as it stood after the event was handled.

Otherwise, the event is not canceled, and the user agent must perform the event's default action, which depends on the exact target as follows:

↪ **If the current target element (page 530) is a text field (e.g. textarea, or an input element with type="text")**

The user agent must insert the data associated with the text/plain format, if any, into the text field in a manner consistent with platform-specific conventions (e.g. inserting it at the current mouse cursor position, or inserting it at the end of the field).

↪ **Otherwise**

Reset the current drag operation (page 530) to "none".

3. Finally, the user agent must fire a dragend event at the source node (page 529), with the dropEffect attribute of the event's dataTransfer object being set to the value corresponding to the current drag operation (page 530).

Note: *The current drag operation (page 530) can change during the processing of the drop event, if one was fired.*

The event is not cancelable. After the event has been handled, the user agent must act as follows:

↪ **If the current target element (page 530) is a text field (e.g. textarea, or an input element with type="text"), and a drop event was fired in the previous step, and the current drag operation (page 530) is "move", and the source of the drag-and-drop operation is a selection in the DOM**

The user agent should delete the range representing the dragged selection from the DOM.

↪ **If the current target element (page 530) is a text field (e.g. textarea, or an input element with type="text"), and a drop event was fired in the previous step, and the current drag operation (page 530) is "move", and the source of the drag-and-drop operation is a selection in a text field**

The user agent should delete the dragged selection from the relevant text field.

↪ **Otherwise**

The event has no default action.

6.8.4.1 When the drag-and-drop operation starts or ends in another document

The model described above is independent of which Document object the nodes involved are from; the events must be fired as described above and the rest of the processing model must be followed as described above, irrespective of how many documents are involved in the operation.

6.8.4.2 When the drag-and-drop operation starts or ends in another application

If the drag is initiated in another application, the source node (page 529) is not a DOM node, and the user agent must use platform-specific conventions instead when the requirements above involve the source node. User agents in this situation must act as if the dragged data had been added to the DataTransfer object when the drag started, even though no dragstart event was actually fired; user agents must similarly use platform-specific conventions when deciding on what drag feedback to use.

If a drag is started in a document but ends in another application, then the user agent must instead replace the parts of the processing model relating to handling the *target* according to platform-specific conventions.

In any case, scripts running in the context of the document must not be able to distinguish the case of a drag-and-drop operation being started or ended in another application from the case of a drag-and-drop operation being started or ended in another document from another domain.

6.8.5 The **draggable** attribute

All elements may have the draggable content attribute set. The draggable attribute is an enumerated attribute (page 51). It has three states. The first state is *true* and it has the keyword *true*. The second state is *false* and it has the keyword *false*. The third state is *auto*; it has no keywords but it is the *missing value default*.

The **draggable** DOM attribute, whose value depends on the content attribute's in the way described below, controls whether or not the element is draggable. Generally, only text selections are draggable, but elements whose draggable DOM attribute is true become draggable as well.

If an element's draggable content attribute has the state *true*, the draggable DOM attribute must return true.

Otherwise, if the element's draggable content attribute has the state *false*, the draggable DOM attribute must return false.

Otherwise, the element's draggable content attribute has the state *auto*. If the element is an *img* element, or, if the element is an *a* element with an *href* content attribute, the draggable DOM attribute must return true.

Otherwise, the draggable DOM must return false.

If the draggable DOM attribute is set to the value *false*, the draggable content attribute must be set to the literal value *false*. If the draggable DOM attribute is set to the value *true*, the draggable content attribute must be set to the literal value *true*.

6.8.6 Copy and paste

Copy-and-paste is a form of drag-and-drop: the "copy" part is equivalent to dragging content to another application (the "clipboard"), and the "paste" part is equivalent to dragging content *from* another application.

Select-and-paste (a model used by mouse operations in the X Window System) is equivalent to a drag-and-drop operation where the source is the selection.

6.8.6.1 Copy to clipboard

When the user invokes a copy operation, the user agent must act as if the user had invoked a drag on the current selection. If the drag-and-drop operation initiates, then the user agent must act as if the user had indicated (as the immediate user selection (page 530)) a

hypothetical application representing the clipboard. Then, the user agent must act as if the user had ended the drag-and-drop operation without canceling it. If the drag-and-drop operation didn't get canceled, the user agent should then follow the relevant platform-specific conventions for copy operations (e.g. updating the clipboard).

6.8.6.2 Cut to clipboard

When the user invokes a cut operation, the user agent must act as if the user had invoked a copy operation (see the previous section), followed, if the copy was completed successfully, by a selection delete operation (page 523).

6.8.6.3 Paste from clipboard

When the user invokes a clipboard paste operation, the user agent must act as if the user had invoked a drag on a hypothetical application representing the clipboard, setting the data associated with the drag as the content on the clipboard (in whatever formats are available).

Then, the user agent must act as if the user had indicated (as the immediate user selection (page 530)) the element with the keyboard focus, and then ended the drag-and-drop operation without canceling it.

6.8.6.4 Paste from selection

When the user invokes a selection paste operation, the user agent must act as if the user had invoked a drag on the current selection, then indicated (as the immediate user selection (page 530)) the element with the keyboard focus, and then ended the drag-and-drop operation without canceling it.

6.8.7 Security risks in the drag-and-drop model

User agents must not make the data added to the DataTransfer object during the dragstart event available to scripts until the drop event, because otherwise, if a user were to drag sensitive information from one document to a second document, crossing a hostile third document in the process, the hostile document could intercept the data.

For the same reason, user agents must consider a drop to be successful only if the user specifically ended the drag operation — if any scripts end the drag operation, it must be considered unsuccessful (canceled) and the drop event must not be fired.

User agents should take care to not start drag-and-drop operations in response to script actions. For example, in a mouse-and-window environment, if a script moves a window while the user has his mouse button depressed, the UA would not consider that to start a drag. This is important because otherwise UAs could cause data to be dragged from sensitive sources and dropped into hostile documents without the user's consent.

6.9 Undo history

** There has got to be a better way of doing this, surely.

The user agent must associate an **undo transaction history** with each `HTMLDocument` object.

The undo transaction history (page 536) is a list of entries. The entries are of two type: DOM changes (page 536) and undo objects (page 536).

Each **DOM changes** entry in the undo transaction history (page 536) consists of batches of one or more of the following:

- Changes to the content attributes (page 23) of an `Element` node.
- Changes to the DOM attributes (page 23) of a `Node`.
- Changes to the DOM hierarchy of nodes that are descendants of the `HTMLDocument` object (`parentNode`, `childNodes`).

Undo object entries consist of objects representing state that scripts running in the document are managing. For example, a Web mail application could use an undo object (page 536) to keep track of the fact that a user has moved an e-mail to a particular folder, so that the user can undo the action and have the e-mail return to its former location.

Broadly speaking, DOM changes (page 536) entries are handled by the UA in response to user edits of form controls and editing hosts on the page, and undo object (page 536) entries are handled by script in response to higher-level user actions (such as interactions with server-side state, or in the implementation of a drawing tool).

6.9.1 The UndoManager interface

** This API sucks. Seriously. It's a terrible API. Really bad. I hate it. Here are the requirements:

- Has to cope with cases where the server has undo state already when the page is loaded, that can be stuffed into the undo buffer onload.
- Has to support undo/redo.
- Has to cope with the "undo" action being "contact the server and tell it to undo", rather than it being the opposite of the "redo" action.
- Has to cope with some undo states expiring from the undo history (e.g. server can only remember one undelete action) but other states not expiring (e.g. client can undo arbitrary amounts of local edits).

To manage undo object (page 536) entries in the undo transaction history (page 536), the `UndoManager` interface can be used:

```

interface UndoManager {
    unsigned long add(in DOMObject data, in DOMString title);
    [XXX] void remove(in unsigned long index);
    void clearUndo();
    void clearRedo();
    [IndexGetter] DOMObject item(in unsigned long index);
    readonly attribute unsigned long length;
    readonly attribute unsigned long position;
};

```

The **undoManager** attribute of the Window interface must return the object implementing the UndoManager interface for that Window object's associated HTMLDocument object.

UndoManager objects represent their document's undo transaction history (page 536). Only undo object (page 536) entries are visible with this API, but this does not mean that DOM changes (page 536) entries are absent from the undo transaction history (page 536).

The **length** attribute must return the number of undo object (page 536) entries in the undo transaction history (page 536).

The **item(n)** method must return the *n*th undo object (page 536) entry in the undo transaction history (page 536).

The undo transaction history (page 536) has a **current position**. This is the position between two entries in the undo transaction history (page 536)'s list where the previous entry represents what needs to happen if the user invokes the "undo" command (the "undo" side, lower numbers), and the next entry represents what needs to happen if the user invokes the "redo" command (the "redo" side, higher numbers).

The **position** attribute must return the index of the undo object (page 536) entry nearest to the undo position (page 537), on the "redo" side. If there are no undo object (page 536) entries on the "redo" side, then the attribute must return the same as the length attribute. If there are no undo object (page 536) entries on the "undo" side of the undo position (page 537), the position attribute returns zero.

Note: Since the undo transaction history (page 536) contains both undo object (page 536) entries and DOM changes (page 536) entries, but the position attribute only returns indices relative to undo object (page 536) entries, it is possible for several "undo" or "redo" actions to be performed without the value of the position attribute changing.

The **add(data, title)** method's behavior depends on the current state. Normally, it must insert the *data* object passed as an argument into the undo transaction history (page 536) immediately before the undo position (page 537), optionally remembering the given *title* to use in the UI. If the method is called during an undo operation (page 538), however, the object must instead be added immediately after the undo position (page 537).

If the method is called and there is neither an undo operation in progress (page 538) nor a redo operation in progress (page 539) then any entries in the undo transaction history (page 536) after the undo position (page 537) must be removed (as if `clearRedo()` had been called).

- ** We could fire events when someone adds something to the undo history -- one event per undo object entry before the position (or after, during redo addition), allowing the script to decide if that entry should remain or not. Or something. Would make it potentially easier to expire server-held state when the server limitations come into play.

The `remove(index)` method must remove the undo object (page 536) entry with the specified `index`. If the index is less than zero or greater than or equal to `length` then the method must raise an `INDEX_SIZE_ERR` exception. DOM changes (page 536) entries are unaffected by this method.

The `clearUndo()` method must remove all entries in the undo transaction history (page 536) before the undo position (page 537), be they DOM changes (page 536) entries or undo object (page 536) entries.

The `clearRedo()` method must remove all entries in the undo transaction history (page 536) after the undo position (page 537), be they DOM changes (page 536) entries or undo object (page 536) entries.

- ** Another idea is to have a way for scripts to say "startBatchingDOMChangesForUndo()" and after that the changes to the DOM go in as if the user had done them.

6.9.2 Undo: moving back in the undo transaction history

When the user invokes an undo operation, or when the `execCommand()` method is called with the `undo` command, the user agent must perform an undo operation.

If the undo position (page 537) is at the start of the undo transaction history (page 536), then the user agent must do nothing.

If the entry immediately before the undo position (page 537) is a DOM changes (page 536) entry, then the user agent must remove that DOM changes (page 536) entry, reverse the DOM changes that were listed in that entry, and, if the changes were reversed with no problems, add a new DOM changes (page 536) entry (consisting of the opposite of those DOM changes) to the undo transaction history (page 536) on the other side of the undo position (page 537).

If the DOM changes cannot be undone (e.g. because the DOM state is no longer consistent with the changes represented in the entry), then the user agent must simply remove the DOM changes (page 536) entry, without doing anything else.

If the entry immediately before the undo position (page 537) is an undo object (page 536) entry, then the user agent must first remove that undo object (page 536) entry from the undo transaction history (page 536), and then must fire an `undo` event on the `Document` object, using the undo object (page 536) entry's associated undo object as the event's data.

Any calls to `add()` while the event is being handled will be used to populate the redo history, and will then be used if the user invokes the "redo" command to undo his undo.

6.9.3 Redo: moving forward in the undo transaction history

When the user invokes a redo operation, or when the execCommand() method is called with the redo command, the user agent must perform a redo operation.

This is mostly the opposite of an undo operation (page 538), but the full definition is included here for completeness.

If the undo position (page 537) is at the end of the undo transaction history (page 536), then the user agent must do nothing.

If the entry immediately after the undo position (page 537) is a DOM changes (page 536) entry, then the user agent must remove that DOM changes (page 536) entry, reverse the DOM changes that were listed in that entry, and, if the changes were reversed with no problems, add a new DOM changes (page 536) entry (consisting of the opposite of those DOM changes) to the undo transaction history (page 536) on the other side of the undo position (page 537).

If the DOM changes cannot be redone (e.g. because the DOM state is no longer consistent with the changes represented in the entry), then the user agent must simply remove the DOM changes (page 536) entry, without doing anything else.

If the entry immediately after the undo position (page 537) is an undo object (page 536) entry, then the user agent must first remove that undo object (page 536) entry from the undo transaction history (page 536), and then must fire a redo event on the Document object, using the undo object (page 536) entry's associated undo object as the event's data.

6.9.4 The UndoManagerEvent interface and the undo and redo events

```
interface UndoManagerEvent : Event {  
    readonly attribute DOMObject data;  
    void initUndoManagerEvent(in DOMString typeArg, in boolean  
        canBubbleArg, in boolean cancelableArg, in DOMObject dataArg);  
    void initUndoManagerEventNS(in DOMString namespaceURIArg, in DOMString  
        typeArg, in boolean canBubbleArg, in boolean cancelableArg, in DOMObject  
        dataArg);  
};
```

The **initUndoManagerEvent()** and **initUndoManagerEventNS()** methods must initialize the event in a manner analogous to the similarly-named methods in the DOM3 Events interfaces. [DOM3EVENTS]

The **data** attribute represents the undo object (page 536) for the event.

The **undo** and **redo** events do not bubble, cannot be canceled, and have no default action. When the user agent fires one of these events it must use the UndoManagerEvent interface, with the data field containing the relevant undo object (page 536).

6.9.5 Implementation notes

How user agents present the above conceptual model to the user is not defined. The undo interface could be a filtered view of the undo transaction history (page 536), it could manipulate the undo transaction history (page 536) in ways not described above, and so forth. For example, it is possible to design a UA that appears to have separate undo transaction histories (page 536) for each form control; similarly, it is possible to design systems where the user has access to more undo information than is present in the official (as described above) undo transaction history (page 536) (such as providing a tree-based approach to document state). Such UI models should be based upon the single undo transaction history (page 536) described in this section, however, such that to a script there is no detectable difference.

6.10 Command APIs

The **execCommand(*commandId*, *showUI*, *value*)** method on the `HTMLDocument` interface allows scripts to perform actions on the current selection (page 517) or at the current caret position. Generally, these commands would be used to implement editor UI, for example having a "delete" button on a toolbar.

There are three variants to this method, with one, two, and three arguments respectively. The *showUI* and *value* parameters, even if specified, are ignored unless otherwise stated.

When `execCommand()` is invoked, the user agent must follow the following steps:

1. If the given *commandId* maps to an entry in the list below whose "Enabled When" entry has a condition that is currently false, do nothing; abort these steps.
2. Otherwise, execute the "Action" listed below for the given *commandId*.

A document is **ready for editing host commands** if it has a selection that is entirely within an editing host (page 522), or if it has no selection but its caret is inside an editing host (page 522).

The **queryCommandEnabled(*commandId*)** method, when invoked, must return true if the condition listed below under "Enabled When" for the given *commandId* is true, and false otherwise.

The **queryCommandIndeterm(*commandId*)** method, when invoked, must return true if the condition listed below under "Indeterminate When" for the given *commandId* is true, and false otherwise.

The **queryCommandState(*commandId*)** method, when invoked, must return the value expressed below under "State" for the given *commandId*.

The **queryCommandSupported(*commandId*)** method, when invoked, must return true if the given *commandId* is in the list below, and false otherwise.

The **queryCommandValue(*commandId*)** method, when invoked, must return the value expressed below under "Value" for the given *commandId*.

The possible values for *commandId*, and their corresponding meanings, are as follows. These values must be compared to the argument in an ASCII case-insensitive (page 31) manner.

bold

Action: The user agent must act as if the user had requested that the selection be wrapped in the semantics (page 524) of the *b* element (or, again, unwrapped, or have that semantic inserted or removed, as defined by the UA).

Enabled When: The document is ready for editing host commands (page 540).

Indeterminate When: Never.

State: True if the selection, or the caret, if there is no selection, is, or is contained within, a *b* element. False otherwise.

Value: The string "true" if the expression given for the "State" above is true, the string "false" otherwise.

createLink

Action: The user agent must act as if the user had requested that the selection be wrapped in the semantics (page 524) of the *a* element (or, again, unwrapped, or have that semantic inserted or removed, as defined by the UA). If the user agent creates an *a* element or modifies an existing *a* element, then if the *showUI* argument is present and has the value false, then the value of the *value* argument must be used as the URL (page 52) of the link. Otherwise, the user agent should prompt the user for the URL of the link.

Enabled When: The document is ready for editing host commands (page 540).

Indeterminate When: Never.

State: Always false.

Value: Always the string "false".

delete

Action: The user agent must act as if the user had performed a backspace operation (page 523).

Enabled When: The document is ready for editing host commands (page 540).

Indeterminate When: Never.

State: Always false.

Value: Always the string "false".

formatBlock

Action: The user agent must run the following steps:

1. If the *value* argument wasn't specified, abort these steps without doing anything.
2. If the *value* argument has a leading U+003C LESS-THAN SIGN character ('<') and a trailing U+003E GREATER-THAN SIGN character ('>'), then remove the first and last characters from *value*.
3. If *value* is (now) an ASCII case-insensitive (page 31) match for the tag name of an element defined by this specification that is defined to be a prose element but not a phrasing element, then, for every position in the selection, take the furthest flow content (page 93) ancestor element of that position that contains only phrasing content (page 94), and, if that element is editable (page 521), and has a content model that allows it to contain prose content other than phrasing content (page 94), and has a parent element whose content model allows that parent to contain any prose content, rename the element (as if the `Element.renameNode()` method had been used) to *value*, using the HTML namespace.

If there is no selection, then, where in the description above refers to the selection, the user agent must act as if the selection was an empty range (with just one position) at the caret position.

Enabled When: The document is ready for editing host commands (page 540).

Indeterminate When: Never.

State: Always false.

Value: Always the string "false".

forwardDelete

Action: The user agent must act as if the user had performed a forward delete operation (page 523).

Enabled When: The document is ready for editing host commands (page 540).

Indeterminate When: Never.

State: Always false.

Value: Always the string "false".

insertImage

Action: The user agent must act as if the user had requested that the selection be wrapped in the semantics (page 524) of the *img* element (or, again, unwrapped, or have that semantic inserted or removed, as defined by the UA). If the user agent creates an *img* element or modifies an existing *img* element, then if the *showUI* argument is present and has the value *false*, then the value of the *value* argument must be used as the URL (page 52) of the image. Otherwise, the user agent should prompt the user for the URL of the image.

Enabled When: The document is ready for editing host commands (page 540).

Indeterminate When: Never.

State: Always false.

Value: Always the string "false".

insertHTML

Action: The user agent must run the following steps:

1. If the document is an XML document, then throw an `INVALID_ACCESS_ERR` exception and abort these steps.
2. If the *value* argument wasn't specified, abort these steps without doing anything.
3. If there is a selection, act as if the user had requested that the selection be deleted (page 523).
4. Invoke the HTML fragment parsing algorithm (page 661) with an arbitrary orphan body element as the *context* (page 262) element and with the *value* argument as *input* (page 318).
5. Insert the nodes returned by the previous step into the document at the location of the caret.

Enabled When: The document is ready for editing host commands (page 540).

Indeterminate When: Never.

State: Always false.

Value: Always the string "false".

insertLineBreak

Action: The user agent must act as if the user had requested a line separator (page 523).

Enabled When: The document is ready for editing host commands (page 540).

Indeterminate When: Never.

State: Always false.

Value: Always the string "false".

insertOrderedList

Action: The user agent must act as if the user had requested that the selection be wrapped in the semantics (page 524) of the ol element (or unwrapped, or, if there is no selection, have that semantic inserted or removed — the exact behavior is UA-defined).

Enabled When: The document is ready for editing host commands (page 540).

Indeterminate When: Never.

State: Always false.

Value: Always the string "false".

insertUnorderedList

Action: The user agent must act as if the user had requested that the selection be wrapped in the semantics (page 524) of the ul element (or unwrapped, or, if there is no selection, have that semantic inserted or removed — the exact behavior is UA-defined).

Enabled When: The document is ready for editing host commands (page 540).

Indeterminate When: Never.

State: Always false.

Value: Always the string "false".

insertParagraph

Action: The user agent must act as if the user had performed a break block (page 523) editing action.

Enabled When: The document is ready for editing host commands (page 540).

Indeterminate When: Never.

State: Always false.

Value: Always the string "false".

insertText

Action: The user agent must act as if the user had inserted text (page 522) corresponding to the *value* parameter.

Enabled When: The document is ready for editing host commands (page 540).

Indeterminate When: Never.

State: Always false.

Value: Always the string "false".

italic

Action: The user agent must act as if the user had requested that the selection be wrapped in the semantics (page 524) of the *i* element (or, again, unwrapped, or have that semantic inserted or removed, as defined by the UA).

Enabled When: The document is ready for editing host commands (page 540).

Indeterminate When: Never.

State: True if the selection, or the caret, if there is no selection, is, or is contained within, a *i* element. False otherwise.

Value: The string "true" if the expression given for the "State" above is true, the string "false" otherwise.

redo

Action: The user agent must move forward one step (page 539) in its undo transaction history (page 536), restoring the associated state. If the undo position (page 537) is at the end of the undo transaction history (page 536), the user agent must do nothing. See the undo history (page 536).

Enabled When: The undo position (page 537) is not at the end of the undo transaction history (page 536).

Indeterminate When: Never.

State: Always false.

Value: Always the string "false".

selectAll

Action: The user agent must change the selection so that all the content in the currently focused editing host (page 522) is selected. If no editing host (page 522) is focused, then the content of the entire document must be selected.

Enabled When: Always.

Indeterminate When: Never.

State: Always false.

Value: Always the string "false".

subscript

Action: The user agent must act as if the user had requested that the selection be wrapped in the semantics (page 524) of the *sub* element (or, again, unwrapped, or have that semantic inserted or removed, as defined by the UA).

Enabled When: The document is ready for editing host commands (page 540).

Indeterminate When: Never.

State: True if the selection, or the caret, if there is no selection, is, or is contained within, a *sub* element. False otherwise.

Value: The string "true" if the expression given for the "State" above is true, the string "false" otherwise.

superscript

Action: The user agent must act as if the user had requested that the selection be wrapped in the semantics (page 524) of the *sup* element (or unwrapped, or, if there is no selection, have that semantic inserted or removed — the exact behavior is UA-defined).

Enabled When: The document is ready for editing host commands (page 540).

Indeterminate When: Never.

State: True if the selection, or the caret, if there is no selection, is, or is contained within, a *sup* element. False otherwise.

Value: The string "true" if the expression given for the "State" above is true, the string "false" otherwise.

undo

Action: The user agent must move back one step (page 538) in its undo transaction history (page 536), restoring the associated state. If the undo position (page 537) is at the start of the undo transaction history (page 536), the user agent must do nothing. See the undo history (page 536).

Enabled When: The undo position (page 537) is not at the start of the undo transaction history (page 536).

Indeterminate When: Never.

State: Always false.

Value: Always the string "false".

unlink

Action: The user agent must remove all elements that have href attributes and that are partially or completely included in the current selection.

Enabled When: The document has a selection that is entirely within an editing host (page 522) and that contains (either partially or completely) at least one element that has an href attribute.

Indeterminate When: Never.

State: Always false.

Value: Always the string "false".

unselect

Action: The user agent must change the selection so that nothing is selected.

Enabled When: Always.

Indeterminate When: Never.

State: Always false.

Value: Always the string "false".

vendorID-customCommandID

Action: User agents may implement vendor-specific extensions to this API. Vendor-specific extensions to the list of commands should use the syntax *vendorID*-*customCommandID* so as to prevent clashes between extensions from different vendors and future additions to this specification.

Enabled When: UA-defined.

Indeterminate When: UA-defined.

State: UA-defined.

Value: UA-defined.

Anything else

Action: User agents must do nothing.

Enabled When: Never.

Indeterminate When: Never.

State: Always false.

Value: Always the string "false".

7 Communication

7.1 Event definitions

Messages in server-sent events (page 547), Web sockets (page 553), cross-document messaging (page 564), and channel messaging (page 567) use the **message** event.

The following interface is defined for this event:

```
interface MessageEvent : Event {  
    readonly attribute DOMString data;  
    readonly attribute DOMString origin;  
    readonly attribute DOMString lastEventId;  
    readonly attribute Window source;  
    readonly attribute MessagePort messagePort;  
    void initMessageEvent(in DOMString typeArg, in boolean canBubbleArg,  
    in boolean cancelableArg, in DOMString dataArg, in DOMString originArg,  
    in DOMString lastEventIdArg, in Window sourceArg, in MessagePort  
    messagePortArg);  
    void initMessageEventNS(in DOMString namespaceURI, in DOMString  
    typeArg, in boolean canBubbleArg, in boolean cancelableArg, in DOMString  
    dataArg, in DOMString originArg, in DOMString lastEventIdArg, in Window  
    sourceArg, in MessagePort messagePortArg);  
};
```

The **initMessageEvent()** and **initMessageEventNS()** methods must initialize the event in a manner analogous to the similarly-named methods in the DOM3 Events interfaces.
[DOM3EVENTS]

The **data** attribute represents the message being sent.

The **origin** attribute represents, in server-sent events (page 547) and cross-document messaging (page 564), the origin (page 423) of the document that sent the message (typically the scheme, hostname, and port of the document, but not its path or fragment identifier).

The **lastEventId** attribute represents, in server-sent events (page 547), the last event ID string of the event source.

The **source** attribute represents, in cross-document messaging (page 564), the Window from which the message came.

The **messagePort** attribute represents, in cross-document messaging (page 564) and channel messaging (page 567) the MessagePort being sent, if any.

Unless otherwise specified, when the user agent creates and dispatches a message event in the algorithms described in the following sections, the **lastEventId** attribute must be the empty string, the **origin** attribute must be the empty string, the **source** attribute must be null, and the **messagePort** attribute must be null.

Tasks (page 429) in server-sent events (page 547) and Web Sockets (page 553) use their own task sources (page 429), but the task source (page 429) for the tasks (page 429) in

cross-document messaging (page 564) and channel messaging (page 567) is the **posted message task source**.

7.2 Server-sent events

This section describes a mechanism for allowing servers to dispatch DOM events into documents that expect it. The `eventsource` element provides a simple interface to this mechanism.

7.2.1 The `RemoteEventTarget` interface

Any object that implements the `EventTarget` interface must also implement the `RemoteEventTarget` interface.

```
interface RemoteEventTarget {
    void addEventSource(in DOMString src);
    void removeEventSource(in DOMString src);
};
```

When the `addEventSource(src)` method is invoked, the user agent must resolve (page 55) the URL (page 52) specified in `src`, and if that succeeds, add the resulting absolute URL (page 56) to the list of event sources (page 547) for that object. The same URL can be registered multiple times. If the URL fails to resolve, then the user agent must raise a `SYNTAX_ERR` exception.

When the `removeEventSource(src)` method is invoked, the user agent must resolve (page 55) the URL (page 52) specified in `src`, and if that succeeds, remove the resulting absolute URL (page 56) from the list of event sources (page 547) for that object. If the same URL has been registered multiple times, removing it must remove only one instance of that URL for each invocation of the `removeEventSource()` method. If the URL fails to resolve, the user agent does nothing.

7.2.2 Connecting to an event stream

Each object implementing the `EventTarget` and `RemoteEventTarget` interfaces has a **list of event sources** that are registered for that object.

When a new absolute URL (page 56) is added to this list, the user agent should queue a task (page 429) to run the following steps with the new absolute URL (page 56):

1. If the entry for the new absolute URL (page 56) has been removed from the list, then abort these steps.
2. Fetch (page 59) the resource identified by that absolute URL (page 56).

As data is received, the tasks (page 429) queued by the networking task source (page 430) to handle the data must consist of following the rules given in the following sections.

When an event source is removed from the list of event sources for an object, if that resource is still being fetched, then the relevant connection must be closed.

Since connections established to remote servers for such resources are expected to be long-lived, UAs should ensure that appropriate buffering is used. In particular, while line buffering may be safe if lines are defined to end with a single U+000A LINE FEED character, block buffering or line buffering with different expected line endings can cause delays in event dispatch.

Each event source in the list must have associated with it the following:

- The **reconnection time**, in milliseconds. This must initially be a user-agent-defined value, probably in the region of a few seconds.
- The **last event ID string**. This must initially be the empty string.

In general, the semantics of the transport protocol specified by the URLs for the event sources must be followed, including HTTP caching rules.

For HTTP connections, the Accept header may be included; if included, it must contain only formats of event framing that are supported by the user agent (one of which must be text/event-stream, as described below).

Other formats of event framing may also be supported in addition to text/event-stream, but this specification does not define how they are to be parsed or processed.

Note: Such formats could include systems like SMS-push; for example servers could use Accept headers and HTTP redirects to an SMS-push mechanism as a kind of protocol negotiation to reduce network load in GSM environments.

User agents should use the Cache-Control: no-cache header in requests to bypass any caches for requests of event sources.

If the event source's last event ID string is not the empty string, then a Last-Event-ID HTTP header must be included with the request, whose value is the value of the event source's last event ID string.

For connections to domains other than the document's domain (page 427), the semantics of the Access-Control HTTP header must be followed. [ACCESSCONTROL]

HTTP 200 OK responses with a Content-Type (page 60) header specifying the type text/event-stream that are either from the document's domain (page 427) or explicitly allowed by the Access-Control HTTP headers must be processed line by line as described below (page 550).

For the purposes of such successfully opened event streams only, user agents should ignore HTTP cache headers, and instead assume that the resource indicates that it does not wish to be cached.

If such a resource completes loading (i.e. the entire HTTP response body is received or the connection itself closes), the user agent should request the event source resource again after a delay equal to the reconnection time of the event source.

HTTP 200 OK responses that have a Content-Type (page 60) other than text/event-stream (or some other supported type), and HTTP responses whose Access-Control headers indicate that the resource are not to be used, must be ignored and must prevent the user agent from refetching the resource for that event source.

HTTP 201 Created, 202 Accepted, 203 Non-Authoritative Information, and 206 Partial Content responses must be treated like HTTP 200 OK responses for the purposes of reopening event source resources. They are, however, likely to indicate an error has occurred somewhere and may cause the user agent to emit a warning.

HTTP 204 No Content, and 205 Reset Content responses must be treated as if they were 200 OK responses with the right MIME type but no content, and should therefore cause the user agent to refetch the resource after a delay equal to the reconnection time of the event source.

HTTP 300 Multiple Choices responses should be handled automatically if possible (treating the responses as if they were 302 Found responses pointing to the appropriate resource), and otherwise must be treated as HTTP 404 responses.

HTTP 301 Moved Permanently responses must cause the user agent to reconnect using the new server specified URL instead of the previously specified URL for all subsequent requests for this event source. (It doesn't affect other event sources with the same URL unless they also receive 301 responses, and it doesn't affect future sessions, e.g. if the page is reloaded.)

HTTP 302 Found, 303 See Other, and 307 Temporary Redirect responses must cause the user agent to connect to the new server-specified URL, but if the user agent needs to again request the resource at a later point, it must return to the previously specified URL for this event source.

HTTP 304 Not Modified responses should be handled like HTTP 200 OK responses, with the content coming from the user agent cache. A new request should then be made after a delay equal to the reconnection time of the event source.

HTTP 305 Use Proxy, HTTP 401 Unauthorized, and 407 Proxy Authentication Required should be treated transparently as for any other subresource.

Any other HTTP response code not listed here should cause the user agent to stop trying to process this event source.

DNS errors must be considered fatal, and cause the user agent to not open any connection for that event source.

For non-HTTP protocols, UAs should act in equivalent ways.

7.2.3 Parsing an event stream

This event stream format's MIME type is text/event-stream.

The event stream format is (in pseudo-BNF):

```
<stream>      ::= <bom>? <event>*
<event>       ::= [ <comment> | <field> ]* <newline>
<comment>     ::= <colon> <any-char>* <newline>
<field>       ::= <name-char>+ [ <colon> <space>? <any-char>* ]?
```

```

<newline>

# characters:
<bom>          ::= a single U+FEFF BYTE ORDER MARK character
<space>         ::= a single U+0020 SPACE character (' ')
<newline>        ::= a U+000D CARRIAGE RETURN character
                      followed by a U+000A LINE FEED character
                      | a single U+000D CARRIAGE RETURN character
                      | a single U+000A LINE FEED character
                      | the end of the file
<colon>         ::= a single U+003A COLON character (':')
<name-char>      ::= a single Unicode character other than
                      U+003A COLON, U+000D CARRIAGE RETURN and U+000A LINE
                      FEED
<any-char>       ::= a single Unicode character other than
                      U+000D CARRIAGE RETURN and U+000A LINE FEED

```

Event streams in this format must always be encoded as UTF-8.

Lines must be separated by either a U+000D CARRIAGE RETURN U+000A LINE FEED (CRLF) character pair, a single U+000A LINE FEED (LF) character, or a single U+000D CARRIAGE RETURN (CR) character.

7.2.4 Interpreting an event stream

Bytes or sequences of bytes that are not valid UTF-8 sequences must be interpreted as the U+FFFD REPLACEMENT CHARACTER.

One leading U+FEFF BYTE ORDER MARK character must be ignored if any are present.

The stream must then be parsed by reading everything line by line, with a U+000D CARRIAGE RETURN U+000A LINE FEED (CRLF) character pair, a single U+000A LINE FEED (LF) character, a single U+000D CARRIAGE RETURN (CR) character, and the end of the file being the four ways in which a line can end.

When a stream is parsed, a *data* buffer and an *event name* buffer must be associated with it. They must be initialized to the empty string

Lines must be processed, in the order they are received, as follows:

- ↪ **If the line is empty (a blank line)**

Dispatch the event (page 551), as defined below.

- ↪ **If the line starts with a U+003A COLON character (':')**

Ignore the line.

- ↪ **If the line contains a U+003A COLON character (':') character**

Collect the characters on the line before the first U+003A COLON character (':'), and let *field* be that string.

Collect the characters on the line after the first U+003A COLON character (':'), and let *value* be that string. If *value* starts with a single U+0020 SPACE character, remove it from *value*.

Process the field (page 551) using the steps described below, using *field* as the field name and *value* as the field value.

↪ **Otherwise, the string is not empty but does not contain a U+003A COLON character (':') character**

Process the field (page 551) using the steps described below, using the whole line as the field name, and the empty string as the field value.

Once the end of the file is reached, the user agent must dispatch the event (page 551) one final time, as defined below.

The steps to **process the field** given a field name and a field value depend on the field name, as given in the following list. Field names must be compared literally, with no case folding performed.

↪ **If the field name is "event"**

Set the *event name* buffer to the field value.

↪ **If the field name is "data"**

If the *data* buffer is not the empty string, then append a single U+000A LINE FEED character to the *data* buffer. Append the field value to the *data* buffer.

↪ **If the field name is "id"**

Set the event stream's last event ID (page 548) to the field value.

↪ **If the field name is "retry"**

If the field value consists of only characters in the range U+0030 DIGIT ZERO ('0') U+0039 DIGIT NINE ('9'), then interpret the field value as an integer in base ten, and set the event stream's reconnection time (page 548) to that integer. Otherwise, ignore the field.

↪ **Otherwise**

The field is ignored.

When the user agent is required to **dispatch the event**, then the user agent must act as follows:

1. If the *data* buffer is an empty string, set the *data* buffer and the *event name* buffer to the empty string and abort these steps.
2. If the *event name* buffer is not the empty string but is also not a valid NCName, set the *data* buffer and the *event name* buffer to the empty string and abort these steps.
3. Otherwise, create an event that uses the MessageEvent interface, with the event name message, which does not bubble, is cancelable, and has no default action. The data attribute must be set to the value of the *data* buffer, the origin attribute must be set to the Unicode serialization (page 425) of the origin (page 423) of the event stream's URL, and the lastEventId attribute must be set to the last event ID string of the event source.
4. If the *event name* buffer has a value other than the empty string, change the type of the newly created event to equal the value of the *event name* buffer.
5. Set the *data* buffer and the *event name* buffer to the empty string.

- Queue a task (page 429) to dispatch the newly created event at the `RemoteEventTarget` object to which the event stream is registered. The task source (page 429) for this task (page 429) is the **remote event task source**.

Note: If an event doesn't have an "id" field, but an earlier event did set the event source's last event ID string, then the event's `LastEventId` field will be set to the value of whatever the last seen "id" field was.

The following event stream, once followed by a blank line:

```
data: YH00
data: -2
data: 10
```

...would cause an event message with the interface `MessageEvent` to be dispatched on the `eventsource` element, whose `data` attribute would contain the string `YH00\n-2\n10` (where `\n` represents a newline).

This could be used as follows:

```
<eventsource src="http://stocks.example.com/ticker.php"
            onmessage="var data = event.data.split('\n');
            updateStocks(data[0], data[1], data[2]);">
```

...where `updateStocks()` is a function defined as:

```
function updateStocks(symbol, delta, value) { ... }
```

...or some such.

The following stream contains four blocks. The first block has just a comment, and will fire nothing. The second block has two fields with names "data" and "id" respectively; an event will be fired for this block, with the data "first event", and will then set the last event ID to "1" so that if the connection died between this block and the next, the server would be sent a `Last-Event-ID` header with the value "1". The third block fires an event with data "second event", and also has an "id" field, this time with no value, which resets the last event ID to the empty string (meaning no `Last-Event-ID` header will now be sent in the event of a reconnection being attempted). Finally the last block just fires an event with the data "third event". Note that the last block doesn't have to end with a blank line, the end of the stream is enough to trigger the dispatch of the last event.

```
: test stream

data: first event
id: 1

data: second event
id

data: third event
```

The following stream fires just one event:

```
data
```

```
data  
data
```

```
data:
```

The first and last blocks do nothing, since they do not contain any actual data (the `data` buffer remains at the empty string, and so nothing gets dispatched). The middle block fires an event with the data set to a single newline character.

The following stream fires two identical events:

```
data:test
```

```
data: test
```

This is because the space after the colon is ignored if present.

7.2.5 Notes

Legacy proxy servers are known to, in certain cases, drop HTTP connections after a short timeout. To protect against such proxy servers, authors can include a comment line (one starting with a '`:`' character) every 15 seconds or so.

Authors wishing to relate event source connections to each other or to specific documents previously served might find that relying on IP addresses doesn't work, as individual clients can have multiple IP addresses (due to having multiple proxy servers) and individual IP addresses can have multiple clients (due to sharing a proxy server). It is better to include a unique identifier in the document when it is served and then pass that identifier as part of the URL in the `src` attribute of the `eventsource` element.

Implementations that support HTTP's per-server connection limitation might run into trouble when opening multiple pages from a site if each page has an `eventsource` to the same domain. Authors can avoid this using the relatively complex mechanism of using unique domain names per connection, or by allowing the user to enable or disable the `eventsource` functionality on a per-page basis.

7.3 Web sockets

To enable Web applications to maintain bidirectional communications with their originating server, this specification introduces the WebSocket interface.

Note: This interface does not allow for raw access to the underlying network. For example, this interface could not be used to implement an IRC client without proxying messages through a custom server.

7.3.1 Introduction

This section is non-normative.

- ** An introduction to the client-side and server-side of using the direct connection APIs.

7.3.2 The WebSocket interface

```
[Constructor(in DOMString url)]
interface WebSocket {
    readonly attribute DOMString URL;

    // ready state
    const unsigned short CONNECTING = 0;
    const unsigned short OPEN = 1;
    const unsigned short CLOSED = 2;
    readonly attribute long readyState;

    // networking
        attribute EventListener onopen;
        attribute EventListener onmessage;
        attribute EventListener onclosed;
    void postMessage(in DOMString data);
    void disconnect();
};
```

WebSocket objects must also implement the EventTarget interface. [DOM3EVENTS]

The **WebSocket(url)** constructor takes one argument, *url*, which specifies the URL (page 52) to which to connect. When a WebSocket object is created, the UA must parse (page 53) this argument and verify that the URL parses without failure and has a <scheme> (page 53) component whose value is either "ws" or "wss", when compared in an ASCII case-insensitive (page 31) manner. If it does, it has, and it is, then the user agent must asynchronously establish a Web Socket connection (page 555) to *url*. Otherwise, the constructor must raise a SYNTAX_ERR exception.

The **URL** attribute must return the value that was passed to the constructor.

The **readyState** attribute represents the state of the connection. It can have the following values:

CONNECTING (numeric value 0)

The connection has not yet been established.

OPEN (numeric value 1)

The Web Socket connection is established (page 560) and communication is possible.

CLOSED (numeric value 2)

The connection has been closed or could not be opened.

When the object is created its readyState must be set to CONNECTING (0).

The **postMessage(data)** method transmits data using the connection. If the connection is not established (readyState is not OPEN), it must raise an INVALID_STATE_ERR exception. If the

connection is established, then the user agent must send *data* using the Web Socket (page 561).

The **disconnect()** method must close the Web Socket connection (page 564) or connection attempt, if any. If the connection is already closed, it must do nothing. Closing the connection causes a **close** event to be fired and the **readyState** attribute's value to change, as described below (page 564).

7.3.3 WebSocket Events

The **open** event is fired when the Web Socket connection is established (page 560).

The **close** event is fired when the connection is closed (whether by the author, calling the **disconnect()** method, or by the server, or by a network error).

Note: *No information regarding why the connection was closed is passed to the application in this version of this specification.*

The message event is fired when when data is received for a connection.

The following are the event handler DOM attributes (page 432) that must be supported by objects implementing the WebSocket interface:

onopen

Must be invoked whenever an open event is targeted at or bubbles through the WebSocket object.

onmessage

Must be invoked whenever a message event is targeted at or bubbles through the WebSocket object.

onclosed

Must be invoked whenever an closed event is targeted at or bubbles through the WebSocket object.

7.3.4 The Web Socket protocol

The task source (page 429) for all tasks (page 429) queued (page 429) by algorithms in this section and its subsections is the **Web Socket task source**.

7.3.4.1 Client-side requirements

This section only applies to user agents.

7.3.4.1.1 Handshake

When the user agent is to **establish a Web Socket connection** to *url*, it must run the following steps, in the background (without blocking scripts or anything like that):

1. Resolve (page 55) the URL (page 52) *url*.
2. If the <scheme> (page 53) component of the resulting absolute URL (page 56) is "ws", set *secure* to false; otherwise, the <scheme> (page 53) component is "wss", set *secure* to true.
3. Let *host* be the value of the <host> (page 53) component in the resulting absolute URL (page 56).
4. If the resulting absolute URL (page 56) has a <port> (page 53) component, then let *port* be that component's value; otherwise, if *secure* is false, let *port* be 81, otherwise let *port* be 815.
5. Let *resource name* be the value of the <path> (page 54) component (which might be empty) in the resulting absolute URL (page 56).
6. If *resource name* is the empty string, set it to a single character U+002F SOLIDUS (/).
7. If the resulting absolute URL (page 56) has a <query> (page 54) component, then append a single 003F QUESTION MARK (?) character to *resource name*, followed by the value of the <query> (page 54) component.
8. If the user agent is configured to use a proxy to connect to port *port*, then connect to that proxy and ask it to open a TCP/IP connection to the host given by *host* and the port given by *port*.

For example, if the user agent uses an HTTP proxy, then if it was to try to connect to port 80 on server example.com, it might send the following lines to the proxy server:

```
CONNECT example.com HTTP/1.1
```

If there was a password, the connection might look like:

```
CONNECT example.com HTTP/1.1
Proxy-authorization: Basic ZWRuYW1vZGU6bm9jYXBlcyE=
```

Otherwise, if the user agent is not configured to use a proxy, then open a TCP/IP connection to the host given by *host* and the port given by *port*.

9. If the connection could not be opened, then fail the Web Socket connection (page 560) and abort these steps.
10. If *secure* is true, perform a TLS handshake over the connection. If this fails (e.g. the server's certificate could not be verified), then fail the Web Socket connection (page 560) and abort these steps. Otherwise, all further communication on this channel must run through the encrypted tunnel. [RFC2246]
11. Send the following bytes to the remote side (the server):

```
47 45 54 20
```

Send the *resource name* value, encoded as US-ASCII.

Send the following bytes:

```
20 48 54 54 50 2f 31 2e 31 0d 0a 55 70 67 72 61  
64 65 3a 20 57 65 62 53 6f 63 6b 65 74 0d 0a 43  
6f 6e 6e 65 63 74 69 6f 6e 3a 20 55 70 67 72 61  
64 65 0d 0a
```

Note: The string "GET ", the path, " HTTP/1.1", CRLF, the string "Upgrade: WebSocket", CRLF, and the string "Connection: Upgrade", CRLF.

12. Send the following bytes:

```
48 6f 73 74 3a 20
```

Send the *host* value, encoded as US-ASCII, if it represents a host name (and not an IP address).

Send the following bytes:

```
0d 0a
```

Note: The string "Host: ", the host, and CRLF.

13. Send the following bytes:

```
4f 72 69 67 69 6e 3a 20
```

Send the ASCII serialization (page 426) of the origin (page 423) of the script that invoked the `WebSocket()` constructor.

Send the following bytes:

```
0d 0a
```

Note: The string "Origin: ", the origin, and CRLF.

14. If the client has any authentication information or cookies that would be relevant to a resource with a URL (page 52) that has a scheme of `http` if `secure` is false and `https` if `secure` is true and is otherwise identical to `url`, then HTTP headers that would be appropriate for that information should be sent at this point. [RFC2616] [RFC2109] [RFC2965]

Each header must be on a line of its own (each ending with a CR LF sequence). For the purposes of this step, each header must not be split into multiple lines (despite HTTP otherwise allowing this with continuation lines).

For example, if the server had a username and password that applied to that URL, it could send:

```
Authorization: Basic d2FsbGU6ZXZl
```

15. Send the following bytes:

```
0d 0a
```

Note: Just a CRLF (a blank line).

16. Read the first 85 bytes from the server. If the connection closes before 85 bytes are received, or if the first 85 bytes aren't exactly equal to the following bytes, then fail the Web Socket connection (page 560) and abort these steps.

```
48 54 54 50 2f 31 2e 31 20 31 30 31 20 57 65 62  
20 53 6f 63 6b 65 74 20 50 72 6f 74 6f 63 6f 6c  
20 48 61 6e 64 73 68 61 6b 65 0d 0a 55 70 67 72  
61 64 65 3a 20 57 65 62 53 6f 63 6b 65 74 0d 0a  
43 6f 6e 6e 65 63 74 69 6f 6e 3a 20 55 70 67 72  
61 64 65 0d 0a
```

**Note: The string "HTTP/1.1 101 Web Socket Protocol Handshake",
CRLF, the string "Upgrade: WebSocket", CRLF, the string
"Connection: Upgrade", CRLF.**

**

What if the response is a 401 asking for credentials?

17. Let *headers* be a list of name-value pairs, initially empty.
18. *Header*: Let *name* and *value* be empty byte arrays.
19. Read a byte from the server.

If the connection closes before this byte is received, then fail the Web Socket connection (page 560) and abort these steps.

Otherwise, handle the byte as described in the appropriate entry below:

↪ **If the byte is 0x0d (ASCII CR)**

If the *name* byte array is empty, then jump to the headers processing (page 559) step. Otherwise, fail the Web Socket connection (page 560) and abort these steps.

↪ **If the byte is 0x0a (ASCII LF)**

Fail the Web Socket connection (page 560) and abort these steps.

↪ **If the byte is 0x3a (ASCII ":")**

Move on to the next step.

↪ **If the byte is in the range 0x41 .. 0x5a (ASCII "A" .. "Z")**

Append a byte whose value is the byte's value plus 0x20 to the *name* byte array and redo this step for the next byte.

↪ **Otherwise**

Append the byte to the *name* byte array and redo this step for the next byte.

Note: This reads a header name, terminated by a colon, converting upper-case ASCII letters to lowercase, and aborting if a stray CR or LF is found.

20. Read a byte from the server.

If the connection closes before this byte is received, then fail the Web Socket connection (page 560) and abort these steps.

Otherwise, handle the byte as described in the appropriate entry below:

↪ **If the byte is 0x20 (ASCII space)**

Ignore the byte and move on to the next step.

↪ **Otherwise**

Treat the byte as described by the list in the next step, then move on to that next step for real.

Note: This skips past a space character after the colon, if necessary.

21. Read a byte from the server.

If the connection closes before this byte is received, then fail the Web Socket connection (page 560) and abort these steps.

Otherwise, handle the byte as described in the appropriate entry below:

↪ **If the byte is 0x0d (ASCII CR)**

Move on to the next step.

↪ **If the byte is 0x0a (ASCII LF)**

Fail the Web Socket connection (page 560) and abort these steps.

↪ **Otherwise**

Append the byte to the *name* byte array and redo this step for the next byte.

Note: This reads a header value, terminated by a CRLF.

22. Read a byte from the server.

If the connection closes before this byte is received, or if the byte is not a 0x0a byte (ASCII LF), then fail the Web Socket connection (page 560) and abort these steps.

Note: This skips past the LF byte of the CRLF after the header.

23. Append an entry to the *headers* list that has the name given by the string obtained by interpreting the *name* byte array as a UTF-8 byte stream and the value given by the string obtained by interpreting the *value* byte array as a UTF-8 byte stream.
24. Return to the header (page 558) step above.
25. *Headers processing*: If there is not exactly one entry in the *headers* list whose name is "websocket-origin", or if there is not exactly one entry in the *headers* list whose name is "websocket-location", or if there are any entries in the *headers* list whose names are the empty string, then fail the Web Socket connection (page 560) and abort these steps.
26. Handle each entry in the *headers* list as follows:

↪ **If the entry's name is "websocket-origin"**

If the value is not exactly equal to the ASCII serialization (page 426) of the origin (page 423) of the script that invoked the `WebSocket()` constructor, then fail the Web Socket connection (page 560) and abort these steps.

↪ **If the entry's name is "websocket-location"**

If the value is not exactly equal to the absolute URL (page 56) that resulted from the first step (page 556) of the algorithm, then fail the Web Socket connection (page 560) and abort these steps.

↪ **If the entry's name is "set-cookie" or "set-cookie2" or another cookie-related header name**

Handle the cookie as defined by the appropriate spec, except pretend that the resource's URL (page 52) actually has a scheme of `http` if `secure` is false and `https` if `secure` is true and is otherwise identical to `url`. [RFC2109] [RFC2965]

↪ **Any other name**

Ignore it.

27. Change the `readyState` attribute's value to `OPEN` (1).
28. Queue a task (page 429) to fire a simple event (page 436) named `open` at the `WebSocket` object.
29. The **Web Socket connection is established**. Now the user agent must send and receive to and from the connection as described in the next section.

To **fail the Web Socket connection**, the user agent must close the Web Socket connection (page 564), and may report the problem to the user (which would be especially useful for developers). However, user agents must not convey the failure information to the script in a way distinguishable from the Web Socket being closed normally.

7.3.4.1.2 Data framing

Once a Web Socket connection is established (page 560), the user agent must run through the following state machine for the bytes sent by the server.

1. Try to read a byte from the server. Let `frame type` be that byte.

If no byte could be read because the Web Socket connection is closed (page 564), then abort.

2. Handle the `frame type` byte as follows:

If the high-order bit of the `frame type` byte is set (i.e. if `frame type` anded with `0x80` returns `0x80`)

Run these steps. If at any point during these steps a read is attempted but fails because the Web Socket connection is closed (page 564), then abort.

1. Let `length` be zero.
2. `Length`: Read a byte, let `b` be that byte.

3. Let b_V be integer corresponding to the low 7 bits of b (the value you would get by *anding* b with 0x7f).
4. Multiply $length$ by 128, add b_V to that result, and store the final result in $length$.
5. If the high-order bit of b is set (i.e. if b *anded* with 0x80 returns 0x80), then return to the step above labeled length (page 560).
6. Read $length$ bytes.
7. Discard the read bytes.

If the high-order bit of the frame type byte is not set (i.e. if frame type anded with 0x80 returns 0x00)

Run these steps. If at any point during these steps a read is attempted but fails because the Web Socket connection is closed (page 564), then abort.

1. Let *raw data* be an empty byte array.
2. *Data*: Read a byte, let b be that byte.
3. If b is not 0xff, then append b to *raw data* and return to the previous step (labeled *data* (page 561)).
4. Interpret *raw data* as a UTF-8 string, and store that string in *data*.
5. If *frame type* is 0x00, create an event that uses the MessageEvent interface, with the event name message, which does not bubble, is cancelable, has no default action, and whose data attribute is set to *data*, and queue a task (page 429) to dispatch it at the WebSocket object. Otherwise, discard the data.
3. Return to the first step to read the next byte.

If the user agent is faced with content that is too large to be handled appropriately, then it must fail the Web Socket connection (page 560).

Once a Web Socket connection is established (page 560), the user agent must use the following steps to **send data using the Web Socket**:

1. Send a 0x00 byte to the server.
2. Encode *data* using UTF-8 and send the resulting byte stream to the server.
3. Send a 0xff byte to the server.

** People often request the ability to send binary blobs over this API; also, once the other postMessage() methods support it, we should look into allowing name/value pairs, arrays, and numbers using postMessage() instead of just strings and binary data.

7.3.4.2 Server-side requirements

This section only applies to servers.

7.3.4.2.1 Minimal handshake

Note: This section describes the minimal requirements for a server-side implementation of Web Sockets.

Listen on a port for TCP/IP. Upon receiving a connection request, open a connection and send the following bytes back to the client:

```
48 54 54 50 2f 31 2e 31 20 31 30 31 20 57 65 62  
20 53 6f 63 6b 65 74 20 50 72 6f 74 6f 63 6f 6c  
20 48 61 6e 64 73 68 61 6b 65 0d 0a 55 70 67 72  
61 64 65 3a 20 57 65 62 53 6f 63 6b 65 74 0d 0a  
43 6f 6e 6e 65 63 74 69 6f 6e 3a 20 55 70 67 72  
61 64 65 0d 0a
```

Send the string "WebSocket-Origin" followed by a U+003A COLON ":" followed by the ASCII serialization (page 426) of the origin from which the server is willing to accept connections, followed by a CRLF pair (0x0d 0x0a).

For instance:

```
WebSocket-Origin: http://example.com
```

Send the string "WebSocket-Location" followed by a U+003A COLON ":" followed by the URL (page 52) of the Web Socket script, followed by a CRLF pair (0x0d 0x0a).

For instance:

```
WebSocket-Location: ws://example.com:80/demo
```

Send another CRLF pair (0x0d 0x0a).

Read (and discard) data from the client until four bytes 0x0d 0x0a 0x0d 0x0a are read.

If the connection isn't dropped at this point, go to the data framing (page 563) section.

7.3.4.2.2 Handshake details

The previous section ignores the data that is transmitted by the client during the handshake.

The data sent by the client consists of a number of fields separated by CR LF pairs (bytes 0x0d 0x0a).

The first field consists of three tokens separated by space characters (byte 0x20). The middle token is the path being opened. If the server supports multiple paths, then the server should echo the value of this field in the initial handshake, as part of the URL (page 52) given on the WebSocket-Location line (after the appropriate scheme and host).

The remaining fields consist of name-value pairs, with the name part separated from the value part by a colon and a space (bytes 0x3a 0x20). Of these, several are interesting:

Host (bytes 48 6f 73 74)

The value gives the hostname that the client intended to use when opening the Web Socket. It would be of interest in particular to virtual hosting environments, where one server might serve multiple hosts, and might therefore want to return different data.

The right host has to be output as part of the URL (page 52) given on the WebSocket-Location line of the handshake described above, to verify that the server knows that it is really representing that host.

Origin (bytes 4f 72 69 67 69 6e)

The value gives the scheme, hostname, and port (if it's not the default port for the given scheme) of the page that asked the client to open the Web Socket. It would be interesting if the server's operator had deals with operators of other sites, since the server could then decide how to respond (or indeed, whether to respond) based on which site was requesting a connection.

If the server supports connections from more than one origin, then the server should echo the value of this field in the initial handshake, on the WebSocket-Origin line.

Other fields

Other fields can be used, such as "Cookie" or "Authorization", for authentication purposes.

7.3.4.2.3 Data framing

Note: This section only describes how to handle content that this specification allows user agents to send (text). It doesn't handle any arbitrary content in the same way that the requirements on user agents defined earlier handle any content including possible future extensions to the protocols.

The server should run through the following steps to process the bytes sent by the client:

1. Read a byte from the client. Assuming everything is going according to plan, it will be a 0x00 byte. Behaviour for the server is undefined if the byte is not 0x00.
2. Let *raw data* be an empty byte array.
3. *Data*: Read a byte, let *b* be that byte.
4. If *b* is not 0xff, then append *b* to *raw data* and return to the previous step (labeled *data* (page 563)).
5. Interpret *raw data* as a UTF-8 string, and apply whatever server-specific processing should occur for the resulting string.
6. Return to the first step to read the next byte.

The server should run through the followin steps to send strings to the client:

1. Send a 0x00 byte to the client to indicate the start of a string.

2. Encode *data* using UTF-8 and send the resulting byte stream to the client.
3. Send a 0xff byte to the client to indicate the end of the message.

7.3.4.3 Closing the connection

To **close the Web Socket connection**, either the user agent or the server closes the TCP/IP connection. There is no closing handshake. Whether the user agent or the server closes the connection, it is said that the **Web Socket connection is closed**.

Servers may close the Web Socket connection (page 564) whenever desired.

User agents should not close the Web Socket connection (page 564) arbitrarily.

When the Web Socket connection is closed (page 564), the readyState attribute's value must be changed to CLOSED (2), and the user agent must queue a task (page 429) to fire a simple event (page 436) named `close` at the `WebSocket` object.

7.4 Cross-document messaging

Web browsers, for security and privacy reasons, prevent documents in different domains from affecting each other; that is, cross-site scripting is disallowed.

While this is an important security feature, it prevents pages from different domains from communicating even when those pages are not hostile. This section introduces a messaging system that allows documents to communicate with each other regardless of their source domain, in a way designed to not enable cross-site scripting attacks.

7.4.1 Introduction

This section is non-normative.

For example, if document A contains an `iframe` element that contains document B, and script in document A calls `postMessage()` on the `Window` object of document B, then a message event will be fired on that object, marked as originating from the `Window` of document A. The script in document A might look like:

```
var o = document.getElementsByTagName('iframe')[0];
o.contentWindow.postMessage('Hello world', 'http://b.example.org/');
```

To register an event handler for incoming events, the script would use `addEventListener()` (or similar mechanisms). For example, the script in document B might look like:

```
window.addEventListener('message', receiver, false);
function receiver(e) {
  if (e.origin == 'http://example.com') {
    if (e.data == 'Hello world') {
      e.source.postMessage('Hello', e.origin);
    } else {
      alert(e.data);
    }
  }
}
```

```
    }  
}
```

This script first checks the domain is the expected domain, and then looks at the message, which it either displays to the user, or responds to by sending a message back to the document which sent the message in the first place.

7.4.2 Security

7.4.2.1 Authors

⚠Warning! Use of this API requires extra care to protect users from hostile entities abusing a site for their own purposes.

Authors should check the origin attribute to ensure that messages are only accepted from domains that they expect to receive messages from. Otherwise, bugs in the author's message handling code could be exploited by hostile sites.

Authors should not use the wildcard keyword ("*") in the `targetOrigin` argument in messages that contain any confidential information, as otherwise there is no way to guarantee that the message is only delivered to the recipient to which it was intended.

7.4.2.2 User agents

The integrity of this API is based on the inability for scripts of one origin (page 423) to post arbitrary events (using `dispatchEvent()` or otherwise) to objects in other origins (those that are not the same (page 426)).

Note: Implementors are urged to take extra care in the implementation of this feature. It allows authors to transmit information from one domain to another domain, which is normally disallowed for security reasons. It also requires that UAs be careful to allow access to certain properties but not others.

7.4.3 Posting text

When a script invokes the `postMessage(message, targetOrigin)` method (with only two arguments) on a Window object, the user agent must follow these steps:

1. If the value of the `targetOrigin` argument is not a single U+002A ASTERISK character ("*"), and parsing (page 53) it as a URL (page 52) fails, then throw a `SYNTAX_ERR` exception and abort the overall set of steps.
2. Return from the `postMessage()` method, but asynchronously continue running these steps.
3. If the `targetOrigin` argument has a value other than a single literal U+002A ASTERISK character ("*"), and the active document (page 414) of the browsing context (page 414) of the Window object on which the method was invoked does not have the same origin (page 426) as `targetOrigin`, then abort these steps silently.

4. Create an event that uses the `MessageEvent` interface, with the event name `message`, which does not bubble, is cancelable, and has no default action. The `data` attribute must be set to the value passed as the `message` argument to the `postMessage()` method, the `origin` attribute must be set to the Unicode serialization (page 425) of the origin (page 423) of the script that invoked the method, and the `source` attribute must be set to the `Window` object of the default view (page 414) of the browsing context (page 414) for which the `Document` object with which the script is associated is the active document (page 414).
5. Queue a task (page 429) to dispatch the event created in the previous step at the `Window` object on which the method was invoked. The task source (page 429) for this task (page 429) is the posted message task source (page 547).

7.4.4 Posting message ports

When a script invokes the `postMessage(message, messagePort, targetOrigin)` method (with three arguments) on a `Window` object, the user agent must follow these steps:

1. If the value of the `targetOrigin` argument is not a single U+002A ASTERISK character ("*"), and parsing (page 53) it as a URL (page 52) fails, then throw a `SYNTAX_ERR` exception and abort the overall set of steps.
2. If the `messagePort` argument is null, then act as if the method had just been called with two arguments (page 565), `message` and `targetOrigin`.
3. Try to obtain a *new port* by cloning (page 568) the `messagePort` argument with the `Window` object on which the method was invoked as the owner of the clone. If this returns an exception, then throw that exception and abort these steps.
4. Return from the `postMessage()` method, but asynchronously continue running these steps.
5. If the `targetOrigin` argument has a value other than a single literal U+002A ASTERISK character ("*"), and the active document (page 414) of the browsing context (page 414) of the `Window` object on which the method was invoked does not have the same origin (page 426) as `targetOrigin`, then abort these steps silently.
6. Create an event that uses the `MessageEvent` interface, with the event name `message`, which does not bubble, is cancelable, and has no default action. The `data` attribute must be set to the value passed as the `message` argument to the `postMessage()` method, the `origin` attribute must be set to the Unicode serialization (page 425) of the origin (page 423) of the script that invoked the method, and the `source` attribute must be set to the `Window` object of the default view (page 414) of the browsing context (page 414) for which the `Document` object with which the script is associated is the active document (page 414).
7. Let the `messagePort` attribute of the event be the *new port*.
8. Queue a task (page 429) to dispatch the event created in the previous step at the `Window` object on which the method was invoked. The task source (page 429) for this task (page 429) is the posted message task source (page 547).

Note: These steps, with the exception of the second and third steps and the penultimate step, are identical to those in the previous section.

7.4.5 Posting structured data

- ** People often request the ability to send name/value pairs, arrays, and numbers using `postMessage()` instead of just strings.

7.5 Channel messaging

7.5.1 Introduction

This section is non-normative.

- ** An introduction to the channel and port APIs.

7.5.2 Message channels

```
[Constructor]
interface MessageChannel {
    readonly attribute MessagePort port1;
    readonly attribute MessagePort port2;
};
```

When the `MessageChannel()` constructor is called, it must run the following algorithm:

1. Create a new `MessagePort` object (page 568) owned by the script execution context (page 428), and let `port1` be that object.
2. Create a new `MessagePort` object (page 568) owned by the script execution context (page 428), and let `port2` be that object.
3. Entangle (page 568) the `port1` and `port2` objects.
4. Instantiate a new `MessageChannel` object, and let `channel` be that object.
5. Let the `port1` attribute of the `channel` object be `port1`.
6. Let the `port2` attribute of the `channel` object be `port2`.
7. Return `channel`.

The `port1` and `port2` attributes must return the values they were assigned when the `MessageChannel` object was created.

7.5.3 Message ports

Each channel has two message ports. Data sent through one port is received by the other port, and vice versa.

```
interface MessagePort {
    readonly attribute boolean active;
    void postMessage(in DOMString message);
    void postMessage(in DOMString message, in MessagePort messagePort);
    MessagePort startConversation(in DOMString message);
    void start();
    void close();

    // event handler attributes
    attribute EventListener onmessage;
    attribute EventListener onclose;
};
```

Objects implementing the MessagePort interface must also implement the EventTarget interface.

Each MessagePort object can be entangled with another (a symmetric relationship). Each MessagePort object also has a **port message queue**, initial empty. A port message queue (page 568) can be open or closed, and is initially closed.

When the user agent is to **create a new MessagePort object** owned by a script execution context (page 428) object *owner*, it must instantiate a new MessagePort object, and let its owner be *owner*.

When the user agent is to **entangle** two MessagePort objects, it must run the following steps:

1. If one of the ports is already entangled, then unentangle it and the port that it was entangled with.

Note: *If those two previously entangled ports were the two parts of a MessageChannel object, then that MessageChannel object no longer represents an actual channel: the two ports in that object are no longer entangled.*

2. Associate the two ports to be entangled, so that they form the two parts of a new channel. (There is no MessageChannel object that represents this channel.)

When the user agent is to **clone a port** *original port*, with the clone being owned by *owner*, it must run the following steps, which return either a new MessagePort object or an exception for the caller to raise:

1. If the *original port* is not entangled without another port, then return an INVALID_STATE_ERR exception and abort all these steps.
2. Let the *remote port* be the port with which the *original port* is entangled.

3. Create a new MessagePort object (page 568) owned by *owner*, and let *new port* be that object.
4. Move all the events in the port message queue (page 568) of *original port* to the port message queue (page 568) of *new port*, if any, leaving the *new port*'s port message queue (page 568) in its initial closed state.
5. Entangle (page 568) the *remote port* and *new port* objects. The *original port* object will be unentangled by this process.
6. Return *new port*. It is the clone.

The **active** attribute must return true if the port is entangled, and false otherwise.

The **postMessage()** method, when called on a port *source port*, must cause the user agent to run the following steps:

1. Let *message* be the method's first argument.
2. Let *data port* be the method's second argument, if any.
3. If the *source port* is not entangled with another port, then return and abort these steps.
4. Let *target port* be the port with which *source port* is entangled.
5. Create an event that uses the MessageEvent interface, with the name *message*, which does not bubble, is cancelable, and has no default action.
6. Let the data attribute of the event have the value of *message*, the method's first argument.
7. If the method was called with a second argument *data port* and that argument isn't null, then run the following substeps:
 1. If the *data port* is the *source port* or the *target port*, then throw an INVALID_ACCESS_ERR exception and abort all these steps.
 2. Try to obtain a *new data port* by cloning (page 568) the *data port* with the owner of the *target port* as the owner of the clone. If this returns an exception, then throw that exception and abort these steps.
 3. Let the messagePort attribute of the event be the *new data port*.
8. Return from the method, but continue with these steps.
9. Add the event to the port message queue (page 568) of *target port*.

** People often request the ability to send name/value pairs, arrays, and numbers using postMessage() instead of just strings.

The **startConversation(message)** method is a convenience method that simplifies creating a new MessageChannel and invoking `postMessage()` with one of the new ports. When invoked on a port *source port*, it must run the following steps:

1. Let *message* be the method's first argument.
2. Create a new MessagePort object (page 568) owned by the script execution context (page 428), and let *port1* be that object.
3. If the *source port* is not entangled with another port, then return *port1* and abort these steps.
4. Let *target port* be the port with which *source port* is entangled.
5. Create a new MessagePort object (page 568) owned by the owner of the *target port*, and let *port2* be that object.
6. Entangle (page 568) the *port1* and *port2* objects.
7. Create an event that uses the MessageEvent interface, with the name *message*, which does not bubble, is cancelable, and has no default action.
8. Let the data attribute of the event have the value of *message*, the method's first argument.
9. Let the messagePort attribute of the event be *port2*.
10. Return *port1* from the method, but continue with these steps.
11. Add the event to the port message queue (page 568) of *target port*.

The **start()** method must open its port's port message queue (page 568), if it is not already open.

When a port's port message queue (page 568) is open, contains an event, and its owner is available (page 570), the user agent must queue a task (page 429) in the event loop (page 429) to dispatch the first event in the queue on the MessagePort object, and remove the event from the queue. The task source (page 429) for this task (page 429) is the posted message task source (page 547).

A MessagePort's owner is **available** if the MessagePort is owned by an object other than a Window object, or if it is owned by a Window object and the Document that was the active document (page 414) in that browsing context (page 414) when the MessagePort was created is fully active (page 415). If that Document is discarded before the port's owner becomes available, then the events are lost.

The **close()** method, when called on a port *local port* that is entangled with another port, must cause the user agents to run the following steps:

1. Unentangle the two ports.
2. Queue a task (page 429) to fire a simple event (page 436) called `close` at the port on which the method was called.

3. Queue a task (page 429) to fire a simple event (page 436) called `close` at the other port.

The task source (page 429) for the two tasks (page 429) above is the posted message task source (page 547).

If the method is called on a port that is not entangled, then the method must do nothing.

The following are the event handler DOM attributes (page 432) that must be supported by objects implementing the `MessagePort` interface:

`onmessage`

Must be invoked whenever a `message` event is targeted at or bubbles through the `MessagePort` object.

The first time a `MessagePort` object's `onmessage` attribute is set, the port's port message queue (page 568) must be opened, as if the `start()` method had been called.

`onclose`

Must be invoked whenever an `close` event is targeted at or bubbles through the `MessagePort` object.

7.5.3.1 Ports and browsing contexts

When a Document is discarded (page 482), if there are any `MessagePort` objects that:

- are entangled, and
- are owned by the browsing context (page 414) that contained that Document, and
- were created while that Document was the active document (page 414) of that browsing context (page 414), and
- are entangled with a port that is either not owned by that browsing context (page 414) or was not created while that Document was the active document (page 414) of that browsing context (page 414),

...then the user agent must run the following steps for each such port:

1. Let *surviving port* be the port with which the `MessagePort` object in question is entangled.
2. Unentangle the two ports.
3. Queue a task (page 429) to fire a simple event (page 436) called `close` at *surviving port*. The task source (page 429) for this task (page 429) is the posted message task source (page 547).

7.5.3.2 Ports and garbage collection

User agents must act as if `MessagePort` objects have a strong reference to their entangled `MessagePort` object.

Thus, a message port can be received, given an event listener, and then forgotten, and so long as that event listener could receive a message, the channel will be maintained.

Of course, if this was to occur on both sides of the channel, then both ports would be garbage collected, since they would not be reachable from live code, despite having a strong reference to each other.

8 The HTML syntax

8.1 Writing HTML documents

This section only applies to documents, authoring tools, and markup generators. In particular, it does not apply to conformance checkers; conformance checkers must use the requirements given in the next section ("parsing HTML documents").

Documents must consist of the following parts, in the given order:

1. Optionally, a single U+FEFF BYTE ORDER MARK (BOM) character.
2. Any number of comments (page 582) and space characters (page 32).
3. A DOCTYPE (page 573).
4. Any number of comments (page 582) and space characters (page 32).
5. The root element, in the form of an `html` element (page 574).
6. Any number of comments (page 582) and space characters (page 32).

The various types of content mentioned above are described in the next few sections.

In addition, there are some restrictions on how character encoding declarations are to be serialized, as discussed in the section on that topic.

Space characters before the root `html` element, and space characters at the start of the `html` element and before the `head` element, will be dropped when the document is parsed; space characters after the root `html` element will be parsed as if they were at the end of the `body` element. Thus, space characters around the root element do not round-trip.

It is suggested that newlines be inserted after the DOCTYPE, after any comments that are before the root element, after the `html` element's start tag (if it is not omitted (page 578)), and after any comments that are inside the `html` element but before the `head` element.

Many strings in the HTML syntax (e.g. the names of elements and their attributes) are case-insensitive, but only for characters in the ranges U+0041 .. U+005A (LATIN CAPITAL LETTER A to LATIN CAPITAL LETTER Z) and U+0061 .. U+007A (LATIN SMALL LETTER A to LATIN SMALL LETTER Z). For convenience, in this section this is just referred to as "case-insensitive".

8.1.1 The DOCTYPE

A **DOCTYPE** is a mostly useless, but required, header.

Note: DOCTYPES are required for legacy reasons. When omitted, browsers tend to use a different rendering mode that is incompatible with some specifications. Including the DOCTYPE in a document ensures that the

browser makes a best-effort attempt at following the relevant specifications.

A DOCTYPE must consist of the following characters, in this order:

1. A U+003C LESS-THAN SIGN (<) character.
2. A U+0021 EXCLAMATION MARK (!) character.
3. A U+0044 LATIN CAPITAL LETTER D or U+0064 LATIN SMALL LETTER D character.
4. A U+004F LATIN CAPITAL LETTER O or U+006F LATIN SMALL LETTER O character.
5. A U+0043 LATIN CAPITAL LETTER C or U+0063 LATIN SMALL LETTER C character.
6. A U+0054 LATIN CAPITAL LETTER T or U+0074 LATIN SMALL LETTER T character.
7. A U+0059 LATIN CAPITAL LETTER Y or U+0079 LATIN SMALL LETTER Y character.
8. A U+0050 LATIN CAPITAL LETTER P or U+0070 LATIN SMALL LETTER P character.
9. A U+0045 LATIN CAPITAL LETTER E or U+0065 LATIN SMALL LETTER E character.
10. One or more space characters (page 32).
11. A U+0048 LATIN CAPITAL LETTER H or U+0068 LATIN SMALL LETTER H character.
12. A U+0054 LATIN CAPITAL LETTER T or U+0074 LATIN SMALL LETTER T character.
13. A U+004D LATIN CAPITAL LETTER M or U+006D LATIN SMALL LETTER M character.
14. A U+004C LATIN CAPITAL LETTER L or U+006C LATIN SMALL LETTER L character.
15. Optionally, a DOCTYPE legacy string (page 574) (defined below).
16. Zero or more space characters (page 32).
17. A U+003E GREATER-THAN SIGN (>) character.

Note: In other words, <!DOCTYPE HTML>, case-insensitively.

For the purposes of XSLT generators that cannot output HTML markup without a DOCTYPE, a **DOCTYPE legacy string** may be inserted into the DOCTYPE (in the position defined above). This string must consist of:

1. One or more space characters (page 32).
2. A U+0050 LATIN CAPITAL LETTER P or U+0070 LATIN SMALL LETTER P character.
3. A U+0055 LATIN CAPITAL LETTER U or U+0075 LATIN SMALL LETTER U character.
4. A U+0042 LATIN CAPITAL LETTER B or U+0062 LATIN SMALL LETTER B character.
5. A U+004C LATIN CAPITAL LETTER L or U+006C LATIN SMALL LETTER L character.
6. A U+0049 LATIN CAPITAL LETTER I or U+0069 LATIN SMALL LETTER I character.
7. A U+0043 LATIN CAPITAL LETTER C or U+0063 LATIN SMALL LETTER C character.
8. One or more space characters (page 32).
9. A U+0022 QUOTATION MARK or U+0027 APOSTROPHE character.
10. The literal string "XSLT-compat".
11. The same character as in item 9 (a matching U+0022 QUOTATION MARK or U+0027 APOSTROPHE character).

Note: In other words, <!DOCTYPE HTML PUBLIC "XSLT-compat"> or <!DOCTYPE HTML PUBLIC 'XSLT-compat'>, case-insensitively except for the bit in quotes.

The DOCTYPE legacy string (page 574) should not be used unless the document is generated from XSLT.

8.1.2 Elements

There are five different kinds of **elements**: void elements, CDATA elements, RCDATA elements, foreign elements, and normal elements.

Void elements

base, command, eventsource, link, meta, hr, br, img, embed, param, area, col, input, source

CDATA elements

style, script

RCDATA elements

title, textarea

Foreign elements

Elements from the MathML namespace (page 658)

Normal elements

All other allowed HTML elements (page 23) are normal elements.

Tags are used to delimit the start and end of elements in the markup. CDATA, RCDATA, and normal elements have a start tag (page 576) to indicate where they begin, and an end tag (page 576) to indicate where they end. The start and end tags of certain normal elements can be omitted (page 578), as described later. Those that cannot be omitted must not be omitted. Void elements (page 574) only have a start tag; end tags must not be specified for void elements. Foreign elements must either have a start tag and an end tag, or a start tag that is marked as self-closing, in which case they must not have an end tag.

The contents of the element must be placed between just after the start tag (which might be implied, in certain cases (page 578)) and just before the end tag (which again, might be implied in certain cases (page 578)). The exact allowed contents of each individual element depends on the content model of that element, as described earlier in this specification. Elements must not contain content that their content model disallows. In addition to the restrictions placed on the contents by those content models, however, the five types of elements have additional *syntactic* requirements.

Void elements (page 574) can't have any contents (since there's no end tag, no content can be put between the start tag and the end tag).

CDATA elements can have text (page 581), though it has restrictions (page 580) described below.

RCDATA elements can have text (page 581) and character references (page 581), but the text must not contain an ambiguous ampersand (page 581). There are also further restrictions (page 580) described below.

Foreign elements whose start tag is marked as self-closing can't have any contents (since, again, as there's no end tag, no content can be put between the start tag and the end tag). Foreign elements whose start tag is *not* marked as self-closing can have text (page 581), character references (page 581), CDATA sections (page 582), other elements (page 574), and comments (page 582), but the text must not contain the character U+003C LESS-THAN SIGN (<) or an ambiguous ampersand (page 581).

Normal elements can have text (page 581), character references (page 581), other elements (page 574), and comments (page 582), but the text must not contain the character U+003C LESS-THAN SIGN (<) or an ambiguous ampersand (page 581). Some normal elements also have yet more restrictions (page 579) on what content they are allowed to hold, beyond the restrictions imposed by the content model and those described in this paragraph. Those restrictions are described below.

Tags contain a **tag name**, giving the element's name. HTML elements all have names that only use characters in the range U+0030 DIGIT ZERO .. U+0039 DIGIT NINE, U+0061 LATIN

SMALL LETTER A .. U+007A LATIN SMALL LETTER Z, U+0041 LATIN CAPITAL LETTER A .. U+005A LATIN CAPITAL LETTER Z, and U+002D HYPHEN-MINUS (-). In the HTML syntax, tag names may be written with any mix of lower- and uppercase letters that, when converted to all-lowercase, matches the element's tag name; tag names are case-insensitive.

8.1.2.1 Start tags

Start tags must have the following format:

1. The first character of a start tag must be a U+003C LESS-THAN SIGN (<).
2. The next few characters of a start tag must be the element's tag name (page 575).
3. If there are to be any attributes in the next step, there must first be one or more space characters (page 32).
4. Then, the start tag may have a number of attributes, the syntax for which (page 576) is described below. Attributes may be separated from each other by one or more space characters (page 32).
5. After the attributes, there may be one or more space characters (page 32). (Some attributes are required to be followed by a space. See the attributes section (page 576) below.)
6. Then, if the element is one of the void elements (page 574), or if the element is a foreign element, then there may be a single U+002F SOLIDUS (/) character. This character has no effect on void elements, but on foreign elements it marks the start tag as self-closing.
7. Finally, start tags must be closed by a U+003E GREATER-THAN SIGN (>) character.

8.1.2.2 End tags

End tags must have the following format:

1. The first character of an end tag must be a U+003C LESS-THAN SIGN (<).
2. The second character of an end tag must be a U+002F SOLIDUS (/).
3. The next few characters of an end tag must be the element's tag name (page 575).
4. After the tag name, there may be one or more space characters (page 32).
5. Finally, end tags must be closed by a U+003E GREATER-THAN SIGN (>) character.

8.1.2.3 Attributes

Attributes for an element are expressed inside the element's start tag.

Attributes have a name and a value. **Attribute names** must consist of one or more characters other than the space characters (page 32), U+0000 NULL, U+0022 QUOTATION MARK ("'), U+0027 APOSTROPHE ('), U+003E GREATER-THAN SIGN (>), U+002F SOLIDUS (/), and U+003D EQUALS SIGN (=) characters, the control characters, and any characters that

are not defined by Unicode. In the HTML syntax, attribute names may be written with any mix of lower- and uppercase letters that are an ASCII case-insensitive (page 31) match for the attribute's name.

Attribute values are a mixture of text (page 581) and character references (page 581), except with the additional restriction that the text cannot contain an ambiguous ampersand (page 581).

Attributes can be specified in four different ways:

Empty attribute syntax

Just the attribute name (page 576).

In the following example, the disabled attribute is given with the empty attribute syntax:

```
<input disabled>
```

If an attribute using the empty attribute syntax is to be followed by another attribute, then there must be a space character (page 32) separating the two.

Unquoted attribute value syntax

The attribute name (page 576), followed by zero or more space characters (page 32), followed by a single U+003D EQUALS SIGN character, followed by zero or more space characters (page 32), followed by the attribute value (page 577), which, in addition to the requirements given above for attribute values, must not contain any literal space characters (page 32), any U+0022 QUOTATION MARK ("") characters, U+0027 APOSTROPHE ('') characters, U+003D EQUALS SIGN (=) characters, or U+003E GREATER-THAN SIGN (>) characters, and must not be the empty string.

In the following example, the value attribute is given with the unquoted attribute value syntax:

```
<input value=yes>
```

If an attribute using the unquoted attribute syntax is to be followed by another attribute or by one of the optional U+002F SOLIDUS (/) characters allowed in step 6 of the start tag syntax above, then there must be a space character (page 32) separating the two.

Single-quoted attribute value syntax

The attribute name (page 576), followed by zero or more space characters (page 32), followed by a single U+003D EQUALS SIGN character, followed by zero or more space characters (page 32), followed by a single U+0027 APOSTROPHE ('') character, followed by the attribute value (page 577), which, in addition to the requirements given above for attribute values, must not contain any literal U+0027 APOSTROPHE ('') characters, and finally followed by a second single U+0027 APOSTROPHE ('') character.

In the following example, the type attribute is given with the single-quoted attribute value syntax:

```
<input type='checkbox'>
```

If an attribute using the single-quoted attribute syntax is to be followed by another attribute, then there must be a space character (page 32) separating the two.

Double-quoted attribute value syntax

The attribute name (page 576), followed by zero or more space characters (page 32), followed by a single U+003D EQUALS SIGN character, followed by zero or more space characters (page 32), followed by a single U+0022 QUOTATION MARK ("") character, followed by the attribute value (page 577), which, in addition to the requirements given above for attribute values, must not contain any literal U+0022 QUOTATION MARK ("") characters, and finally followed by a second single U+0022 QUOTATION MARK ("") character.

In the following example, the name attribute is given with the double-quoted attribute value syntax:

```
<input name="be evil">
```

If an attribute using the double-quoted attribute syntax is to be followed by another attribute, then there must be a space character (page 32) separating the two.

There must never be two or more attributes on the same start tag whose names are an ASCII case-insensitive (page 31) match for each other.

8.1.2.4 Optional tags

Certain tags can be **omitted**.

An `html` element's start tag may be omitted if the first thing inside the `html` element is not a comment (page 582).

An `html` element's end tag may be omitted if the `html` element is not immediately followed by a comment (page 582) and the element contains a `body` element that is either not empty or whose start tag has not been omitted.

A `head` element's start tag may be omitted if the first thing inside the `head` element is an element.

A `head` element's end tag may be omitted if the `head` element is not immediately followed by a space character (page 32) or a comment (page 582).

A `body` element's start tag may be omitted if the first thing inside the `body` element is not a space character (page 32) or a comment (page 582), except if the first thing inside the `body` element is a `script` or `style` element.

A `body` element's end tag may be omitted if the `body` element is not immediately followed by a comment (page 582) and the element is either not empty or its start tag has not been omitted.

A `li` element's end tag may be omitted if the `li` element is immediately followed by another `li` element or if there is no more content in the parent element.

A `dt` element's end tag may be omitted if the `dt` element is immediately followed by another `dt` element or a `dd` element.

A `dd` element's end tag may be omitted if the `dd` element is immediately followed by another `dd` element or a `dt` element, or if there is no more content in the parent element.

A `p` element's end tag may be omitted if the `p` element is immediately followed by an `address`, `article`, `aside`, `blockquote`, `datagrid`, `dialog`, `dir`, `div`, `dl`, `fieldset`, `footer`, `form`, `h1`, `h2`, `h3`, `h4`, `h5`, `h6`, `header`, `hr`, `menu`, `nav`, `ol`, `p`, `pre`, `section`, `table`, or `ul`, element, or if there is no more content in the parent element and the parent element is not an `a` element.

An `rt` element's end tag may be omitted if the `rt` element is immediately followed by an `rt` or `rp` element, or if there is no more content in the parent element.

An `rp` element's end tag may be omitted if the `rp` element is immediately followed by an `rt` or `rp` element, or if there is no more content in the parent element.

An `optgroup` element's end tag may be omitted if the `optgroup` element is immediately followed by another `optgroup` element, or if there is no more content in the parent element.

An `option` element's end tag may be omitted if the `option` element is immediately followed by another `option` element, or if it is immediately followed by an `optgroup` element, or if there is no more content in the parent element.

A `colgroup` element's start tag may be omitted if the first thing inside the `colgroup` element is a `col` element, and if the element is not immediately preceded by another `colgroup` element whose end tag has been omitted.

A `colgroup` element's end tag may be omitted if the `colgroup` element is not immediately followed by a space character (page 32) or a comment (page 582).

A `thead` element's end tag may be omitted if the `thead` element is immediately followed by a `tbody` or `tfoot` element.

A `tbody` element's start tag may be omitted if the first thing inside the `tbody` element is a `tr` element, and if the element is not immediately preceded by a `tbody`, `thead`, or `tfoot` element whose end tag has been omitted.

A `tbody` element's end tag may be omitted if the `tbody` element is immediately followed by a `tbody` or `tfoot` element, or if there is no more content in the parent element.

A `tfoot` element's end tag may be omitted if the `tfoot` element is immediately followed by a `tbody` element, or if there is no more content in the parent element.

A `tr` element's end tag may be omitted if the `tr` element is immediately followed by another `tr` element, or if there is no more content in the parent element.

A `td` element's end tag may be omitted if the `td` element is immediately followed by a `td` or `th` element, or if there is no more content in the parent element.

A `th` element's end tag may be omitted if the `th` element is immediately followed by a `td` or `th` element, or if there is no more content in the parent element.

However, a start tag must never be omitted if it has any attributes.

8.1.2.5 Restrictions on content models

For historical reasons, certain elements have extra restrictions beyond even the restrictions given by their content model.

An optgroup element must not contain optgroup elements, even though these elements are technically allowed to be nested according to the content models described in this specification. (If an optgroup element is put inside another in the markup, it will in fact imply an optgroup end tag before it.)

A table element must not contain tr elements, even though these elements are technically allowed inside table elements according to the content models described in this specification. (If a tr element is put inside a table in the markup, it will in fact imply a tbody start tag before it.)

A single U+000A LINE FEED (LF) character may be placed immediately after the start tag of pre and textarea elements. This does not affect the processing of the element. The otherwise optional U+000A LINE FEED (LF) character *must* be included if the element's contents start with that character (because otherwise the leading newline in the contents would be treated like the optional newline, and ignored).

The following two pre blocks are equivalent:

```
<pre>Hello</pre>
<pre>
Hello</pre>
```

8.1.2.6 Restrictions on the contents of CDATA and RCDATA elements

The text in CDATA and RCDATA elements must not contain any occurrences of the string "</" (U+003C LESS-THAN SIGN, U+002F SOLIDUS) followed by characters that case-insensitively match the tag name of the element followed by one of U+0009 CHARACTER TABULATION, U+000A LINE FEED (LF), U+000C FORM FEED (FF), U+0020 SPACE, U+003E GREATER-THAN SIGN (>), or U+002F SOLIDUS (/), unless that string is part of an escaping text span (page 580).

An **escaping text span** is a span of text (page 581) that starts with an escaping text span start (page 580) that is not itself in an escaping text span (page 580), and ends at the next escaping text span end (page 580). There cannot be any character references (page 581) inside an escaping text span (page 580).

An **escaping text span start** is a part of text (page 581) that consists of the four character sequence "<!--" (U+003C LESS-THAN SIGN, U+0021 EXCLAMATION MARK, U+002D HYPHEN-MINUS, U+002D HYPHEN-MINUS).

An **escaping text span end** is a part of text (page 581) that consists of the three character sequence "-->" (U+002D HYPHEN-MINUS, U+002D HYPHEN-MINUS, U+003E GREATER-THAN SIGN) whose U+003E GREATER-THAN SIGN (>).

An escaping text span start (page 580) may share its U+002D HYPHEN-MINUS characters with its corresponding escaping text span end (page 580).

The text in CDATA and RCDATA elements must not have an escaping text span start (page 580) that is not followed by an escaping text span end (page 580).

8.1.3 Text

Text is allowed inside elements, attributes, and comments. Text must consist of Unicode characters. Text must not contain U+0000 characters. Text must not contain permanently undefined Unicode characters. Text must not contain control characters other than space characters (page 32). Extra constraints are placed on what is and what is not allowed in text based on where the text is to be put, as described in the other sections.

8.1.3.1 Newlines

Newlines in HTML may be represented either as U+000D CARRIAGE RETURN (CR) characters, U+000A LINE FEED (LF) characters, or pairs of U+000D CARRIAGE RETURN (CR), U+000A LINE FEED (LF) characters in that order.

8.1.4 Character references

In certain cases described in other sections, text (page 581) may be mixed with **character references**. These can be used to escape characters that couldn't otherwise legally be included in text (page 581).

Character references must start with a U+0026 AMPERSAND (&). Following this, there are three possible kinds of character references:

Named character references

The ampersand must be followed by one of the names given in the named character references (page 662) section, using the same case. The name must be one that is terminated by a U+003B SEMICOLON (;) character.

Decimal numeric character reference

The ampersand must be followed by a U+0023 NUMBER SIGN (#) character, followed by one or more digits in the range U+0030 DIGIT ZERO .. U+0039 DIGIT NINE, representing a base-ten integer that itself is a Unicode code point that is not U+0000, U+000D, in the range U+0080 .. U+009F, or in the range 0xD800 .. 0xDFFF (surrogates). The digits must then be followed by a U+003B SEMICOLON character (;).

Hexadecimal numeric character reference

The ampersand must be followed by a U+0023 NUMBER SIGN (#) character, which must be followed by either a U+0078 LATIN SMALL LETTER X or a U+0058 LATIN CAPITAL LETTER X character, which must then be followed by one or more digits in the range U+0030 DIGIT ZERO .. U+0039 DIGIT NINE, U+0061 LATIN SMALL LETTER A .. U+0066 LATIN SMALL LETTER F, and U+0041 LATIN CAPITAL LETTER A .. U+0046 LATIN CAPITAL LETTER F, representing a base-sixteen integer that itself is a Unicode code point that is not U+0000, U+000D, in the range U+0080 .. U+009F, or in the range 0xD800 .. 0xDFFF (surrogates). The digits must then be followed by a U+003B SEMICOLON character (;).

An **ambiguous ampersand** is a U+0026 AMPERSAND (&) character that is followed by some text (page 581) other than a space character (page 32), a U+003C LESS-THAN SIGN character ('<'), or another U+0026 AMPERSAND (&) character.

8.1.5 CDATA sections

CDATA sections must start with the character sequence U+003C LESS-THAN SIGN, U+0021 EXCLAMATION MARK, U+005B LEFT SQUARE BRACKET, U+0043 LATIN CAPITAL LETTER C, U+0044 LATIN CAPITAL LETTER D, U+0041 LATIN CAPITAL LETTER A, U+0054 LATIN CAPITAL LETTER T, U+0041 LATIN CAPITAL LETTER A, U+005B LEFT SQUARE BRACKET (<! [CDATA[). Following this sequence, the CDATA section may have text (page 581), with the additional restriction that the text must not contain the three character sequence U+005D RIGHT SQUARE BRACKET, U+005D RIGHT SQUARE BRACKET, U+003E GREATER-THAN SIGN ()]). Finally, the CDATA section must be ended by the three character sequence U+005D RIGHT SQUARE BRACKET, U+005D RIGHT SQUARE BRACKET, U+003E GREATER-THAN SIGN ()]).

8.1.6 Comments

Comments must start with the four character sequence U+003C LESS-THAN SIGN, U+0021 EXCLAMATION MARK, U+002D HYPHEN-MINUS, U+002D HYPHEN-MINUS (<! --). Following this sequence, the comment may have text (page 581), with the additional restriction that the text must not start with a single U+003E GREATER-THAN SIGN ('>') character, nor start with a U+002D HYPHEN-MINUS (-) character followed by a U+003E GREATER-THAN SIGN ('>') character, nor contain two consecutive U+002D HYPHEN-MINUS (-) characters, nor end with a U+002D HYPHEN-MINUS (-) character. Finally, the comment must be ended by the three character sequence U+002D HYPHEN-MINUS, U+002D HYPHEN-MINUS, U+003E GREATER-THAN SIGN (-->).

8.2 Parsing HTML documents

This section only applies to user agents, data mining tools, and conformance checkers.

The rules for parsing XML documents (page 76) (and thus XHTML (page 28) documents) into DOM trees are covered by the XML and Namespaces in XML specifications, and are out of scope of this specification. [XML] [XMLNS]

For HTML documents (page 76), user agents must use the parsing rules described in this section to generate the DOM trees. Together, these rules define what is referred to as the **HTML parser**.

While the HTML form of HTML5 bears a close resemblance to SGML and XML, it is a separate language with its own parsing rules.

Some earlier versions of HTML (in particular from HTML2 to HTML4) were based on SGML and used SGML parsing rules. However, few (if any) web browsers ever implemented true SGML parsing for HTML documents; the only user agents to strictly handle HTML as an SGML application have historically been validators. The resulting confusion — with validators claiming documents to have one representation while widely deployed Web browsers interoperably implemented a different representation — has wasted decades of productivity. This version of HTML thus returns to a non-SGML basis.

Authors interested in using SGML tools in their authoring pipeline are encouraged to use XML tools and the XML serialization of HTML5.

This specification defines the parsing rules for HTML documents, whether they are syntactically correct or not. Certain points in the parsing algorithm are said to be **parse errors**. The error handling for parse errors is well-defined: user agents must either act as described below when encountering such problems, or must abort processing at the first error that they encounter for which they do not wish to apply the rules described below.

Conformance checkers must report at least one parse error condition to the user if one or more parse error conditions exist in the document and must not report parse error conditions if none exist in the document. Conformance checkers may report more than one parse error condition if more than one parse error conditions exist in the document. Conformance checkers are not required to recover from parse errors.

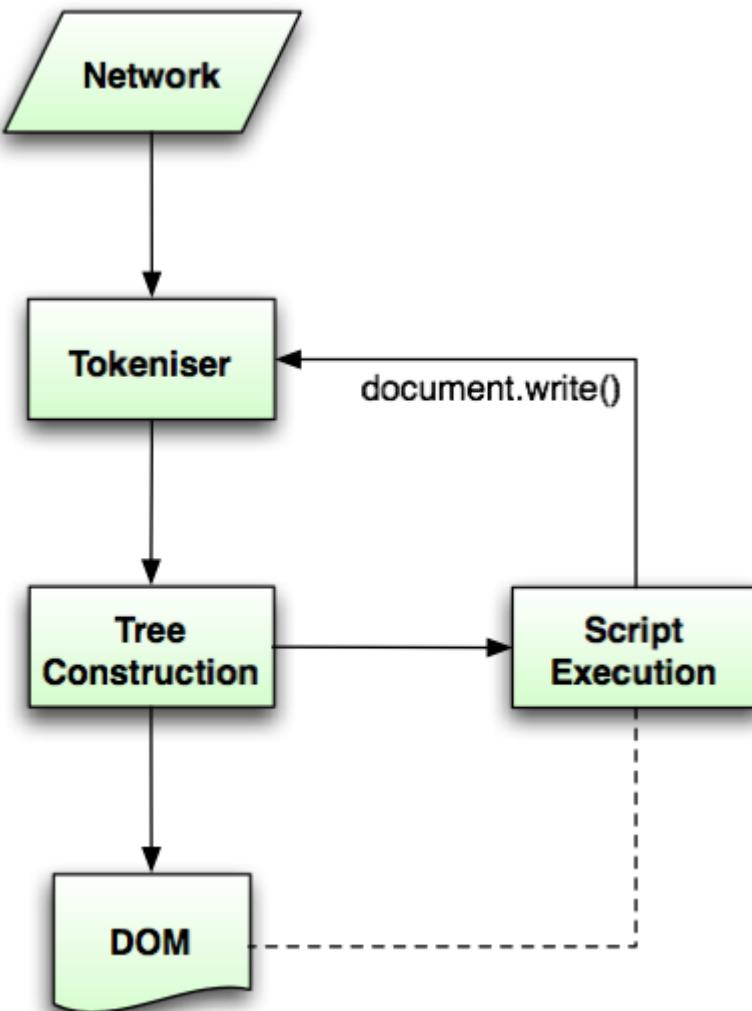
Note: *Parse errors are only errors with the syntax of HTML. In addition to checking for parse errors, conformance checkers will also verify that the document obeys all the other conformance requirements described in this specification.*

8.2.1 Overview of the parsing model

The input to the HTML parsing process consists of a stream of Unicode characters, which is passed through a tokenization (page 597) stage (lexical analysis) followed by a tree construction (page 617) stage (semantic analysis). The output is a Document object.

Note: *Implementations that do not support scripting (page 26) do not have to actually create a DOM Document object, but the DOM tree in such cases is still used as the model for the rest of the specification.*

In the common case, the data handled by the tokenization stage comes from the network, but it can also come from script (page 99), e.g. using the `document.write()` API.



There is only one set of state for the tokeniser stage and the tree construction stage, but the tree construction stage is reentrant, meaning that while the tree construction stage is handling one token, the tokeniser might be resumed, causing further tokens to be emitted and processed before the first token's processing is complete.

In the following example, the tree construction stage will be called upon to handle a "p" start tag token while handling the "script" start tag token:

```

...
<script>
  document.write('<p>');
</script>
...

```

8.2.2 The input stream

The stream of Unicode characters that consists the input to the tokenization stage will be initially seen by the user agent as a stream of bytes (typically coming over the network or

from the local file system). The bytes encode the actual characters according to a particular *character encoding*, which the user agent must use to decode the bytes into characters.

Note: For XML documents, the algorithm user agents must use to determine the character encoding is given by the XML specification. This section does not apply to XML documents. [XML]

8.2.2.1 Determining the character encoding

In some cases, it might be impractical to unambiguously determine the encoding before parsing the document. Because of this, this specification provides for a two-pass mechanism with an optional pre-scan. Implementations are allowed, as described below, to apply a simplified parsing algorithm to whatever bytes they have available before beginning to parse the document. Then, the real parser is started, using a tentative encoding derived from this pre-parse and other out-of-band metadata. If, while the document is being loaded, the user agent discovers an encoding declaration that conflicts with this information, then the parser can get reinvoked to perform a parse of the document with the real encoding.

User agents must use the following algorithm (the **encoding sniffing algorithm**) to determine the character encoding to use when decoding a document in the first pass. This algorithm takes as input any out-of-band metadata available to the user agent (e.g. the Content-Type metadata (page 60) of the document) and all the bytes available so far, and returns an encoding and a **confidence**. The confidence is either *tentative* or *certain*. The encoding used, and whether the confidence in that encoding is *tentative* or *confident*, is used during the parsing (page 625) to determine whether to change the encoding (page 592).

1. If the transport layer specifies an encoding, and it is supported, return that encoding with the confidence (page 585) *certain*, and abort these steps.
2. The user agent may wait for more bytes of the resource to be available, either in this step or at any later step in this algorithm. For instance, a user agent might wait 500ms or 512 bytes, whichever came first. In general preparsing the source to find the encoding improves performance, as it reduces the need to throw away the data structures used when parsing upon finding the encoding information. However, if the user agent delays too long to obtain data to determine the encoding, then the cost of the delay could outweigh any performance improvements from the preparse.
3. For each of the rows in the following table, starting with the first one and going down, if there are as many or more bytes available than the number of bytes in the first column, and the first bytes of the file match the bytes given in the first column, then return the encoding given in the cell in the second column of that row, with the confidence (page 585) *certain*, and abort these steps:

Bytes in Hexadecimal	Encoding
FE FF	UTF-16BE
FF FE	UTF-16LE
EF BB BF	UTF-8

Note: This step looks for Unicode Byte Order Marks (BOMs).

4. Otherwise, the user agent will have to search for explicit character encoding information in the file itself. This should proceed as follows:

Let *position* be a pointer to a byte in the input stream, initially pointing at the first byte. If at any point during these substeps the user agent either runs out of bytes or decides that scanning further bytes would not be efficient, then skip to the next step of the overall character encoding detection algorithm. User agents may decide that scanning *any* bytes is not efficient, in which case these substeps are entirely skipped.

Now, repeat the following "two" steps until the algorithm aborts (either because user agent aborts, as described above, or because a character encoding is found):

1. If *position* points to:

↪ **A sequence of bytes starting with: 0x3C 0x21 0x2D 0x2D (ASCII '<!--')**

Advance the *position* pointer so that it points at the first 0x3E byte which is preceded by two 0x2D bytes (i.e. at the end of an ASCII '-->' sequence) and comes after the 0x3C byte that was found. (The two 0x2D bytes can be the same as those in the '<!--' sequence.)

↪ **A sequence of bytes starting with: 0x3C, 0x4D or 0x6D, 0x45 or 0x65, 0x54 or 0x74, 0x41 or 0x61, and finally one of 0x09, 0x0A, 0x0C, 0x0D, 0x20, 0x2F (case-insensitive ASCII '<meta' followed by a space or slash)**

1. Advance the *position* pointer so that it points at the next 0x09, 0x0A, 0x0C, 0x0D, 0x20, or 0x2F byte (the one in sequence of characters matched above).
2. Get an attribute (page 587) and its value. If no attribute was sniffed, then skip this inner set of steps, and jump to the second step in the overall "two step" algorithm.
3. If the attribute's name is neither "charset" nor "content", then return to step 2 in these inner steps.
4. If the attribute's name is "charset", let *charset* be the attribute's value, interpreted as a character encoding.
5. Otherwise, the attribute's name is "content": apply the algorithm for extracting an encoding from a Content-Type (page 60), giving the attribute's value as the string to parse. If an encoding is returned, let *charset* be that encoding. Otherwise, return to step 2 in these inner steps.
6. If *charset* is a UTF-16 encoding, change it to UTF-8.
7. If *charset* is a supported character encoding, then return the given encoding, with confidence (page 585) *tentative*, and abort all these steps.
8. Otherwise, return to step 2 in these inner steps.

↪ A sequence of bytes starting with a 0x3C byte (ASCII '<'), optionally a 0x2F byte (ASCII '/'), and finally a byte in the range 0x41-0x5A or 0x61-0x7A (an ASCII letter)

1. Advance the *position* pointer so that it points at the next 0x09 (ASCII TAB), 0x0A (ASCII LF), 0x0C (ASCII FF), 0x0D (ASCII CR), 0x20 (ASCII space), or 0x3E (ASCII '>') byte.
2. Repeatedly get an attribute (page 587) until no further attributes can be found, then jump to the second step in the overall "two step" algorithm.

↪ A sequence of bytes starting with: 0x3C 0x21 (ASCII '<!')

↪ A sequence of bytes starting with: 0x3C 0x2F (ASCII '</')

↪ A sequence of bytes starting with: 0x3C 0x3F (ASCII '<?')

Advance the *position* pointer so that it points at the first 0x3E byte (ASCII '>') that comes after the 0x3C byte that was found.

↪ Any other byte

Do nothing with that byte.

2. Move *position* so it points at the next byte in the input stream, and return to the first step of this "two step" algorithm.

When the above "two step" algorithm says to **get an attribute**, it means doing this:

1. If the byte at *position* is one of 0x09 (ASCII TAB), 0x0A (ASCII LF), 0x0C (ASCII FF), 0x0D (ASCII CR), 0x20 (ASCII space), or 0x2F (ASCII '/') then advance *position* to the next byte and redo this substep.
2. If the byte at *position* is 0x3E (ASCII '>'), then abort the "get an attribute" algorithm. There isn't one.
3. Otherwise, the byte at *position* is the start of the attribute name. Let *attribute name* and *attribute value* be the empty string.
4. *Attribute name*: Process the byte at *position* as follows:

↪ If it is 0x3D (ASCII '='), and the **attribute name** is longer than the empty string

Advance *position* to the next byte and jump to the step below labeled *value*.

↪ If it is 0x09 (ASCII TAB), 0x0A (ASCII LF), 0x0C (ASCII FF), 0x0D (ASCII CR), or 0x20 (ASCII space)

Jump to the step below labeled *spaces*.

↪ If it is 0x2F (ASCII '/') or 0x3E (ASCII '>')

Abort the "get an attribute" algorithm. The attribute's name is the value of *attribute name*, its value is the empty string.

↪ If it is in the range 0x41 (ASCII 'A') to 0x5A (ASCII 'Z')

Append the Unicode character with codepoint *b*+0x20 to *attribute name* (where *b* is the value of the byte at *position*).

↪ Anything else

Append the Unicode character with the same codepoint as the value of the byte at *position*) to *attribute name*. (It doesn't actually matter how bytes outside the ASCII range are handled here, since only ASCII characters can contribute to the detection of a character encoding.)

5. Advance *position* to the next byte and return to the previous step.
6. *Spaces*. If the byte at *position* is one of 0x09 (ASCII TAB), 0x0A (ASCII LF), 0x0C (ASCII FF), 0x0D (ASCII CR), or 0x20 (ASCII space) then advance *position* to the next byte, then, repeat this step.
7. If the byte at *position* is not 0x3D (ASCII '='), abort the "get an attribute" algorithm. The attribute's name is the value of *attribute name*, its value is the empty string.
8. Advance *position* past the 0x3D (ASCII '=') byte.
9. *Value*. If the byte at *position* is one of 0x09 (ASCII TAB), 0x0A (ASCII LF), 0x0C (ASCII FF), 0x0D (ASCII CR), or 0x20 (ASCII space) then advance *position* to the next byte, then, repeat this step.
10. Process the byte at *position* as follows:

↪ If it is 0x22 (ASCII "") or 0x27 ("")

1. Let *b* be the value of the byte at *position*.
2. Advance *position* to the next byte.
3. If the value of the byte at *position* is the value of *b*, then advance *position* to the next byte and abort the "get an attribute" algorithm. The attribute's name is the value of *attribute name*, and its value is the value of *attribute value*.
4. Otherwise, if the value of the byte at *position* is in the range 0x41 (ASCII 'A') to 0x5A (ASCII 'Z'), then append a Unicode character to *attribute value* whose codepoint is 0x20 more than the value of the byte at *position*.
5. Otherwise, append a Unicode character to *attribute value* whose codepoint is the same as the value of the byte at *position*.
6. Return to the second step in these substeps.

↪ If it is 0x3E (ASCII '>')

Abort the "get an attribute" algorithm. The attribute's name is the value of *attribute name*, its value is the empty string.

↪ If it is in the range 0x41 (ASCII 'A') to 0x5A (ASCII 'Z')

Append the Unicode character with codepoint *b*+0x20 to *attribute value* (where *b* is the value of the byte at *position*). Advance *position* to the next byte.

↪ **Anything else**

Append the Unicode character with the same codepoint as the value of the byte at *position*) to *attribute value*. Advance *position* to the next byte.

11. Process the byte at *position* as follows:

↪ **If it is 0x09 (ASCII TAB), 0x0A (ASCII LF), 0x0C (ASCII FF), 0x0D (ASCII CR), 0x20 (ASCII space), or 0x3E (ASCII '>')**

Abort the "get an attribute" algorithm. The attribute's name is the value of *attribute name* and its value is the value of *attribute value*.

↪ **If it is in the range 0x41 (ASCII 'A') to 0x5A (ASCII 'Z')**

Append the Unicode character with codepoint *b*+0x20 to *attribute value* (where *b* is the value of the byte at *position*).

↪ **Anything else**

Append the Unicode character with the same codepoint as the value of the byte at *position*) to *attribute value*.

12. Advance *position* to the next byte and return to the previous step.

For the sake of interoperability, user agents should not use a pre-scan algorithm that returns different results than the one described above. (But, if you do, please at least let us know, so that we can improve this algorithm and benefit everyone...)

5. If the user agent has information on the likely encoding for this page, e.g. based on the encoding of the page when it was last visited, then return that encoding, with the confidence (page 585) *tentative*, and abort these steps.
6. The user agent may attempt to autodetect the character encoding from applying frequency analysis or other algorithms to the data stream. If autodetection succeeds in determining a character encoding, then return that encoding, with the confidence (page 585) *tentative*, and abort these steps. [UNIVCHARDET]
7. Otherwise, return an implementation-defined or user-specified default character encoding, with the confidence (page 585) *tentative*. In non-legacy environments, the more comprehensive UTF-8 encoding is recommended. Due to its use in legacy content, windows-1252 is recommended as a default in predominantly Western demographics instead. Since these encodings can in many cases be distinguished by inspection, a user agent may heuristically decide which to use as a default.

The document's character encoding (page 79) must immediately be set to the value returned from this algorithm, at the same time as the user agent uses the returned value to select the decoder to use for the input stream.

8.2.2.2 Character encoding requirements

User agents must at a minimum support the UTF-8 and Windows-1252 encodings, but may support more.

Note: It is not unusual for Web browsers to support dozens if not upwards of a hundred distinct character encodings.

User agents must support the preferred MIME name of every character encoding they support that has a preferred MIME name, and should support all the IANA-registered aliases. [IANACHARSET]

When comparing a string specifying a character encoding with the name or alias of a character encoding to determine if they are equal, user agents must ignore all characters in the ranges U+0009 to U+000D, U+0020 to U+002F, U+003A to U+0040, U+005B to U+0060, and U+007B to U+007E (all whitespace and punctuation characters in ASCII) in both names, and then perform the comparison in an ASCII case-insensitive (page 31) manner.

For instance, "GB_2312-80" and "g.b.2312(80)" are considered equivalent names.

When a user agent would otherwise use an encoding given in the first column of the following table, it must instead use the encoding given in the cell in the second column of the same row. Any bytes that are treated differently due to this encoding aliasing must be considered parse errors (page 583).

Character encoding overrides		
Input encoding	Replacement encoding	References
EUC-KR	Windows-949	[EUCKR] [WIN949]
GB2312	GBK	[GB2312] [GBK]
GB_2312-80	GBK	[RFC1345] [GBK]
ISO-8859-1	Windows-1252	[RFC1345] [WIN1252]
ISO-8859-9	Windows-1254	[RFC1345] [WIN1254]
ISO-8859-11	Windows-874	[ISO885911] [WIN874]
KS_C_5601-1987	Windows-949	[RFC1345] [WIN949]
TIS-620	Windows-874	[TIS620] [WIN874]
US-ASCII	Windows-1252	[RFC1345] [WIN1252]
x-x-big5	Big5	[BIG5]

Note: The requirement to treat certain encodings as other encodings according to the table above is a willful violation of the W3C Character Model specification. [CHARMOD]

User agents must not support the CESU-8, UTF-7, BOCU-1 and SCSU encodings. [CESU8] [UTF7] [BOCU1] [SCSU]

Support for encodings based on EBCDIC is not recommended. This encoding is rarely used for publicly-facing Web content.

Support for UTF-32 is not recommended. This encoding is rarely used, and frequently misimplemented.

Note: This specification does not make any attempt to support EBCDIC-based encodings and UTF-32 in its algorithms; support and use of these encodings can thus lead to unexpected behavior in implementations of this specification.

8.2.2.3 Preprocessing the input stream

Given an encoding, the bytes in the input stream must be converted to Unicode characters for the tokeniser, as described by the rules for that encoding, except that the leading U+FEFF BYTE ORDER MARK character, if any, must not be stripped by the encoding layer (it is stripped by the rule below).

Bytes or sequences of bytes in the original byte stream that could not be converted to Unicode characters must be converted to U+FFFD REPLACEMENT CHARACTER code points.

Note: Bytes or sequences of bytes in the original byte stream that did not conform to the encoding specification (e.g. invalid UTF-8 byte sequences in a UTF-8 input stream) are errors that conformance checkers are expected to report.

One leading U+FEFF BYTE ORDER MARK character must be ignored if any are present.

All U+0000 NULL characters in the input must be replaced by U+FFFD REPLACEMENT CHARACTERs. Any occurrences of such characters is a parse error (page 583).

Any occurrences of any characters in the ranges U+0001 to U+0008, U+000B, U+000E to U+001F, U+007F to U+009F, U+D800 to U+DFFF, U+FDD0 to U+FDDF, and characters U+FFFE, U+FFFF, U+1FFE, U+1FFF, U+2FFE, U+2FFF, U+3FFE, U+3FFF, U+4FFE, U+4FFF, U+5FFE, U+5FFF, U+6FFE, U+6FFF, U+7FFE, U+7FFF, U+8FFE, U+8FFF, U+9FFE, U+9FFF, U+AFFE, U+AFFF, U+BFFE, U+BFFF, U+CFFE, U+CFEE, U+DFFF, U+EFFE, U+EFFF, U+FFFF, U+FFFFF, U+10FFE, and U+10FFF are parse errors (page 583). (These are all control characters or permanently undefined Unicode characters.)

U+000D CARRIAGE RETURN (CR) characters and U+000A LINE FEED (LF) characters are treated specially. Any CR characters that are followed by LF characters must be removed, and any CR characters not followed by LF characters must be converted to LF characters. Thus, newlines in HTML DOMs are represented by LF characters, and there are never any CR characters in the input to the tokenization (page 597) stage.

The **next input character** is the first character in the input stream that has not yet been **consumed**. Initially, the *next input character* is the first character in the input.

The **insertion point** is the position (just before a character or just before the end of the input stream) where content inserted using `document.write()` is actually inserted. The insertion point is relative to the position of the character immediately after it, it is not an absolute offset into the input stream. Initially, the insertion point is uninitialized.

The "EOF" character in the tables below is a conceptual character representing the end of the input stream (page 584). If the parser is a script-created parser (page 100), then the end of the input stream (page 584) is reached when an **explicit "EOF" character** (inserted by the `document.close()` method) is consumed. Otherwise, the "EOF" character is not a real character in the stream, but rather the lack of any further characters.

8.2.2.4 Changing the encoding while parsing

When the parser requires the user agent to **change the encoding**, it must run the following steps. This might happen if the encoding sniffing algorithm (page 585) described above failed to find an encoding, or if it found an encoding that was not the actual encoding of the file.

1. If the new encoding is a UTF-16 encoding, change it to UTF-8.
2. If the new encoding is identical or equivalent to the encoding that is already being used to interpret the input stream, then set the confidence (page 585) to *confident* and abort these steps. This happens when the encoding information found in the file matches what the encoding sniffing algorithm (page 585) determined to be the encoding, and in the second pass through the parser if the first pass found that the encoding sniffing algorithm described in the earlier section failed to find the right encoding.
3. If all the bytes up to the last byte converted by the current decoder have the same Unicode interpretations in both the current encoding and the new encoding, and if the user agent supports changing the converter on the fly, then the user agent may change to the new converter for the encoding on the fly. Set the document's character encoding (page 79) and the encoding used to convert the input stream to the new encoding, set the confidence (page 585) to *confident*, and abort these steps.
4. Otherwise, navigate (page 473) to the document again, with replacement enabled (page 476), and using the same source browsing context (page 473), but this time skip the encoding sniffing algorithm (page 585) and instead just set the encoding to the new encoding and the confidence (page 585) to *confident*. Whenever possible, this should be done without actually contacting the network layer (the bytes should be re-parsed from memory), even if, e.g., the document is marked as not being cacheable.

8.2.3 Parse state

8.2.3.1 The insertion mode

The **insertion mode** is a flag that controls the primary operation of the tree construction stage.

Initially the insertion mode (page 592) is "initial (page 620)". It can change to "before html (page 623)", "before head (page 624)", "in head (page 625)", "in head noscript (page 626)", "after head (page 627)", "in body (page 628)", "in CDATA/RCDATA (page 640)", "in table (page 642)", "in caption (page 644)", "in column group (page 645)", "in table body (page 646)", "in row (page 647)", "in cell (page 648)", "in select (page 649)", "in select in table (page 651)", "in foreign content (page 651)", "after body (page 653)", "in frameset (page 653)", "after frameset (page 654)", "after after body (page 655)", and "after after frameset (page 656)" during the course of the parsing, as described in the tree construction (page 617) stage. The insertion mode affects how tokens are processed and whether CDATA sections are supported.

Seven of these modes, namely "in head (page 625)", "in body (page 628)", "in CDATA/RCDATA (page 640)", "in table (page 642)", "in table body (page 646)", "in row (page 647)", "in cell (page 648)", and "in select (page 649)", are special, in that the other modes defer to them at various times. When the algorithm below says that the user agent is to do something

"using the rules for the m insertion mode", where m is one of these modes, the user agent must use the rules described under the m insertion mode (page 592)'s section, but must leave the insertion mode (page 592) unchanged unless the rules in m themselves switch the insertion mode (page 592) to a new value.

When the insertion mode is switched to "in CDATA/RCDATA (page 640)", the **original insertion mode** is also set. This is the insertion mode to which the tree construction stage will return when the corresponding end tag is parsed.

When the insertion mode is switched to "in foreign content (page 651)", the **secondary insertion mode** is also set. This secondary mode is used within the rules for the "in foreign content (page 651)" mode to handle HTML (i.e. not foreign) content.

When the steps below require the UA to **reset the insertion mode appropriately**, it means the UA must follow these steps:

1. Let $last$ be false.
2. Let $node$ be the last node in the stack of open elements (page 594).
3. If $node$ is the first node in the stack of open elements, then set $last$ to true and set $node$ to the context element. (fragment case (page 661))
4. If $node$ is a select element, then switch the insertion mode (page 592) to "in select (page 649)" and abort these steps. (fragment case (page 661))
5. If $node$ is a td or th element and $last$ is false, then switch the insertion mode (page 592) to "in cell (page 648)" and abort these steps.
6. If $node$ is a tr element, then switch the insertion mode (page 592) to "in row (page 647)" and abort these steps.
7. If $node$ is a tbody, thead, or tfoot element, then switch the insertion mode (page 592) to "in table body (page 646)" and abort these steps.
8. If $node$ is a caption element, then switch the insertion mode (page 592) to "in caption (page 644)" and abort these steps.
9. If $node$ is a colgroup element, then switch the insertion mode (page 592) to "in column group (page 645)" and abort these steps. (fragment case (page 661))
10. If $node$ is a table element, then switch the insertion mode (page 592) to "in table (page 642)" and abort these steps.
11. If $node$ is an element from the MathML namespace (page 658), then switch the insertion mode (page 592) to "in foreign content (page 651)", let the secondary insertion mode (page 593) be "in body (page 628)", and abort these steps.
12. If $node$ is a head element, then switch the insertion mode (page 592) to "in body (page 628)" ("in body (page 628)"! *not* "in head (page 625)"!) and abort these steps. (fragment case (page 661))
13. If $node$ is a body element, then switch the insertion mode (page 592) to "in body (page 628)" and abort these steps.

14. If *node* is a frameset element, then switch the insertion mode (page 592) to "in frameset (page 653)" and abort these steps. (fragment case (page 661))
15. If *node* is an html element, then: if the head element pointer (page 597) is null, switch the insertion mode (page 592) to "before head (page 624)", otherwise, switch the insertion mode (page 592) to "after head (page 627)". In either case, abort these steps. (fragment case (page 661))
16. If *last* is true, then switch the insertion mode (page 592) to "in body (page 628)" and abort these steps. (fragment case (page 661))
17. Let *node* now be the node before *node* in the stack of open elements (page 594).
18. Return to step 3.

8.2.3.2 The stack of open elements

Initially the **stack of open elements** is empty. The stack grows downwards; the topmost node on the stack is the first one added to the stack, and the bottommost node of the stack is the most recently added node in the stack (notwithstanding when the stack is manipulated in a random access fashion as part of the handling for misnested tags (page 633)).

The "before html (page 623)" insertion mode (page 592) creates the html root element node, which is then added to the stack.

In the fragment case (page 661), the stack of open elements (page 594) is initialized to contain an html element that is created as part of that algorithm (page 661). (The fragment case (page 661) skips the "before html (page 623)" insertion mode (page 592).)

The html node, however it is created, is the topmost node of the stack. It never gets popped off the stack.

The **current node** is the bottommost node in this stack.

The **current table** is the last table element in the stack of open elements (page 594), if there is one. If there is no table element in the stack of open elements (page 594) (fragment case (page 661)), then the current table (page 594) is the first element in the stack of open elements (page 594) (the html element).

Elements in the stack fall into the following categories:

Special

The following HTML elements have varying levels of special parsing rules: address, area, article, aside, base, basefont, bgsound,blockquote, body, br, center, col, colgroup, command, datagrid, dd, details, dialog, dir, div, dl, dt, embed, eventsource fieldset, figure, footer, form, frame, frameset, h1, h2, h3, h4, h5, h6, head, header, hr, iframe, img, input, isindex, li, link, listing, menu, meta, nav, noembed,noframes, noscript, ol, p, param, plaintext, pre, script, section, select, spacer, style, tbody, textarea, tfoot, thead, title, tr, ul, and wbr.

Scoping

The following HTML elements introduce new scopes (page 595) for various parts of the parsing: applet, button, caption, html, marquee, object, table, td, th.

Formatting

The following HTML elements are those that end up in the list of active formatting elements (page 596): a, b, big, em, font, i, nobr, s, small, strike, strong, tt, and u.

Phrasing

All other elements found while parsing an HTML document.

The stack of open elements (page 594) is said to **have an element in scope** when the following algorithm terminates in a match state:

1. Initialize *node* to be the current node (page 594) (the bottommost node of the stack).
2. If *node* is the target node, terminate in a match state.
3. Otherwise, if *node* is one of the following elements, terminate in a failure state:
 - applet in the HTML namespace
 - caption in the HTML namespace
 - html in the HTML namespace
 - table in the HTML namespace
 - td in the HTML namespace
 - th in the HTML namespace
 - button in the HTML namespace
 - marquee in the HTML namespace
 - object in the HTML namespace
4. Otherwise, set *node* to the previous entry in the stack of open elements (page 594) and return to step 2. (This will never fail, since the loop will always terminate in the previous step if the top of the stack — an html element — is reached.)

The stack of open elements (page 594) is said to **have an element in table scope** when the following algorithm terminates in a match state:

1. Initialize *node* to be the current node (page 594) (the bottommost node of the stack).
2. If *node* is the target node, terminate in a match state.
3. Otherwise, if *node* is one of the following elements, terminate in a failure state:
 - html in the HTML namespace
 - table in the HTML namespace
4. Otherwise, set *node* to the previous entry in the stack of open elements (page 594) and return to step 2. (This will never fail, since the loop will always terminate in the previous step if the top of the stack — an html element — is reached.)

Nothing happens if at any time any of the elements in the stack of open elements (page 594) are moved to a new location in, or removed from, the Document tree. In particular, the stack is not changed in this situation. This can cause, amongst other strange effects, content to be appended to nodes that are no longer in the DOM.

Note: In some cases (namely, when closing misnested formatting elements (page 633)), the stack is manipulated in a random-access fashion.

8.2.3.3 The list of active formatting elements

Initially the **list of active formatting elements** is empty. It is used to handle mis-nested formatting element tags (page 595).

The list contains elements in the formatting (page 595) category, and scope markers. The scope markers are inserted when entering applet elements, buttons, object elements, marquees, table cells, and table captions, and are used to prevent formatting from "leaking" into applet elements, buttons, object elements, marquees, and tables.

When the steps below require the UA to **reconstruct the active formatting elements**, the UA must perform the following steps:

1. If there are no entries in the list of active formatting elements (page 596), then there is nothing to reconstruct; stop this algorithm.
2. If the last (most recently added) entry in the list of active formatting elements (page 596) is a marker, or if it is an element that is in the stack of open elements (page 594), then there is nothing to reconstruct; stop this algorithm.
3. Let *entry* be the last (most recently added) element in the list of active formatting elements (page 596).
4. If there are no entries before *entry* in the list of active formatting elements (page 596), then jump to step 8.
5. Let *entry* be the entry one earlier than *entry* in the list of active formatting elements (page 596).
6. If *entry* is neither a marker nor an element that is also in the stack of open elements (page 594), go to step 4.
7. Let *entry* be the element one later than *entry* in the list of active formatting elements (page 596).
8. Perform a shallow clone of the element *entry* to obtain *clone*. [DOM3CORE]
9. Append *clone* to the current node (page 594) and push it onto the stack of open elements (page 594) so that it is the new current node (page 594).
10. Replace the entry for *entry* in the list with an entry for *clone*.
11. If the entry for *clone* in the list of active formatting elements (page 596) is not the last entry in the list, return to step 7.

This has the effect of reopening all the formatting elements that were opened in the current body, cell, or caption (whichever is youngest) that haven't been explicitly closed.

Note: The way this specification is written, the list of active formatting elements (page 596) always consists of elements in chronological order with the least recently added element first and the most recently added element last (except for while steps 8 to 11 of the above algorithm are being executed, of course).

When the steps below require the UA to **clear the list of active formatting elements up to the last marker**, the UA must perform the following steps:

1. Let *entry* be the last (most recently added) entry in the list of active formatting elements (page 596).
2. Remove *entry* from the list of active formatting elements (page 596).
3. If *entry* was a marker, then stop the algorithm at this point. The list has been cleared up to the last marker.
4. Go to step 1.

8.2.3.4 The element pointers

Initially the **head element pointer** and the **form element pointer** are both null.

Once a head element has been parsed (whether implicitly or explicitly) the head element pointer (page 597) gets set to point to this node.

The form element pointer (page 597) points to the last form element that was opened and whose end tag has not yet been seen. It is used to make form controls associate with forms in the face of dramatically bad markup, for historical reasons.

8.2.3.5 The scripting state

The **scripting flag** is set to "enabled" if the Document with which the parser is associated was with script (page 428) when the parser was created, and "disabled" otherwise.

8.2.4 Tokenization

Implementations must act as if they used the following state machine to tokenise HTML. The state machine must start in the data state (page 598). Most states consume a single character, which may have various side-effects, and either switches the state machine to a new state to *reconsume* the same character, or switches it to a new state (to consume the next character), or repeats the same state (to consume the next character). Some states have more complicated behavior and can consume several characters before switching to another state.

The exact behavior of certain states depends on a **content model flag** that is set after certain tokens are emitted. The flag has several states: *PCDATA*, *RCDATA*, *CDATA*, and *PLAINTEXT*. Initially it must be in the *PCDATA* state. In the *RCDATA* and *CDATA* states, a further **escape flag** is used to control the behavior of the tokeniser. It is either true or false, and initially must be set to the false state. The insertion mode (page 592) and the stack of open elements (page 594) also affects tokenization.

The output of the tokenization step is a series of zero or more of the following tokens: DOCTYPE, start tag, end tag, comment, character, end-of-file. DOCTYPE tokens have a name, a public identifier, a system identifier, and a *force-quirks flag*. When a DOCTYPE token is created, its name, public identifier, and system identifier must be marked as missing (which is a distinct state from the empty string), and the *force-quirks flag* must be set to *off* (its other state is *on*). Start and end tag tokens have a tag name, a *self-closing flag*, and a list of

attributes, each of which has a name and a value. When a start or end tag token is created, its *self-closing flag* must be unset (its other state is that it be set), and its attributes list must be empty. Comment and character tokens have data.

When a token is emitted, it must immediately be handled by the tree construction (page 617) stage. The tree construction stage can affect the state of the content model flag (page 597), and can insert additional characters into the stream. (For example, the script element can result in scripts executing and using the dynamic markup insertion (page 99) APIs to insert characters into the stream being tokenised.)

When a start tag token is emitted with its *self-closing flag* set, if the flag is not **acknowledged** when it is processed by the tree construction stage, that is a parse error (page 583).

When an end tag token is emitted, the content model flag (page 597) must be switched to the PCDATA state.

When an end tag token is emitted with attributes, that is a parse error (page 583).

When an end tag token is emitted with its *self-closing flag* set, that is a parse error (page 583).

Before each step of the tokeniser, the user agent may check to see if either one of the scripts in the list of scripts that will execute as soon as possible (page 127) or the first script in the list of scripts that will execute asynchronously (page 127), has completed loading (page 125). If one has, then it must be executed (page 127) and removed from its list.

The tokeniser state machine consists of the states defined in the following subsections.

8.2.4.1 Data state

Consume the next input character (page 591):

↪ U+0026 AMPERSAND (&)

When the content model flag (page 597) is set to one of the PCDATA or RCDATA states and the escape flag (page 597) is false: switch to the character reference data state (page 599).

Otherwise: treat it as per the "anything else" entry below.

↪ U+002D HYPHEN-MINUS (-)

If the content model flag (page 597) is set to either the RCDATA state or the CDATA state, and the escape flag (page 597) is false, and there are at least three characters before this one in the input stream, and the last four characters in the input stream, including this one, are U+003C LESS-THAN SIGN, U+0021 EXCLAMATION MARK, U+002D HYPHEN-MINUS, and U+002D HYPHEN-MINUS ("<!--"), then set the escape flag (page 597) to true.

In any case, emit the input character as a character token. Stay in the data state (page 598).

↪ U+003C LESS-THAN SIGN (<)

When the content model flag (page 597) is set to the PCDATA state: switch to the tag open state (page 599).

When the content model flag (page 597) is set to either the RCDATA state or the CDATA state and the escape flag (page 597) is false: switch to the tag open state (page 599).

Otherwise: treat it as per the "anything else" entry below.

↪ **U+003E GREATER-THAN SIGN (>)**

If the content model flag (page 597) is set to either the RCDATA state or the CDATA state, and the escape flag (page 597) is true, and the last three characters in the input stream including this one are U+002D HYPHEN-MINUS, U+002D HYPHEN-MINUS, U+003E GREATER-THAN SIGN ("-->"), set the escape flag (page 597) to false.

In any case, emit the input character as a character token. Stay in the data state (page 598).

↪ **EOF**

Emit an end-of-file token.

↪ **Anything else**

Emit the input character as a character token. Stay in the data state (page 598).

8.2.4.2 Character reference data state

(*This cannot happen if the content model flag (page 597) is set to the CDATA state.*)

Attempt to consume a character reference (page 615), with no additional allowed character (page 615).

If nothing is returned, emit a U+0026 AMPERSAND character token.

Otherwise, emit the character token that was returned.

Finally, switch to the data state (page 598).

8.2.4.3 Tag open state

The behavior of this state depends on the content model flag (page 597).

If the content model flag (page 597) is set to the RCDATA or CDATA states

Consume the next input character (page 591). If it is a U+002F SOLIDUS (/) character, switch to the close tag open state (page 600). Otherwise, emit a U+003C LESS-THAN SIGN character token and reconsume the current input character in the data state (page 598).

If the content model flag (page 597) is set to the PCDATA state

Consume the next input character (page 591):

↪ **U+0021 EXCLAMATION MARK (!)**

Switch to the markup declaration open state (page 607).

↪ **U+002F SOLIDUS (/)**

Switch to the close tag open state (page 600).

↪ **U+0041 LATIN CAPITAL LETTER A through to U+005A LATIN CAPITAL LETTER Z**

Create a new start tag token, set its tag name to the lowercase version of the input character (add 0x0020 to the character's code point), then switch to the tag name state (page 601). (Don't emit the token yet; further details will be filled in before it is emitted.)

↪ **U+0061 LATIN SMALL LETTER A through to U+007A LATIN SMALL LETTER Z**

Create a new start tag token, set its tag name to the input character, then switch to the tag name state (page 601). (Don't emit the token yet; further details will be filled in before it is emitted.)

↪ **U+003E GREATER-THAN SIGN (>)**

Parse error (page 583). Emit a U+003C LESS-THAN SIGN character token and a U+003E GREATER-THAN SIGN character token. Switch to the data state (page 598).

↪ **U+003F QUESTION MARK (?)**

Parse error (page 583). Switch to the bogus comment state (page 606).

↪ **Anything else**

Parse error (page 583). Emit a U+003C LESS-THAN SIGN character token and reconsume the current input character in the data state (page 598).

8.2.4.4 Close tag open state

If the content model flag (page 597) is set to the RCDATA or CDATA states but no start tag token has ever been emitted by this instance of the tokeniser (fragment case (page 661)), or if the content model flag (page 597) is set to the RCDATA or CDATA states and the next few characters do not match the tag name of the last start tag token emitted (compared in an ASCII case insensitive manner), or if they do but they are not immediately followed by one of the following characters:

- U+0009 CHARACTER TABULATION
- U+000A LINE FEED (LF)
- U+000C FORM FEED (FF)
- U+0020 SPACE
- U+003E GREATER-THAN SIGN (>)
- U+002F SOLIDUS (/)
- EOF

...then emit a U+003C LESS-THAN SIGN character token, a U+002F SOLIDUS character token, and switch to the data state (page 598) to process the next input character (page 591).

Otherwise, if the content model flag (page 597) is set to the PCDATA state, or if the next few characters *do* match that tag name, consume the next input character (page 591):

↪ **U+0041 LATIN CAPITAL LETTER A through to U+005A LATIN CAPITAL LETTER Z**

Create a new end tag token, set its tag name to the lowercase version of the input character (add 0x0020 to the character's code point), then switch to the tag name state (page 601). (Don't emit the token yet; further details will be filled in before it is emitted.)

↪ **U+0061 LATIN SMALL LETTER A through to U+007A LATIN SMALL LETTER Z**

Create a new end tag token, set its tag name to the input character, then switch to the tag name state (page 601). (Don't emit the token yet; further details will be filled in before it is emitted.)

↪ **U+003E GREATER-THAN SIGN (>)**

Parse error (page 583). Switch to the data state (page 598).

↪ **EOF**

Parse error (page 583). Emit a U+003C LESS-THAN SIGN character token and a U+002F SOLIDUS character token. Reconsume the EOF character in the data state (page 598).

↪ **Anything else**

Parse error (page 583). Switch to the bogus comment state (page 606).

8.2.4.5 Tag name state

Consume the next input character (page 591):

↪ **U+0009 CHARACTER TABULATION**

↪ **U+000A LINE FEED (LF)**

↪ **U+000C FORM FEED (FF)**

↪ **U+0020 SPACE**

Switch to the before attribute name state (page 601).

↪ **U+002F SOLIDUS (/)**

Switch to the self-closing start tag state (page 606).

↪ **U+003E GREATER-THAN SIGN (>)**

Emit the current tag token. Switch to the data state (page 598).

↪ **U+0041 LATIN CAPITAL LETTER A through to U+005A LATIN CAPITAL LETTER Z**

Append the lowercase version of the current input character (add 0x0020 to the character's code point) to the current tag token's tag name. Stay in the tag name state (page 601).

↪ **EOF**

Parse error (page 583). Emit the current tag token. Reconsume the EOF character in the data state (page 598).

↪ **Anything else**

Append the current input character to the current tag token's tag name. Stay in the tag name state (page 601).

8.2.4.6 Before attribute name state

Consume the next input character (page 591):

↪ **U+0009 CHARACTER TABULATION**

↪ **U+000A LINE FEED (LF)**

↪ **U+000C FORM FEED (FF)**

↪ **U+0020 SPACE**

Stay in the before attribute name state (page 601).

↪ **U+002F SOLIDUS (/)**

Switch to the self-closing start tag state (page 606).

↪ **U+003E GREATER-THAN SIGN (>)**

Emit the current tag token. Switch to the data state (page 598).

↪ **U+0041 LATIN CAPITAL LETTER A through to U+005A LATIN CAPITAL LETTER Z**

Start a new attribute in the current tag token. Set that attribute's name to the lowercase version of the current input character (add 0x0020 to the character's code point), and its value to the empty string. Switch to the attribute name state (page 602).

↪ **U+0022 QUOTATION MARK ("")**

↪ **U+0027 APOSTROPHE ('')**

↪ **U+003D EQUALS SIGN (=)**

Parse error (page 583). Treat it as per the "anything else" entry below.

↪ **EOF**

Parse error (page 583). Emit the current tag token. Reconsume the EOF character in the data state (page 598).

↪ **Anything else**

Start a new attribute in the current tag token. Set that attribute's name to the current input character, and its value to the empty string. Switch to the attribute name state (page 602).

8.2.4.7 Attribute name state

Consume the next input character (page 591):

↪ **U+0009 CHARACTER TABULATION**

↪ **U+000A LINE FEED (LF)**

↪ **U+000C FORM FEED (FF)**

↪ **U+0020 SPACE**

Switch to the after attribute name state (page 603).

↪ **U+002F SOLIDUS (/)**

Switch to the self-closing start tag state (page 606).

↪ **U+003D EQUALS SIGN (=)**

Switch to the before attribute value state (page 604).

↪ **U+003E GREATER-THAN SIGN (>)**

Emit the current tag token. Switch to the data state (page 598).

↪ **U+0041 LATIN CAPITAL LETTER A through to U+005A LATIN CAPITAL LETTER Z**

Append the lowercase version of the current input character (add 0x0020 to the character's code point) to the current attribute's name. Stay in the attribute name state (page 602).

↪ **U+0022 QUOTATION MARK ("")**

↪ **U+0027 APOSTROPHE ('')**

Parse error (page 583). Treat it as per the "anything else" entry below.

↪ **EOF**

Parse error (page 583). Emit the current tag token. Reconsume the EOF character in the data state (page 598).

↪ **Anything else**

Append the current input character to the current attribute's name. Stay in the attribute name state (page 602).

When the user agent leaves the attribute name state (and before emitting the tag token, if appropriate), the complete attribute's name must be compared to the other attributes on the same token; if there is already an attribute on the token with the exact same name, then this is a parse error (page 583) and the new attribute must be dropped, along with the value that gets associated with it (if any).

8.2.4.8 After attribute name state

Consume the next input character (page 591):

↪ **U+0009 CHARACTER TABULATION**

↪ **U+000A LINE FEED (LF)**

↪ **U+000C FORM FEED (FF)**

↪ **U+0020 SPACE**

Stay in the after attribute name state (page 603).

↪ **U+002F SOLIDUS (/)**

Switch to the self-closing start tag state (page 606).

↪ **U+003D EQUALS SIGN (=)**

Switch to the before attribute value state (page 604).

↪ **U+003E GREATER-THAN SIGN (>)**

Emit the current tag token. Switch to the data state (page 598).

↪ **U+0041 LATIN CAPITAL LETTER A through to U+005A LATIN CAPITAL LETTER Z**

Start a new attribute in the current tag token. Set that attribute's name to the lowercase version of the current input character (add 0x0020 to the character's code point), and its value to the empty string. Switch to the attribute name state (page 602).

↪ **U+0022 QUOTATION MARK ("")**

↪ **U+0027 APOSTROPHE ('')**

Parse error (page 583). Treat it as per the "anything else" entry below.

↪ **EOF**

Parse error (page 583). Emit the current tag token. Reconsume the EOF character in the data state (page 598).

↪ **Anything else**

Start a new attribute in the current tag token. Set that attribute's name to the current input character, and its value to the empty string. Switch to the attribute name state (page 602).

8.2.4.9 Before attribute value state

Consume the next input character (page 591):

- ↪ **U+0009 CHARACTER TABULATION**
- ↪ **U+000A LINE FEED (LF)**
- ↪ **U+000C FORM FEED (FF)**
- ↪ **U+0020 SPACE**

Stay in the before attribute value state (page 604).

- ↪ **U+0022 QUOTATION MARK ("")**

Switch to the attribute value (double-quoted) state (page 604).

- ↪ **U+0026 AMPERSAND (&)**

Switch to the attribute value (unquoted) state (page 605) and reconsume this input character.

- ↪ **U+0027 APOSTROPHE ('')**

Switch to the attribute value (single-quoted) state (page 605).

- ↪ **U+003E GREATER-THAN SIGN (>)**

Parse error (page 583). Emit the current tag token. Switch to the data state (page 598).

- ↪ **U+003D EQUALS SIGN (=)**

Parse error (page 583). Treat it as per the "anything else" entry below.

- ↪ **EOF**

Parse error (page 583). Emit the current tag token. Reconsume the character in the data state (page 598).

- ↪ **Anything else**

Append the current input character to the current attribute's value. Switch to the attribute value (unquoted) state (page 605).

8.2.4.10 Attribute value (double-quoted) state

Consume the next input character (page 591):

- ↪ **U+0022 QUOTATION MARK ("")**

Switch to the after attribute value (quoted) state (page 606).

- ↪ **U+0026 AMPERSAND (&)**

Switch to the character reference in attribute value state (page 605), with the additional allowed character (page 615) being U+0022 QUOTATION MARK ("").

- ↪ **EOF**

Parse error (page 583). Emit the current tag token. Reconsume the character in the data state (page 598).

- ↪ **Anything else**

Append the current input character to the current attribute's value. Stay in the attribute value (double-quoted) state (page 604).

8.2.4.11 Attribute value (single-quoted) state

Consume the next input character (page 591):

↪ **U+0027 APOSTROPHE (')**

Switch to the after attribute value (quoted) state (page 606).

↪ **U+0026 AMPERSAND (&)**

Switch to the character reference in attribute value state (page 605), with the additional allowed character (page 615) being U+0027 APOSTROPHE (').

↪ **EOF**

Parse error (page 583). Emit the current tag token. Reconsume the character in the data state (page 598).

↪ **Anything else**

Append the current input character to the current attribute's value. Stay in the attribute value (single-quoted) state (page 605).

8.2.4.12 Attribute value (unquoted) state

Consume the next input character (page 591):

↪ **U+0009 CHARACTER TABULATION**

↪ **U+000A LINE FEED (LF)**

↪ **U+000C FORM FEED (FF)**

↪ **U+0020 SPACE**

Switch to the before attribute name state (page 601).

↪ **U+0026 AMPERSAND (&)**

Switch to the character reference in attribute value state (page 605), with no additional allowed character (page 615).

↪ **U+003E GREATER-THAN SIGN (>)**

Emit the current tag token. Switch to the data state (page 598).

↪ **U+0022 QUOTATION MARK ("")**

↪ **U+0027 APOSTROPHE ('')**

↪ **U+003D EQUALS SIGN (=)**

Parse error (page 583). Treat it as per the "anything else" entry below.

↪ **EOF**

Parse error (page 583). Emit the current tag token. Reconsume the character in the data state (page 598).

↪ **Anything else**

Append the current input character to the current attribute's value. Stay in the attribute value (unquoted) state (page 605).

8.2.4.13 Character reference in attribute value state

Attempt to consume a character reference (page 615).

If nothing is returned, append a U+0026 AMPERSAND character to the current attribute's value.

Otherwise, append the returned character token to the current attribute's value.

Finally, switch back to the attribute value state that you were in when were switched into this state.

8.2.4.14 After attribute value (quoted) state

Consume the next input character (page 591):

- ↪ **U+0009 CHARACTER TABULATION**
- ↪ **U+000A LINE FEED (LF)**
- ↪ **U+000C FORM FEED (FF)**
- ↪ **U+0020 SPACE**

Switch to the before attribute name state (page 601).

- ↪ **U+002F SOLIDUS (/)**

Switch to the self-closing start tag state (page 606).

- ↪ **U+003E GREATER-THAN SIGN (>)**

Emit the current tag token. Switch to the data state (page 598).

- ↪ **EOF**

Parse error (page 583). Emit the current tag token. Reconsume the EOF character in the data state (page 598).

- ↪ **Anything else**

Parse error (page 583). Reconsume the character in the before attribute name state (page 601).

8.2.4.15 Self-closing start tag state

Consume the next input character (page 591):

- ↪ **U+003E GREATER-THAN SIGN (>)**

Set the *self-closing flag* of the current tag token. Emit the current tag token. Switch to the data state (page 598).

- ↪ **EOF**

Parse error (page 583). Emit the current tag token. Reconsume the EOF character in the data state (page 598).

- ↪ **Anything else**

Parse error (page 583). Reconsume the character in the before attribute name state (page 601).

8.2.4.16 Bogus comment state

(This can only happen if the content model flag (page 597) is set to the PCDATA state.)

Consume every character up to and including the first U+003E GREATER-THAN SIGN character (>) or the end of the file (EOF), whichever comes first. Emit a comment token whose data is the concatenation of all the characters starting from and including the character that caused the state machine to switch into the bogus comment state, up to and including the character immediately before the last consumed character (i.e. up to the character just before the U+003E or EOF character). (If the comment was started by the end of the file (EOF), the token is empty.)

Switch to the data state (page 598).

If the end of the file was reached, reconsume the EOF character.

8.2.4.17 Markup declaration open state

(This can only happen if the content model flag (page 597) is set to the PCDATA state.)

If the next two characters are both U+002D HYPHEN-MINUS (-) characters, consume those two characters, create a comment token whose data is the empty string, and switch to the comment start state (page 607).

Otherwise, if the next seven characters are an ASCII case-insensitive (page 31) match for the word "DOCTYPE", then consume those characters and switch to the DOCTYPE state (page 609).

Otherwise, if the insertion mode (page 592) is "in foreign content (page 651)" and the current node (page 594) is not an element in the HTML namespace (page 658) and the next seven characters are an ASCII case-sensitive match for the string "[CDATA[" (the five uppercase letters "CDATA" with a U+005B LEFT SQUARE BRACKET character before and after), then consume those characters and switch to the CDATA section state (page 614) (which is unrelated to the content model flag (page 597)'s CDATA state).

Otherwise, this is a parse error (page 583). Switch to the bogus comment state (page 606). The next character that is consumed, if any, is the first character that will be in the comment.

8.2.4.18 Comment start state

Consume the next input character (page 591):

↪ **U+002D HYPHEN-MINUS (-)**

Switch to the comment start dash state (page 608).

↪ **U+003E GREATER-THAN SIGN (>)**

Parse error (page 583). Emit the comment token. Switch to the data state (page 598).

↪ **EOF**

Parse error (page 583). Emit the comment token. Reconsume the EOF character in the data state (page 598).

↪ **Anything else**

Append the input character to the comment token's data. Switch to the comment state (page 608).

8.2.4.19 Comment start dash state

Consume the next input character (page 591):

↪ **U+002D HYPHEN-MINUS (-)**

Switch to the comment end state (page 608)

↪ **U+003E GREATER-THAN SIGN (>)**

Parse error (page 583). Emit the comment token. Switch to the data state (page 598).

↪ **EOF**

Parse error (page 583). Emit the comment token. Reconsume the EOF character in the data state (page 598).

↪ **Anything else**

Append a U+002D HYPHEN-MINUS (-) character and the input character to the comment token's data. Switch to the comment state (page 608).

8.2.4.20 Comment state

Consume the next input character (page 591):

↪ **U+002D HYPHEN-MINUS (-)**

Switch to the comment end dash state (page 608)

↪ **EOF**

Parse error (page 583). Emit the comment token. Reconsume the EOF character in the data state (page 598).

↪ **Anything else**

Append the input character to the comment token's data. Stay in the comment state (page 608).

8.2.4.21 Comment end dash state

Consume the next input character (page 591):

↪ **U+002D HYPHEN-MINUS (-)**

Switch to the comment end state (page 608)

↪ **EOF**

Parse error (page 583). Emit the comment token. Reconsume the EOF character in the data state (page 598).

↪ **Anything else**

Append a U+002D HYPHEN-MINUS (-) character and the input character to the comment token's data. Switch to the comment state (page 608).

8.2.4.22 Comment end state

Consume the next input character (page 591):

↪ **U+003E GREATER-THAN SIGN (>)**

Emit the comment token. Switch to the data state (page 598).

↪ **U+002D HYPHEN-MINUS (-)**

Parse error (page 583). Append a U+002D HYPHEN-MINUS (-) character to the comment token's data. Stay in the comment end state (page 608).

↪ **EOF**

Parse error (page 583). Emit the comment token. Reconsume the EOF character in the data state (page 598).

↪ **Anything else**

Parse error (page 583). Append two U+002D HYPHEN-MINUS (-) characters and the input character to the comment token's data. Switch to the comment state (page 608).

8.2.4.23 DOCTYPE state

Consume the next input character (page 591):

↪ **U+0009 CHARACTER TABULATION**

↪ **U+000A LINE FEED (LF)**

↪ **U+000C FORM FEED (FF)**

↪ **U+0020 SPACE**

Switch to the before DOCTYPE name state (page 609).

↪ **Anything else**

Parse error (page 583). Reconsume the current character in the before DOCTYPE name state (page 609).

8.2.4.24 Before DOCTYPE name state

Consume the next input character (page 591):

↪ **U+0009 CHARACTER TABULATION**

↪ **U+000A LINE FEED (LF)**

↪ **U+000C FORM FEED (FF)**

↪ **U+0020 SPACE**

Stay in the before DOCTYPE name state (page 609).

↪ **U+003E GREATER-THAN SIGN (>)**

Parse error (page 583). Create a new DOCTYPE token. Set its *force-quirks flag* to *on*. Emit the token. Switch to the data state (page 598).

↪ **EOF**

Parse error (page 583). Create a new DOCTYPE token. Set its *force-quirks flag* to *on*. Emit the token. Reconsume the EOF character in the data state (page 598).

↪ **Anything else**

Create a new DOCTYPE token. Set the token's name to the current input character. Switch to the DOCTYPE name state (page 610).

8.2.4.25 DOCTYPE name state

First, consume the next input character (page 591):

- ↪ **U+0009 CHARACTER TABULATION**
- ↪ **U+000A LINE FEED (LF)**
- ↪ **U+000C FORM FEED (FF)**
- ↪ **U+0020 SPACE**

Switch to the after DOCTYPE name state (page 610).

- ↪ **U+003E GREATER-THAN SIGN (>)**

Emit the current DOCTYPE token. Switch to the data state (page 598).

- ↪ **EOF**

Parse error (page 583). Set the DOCTYPE token's *force-quirks flag* to *on*. Emit that DOCTYPE token. Reconsume the EOF character in the data state (page 598).

- ↪ **Anything else**

Append the current input character to the current DOCTYPE token's name. Stay in the DOCTYPE name state (page 610).

8.2.4.26 After DOCTYPE name state

Consume the next input character (page 591):

- ↪ **U+0009 CHARACTER TABULATION**
- ↪ **U+000A LINE FEED (LF)**
- ↪ **U+000C FORM FEED (FF)**
- ↪ **U+0020 SPACE**

Stay in the after DOCTYPE name state (page 610).

- ↪ **U+003E GREATER-THAN SIGN (>)**

Emit the current DOCTYPE token. Switch to the data state (page 598).

- ↪ **EOF**

Parse error (page 583). Set the DOCTYPE token's *force-quirks flag* to *on*. Emit that DOCTYPE token. Reconsume the EOF character in the data state (page 598).

- ↪ **Anything else**

If the next six characters are an ASCII case-insensitive (page 31) match for the word "PUBLIC", then consume those characters and switch to the before DOCTYPE public identifier state (page 611).

Otherwise, if the next six characters are an ASCII case-insensitive (page 31) match for the word "SYSTEM", then consume those characters and switch to the before DOCTYPE system identifier state (page 612).

Otherwise, this is the parse error (page 583). Set the DOCTYPE token's *force-quirks flag* to *on*. Switch to the bogus DOCTYPE state (page 614).

8.2.4.27 Before DOCTYPE public identifier state

Consume the next input character (page 591):

- ↪ **U+0009 CHARACTER TABULATION**
- ↪ **U+000A LINE FEED (LF)**
- ↪ **U+000C FORM FEED (FF)**
- ↪ **U+0020 SPACE**

Stay in the before DOCTYPE public identifier state (page 611).

- ↪ **U+0022 QUOTATION MARK ("")**

Set the DOCTYPE token's public identifier to the empty string (not missing), then switch to the DOCTYPE public identifier (double-quoted) state (page 611).

- ↪ **U+0027 APOSTROPHE ('')**

Set the DOCTYPE token's public identifier to the empty string (not missing), then switch to the DOCTYPE public identifier (single-quoted) state (page 611).

- ↪ **U+003E GREATER-THAN SIGN (>)**

Parse error (page 583). Set the DOCTYPE token's *force-quirks flag* to *on*. Emit that DOCTYPE token. Switch to the data state (page 598).

- ↪ **EOF**

Parse error (page 583). Set the DOCTYPE token's *force-quirks flag* to *on*. Emit that DOCTYPE token. Reconsume the EOF character in the data state (page 598).

- ↪ **Anything else**

Parse error (page 583). Set the DOCTYPE token's *force-quirks flag* to *on*. Switch to the bogus DOCTYPE state (page 614).

8.2.4.28 DOCTYPE public identifier (double-quoted) state

Consume the next input character (page 591):

- ↪ **U+0022 QUOTATION MARK ("")**

Switch to the after DOCTYPE public identifier state (page 612).

- ↪ **U+003E GREATER-THAN SIGN (>)**

Parse error (page 583). Set the DOCTYPE token's *force-quirks flag* to *on*. Emit that DOCTYPE token. Switch to the data state (page 598).

- ↪ **EOF**

Parse error (page 583). Set the DOCTYPE token's *force-quirks flag* to *on*. Emit that DOCTYPE token. Reconsume the EOF character in the data state (page 598).

- ↪ **Anything else**

Append the current input character to the current DOCTYPE token's public identifier. Stay in the DOCTYPE public identifier (double-quoted) state (page 611).

8.2.4.29 DOCTYPE public identifier (single-quoted) state

Consume the next input character (page 591):

↪ **U+0027 APOSTROPHE ('')**

Switch to the after DOCTYPE public identifier state (page 612).

↪ **U+003E GREATER-THAN SIGN (>)**

Parse error (page 583). Set the DOCTYPE token's *force-quirks flag* to *on*. Emit that DOCTYPE token. Switch to the data state (page 598).

↪ **EOF**

Parse error (page 583). Set the DOCTYPE token's *force-quirks flag* to *on*. Emit that DOCTYPE token. Reconsume the EOF character in the data state (page 598).

↪ **Anything else**

Append the current input character to the current DOCTYPE token's public identifier. Stay in the DOCTYPE public identifier (single-quoted) state (page 611).

8.2.4.30 After DOCTYPE public identifier state

Consume the next input character (page 591):

↪ **U+0009 CHARACTER TABULATION**

↪ **U+000A LINE FEED (LF)**

↪ **U+000C FORM FEED (FF)**

↪ **U+0020 SPACE**

Stay in the after DOCTYPE public identifier state (page 612).

↪ **U+0022 QUOTATION MARK ("")**

Set the DOCTYPE token's system identifier to the empty string (not missing), then switch to the DOCTYPE system identifier (double-quoted) state (page 613).

↪ **U+0027 APOSTROPHE ('')**

Set the DOCTYPE token's system identifier to the empty string (not missing), then switch to the DOCTYPE system identifier (single-quoted) state (page 613).

↪ **U+003E GREATER-THAN SIGN (>)**

Emit the current DOCTYPE token. Switch to the data state (page 598).

↪ **EOF**

Parse error (page 583). Set the DOCTYPE token's *force-quirks flag* to *on*. Emit that DOCTYPE token. Reconsume the EOF character in the data state (page 598).

↪ **Anything else**

Parse error (page 583). Set the DOCTYPE token's *force-quirks flag* to *on*. Switch to the bogus DOCTYPE state (page 614).

8.2.4.31 Before DOCTYPE system identifier state

Consume the next input character (page 591):

↪ **U+0009 CHARACTER TABULATION**

↪ **U+000A LINE FEED (LF)**

↪ **U+000C FORM FEED (FF)**

↪ **U+0020 SPACE**

Stay in the before DOCTYPE system identifier state (page 612).

↪ **U+0022 QUOTATION MARK ("")**

Set the DOCTYPE token's system identifier to the empty string (not missing), then switch to the DOCTYPE system identifier (double-quoted) state (page 613).

↪ **U+0027 APOSTROPHE ('')**

Set the DOCTYPE token's system identifier to the empty string (not missing), then switch to the DOCTYPE system identifier (single-quoted) state (page 613).

↪ **U+003E GREATER-THAN SIGN (>)**

Parse error (page 583). Set the DOCTYPE token's *force-quirks* flag to *on*. Emit that DOCTYPE token. Switch to the data state (page 598).

↪ **EOF**

Parse error (page 583). Set the DOCTYPE token's *force-quirks* flag to *on*. Emit that DOCTYPE token. Reconsume the EOF character in the data state (page 598).

↪ **Anything else**

Parse error (page 583). Set the DOCTYPE token's *force-quirks* flag to *on*. Switch to the bogus DOCTYPE state (page 614).

8.2.4.32 DOCTYPE system identifier (double-quoted) state

Consume the next input character (page 591):

↪ **U+0022 QUOTATION MARK ("")**

Switch to the after DOCTYPE system identifier state (page 614).

↪ **U+003E GREATER-THAN SIGN (>)**

Parse error (page 583). Set the DOCTYPE token's *force-quirks* flag to *on*. Emit that DOCTYPE token. Switch to the data state (page 598).

↪ **EOF**

Parse error (page 583). Set the DOCTYPE token's *force-quirks* flag to *on*. Emit that DOCTYPE token. Reconsume the EOF character in the data state (page 598).

↪ **Anything else**

Append the current input character to the current DOCTYPE token's system identifier. Stay in the DOCTYPE system identifier (double-quoted) state (page 613).

8.2.4.33 DOCTYPE system identifier (single-quoted) state

Consume the next input character (page 591):

↪ **U+0027 APOSTROPHE ('')**

Switch to the after DOCTYPE system identifier state (page 614).

↪ **U+003E GREATER-THAN SIGN (>)**

Parse error (page 583). Set the DOCTYPE token's *force-quirks flag* to *on*. Emit that DOCTYPE token. Switch to the data state (page 598).

↪ **EOF**

Parse error (page 583). Set the DOCTYPE token's *force-quirks flag* to *on*. Emit that DOCTYPE token. Reconsume the EOF character in the data state (page 598).

↪ **Anything else**

Append the current input character to the current DOCTYPE token's system identifier. Stay in the DOCTYPE system identifier (single-quoted) state (page 613).

8.2.4.34 After DOCTYPE system identifier state

Consume the next input character (page 591):

↪ **U+0009 CHARACTER TABULATION**

↪ **U+000A LINE FEED (LF)**

↪ **U+000C FORM FEED (FF)**

↪ **U+0020 SPACE**

Stay in the after DOCTYPE system identifier state (page 614).

↪ **U+003E GREATER-THAN SIGN (>)**

Emit the current DOCTYPE token. Switch to the data state (page 598).

↪ **EOF**

Parse error (page 583). Set the DOCTYPE token's *force-quirks flag* to *on*. Emit that DOCTYPE token. Reconsume the EOF character in the data state (page 598).

↪ **Anything else**

Parse error (page 583). Switch to the bogus DOCTYPE state (page 614). (This does *not* set the DOCTYPE token's *force-quirks flag* to *on*.)

8.2.4.35 Bogus DOCTYPE state

Consume the next input character (page 591):

↪ **U+003E GREATER-THAN SIGN (>)**

Emit the DOCTYPE token. Switch to the data state (page 598).

↪ **EOF**

Emit the DOCTYPE token. Reconsume the EOF character in the data state (page 598).

↪ **Anything else**

Stay in the bogus DOCTYPE state (page 614).

8.2.4.36 CDATA section state

(*This can only happen if the content model flag (page 597) is set to the PCDATA state, and is unrelated to the content model flag (page 597)'s CDATA state.*)

Consume every character up to the next occurrence of the three character sequence U+005D RIGHT SQUARE BRACKET U+005D RIGHT SQUARE BRACKET U+003E GREATER-THAN SIGN ([]>), or the end of the file (EOF), whichever comes first. Emit a series of text tokens consisting of all the characters consumed except the matching three character sequence at the end (if one was found before the end of the file).

Switch to the data state (page 598).

If the end of the file was reached, reconsume the EOF character.

8.2.4.37 Tokenizing character references

This section defines how to **consume a character reference**. This definition is used when parsing character references in text (page 599) and in attributes (page 605).

The behavior depends on the identity of the next character (the one immediately after the U+0026 AMPERSAND character):

- ↪ **U+0009 CHARACTER TABULATION**
- ↪ **U+000A LINE FEED (LF)**
- ↪ **U+000C FORM FEED (FF)**
- ↪ **U+0020 SPACE**
- ↪ **U+003C LESS-THAN SIGN**
- ↪ **U+0026 AMPERSAND**
- ↪ **EOF**
- ↪ **The additional allowed character, if there is one**

Not a character reference. No characters are consumed, and nothing is returned.
(This is not an error, either.)

- ↪ **U+0023 NUMBER SIGN (#)**

Consume the U+0023 NUMBER SIGN.

The behavior further depends on the character after the U+0023 NUMBER SIGN:

- ↪ **U+0078 LATIN SMALL LETTER X**
- ↪ **U+0058 LATIN CAPITAL LETTER X**

Consume the X.

Follow the steps below, but using the range of characters U+0030 DIGIT ZERO through to U+0039 DIGIT NINE, U+0061 LATIN SMALL LETTER A through to U+0066 LATIN SMALL LETTER F, and U+0041 LATIN CAPITAL LETTER A, through to U+0046 LATIN CAPITAL LETTER F (in other words, 0-9, A-F, a-f).

When it comes to interpreting the number, interpret it as a hexadecimal number.

- ↪ **Anything else**

Follow the steps below, but using the range of characters U+0030 DIGIT ZERO through to U+0039 DIGIT NINE (i.e. just 0-9).

When it comes to interpreting the number, interpret it as a decimal number.

Consume as many characters as match the range of characters given above.

If no characters match the range, then don't consume any characters (and unconsume the U+0023 NUMBER SIGN character and, if appropriate, the X character). This is a parse error (page 583); nothing is returned.

Otherwise, if the next character is a U+003B SEMICOLON, consume that too. If it isn't, there is a parse error (page 583).

If one or more characters match the range, then take them all and interpret the string of characters as a number (either hexadecimal or decimal as appropriate).

If that number is one of the numbers in the first column of the following table, then this is a parse error (page 583). Find the row with that number in the first column, and return a character token for the Unicode character given in the second column of that row.

Number	Unicode character
0x0D	U+000A LINE FEED (LF)
0x80	U+20AC EURO SIGN ('€')
0x81	U+FFFD REPLACEMENT CHARACTER
0x82	U+201A SINGLE LOW-9 QUOTATION MARK ('„')
0x83	U+0192 LATIN SMALL LETTER F WITH HOOK ('ƒ')
0x84	U+201E DOUBLE LOW-9 QUOTATION MARK ('„„')
0x85	U+2026 HORIZONTAL ELLIPSIS ('...')
0x86	U+2020 DAGGER ('†')
0x87	U+2021 DOUBLE DAGGER ('‡')
0x88	U+02C6 MODIFIER LETTER CIRCUMFLEX ACCENT ('^')
0x89	U+2030 PER MILLE SIGN ('‰')
0x8A	U+0160 LATIN CAPITAL LETTER S WITH CARON ('Š')
0x8B	U+2039 SINGLE LEFT-POINTING ANGLE QUOTATION MARK ('<')
0x8C	U+0152 LATIN CAPITAL LIGATURE OE ('Œ')
0x8D	U+FFFD REPLACEMENT CHARACTER
0x8E	U+017D LATIN CAPITAL LETTER Z WITH CARON ('Ž')
0x8F	U+FFFD REPLACEMENT CHARACTER
0x90	U+FFFD REPLACEMENT CHARACTER
0x91	U+2018 LEFT SINGLE QUOTATION MARK ('“')
0x92	U+2019 RIGHT SINGLE QUOTATION MARK ('”')
0x93	U+201C LEFT DOUBLE QUOTATION MARK ('““')
0x94	U+201D RIGHT DOUBLE QUOTATION MARK ('””')
0x95	U+2022 BULLET ('•')
0x96	U+2013 EN DASH ('–')
0x97	U+2014 EM DASH ('—')
0x98	U+02DC SMALL TILDE ('˜')
0x99	U+2122 TRADE MARK SIGN ('™')
0x9A	U+0161 LATIN SMALL LETTER S WITH CARON ('š')
0x9B	U+203A SINGLE RIGHT-POINTING ANGLE QUOTATION MARK ('>')
0x9C	U+0153 LATIN SMALL LIGATURE OE ('œ')
0x9D	U+FFFD REPLACEMENT CHARACTER

Number		Unicode character
0x9E	U+017E	LATIN SMALL LETTER Z WITH CARON ('Ž')
0x9F	U+0178	LATIN CAPITAL LETTER Y WITH DIAERESIS ('Ŷ')

Otherwise, if the number is in the range 0x0000 to 0x0008, U+000B, U+000E to 0x001F, 0x007F to 0x009F, 0xD800 to 0xDFFF, 0xFDD0 to 0xFDDF, or is one of 0xFFFFE, 0xFFFFF, 0xFFFFE, 0xFFFFF, 0x2FFFF, 0x3FFFF, 0x3FFFF, 0x4FFFF, 0x4FFFF, 0x5FFFF, 0x6FFFF, 0x6FFFF, 0x7FFFF, 0x7FFFF, 0x8FFFF, 0x8FFFF, 0x9FFFF, 0x9FFFF, 0xAFFE, 0xAFEE, 0xBFFE, 0xBFFF, 0xCFFE, 0xCFFF, 0xDFFF, 0xEFFF, 0xEFFF, 0xFFFFF, 0x10FFFF, or 0x10FFFF, or is higher than 0x10FFFF, then this is a parse error (page 583); return a character token for the U+FFFD REPLACEMENT CHARACTER character instead.

Otherwise, return a character token for the Unicode character whose code point is that number.

↳ Anything else

Consume the maximum number of characters possible, with the consumed characters matching one of the identifiers in the first column of the named character references (page 662) table (in a case-sensitive (page 31) manner).

If no match can be made, then this is a parse error (page 583). No characters are consumed, and nothing is returned.

If the last character matched is not a U+003B SEMICOLON (;), there is a parse error (page 583).

If the character reference is being consumed as part of an attribute (page 605), and the last character matched is not a U+003B SEMICOLON (;), and the next character is in the range U+0030 DIGIT ZERO to U+0039 DIGIT NINE, U+0041 LATIN CAPITAL LETTER A to U+005A LATIN CAPITAL LETTER Z, or U+0061 LATIN SMALL LETTER A to U+007A LATIN SMALL LETTER Z, then, for historical reasons, all the characters that were matched after the U+0026 AMPERSAND (&) must be unconsumed, and nothing is returned.

Otherwise, return a character token for the character corresponding to the character reference name (as given by the second column of the named character references (page 662) table).

If the markup contains I'm ¬it; I tell you, the character reference is parsed as "not", as in, I'm -it; I tell you. But if the markup was I'm ∉ I tell you, the character reference would be parsed as "notin;", resulting in I'm # I tell you.

8.2.5 Tree construction

The input to the tree construction stage is a sequence of tokens from the tokenization (page 597) stage. The tree construction stage is associated with a DOM Document object when a parser is created. The "output" of this stage consists of dynamically modifying or extending that document's DOM tree.

This specification does not define when an interactive user agent has to render the Document so that it is available to the user, or when it has to begin accepting user input.

As each token is emitted from the tokeniser, the user agent must process the token according to the rules given in the section corresponding to the current insertion mode (page 592).

When the steps below require the UA to **insert a character** into a node, if that node has a child immediately before where the character is to be inserted, and that child is a Text node, and that Text node was the last node that the parser inserted into the document, then the character must be appended to that Text node; otherwise, a new Text node whose data is just that character must be inserted in the appropriate place.

DOM mutation events must not fire for changes caused by the UA parsing the document. (Conceptually, the parser is not mutating the DOM, it is constructing it.) This includes the parsing of any content inserted using `document.write()` and `document.writeln()` calls. [DOM3EVENTS]

Note: Not all of the tag names mentioned below are conformant tag names in this specification; many are included to handle legacy content. They still form part of the algorithm that implementations are required to implement to claim conformance.

Note: The algorithm described below places no limit on the depth of the DOM tree generated, or on the length of tag names, attribute names, attribute values, text nodes, etc. While implementors are encouraged to avoid arbitrary limits, it is recognized that practical concerns (page 28) will likely force user agents to impose nesting depths.

8.2.5.1 Creating and inserting elements

When the steps below require the UA to **create an element for a token** in a particular namespace, the UA must create a node implementing the interface appropriate for the element type corresponding to the tag name of the token in the given namespace (as given in the specification that defines that element, e.g. for an element in the HTML namespace (page 658), this specification defines it to be the `HTMLAnchorElement` interface), with the tag name being the name of that element, with the node being in the given namespace, and with the attributes on the node being those given in the given token.

The interface appropriate for an element in the HTML namespace (page 658) that is not defined in this specification is `HTMLElement`. The interface appropriate for an element in another namespace that is not defined by that namespace's specification is `Element`.

When a resettable (page 314) element is created in this manner, its reset algorithm (page 376) must be invoked once the attributes are set. (This initializes the element's value (page 362) and checkedness (page 362) based on the element's attributes.)

When the steps below require the UA to **insert an HTML element** for a token, the UA must first create an element for the token (page 618) in the HTML namespace (page 658), and then append this node to the current node (page 594), and push it onto the stack of open elements (page 594) so that it is the new current node (page 594).

The steps below may also require that the UA insert an HTML element in a particular place, in which case the UA must follow the same steps except that it must insert or append the new node in the location specified instead of appending it to the current node (page 594). (This happens in particular during the parsing of tables with invalid content.)

When the steps below require the UA to **insert a foreign element** for a token, the UA must first create an element for the token (page 618) in the given namespace, and then append this node to the current node (page 594), and push it onto the stack of open elements (page 594) so that it is the new current node (page 594). If the newly created element has an `xmlns` attribute in the XMLNS namespace (page 658) whose value is not exactly the same as the element's namespace, that is a parse error (page 583).

When the steps below require the user agent to **adjust MathML attributes** for a token, then, if the token has an attribute named `definitionurl`, change its name to `definitionURL` (note the case difference).

When the steps below require the user agent to **adjust foreign attributes** for a token, then, if any of the attributes on the token match the strings given in the first column of the following table, let the attribute be a namespaced attribute, with the prefix being the string given in the corresponding cell in the second column, the local name being the string given in the corresponding cell in the third column, and the namespace being the namespace given in the corresponding cell in the fourth column. (This fixes the use of namespaced attributes, in particular `xml:lang`.)

Attribute name	Prefix	Local name	Namespace
<code>xlink:actuate</code>	<code>xlink</code>	<code>actuate</code>	XLink namespace (page 658)
<code>xlink:arcrole</code>	<code>xlink</code>	<code>arcrole</code>	XLink namespace (page 658)
<code>xlink:href</code>	<code>xlink</code>	<code>href</code>	XLink namespace (page 658)
<code>xlink:role</code>	<code>xlink</code>	<code>role</code>	XLink namespace (page 658)
<code>xlink:show</code>	<code>xlink</code>	<code>show</code>	XLink namespace (page 658)
<code>xlink:title</code>	<code>xlink</code>	<code>title</code>	XLink namespace (page 658)
<code>xlink:type</code>	<code>xlink</code>	<code>type</code>	XLink namespace (page 658)
<code>xml:base</code>	<code>xml</code>	<code>base</code>	XML namespace (page 658)
<code>xml:lang</code>	<code>xml</code>	<code>lang</code>	XML namespace (page 658)
<code>xml:space</code>	<code>xml</code>	<code>space</code>	XML namespace (page 658)
<code>xmlns</code>	(none)	<code>xmlns</code>	XMLNS namespace (page 658)
<code>xmlns:xlink</code>	<code>xmlns</code>	<code>xlink</code>	XMLNS namespace (page 658)

The **generic CDATA element parsing algorithm** and the **generic RCDATA element parsing algorithm** consist of the following steps. These algorithms are always invoked in response to a start tag token.

1. Insert an HTML element (page 618) for the token.
2. If the algorithm that was invoked is the generic CDATA element parsing algorithm (page 619), switch the tokeniser's content model flag (page 597) to the CDATA state; otherwise the algorithm invoked was the generic RCDATA element parsing algorithm (page 619), switch the tokeniser's content model flag (page 597) to the RCDATA state.

3. Let the original insertion mode (page 593) be the current insertion mode (page 592).
4. Then, switch the insertion mode (page 592) to "in CDATA/RCDATA (page 640)".

8.2.5.2 Closing elements that have implied end tags

When the steps below require the UA to **generate implied end tags**, then, while the current node (page 594) is a dd element, a dt element, an li element, an option element, an optgroup element, a p element, an rp element, or an rt element, the UA must pop the current node (page 594) off the stack of open elements (page 594).

If a step requires the UA to generate implied end tags but lists an element to exclude from the process, then the UA must perform the above steps as if that element was not in the above list.

8.2.5.3 Foster parenting

Foster parenting happens when content is misnested in tables.

When a node *node* is to be **foster parented**, the node *node* must be inserted into the *foster parent element* (page 620), and the current table (page 594) must be marked as **tainted**. (Once the current table (page 594) has been tainted (page 620), whitespace characters are inserted into the *foster parent element* (page 620) instead of the current node (page 594).)

The **foster parent element** is the parent element of the last table element in the stack of open elements (page 594), if there is a table element and it has such a parent element. If there is no table element in the stack of open elements (page 594) (fragment case (page 661)), then the *foster parent element* (page 620) is the first element in the stack of open elements (page 594) (the *html* element). Otherwise, if there is a table element in the stack of open elements (page 594), but the last table element in the stack of open elements (page 594) has no parent, or its parent node is not an element, then the *foster parent element* (page 620) is the element before the last table element in the stack of open elements (page 594).

If the *foster parent element* (page 620) is the parent element of the last table element in the stack of open elements (page 594), then *node* must be inserted immediately *before* the last table element in the stack of open elements (page 594) in the *foster parent element* (page 620); otherwise, *node* must be *appended* to the *foster parent element* (page 620).

8.2.5.4 The "initial" insertion mode

When the insertion mode (page 592) is "initial (page 620)", tokens must be handled as follows:

- ↪ **A character token that is one of one of U+0009 CHARACTER TABULATION, U+000A LINE FEED (LF), U+000C FORM FEED (FF), or U+0020 SPACE**

Ignore the token.

- ↪ **A comment token**

Append a Comment node to the Document object with the data attribute set to the data given in the comment token.

↪ A DOCTYPE token

If the DOCTYPE token's name is not an ASCII case-insensitive (page 31) match for the string "HTML", or if the token's public identifier is neither missing nor equal to the string "XSLT-compat", or if the token's system identifier is not missing, then there is a parse error (page 583). Conformance checkers may, instead of reporting this error, switch to a conformance checking mode for another language (e.g. based on the DOCTYPE token a conformance checker could recognize that the document is an HTML4-era document, and defer to an HTML4 conformance checker.)

Append a DocumentType node to the Document node, with the name attribute set to the name given in the DOCTYPE token; the publicId attribute set to the public identifier given in the DOCTYPE token, or the empty string if the public identifier was missing; the systemId attribute set to the system identifier given in the DOCTYPE token, or the empty string if the system identifier was missing; and the other attributes specific to DocumentType objects set to null and empty lists as appropriate. Associate the DocumentType node with the Document object so that it is returned as the value of the doctype attribute of the Document object.

Then, if the DOCTYPE token matches one of the conditions in the following list, then set the document to quirks mode (page 79):

- The *force-quirks flag* is set to *on*.
- The name is set to anything other than "HTML".
- The public identifier starts with: "+//Silmaril//dtd html Pro v0r11 19970101//"
- The public identifier starts with: "-//AdvaSoft Ltd//DTD HTML 3.0 asWedit + extensions//"
- The public identifier starts with: "-//AS//DTD HTML 3.0 asWedit + extensions//"
- The public identifier starts with: "-//IETF//DTD HTML 2.0 Level 1//"
- The public identifier starts with: "-//IETF//DTD HTML 2.0 Level 2//"
- The public identifier starts with: "-//IETF//DTD HTML 2.0 Strict Level 1//"
- The public identifier starts with: "-//IETF//DTD HTML 2.0 Strict Level 2//"
- The public identifier starts with: "-//IETF//DTD HTML 2.0 Strict//"
- The public identifier starts with: "-//IETF//DTD HTML 2.0//"
- The public identifier starts with: "-//IETF//DTD HTML 2.1E//"
- The public identifier starts with: "-//IETF//DTD HTML 3.0//"
- The public identifier starts with: "-//IETF//DTD HTML 3.2 Final//"
- The public identifier starts with: "-//IETF//DTD HTML 3.2//"
- The public identifier starts with: "-//IETF//DTD HTML 3//"
- The public identifier starts with: "-//IETF//DTD HTML Level 0//"
- The public identifier starts with: "-//IETF//DTD HTML Level 1//"
- The public identifier starts with: "-//IETF//DTD HTML Level 2//"
- The public identifier starts with: "-//IETF//DTD HTML Level 3//"
- The public identifier starts with: "-//IETF//DTD HTML Strict Level 0//"
- The public identifier starts with: "-//IETF//DTD HTML Strict Level 1//"
- The public identifier starts with: "-//IETF//DTD HTML Strict Level 2//"
- The public identifier starts with: "-//IETF//DTD HTML Strict Level 3//"
- The public identifier starts with: "-//IETF//DTD HTML Strict//"
- The public identifier starts with: "-//IETF//DTD HTML//"
- The public identifier starts with: "-//Metrius//DTD Metrius Presentational//"
- The public identifier starts with: "-//Microsoft//DTD Internet Explorer 2.0 HTML Strict//"
- The public identifier starts with: "-//Microsoft//DTD Internet Explorer 2.0 HTML//"

- The public identifier starts with: "-//Microsoft//DTD Internet Explorer 2.0 Tables//"
- The public identifier starts with: "-//Microsoft//DTD Internet Explorer 3.0 HTML Strict//"
- The public identifier starts with: "-//Microsoft//DTD Internet Explorer 3.0 HTML//"
- The public identifier starts with: "-//Microsoft//DTD Internet Explorer 3.0 Tables//"
- The public identifier starts with: "-//Netscape Comm. Corp./DTD HTML//"
- The public identifier starts with: "-//Netscape Comm. Corp./DTD Strict HTML//"
- The public identifier starts with: "-//O'Reilly and Associates//DTD HTML 2.0//"
- The public identifier starts with: "-//O'Reilly and Associates//DTD HTML Extended 1.0//"
- The public identifier starts with: "-//O'Reilly and Associates//DTD HTML Extended Relaxed 1.0//"
- The public identifier starts with: "-//SoftQuad Software//DTD HoTMetal PRO 6.0::19990601::extensions to HTML 4.0//"
- The public identifier starts with: "-//SoftQuad//DTD HoTMetal PRO 4.0::19971010::extensions to HTML 4.0//"
- The public identifier starts with: "-//Spyglass//DTD HTML 2.0 Extended//"
- The public identifier starts with: "-//SQ//DTD HTML 2.0 HoTMetal + extensions//"
- The public identifier starts with: "-//Sun Microsystems Corp./DTD HotJava HTML//"
- The public identifier starts with: "-//Sun Microsystems Corp./DTD HotJava Strict HTML//"
- The public identifier starts with: "-//W3C//DTD HTML 3 1995-03-24//"
- The public identifier starts with: "-//W3C//DTD HTML 3.2 Draft//"
- The public identifier starts with: "-//W3C//DTD HTML 3.2 Final//"
- The public identifier starts with: "-//W3C//DTD HTML 3.2//"
- The public identifier starts with: "-//W3C//DTD HTML 3.2S Draft//"
- The public identifier starts with: "-//W3C//DTD HTML 4.0 Frameset//"
- The public identifier starts with: "-//W3C//DTD HTML 4.0 Transitional//"
- The public identifier starts with: "-//W3C//DTD HTML Experimental 19960712//"
- The public identifier starts with: "-//W3C//DTD HTML Experimental 970421//"
- The public identifier starts with: "-//W3C//DTD W3 HTML//"
- The public identifier starts with: "-//W30//DTD W3 HTML 3.0//"
- The public identifier is set to: "-//W30//DTD W3 HTML Strict 3.0//EN//"
- The public identifier starts with: "-//WebTechs//DTD Mozilla HTML 2.0//"
- The public identifier starts with: "-//WebTechs//DTD Mozilla HTML//"
- The public identifier is set to: "-//W3C/DTD HTML 4.0 Transitional/EN"
- The public identifier is set to: "HTML"
- The system identifier is set to: "<http://www.ibm.com/data/dtd/v11/ibmxhtml1-transitional.dtd>"
- The system identifier is missing and the public identifier starts with: "-//W3C//DTD HTML 4.01 Frameset//"
- The system identifier is missing and the public identifier starts with: "-//W3C//DTD HTML 4.01 Transitional//"

Otherwise, if the DOCTYPE token matches one of the conditions in the following list, then set the document to limited quirks mode (page 79):

- The public identifier starts with: "-//W3C//DTD XHTML 1.0 Frameset//"
- The public identifier starts with: "-//W3C//DTD XHTML 1.0 Transitional//"
- The system identifier is not missing and the public identifier starts with: "-//W3C//DTD HTML 4.01 Frameset//"
- The system identifier is not missing and the public identifier starts with: "-//W3C//DTD HTML 4.01 Transitional//"

The name, system identifier, and public identifier strings must be compared to the values given in the lists above in an ASCII case-insensitive (page 31) manner. A system identifier whose value is the empty string is not considered missing for the purposes of the conditions above.

Then, switch the insertion mode (page 592) to "before html (page 623)".

↪ **Anything else**

Parse error (page 583).

Set the document to quirks mode (page 79).

Switch the insertion mode (page 592) to "before html (page 623)", then reprocess the current token.

8.2.5.5 The "before html" insertion mode

When the insertion mode (page 592) is "before html (page 623)", tokens must be handled as follows:

↪ **A DOCTYPE token**

Parse error (page 583). Ignore the token.

↪ **A comment token**

Append a Comment node to the Document object with the data attribute set to the data given in the comment token.

↪ **A character token that is one of one of U+0009 CHARACTER TABULATION, U+000A LINE FEED (LF), U+000C FORM FEED (FF), or U+0020 SPACE**

Ignore the token.

↪ **A start tag whose tag name is "html"**

Create an element for the token (page 618) in the HTML namespace (page 658). Append it to the Document object. Put this element in the stack of open elements (page 594).

If the token has an attribute "manifest", then resolve (page 55) the value of that attribute to an absolute URL (page 56), and if that is successful, run the application cache selection algorithm (page 460) with the resulting absolute URL (page 56). Otherwise, if there is no such attribute or resolving it fails, run the application cache selection algorithm (page 461) with no manifest.

Switch the insertion mode (page 592) to "before head (page 624)".

↪ **Anything else**

Create an HTMLElement node with the tag name html, in the HTML namespace (page 658). Append it to the Document object. Put this element in the stack of open elements (page 594).

Run the application cache selection algorithm (page 461) with no manifest.

Switch the insertion mode (page 592) to "before head (page 624)", then reprocess the current token.

**
**

Should probably make end tags be ignored, so that "</head><!-- --><html>" puts the comment before the root node (or should we?)

The root element can end up being removed from the Document object, e.g. by scripts; nothing in particular happens in such cases, content continues being appended to the nodes as described in the next section.

8.2.5.6 The "before head" insertion mode

When the insertion mode (page 592) is "before head (page 624)", tokens must be handled as follows:

↪ **A character token that is one of one of U+0009 CHARACTER TABULATION, U+000A LINE FEED (LF), U+000C FORM FEED (FF), or U+0020 SPACE**

Ignore the token.

↪ **A comment token**

Append a Comment node to the current node (page 594) with the data attribute set to the data given in the comment token.

↪ **A DOCTYPE token**

Parse error (page 583). Ignore the token.

↪ **A start tag whose tag name is "html"**

Process the token using the rules for (page 593) the "in body (page 628)" insertion mode (page 592).

↪ **A start tag whose tag name is "head"**

Insert an HTML element (page 618) for the token.

Set the head element pointer (page 597) to the newly created head element.

Switch the insertion mode (page 592) to "in head (page 625)".

↪ **An end tag whose tag name is one of: "head", "br"**

Act as if a start tag token with the tag name "head" and no attributes had been seen, then reprocess the current token.

↪ **Any other end tag**

Parse error (page 583). Ignore the token.

↪ **Anything else**

Act as if a start tag token with the tag name "head" and no attributes had been seen, then reprocess the current token.

Note: This will result in an empty head element being generated, with the current token being reprocessed in the "after head (page 627)" insertion mode (page 592).

8.2.5.7 The "in head" insertion mode

When the insertion mode (page 592) is "in head (page 625)", tokens must be handled as follows:

↪ **A character token that is one of one of U+0009 CHARACTER TABULATION, U+000A LINE FEED (LF), U+000C FORM FEED (FF), or U+0020 SPACE**

Insert the character (page 618) into the current node (page 594).

↪ **A comment token**

Append a Comment node to the current node (page 594) with the data attribute set to the data given in the comment token.

↪ **A DOCTYPE token**

Parse error (page 583). Ignore the token.

↪ **A start tag whose tag name is "html"**

Process the token using the rules for (page 593) the "in body (page 628)" insertion mode (page 592).

↪ **A start tag whose tag name is one of: "base", "command", "eventsouce", "link"**

Insert an HTML element (page 618) for the token. Immediately pop the current node (page 594) off the stack of open elements (page 594).

Acknowledge the token's *self-closing flag* (page 598), if it is set.

↪ **A start tag whose tag name is "meta"**

Insert an HTML element (page 618) for the token. Immediately pop the current node (page 594) off the stack of open elements (page 594).

Acknowledge the token's *self-closing flag* (page 598), if it is set.

If the element has a charset attribute, and its value is a supported encoding, and the confidence (page 585) is currently *tentative*, then change the encoding (page 592) to the encoding given by the value of the charset attribute.

Otherwise, if the element has a content attribute, and applying the algorithm for extracting an encoding from a Content-Type (page 60) to its value returns a supported encoding *encoding*, and the confidence (page 585) is currently *tentative*, then change the encoding (page 592) to the encoding *encoding*.

↪ **A start tag whose tag name is "title"**

Follow the generic RCDATA element parsing algorithm (page 619).

↪ **A start tag whose tag name is "noscript", if the scripting flag (page 597) is enabled**

↪ **A start tag whose tag name is one of: "noframes", "style"**

Follow the generic CDATA element parsing algorithm (page 619).

↪ **A start tag whose tag name is "noscript", if the scripting flag (page 597) is disabled**

Insert an HTML element (page 618) for the token.

Switch the insertion mode (page 592) to "in head noscript (page 626)".

↪ **A start tag whose tag name is "script"**

1. Create an element for the token (page 618) in the HTML namespace (page 658).
2. Mark the element as being "parser-inserted" (page 124).

Note: This ensures that, if the script is external, any `document.write()` calls in the script will execute in-line, instead of blowing the document away, as would happen in most other cases. It also prevents the script from executing until the end tag is seen.

3. If the parser was originally created for the HTML fragment parsing algorithm (page 661), then mark the script element as "already executed" (page 124). (fragment case (page 661))
4. Append the new element to the current node (page 594).
5. Switch the tokeniser's content model flag (page 597) to the CDATA state.
6. Let the original insertion mode (page 593) be the current insertion mode (page 592).
7. Switch the insertion mode (page 592) to "in CDATA/RCDATA (page 640)".

↪ **An end tag whose tag name is "head"**

Pop the current node (page 594) (which will be the head element) off the stack of open elements (page 594).

Switch the insertion mode (page 592) to "after head (page 627)".

↪ **An end tag whose tag name is "br"**

Act as described in the "anything else" entry below.

↪ **A start tag whose tag name is "head"**

↪ **Any other end tag**

Parse error (page 583). Ignore the token.

↪ **Anything else**

Act as if an end tag token with the tag name "head" had been seen, and reprocess the current token.

**

In certain UAs, some elements don't trigger the "in body" mode straight away, but instead get put into the head. Do we want to copy that?

**

8.2.5.8 The "in head noscript" insertion mode

When the insertion mode (page 592) is "in head noscript (page 626)", tokens must be handled as follows:

↪ **A DOCTYPE token**

Parse error (page 583). Ignore the token.

↪ **A start tag whose tag name is "html"**

Process the token using the rules for (page 593) the "in body (page 628)" insertion mode (page 592).

↪ **An end tag whose tag name is "noscript"**

Pop the current node (page 594) (which will be a noscript element) from the stack of open elements (page 594); the new current node (page 594) will be a head element.

Switch the insertion mode (page 592) to "in head (page 625)".

↪ **A character token that is one of one of U+0009 CHARACTER TABULATION, U+000A LINE FEED (LF), U+000C FORM FEED (FF), or U+0020 SPACE**

↪ **A comment token**

↪ **A start tag whose tag name is one of: "link", "meta", "noframes", "style"**

Process the token using the rules for (page 593) the "in head (page 625)" insertion mode (page 592).

↪ **An end tag whose tag name is one of: "br"**

Act as described in the "anything else" entry below.

↪ **A start tag whose tag name is one of: "head", "noscript"**

↪ **Any other end tag**

Parse error (page 583). Ignore the token.

↪ **Anything else**

Parse error (page 583). Act as if an end tag with the tag name "noscript" had been seen and reprocess the current token.

8.2.5.9 The "after head" insertion mode

When the insertion mode (page 592) is "after head (page 627)", tokens must be handled as follows:

↪ **A character token that is one of one of U+0009 CHARACTER TABULATION, U+000A LINE FEED (LF), U+000C FORM FEED (FF), or U+0020 SPACE**

Insert the character (page 618) into the current node (page 594).

↪ **A comment token**

Append a Comment node to the current node (page 594) with the data attribute set to the data given in the comment token.

↪ **A DOCTYPE token**

Parse error (page 583). Ignore the token.

↪ **A start tag whose tag name is "html"**

Process the token using the rules for (page 593) the "in body (page 628)" insertion mode (page 592).

↪ **A start tag whose tag name is "body"**

Insert an HTML element (page 618) for the token.

Switch the insertion mode (page 592) to "in body (page 628)".

↪ **A start tag whose tag name is "frameset"**

Insert an HTML element (page 618) for the token.

Switch the insertion mode (page 592) to "in frameset (page 653)".

↪ **A start tag token whose tag name is one of: "base", "link", "meta", "noframes", "script", "style", "title"**

Parse error (page 583).

Push the node pointed to by the head element pointer (page 597) onto the stack of open elements (page 594).

Process the token using the rules for (page 593) the "in head (page 625)" insertion mode (page 592).

Pop the current node (page 594) (which will be the node pointed to by the head element pointer (page 597)) off the stack of open elements (page 594).

↪ **A start tag whose tag name is "head"**

↪ **Any end tag**

Parse error (page 583). Ignore the token.

↪ **Anything else**

Act as if a start tag token with the tag name "body" and no attributes had been seen, and then reprocess the current token.

8.2.5.10 The "in body" insertion mode

When the insertion mode (page 592) is "in body (page 628)", tokens must be handled as follows:

↪ **A character token**

Reconstruct the active formatting elements (page 596), if any.

Insert the token's character (page 618) into the current node (page 594).

↪ **A comment token**

Append a Comment node to the current node (page 594) with the data attribute set to the data given in the comment token.

↪ **A DOCTYPE token**

Parse error (page 583). Ignore the token.

↪ **A start tag whose tag name is "html"**

Parse error (page 583). For each attribute on the token, check to see if the attribute is already present on the top element of the stack of open elements (page 594). If it is not, add the attribute and its corresponding value to that element.

↪ **A start tag token whose tag name is one of: "base", "command", "eventsource", "link", "meta", "noframes", "script", "style", "title"**

Process the token using the rules for (page 593) the "in head (page 625)" insertion mode (page 592).

↪ **A start tag whose tag name is "body"**

Parse error (page 583).

If the second element on the stack of open elements (page 594) is not a body element, or, if the stack of open elements (page 594) has only one node on it, then ignore the token. (fragment case (page 661))

Otherwise, for each attribute on the token, check to see if the attribute is already present on the body element (the second element) on the stack of open elements (page 594). If it is not, add the attribute and its corresponding value to that element.

↪ **An end-of-file token**

If there is a node in the stack of open elements (page 594) that is not either a dd element, a dt element, an li element, a p element, a tbody element, a td element, a tfoot element, a th element, a thead element, a tr element, the body element, or the html element, then this is a parse error (page 583).

Stop parsing (page 656).

↪ **An end tag whose tag name is "body"**

If the stack of open elements (page 594) does not have a body element in scope (page 595), this is a parse error (page 583); ignore the token.

Otherwise, if there is a node in the stack of open elements (page 594) that is not either a dd element, a dt element, an li element, a p element, a tbody element, a td element, a tfoot element, a th element, a thead element, a tr element, the body element, or the html element, then this is a parse error (page 583).

Switch the insertion mode (page 592) to "after body (page 653)". Otherwise, ignore the token.

↪ **An end tag whose tag name is "html"**

Act as if an end tag with tag name "body" had been seen, then, if that token wasn't ignored, reprocess the current token.

Note: *The fake end tag token here can only be ignored in the fragment case (page 661).*

↪ **A start tag whose tag name is one of: "address", "article", "aside", "blockquote", "center", "datagrid", "details", "dialog", "dir", "div", "dl", "figure", "footer", "h1", "h2", "h3", "h4", "h5", "h6", "header", "menu", "nav", "ol", "p", "section", "ul"**

If the stack of open elements (page 594) has a p element in scope (page 595), then act as if an end tag with the tag name "p" had been seen.

Insert an HTML element (page 618) for the token.

↪ A start tag whose tag name is one of: "pre", "listing"

If the stack of open elements (page 594) has a p element in scope (page 595), then act as if an end tag with the tag name "p" had been seen.

Insert an HTML element (page 618) for the token.

If the next token is a U+000A LINE FEED (LF) character token, then ignore that token and move on to the next one. (Newlines at the start of pre blocks are ignored as an authoring convenience.)

↪ A start tag whose tag name is "form"

If the form element pointer (page 597) is not null, then this is a parse error (page 583); ignore the token.

Otherwise:

If the stack of open elements (page 594) has a p element in scope (page 595), then act as if an end tag with the tag name "p" had been seen.

Insert an HTML element (page 618) for the token, and set the form element pointer to point to the element created.

↪ A start tag whose tag name is "fieldset"

If the stack of open elements (page 594) has a p element in scope (page 595), then act as if an end tag with the tag name "p" had been seen.

Insert an HTML element (page 618) for the token.

If the form element pointer (page 597) is not null and the newly created element doesn't have a form attribute, then associate (page 361) the newly created fieldset element with the form element pointed to by the form element pointer (page 597).

↪ A start tag whose tag name is "li"

Run the following algorithm:

1. Initialize *node* to be the current node (page 594) (the bottommost node of the stack).
2. If *node* is an li element, then act as if an end tag with the tag name "li" had been seen, then jump to the last step.
3. If *node* is not in the formatting (page 595) category, and is not in the phrasing (page 595) category, and is not an address, div, or p element, then jump to the last step.
4. Otherwise, set *node* to the previous entry in the stack of open elements (page 594) and return to step 2.
5. If the stack of open elements (page 594) has a p element in scope (page 595), then act as if an end tag with the tag name "p" had been seen.

Finally, insert an HTML element (page 618) for the token.

↪ A start tag whose tag name is one of: "dd", "dt"

Run the following algorithm:

1. Initialize *node* to be the current node (page 594) (the bottommost node of the stack).
2. If *node* is a dd or dt element, then act as if an end tag with the same tag name as *node* had been seen, then jump to the last step.
3. If *node* is not in the formatting (page 595) category, and is not in the phrasing (page 595) category, and is not an address, div, or p element, then jump to the last step.
4. Otherwise, set *node* to the previous entry in the stack of open elements (page 594) and return to step 2.
5. If the stack of open elements (page 594) has a p element in scope (page 595), then act as if an end tag with the tag name "p" had been seen.

Finally, insert an HTML element (page 618) for the token.

↪ A start tag whose tag name is "plaintext"

If the stack of open elements (page 594) has a p element in scope (page 595), then act as if an end tag with the tag name "p" had been seen.

Insert an HTML element (page 618) for the token.

Switch the content model flag (page 597) to the PLAINTEXT state.

Note: Once a start tag with the tag name "plaintext" has been seen, that will be the last token ever seen other than character tokens (and the end-of-file token), because there is no way to switch the content model flag (page 597) out of the PLAINTEXT state.

↪ An end tag whose tag name is one of: "address", "article", "aside", "blockquote", "center", "datagrid", "details", "dialog", "dir", "div", "dl", "fieldset", "figure", "footer", "header", "listing", "menu", "nav", "ol", "pre", "section", "ul"

If the stack of open elements (page 594) does not have an element in scope (page 595) with the same tag name as that of the token, then this is a parse error (page 583); ignore the token.

Otherwise, run these steps:

1. Generate implied end tags (page 620).
2. If the current node (page 594) is not an element with the same tag name as that of the token, then this is a parse error (page 583).
3. Pop elements from the stack of open elements (page 594) until an element with the same tag name as the token has been popped from the stack.

↪ An end tag whose tag name is "form"

Set the form element pointer (page 597) to null.

If the stack of open elements (page 594) does not have an element in scope (page 595) with the same tag name as that of the token, then this is a parse error (page 583); ignore the token.

Otherwise, run these steps:

1. Generate implied end tags (page 620).
2. If the current node (page 594) is not an element with the same tag name as that of the token, then this is a parse error (page 583).
3. Pop elements from the stack of open elements (page 594) until an element with the same tag name as the token has been popped from the stack.

↪ **An end tag whose tag name is "p"**

If the stack of open elements (page 594) does not have an element in scope (page 595) with the same tag name as that of the token, then this is a parse error (page 583); act as if a start tag with the tag name p had been seen, then reprocess the current token.

Otherwise, run these steps:

1. Generate implied end tags (page 620), except for elements with the same tag name as the token.
2. If the current node (page 594) is not an element with the same tag name as that of the token, then this is a parse error (page 583).
3. Pop elements from the stack of open elements (page 594) until an element with the same tag name as the token has been popped from the stack.

↪ **An end tag whose tag name is one of: "dd", "dt", "li"**

If the stack of open elements (page 594) does not have an element in scope (page 595) with the same tag name as that of the token, then this is a parse error (page 583); ignore the token.

Otherwise, run these steps:

1. Generate implied end tags (page 620), except for elements with the same tag name as the token.
2. If the current node (page 594) is not an element with the same tag name as that of the token, then this is a parse error (page 583).
3. Pop elements from the stack of open elements (page 594) until an element with the same tag name as the token has been popped from the stack.

↪ **An end tag whose tag name is one of: "h1", "h2", "h3", "h4", "h5", "h6"**

If the stack of open elements (page 594) does not have an element in scope (page 595) whose tag name is one of "h1", "h2", "h3", "h4", "h5", or "h6", then this is a parse error (page 583); ignore the token.

Otherwise, run these steps:

1. Generate implied end tags (page 620).

2. If the current node (page 594) is not an element with the same tag name as that of the token, then this is a parse error (page 583).
3. Pop elements from the stack of open elements (page 594) until an element whose tag name is one of "h1", "h2", "h3", "h4", "h5", or "h6" has been popped from the stack.

↪ **An end tag whose tag name is "sarcasm"**

Take a deep breath, then act as described in the "any other end tag" entry below.

↪ **A start tag whose tag name is "a"**

If the list of active formatting elements (page 596) contains an element whose tag name is "a" between the end of the list and the last marker on the list (or the start of the list if there is no marker on the list), then this is a parse error (page 583); act as if an end tag with the tag name "a" had been seen, then remove that element from the list of active formatting elements (page 596) and the stack of open elements (page 594) if the end tag didn't already remove it (it might not have if the element is not in table scope (page 595)).

In the non-conforming stream

`a<table>b</table>x`, the first a element would be closed upon seeing the second one, and the "x" character would be inside a link to "b", not to "a". This is despite the fact that the outer a element is not in table scope (meaning that a regular end tag at the start of the table wouldn't close the outer a element).

Reconstruct the active formatting elements (page 596), if any.

Insert an HTML element (page 618) for the token. Add that element to the list of active formatting elements (page 596).

↪ **A start tag whose tag name is one of: "b", "big", "em", "font", "i", "s", "small", "strike", "strong", "tt", "u"**

Reconstruct the active formatting elements (page 596), if any.

Insert an HTML element (page 618) for the token. Add that element to the list of active formatting elements (page 596).

↪ **A start tag whose tag name is "nobr"**

Reconstruct the active formatting elements (page 596), if any.

If the stack of open elements (page 594) has a nobr element in scope (page 595), then this is a parse error (page 583); act as if an end tag with the tag name "nobr" had been seen, then once again reconstruct the active formatting elements (page 596), if any.

Insert an HTML element (page 618) for the token. Add that element to the list of active formatting elements (page 596).

↪ **An end tag whose tag name is one of: "a", "b", "big", "em", "font", "i", "nobr", "s", "small", "strike", "strong", "tt", "u"**

Follow these steps:

- Let the *formatting element* be the last element in the list of active formatting elements (page 596) that:
 - is between the end of the list and the last scope marker in the list, if any, or the start of the list otherwise, and
 - has the same tag name as the token.

If there is no such node, or, if that node is also in the stack of open elements (page 594) but the element is not in scope (page 595), then this is a parse error (page 583); ignore the token, and abort these steps.

Otherwise, if there is such a node, but that node is not in the stack of open elements (page 594), then this is a parse error (page 583); remove the element from the list, and abort these steps.

Otherwise, there is a *formatting element* and that element is in the stack (page 594) and is in scope (page 595). If the element is not the current node (page 594), this is a parse error (page 583). In any case, proceed with the algorithm as written in the following steps.

- Let the *furthest block* be the topmost node in the stack of open elements (page 594) that is lower in the stack than the *formatting element*, and is not an element in the phrasing (page 595) or formatting (page 595) categories. There might not be one.
- If there is no *furthest block*, then the UA must skip the subsequent steps and instead just pop all the nodes from the bottom of the stack of open elements (page 594), from the current node (page 594) up to and including the *formatting element*, and remove the *formatting element* from the list of active formatting elements (page 596).
- Let the *common ancestor* be the element immediately above the *formatting element* in the stack of open elements (page 594).
- If the *furthest block* has a parent node, then remove the *furthest block* from its parent node.
- Let a bookmark note the position of the *formatting element* in the list of active formatting elements (page 596) relative to the elements on either side of it in the list.
- Let *node* and *last node* be the *furthest block*. Follow these steps:
 - Let *node* be the element immediately above *node* in the stack of open elements (page 594).
 - If *node* is not in the list of active formatting elements (page 596), then remove *node* from the stack of open elements (page 594) and then go back to step 1.
 - Otherwise, if *node* is the *formatting element*, then go to the next step in the overall algorithm.

4. Otherwise, if *last node* is the *furthest block*, then move the aforementioned bookmark to be immediately after the *node* in the list of active formatting elements (page 596).
 5. If *node* has any children, perform a shallow clone of *node*, replace the entry for *node* in the list of active formatting elements (page 596) with an entry for the clone, replace the entry for *node* in the stack of open elements (page 594) with an entry for the clone, and let *node* be the clone.
 6. Insert *last node* into *node*, first removing it from its previous parent node if any.
 7. Let *last node* be *node*.
 8. Return to step 1 of this inner set of steps.
 8. If the *common ancestor* node is a *table*, *tbody*, *tfoot*, *thead*, or *tr* element, then, foster parent (page 620) whatever *last node* ended up being in the previous step.
- Otherwise, append whatever *last node* ended up being in the previous step to the *common ancestor* node, first removing it from its previous parent node if any.
9. Perform a shallow clone of the *formatting element*.
 10. Take all of the child nodes of the *furthest block* and append them to the clone created in the last step.
 11. Append that clone to the *furthest block*.
 12. Remove the *formatting element* from the list of active formatting elements (page 596), and insert the clone into the list of active formatting elements (page 596) at the position of the aforementioned bookmark.
 13. Remove the *formatting element* from the stack of open elements (page 594), and insert the clone into the stack of open elements (page 594) immediately below the position of the *furthest block* in that stack.
 14. Jump back to step 1 in this series of steps.

Note: The way these steps are defined, only elements in the formatting (page 595) category ever get cloned by this algorithm.

Note: Because of the way this algorithm causes elements to change parents, it has been dubbed the "adoption agency algorithm" (in contrast with other possibly algorithms for dealing with misnested content, which included the "incest algorithm", the "secret affair algorithm", and the "Heisenberg algorithm").

↪ **A start tag whose tag name is "button"**

If the stack of open elements (page 594) has a button element in scope (page 595), then this is a parse error (page 583); act as if an end tag with the tag name "button" had been seen, then reprocess the token.

Otherwise:

Reconstruct the active formatting elements (page 596), if any.

Insert an HTML element (page 618) for the token.

If the form element pointer (page 597) is not null and the newly created element doesn't have a form attribute, then associate (page 361) the button element with the form element pointed to by the form element pointer (page 597).

Insert a marker at the end of the list of active formatting elements (page 596).

↪ **A start tag token whose tag name is one of: "applet", "marquee", "object"**

Reconstruct the active formatting elements (page 596), if any.

Insert an HTML element (page 618) for the token.

Insert a marker at the end of the list of active formatting elements (page 596).

↪ **An end tag token whose tag name is one of: "applet", "button", "marquee", "object"**

If the stack of open elements (page 594) does not have an element in scope (page 595) with the same tag name as that of the token, then this is a parse error (page 583); ignore the token.

Otherwise, run these steps:

1. Generate implied end tags (page 620).
2. If the current node (page 594) is not an element with the same tag name as that of the token, then this is a parse error (page 583).
3. Pop elements from the stack of open elements (page 594) until an element with the same tag name as the token has been popped from the stack.
4. Clear the list of active formatting elements up to the last marker (page 597).

↪ **A start tag whose tag name is "xmp"**

Reconstruct the active formatting elements (page 596), if any.

Follow the generic CDATA element parsing algorithm (page 619).

↪ **A start tag whose tag name is "table"**

If the stack of open elements (page 594) has a p element in scope (page 595), then act as if an end tag with the tag name "p" had been seen.

Insert an HTML element (page 618) for the token.

Switch the insertion mode (page 592) to "in table (page 642)".

↪ **A start tag whose tag name is one of: "area", "basefont", "bgsound", "br", "embed", "img", "spacer", "wbr"**

Reconstruct the active formatting elements (page 596), if any.

Insert an HTML element (page 618) for the token. Immediately pop the current node (page 594) off the stack of open elements (page 594).

Acknowledge the token's *self-closing flag* (page 598), if it is set.

↪ **A start tag whose tag name is one of: "param", "source"**

Insert an HTML element (page 618) for the token. Immediately pop the current node (page 594) off the stack of open elements (page 594).

Acknowledge the token's *self-closing flag* (page 598), if it is set.

↪ **A start tag whose tag name is "hr"**

If the stack of open elements (page 594) has a p element in scope (page 595), then act as if an end tag with the tag name "p" had been seen.

Insert an HTML element (page 618) for the token. Immediately pop the current node (page 594) off the stack of open elements (page 594).

Acknowledge the token's *self-closing flag* (page 598), if it is set.

↪ **A start tag whose tag name is "image"**

Parse error (page 583). Change the token's tag name to "img" and reprocess it. (Don't ask.)

↪ **A start tag whose tag name is "input"**

Reconstruct the active formatting elements (page 596), if any.

Insert an HTML element (page 618) for the token. Immediately pop the current node (page 594) off the stack of open elements (page 594).

Acknowledge the token's *self-closing flag* (page 598), if it is set.

If the form element pointer (page 597) is not null and the newly created element doesn't have a form attribute, then associate (page 361) the newly created input element with the form element pointed to by the form element pointer (page 597).

↪ **A start tag whose tag name is "label"**

Reconstruct the active formatting elements (page 596), if any.

Insert an HTML element (page 618) for the token.

If the form element pointer (page 597) is not null and the newly created element doesn't have a form attribute, then associate (page 361) the newly created label element with the form element pointed to by the form element pointer (page 597).

↪ **A start tag whose tag name is "isindex"**

Parse error (page 583).

If the form element pointer (page 597) is not null, then ignore the token.

Otherwise:

Acknowledge the token's *self-closing flag* (page 598), if it is set.

Act as if a start tag token with the tag name "form" had been seen.

If the token has an attribute called "action", set the action attribute on the resulting form element to the value of the "action" attribute of the token.

Act as if a start tag token with the tag name "hr" had been seen.

Act as if a start tag token with the tag name "p" had been seen.

Act as if a start tag token with the tag name "label" had been seen.

Act as if a stream of character tokens had been seen (see below for what they should say).

Act as if a start tag token with the tag name "input" had been seen, with all the attributes from the "isindex" token except "name", "action", and "prompt". Set the name attribute of the resulting input element to the value "isindex".

Act as if a stream of character tokens had been seen (see below for what they should say).

Act as if an end tag token with the tag name "label" had been seen.

Act as if an end tag token with the tag name "p" had been seen.

Act as if a start tag token with the tag name "hr" had been seen.

Act as if an end tag token with the tag name "form" had been seen.

If the token has an attribute with the name "prompt", then the first stream of characters must be the same string as given in that attribute, and the second stream of characters must be empty. Otherwise, the two streams of character tokens together should, together with the input element, express the equivalent of "This is a searchable index. Insert your search keywords here: (input field)" in the user's preferred language.

↪ A start tag whose tag name is "textarea"

1. Insert an HTML element (page 618) for the token.
2. If the form element pointer (page 597) is not null and the newly created element doesn't have a form attribute, then associate (page 361) the newly created textarea element with the form element pointed to by the form element pointer (page 597).
3. If the next token is a U+000A LINE FEED (LF) character token, then ignore that token and move on to the next one. (Newlines at the start of textarea elements are ignored as an authoring convenience.)
4. Append the new element to the current node (page 594).
5. Switch the tokeniser's content model flag (page 597) to the RCDATA state.

6. Let the original insertion mode (page 593) be the current insertion mode (page 592).
7. Switch the insertion mode (page 592) to "in CDATA/RCDATA (page 640)".

↪ **A start tag whose tag name is one of: "iframe", "noembed"**

↪ **A start tag whose tag name is "noscript", if the scripting flag (page 597) is enabled**

Follow the generic CDATA element parsing algorithm (page 619).

↪ **A start tag whose tag name is "select"**

Reconstruct the active formatting elements (page 596), if any.

Insert an HTML element (page 618) for the token.

If the form element pointer (page 597) is not null and the newly created element doesn't have a form attribute, then associate (page 361) the select element with the form element pointed to by the form element pointer (page 597).

If the insertion mode (page 592) is one of in table (page 642)", "in caption (page 644)", "in column group (page 645)", "in table body (page 646)", "in row (page 647)", or "in cell (page 648)", then switch the insertion mode (page 592) to "in select in table (page 651)". Otherwise, switch the insertion mode (page 592) to "in select (page 649)".

↪ **A start tag whose tag name is one of: "optgroup", "option"**

If the stack of open elements (page 594) has an option element in scope (page 595), then act as if an end tag with the tag name "option" had been seen.

Reconstruct the active formatting elements (page 596), if any.

Insert an HTML element (page 618) for the token.

↪ **A start tag whose tag name is one of: "rp", "rt"**

If the stack of open elements (page 594) has a ruby element in scope (page 595), then generate implied end tags (page 620). If the current node (page 594) is not then a ruby element, this is a parse error (page 583); pop all the nodes from the current node (page 594) up to the node immediately before the bottommost ruby element on the stack of open elements (page 594).

Insert an HTML element (page 618) for the token.

↪ **An end tag whose tag name is "br"**

Parse error (page 583). Act as if a start tag token with the tag name "br" had been seen. Ignore the end tag token.

↪ **A start tag whose tag name is "math"**

Reconstruct the active formatting elements (page 596), if any.

Adjust MathML attributes (page 619) for the token. (This fixes the case of MathML attributes that are not all lowercase.)

Adjust foreign attributes (page 619) for the token. (This fixes the use of namespaced attributes, in particular XLink.)

Insert a foreign element (page 619) for the token, in the MathML namespace (page 658).

If the token has its *self-closing flag* set, pop the current node (page 594) off the stack of open elements (page 594) and acknowledge the token's *self-closing flag* (page 598).

Otherwise, let the secondary insertion mode (page 593) be the current insertion mode (page 592), and then switch the insertion mode (page 592) to "in foreign content (page 651)".

↪ **A start tag whose tag name is one of: "caption", "col", "colgroup", "frame", "frameset", "head", "tbody", "td", "tfoot", "th", "thead", "tr"**

Parse error (page 583). Ignore the token.

↪ **Any other start tag**

Reconstruct the active formatting elements (page 596), if any.

Insert an HTML element (page 618) for the token.

Note: This element will be a phrasing (page 595) element.

↪ **Any other end tag**

Run the following steps:

1. Initialize *node* to be the current node (page 594) (the bottommost node of the stack).
2. If *node* has the same tag name as the end tag token, then:
 1. Generate implied end tags (page 620).
 2. If the tag name of the end tag token does not match the tag name of the current node (page 594), this is a parse error (page 583).
 3. Pop all the nodes from the current node (page 594) up to *node*, including *node*, then stop these steps.
3. Otherwise, if *node* is in neither the formatting (page 595) category nor the phrasing (page 595) category, then this is a parse error (page 583); ignore the token, and abort these steps.
4. Set *node* to the previous entry in the stack of open elements (page 594).
5. Return to step 2.

8.2.5.11 The "in CDATA/RCDATA" insertion mode

When the insertion mode (page 592) is "in CDATA/RCDATA (page 640)", tokens must be handled as follows:

↪ **A character token**

Insert the token's character (page 618) into the current node (page 594).

↪ **An end-of-file token**

Parse error (page 583).

If the current node (page 594) is a script element, mark the script element as "already executed" (page 124).

Pop the current node (page 594) off the stack of open elements (page 594).

Switch the insertion mode (page 592) to the original insertion mode (page 593) and reprocess the current token.

↪ **An end tag whose tag name is "script"**

Let *script* be the current node (page 594) (which will be a script element).

Pop the current node (page 594) off the stack of open elements (page 594).

Switch the insertion mode (page 592) to the original insertion mode (page 593).

Let the *old insertion point* have the same value as the current insertion point (page 591). Let the insertion point (page 591) be just before the next input character (page 591).

Run (page 125) the *script*. This might cause some script to execute, which might cause new characters to be inserted into the tokeniser (page 101), and might cause the tokeniser to output more tokens, resulting in a reentrant invocation of the parser (page 584).

Let the insertion point (page 591) have the value of the *old insertion point*. (In other words, restore the insertion point (page 591) to the value it had before the previous paragraph. This value might be the "undefined" value.)

At this stage, if there is a pending external script, then:

↪ **If the tree construction stage is being called reentrantly (page 584), say from a call to document.write():**

Abort the processing of any nested invocations of the tokeniser, yielding control back to the caller. (Tokenization will resume when the caller returns to the "outer" tree construction stage.)

↪ **Otherwise:**

Follow these steps:

1. Let the *script* be the pending external script. There is no longer a pending external script.
2. Pause (page 23) until the script has completed loading (page 125).
3. Let the insertion point (page 591) be just before the next input character (page 591).
4. Execute the script (page 127).
5. Let the insertion point (page 591) be undefined again.
6. If there is once again a pending external script, then repeat these steps from step 1.

↪ **Any other end tag**

Pop the current node (page 594) off the stack of open elements (page 594).

Switch the insertion mode (page 592) to the original insertion mode (page 593).

8.2.5.12 The "in table" insertion mode

When the insertion mode (page 592) is "in table (page 642)", tokens must be handled as follows:

↪ **A character token that is one of one of U+0009 CHARACTER TABULATION, U+000A LINE FEED (LF), U+000C FORM FEED (FF), or U+0020 SPACE**

If the current table (page 594) is tainted (page 620), then act as described in the "anything else" entry below.

Otherwise, insert the character (page 618) into the current node (page 594).

↪ **A comment token**

Append a Comment node to the current node (page 594) with the data attribute set to the data given in the comment token.

↪ **A DOCTYPE token**

Parse error (page 583). Ignore the token.

↪ **A start tag whose tag name is "caption"**

Clear the stack back to a table context (page 644). (See below.)

Insert a marker at the end of the list of active formatting elements (page 596).

Insert an HTML element (page 618) for the token, then switch the insertion mode (page 592) to "in caption (page 644)".

↪ **A start tag whose tag name is "colgroup"**

Clear the stack back to a table context (page 644). (See below.)

Insert an HTML element (page 618) for the token, then switch the insertion mode (page 592) to "in column group (page 645)".

↪ **A start tag whose tag name is "col"**

Act as if a start tag token with the tag name "colgroup" had been seen, then reprocess the current token.

↪ **A start tag whose tag name is one of: "tbody", "tfoot", "thead"**

Clear the stack back to a table context (page 644). (See below.)

Insert an HTML element (page 618) for the token, then switch the insertion mode (page 592) to "in table body (page 646)".

↪ **A start tag whose tag name is one of: "td", "th", "tr"**

Act as if a start tag token with the tag name "tbody" had been seen, then reprocess the current token.

↪ **A start tag whose tag name is "table"**

Parse error (page 583). Act as if an end tag token with the tag name "table" had been seen, then, if that token wasn't ignored, reprocess the current token.

Note: *The fake end tag token here can only be ignored in the fragment case (page 661).*

↪ **An end tag whose tag name is "table"**

If the stack of open elements (page 594) does not have an element in table scope (page 595) with the same tag name as the token, this is a parse error (page 583). Ignore the token. (fragment case (page 661))

Otherwise:

Pop elements from this stack until a table element has been popped from the stack.

Reset the insertion mode appropriately (page 593).

↪ **An end tag whose tag name is one of: "body", "caption", "col", "colgroup", "html", "tbody", "td", "tfoot", "th", "thead", "tr"**

Parse error (page 583). Ignore the token.

↪ **A start tag whose tag name is one of: "style", "script"**

If the current table (page 594) is tainted (page 620) then act as described in the "anything else" entry below.

Otherwise, process the token using the rules for (page 593) the "in head (page 625)" insertion mode (page 592).

↪ **A start tag whose tag name is "input"**

If the token does not have an attribute with the name "type", or if it does, but that attribute's value is not an ASCII case-insensitive (page 31) match for the string "hidden", or, if the current table (page 594) is tainted (page 620), then: act as described in the "anything else" entry below.

Otherwise:

Parse error (page 583).

Insert an HTML element (page 618) for the token.

If the form element pointer (page 597) is not null, then associate the input element with the form element pointed to by the form element pointer (page 597).

Pop that input element off the stack of open elements (page 594).

↪ **An end-of-file token**

If the current node (page 594) is not the root html element, then this is a parse error (page 583).

Note: *It can only be the current node (page 594) in the fragment case (page 661).*

Stop parsing (page 656).

↪ **Anything else**

Parse error (page 583). Process the token using the rules for (page 593) the "in body (page 628)" insertion mode (page 592), except that if the current node (page 594) is a table, tbody, tfoot, thead, or tr element, then, whenever a node would be inserted into the current node (page 594), it must instead be foster parented (page 620).

When the steps above require the UA to **clear the stack back to a table context**, it means that the UA must, while the current node (page 594) is not a table element or an html element, pop elements from the stack of open elements (page 594).

Note: *The current node (page 594) being an html element after this process is a fragment case (page 661).*

8.2.5.13 The "in caption" insertion mode

When the insertion mode (page 592) is "in caption (page 644)", tokens must be handled as follows:

↪ **An end tag whose tag name is "caption"**

If the stack of open elements (page 594) does not have an element in table scope (page 595) with the same tag name as the token, this is a parse error (page 583). Ignore the token. (fragment case (page 661))

Otherwise:

Generate implied end tags (page 620).

Now, if the current node (page 594) is not a caption element, then this is a parse error (page 583).

Pop elements from this stack until a caption element has been popped from the stack.

Clear the list of active formatting elements up to the last marker (page 597).

Switch the insertion mode (page 592) to "in table (page 642)".

↪ **A start tag whose tag name is one of: "caption", "col", "colgroup", "tbody", "td", "tfoot", "th", "thead", "tr"**

↪ **An end tag whose tag name is "table"**

Parse error (page 583). Act as if an end tag with the tag name "caption" had been seen, then, if that token wasn't ignored, reprocess the current token.

Note: *The fake end tag token here can only be ignored in the fragment case (page 661).*

↪ **An end tag whose tag name is one of: "body", "col", "colgroup", "html", "tbody", "td", "tfoot", "th", "thead", "tr"**

Parse error (page 583). Ignore the token.

↪ **Anything else**

Process the token using the rules for (page 593) the "in body (page 628)" insertion mode (page 592).

8.2.5.14 The "in column group" insertion mode

When the insertion mode (page 592) is "in column group (page 645)", tokens must be handled as follows:

↪ **A character token that is one of one of U+0009 CHARACTER TABULATION, U+000A LINE FEED (LF), U+000C FORM FEED (FF), or U+0020 SPACE**

Insert the character (page 618) into the current node (page 594).

↪ **A comment token**

Append a Comment node to the current node (page 594) with the data attribute set to the data given in the comment token.

↪ **A DOCTYPE token**

Parse error (page 583). Ignore the token.

↪ **A start tag whose tag name is "html"**

Process the token using the rules for (page 593) the "in body (page 628)" insertion mode (page 592).

↪ **A start tag whose tag name is "col"**

Insert an HTML element (page 618) for the token. Immediately pop the current node (page 594) off the stack of open elements (page 594).

Acknowledge the token's *self-closing flag* (page 598), if it is set.

↪ **An end tag whose tag name is "colgroup"**

If the current node (page 594) is the root html element, then this is a parse error (page 583); ignore the token. (fragment case (page 661))

Otherwise, pop the current node (page 594) (which will be a colgroup element) from the stack of open elements (page 594). Switch the insertion mode (page 592) to "in table (page 642)".

↪ **An end tag whose tag name is "col"**

Parse error (page 583). Ignore the token.

↪ **An end-of-file token**

If the current node (page 594) is the root html element, then stop parsing (page 656). (fragment case (page 661))

Otherwise, act as described in the "anything else" entry below.

↪ **Anything else**

Act as if an end tag with the tag name "colgroup" had been seen, and then, if that token wasn't ignored, reprocess the current token.

Note: The fake end tag token here can only be ignored in the fragment case (page 661).

8.2.5.15 The "in table body" insertion mode

When the insertion mode (page 592) is "in table body (page 646)", tokens must be handled as follows:

↪ **A start tag whose tag name is "tr"**

Clear the stack back to a table body context (page 647). (See below.)

Insert an HTML element (page 618) for the token, then switch the insertion mode (page 592) to "in row (page 647)".

↪ **A start tag whose tag name is one of: "th", "td"**

Parse error (page 583). Act as if a start tag with the tag name "tr" had been seen, then reprocess the current token.

↪ **An end tag whose tag name is one of: "tbody", "tfoot", "thead"**

If the stack of open elements (page 594) does not have an element in table scope (page 595) with the same tag name as the token, this is a parse error (page 583). Ignore the token.

Otherwise:

Clear the stack back to a table body context (page 647). (See below.)

Pop the current node (page 594) from the stack of open elements (page 594). Switch the insertion mode (page 592) to "in table (page 642)".

↪ **A start tag whose tag name is one of: "caption", "col", "colgroup", "tbody", "tfoot", "thead"**

↪ **An end tag whose tag name is "table"**

If the stack of open elements (page 594) does not have a tbody, thead, or tfoot element in table scope (page 595), this is a parse error (page 583). Ignore the token. (fragment case (page 661))

Otherwise:

Clear the stack back to a table body context (page 647). (See below.)

Act as if an end tag with the same tag name as the current node (page 594) ("tbody", "tfoot", or "thead") had been seen, then reprocess the current token.

↪ **An end tag whose tag name is one of: "body", "caption", "col", "colgroup", "html", "td", "th", "tr"**

Parse error (page 583). Ignore the token.

↪ **Anything else**

Process the token using the rules for (page 593) the "in table (page 642)" insertion mode (page 592).

When the steps above require the UA to **clear the stack back to a table body context**, it means that the UA must, while the current node (page 594) is not a tbody, tfoot, thead, or html element, pop elements from the stack of open elements (page 594).

Note: *The current node (page 594) being an html element after this process is a fragment case (page 661).*

8.2.5.16 The "in row" insertion mode

When the insertion mode (page 592) is "in row (page 647)", tokens must be handled as follows:

↪ **A start tag whose tag name is one of: "th", "td"**

Clear the stack back to a table row context (page 648). (See below.)

Insert an HTML element (page 618) for the token, then switch the insertion mode (page 592) to "in cell (page 648)".

Insert a marker at the end of the list of active formatting elements (page 596).

↪ **An end tag whose tag name is "tr"**

If the stack of open elements (page 594) does not have an element in table scope (page 595) with the same tag name as the token, this is a parse error (page 583). Ignore the token. (fragment case (page 661))

Otherwise:

Clear the stack back to a table row context (page 648). (See below.)

Pop the current node (page 594) (which will be a tr element) from the stack of open elements (page 594). Switch the insertion mode (page 592) to "in table body (page 646)".

↪ **A start tag whose tag name is one of: "caption", "col", "colgroup", "tbody", "tfoot", "thead", "tr"**

↪ **An end tag whose tag name is "table"**

Act as if an end tag with the tag name "tr" had been seen, then, if that token wasn't ignored, reprocess the current token.

Note: *The fake end tag token here can only be ignored in the fragment case (page 661).*

↪ **An end tag whose tag name is one of: "tbody", "tfoot", "thead"**

If the stack of open elements (page 594) does not have an element in table scope (page 595) with the same tag name as the token, this is a parse error (page 583). Ignore the token.

Otherwise, act as if an end tag with the tag name "tr" had been seen, then reprocess the current token.

- ↪ **An end tag whose tag name is one of: "body", "caption", "col", "colgroup", "html", "td", "th"**

Parse error (page 583). Ignore the token.

- ↪ **Anything else**

Process the token using the rules for (page 593) the "in table (page 642)" insertion mode (page 592).

When the steps above require the UA to **clear the stack back to a table row context**, it means that the UA must, while the current node (page 594) is not a tr element or an html element, pop elements from the stack of open elements (page 594).

Note: *The current node (page 594) being an html element after this process is a fragment case (page 661).*

8.2.5.17 The "in cell" insertion mode

When the insertion mode (page 592) is "in cell (page 648)", tokens must be handled as follows:

- ↪ **An end tag whose tag name is one of: "td", "th"**

If the stack of open elements (page 594) does not have an element in table scope (page 595) with the same tag name as that of the token, then this is a parse error (page 583) and the token must be ignored.

Otherwise:

Generate implied end tags (page 620).

Now, if the current node (page 594) is not an element with the same tag name as the token, then this is a parse error (page 583).

Pop elements from this stack until an element with the same tag name as the token has been popped from the stack.

Clear the list of active formatting elements up to the last marker (page 597).

Switch the insertion mode (page 592) to "in row (page 647)". (The current node (page 594) will be a tr element at this point.)

- ↪ **A start tag whose tag name is one of: "caption", "col", "colgroup", "tbody", "td", "tfoot", "th", "thead", "tr"**

If the stack of open elements (page 594) does not have a td or th element in table scope (page 595), then this is a parse error (page 583); ignore the token. (fragment case (page 661))

Otherwise, close the cell (page 649) (see below) and reprocess the current token.

- ↪ **An end tag whose tag name is one of: "body", "caption", "col", "colgroup", "html"**

Parse error (page 583). Ignore the token.

↪ **An end tag whose tag name is one of: "table", "tbody", "tfoot", "thead", "tr"**

If the stack of open elements (page 594) does not have an element in table scope (page 595) with the same tag name as that of the token (which can only happen for "tbody", "tfoot" and "thead", or, in the fragment case (page 661)), then this is a parse error (page 583) and the token must be ignored.

Otherwise, close the cell (page 649) (see below) and reprocess the current token.

↪ **Anything else**

Process the token using the rules for (page 593) the "in body (page 628)" insertion mode (page 592).

Where the steps above say to **close the cell**, they mean to run the following algorithm:

1. If the stack of open elements (page 594) has a td element in table scope (page 595), then act as if an end tag token with the tag name "td" had been seen.
2. Otherwise, the stack of open elements (page 594) will have a th element in table scope (page 595); act as if an end tag token with the tag name "th" had been seen.

Note: *The stack of open elements (page 594) cannot have both a td and a th element in table scope (page 595) at the same time, nor can it have neither when the insertion mode (page 592) is "in cell (page 648)".*

8.2.5.18 The "in select" insertion mode

When the insertion mode (page 592) is "in select (page 649)", tokens must be handled as follows:

↪ **A character token**

Insert the token's character (page 618) into the current node (page 594).

↪ **A comment token**

Append a Comment node to the current node (page 594) with the data attribute set to the data given in the comment token.

↪ **A DOCTYPE token**

Parse error (page 583). Ignore the token.

↪ **A start tag whose tag name is "html"**

Process the token using the rules for (page 593) the "in body (page 628)" insertion mode (page 592).

↪ **A start tag whose tag name is "option"**

If the current node (page 594) is an option element, act as if an end tag with the tag name "option" had been seen.

Insert an HTML element (page 618) for the token.

↪ **A start tag whose tag name is "optgroup"**

If the current node (page 594) is an option element, act as if an end tag with the tag name "option" had been seen.

If the current node (page 594) is an optgroup element, act as if an end tag with the tag name "optgroup" had been seen.

Insert an HTML element (page 618) for the token.

↪ **An end tag whose tag name is "optgroup"**

First, if the current node (page 594) is an option element, and the node immediately before it in the stack of open elements (page 594) is an optgroup element, then act as if an end tag with the tag name "option" had been seen.

If the current node (page 594) is an optgroup element, then pop that node from the stack of open elements (page 594). Otherwise, this is a parse error (page 583); ignore the token.

↪ **An end tag whose tag name is "option"**

If the current node (page 594) is an option element, then pop that node from the stack of open elements (page 594). Otherwise, this is a parse error (page 583); ignore the token.

↪ **An end tag whose tag name is "select"**

If the stack of open elements (page 594) does not have an element in table scope (page 595) with the same tag name as the token, this is a parse error (page 583). Ignore the token. (fragment case (page 661))

Otherwise:

Pop elements from the stack of open elements (page 594) until a select element has been popped from the stack.

Reset the insertion mode appropriately (page 593).

↪ **A start tag whose tag name is "select"**

Parse error (page 583). Act as if the token had been an end tag with the tag name "select" instead.

↪ **A start tag whose tag name is one of: "input", "textarea"**

Parse error (page 583). Act as if an end tag with the tag name "select" had been seen, and reprocess the token.

↪ **An end-of-file token**

If the current node (page 594) is not the root html element, then this is a parse error (page 583).

Note: It can only be the current node (page 594) in the fragment case (page 661).

Stop parsing (page 656).

↪ **Anything else**

Parse error (page 583). Ignore the token.

8.2.5.19 The "in select in table" insertion mode

When the insertion mode (page 592) is "in select in table (page 651)", tokens must be handled as follows:

- ↪ **A start tag whose tag name is one of: "caption", "table", "tbody", "tfoot", "thead", "tr", "td", "th"**

Parse error (page 583). Act as if an end tag with the tag name "select" had been seen, and reprocess the token.

- ↪ **An end tag whose tag name is one of: "caption", "table", "tbody", "tfoot", "thead", "tr", "td", "th"**

Parse error (page 583).

If the stack of open elements (page 594) has an element in table scope (page 595) with the same tag name as that of the token, then act as if an end tag with the tag name "select" had been seen, and reprocess the token. Otherwise, ignore the token.

- ↪ **Anything else**

Process the token using the rules for (page 593) the "in select (page 649)" insertion mode (page 592).

8.2.5.20 The "in foreign content" insertion mode

When the insertion mode (page 592) is "in foreign content (page 651)", tokens must be handled as follows:

- ↪ **A character token**

Insert the token's character (page 618) into the current node (page 594).

- ↪ **A comment token**

Append a Comment node to the current node (page 594) with the data attribute set to the data given in the comment token.

- ↪ **A DOCTYPE token**

Parse error (page 583). Ignore the token.

- ↪ A start tag whose tag name is neither "mglyph" nor "malignmark", if the current node (page 594) is an `mi` element in the MathML namespace (page 658).
- ↪ A start tag whose tag name is neither "mglyph" nor "malignmark", if the current node (page 594) is an `mo` element in the MathML namespace (page 658).
- ↪ A start tag whose tag name is neither "mglyph" nor "malignmark", if the current node (page 594) is an `mn` element in the MathML namespace (page 658).
- ↪ A start tag whose tag name is neither "mglyph" nor "malignmark", if the current node (page 594) is an `ms` element in the MathML namespace (page 658).
- ↪ A start tag whose tag name is neither "mglyph" nor "malignmark", if the current node (page 594) is an `mtext` element in the MathML namespace (page 658).
- ↪ A start tag, if the current node (page 594) is an element in the HTML namespace (page 658).
- ↪ An end tag

Process the token using the rules for (page 593) the secondary insertion mode (page 593).

If, after doing so, the insertion mode (page 592) is still "in foreign content (page 561)", but there is no element in scope that has a namespace other than the HTML namespace (page 658), switch the insertion mode (page 592) to the secondary insertion mode (page 593).

- ↪ A start tag whose tag name is one of: "b", "big", "blockquote", "body", "br", "center", "code", "dd", "div", "dl", "dt", "em", "embed", "font", "h1", "h2", "h3", "h4", "h5", "h6", "head", "hr", "i", "img", "li", "listing", "menu", "meta", "nobr", "ol", "p", "pre", "ruby", "s", "small", "span", "strong", "strike", "sub", "sup", "table", "tt", "u", "ul", "var"

- ↪ An end-of-file token

Parse error (page 583).

Pop elements from the stack of open elements (page 594) until the current node (page 594) is in the HTML namespace (page 658).

Switch the insertion mode (page 592) to the secondary insertion mode (page 593), and reprocess the token.

- ↪ Any other start tag

If the current node (page 594) is an element in the MathML namespace (page 658), adjust MathML attributes (page 619) for the token. (This fixes the case of MathML attributes that are not all lowercase.)

Adjust foreign attributes (page 619) for the token. (This fixes the use of namespaced attributes, in particular XLink in SVG.)

Insert a foreign element (page 619) for the token, in the same namespace as the current node (page 594).

If the token has its *self-closing flag* set, pop the current node (page 594) off the stack of open elements (page 594) and acknowledge the token's *self-closing flag* (page 598).

8.2.5.21 The "after body" insertion mode

When the insertion mode (page 592) is "after body (page 653)", tokens must be handled as follows:

↪ **A character token that is one of one of U+0009 CHARACTER TABULATION, U+000A LINE FEED (LF), U+000C FORM FEED (FF), or U+0020 SPACE**

Process the token using the rules for (page 593) the "in body (page 628)" insertion mode (page 592).

↪ **A comment token**

Append a Comment node to the first element in the stack of open elements (page 594) (the `html` element), with the `data` attribute set to the data given in the comment token.

↪ **A DOCTYPE token**

Parse error (page 583). Ignore the token.

↪ **A start tag whose tag name is "html"**

Process the token using the rules for (page 593) the "in body (page 628)" insertion mode (page 592).

↪ **An end tag whose tag name is "html"**

If the parser was originally created as part of the HTML fragment parsing algorithm (page 661), this is a parse error (page 583); ignore the token. (fragment case (page 661))

Otherwise, switch the insertion mode (page 592) to "after after body (page 655)".

↪ **An end-of-file token**

Stop parsing (page 656).

↪ **Anything else**

Parse error (page 583). Switch the insertion mode (page 592) to "in body (page 628)" and reprocess the token.

8.2.5.22 The "in frameset" insertion mode

When the insertion mode (page 592) is "in frameset (page 653)", tokens must be handled as follows:

↪ **A character token that is one of one of U+0009 CHARACTER TABULATION, U+000A LINE FEED (LF), U+000C FORM FEED (FF), or U+0020 SPACE**

Insert the character (page 618) into the current node (page 594).

↪ **A comment token**

Append a Comment node to the current node (page 594) with the data attribute set to the data given in the comment token.

↪ **A DOCTYPE token**

Parse error (page 583). Ignore the token.

↪ **A start tag whose tag name is "html"**

Process the token using the rules for (page 593) the "in body (page 628)" insertion mode (page 592).

↪ **A start tag whose tag name is "frameset"**

Insert an HTML element (page 618) for the token.

↪ **An end tag whose tag name is "frameset"**

If the current node (page 594) is the root `html` element, then this is a parse error (page 583); ignore the token. (fragment case (page 661))

Otherwise, pop the current node (page 594) from the stack of open elements (page 594).

If the parser was *not* originally created as part of the HTML fragment parsing algorithm (page 661) (fragment case (page 661)), and the current node (page 594) is no longer a `frameset` element, then switch the insertion mode (page 592) to "after frameset (page 654)".

↪ **A start tag whose tag name is "frame"**

Insert an HTML element (page 618) for the token. Immediately pop the current node (page 594) off the stack of open elements (page 594).

Acknowledge the token's *self-closing flag* (page 598), if it is set.

↪ **A start tag whose tag name is "noframes"**

Process the token using the rules for (page 593) the "in head (page 625)" insertion mode (page 592).

↪ **An end-of-file token**

If the current node (page 594) is not the root `html` element, then this is a parse error (page 583).

Note: *It can only be the current node (page 594) in the fragment case (page 661).*

Stop parsing (page 656).

↪ **Anything else**

Parse error (page 583). Ignore the token.

8.2.5.23 The "after frameset" insertion mode

When the insertion mode (page 592) is "after frameset (page 654)", tokens must be handled as follows:

- ↪ **A character token that is one of one of U+0009 CHARACTER TABULATION, U+000A LINE FEED (LF), U+000C FORM FEED (FF), or U+0020 SPACE**
Insert the character (page 618) into the current node (page 594).
- ↪ **A comment token**
Append a Comment node to the current node (page 594) with the data attribute set to the data given in the comment token.
- ↪ **A DOCTYPE token**
Parse error (page 583). Ignore the token.
- ↪ **A start tag whose tag name is "html"**
Process the token using the rules for (page 593) the "in body (page 628)" insertion mode (page 592).
- ↪ **An end tag whose tag name is "html"**
Switch the insertion mode (page 592) to "after after frameset (page 656)".
- ↪ **A start tag whose tag name is "noframes"**
Process the token using the rules for (page 593) the "in head (page 625)" insertion mode (page 592).
- ↪ **An end-of-file token**
Stop parsing (page 656).
- ↪ **Anything else**
Parse error (page 583). Ignore the token.

** This doesn't handle UAs that don't support frames, or that do support frames but want to show the NOFRAMES content. Supporting the former is easy; supporting the latter is harder.

8.2.5.24 The "after after body" insertion mode

When the insertion mode (page 592) is "after after body (page 655)", tokens must be handled as follows:

- ↪ **A comment token**
Append a Comment node to the Document object with the data attribute set to the data given in the comment token.
- ↪ **A DOCTYPE token**
- ↪ **A character token that is one of one of U+0009 CHARACTER TABULATION, U+000A LINE FEED (LF), U+000C FORM FEED (FF), or U+0020 SPACE**
- ↪ **A start tag whose tag name is "html"**
Process the token using the rules for (page 593) the "in body (page 628)" insertion mode (page 592).
- ↪ **An end-of-file token**
Stop parsing (page 656).

↪ **Anything else**

Parse error (page 583). Switch the insertion mode (page 592) to "in body (page 628)" and reprocess the token.

8.2.5.25 The "after after frameset" insertion mode

When the insertion mode (page 592) is "after after frameset (page 656)", tokens must be handled as follows:

↪ **A comment token**

Append a Comment node to the Document object with the data attribute set to the data given in the comment token.

↪ **A DOCTYPE token**

↪ **A character token that is one of one of U+0009 CHARACTER TABULATION, U+000A LINE FEED (LF), U+000C FORM FEED (FF), or U+0020 SPACE**

↪ **A start tag whose tag name is "html"**

Process the token using the rules for (page 593) the "in body (page 628)" insertion mode (page 592).

↪ **An end-of-file token**

Stop parsing (page 656).

↪ **A start tag whose tag name is "noframes"**

Process the token using the rules for (page 593) the "in head (page 625)" insertion mode (page 592).

↪ **Anything else**

Parse error (page 583). Ignore the token.

8.2.6 The end

Once the user agent **stops parsing** the document, the user agent must follow the steps in this section.

First, the current document readiness (page 80) must be set to "interactive".

Then, the rules for when a script completes loading (page 126) start applying (script execution is no longer managed by the parser).

If any of the scripts in the list of scripts that will execute as soon as possible (page 127) have completed loading (page 125), or if the list of scripts that will execute asynchronously (page 127) is not empty and the first script in that list has completed loading (page 125), then the user agent must act as if those scripts just completed loading, following the rules given for that in the script element definition.

Then, if the list of scripts that will execute when the document has finished parsing (page 126) is not empty, and the first item in this list has already completed loading (page 125), then the user agent must act as if that script just finished loading.

By this point, there will be no scripts that have loaded but have not yet been executed.

The user agent must then fire a simple event (page 436) called `DOMContentLoaded` at the Document.

Once everything that **delays the load event** has completed, the user agent must set the current document readiness (page 80) to "complete", and then fire a `load` event (page 436) at the body element (page 81).

- ** delaying the load event for things like image loads allows for intranet port scans (even without javascript!). Should we really encode that into the spec?

8.2.7 Coercing an HTML DOM into an infoset

When an application uses an HTML parser (page 582) in conjunction with an XML pipeline, it is possible that the constructed DOM is not compatible with the XML tool chain in certain subtle ways. For example, an XML toolchain might not be able to represent attributes with the name `xmlns`, since they conflict with the Namespaces in XML syntax. There is also some data that the HTML parser (page 582) generates that isn't included in the DOM itself. This section specifies some rules for handling these issues.

If the XML API being used doesn't support DOCTYPEs, the tool may drop DOCTYPEs altogether.

If the XML API doesn't support attributes in no namespace that are named "`xmlns`", attributes whose names start with "`xmlns :`", or attributes in the XMLNS namespace (page 658), then the tool may drop such attributes.

The tool may annotate the output with any namespace declarations required for proper operation.

If the XML API being used restricts the allowable characters in the local names of elements and attributes, then the tool may map all element and attribute local names that the API wouldn't support to a set of names that *are* allowed, by replacing any character that isn't supported with the uppercase letter U and the five digits of the character's Unicode codepoint when expressed in hexadecimal, using digits 0-9 and capital letters A-F as the symbols, in increasing numeric order.

For example, the element name `foo<bar`, which can be output by the HTML parser (page 582), though it is neither a legal HTML element name nor a well-formed XML element name, would be converted into `fooU0003Cbar`, which *is* a well-formed XML element name (though it's still not legal in HTML by any means).

As another example, consider the attribute `xlink:href`. Used on a MathML element, it becomes, after being adjusted (page 619), an attribute with a prefix "xlink" and a local name "href". However, used on an HTML element, it becomes an attribute with no prefix and the local name "xlink:href", which is not a valid NCName, and thus might not be accepted by an XML API. It could thus get converted, becoming "`xlinkU0003Ahref`".

Note: *The resulting names from this conversion conveniently can't clash with any attribute generated by the HTML parser (page 582), since those are all either lowercase or those listed in the adjust foreign attributes (page 619) algorithm's table.*

If the XML API restricts comments from having two consecutive U+002D HYPHEN-MINUS characters (--), the tool may insert a single U+0020 SPACE character between any such offending characters.

If the XML API restricts comments from ending in a U+002D HYPHEN-MINUS character (-), the tool may insert a single U+0020 SPACE character at the end of such comments.

If the XML API restricts allowed characters in character data, the tool may replace any U+000C FORM FEED (FF) character with a U+0020 SPACE character, and any other literal non-XML character with a U+FFFFD REPLACEMENT CHARACTER.

If the tool has no way to convey out-of-band information, then the tool may drop the following information:

- Whether the document is set to *no quirks mode* (page 79), *limited quirks mode* (page 79), or *quirks mode* (page 79)
- The association between form controls and forms that aren't their nearest form element ancestor (use of the form element pointer (page 597) in the parser)

Note: *The mutations allowed by this section apply after the HTML parser (page 582)'s rules have been applied. For example, a <a:> start tag will be closed by a </a:> end tag, and never by a </aU0003AU0003A> end tag, even if the user agent is using the rules above to then generate an actual element in the DOM with the name aU0003AU0003A for that start tag.*

8.3 Namespaces

The **HTML namespace** is: <http://www.w3.org/1999/xhtml>

The **MathML namespace** is: <http://www.w3.org/1998/Math/MathML>

The **SVG namespace** is: <http://www.w3.org/2000/svg>

The **XLink namespace** is: <http://www.w3.org/1999/xlink>

The **XML namespace** is: <http://www.w3.org/XML/1998/namespace>

The **XMLNS namespace** is: <http://www.w3.org/2000/xmlns/>

8.4 Serializing HTML fragments

The following steps form the **HTML fragment serialization algorithm**. The algorithm takes as input a DOM Element or Document, referred to as *the node*, and either returns a string or raises an exception.

Note: *This algorithm serializes the children of the node being serialized, not the node itself.*

1. Let s be a string, and initialize it to the empty string.

2. For each child node of *the node*, in tree order (page 24), run the following steps:

1. Let *current node* be the child node being processed.
2. Append the appropriate string from the following list to *s*:

↪ **If *current node* is an Element**

Append a U+003C LESS-THAN SIGN (<) character, followed by the element's tag name. (For nodes created by the HTML parser (page 582), `Document.createElement()`, or `Document.renameNode()`, the tag name will be lowercase.)

For each attribute that the element has, append a U+0020 SPACE character, the attribute's name (which, for attributes set by the HTML parser (page 582) or by `Element.setAttributeNode()` or `Element.setAttribute()`, will be lowercase), a U+003D EQUALS SIGN (=) character, a U+0022 QUOTATION MARK ("") character, the attribute's value, escaped as described below (page 660) in *attribute mode*, and a second U+0022 QUOTATION MARK ("") character.

While the exact order of attributes is UA-defined, and may depend on factors such as the order that the attributes were given in the original markup, the sort order must be stable, such that consecutive invocations of this algorithm serialize an element's attributes in the same order.

Append a U+003E GREATER-THAN SIGN (>) character.

If *current node* is an area, base, basefont, bgsound, br, col, embed, frame, hr, img, input, link, meta, param, spacer, or wbr element, then continue on to the next child node at this point.

If *current node* is a pre textarea, or listing element, append a U+000A LINE FEED (LF) character.

Append the value of running the HTML fragment serialization algorithm (page 658) on the *current node* element (thus recursing into this algorithm for that element), followed by a U+003C LESS-THAN SIGN (<) character, a U+002F SOLIDUS (/) character, the element's tag name again, and finally a U+003E GREATER-THAN SIGN (>) character.

↪ **If *current node* is a Text or CDATASEction node**

If one of the ancestors of *current node* is a style, script, xmp, iframe, noembed, noframes, noscript, or plaintext element, then append the value of *current node*'s data DOM attribute literally.

Otherwise, append the value of *current node*'s data DOM attribute, escaped as described below (page 660).

↪ **If *current node* is a Comment**

Append the literal string <! -- (U+003C LESS-THAN SIGN, U+0021 EXCLAMATION MARK, U+002D HYPHEN-MINUS, U+002D HYPHEN-MINUS), followed by the value of *current node*'s data DOM

attribute, followed by the literal string `-->` (U+002D HYPHEN-MINUS, U+002D HYPHEN-MINUS, U+003E GREATER-THAN SIGN).

↪ **If current node is a ProcessingInstruction**

Append the literal string `<? (U+003C LESS-THAN SIGN, U+003F QUESTION MARK), followed by the value of current node's target DOM attribute, followed by a single U+0020 SPACE character, followed by the value of current node's data DOM attribute, followed by a single U+003E GREATER-THAN SIGN character ('>').`

↪ **If current node is a DocumentType**

Append the literal string `<!DOCTYPE (U+003C LESS-THAN SIGN, U+0021 EXCLAMATION MARK, U+0044 LATIN CAPITAL LETTER D, U+004F LATIN CAPITAL LETTER O, U+0043 LATIN CAPITAL LETTER C, U+0054 LATIN CAPITAL LETTER T, U+0059 LATIN CAPITAL LETTER Y, U+0050 LATIN CAPITAL LETTER P, U+0045 LATIN CAPITAL LETTER E), followed by a space (U+0020 SPACE), followed by the value of current node's name DOM attribute, followed by the literal string > (U+003E GREATER-THAN SIGN).`

Other node types (e.g. Attr) cannot occur as children of elements. If, despite this, they somehow do occur, this algorithm must raise an `INVALID_STATE_ERR` exception.

3. The result of the algorithm is the string `s`.

Escaping a string (for the purposes of the algorithm above) consists of replacing any occurrences of the "&" character by the string "&", any occurrences of the "<" character by the string "<", any occurrences of the ">" character by the string ">", any occurrences of the U+00A0 NO-BREAK SPACE character by the string " ", and, if the algorithm was invoked in the *attribute mode*, any occurrences of the """" character by the string """.

Note: Entity reference nodes are assumed to be expanded (page 29) by the user agent, and are therefore not covered in the algorithm above.

Note: It is possible that the output of this algorithm, if parsed with an HTML parser (page 582), will not return the original tree structure. For instance, if a `textarea` element to which a Comment node has been appended is serialized and the output is then reparsed, the comment will end up being displayed in the text field. Similarly, if, as a result of DOM manipulation, an element contains a comment that contains the literal string `-->`, then when the result of serializing the element is parsed, the comment will be truncated at that point and the rest of the comment will be interpreted as markup. More examples would be making a `script` element contain a text node with the text string `</script>`, or having a `p` element that contains a `ul` element (as the `ul` element's start tag would imply the end tag for the `p`).

8.5 Parsing HTML fragments

The following steps form the **HTML fragment parsing algorithm**. The algorithm takes as input a DOM Element, referred to as the *context* element, which gives the context for the parser, as well as *input*, a string to parse, and returns a list of zero or more nodes.

Note: *Parts marked fragment case in algorithms in the parser section are parts that only occur if the parser was created for the purposes of this algorithm. The algorithms have been annotated with such markings for informational purposes only; such markings have no normative weight. If it is possible for a condition described as a fragment case (page 661) to occur even when the parser wasn't created for the purposes of handling this algorithm, then that is an error in the specification.*

1. Create a new Document node, and mark it as being an HTML document (page 76).
2. Create a new HTML parser (page 582), and associate it with the just created Document node.
3. Set the HTML parser (page 582)'s tokenization (page 597) stage's content model flag (page 597) according to the *context* element, as follows:
 - ↪ **If it is a title or textarea element**
Set the content model flag (page 597) to the RCDATA state.
 - ↪ **If it is a style, script, xmp, iframe, noembed, or noframes element**
Set the content model flag (page 597) to the CDATA state.
 - ↪ **If it is a noscript element**
If the scripting flag (page 597) is enabled, set the content model flag (page 597) to the CDATA state. Otherwise, set the content model flag (page 597) to the PCDATA state.
 - ↪ **If it is a plaintext element**
Set the content model flag (page 597) to PLAINTEXT.
 - ↪ **Otherwise**
Set the content model flag (page 597) to the PCDATA state.
4. Let *root* be a new html element with no attributes.
5. Append the element *root* to the Document node created above.
6. Set up the parser's stack of open elements (page 594) so that it contains just the single element *root*.
7. Reset the parser's insertion mode appropriately (page 593).

Note: *The parser will reference the context element as part of that algorithm.*

8. Set the parser's form element pointer (page 597) to the nearest node to the *context* element that is a form element (going straight up the ancestor chain, and including

the element itself, if it is a form element), or, if there is no such form element, to null.

9. Place into the input stream (page 584) for the HTML parser (page 582) just created the *input*.
10. Start the parser and let it run until it has consumed all the characters just inserted into the input stream.
11. Return all the child nodes of *root*, preserving the document order.

8.6 Named character references

This table lists the character reference names that are supported by HTML, and the code points to which they refer. It is referenced by the previous sections.

Name	Character	Name	Character	Name	Character
AElig;	U+00C6	Cdot;	U+0010A	DoubleLeftArrow;	U+021D0
AElig	U+00C6	Cedilla;	U+000B8	DoubleLeftRightArrow;	U+021D4
AMP;	U+0026	CenterDot;	U+000B7	DoubleLeftTee;	U+02AE4
AMP	U+0026	Cfr;	U+0212D	DoubleLongLeftArrow;	U+027F8
Aacute;	U+00C1	Chi;	U+003A7	DoubleLongLeftRightArrow;	U+027FA
Aacute	U+00C1	CircleDot;	U+02299	DoubleLongRightArrow;	U+027F9
Abreve;	U+00102	CircleMinus;	U+02296	DoubleRightArrow;	U+021D2
Acirc;	U+00C2	CirclePlus;	U+02295	DoubleRightTee;	U+02A8
Acirc	U+00C2	CircleTimes;	U+02297	DoubleUpArrow;	U+021D1
Acy;	U+00410	ClockwiseContourIntegral;	U+02232	DoubleUpDownArrow;	U+021D5
Afr;	U+1D504	CloseCurlyDoubleQuote;	U+0201D	DoubleVerticalBar;	U+02225
Agrave;	U+00C0	CloseCurlyQuote;	U+02019	DownArrow;	U+02193
Agrave	U+00C0	Colon;	U+02237	DownArrowBar;	U+02913
Alpha;	U+00391	Colone;	U+02A74	DownArrowUpArrow;	U+021F5
Amacr;	U+00100	Congruent;	U+02261	DownBreve;	U+00311
And;	U+02A53	Conint;	U+0222F	DownLeftRightVector;	U+0295F
Aogon;	U+00104	ContourIntegral;	U+0222E	DownLeftTeeVector;	U+0295E
Aopf;	U+1D538	Copf;	U+02102	DownLeftVector;	U+021BD
ApplyFunction;	U+02061	Coproduct;	U+02210	DownLeftVectorBar;	U+02956
Aring;	U+00C5	CounterClockwiseContourIntegral;	U+02233	DownRightTeeVector;	U+0295F
Aring	U+00C5	Cross;	U+02A2F	DownRightVector;	U+021C1
Ascr;	U+1D49C	Cscr;	U+1D49E	DownRightVectorBar;	U+02957
Assign;	U+02254	Cup;	U+022D3	DownTee;	U+022A4
Atilde;	U+000C3	CupCap;	U+0224D	DownTeeArrow;	U+021A7
Atilde	U+000C3	DD;	U+02145	Downarrow;	U+021D3
Auml;	U+000C4	DDotrahd;	U+02911	Dscr;	U+1D49F
Auml	U+000C4	DJcy;	U+00402	Dstrok;	U+00110
Backslash;	U+02216	DScy;	U+00405	ENG;	U+0014A
Barv;	U+02AE7	DZcy;	U+0040F	ETH;	U+000D0
Barwed;	U+02306	Dagger;	U+02021	ETH	U+000D0
Bcy;	U+00411	Darr;	U+021A1	Eacute;	U+000C9
Because;	U+02235	Dashv;	U+02AE4	Eacute	U+000C9
Bernoullis;	U+0212C	Dcaron;	U+0010E	Ecaron;	U+0011A
Beta;	U+00392	Dcy;	U+00414	Ecirc;	U+000CA
Bfr;	U+1D505	Del;	U+02207	Ecirc	U+000CA
Bopf;	U+1D539	Delta;	U+00394	Ecy;	U+0042D
Breve;	U+002D8	Dfr;	U+1D507	Edot;	U+00116
Bscr;	U+0212C	DiacriticalAcute;	U+000B4	Efr;	U+1D508
Bumpeq;	U+0224E	DiacriticalDot;	U+002D9	Egrave;	U+000C8
CHcy;	U+00427	DiacriticalDoubleAcute;	U+002DD	Egrave	U+000C8
COPY;	U+000A9	DiacriticalGrave;	U+00060	Element;	U+02208
COPY	U+000A9	DiacriticalTilde;	U+002DC	Emacr;	U+00112
Acute;	U+00106	Diamond;	U+022C4	EmptySmallSquare;	U+025FB
Cap;	U+022D2	DifferentialD;	U+02146	EmptyVerySmallSquare;	U+025AB
CapitalDifferentialD;	U+02145	Dopf;	U+1D53B	Egon;	U+00118
Cayleys;	U+0212D	Dot;	U+000A8	Eopf;	U+1D53C
Ccaron;	U+0010C	DotDot;	U+020DC	Epsilon;	U+00395
Ccedil;	U+000C7	DotEqual;	U+02250	Equal;	U+02A75
Ccedil	U+000C7	DoubleContourIntegral;	U+0222F	EqualTilde;	U+02242
Ccirc;	U+00108	DoubleDot;	U+000A8	Equilibrium;	U+021CC
Cconint;	U+02230	DoubleDownArrow;	U+021D3	Escr;	U+02130

Name	Character	Name	Character	Name	Character
Esim;	U+02A73	Iuml;	U+000CF	Lstrok;	U+00141
Eta;	U+00397	Iuml;	U+000CF	Lt;	U+0226A
Euml;	U+000CB	Jcirc;	U+00134	Map;	U+02905
Euml	U+000CB	Jcy;	U+00419	Mcy;	U+0041C
Exists;	U+02203	Jfr;	U+1D50D	MediumSpace;	U+0205F
ExponentialE;	U+02147	Jopf;	U+1D541	Mellinrf;	U+02133
Fcy;	U+00424	Jscr;	U+1D445	Mfr;	U+1D510
Ffr;	U+1D509	Jsercy;	U+00408	MinusPlus;	U+02213
FilledSmallSquare;	U+025FC	Jukcy;	U+00404	Mopf;	U+1D544
FilledVerySmallSquare;	U+025AA	KHcy;	U+00425	Mscr;	U+02133
Fopf;	U+1D53D	KJcy;	U+0040C	Mu;	U+0039C
ForAll;	U+02200	Kappa;	U+0039A	NJcy;	U+0040A
Fouriertrf;	U+02131	Kcedil;	U+00136	Nacute;	U+00143
Fscr;	U+02131	Kcy;	U+0041A	Ncaron;	U+00147
GJcy;	U+00403	Kfr;	U+1D50E	Ncedil;	U+00145
GT;	U+0003E	Kopf;	U+1D542	Ncy;	U+0041D
GT	U+0003E	Kscr;	U+1D4A6	NegativeMediumSpace;	U+0200B
Gamma;	U+00393	LJcy;	U+00409	NegativeThickSpace;	U+0200B
Gammad;	U+003DC	LT;	U+0003C	NegativeThinSpace;	U+0200B
Gbreve;	U+0011E	LT	U+0003C	NegativeVeryThinSpace;	U+0200B
Gcedil;	U+00122	Lacute;	U+00139	NestedGreaterGreater;	U+0226B
Geirc;	U+0011C	Lambda;	U+0039B	NestedLessLess;	U+0226A
Gcy;	U+00413	Lang;	U+027EA	NewLine;	U+0000A
Gdot;	U+00120	Laplacefrf;	U+02112	Nfr;	U+1D511
Gfr;	U+1D50A	Larr;	U+0219E	NoBreak;	U+02060
Gg;	U+022D9	Lcaron;	U+0013D	NonBreakingSpace;	U+0000A
Gopf;	U+1D53E	Lcedil;	U+0013B	Nopf;	U+02115
GreaterEqual;	U+02265	Lcy;	U+0041B	Not;	U+02AEC
GreaterEqualLess;	U+022DB	LeftAngleBracket;	U+027E8	NotCongruent;	U+02262
GreaterFullEqual;	U+02267	LeftArrow;	U+02190	NotCupCap;	U+0226D
GreaterGreater;	U+02AA2	LeftArrowBar;	U+021E4	NotDoubleVerticalBar;	U+02226
GreaterLess;	U+02277	LeftArrowRightArrow;	U+021C6	NotElement;	U+02209
GreaterSlantEqual;	U+02A7E	LeftCeiling;	U+02308	NotEqual;	U+02260
GreaterTilde;	U+02273	LeftDoubleBracket;	U+027E6	NotExists;	U+02204
Gscr;	U+1D4A2	LeftDownTeeVector;	U+02961	NotGreater;	U+0226F
Gt;	U+0226B	LeftDownVector;	U+021C3	NotGreaterEqual;	U+02271
HARDcy;	U+0042A	LeftDownVectorBar;	U+02959	NotGreaterLess;	U+02279
Hacek;	U+002C7	LeftFloor;	U+0230A	NotGreaterTilde;	U+02275
Hat;	U+0005E	LeftRightArrow;	U+02194	NotLeftTriangle;	U+022EA
Hcirc;	U+00124	LeftRightVector;	U+0294E	NotLeftTriangleEqual;	U+022EC
Hfr;	U+0210C	LeftTee;	U+022A3	NotLess;	U+0226E
HilbertSpace;	U+0210B	LeftTeeArrow;	U+021A4	NotLessEqual;	U+02270
Hopf;	U+0210D	LeftTeeVector;	U+0295A	NotLessGreater;	U+02278
HorizontalLine;	U+02500	LeftTriangle;	U+022B2	NotLessTilde;	U+02274
Hscr;	U+0210B	LeftTriangleBar;	U+029CF	NotPrecedes;	U+02280
Hstrok;	U+00126	LeftTriangleEqual;	U+022B4	NotPrecedesSlantEqual;	U+022E0
HumpDownHump;	U+0224E	LeftUpDownVector;	U+02951	NotReverseElement;	U+0220C
HumpEqual;	U+0224F	LeftUpTeeVector;	U+02960	NotRightTriangle;	U+022EB
IEcy;	U+00415	LeftUpVector;	U+021BF	NotRightTriangleEqual;	U+022ED
IJlig;	U+00132	LeftUpVectorBar;	U+02958	NotSubsetEqual;	U+022E2
IOcy;	U+00401	LeftVector;	U+021BC	NotSquareSubsetEqual;	U+022E3
Iacute;	U+000CD	LeftVectorBar;	U+02952	NotSquareSupersetEqual;	U+022E3
Iacute;	U+000CD	Leftarrow;	U+021D0	NotSubsetEqual;	U+02288
Icirc;	U+000CE	Leftrightarrow;	U+021D4	NotSucceeds;	U+02281
Icirc;	U+000CE	LessEqualGreater;	U+022DA	NotSucceedsSlantEqual;	U+022E1
Icy;	U+00418	LessFullEqual;	U+02266	NotSupersetEqual;	U+02289
Idot;	U+00130	LessGreater;	U+02276	NotTilde;	U+02241
Ifr;	U+02111	LessLess;	U+02AA1	NotTildeEqual;	U+02244
Igrave;	U+000CC	LessSlantEqual;	U+02A7D	NotTildeFullEqual;	U+02247
Igrave;	U+000CC	LessTilde;	U+02272	NotTildeTilde;	U+02249
Im;	U+02111	Lfr;	U+1D50F	NotVerticalBar;	U+02224
Imacr;	U+0012A	Ll;	U+022D8	Nscr;	U+1D4A9
ImaginaryI;	U+02148	Lleftarrow;	U+021DA	Ntilde;	U+000D1
Implies;	U+021D2	Lmidot;	U+0013F	Ntilde;	U+000D1
Int;	U+0222C	LongLeftArrow;	U+027F5	Nu;	U+0039D
Integral;	U+022B	LongLeftRightArrow;	U+027F7	OElig;	U+00152
Intersection;	U+022C2	LongRightArrow;	U+027F6	Oacute;	U+000D3
InvisibleComma;	U+02063	Longleftarrow;	U+027F8	Ocirc;	U+000D3
InvisibleTimes;	U+02062	Longleftrightharrow;	U+027FA	Ocirc;	U+000D4
Iogon;	U+0012E	Longrightarrow;	U+027F9	Ocy;	U+0041E
Iopf;	U+1D540	Lopf;	U+1D543	Odblac;	U+00150
Iota;	U+00399	LowerLeftArrow;	U+02199	Ofr;	U+1D512
Iscr;	U+02110	LowerRightArrow;	U+02198	Ograve;	U+000D2
Itilde;	U+00128	Lscr;	U+02112	Ograve;	U+000D2
Iukcy;	U+00406	Lsh;	U+021B0	Omacr;	U+0014C

Name	Character	Name	Character	Name	Character
Omega;	U+003A9	RightUpTeeVector;	U+0295C	Aacute	U+000DA
Omicron;	U+0039F	RightUpVector;	U+021BE	Uarr;	U+0219F
Oopf;	U+1D546	RightUpVectorBar;	U+02954	Uarrocir;	U+02949
OpenCurlyDoubleQuote;	U+0201C	RightVector;	U+021C0	Ubrcy;	U+0040E
OpenCurlyQuote;	U+02018	RightVectorBar;	U+02953	Ubreve;	U+0016C
Or;	U+02A54	Rightarrow;	U+021D2	Ucirc;	U+000DB
Oscr;	U+1D4AA	Ropf;	U+0211D	Ucirc;	U+000DB
Oslash;	U+000D8	RoundImplies;	U+02970	Ucy;	U+00423
Oslash	U+000D8	Rrightarrow;	U+021DB	Udblac;	U+00170
Otilde;	U+000D5	Rscr;	U+0211B	Ufr;	U+1D518
Otilde	U+000D5	Rsh;	U+021B1	Ugrave;	U+000D9
Otimes;	U+02A37	RuleDelayed;	U+029F4	Ugrave;	U+000D9
Ouml;	U+000D6	SHCHcy;	U+00429	Umacr;	U+0016A
Ouml	U+000D6	SHcy;	U+00428	UnderBar;	U+00332
OverBar;	U+000AF	SOFTcy;	U+0042C	UnderBrace;	U+023DF
OverBrace;	U+023DE	Sacute;	U+0015A	UnderBracket;	U+023B5
OverBracket;	U+023B4	Sc;	U+02ABC	UnderParenthesis;	U+023DD
OverParenthesis;	U+023DC	Scaron;	U+00160	Union;	U+022C3
PartialD;	U+02202	Scedil;	U+0015E	UnionPlus;	U+0228E
Pcy;	U+0041F	Scirc;	U+0015C	Uogon;	U+00172
Pfr;	U+1D513	Scy;	U+00421	Uopf;	U+1D54C
Phi;	U+003A6	Sfr;	U+1D516	UpArrow;	U+02191
Pi;	U+003A0	ShortDownArrow;	U+02193	UpArrowBar;	U+02912
PlusMinus;	U+000B1	ShortLeftArrow;	U+02190	UpArrowDownArrow;	U+021C5
Poincareplane;	U+0210C	ShortRightArrow;	U+02192	UpDownArrow;	U+02195
Popf;	U+02119	ShortUpArrow;	U+02191	UpEquilibrium;	U+0296E
Pr;	U+02ABB	Sigma;	U+003A3	UpTee;	U+022A5
Precedes;	U+0227A	SmallCircle;	U+02218	UpTeeArrow;	U+021A5
PrecedesEqual;	U+02AAF	Sopf;	U+1D54A	Uparrow;	U+021D1
PrecedesSlantEqual;	U+0227C	Sqrt;	U+0221A	Updownarrow;	U+021D5
PrecedesTilde;	U+0227E	Square;	U+025A1	UpperLeftArrow;	U+02196
Prime;	U+02033	SquareIntersection;	U+02293	UpperRightArrow;	U+02197
Product;	U+0220F	SquareSubset;	U+0228F	Upsi;	U+003D2
Proportion;	U+02237	SquareSubsetEqual;	U+02291	Upsilon;	U+003A5
Proportional;	U+0221D	SquareSuperset;	U+02290	Uring;	U+0016E
Pscr;	U+1D4AB	SquareSupersetEqual;	U+02292	Uscr;	U+1D4B0
Psi;	U+003A8	SquareUnion;	U+02294	Utilde;	U+00168
QUOT;	U+00022	Sscr;	U+1D4AE	Uuml;	U+000DC
QUOT	U+00022	Star;	U+022C6	Uuml	U+000DC
Qfr;	U+1D514	Sub;	U+022D0	Vdash;	U+022AB
Qopf;	U+0211A	Subset;	U+022D0	Vbar;	U+02AEB
Qscr;	U+1D4AC	SubsetEqual;	U+02286	Vcy;	U+00412
RBarr;	U+02910	Succeeds;	U+0227B	Vdash;	U+022A9
REG;	U+000AE	SucceedsEqual;	U+02AB0	Vdashl;	U+02AE6
REG	U+000AE	SucceedsSlantEqual;	U+0227D	Vee;	U+022C1
Racute;	U+00154	SucceedsTilde;	U+0227F	Verbar;	U+02016
Rang;	U+027EB	SuchThat;	U+0220B	Vert;	U+02016
Rarr;	U+021A0	Sum;	U+02211	VerticalBar;	U+02223
Rarrtl;	U+02916	Sup;	U+022D1	VerticalLine;	U+0007C
Rcaron;	U+00158	Superset;	U+02283	VerticalSeparator;	U+02758
Rcedil;	U+00156	SupersetEqual;	U+02287	VerticalTilde;	U+02240
Rcy;	U+00420	Supset;	U+022D1	VeryThinSpace;	U+0200A
Re;	U+0211C	THORN;	U+000DE	Vfr;	U+1D519
ReverseElement;	U+0220B	THORN	U+000DE	Vopf;	U+1D54D
ReverseEquilibrium;	U+021CB	TRADE;	U+02122	Vscr;	U+1D4B1
ReverseUpEquilibrium;	U+0296F	TShcy;	U+0040B	Vvdash;	U+022AA
Rfr;	U+0211C	TScy;	U+00426	Wirc;	U+00174
Rho;	U+003A1	Tab;	U+00009	Wedge;	U+022C0
RightAngleBracket;	U+027E9	Tau;	U+003A4	Wfr;	U+1D51A
RightArrow;	U+02192	Tcaron;	U+00164	Wopf;	U+1D54E
RightArrowBar;	U+021E5	Tcedil;	U+00162	Wscr;	U+1D4B2
RightArrowLeftArrow;	U+021C4	Tcy;	U+00422	Xfr;	U+1D51B
RightCeiling;	U+02309	Tfr;	U+1D517	Xi;	U+0039E
RightDoubleBracket;	U+027E7	Therefore;	U+02234	Xopf;	U+1D54F
RightDownTeeVector;	U+0295D	Theta;	U+00398	Xscr;	U+1D4B3
RightDownVector;	U+021C2	ThinSpace;	U+02009	YAc;	U+0042F
RightDownVectorBar;	U+02955	Tilde;	U+0223C	YIcy;	U+00407
RightFloor;	U+0230B	TildeEqual;	U+02243	YUcy;	U+0042E
RightTee;	U+022A2	TildeFullEqual;	U+02245	Yacute;	U+000DD
RightTeeArrow;	U+021A6	TildeTilde;	U+02248	Ycirc;	U+000DD
RightTeeVector;	U+0295B	Topf;	U+1D54B	Ycy;	U+0042B
RightTriangle;	U+022B3	TripleDot;	U+020DB	Yfr;	U+1D51C
RightTriangleBar;	U+029D0	Tscr;	U+1D4AF	Yopf;	U+1D550
RightTriangleEqual;	U+022B5	Tstrok;	U+00166	Yscr;	U+1D4B4
RightUpDownVector;	U+0294F	Uacute;	U+000DA		

Name	Character	Name	Character	Name	Character
Yuml;	U+00178	auml;	U+000E4	boxVL;	U+02562
ZHcy;	U+00416	auml	U+000E4	boxVr;	U+0255F
Zacute;	U+00179	awconint;	U+02233	boxbox;	U+029C9
Zcaron;	U+0017D	awint;	U+02A11	boxdL;	U+02555
Zcy;	U+00417	bNot;	U+02AE0	boxdR;	U+02552
Zdot;	U+0017B	backcong;	U+0224C	boxdl;	U+02510
ZerowidthSpace;	U+0200B	backepsilon;	U+003F6	boxdr;	U+0250C
Zeta;	U+00396	backprime;	U+02035	boxh;	U+02500
Zfr;	U+02128	backsimeq;	U+0223D	boxhD;	U+02565
Zopf;	U+02124	barvee;	U+022BD	boxhU;	U+02568
Zscr;	U+1D4B5	barwed;	U+02305	boxhd;	U+0252C
aacute;	U+000E1	barwedge;	U+02305	boxhu;	U+02534
aacute	U+000E1	bbrk;	U+023B5	boxminus;	U+0229F
abreve;	U+00103	bbrktbrk;	U+023B6	boxplus;	U+0229E
ac;	U+0223E	bcong;	U+0224C	boxtimes;	U+022A0
acd;	U+0223F	bcy;	U+00431	boxUL;	U+0255B
acirc;	U+000E2	bdquo;	U+0201E	boxUR;	U+02558
acirc;	U+000E2	becaus;	U+02235	boxUL;	U+02518
acute;	U+000B4	because;	U+02235	boxur;	U+02514
acute	U+000B4	bemptyv;	U+029B0	boxv;	U+02502
acy;	U+00430	bepsi;	U+003F6	boxvH;	U+0256A
aelig;	U+000E6	bernonu;	U+0212C	boxVL;	U+02561
aelig	U+000E6	beta;	U+003B2	boxvR;	U+0255E
af;	U+02061	beth;	U+02136	boxvh;	U+0253C
afr;	U+1D51E	between;	U+0226C	boxvl;	U+02524
agrave;	U+000E0	bfr;	U+1D51F	boxvr;	U+0251C
grave	U+000E0	bigcap;	U+022C2	bprime;	U+02035
alefsym;	U+02135	bigcirc;	U+025EF	breve;	U+002D8
aleph;	U+02135	bigcup;	U+022C3	brvbar;	U+000A6
alpha;	U+003B1	bigodot;	U+02A00	brvbar;	U+000A6
amacr;	U+00101	bigoplus;	U+02A01	bscr;	U+1D4B7
amalg;	U+02A3F	bigotimes;	U+02A02	bsemi;	U+0204F
amp;	U+00026	bigscup;	U+02A06	bsim;	U+0223D
amp	U+00026	bigstar;	U+02605	bsime;	U+022CD
and;	U+02227	bigtriangledown;	U+025BD	bsol;	U+0005C
andand;	U+02A55	bigtriangleup;	U+025B3	bsolb;	U+029C5
andd;	U+02A5C	biguplus;	U+02A04	bull;	U+02022
andslope;	U+02A58	bigvee;	U+022C1	bullet;	U+02022
andv;	U+02A5A	bigwedge;	U+022C0	bump;	U+0224E
ang;	U+02220	bkarow;	U+0290D	bumpE;	U+02AAE
ange;	U+029A4	blacklozenge;	U+029EB	bumpE;	U+0224F
angle;	U+02220	blacksquare;	U+025AA	bumpeq;	U+0224F
angmsd;	U+02221	blacktriangle;	U+025B4	cacute;	U+00107
angmsdaa;	U+029A8	blacktriangledown;	U+025BE	cap;	U+02229
angmsdab;	U+029A9	blacktriangleleft;	U+025C2	capand;	U+02A44
angmsdac;	U+029AA	blacktriangleright;	U+025B8	caprcup;	U+02A49
angmsdad;	U+029AB	blank;	U+02423	capcap;	U+02A4B
angmsdae;	U+029AC	blk12;	U+02592	capcup;	U+02A47
angmsdaf;	U+029AD	blk14;	U+02591	capdot;	U+02A40
angmsdag;	U+029AE	blk34;	U+02593	caret;	U+02041
angmsdah;	U+029AF	block;	U+02588	caron;	U+002C7
angrt;	U+0221F	bnot;	U+02310	ccaps;	U+02A4D
angrvb;	U+022BE	bopf;	U+1D553	ccaron;	U+0010D
angrvbd;	U+0299D	bot;	U+022A5	ccedil;	U+000E7
angsph;	U+02222	bottom;	U+022A5	ccedil;	U+000E7
angst;	U+0212B	bowtie;	U+022C8	ccirc;	U+00109
angzarr;	U+0237C	boxDL;	U+02557	ccups;	U+02AAC
aogon;	U+00105	boxDR;	U+02554	ccupssm;	U+02A50
aopf;	U+1D552	boxDL;	U+02556	cdot;	U+0010B
ap;	U+02248	boxDR;	U+02553	cedil;	U+000B8
apE;	U+02A70	boxH;	U+02550	cedil;	U+00088
apacir;	U+02A6F	boxHD;	U+02566	emptpyv;	U+029B2
ape;	U+0224A	boxHU;	U+02569	cent;	U+000A2
apid;	U+0224B	boxHd;	U+02564	cent	U+000A2
apos;	U+00027	boxHu;	U+02567	centerdot;	U+000B7
approx;	U+02248	boxUL;	U+0255D	cfr;	U+1D520
approxeq;	U+0224A	boxUR;	U+0255A	chcy;	U+00447
aring;	U+000E5	boxUl;	U+0255C	check;	U+02713
aring	U+000E5	boxUr;	U+02559	checkmark;	U+02713
ascr;	U+1D4B6	boxV;	U+02551	chi;	U+003C7
ast;	U+0002A	boxVH;	U+0256C	cir;	U+025CB
asymp;	U+02248	boxVL;	U+02563	cirE;	U+029C3
asympeq;	U+0224D	boxVR;	U+02560	circ;	U+002C6
atilde;	U+000E3	boxVh;	U+0256B	circeq;	U+02257
atilde	U+000E3			circlearrowleft;	U+021BA

Name	Character	Name	Character	Name	Character
circlearrowright;	U+021BB	dd;	U+02146	el;	U+02A99
circledR;	U+000AE	ddagger;	U+02021	elinters;	U+023E7
circledS;	U+024C8	ddarr;	U+021CA	ell;	U+02113
circledast;	U+0229B	ddotseq;	U+02A77	els;	U+02A95
circledcirc;	U+0229A	deg;	U+000B0	elsdot;	U+02A97
circleddash;	U+0229D	deg;	U+000B0	emacr;	U+00113
cire;	U+02257	delta;	U+003B4	empty;	U+02205
cirfnint;	U+02A10	demptyv;	U+029B1	emptyset;	U+02205
cirmid;	U+02AEF	dfish;	U+0297F	emptyv;	U+02205
circcir;	U+029C2	dfr;	U+1D521	emsp13;	U+02004
clubs;	U+02663	dharl;	U+021C3	emsp14;	U+02005
clubsuit;	U+02663	dharr;	U+021C2	emsp;	U+02003
colon;	U+0003A	diam;	U+022C4	eng;	U+0014B
colone;	U+02254	diamond;	U+022C4	ensp;	U+02002
coloneq;	U+02254	diamondsuit;	U+02666	eogon;	U+00119
comma;	U+0002C	diams;	U+02666	eopf;	U+1D556
commat;	U+00040	die;	U+000A8	epar;	U+022D5
comp;	U+02201	digamma;	U+003DD	eparsl;	U+029E3
compfn;	U+02218	disin;	U+022F2	eplus;	U+02A71
complement;	U+02201	div;	U+000F7	epsi;	U+003F5
complexes;	U+02102	divide;	U+000F7	epsilon;	U+003B5
cong;	U+02245	divide	U+000F7	epsiv;	U+003B5
congdot;	U+02A6D	divideontimes;	U+022C7	eqcirc;	U+02256
conint;	U+0222E	divonx;	U+022C7	eqcolon;	U+02255
copf;	U+1D554	djcy;	U+00452	eqsim;	U+02242
coprod;	U+02210	dlcorn;	U+0231E	eqslantgr;	U+02A96
copy;	U+000A9	dlcrop;	U+0230D	eqslantless;	U+02A95
copy	U+000A9	dollar;	U+00024	equals;	U+0003D
copyrs;	U+02117	dopf;	U+1D555	equest;	U+0225F
crarr;	U+021B5	dot;	U+002D9	equiv;	U+02261
cross;	U+02717	doteq;	U+02250	equivDD;	U+02A78
cscr;	U+1D4B8	doteqdot;	U+02251	eqvparls;	U+029E5
csub;	U+02ACF	dotminus;	U+02238	erDot;	U+02253
csube;	U+02AD1	dotplus;	U+02214	erarr;	U+02971
csup;	U+02AD0	dotsquare;	U+022A1	escr;	U+0212F
csupe;	U+02AD2	doublebarwedge;	U+02306	esdot;	U+02250
ctdot;	U+022EF	downarrow;	U+02193	esim;	U+02242
cudarrl;	U+02938	downdownarrows;	U+021CA	eta;	U+003B7
cudarr;	U+02935	downharpoonleft;	U+021C3	eth;	U+000F0
cuepr;	U+022DE	downharpoonright;	U+021C2	eth	U+000F0
cuesc;	U+022DF	drbkarow;	U+02910	euml;	U+000EB
cularr;	U+021B6	drcorn;	U+0231F	euml	U+000EB
cularrp;	U+0293D	drcrop;	U+0230C	euro;	U+020AC
cup;	U+0222A	dscr;	U+1D4B9	excl;	U+00021
cupbrcap;	U+02A48	dscy;	U+00455	exist;	U+02203
cupcap;	U+02A46	dsol;	U+029F6	expectation;	U+02130
cupcup;	U+02A4A	dstrok;	U+00111	exponentiale;	U+02147
cupdot;	U+0228D	dtdot;	U+022F1	fallingdotseq;	U+02252
cupor;	U+02A45	dtri;	U+025BF	fcy;	U+00444
curarr;	U+021B7	dtrif;	U+025BE	female;	U+02640
curarrm;	U+0293C	duarr;	U+021F5	ffilig;	U+0FB03
curlyeqprec;	U+022DE	duhar;	U+0296F	fflig;	U+0FB00
curlyeqsucc;	U+022DF	dwangle;	U+029A6	ffrl;	U+1D523
curlyvee;	U+022CE	dzcy;	U+0045F	filig;	U+0FB01
curlywedge;	U+022CF	dzigrarr;	U+027FF	flat;	U+0266D
curren;	U+000A4	eDDot;	U+02A77	fliig;	U+0FB02
curren	U+000A4	eDot;	U+02251	fltns;	U+025B1
curvearrowleft;	U+021B6	acute;	U+000E9	fnof;	U+00192
curvearrowright;	U+021B7	acute;	U+000E9	fopf;	U+1D557
cuvee;	U+022CE	easter;	U+02A6E	forall;	U+02200
cuwed;	U+022CF	ecaron;	U+0011B	fork;	U+022D4
cwconint;	U+02232	ecir;	U+02256	forkv;	U+02AD9
cwint;	U+02231	ecirc;	U+000EA	fpartint;	U+02A0D
cylcty;	U+0232D	ecirc;	U+000EA	frac12;	U+000BD
dArr;	U+021D3	ecolon;	U+02255	frac12;	U+000BD
dHar;	U+02965	ecy;	U+0044D	frac13;	U+02153
dagger;	U+02020	edot;	U+00117	frac14;	U+000BC
daleth;	U+02138	ee;	U+02147	frac14;	U+000BC
darr;	U+02193	efDot;	U+02252	frac15;	U+02155
dash;	U+02010	efr;	U+1D522	frac16;	U+02159
dashv;	U+022A3	eg;	U+02A9A	frac18;	U+0215B
dbkarow;	U+0290F	egrave;	U+000E8	frac23;	U+02154
dblac;	U+002DD	egrave;	U+000E8	frac25;	U+02156
dcaron;	U+0010F	egs;	U+02A96	frac34;	U+000BE
dcy;	U+00434	egsdot;	U+02A98		

Name	Character	Name	Character	Name	Character
frac34;	U+000BE	hbar;	U+0210F	iuml;	U+000EF
frac35;	U+02157	hcirc;	U+00125	jcirc;	U+00135
frac38;	U+0215C	hearts;	U+02665	jcy;	U+00439
frac45;	U+02158	heartsuit;	U+02665	jfr;	U+1D527
frac56;	U+0215A	hellip;	U+02026	jmath;	U+00237
frac58;	U+0215D	hercon;	U+022B9	jopf;	U+1D55B
frac78;	U+0215E	hfr;	U+1D525	jscr;	U+1D48F
frasl;	U+02044	hksearrow;	U+02925	jsercy;	U+00458
frown;	U+02322	hkswarow;	U+02926	jukcy;	U+00454
fscr;	U+1D4BB	hoarr;	U+021FF	kappa;	U+003BA
gE;	U+02267	homtht;	U+0223B	kappav;	U+003F0
gEl;	U+02A8C	hookleftarrow;	U+021A9	kcedil;	U+00137
gacute;	U+001F5	hookrightarrow;	U+021AA	kcy;	U+0043A
gamma;	U+003B3	hopf;	U+1D559	kfr;	U+1D528
gammad;	U+003DD	horbar;	U+02015	kgreen;	U+00138
gap;	U+02A86	hsqr;	U+1D4BD	khcy;	U+00445
gbreve;	U+0011F	hslash;	U+0210F	kjcy;	U+0045C
gcirc;	U+0011D	hstrok;	U+00127	kopf;	U+1D55C
gcy;	U+00433	hybull;	U+02043	kscr;	U+1D4C0
gdot;	U+00121	hyphen;	U+02010	lAarr;	U+021DA
ge;	U+02265	iacute;	U+000ED	lArr;	U+021D0
gel;	U+022DB	iacute;	U+000ED	ltaill;	U+0291B
geq;	U+02265	ic;	U+02063	lbarr;	U+0290E
geqq;	U+02267	icirc;	U+000EE	lE;	U+02266
geqslant;	U+02A7E	icirc;	U+000EE	lEg;	U+02A8B
ges;	U+02A7E	icy;	U+00438	lHar;	U+02962
gescc;	U+02AA9	iecy;	U+00435	lacute;	U+0013A
gesdot;	U+02A80	iecl;	U+000A1	laemptyv;	U+029B4
gesdoto;	U+02A82	iecl;	U+000A1	lagran;	U+02112
gesdotol;	U+02A84	iff;	U+021D4	lambda;	U+003BB
gesles;	U+02A94	ifr;	U+1D526	lang;	U+027E8
gfr;	U+1D524	igrave;	U+000EC	langd;	U+02991
gg;	U+0226B	igrave;	U+000EC	langl;	U+027E8
ggg;	U+022D9	ii;	U+02148	lap;	U+02A85
gimel;	U+02137	iiint;	U+02A0C	laquo;	U+000AB
gjcy;	U+00453	iiint;	U+0222D	laquo;	U+000AB
gl;	U+02277	iinfin;	U+029DC	larr;	U+02190
glE;	U+02A92	iiota;	U+02129	larrb;	U+021E4
gla;	U+02AA5	ijlig;	U+00133	larrbfs;	U+0291F
glj;	U+02AA4	imacr;	U+0012B	larrfs;	U+0291D
gnE;	U+02269	image;	U+02111	larrhk;	U+021A9
gnap;	U+02A8A	imagline;	U+02110	larrlp;	U+021AB
gnapprox;	U+02A8A	imagpart;	U+02111	larrpl;	U+02939
gne;	U+02A88	imath;	U+00131	larrsim;	U+02973
gneq;	U+02A88	imof;	U+022B7	larrtl;	U+021A2
gneqq;	U+02269	imped;	U+001B5	lat;	U+02AAB
gnsim;	U+022E7	in;	U+02208	lataill;	U+02919
gopf;	U+1D558	incare;	U+02105	late;	U+02AAD
grave;	U+00060	infin;	U+0221E	lbarr;	U+0290C
gscr;	U+0210A	infintie;	U+029DD	lbbrk;	U+02772
gsim;	U+02273	inodot;	U+00131	lbrace;	U+0007B
gsime;	U+02A8E	int;	U+0222B	lbrack;	U+0005B
gsiml;	U+02A90	intcal;	U+022BA	lbrke;	U+0298B
gt;	U+0003E	integers;	U+02124	lbrksld;	U+0298F
gt	U+0003E	intercal;	U+022BA	lbrkslu;	U+0298D
gtcc;	U+02AA7	intlarhk;	U+02A17	lcaron;	U+0013E
gtcir;	U+02A7A	intprod;	U+02A3C	lcedil;	U+0013C
gtdot;	U+022D7	ioct;	U+00451	lceil;	U+02308
gtlPar;	U+02995	iogon;	U+0012F	lcub;	U+0007B
gtquest;	U+02A7C	iopf;	U+1D55A	lcy;	U+0043B
grapprox;	U+02A86	iota;	U+003B9	ldca;	U+02936
gtrarr;	U+02978	iprod;	U+02A3C	ldquo;	U+0201C
gtrdot;	U+022D7	iquest;	U+000BF	ldquor;	U+0201E
gtreqless;	U+022DB	iquest;	U+000BF	ldrddhar;	U+02967
gtreqless;	U+02A8C	iscr;	U+1D4BE	ldrushar;	U+0294B
gtrless;	U+02277	isin;	U+02208	ldsh;	U+021B2
gtrsim;	U+02273	isinE;	U+022F9	le;	U+02264
hArr;	U+021D4	isindot;	U+022F5	leftarrow;	U+02190
hairsp;	U+0200A	isins;	U+022F4	leftarrowtail;	U+021A2
half;	U+000BD	isinsv;	U+022F3	leftharpoondown;	U+021BD
hamilt;	U+0210B	isinv;	U+02208	leftharpoonup;	U+021BC
hardcy;	U+0044A	it;	U+02062	leftleftarrows;	U+021C7
harr;	U+02194	itilde;	U+00129	leftrigharrows;	U+02194
harrcir;	U+02948	iukcy;	U+00456	leftrigharpoons;	U+021C6
harrw;	U+021AD	iuml;	U+000EF	leftrigharpoons;	U+021CB

Name	Character	Name	Character	Name	Character
leftrightsquigarrow;	U+021AD	lsqb;	U+0005B	naron;	U+00148
leftthreetimes;	U+022CB	lsquo;	U+02018	ncedil;	U+00146
leg;	U+022DA	lsquor;	U+0201A	ncong;	U+02247
leq;	U+02264	lstrok;	U+00142	ncup;	U+02A42
leqq;	U+02266	lt;	U+0003C	ncy;	U+0043D
leqlant;	U+02A7D	lt;	U+0003C	ndash;	U+02013
les;	U+02A7D	ltcc;	U+02AA6	ne;	U+02260
lescc;	U+02AA8	ltcir;	U+02A79	neArr;	U+021D7
lesdot;	U+02A7F	ltdot;	U+022D6	nearhk;	U+02924
lesdoto;	U+02A81	lthree;	U+022CB	nearrr;	U+02197
lesdotor;	U+02A83	ltimes;	U+022C9	nearrow;	U+02197
lesges;	U+02A93	ltlarr;	U+02976	nequiv;	U+02262
lessapprox;	U+02A85	ltquest;	U+02A7B	nesear;	U+02928
lessdot;	U+022D6	ltrPar;	U+02996	nexist;	U+02204
lesseqgtr;	U+022DA	ltri;	U+025C3	nexists;	U+02204
lesseqqtr;	U+02A8B	ltrie;	U+022B4	nfr;	U+1D52B
lessgrt;	U+02276	ltrif;	U+025C2	nge;	U+02271
lesssim;	U+02272	lurdshar;	U+0294A	ngeq;	U+02271
lfisht;	U+0297C	luruhar;	U+02966	ngsim;	U+02275
lfloor;	U+0230A	mDDot;	U+0223A	ngt;	U+0226F
lfr;	U+1D529	macr;	U+000AF	ngtr;	U+0226F
lg;	U+02276	macr;	U+000AF	nhArr;	U+021CE
lgE;	U+02A91	male;	U+02642	nharr;	U+021AE
lhard;	U+021BD	malt;	U+02720	nhpar;	U+02AF2
lharu;	U+021BC	maltese;	U+02720	ni;	U+0220B
lharul;	U+0296A	map;	U+021A6	nis;	U+022FC
lblblk;	U+02584	mapsto;	U+021A6	nisd;	U+022FA
ljcy;	U+00459	mapstodown;	U+021A7	niv;	U+0220B
ll;	U+0226A	mapstoleft;	U+021A4	njcy;	U+0045A
llarr;	U+021C7	mapstoup;	U+021A5	nlArr;	U+021CD
llcorner;	U+0231E	marker;	U+025AE	nlarr;	U+0219A
llhard;	U+0296B	mcomma;	U+02A29	nldr;	U+02025
lltri;	U+025FA	mcy;	U+0043C	nle;	U+02270
lmidot;	U+00140	mdash;	U+02014	nleftarrow;	U+0219A
lmoust;	U+02380	measuredangle;	U+02221	nleftrightarrow;	U+021AE
lmoustache;	U+02380	mfr;	U+1D52A	nleq;	U+02270
lnE;	U+02268	mho;	U+02127	nless;	U+0226E
lnap;	U+02A89	micro;	U+000B5	nlsim;	U+02274
lnapprox;	U+02A89	micro	U+000B5	nlt;	U+0226E
lne;	U+02A87	mid;	U+02223	ntri;	U+022EA
lneq;	U+02A87	midast;	U+0002A	ntrier;	U+022EC
lneqq;	U+02268	midcir;	U+02AF0	nmid;	U+02244
lnsim;	U+022E6	middot;	U+000B7	nopf;	U+1D55F
loang;	U+027EC	middot	U+000B7	not;	U+000AC
loarr;	U+021FD	minus;	U+02212	notin;	U+02209
lobrk;	U+027E6	minusb;	U+0229F	notinva;	U+02209
longleftarrow;	U+027F5	minusd;	U+02238	notinvib;	U+022F7
longleftrightarrow;	U+027F7	minusdu;	U+02A2A	notinvc;	U+022F6
longmapsto;	U+027FC	mlcp;	U+02ADB	notni;	U+0220C
longrightarrow;	U+027F6	mldr;	U+02026	notniva;	U+0220C
looparrowleft;	U+021AB	mnplus;	U+02213	notnivb;	U+022FF
looparrowright;	U+021AC	models;	U+022A7	notnivc;	U+022FD
lopar;	U+02985	mopf;	U+1D55E	npar;	U+02226
lopf;	U+1D55D	mp;	U+02213	nparallel;	U+02226
loplus;	U+02A2D	mscr;	U+1D4C2	npolint;	U+02A14
lotimes;	U+02A34	mstpos;	U+0223E	npr;	U+02280
lowast;	U+02217	mu;	U+003BC	nrcue;	U+022E0
lowbar;	U+0005F	multimap;	U+022B8	nprec;	U+02280
loz;	U+025CA	mumap;	U+022B8	nrArr;	U+021CF
lozeng;	U+025CA	nLeftarrow;	U+021CD	nrarr;	U+0219B
lozf;	U+029EB	nLeftrightarrow;	U+021CE	nrightarrow;	U+0219B
lpar;	U+00028	nRightarrow;	U+021CF	nrtri;	U+022EB
lparlt;	U+02993	nVDash;	U+022AF	nrtrie;	U+022ED
lrarr;	U+021C6	nVdash;	U+022AE	nsc;	U+02281
lrcorner;	U+0231F	nabla;	U+02207	nsccue;	U+022E1
lrhar;	U+021CB	nacute;	U+00144	nscri;	U+1D4C3
lrhard;	U+0296D	nap;	U+02249	nshortmid;	U+02224
lrm;	U+0200E	napos;	U+00149	nshortparallel;	U+02226
lrtri;	U+022BF	napprox;	U+02249	nsim;	U+02241
lsaquo;	U+02039	natur;	U+0266E	nsime;	U+02244
lscr;	U+1D4C1	natural;	U+0266E	nsimeq;	U+02244
lsh;	U+021B0	naturals;	U+02115	nsmid;	U+02224
lsim;	U+02272	nnbsp;	U+000A0	nspat;	U+02226
lsime;	U+02A8D	nnbsp;	U+000A0	nsqsube;	U+022E2
lsimg;	U+02A8F	ncap;	U+02A43		

Name	Character	Name	Character	Name	Character
nsqsup;	U+022E3	ordm;	U+000BA	prod;	U+0220F
nsub;	U+02284	ordm	U+000BA	profalar;	U+0232E
nsube;	U+02288	origof;	U+022B6	proline;	U+02312
nsubseteq;	U+02288	oror;	U+02A56	profsurf;	U+02313
nsucc;	U+02281	orslope;	U+02A57	prop;	U+0221D
nsup;	U+02285	orv;	U+02A5B	proto;	U+0221D
nsupe;	U+02289	oscr;	U+02134	prsim;	U+0227E
nsubseteq;	U+02289	oslash;	U+000F8	prurel;	U+02280
ntgl;	U+02279	oslash	U+000F8	pscr;	U+1D4C5
ntilde;	U+000F1	osol;	U+02298	psi;	U+003C8
ntilde	U+000F1	otilde;	U+000F5	puncsp;	U+02008
ntlg;	U+02278	otilde	U+000F5	qfr;	U+1D52E
ntriangleleft;	U+022EA	otimes;	U+02297	qint;	U+02A0C
ntrianglelefteq;	U+022EC	otimesas;	U+02A36	qpof;	U+1D562
ntriangleright;	U+022EB	ouml;	U+000F6	qprime;	U+02057
ntrianglerighteq;	U+022ED	ouml	U+000F6	qscr;	U+1D4C6
nu;	U+003BD	ovbar;	U+0233D	quaternions;	U+0210D
num;	U+00023	par;	U+02225	quatint;	U+02A16
numero;	U+02116	para;	U+000B6	quest;	U+0003F
numsp;	U+02007	para	U+000B6	questeq;	U+0225F
nvDash;	U+022AD	parallel;	U+02225	quot;	U+00022
nvHarr;	U+02904	parsim;	U+02AF3	quot	U+00022
nvdash;	U+022AC	parsl;	U+02AFD	rArr;	U+021DB
nvinfin;	U+029DE	part;	U+02202	rArr;	U+021D2
nvlArr;	U+02902	pcy;	U+0043F	rAtail;	U+0291C
nvrArr;	U+02903	percnt;	U+00025	rBarr;	U+0290F
nwArr;	U+02106	period;	U+0002E	rHar;	U+02964
nwarhk;	U+02923	permil;	U+02030	race;	U+029DA
nwarr;	U+02196	perp;	U+022A5	racute;	U+00155
nwarrow;	U+02196	pertenk;	U+02031	radic;	U+0221A
nwnear;	U+02927	pfr;	U+1D52D	raemptyv;	U+02983
oS;	U+024C8	phi;	U+003C6	rang;	U+027E9
oacute;	U+000F3	phiv;	U+003C6	rangd;	U+02992
oacute;	U+000F3	phmmat;	U+02133	range;	U+029A5
oast;	U+0229B	phone;	U+0260E	rangle;	U+027E9
ocir;	U+0229A	pi;	U+003C0	raquo;	U+000BB
ocirc;	U+000F4	pitchfork;	U+022D4	raquo	U+000BB
ocirc;	U+000F4	piv;	U+003D6	rarr;	U+02192
ocy;	U+0043E	planck;	U+0210F	rarrap;	U+02975
odash;	U+0229D	planckh;	U+0210E	rarrb;	U+021E5
odblac;	U+00151	plankv;	U+0210F	rarrbfs;	U+02920
odiv;	U+02A38	plus;	U+0002B	rarrc;	U+02933
odot;	U+02299	plusacir;	U+02A23	rarrfs;	U+0291E
odsold;	U+029BC	plusb;	U+0229E	rarrhk;	U+021AA
oelig;	U+00153	pluscir;	U+02A22	rarrlp;	U+021AC
ofcir;	U+029BF	plusdo;	U+02214	rarrpl;	U+02945
ofr;	U+1D52C	plusdu;	U+02A25	rarrsim;	U+02974
ogon;	U+002DB	pluse;	U+02A72	rarrtl;	U+021A3
ograve;	U+000F2	plusmn;	U+000B1	rarrw;	U+0219D
ograve	U+000F2	plusmn	U+000B1	ratail;	U+0291A
ogt;	U+029C1	plussim;	U+02A26	ratio;	U+02236
ohbar;	U+02985	plustwo;	U+02A27	rationals;	U+0211A
ohm;	U+02126	pm;	U+000B1	rbarr;	U+0290D
oint;	U+0222E	pointint;	U+02A15	rbbrk;	U+02773
olarr;	U+021BA	popf;	U+1D561	rbrace;	U+0007D
olcir;	U+029BE	pound;	U+000A3	rbrack;	U+0005D
olcross;	U+029BB	pound	U+000A3	rbrke;	U+0298C
oline;	U+0203E	pr;	U+0227A	rbrksld;	U+0298E
olt;	U+029C0	prE;	U+02AB3	rbrkslu;	U+02990
omacr;	U+0014D	prap;	U+02AB7	rcaron;	U+00159
omega;	U+003C9	prcue;	U+0227C	rcedil;	U+00157
omicron;	U+003BF	pre;	U+02AAF	rceil;	U+02309
omid;	U+029B6	prec;	U+0227A	rcub;	U+0007D
ominus;	U+02296	precapprox;	U+02AB7	rCY;	U+00440
oopf;	U+1D560	preccurlyeq;	U+0227C	rdca;	U+02937
opar;	U+02987	preceq;	U+02AAF	rdlhar;	U+02969
operp;	U+029B9	precnapprox;	U+02AB9	rdquo;	U+0201D
oplus;	U+02295	precneqq;	U+02AB5	rdquor;	U+0201D
or;	U+02228	precnsim;	U+022E8	rdsh;	U+021B3
orarr;	U+021BB	precsim;	U+0227E	real;	U+0211C
ord;	U+02A5D	prime;	U+02032	realine;	U+0211B
order;	U+02134	primes;	U+02119	realpart;	U+0211C
orderof;	U+02134	prnE;	U+02AB5	reals;	U+0211D
ordf;	U+000AA	prnap;	U+02AB9	rect;	U+025AD
ordf;	U+000AA	prnsim;	U+022E8	reg;	U+000AE

Name	Character	Name	Character	Name	Character
reg;	U+000AE	sect;	U+000A7	subnE;	U+02ACB
rfisht;	U+0297D	semi;	U+0003B	subne;	U+0228A
rfloor;	U+0230B	seswar;	U+02929	subplus;	U+02ABF
rfr;	U+1D52F	setminus;	U+02216	subrarr;	U+02979
rhard;	U+021C1	setmn;	U+02216	subset;	U+02282
rharu;	U+021C0	sext;	U+02736	subseteq;	U+02286
rharul;	U+0296C	sfr;	U+1D530	subseteqq;	U+02AC5
rho;	U+003C1	srown;	U+02322	subsetneq;	U+0228A
rhol;	U+003F1	sharp;	U+0266F	subsetneqq;	U+02ACB
rightarrow;	U+02192	shchcy;	U+00449	subsim;	U+02AC7
rightarrowtail;	U+021A3	shchy;	U+00448	subsub;	U+02AD5
rightharpoondown;	U+021C1	shortmid;	U+02223	subsup;	U+02AD3
rightharpoonup;	U+021C0	shortparallel;	U+02225	succ;	U+0227B
rightleftarrows;	U+021C4	shy;	U+000AD	succapprox;	U+02AB8
rightleftharpoons;	U+021CC	shy;	U+000AD	succcurlyeq;	U+0227D
rightrightarrow;	U+021C9	sigma;	U+003C3	succeq;	U+02AB0
rightsquigarrow;	U+0219D	sigmaf;	U+003C2	succnapprox;	U+02ABA
rightthreetimes;	U+022CC	sigmav;	U+003C2	succneqq;	U+02AB6
ring;	U+002DA	sim;	U+0223C	succnsim;	U+022E9
risingdotseq;	U+02253	simdot;	U+02A6A	succsim;	U+0227F
rlarr;	U+021C4	sime;	U+02243	sum;	U+02211
rlhar;	U+021CC	simeq;	U+02243	sung;	U+0266A
rlm;	U+0200F	simg;	U+02A9E	sup1;	U+000B9
rmoust;	U+023B1	simgE;	U+02AA0	sup1	U+000B9
rmoustache;	U+023B1	siml;	U+02A9D	sup2;	U+000B2
rnmid;	U+02AEE	simlE;	U+02A9F	sup2	U+000B2
roang;	U+027ED	simne;	U+02246	sup3;	U+000B3
roarr;	U+021FE	simplus;	U+02A24	sup3	U+000B3
robrik;	U+027E7	simrarr;	U+02972	sup;	U+02283
ropar;	U+02986	slarr;	U+02190	supE;	U+02AC6
ropf;	U+1D563	smallsetminus;	U+02216	supdot;	U+02ABE
roplus;	U+02A2E	smashp;	U+02A33	supdsub;	U+02AD8
rotimes;	U+02A35	smeparsl;	U+029E4	supe;	U+02287
rpar;	U+00029	smid;	U+02223	supedot;	U+02AC4
rpargt;	U+02994	smile;	U+02323	suphsub;	U+02AD7
rppolint;	U+02A12	smt;	U+02AAA	suplarr;	U+0297B
rrarr;	U+021C9	smte;	U+02AAC	supmult;	U+02AC2
rsaquo;	U+0203A	softcy;	U+0044C	supnE;	U+02ACC
rscr;	U+1D4C7	sol;	U+0002F	supne;	U+0228B
rsh;	U+021B1	solb;	U+029C4	supplus;	U+02AC0
rsqb;	U+0005D	solbar;	U+0233F	supset;	U+02283
rsquo;	U+02019	sopf;	U+1D564	supseteq;	U+02287
rsquor;	U+02019	spades;	U+02660	supseteqq;	U+02AC6
rthree;	U+022CC	spadesuit;	U+02660	supsetneq;	U+0228B
rtimes;	U+022CA	spar;	U+02225	supsetneqq;	U+02ACC
rtri;	U+025B9	sqcap;	U+02293	supsim;	U+02AC8
rtrie;	U+022B5	sqcup;	U+02294	supsub;	U+02AD4
rtrif;	U+025B8	sqsub;	U+0228F	supsup;	U+02AD6
rtriltri;	U+029CE	sqsube;	U+02291	swArr;	U+021D9
ruluhar;	U+02968	sqsubset;	U+0228F	swarhk;	U+02926
rx;	U+0211E	sqsubseteq;	U+02291	swarr;	U+02199
sacute;	U+0015B	sqsup;	U+02290	swarw;	U+02199
sbquo;	U+0201A	sqsupe;	U+02292	swmwar;	U+0292A
sc;	U+0227B	sqsupset;	U+02290	szlig;	U+000DF
scE;	U+02AB4	sqsupseteq;	U+02292	szlig;	U+000DF
scap;	U+02AB8	squ;	U+025A1	target;	U+02316
scaron;	U+00161	square;	U+025A1	tau;	U+003C4
sccue;	U+0227D	squarf;	U+025AA	tbrk;	U+023B4
sce;	U+02AB0	squf;	U+025AA	tcaron;	U+00165
scedil;	U+0015F	srarr;	U+02192	tcedil;	U+00163
scirc;	U+0015D	sscr;	U+1D4C8	tcy;	U+00442
scnE;	U+02AB6	ssetmn;	U+02216	tdot;	U+020DB
scsnap;	U+02ABA	ssmile;	U+02323	telrec;	U+02315
scnsim;	U+022E9	sstarf;	U+022C6	tfr;	U+1D531
scpolint;	U+02A13	star;	U+02606	there4;	U+02234
scsim;	U+0227F	starf;	U+02605	therefore;	U+02234
scy;	U+00441	straightepsilon;	U+003F5	theta;	U+003B8
sdot;	U+022C5	straightphi;	U+003D5	thetasym;	U+003D1
sdtob;	U+022A1	strns;	U+000AF	thetav;	U+003D1
sdate;	U+02A66	sub;	U+02282	thickapprox;	U+02248
seArr;	U+021D8	subE;	U+02AC5	thicksim;	U+0223C
searhk;	U+02925	subdot;	U+02ABD	thinsp;	U+02009
searr;	U+02198	sube;	U+02286	thkap;	U+02248
searrow;	U+02198	subedot;	U+02AC3	thksim;	U+0223C
sect;	U+000A7	submult;	U+02AC1	thorn;	U+000FE

Name	Character	Name	Character	Name	Character
thorn	U+000FE	umacr;	U+0016B	wedge;	U+02227
tilde;	U+002DC	uml;	U+000A8	wedgeq;	U+02259
times;	U+000D7	uml;	U+000A8	weierp;	U+02118
times	U+000D7	uogon;	U+00173	wfr;	U+1D534
timesb;	U+022A0	uopf;	U+1D566	wopf;	U+1D568
timesbar;	U+02A31	uparrow;	U+02191	wp;	U+02118
timesd;	U+02A30	updownarrow;	U+02195	wr;	U+02240
tint;	U+0222D	upharpoonleft;	U+021BF	wreath;	U+02240
toea;	U+02928	upharpoonright;	U+021BE	wscr;	U+1D4CC
top;	U+022A4	uplus;	U+0228E	xcap;	U+022C2
topbot;	U+02336	upsi;	U+003C5	xcirc;	U+025EF
topcir;	U+02AF1	upsih;	U+003D2	xcup;	U+022C3
topf;	U+1D565	upsilon;	U+003C5	xdtri;	U+025BD
topfork;	U+02ADA	upuparrows;	U+021C8	xfr;	U+1D535
tosa;	U+02929	urcorn;	U+0231D	xhArr;	U+027FA
tprime;	U+02034	urcorner;	U+0231D	xharr;	U+027F7
trade;	U+02122	urcrop;	U+0230E	xi;	U+003BE
triangle;	U+025B5	uring;	U+0016F	xiArr;	U+027F8
triangledown;	U+025BF	urtri;	U+025F9	xlarr;	U+027F5
triangleleft;	U+025C3	uscr;	U+1D4CA	xmap;	U+027FC
trianglelefteq;	U+022B4	utdot;	U+022F0	xnis;	U+022FB
triangleright;	U+0225C	utilde;	U+00169	xodot;	U+02A00
trianglerighteq;	U+025B9	utri;	U+025B5	xopf;	U+1D569
tridot;	U+025E5	utrif;	U+025B4	xoplus;	U+02A01
trie;	U+0225C	uarr;	U+021C8	xotime;	U+02A02
triminus;	U+02A3A	uuml;	U+000FC	xrArr;	U+027F9
triplus;	U+02A39	uuml;	U+000FC	xrarr;	U+027F6
trisb;	U+029CD	uwangle;	U+029A7	xscr;	U+1D4CD
tritime;	U+02A3B	vArr;	U+021D5	xsqcup;	U+02A06
trpezium;	U+023E2	vBar;	U+02AE8	xuplus;	U+02A04
tscr;	U+1D4C9	vBarv;	U+02AE9	xutri;	U+025B3
tscy;	U+00446	vDash;	U+022A8	xvee;	U+022C1
tshcy;	U+0045B	vangrt;	U+0299C	xwedge;	U+022C0
tstrok;	U+00167	varepsilon;	U+003B5	yacute;	U+000FD
twixt;	U+0226C	varkappa;	U+003F0	yacy;	U+0044F
twoheadleftarrow;	U+0219E	varnothing;	U+02205	ycirc;	U+00177
twoheadrightarrow;	U+021A0	varphi;	U+003C6	ycy;	U+0044B
uArr;	U+021D1	varpi;	U+003D6	yen;	U+000A5
uHar;	U+02963	varproto;	U+0221D	yen	U+000A5
uacute;	U+000FA	varr;	U+02195	yfr;	U+1D536
uacute;	U+000FA	varrho;	U+003F1	yicy;	U+00457
uarr;	U+02191	varsigma;	U+003C2	yopf;	U+1D56A
ubrcy;	U+0045E	vartheta;	U+003D1	yscr;	U+1D4CE
ubreve;	U+0016D	vartriangleleft;	U+022B2	yucy;	U+0044E
ucirc;	U+000FB	vartriangleright;	U+022B3	yuml;	U+000FF
ucirc;	U+000FB	vcy;	U+00432	yuml;	U+000FF
ucy;	U+00443	vdash;	U+022A2	zacute;	U+0017A
udarr;	U+021C5	vee;	U+02228	zcaron;	U+0017E
udblac;	U+00171	veebar;	U+022BB	zcy;	U+00437
udhar;	U+0296E	veeq;	U+0225A	zdot;	U+0017C
ufisht;	U+0297E	vellip;	U+022EE	zeetrf;	U+02128
ufr;	U+1D532	verbar;	U+0007C	zeta;	U+003B6
ugrave;	U+000F9	vfr;	U+1D533	zfr;	U+1D537
ugrave	U+000F9	vltri;	U+022B2	zhcy;	U+00436
uharl;	U+021BF	vopf;	U+1D567	zigrarr;	U+021DD
uharr;	U+021BE	vprop;	U+0221D	zopf;	U+1D56B
uhblk;	U+02580	vrtri;	U+022B3	zscr;	U+1D4CF
ulcorn;	U+0231C	vscr;	U+1D4CB	zwj;	U+0200D
ulcorner;	U+0231C	vzigzag;	U+0299A	wnj;	U+0200C
ulcrop;	U+0230F	wcirc;	U+00175		
ultri;	U+025F8	wedbar;	U+02A5F		

9 Rendering and user-agent behavior

- ** This section will probably include details on how to render DATAGRID (including its pseudo-elements), drag-and-drop, etc, in a visual medium, in concert with CSS. Terms that need to be defined include: **sizing of embedded content**

CSS UAs in visual media must, when scrolling a page to a fragment identifier, align the top of the viewport with the target element's top border edge.

- ** must define letting the user "**obtain a physical form** (or a representation of a physical form)" of a document (printing) and what this means for the UA, in particular creating a new view for the print media.
- ** Must define that in CSS, tag and attribute names in HTML documents, and class names in quirks mode documents, are case-insensitive, as well as saying which attribute values must be compared case-insensitively.

9.1 Rendering and the DOM

- ** This section is wrong. mediaMode will end up on Window, I think. All views implement Window.

Any object implement the AbstractView interface must also implement the MediaModeAbstractView interface.

```
interface MediaModeAbstractView {  
    readonly attribute DOMString mediaMode;  
};
```

The **mediaMode** attribute on objects implementing the MediaModeAbstractView interface must return the string that represents the canvas' current rendering mode (screen, print, etc). This is a lowercase string, as defined by the CSS specification. [CSS21]

Some user agents may support multiple media, in which case there will exist multiple objects implementing the AbstractView interface. Only the default view implements the Window interface. The other views can be reached using the view attribute of the UIEvent interface, during event propagation. There is no way currently to enumerate all the views.

9.2 Rendering and menus/toolbars

9.2.1 The 'icon' property

UAs should use the command's Icon as the default generic icon provided by the user agent when the 'icon' property computes to 'auto' on an element that either defines a command or refers to one using the command attribute, but when the property computes to an actual image, it should use that image instead.

9.3 Obsolete elements, attributes, and APIs

9.3.1 The body element

** Need to define the content attributes in terms of CSS or something.

```
[XXX] interface HTMLDocument {  
    attribute DOMString fgColor;  
    attribute DOMString bgColor;  
    attribute DOMString linkColor;  
    attribute DOMString vlinkColor;  
    attribute DOMString alinkColor;  
};
```

The **fgColor** attribute on the Document object must reflect (page 67) the text attribute on the body element (page 81).

The **bgColor** attribute on the Document object must reflect (page 67) the bgcolor attribute on the body element (page 81).

The **linkColor** attribute on the Document object must reflect (page 67) the link attribute on the body element (page 81).

The **vLinkColor** attribute on the Document object must reflect (page 67) the vlink attribute on the body element (page 81).

The **aLinkColor** attribute on the Document object must reflect (page 67) the alink attribute on the body element (page 81).

```
[XXX] interface HTMLBodyElement {  
    attribute DOMString text;  
    attribute DOMString bgColor;  
    attribute DOMString background;  
    attribute DOMString link;  
    attribute DOMString vLink;  
    attribute DOMString aLink;  
};
```

The **text** DOM attribute of the body element must reflect (page 67) the element's text content attribute.

The **bgColor** DOM attribute of the body element must reflect (page 67) the element's bgcolor content attribute.

The **background** DOM attribute of the body element must reflect (page 67) the element's background content attribute.

The **link** DOM attribute of the body element must reflect (page 67) the element's link content attribute.

The **aLink** DOM attribute of the body element must reflect (page 67) the element's alink content attribute.

The **vLink** DOM attribute of the body element must reflect (page 67) the element's vlink content attribute.

9.3.2 The applet element

The applet element is a Java-specific variant of the embed element. In HTML5 the applet element is obsoleted so that all extension frameworks (Java, .NET, Flash, etc) are handled in a consistent manner.

If the sandboxed plugins browsing context flag (page 218) is set on the browsing context (page 414) for which the applet element's document is the active document (page 414), then the element must be ignored (it represents nothing).

** Otherwise, define how the element works, if supported.

```
[XXX] interface HTMLDocument {  
    readonly attribute HTMLCollection applets;  
};
```

The **applets** attribute must return an HTMLCollection rooted at the Document node, whose filter matches only applet elements.

10 Things that you can't do with this specification because they are better handled using other technologies that are further described herein

This section is non-normative.

There are certain features that are not handled by this specification because a client side markup language is not the right level for them, or because the features exist in other languages that can be integrated into this one. This section covers some of the more common requests.

10.1 Localization

If you wish to create localized versions of an HTML application, the best solution is to preprocess the files on the server, and then use HTTP content negotiation to serve the appropriate language.

10.2 Declarative 2D vector graphics and animation

Embedding vector graphics into XHTML documents is the domain of SVG.

10.3 Declarative 3D scenes

Embedding 3D imagery into XHTML documents is the domain of X3D, or technologies based on X3D that are namespace-aware.

10.4 Timers

- ** This section is expected to be moved to its own specification in due course. It needs a lot of work to actually make it into a semi-decent spec.

Objects that implement the Window interface must also implement the WindowTimers interface:

```
[NoInterfaceObject] interface WindowTimers {  
    // timers  
    long setTimeout(in TimeoutHandler handler, in long timeout);  
    long setTimeout(in TimeoutHandler handler, in long timeout,  
    arguments...);  
    long setTimeout(in DOMString code, in long timeout);  
    long setTimeout(in DOMString code, in long timeout, in DOMString  
    language);  
    void clearTimeout(in long handle);  
    long setInterval(in TimeoutHandler handler, in long timeout);  
    long setInterval(in TimeoutHandler handler, in long timeout,
```

```

    arguments...);
    long setInterval(in DOMString code, in long timeout);
    long setInterval(in DOMString code, in long timeout, in DOMString
language);
    void clearInterval(in long handle);
};

interface TimeoutHandler {
    void handleEvent([Variadic] in any args);
};

```

The `setTimeout` and `setInterval` methods allow authors to schedule timer-based events.

The `setTimeout(handler, timeout[, arguments...])` method takes a reference to a `TimeoutHandler` object and a length of time in milliseconds. It must return a handle to the timeout created, and then asynchronously wait `timeout` milliseconds and then queue a task (page 429) to invoke `handleEvent()` on the `handler` object. If any `arguments...` were provided, they must be passed to the `handler` as arguments to the `handleEvent()` function.

Alternatively, `setTimeout(code, timeout[, language])` may be used. This variant takes a

- ** string instead of a `TimeoutHandler` object. define the actual requirements for this method,
- ** as with the previous one. That string must be parsed using the specified `language` (defaulting to ECMAScript if the third argument is omitted) and executed in the scope of the browsing context (page 414) associated with the `Window` object on which the `setTimeout()` method was invoked.
- ** Need to define `language` values.

The `setInterval(...)` variants must work in the same way as the `setTimeout` variants except that if `timeout` is a value greater than zero, the task (page 429) that invokes the `handler` or `code` must be queued (page 429) again every `timeout` milliseconds, not just the once.

The `clearTimeout()` and `clearInterval()` methods take one integer (the value returned by `setTimeout()` and `setInterval()` respectively) and must cancel the specified timeout. When called with a value that does not correspond to an active timeout or interval, the methods must return without doing anything.

Index

This section is non-normative.

** List of elements

** List of attributes

** List of interfaces

** List of events

References

** This section will be written in a future draft.

Acknowledgements

Thanks to Aankhen, Aaron Boodman, Aaron Leventhal, Adam Barth, Adam Roben, Addison Phillips, Adele Peterson, Adrian Sutton, Agustín Fernández, Alastair Campbell, Alexey Feldgendler, Anders Carlsson, Andrew Gove, Andrew Sidwell, Anne van Kesteren, Anthony Hickson, Anthony Ricaud, Antti Koivisto, Arphen Lin, Asbjørn Ulsberg, Ashley Sheridan, Aurelien Levy, Ben Boyle, Ben Godfrey, Ben Meadowcroft, Ben Millard, Benjamin Hawkes-Lewis, Bert Bos, Bill Mason, Billy Wong, Bjoern Hoehrmann, Boris Zbarsky, Brad Fults, Brad Neuberg, Brady Eidson, Brendan Eich, Brett Wilson, Brian Campbell, Brian Smith, Bruce Miller, Cameron McCormack, Carlos Perelló Marín, Chao Cai, ??? (Channy Yun), Charl van Niekerk, Charles Iliya Krempeaux, Charles McCathieNevile, Christian Biesinger, Christian Johansen, Chriswa, Cole Robison, Collin Jackson, Daniel Barclay, Daniel Brumbaugh Keeney, Daniel Glazman, Daniel Peng, Daniel Spång, Daniel Steinberg, Danny Sullivan, Darin Adler, Darin Fisher, Dave Camp, Dave Singer, Dave Townsend, David Baron, David Bloom, David Carlisle, David Flanagan, David Håsäther, David Hyatt, David Smith, Dean Edridge, Debi Orton, Derek Featherstone, DeWitt Clinton, Dimitri Glazkov, dolphinling, Doron Rosenberg, Doug Kramer, Edward O'Connor, Eira Monstad, Elliotte Harold, Eric Carlson, Eric Law, Erik Arvidsson, Evan Martin, Evan Prodromou, fantasai, Felix Sasaki, Franck 'Shift' Quélain, Garrett Smith, Geoffrey Garen, Geoffrey Sneddon, Håkon Wium Lie, Henri Sivonen, Henrik Lied, Henry Mason, Hugh Winkler, Ignacio Javier, Ivo Emanuel Gonçalves, J. King, Jacques Distler, James Graham, James Justin Harrell, James M Snell, James Perrett, Jan-Klaas Kollhof, Jason White, Jasper Bryant-Greene, Jeff Cutsinger, Jeff Schiller, Jeff Walden, Jens Bannmann, Jens Fendler, Jeroen van der Meer, Jim Jewett, Jim Meehan, Joe Clark, Joseph Kesselman, Jjgod Jiang, Joel Spolsky, Johan Herland, John Boyer, John Bussjaeger, John Harding, Johnny Stenback, Jon Perlow, Jonathan Worent, Jorgen Horstink, Josh Levenberg, Joshua Randall, Jukka K. Korpela, Jules Clément-Ripoche, Julian Reschke, Kai Hendry, Kornel Lesinski, ???? (KUROSAWA Takeshi), Kristof Zelechovski, Lachlan Hunt, Larry Page, Lars Gunther, Laura L. Carlson, Laura Wisewell, Laurens Holst, Lee Kowalkowski, Leif Halvard Silli, Lenny Domnitser, Léonard Bouchet, Leons Petrazickis, Logan, Loune, Maciej Stachowiak, Magnus Kristiansen, Malcolm Rowe, Mark Nottingham, Mark Rowe, Mark Schenk, Martijn Wargers, Martin Atkins, Martin Dürst, Martin Honnen, Masataka Yakura, Mathieu Henri, Matthew Mastracci, Matthew Raymond, Matthew Thomas, Mattias Waldau, Max Romantschuk, Michael 'Ratt' Iannarelli, Michael A. Nachbaur, Michael A. Puls II, Michael Carter, Michael Gratton, Michael Nordman, Michael Powers, Michael(tm) Smith, Michel Fortin, Michiel van der Blonk, Mihai Şucan, Mike Brown, Mike Dierken, Mike Dixon, Mike Schinkel, Mike Shaver, Mikko Rantalainen, Neil Deakin, Neil Soiffer, Olaf Hoffmann, Olav Junker Kjær, Oliver Hunt, Peter Karlsson, Peter Kasting, Philip Jägenstedt, Philip Taylor, Philip TAYLOR, Rachid Finge, Rajas Moonka, Ralf Stoltze, Ralph Giles, Raphael Champeimont, Rene Saarsoo, Richard Ishida, Rimantas Liubertas, Robert Blaut, Robert O'Callahan, Robert Sayre, Roman Ivanov, Ryan King, S. Mike Dierken, Sam Ruby, Sam Weinig, Scott Hess, Sean Knapp, Shaun Inman, Silvia Pfeiffer, Simon Pieters, Stefan Haustein, Steffen Meschkat, Stephen Ma, Steve Faulkner, Steve Runyon, Steven Garrity, Stewart Brodie, Stuart Parmenter, Sunava Dutta, Tantek Çelik, Terrence Wood, Thomas Broyer, Thomas O'Connor, Tim Altman, Tim Johansson, Travis Leithead, Tyler Close, Vladimir Vukićević, Wakaba, Wayne Pollock, William Swanson, Yi-An Huang, and Øistein E. Andersen, for their useful and substantial comments.

Thanks also to everyone who has ever posted about HTML5 to their blogs, public mailing lists, or forums, including the W3C public-html list and the various WHATWG lists.

Special thanks to Richard Williamson for creating the first implementation of canvas in Safari, from which the canvas feature was designed.

Special thanks also to the Microsoft employees who first implemented the event-based drag-and-drop mechanism, `contenteditable`, and other features first widely deployed by the Windows Internet Explorer browser.

Special thanks and \$10,000 to David Hyatt who came up with a broken implementation of the adoption agency algorithm (page 633) that the editor had to reverse engineer and fix before using it in the parsing section.

Thanks to the many sources that provided inspiration for the examples used in the specification.

Thanks also to the Microsoft blogging community for some ideas, to the attendees of the W3C Workshop on Web Applications and Compound Documents for inspiration, to the #mrt crew, the #mrt.no crew, and the #whatwg crew, and to Pillar and Hedral for their ideas and support.